



# Linear Systems





Some visual areas...


$$\mathbf{x} [n, m] \rightarrow f \rightarrow \mathbf{y} [n, m]$$



From M. Lewicky

# Retinal Receptive Fields


## Receptive field structure in bipolar cells





# A bit of history:



## Neurocognitron [Fukushima 1980]




“sandwich” architecture (SCSCSC...)

simple cells: modifiable parameters

complex cells: perform pooling



# Linear Systems



One important class of systems is the set of linear systems.

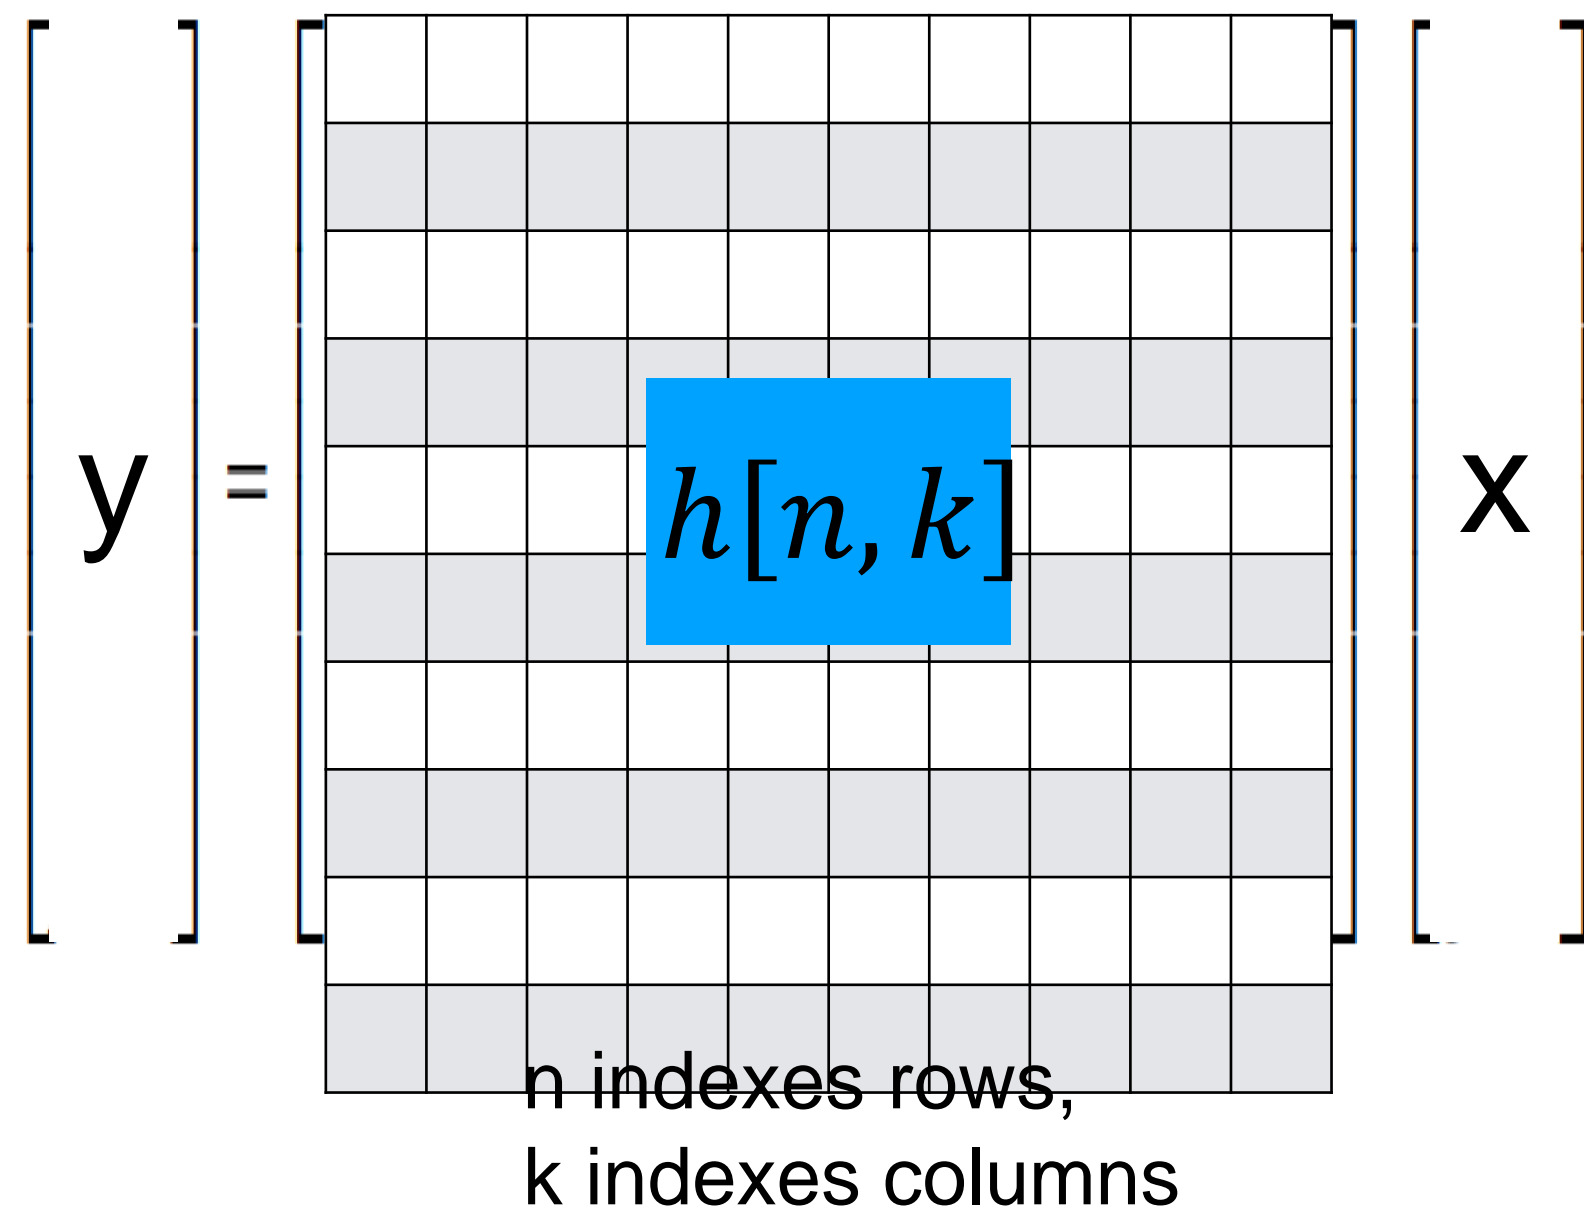
A function  $f$  is linear if it satisfies:

$$f(\alpha \mathbf{x}) = \alpha f(\mathbf{x})$$

$$f(\mathbf{x} + \mathbf{y}) = f(\mathbf{x}) + f(\mathbf{y})$$

# Linear system: $y = f(\mathbf{x})$

A linear function  $f$  can be written as a matrix multiplication:




The diagram illustrates the matrix multiplication  $y = Hx$ . On the left, a vertical vector  $y$  is shown with a bracket. In the center is a 10x10 grid representing the matrix  $H$ . The grid is composed of alternating light gray and white squares in a checkerboard pattern. A blue square in the center of the grid contains the text  $h[n, k]$ . To the right of the grid is a vertical vector  $x$  with a bracket. Below the grid, the text "n indexes rows," and "k indexes columns" is displayed.

$$y = Hx$$


n indexes rows,  
k indexes columns

# Linear system: $y = f(\mathbf{x})$

A linear function  $f$  can be written as a matrix multiplication:



It can also be represented as a fully connected linear neural network




$h[n, k]$  Is the strength of the connection between  $x[k]$  and  $y[n]$




# Linear system: $y = f(\mathbf{x})$

A linear function  $f$  can be written as a matrix multiplication:



It can also be represented as a fully connected linear neural network




$h[n, k]$  Is the strength of the connection between  $x[k]$  and  $y[n]$


$$y[n] = \sum_{k=0}^{N-1} h[n, k]x[k]$$




Images are turned into column vectors  
by concatenating all image columns






Column vector of length 16






=







4x4 image





Quiz: what operation is linear?


# Quiz: what operation is linear?











We need translation invariance










**Classifier**





**“Bird”**








							Bird






Sky	Sky	Sky	Sky	Sky	Sky	Sky	Bird
Sky	Sky	Sky	Sky	Sky	Sky	Sky	Sky
Sky	Sky	Sky	Sky	Sky	Sky	Sky	Sky
Bird	Bird	Bird	Sky	Bird	Sky	Sky	Sky
Sky	Sky	Sky	Bird	Sky	Sky	Sky	Sky






# Moving Average

- basic idea: define a new function by averaging over a sliding window
- a simple example to start off: smoothing



# Moving Average

- Can add weights to our moving average
- *Weights* [..., 0, 1, 1, 1, 1, 1, 0, ...] / 5



# In 2D: box filter

$$\frac{1}{9} \begin{matrix} & & h[\cdot, \cdot] \\ \begin{matrix} 1 \\ | \\ 9 \end{matrix} & \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \end{matrix}$$



# Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g[\cdot, \cdot]$


$$g[m, n] = \sum_{k, l} h[k, l] f[m + k, n + l]$$

# Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g[\cdot, \cdot]$

	0	10							

$$g[m, n] = \sum_{k, l} h[k, l] f[m + k, n + l]$$

# Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20						

# Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20	30					



# Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20	30	30				

# Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g[\cdot, \cdot]$

	0	10	20	30	30				

# Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20	30	30				
						?			
				50					

# Image filtering

$$h[\cdot, \cdot] \quad \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g[\cdot, \cdot]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$g[m, n] = \sum_{k, l} h[k, l] f[m + k, n + l]$$



# Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$$\frac{1}{9} h[\cdot, \cdot]$$


1	1	1
1	1	1
1	1	1

# Linear filters: examples



Original

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} =$$



Blur (with a mean filter)

# Cross-correlation

Let  $F$  be the image,  $H$  be the kernel (of size  $2k+1 \times 2k+1$ ), and  $G$  be the output image


$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called a **cross-correlation** operation:

$$G = H \otimes F$$

- Can think of as a “dot product” between local neighborhood and kernel for each pixel

# Other filters



1	0	-1
2	0	-2
1	0	-1


Sobel



Vertical Edge  
(absolute value)



# Other filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge  
(absolute value)

# Cross-correlation vs. Convolution

---

**cross-correlation:**

$$G = H \otimes F$$

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$


A **convolution** operation is a cross-correlation where the filter is flipped both horizontally and vertically before being applied to the image:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

It is written:

$$G = H \star F$$

# Convolution







# Convolution is nice!

- Notation:  $b = c \star a$
- Convolution is a multiplication-like operation
  - commutative  $a \star b = b \star a$
  - associative  $a \star (b \star c) = (a \star b) \star c$
  - distributes over addition  $a \star (b + c) = a \star b + a \star c$
  - scalars factor out  $\alpha a \star b = a \star \alpha b = \alpha(a \star b)$
  - identity: unit impulse  $e = [\dots, 0, 0, 1, 0, 0, \dots]$ 
$$a \star e = a$$
- Conceptually no distinction between filter and signal
- Usefulness of associativity
  - often apply several filters one after another:  $((a \star b_1) \star b_2) \star b_3$
  - this is equivalent to applying one filter:  $a \star (b_1 \star b_2 \star b_3)$


Quiz: what operation is the result of a convolution?




# Quiz: what operation is the result of a convolution?



# Examples




# Defocus/blurring





Computes the local average over windows of size 5 x 5 pixels

# Examples





# Examples







# Examples




# Examples




# Examples



# Examples




# Examples





# Examples





# Convolution as matrix multiplication

In the 1D case, it helps to make explicit the structure of the matrix:

$$\begin{bmatrix} 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \end{bmatrix} \circ \begin{bmatrix} 1/3 & 1/3 & 1/3 \end{bmatrix} = \begin{bmatrix} & 0 & 30 & 60 & 90 & 90 & 90 & 60 & 30 & \end{bmatrix}$$

# Convolution as matrix multiplication

In the 1D case, it helps to make explicit the structure of the matrix:

$$\begin{bmatrix} 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \end{bmatrix} \circ \begin{bmatrix} 1/3 & 1/3 & 1/3 \end{bmatrix} = \begin{bmatrix} \phantom{0} & 0 & 30 & 60 & 90 & 90 & 90 & 60 & 30 & \phantom{0} \end{bmatrix}$$

In the 1D case, it helps to make explicit the structure of the matrix:

$$\begin{bmatrix} 0 \\ 30 \\ 60 \\ 90 \\ 90 \\ 90 \\ 90 \\ 60 \\ 30 \end{bmatrix} = \begin{bmatrix} & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 0 \\ 90 \\ 90 \\ 90 \\ 90 \\ 90 \\ 0 \\ 0 \end{bmatrix}$$

# Convolution as matrix multiplication

In the 1D case, it helps to make explicit the structure of the matrix:

$$\begin{bmatrix} 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \end{bmatrix} \circ \begin{bmatrix} 1/3 & 1/3 & 1/3 \end{bmatrix} = \begin{bmatrix} \phantom{0} & 0 & 30 & 60 & 90 & 90 & 90 & 60 & 30 & \phantom{0} \end{bmatrix}$$

In the 1D case, it helps to make explicit the structure of the matrix:

$$\begin{bmatrix} 0 \\ 30 \\ 60 \\ 90 \\ 90 \\ 90 \\ 90 \\ 60 \\ 30 \end{bmatrix} = \begin{bmatrix} 1/3 & 1/3 & 1/3 & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \\ & & & & & & & & & \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 0 \\ 90 \\ 90 \\ 90 \\ 90 \\ 90 \\ 0 \\ 0 \end{bmatrix}$$

# Convolution as matrix multiplication

In the 1D case, it helps to make explicit the structure of the matrix:

0

0

0

90

90

90

90

90

0

0

o

1/3

1/3

1/3

=

0

30

60

90

90

90

60

30

In the 1D case, it helps to make explicit the structure of the matrix:

0

30

60

90

90

90

90

60

30

=

1/3

1/3

1/3

1/3

1/3

1/3

0

0

0

90

90

90

90

90

0

0

# Convolution as matrix multiplication

In the 1D case, it helps to make explicit the structure of the matrix:

$$\begin{bmatrix} 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \end{bmatrix} \otimes \begin{bmatrix} 1/3 & 1/3 & 1/3 \end{bmatrix} = \begin{bmatrix} & 0 & 30 & 60 & 90 & 90 & 90 & 60 & 30 & \end{bmatrix}$$


In the 1D case, it helps to make explicit the structure of the matrix:

$$\begin{bmatrix} 0 \\ 30 \\ 60 \\ 90 \\ 90 \\ 90 \\ 90 \\ 60 \\ 30 \end{bmatrix} = \begin{bmatrix} 1/3 & 1/3 & 1/3 & & & & & & \\ & 1/3 & 1/3 & 1/3 & & & & & \\ & & 1/3 & 1/3 & 1/3 & & & & \\ & & & 1/3 & 1/3 & 1/3 & & & \\ & & & & 1/3 & 1/3 & 1/3 & & \\ & & & & & 1/3 & 1/3 & 1/3 & \\ & & & & & & 1/3 & 1/3 & 1/3 \\ & & & & & & & 1/3 & 1/3 & 1/3 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 0 \\ 90 \\ 90 \\ 90 \\ 90 \\ 90 \\ 0 \\ 0 \end{bmatrix}$$




# Linear translation invariant system:

A LTI function  $f$  can be written as a matrix multiplication:



It can also be represented as a convolutional layer of neural net:



$$y[n] = \sum_{k=-1}^1 h[k]x[n-k]$$

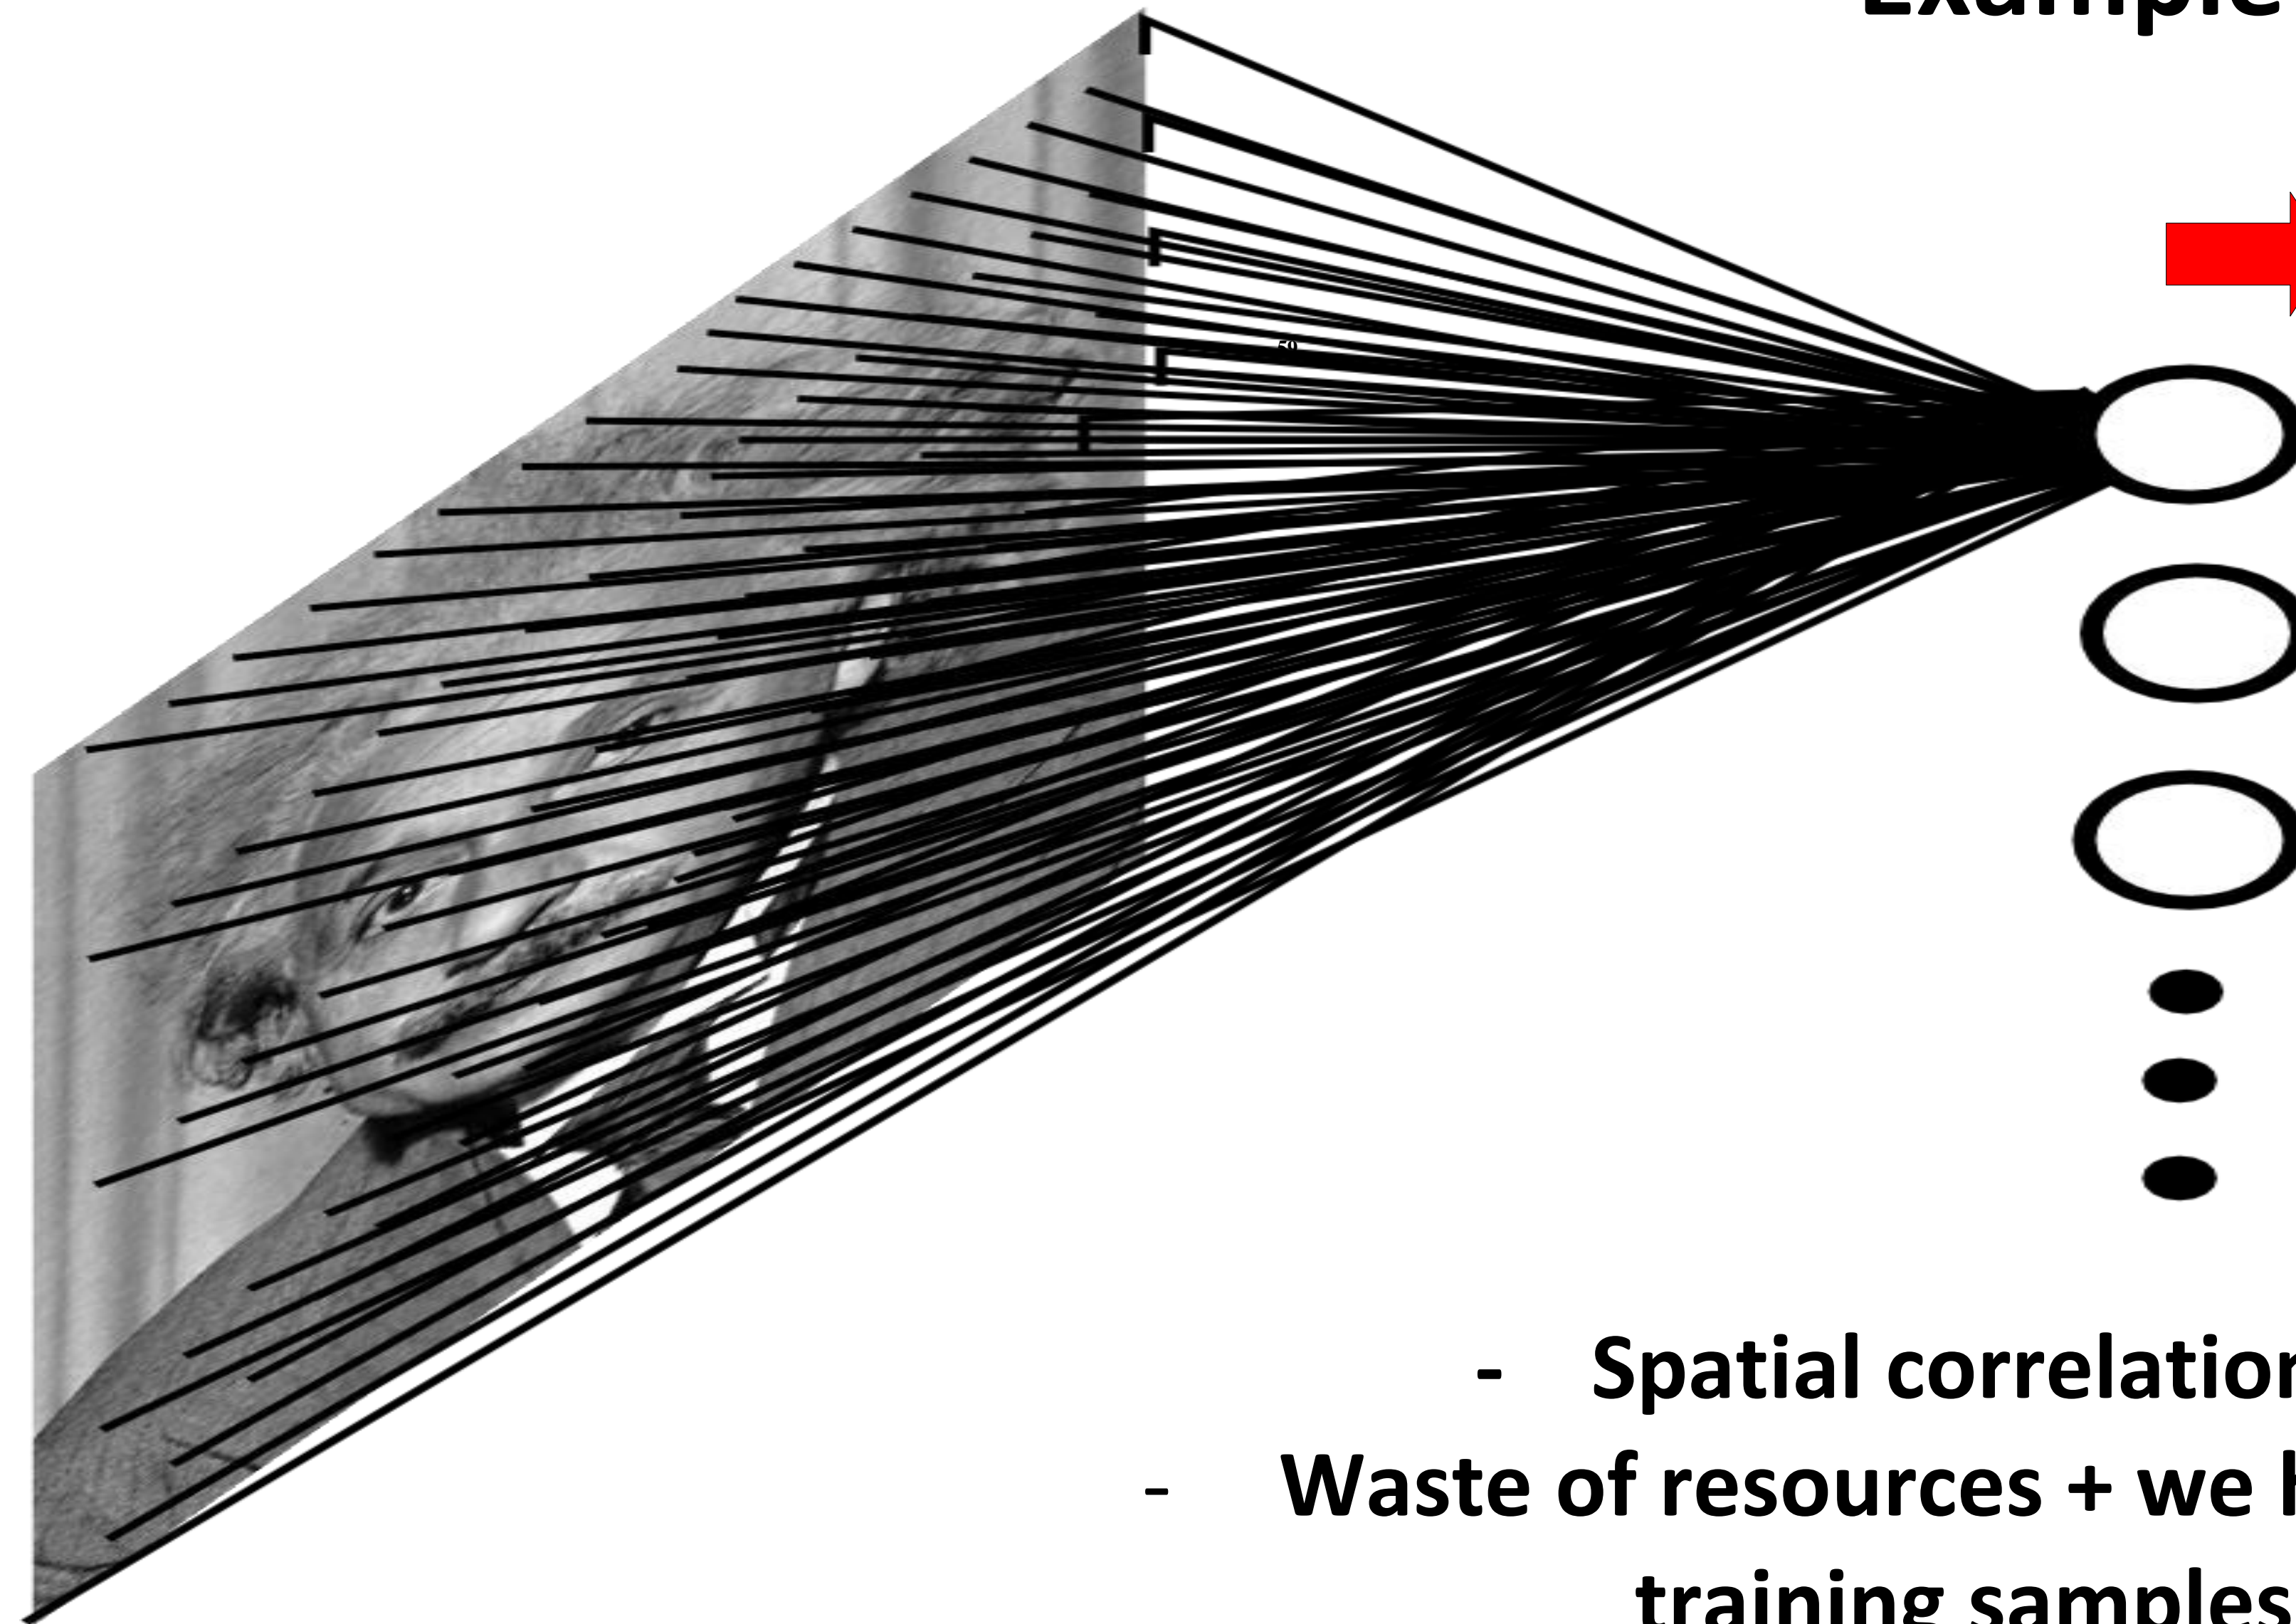
$h[n-k]$  is the strength of the connection between  $x[k]$  and  $y[n]$

# Fully Connected Layer

Example: 200x200 image

40K hidden units

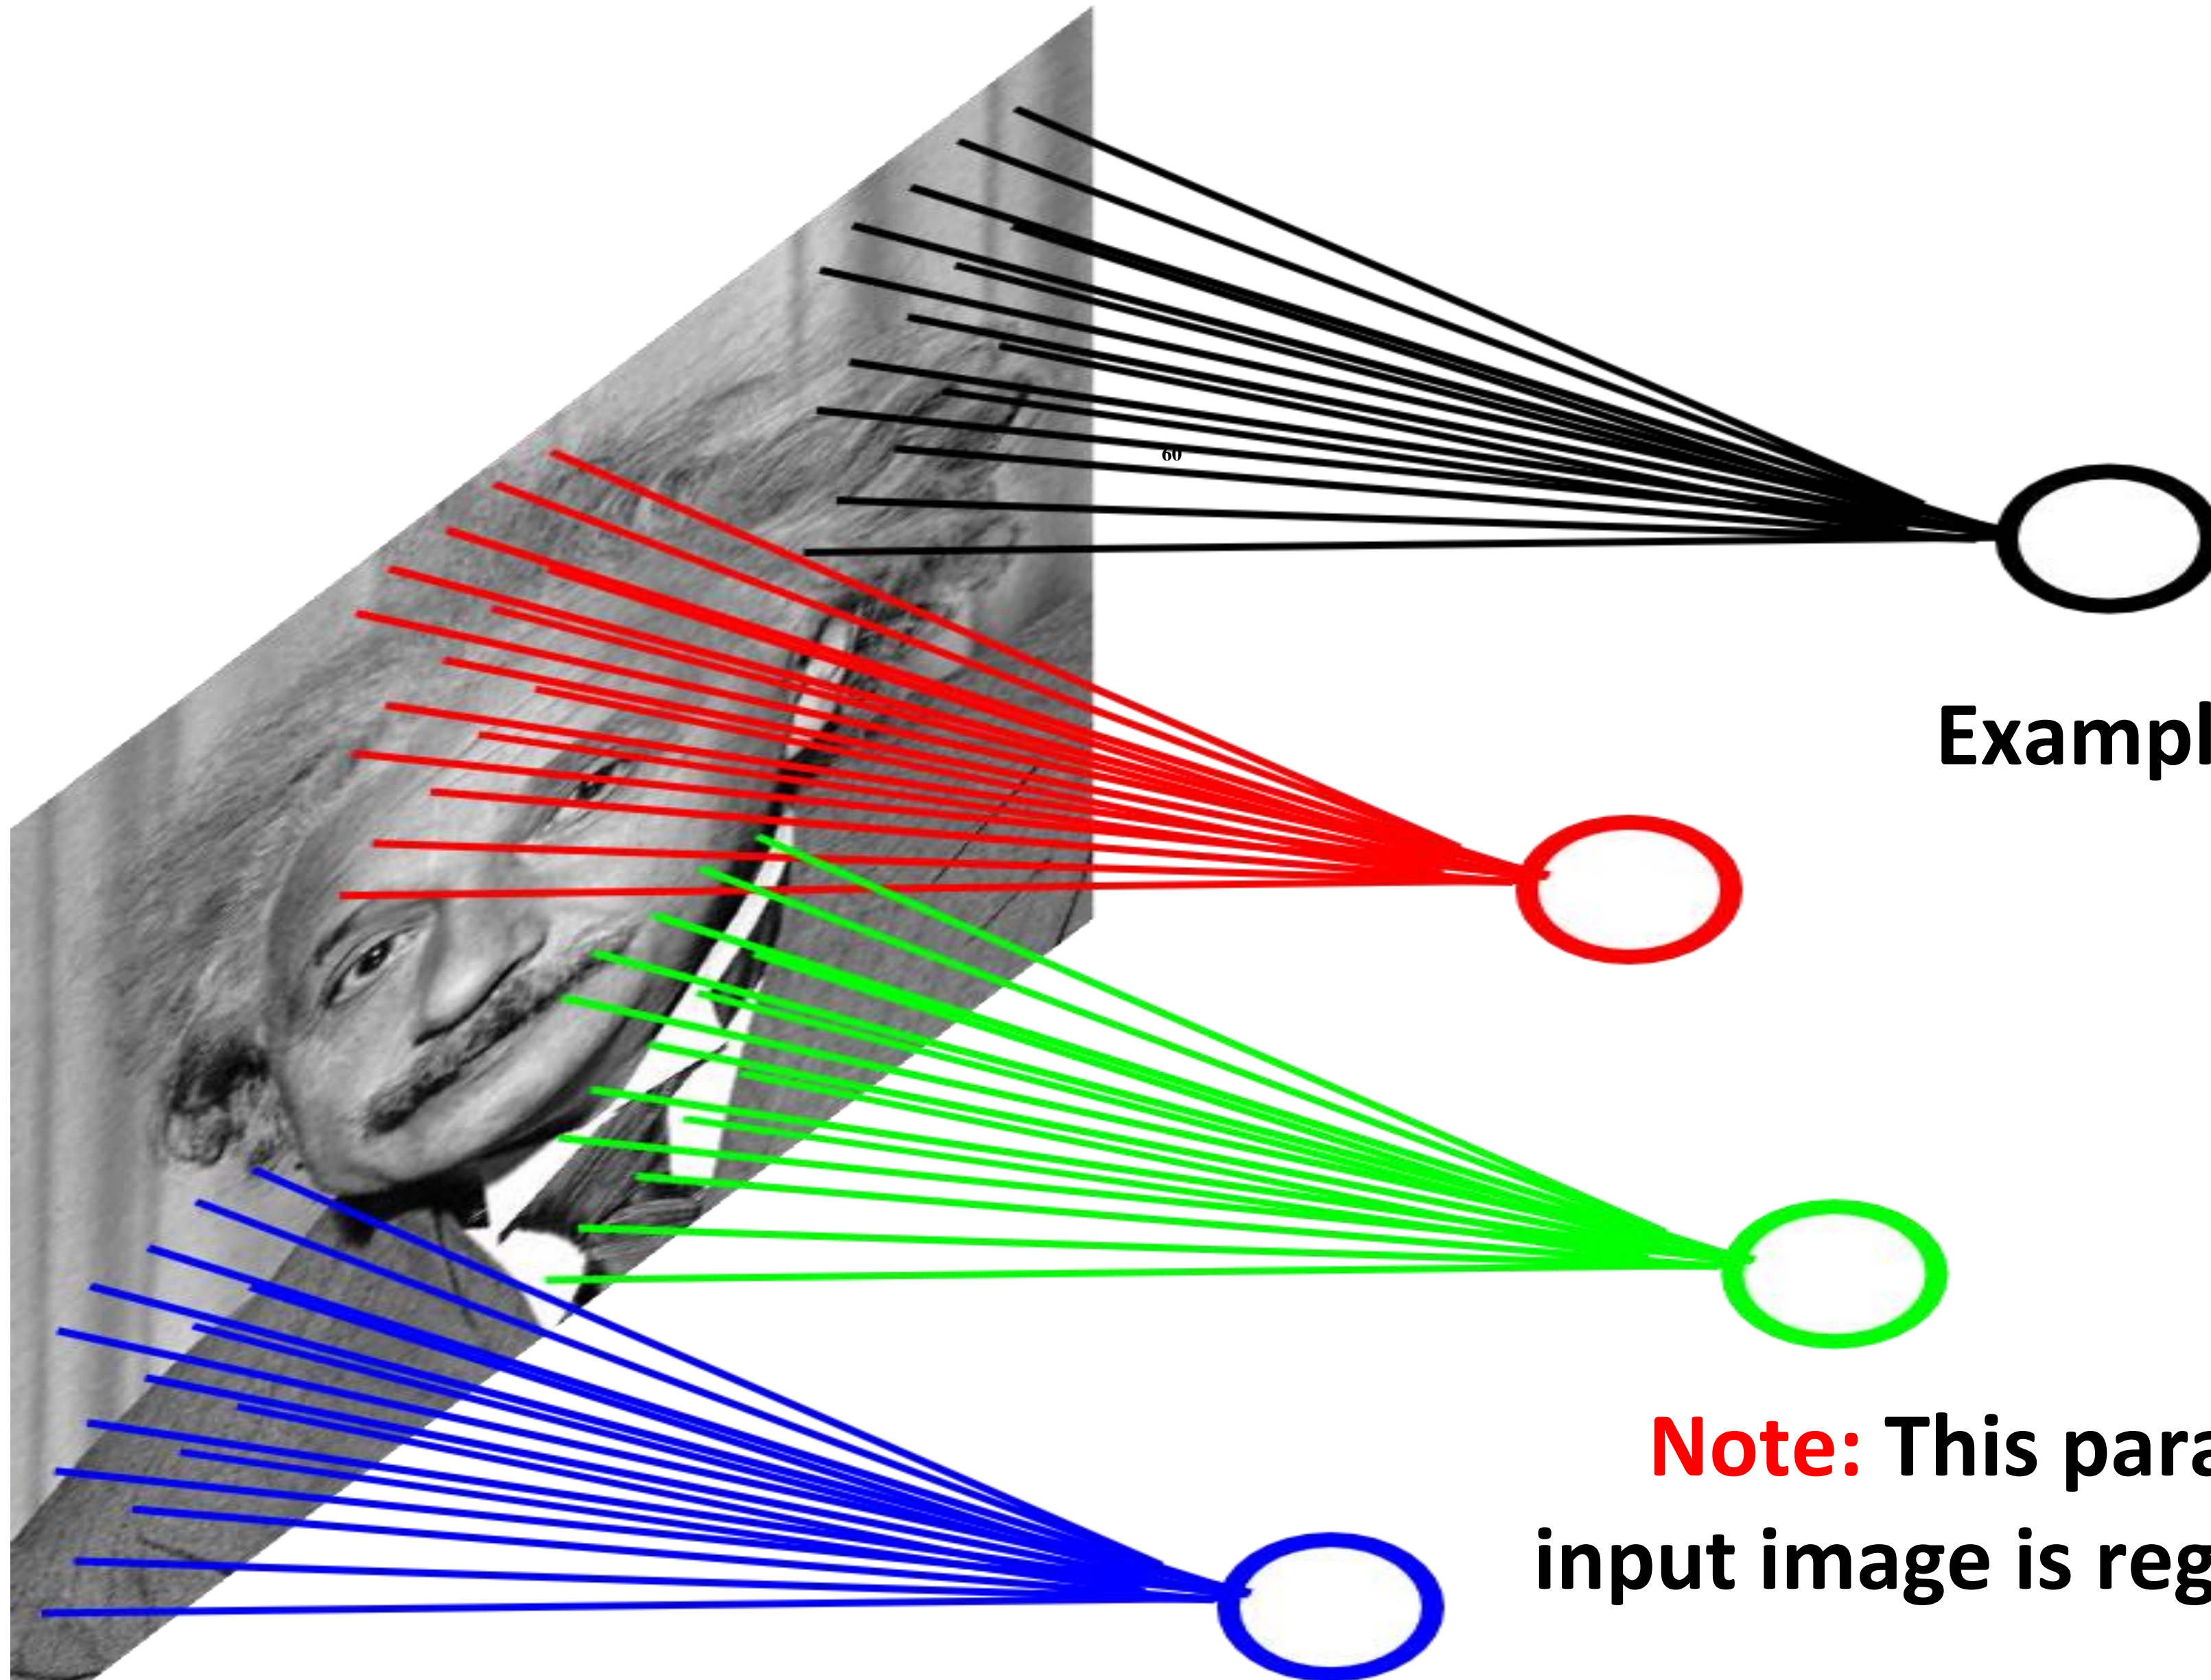
→ **~2B parameters!!!**



- Spatial correlation is local
- Waste of resources + we have not enough training samples anyway..



# Locally Connected Layer




Example: 200x200 image  
40K hidden units  
Filter size: 10x10  
4M parameters

**Note:** This parameterization is good when input image is registered (e.g., face recognition).



# Locally Connected Layer


**STATIONARITY?** Statistics is similar at different locations



Example: 200x200 image  
40K hidden units  
Filter size: 10x10  
4M parameters



# Convolutional Layer



Share the same parameters across different locations (assuming input is stationary):


**Convolutions with learned kernels**



# Rectangular filter



$g[m,n]$



$h[m,n]$



$f[m,n]$

# Rectangular filter



$g[m,n]$

$\otimes$



$h[m,n]$

$=$



$f[m,n]$

# Rectangular filter



$g[m,n]$

$\otimes$




$h[m,n]$

=




$f[m,n]$

# The identity



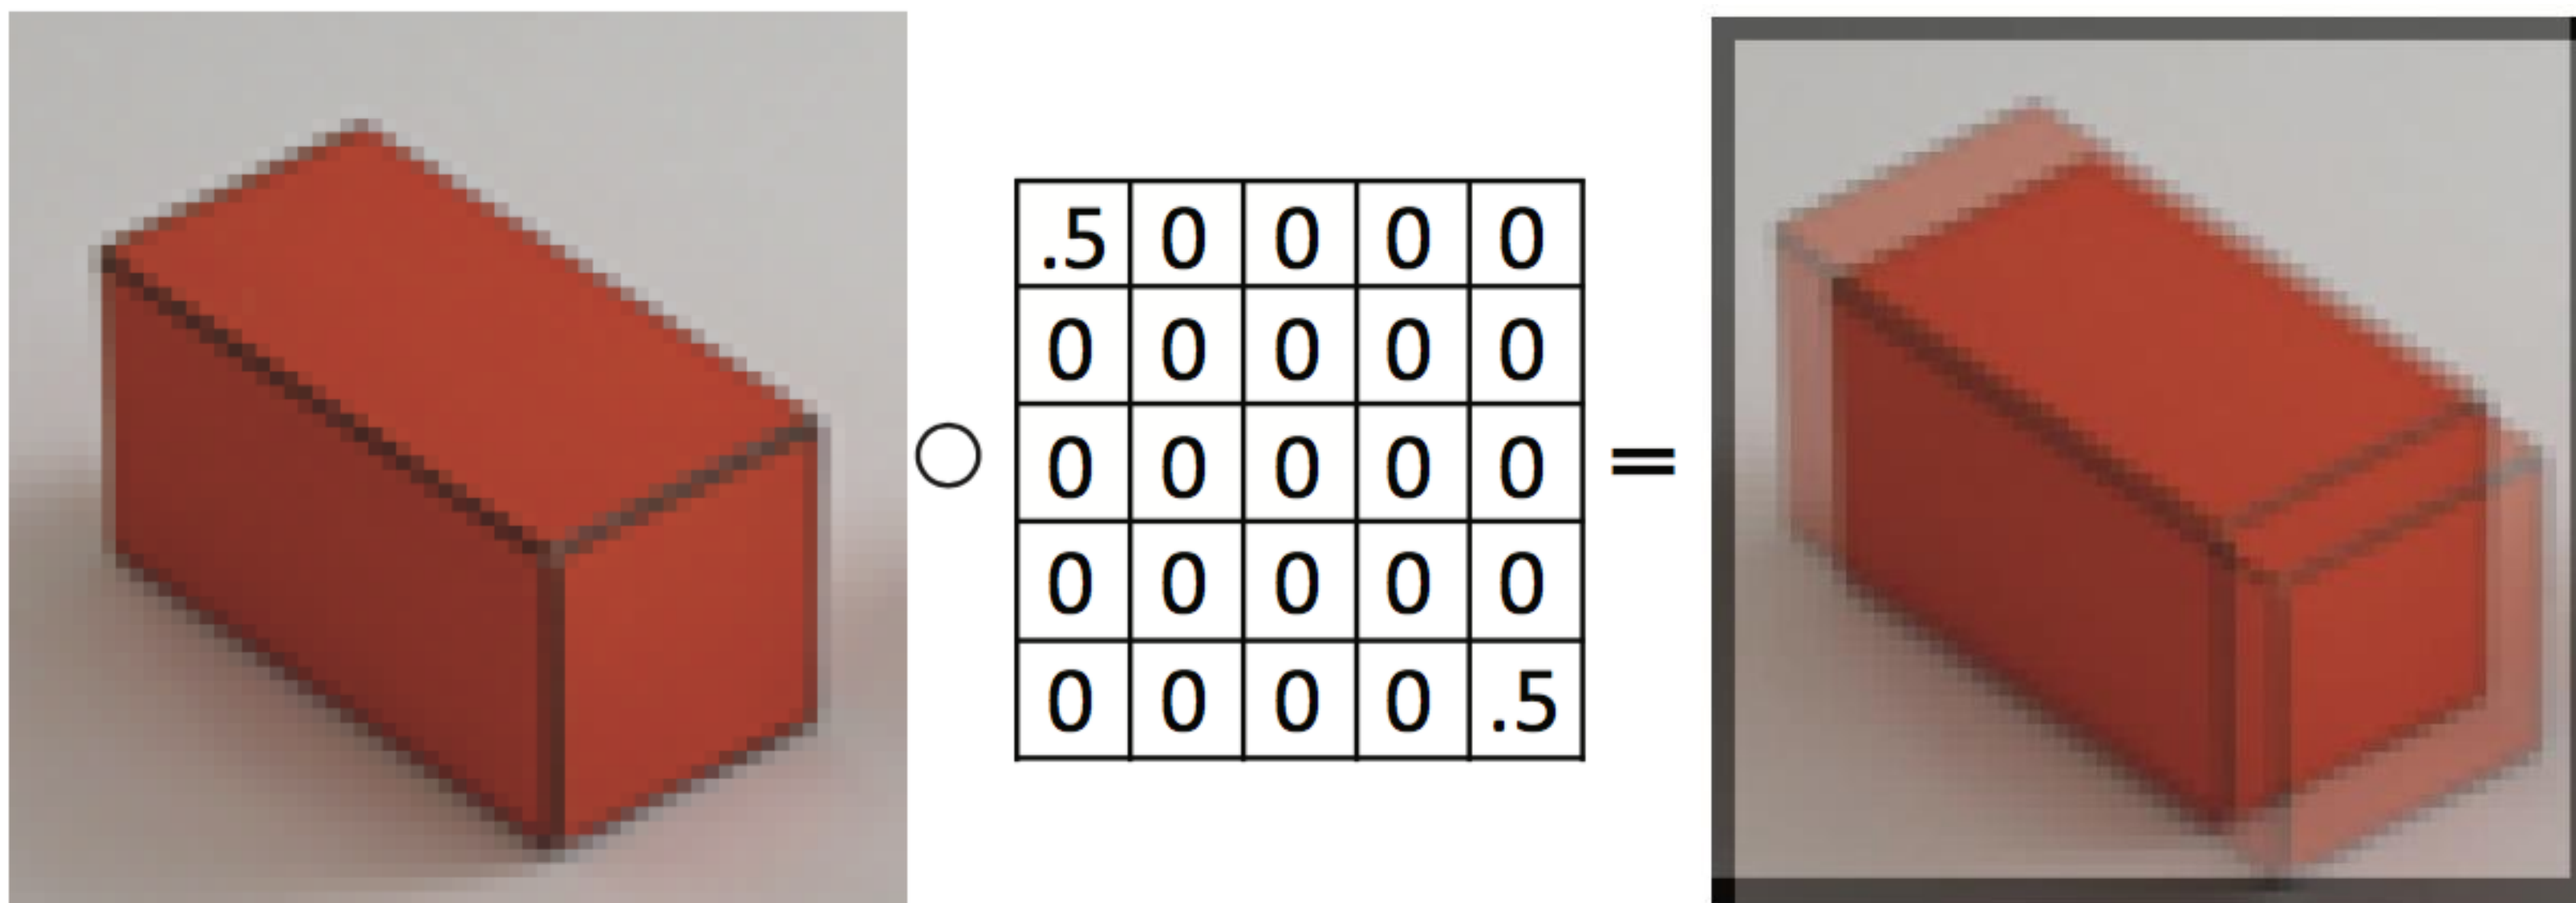


# A shift





# Examples



# “naturally” occurring filters

When we take a picture from a moving car, the resulting picture can be affected by motion blur



Input image



# “naturally” occurring filters

When we take a picture from a moving car, the resulting picture can be affected by motion blur



Input image



Motion blur

# Handling boundaries




# Handling boundaries

Zero padding

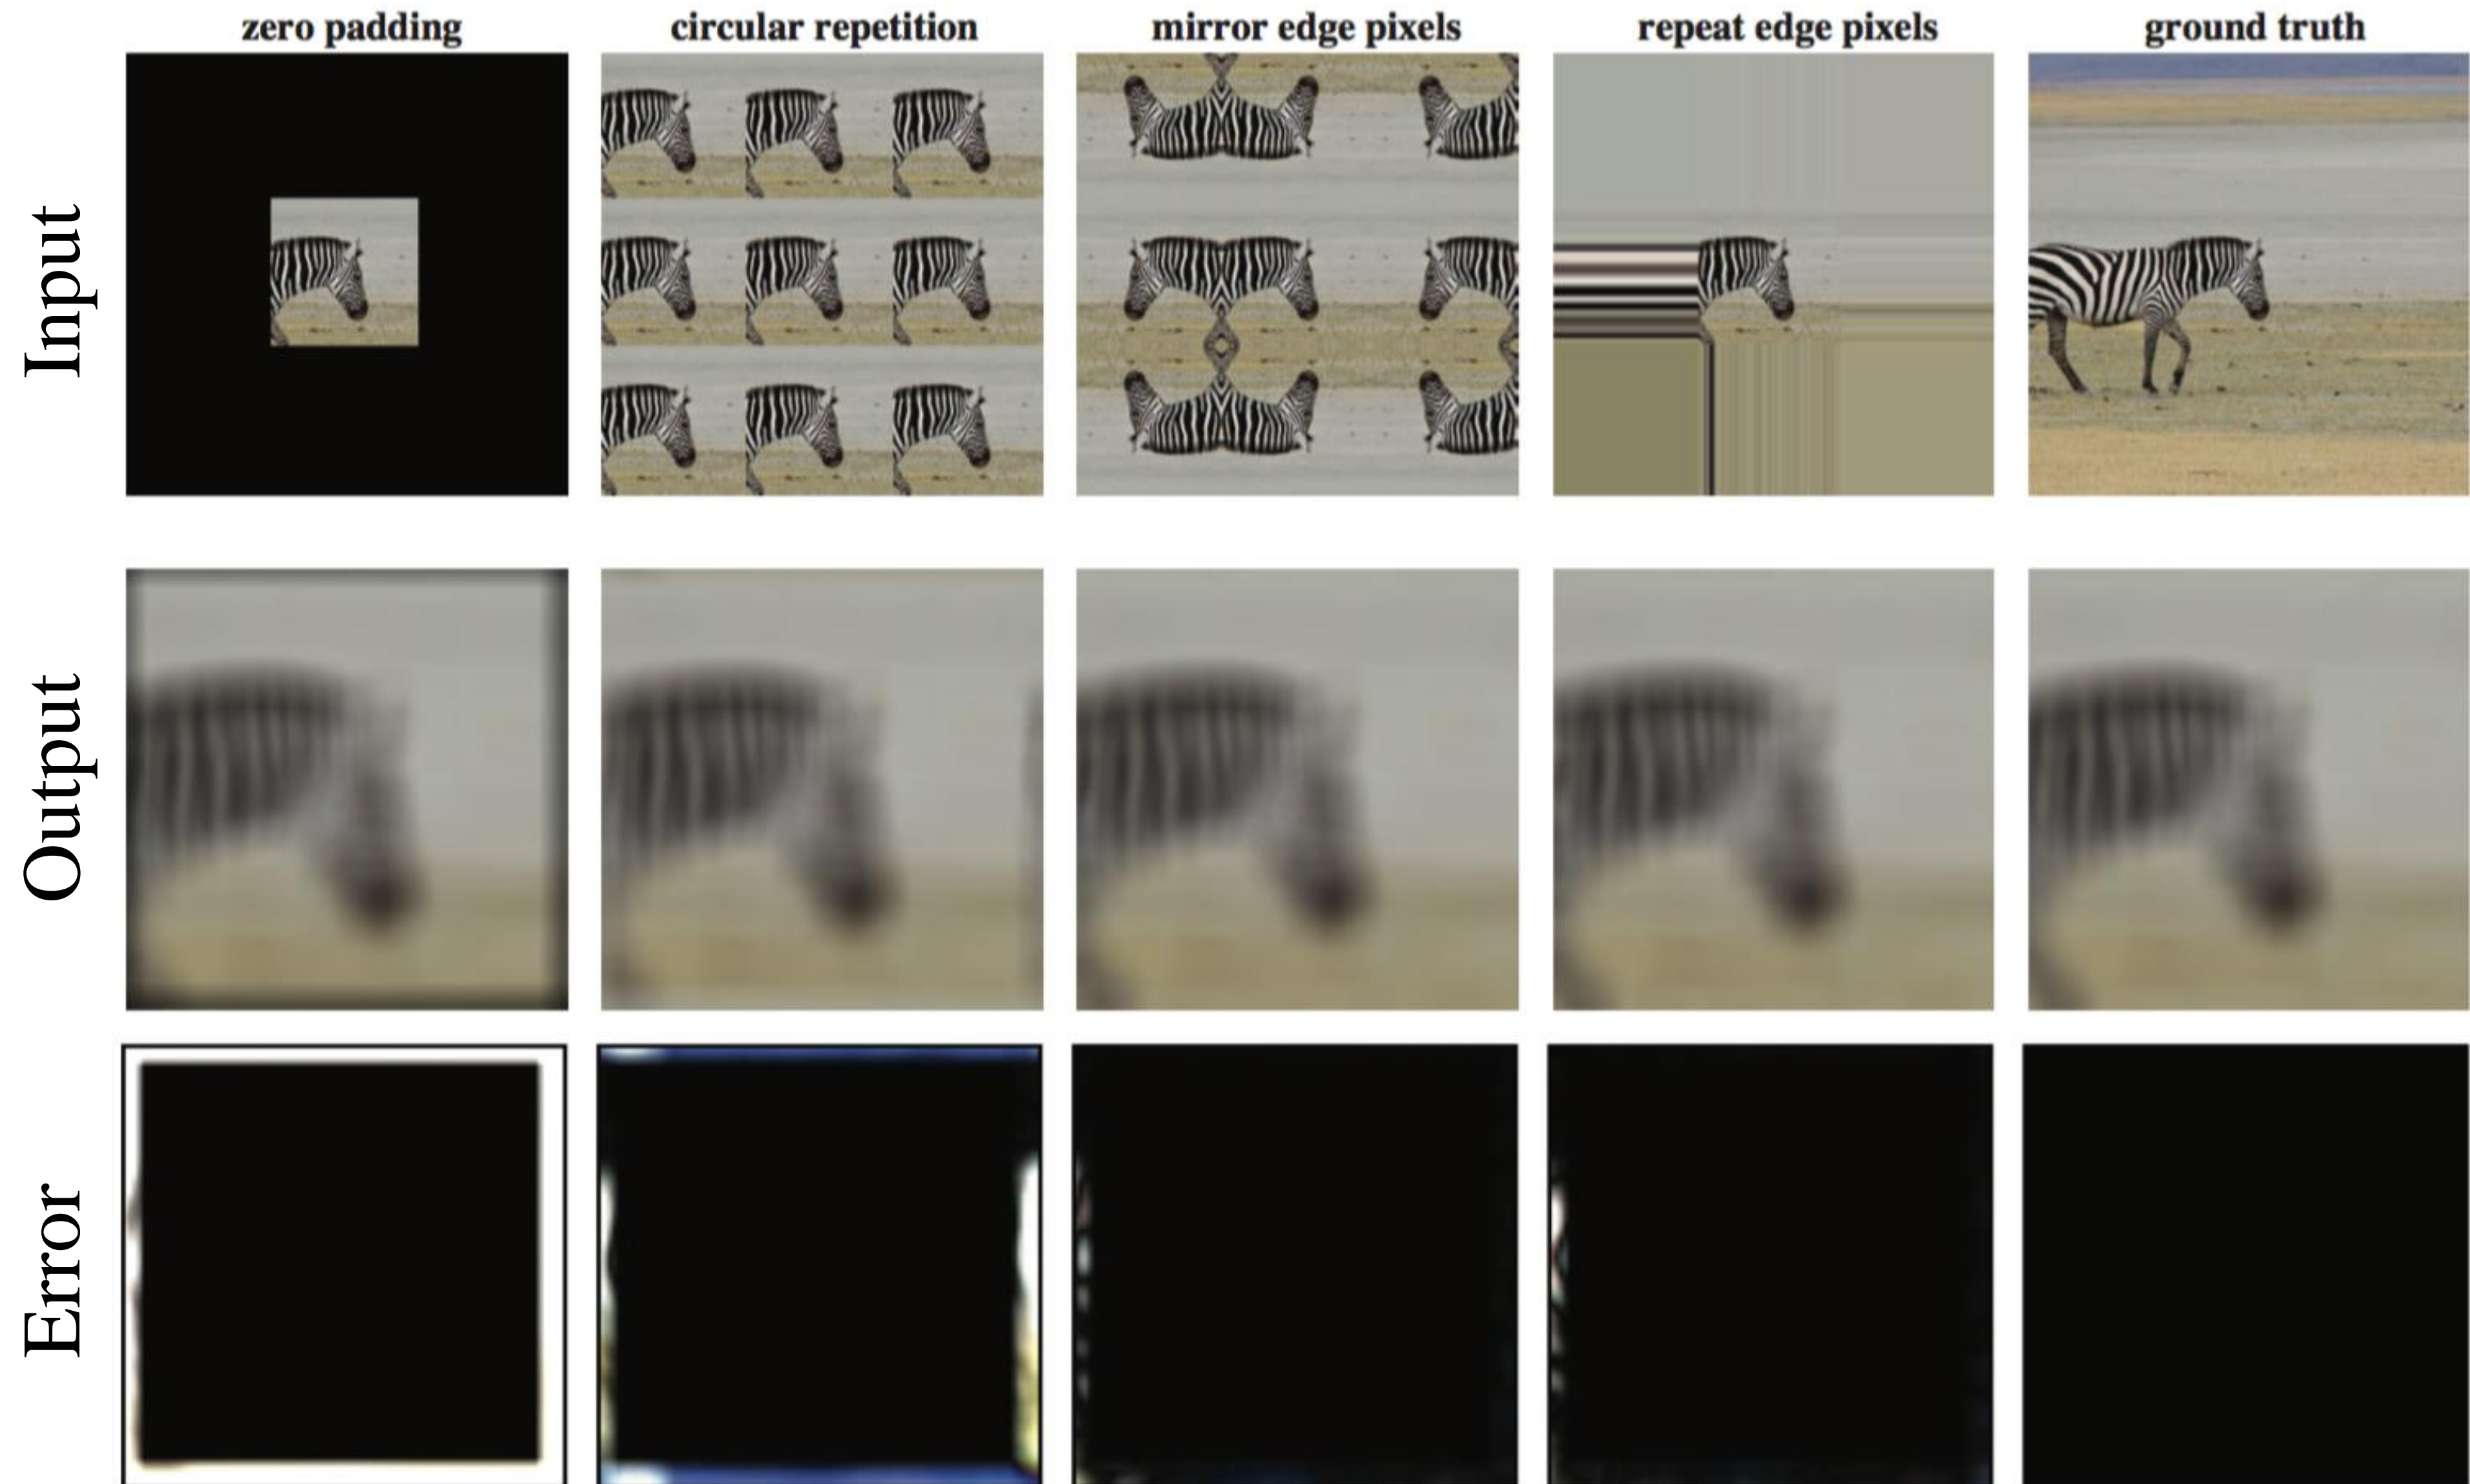


$$\bigcirc \quad \begin{array}{c} \blacksquare \\ \uparrow \\ 11 \times 11 \text{ ones} \end{array} =$$






# Handling boundaries







Some visual areas...



From M. Lewicky





**Figure 1.** Stimulus presentation scheme. The stimuli were originally calibrated to be seen at a distance of 150 cm in a 19" display.

# Campbell & Robson chart



Let's define the following image:

$$\mathbf{I}[n, m] = A[n] \sin(2\pi f[m] m/M)$$

With:



$$A[n] = A_{min} \left( \frac{A_{max}}{A_{min}} \right)^{n/N}$$

$$f[m] = f_{min} \left( \frac{f_{max}}{f_{min}} \right)^{m/M}$$




What do you think you should see when looking at this image?

$$\mathbf{I}[n, m] = A[n] \sin(2\pi f[m] m/M)$$



$$\mathbf{I}[n, m] = A[n] \sin(2\pi f[m] m/M)$$






# Contrast Sensitivity Function

Blackmore & Campbell (1969)

Maximum sensitivity

~ **6** cycles / degree of visual angle



Things that are very close  
and/or large are hard to see

Things far away  
are hard to see