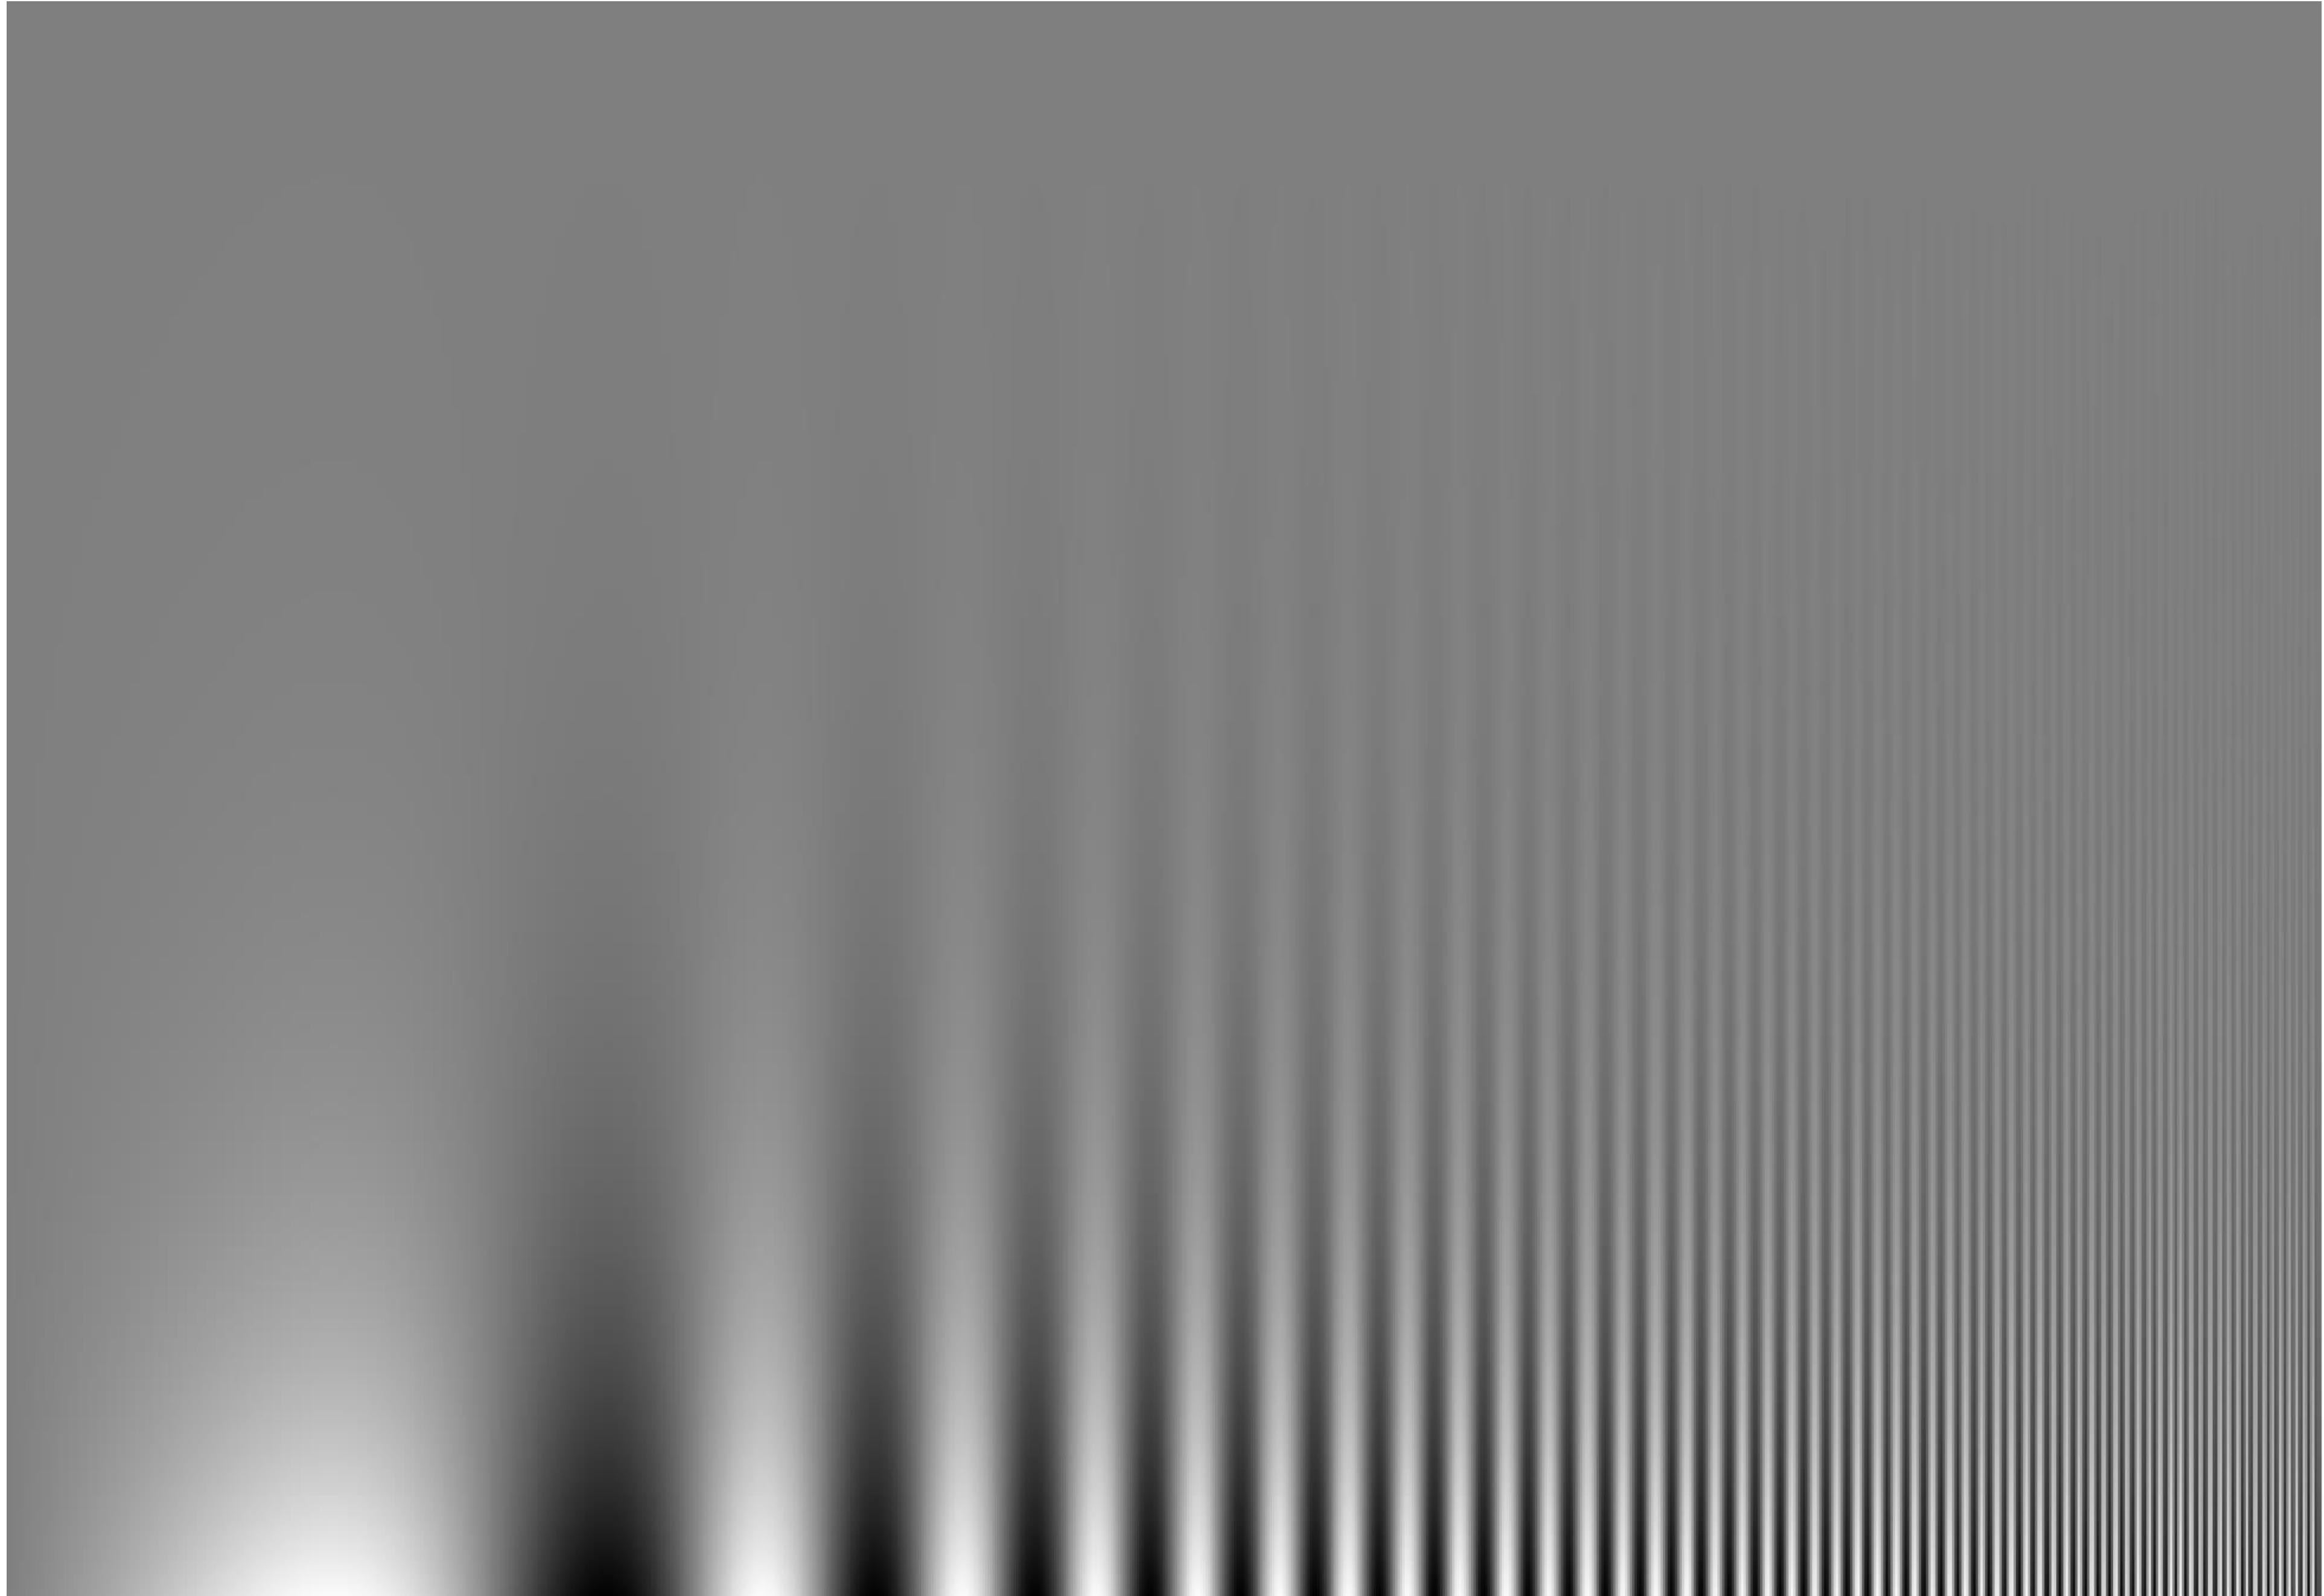


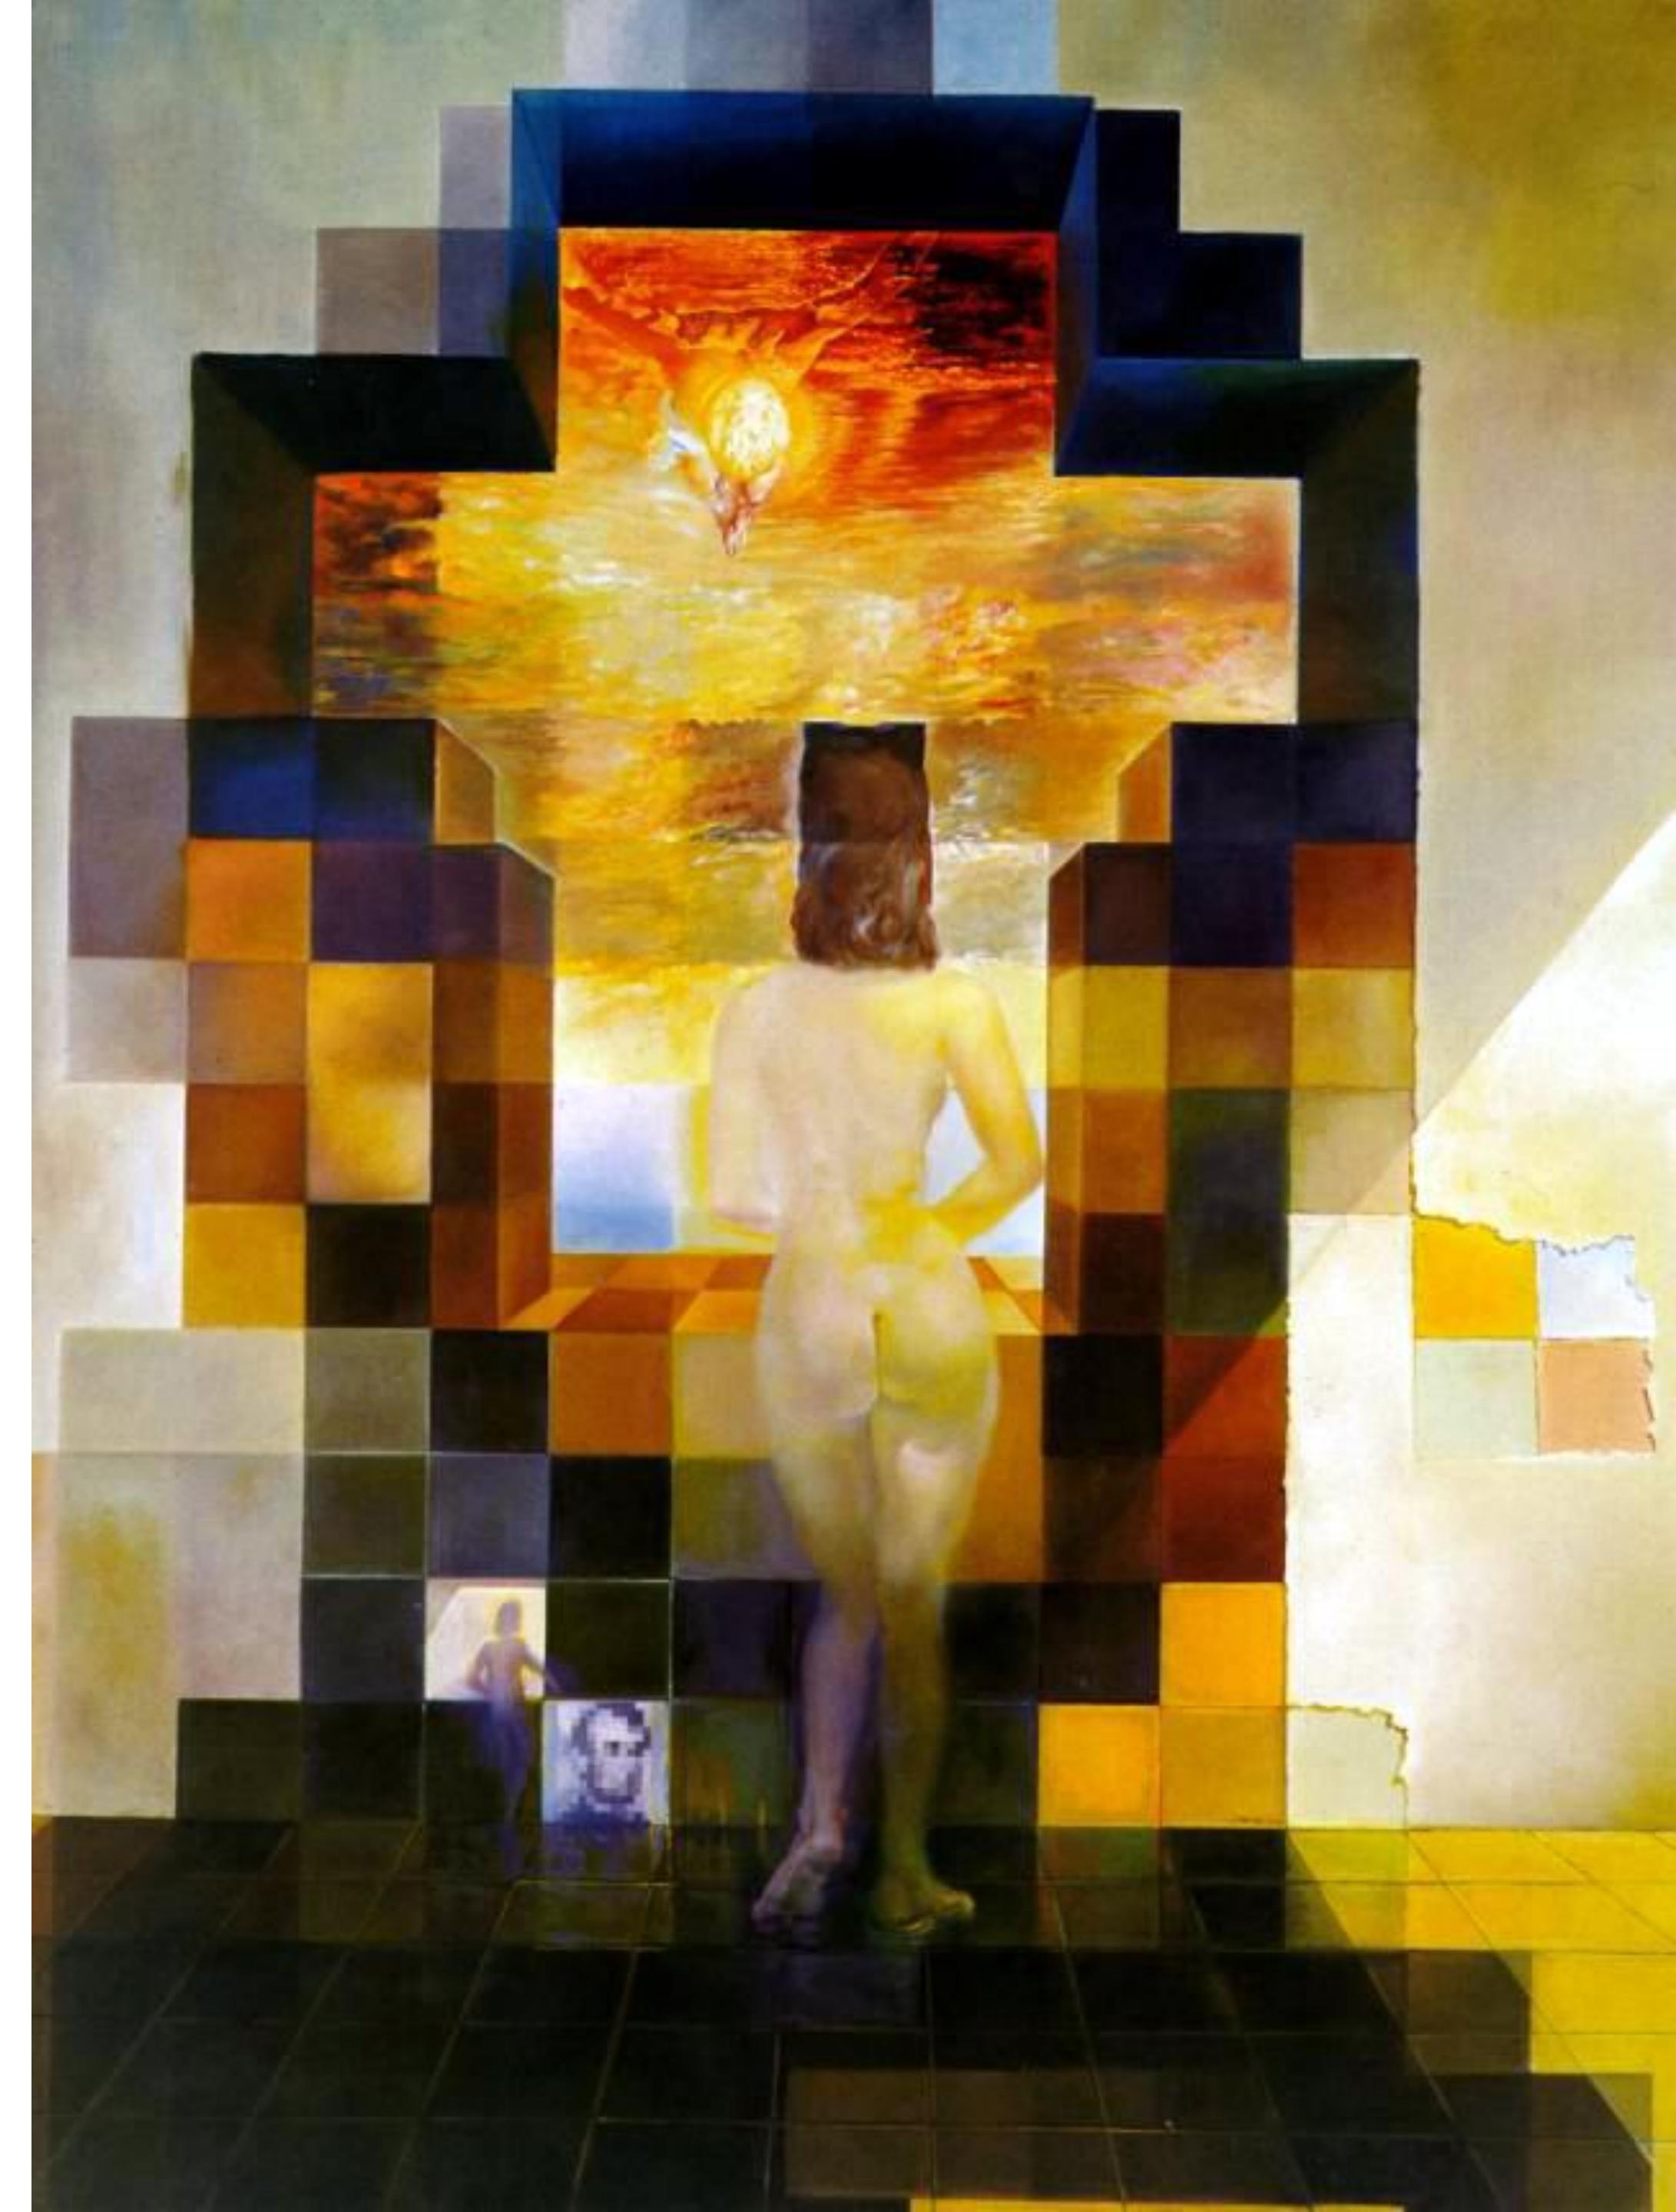
Spatial & Temporal Filtering

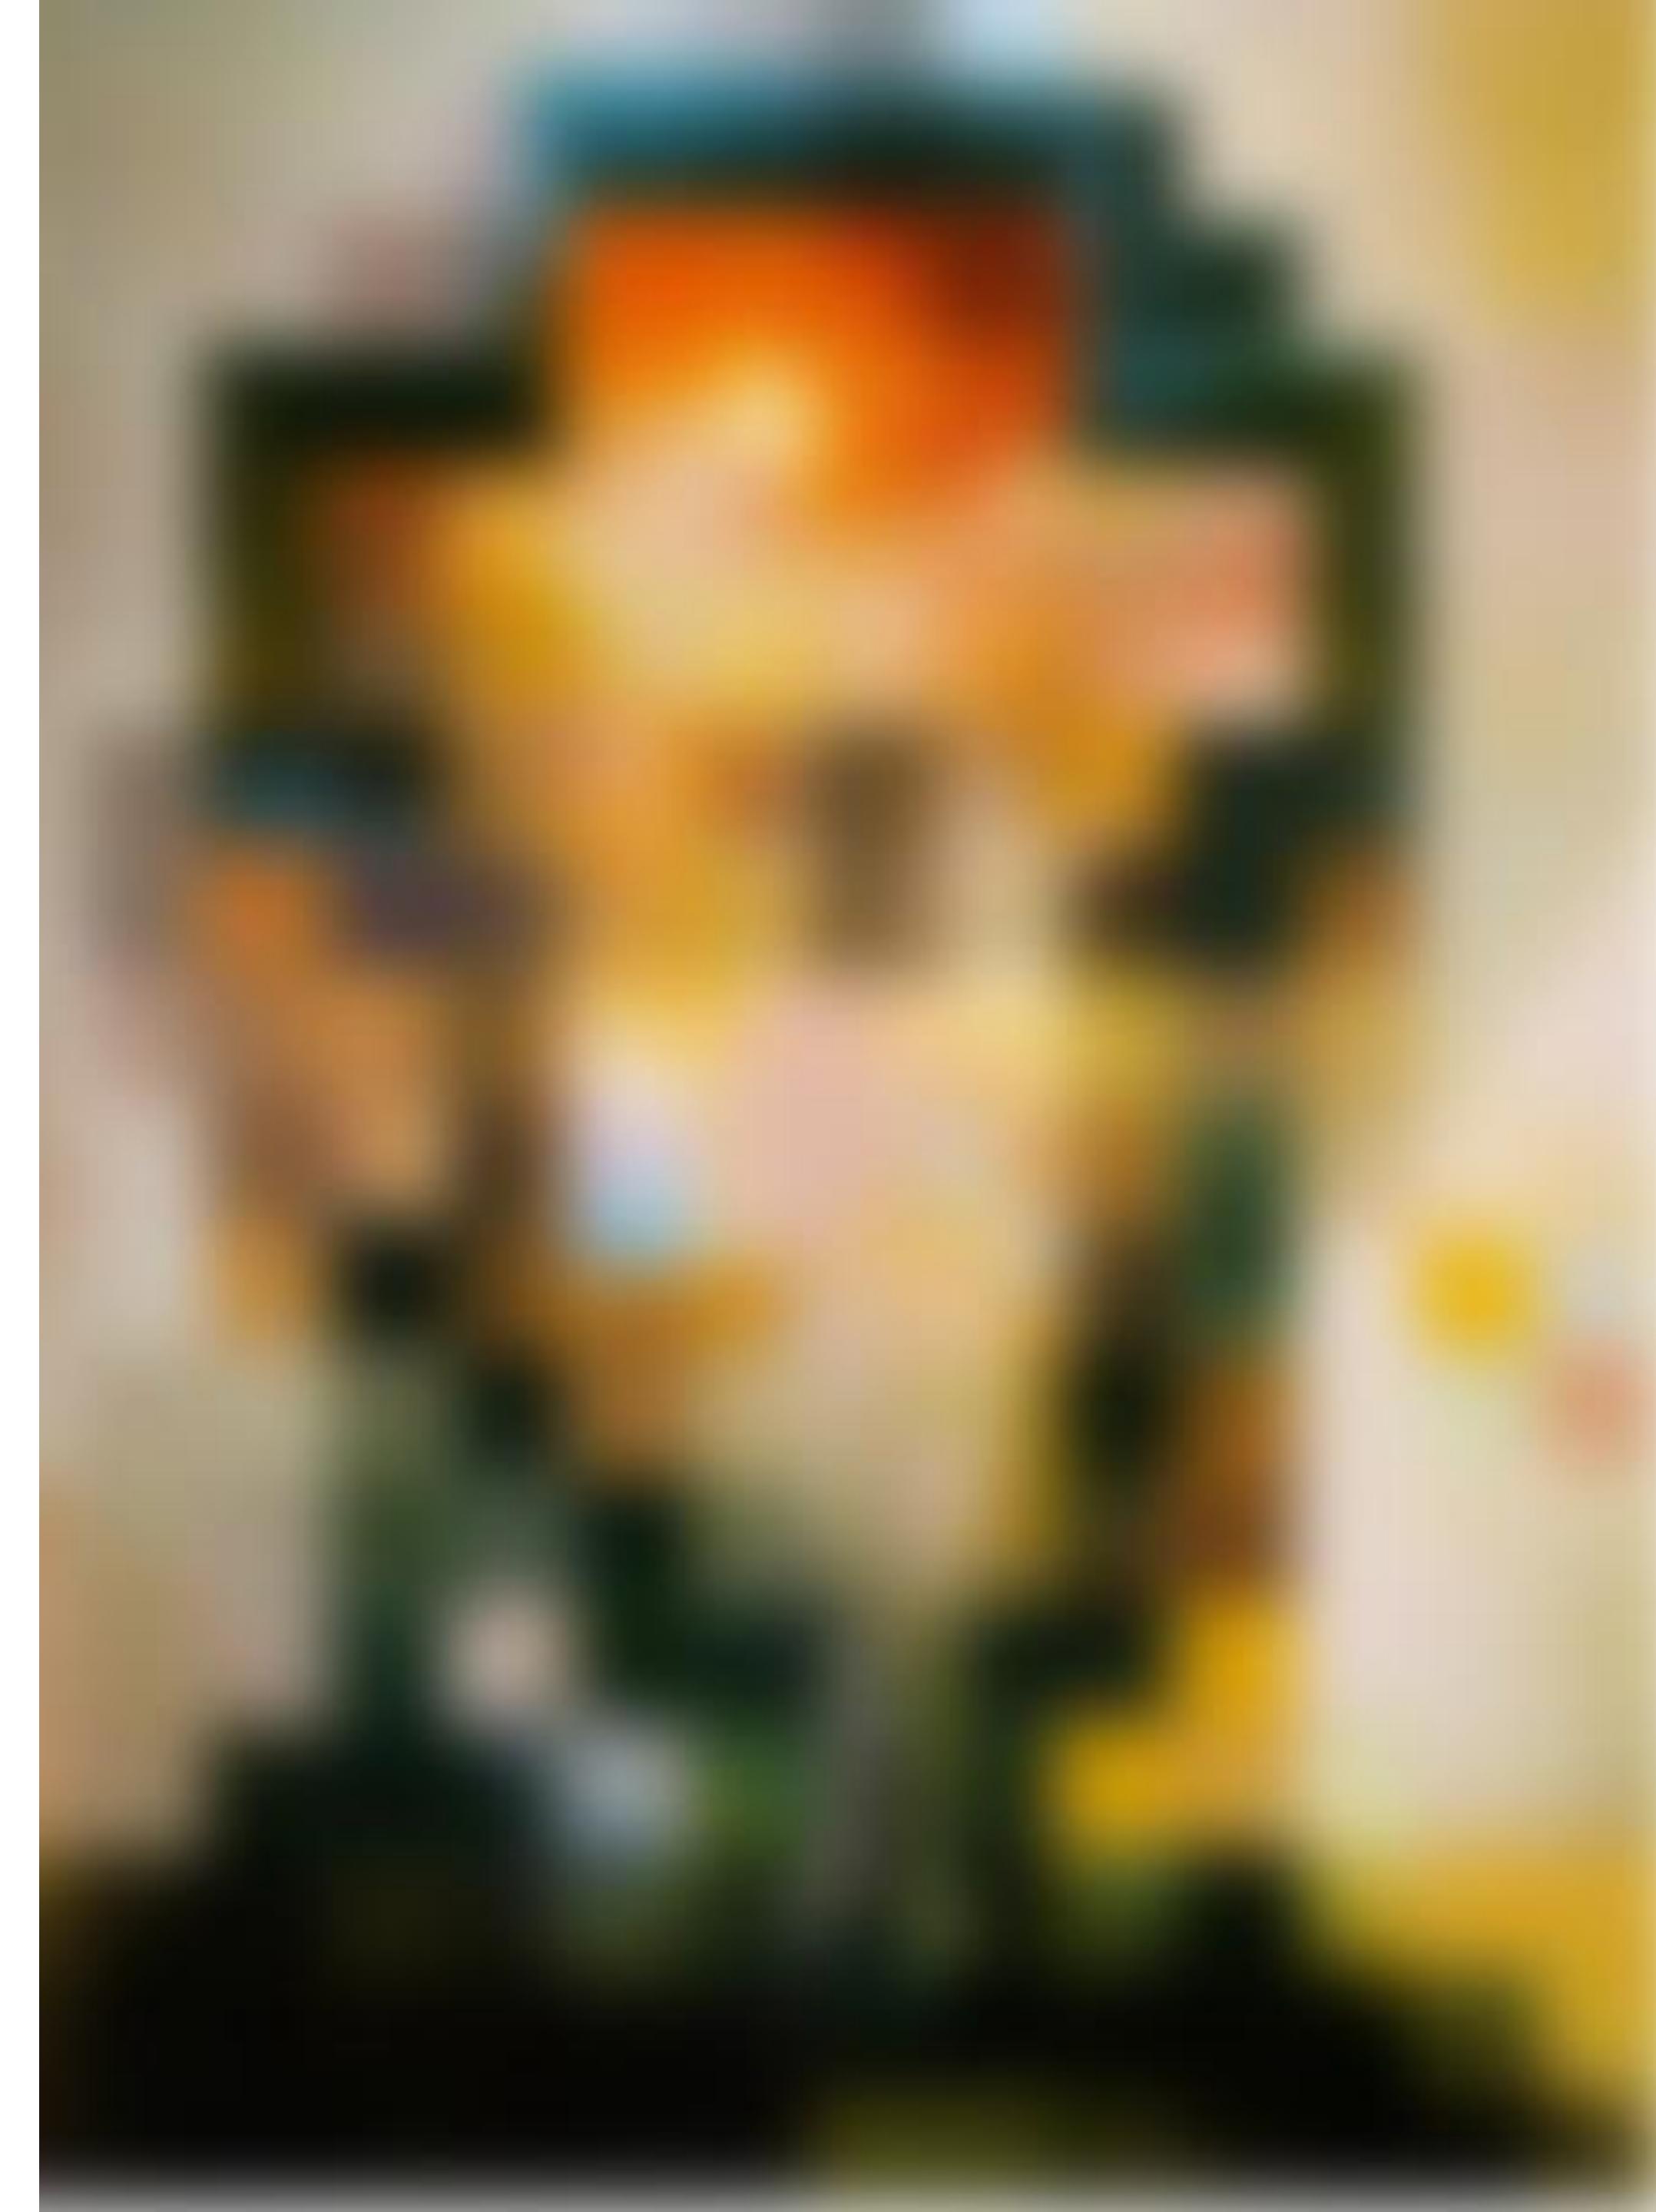


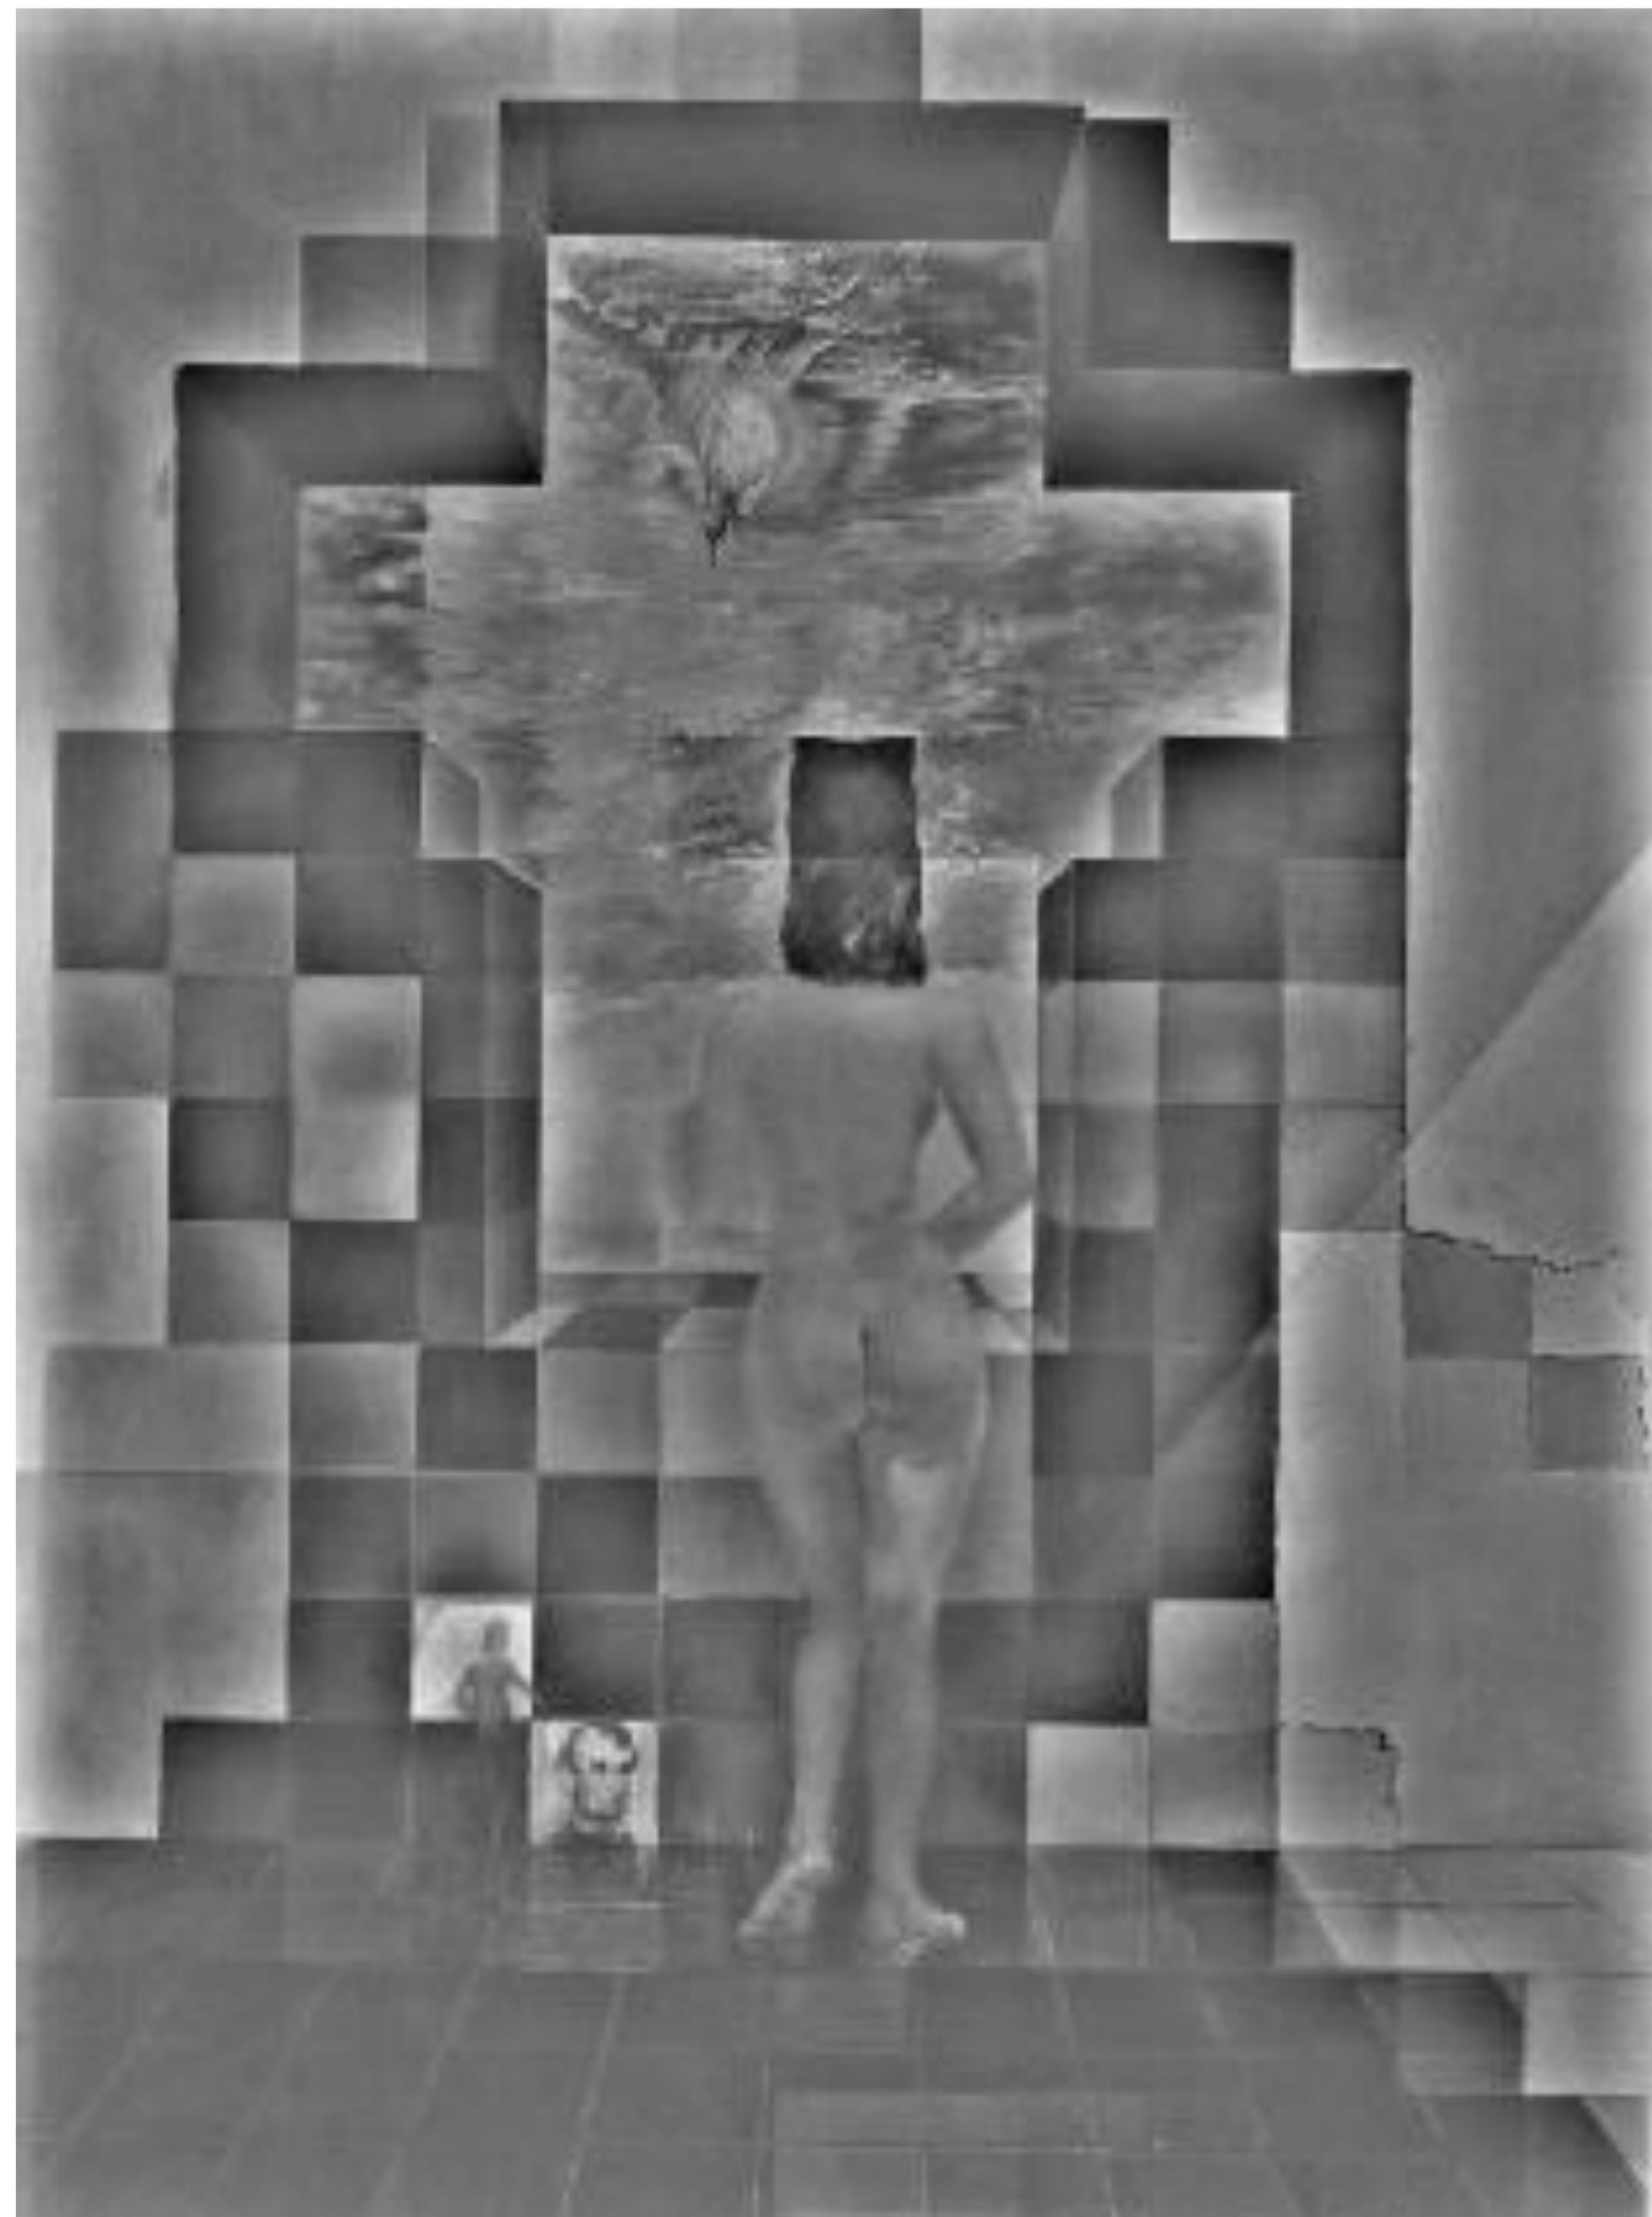


Salvador Dali

*“Gala Contemplating the Mediterranean Sea,
which at 30 meters becomes the portrait
of Abraham Lincoln”, 1976*









Low-pass filters



High/Band-pass filters



Low pass-filters

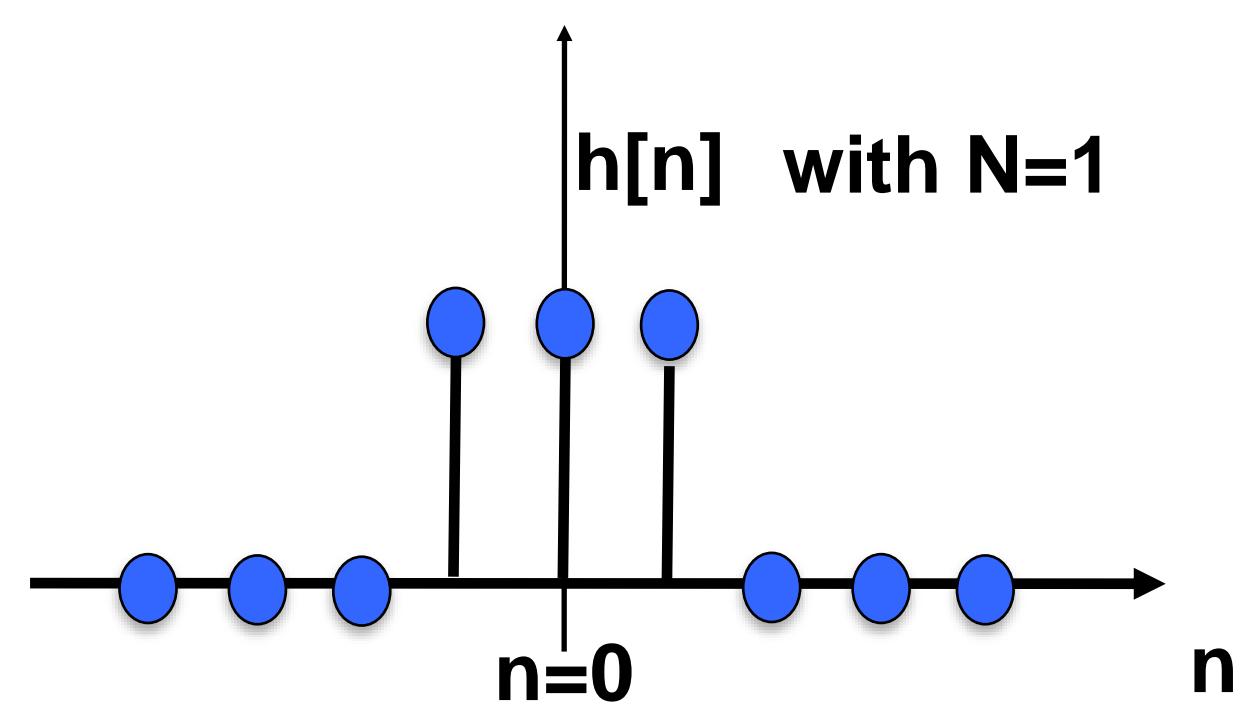
Box filter

$2N+1$

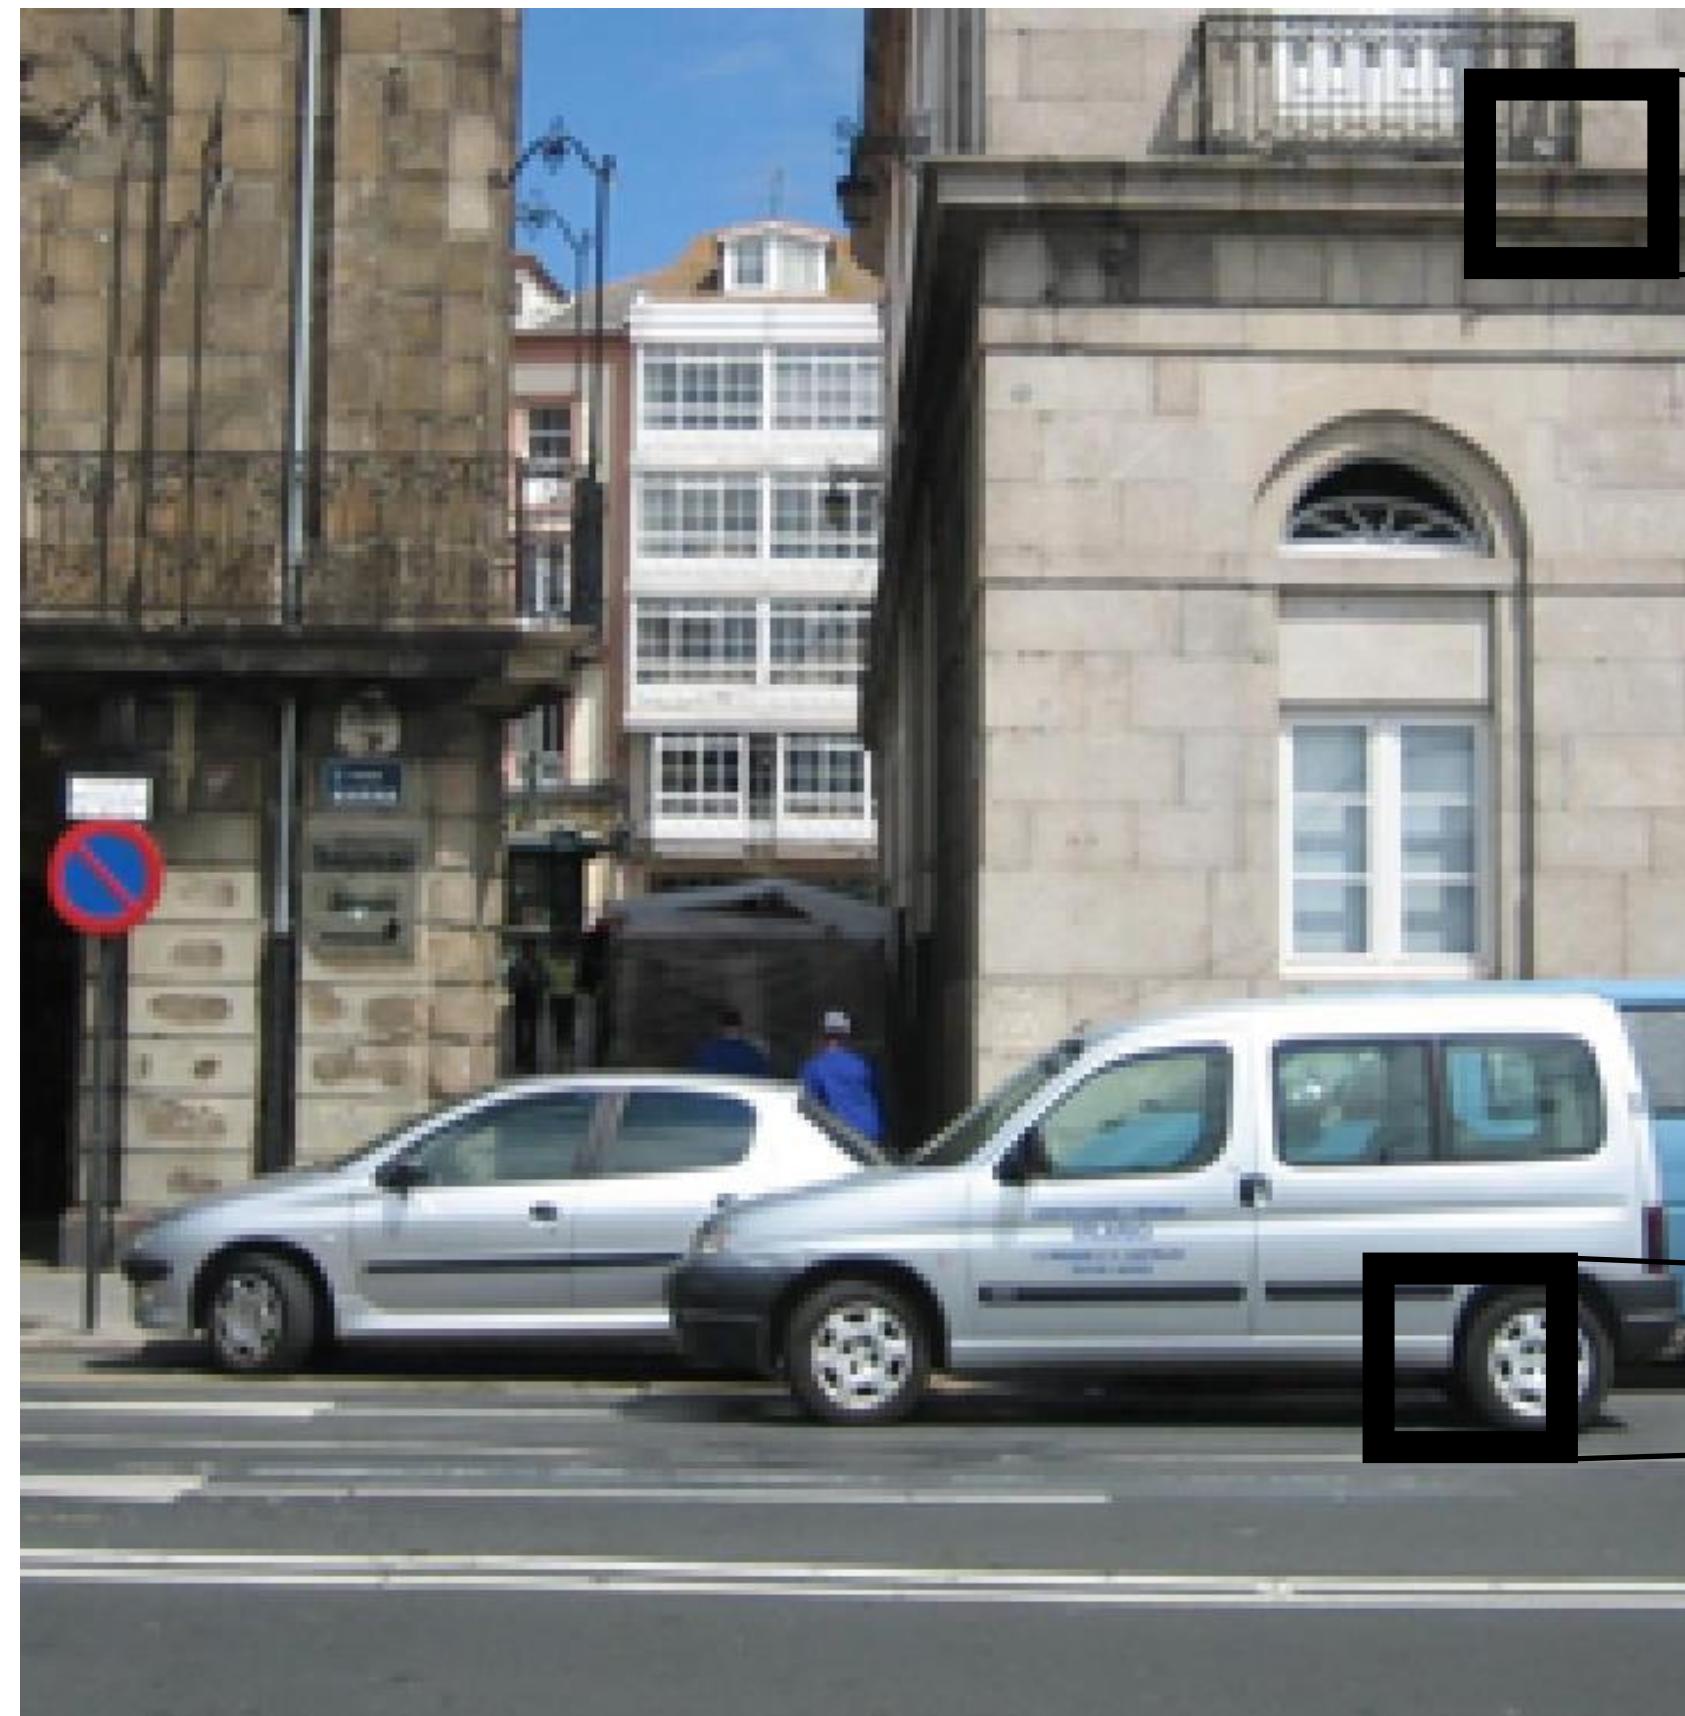
1	1	...	1
1	1		1
1	1		1
...			
1	1	1	1

$2M+1$

$$h_{N,M} [n, m] = \begin{cases} 1 & \text{if } -N \leq n \leq N \text{ and } -M \leq m \leq M \\ 0 & \text{otherwise} \end{cases}$$



Box filter



256X256

$$\bigcirc \frac{1}{21 \times 21} \boxed{} =$$



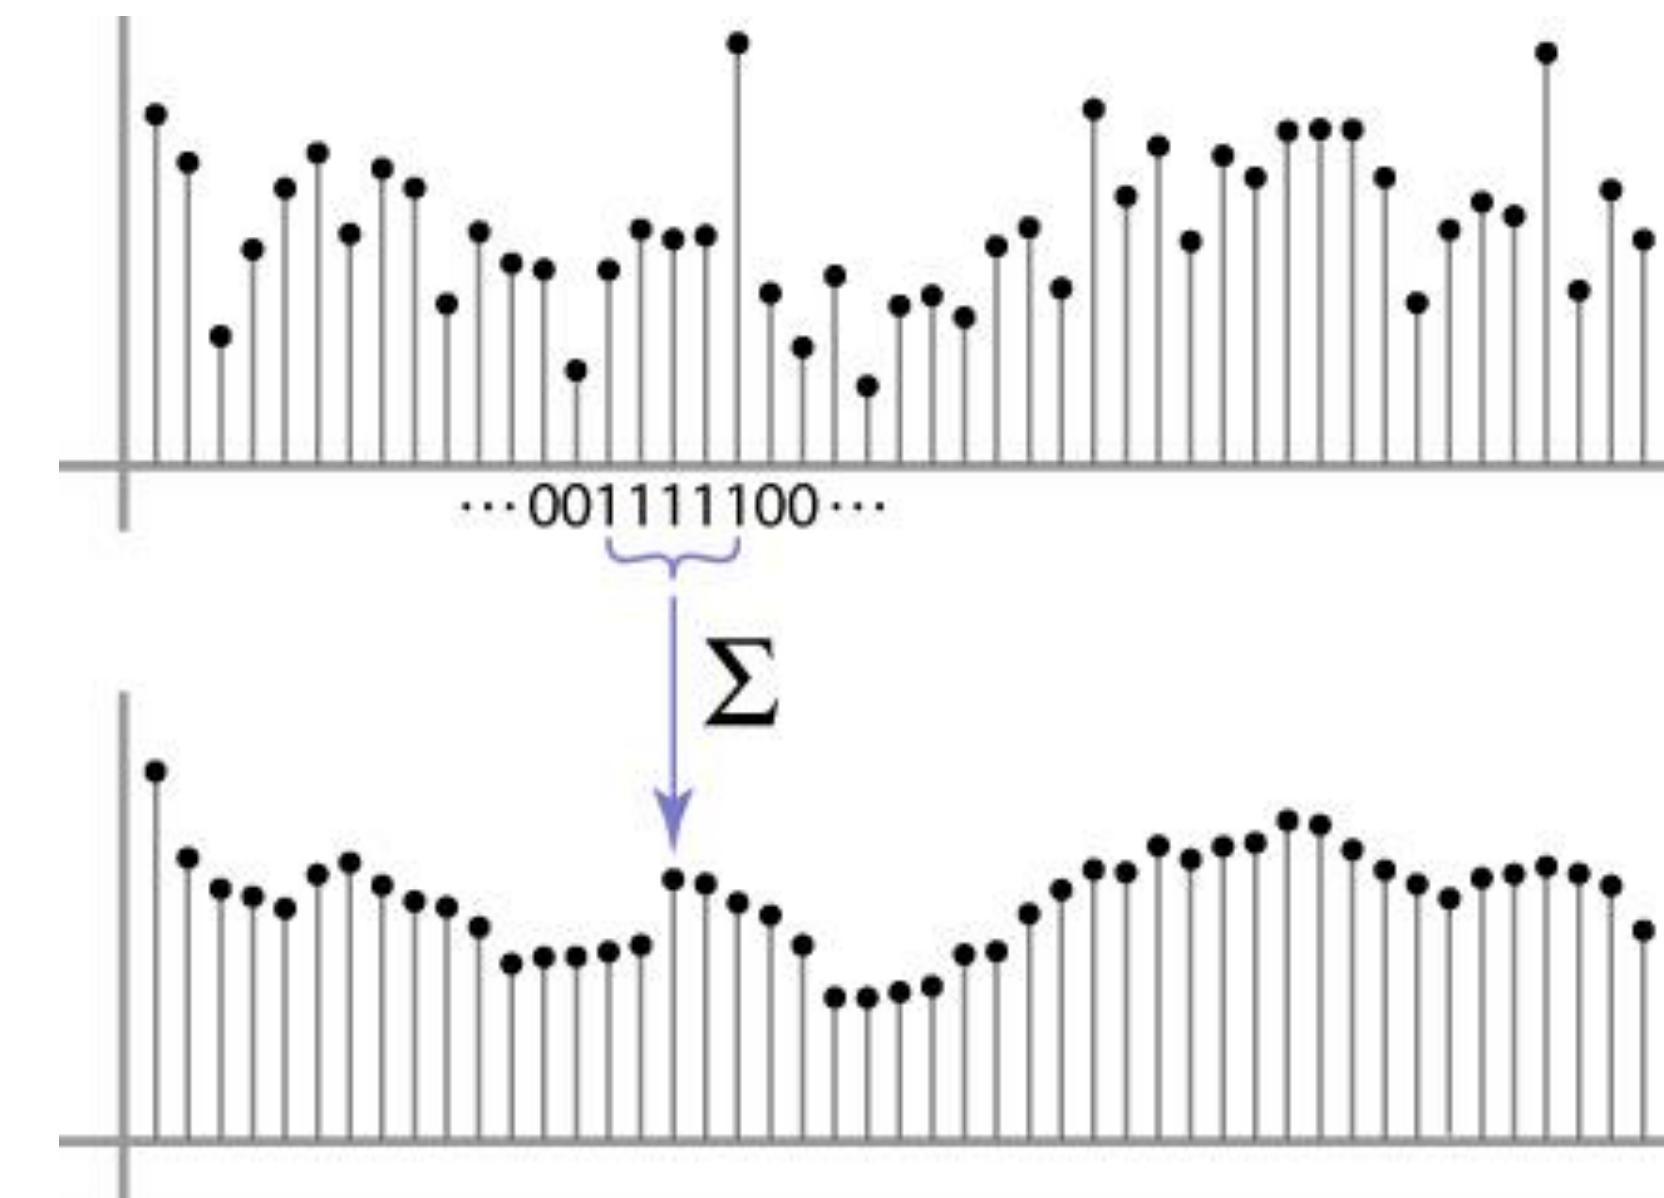
256X256

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

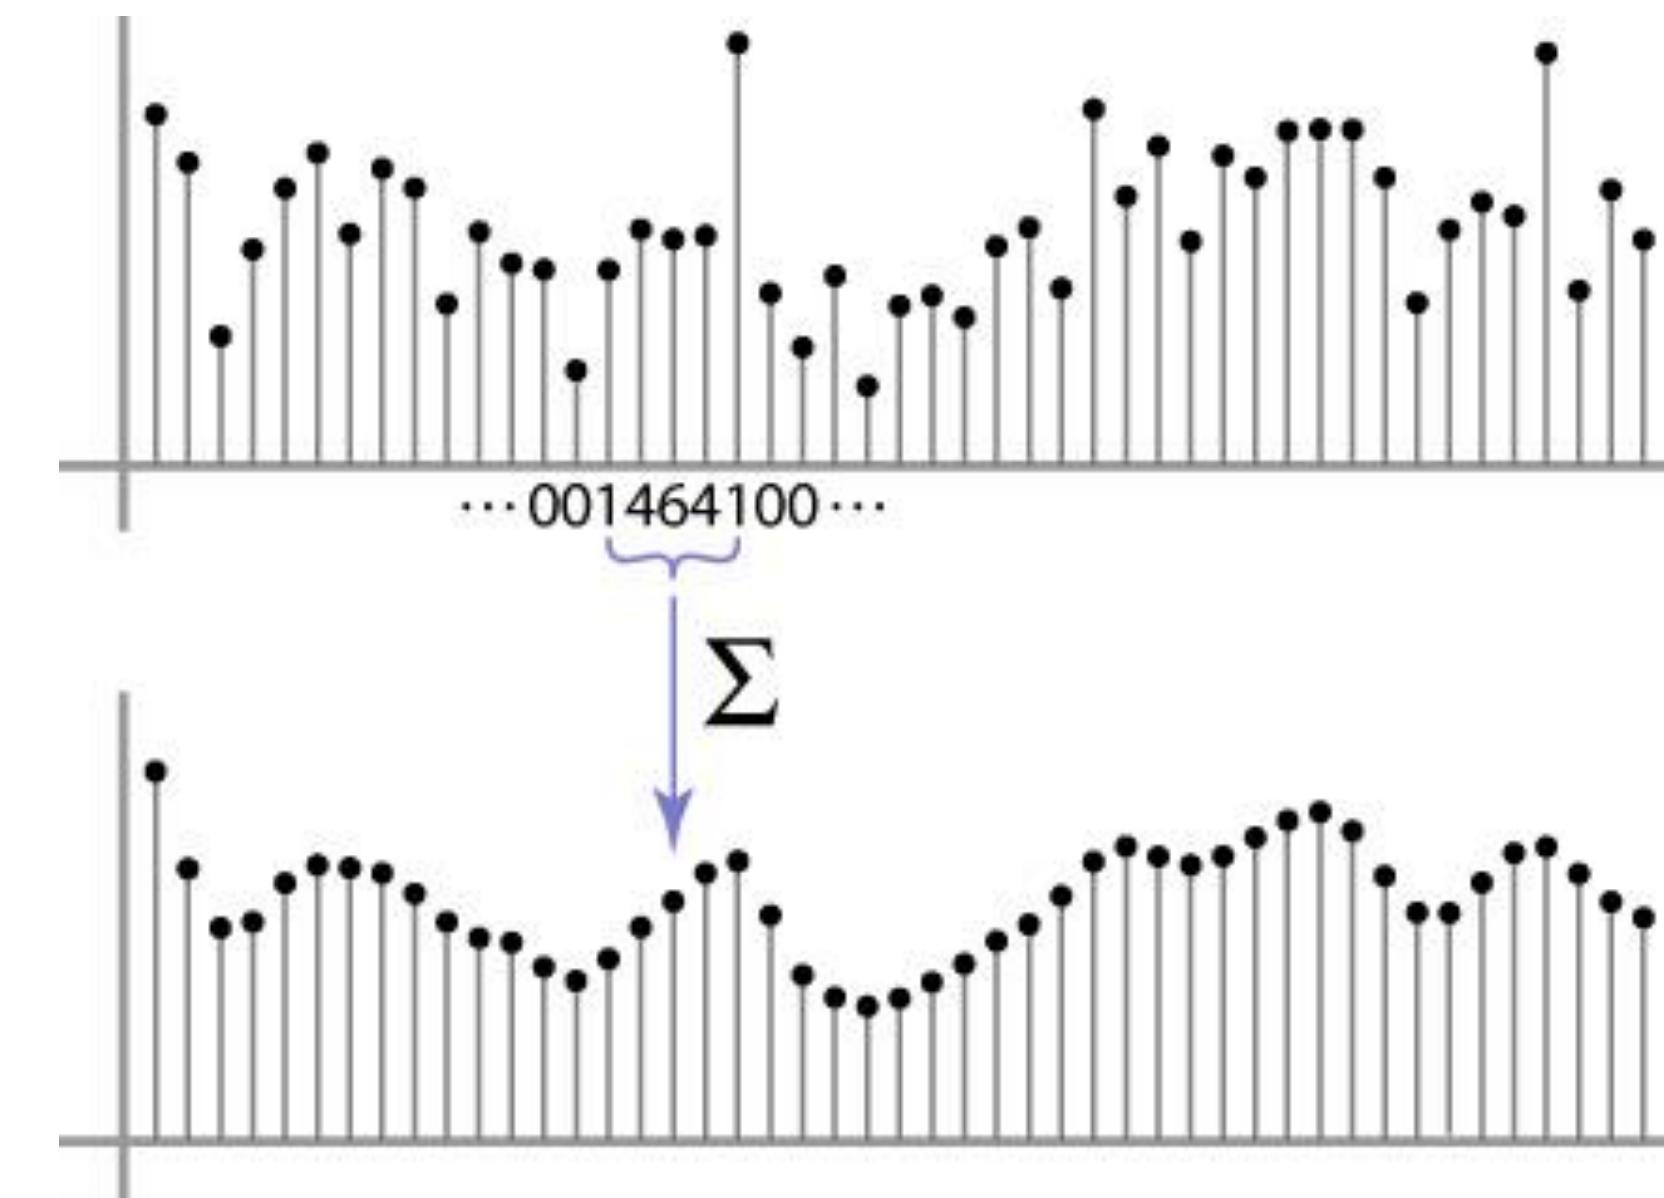
Moving Average

- Can add weights to our moving average
- $Weights \ [..., 0, 1, 1, 1, 1, 1, 0, ...] / 5$



Weighted Moving Average

- bell curve (gaussian-like) weights $[..., 1, 4, 6, 4, 1, ...]$

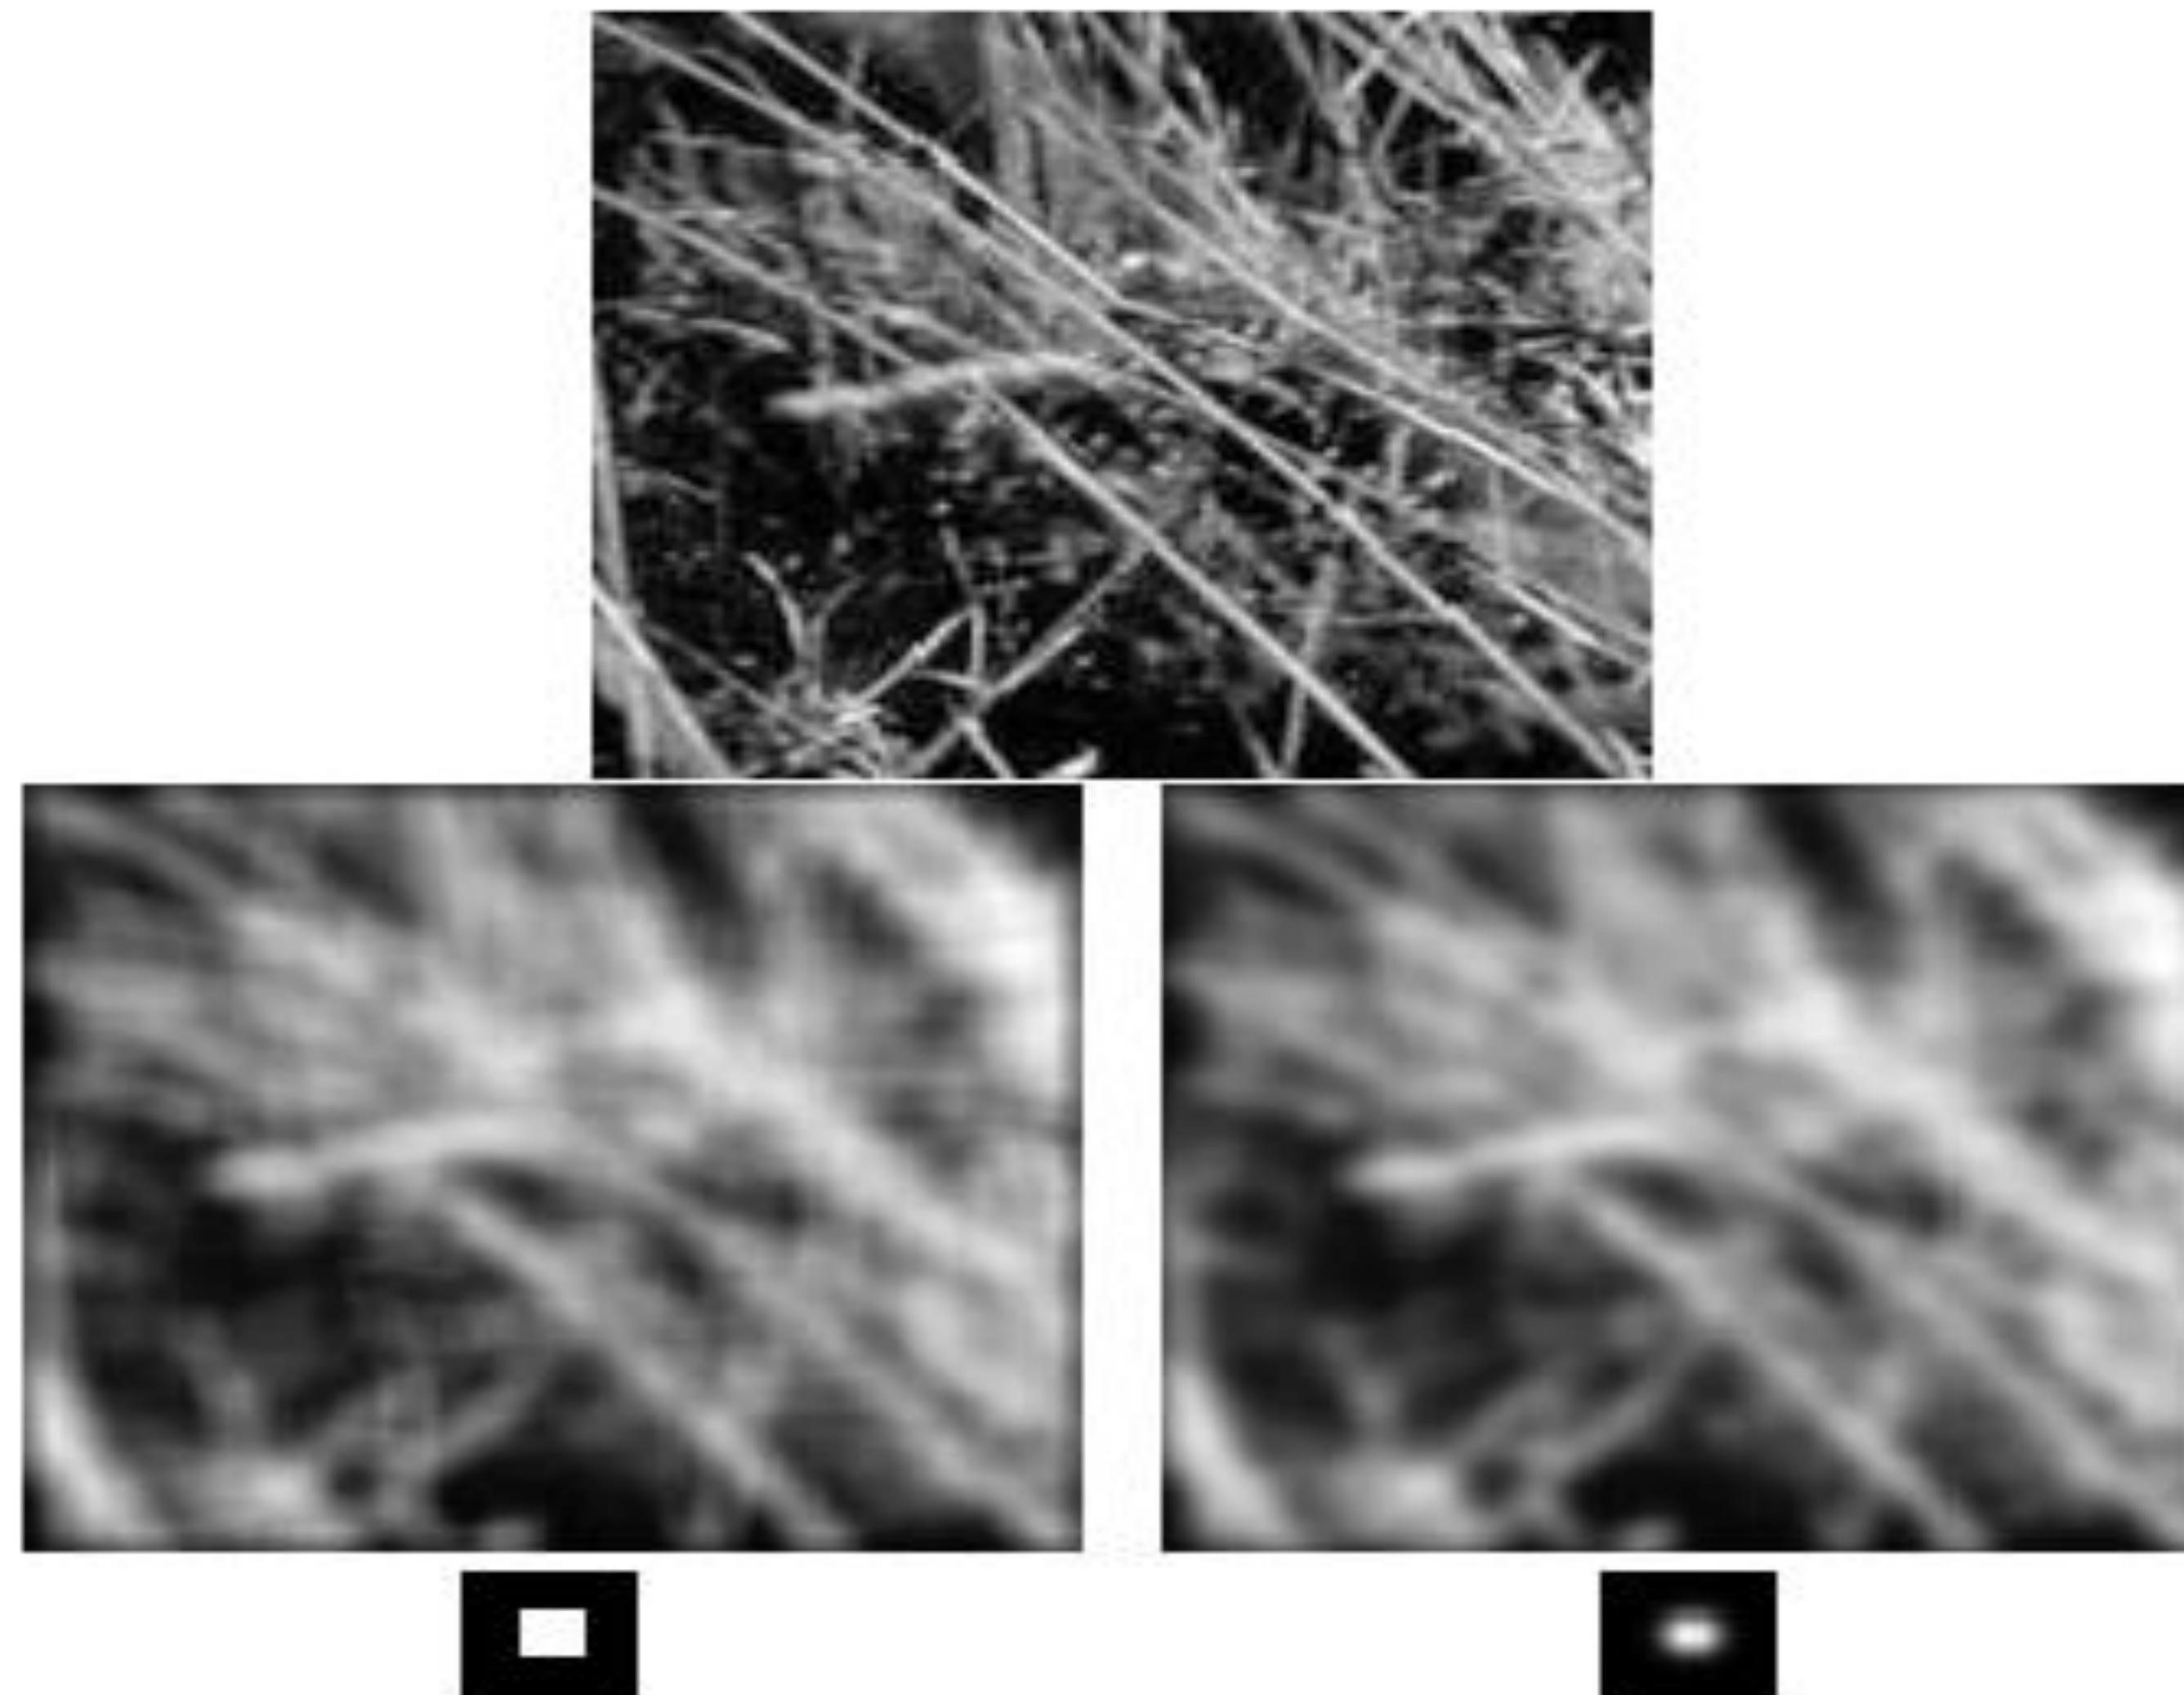


Why Box is bad

$$[1, 1, 1] \circ [\dots, 1, -1, 1, -1, 1, -1, \dots] = [\dots, -1, 1, -1, 1, -1, 1, \dots]$$

$$[1, 2, 1] \circ [\dots, 1, -1, 1, -1, 1, -1, \dots] = [\dots, 0, 0, 0, 0, 0, 0, \dots]$$

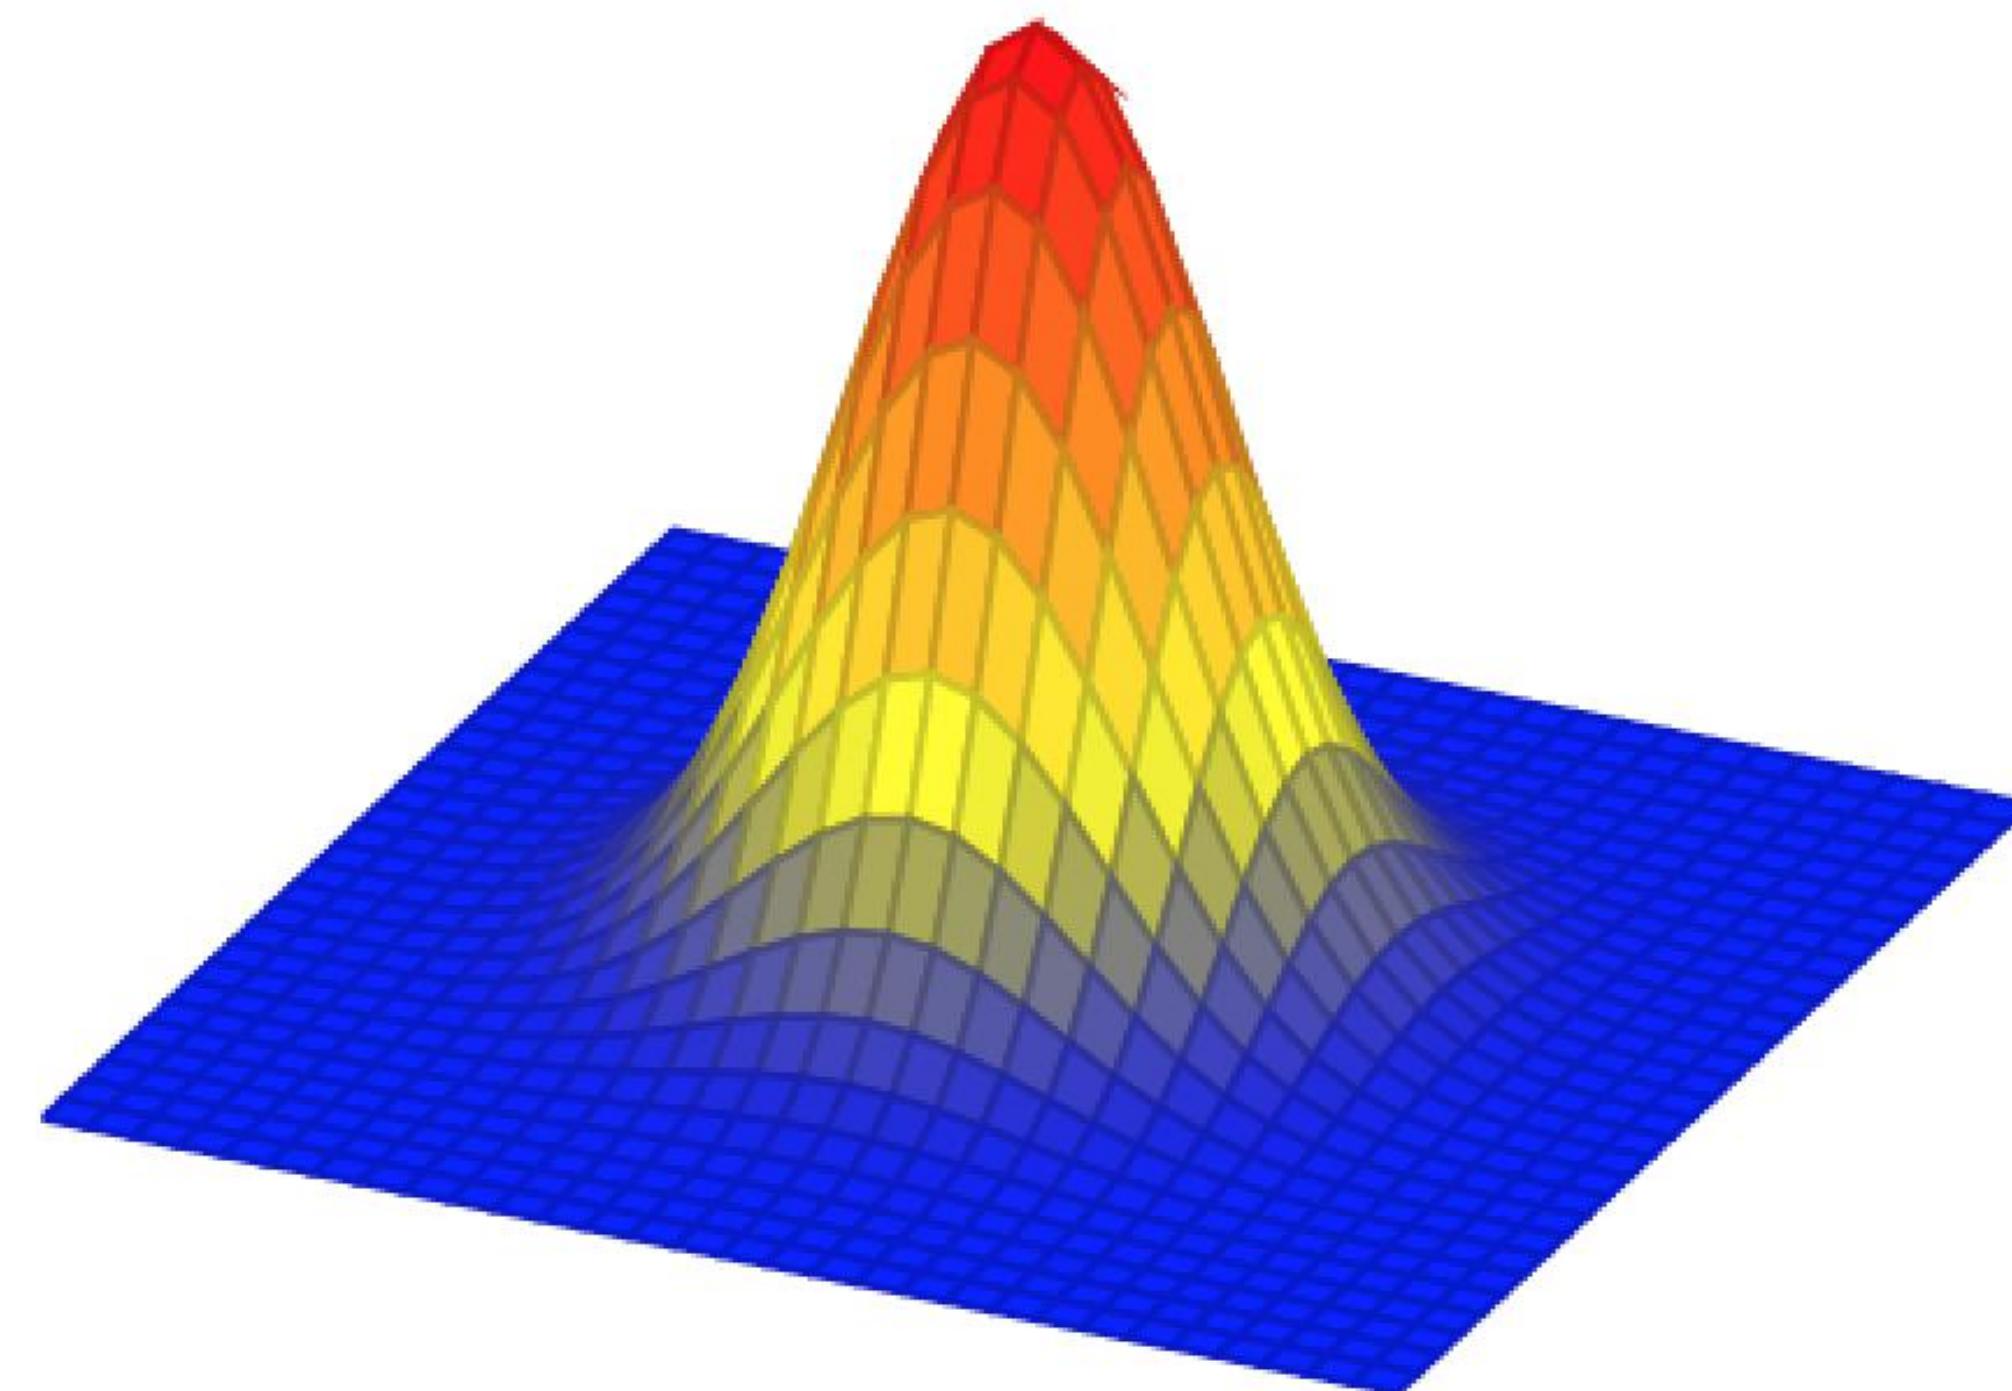
Mean vs. Gaussian filtering



Gaussian filter

In the continuous domain:

$$g(x, y; \sigma) = \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2}$$



Gaussian filter

$$g(x, y; \sigma) = \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2}$$

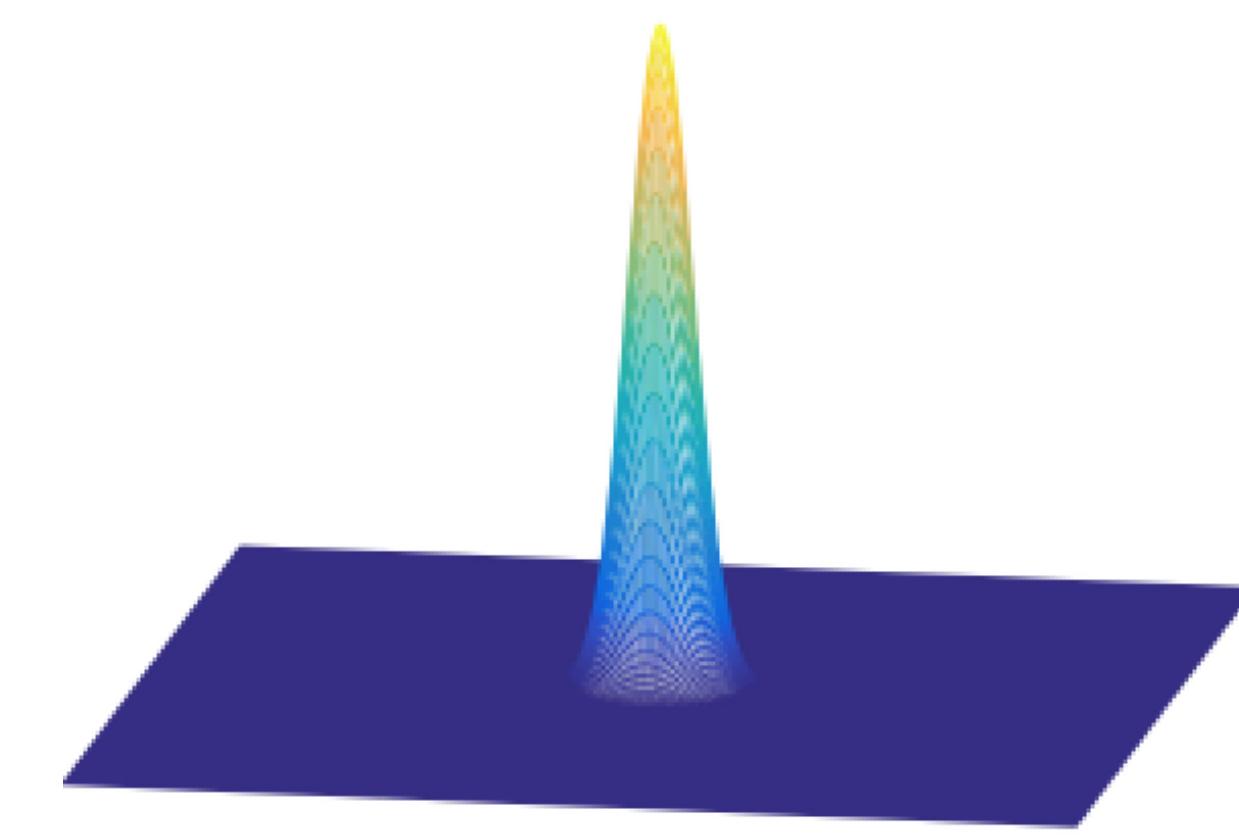
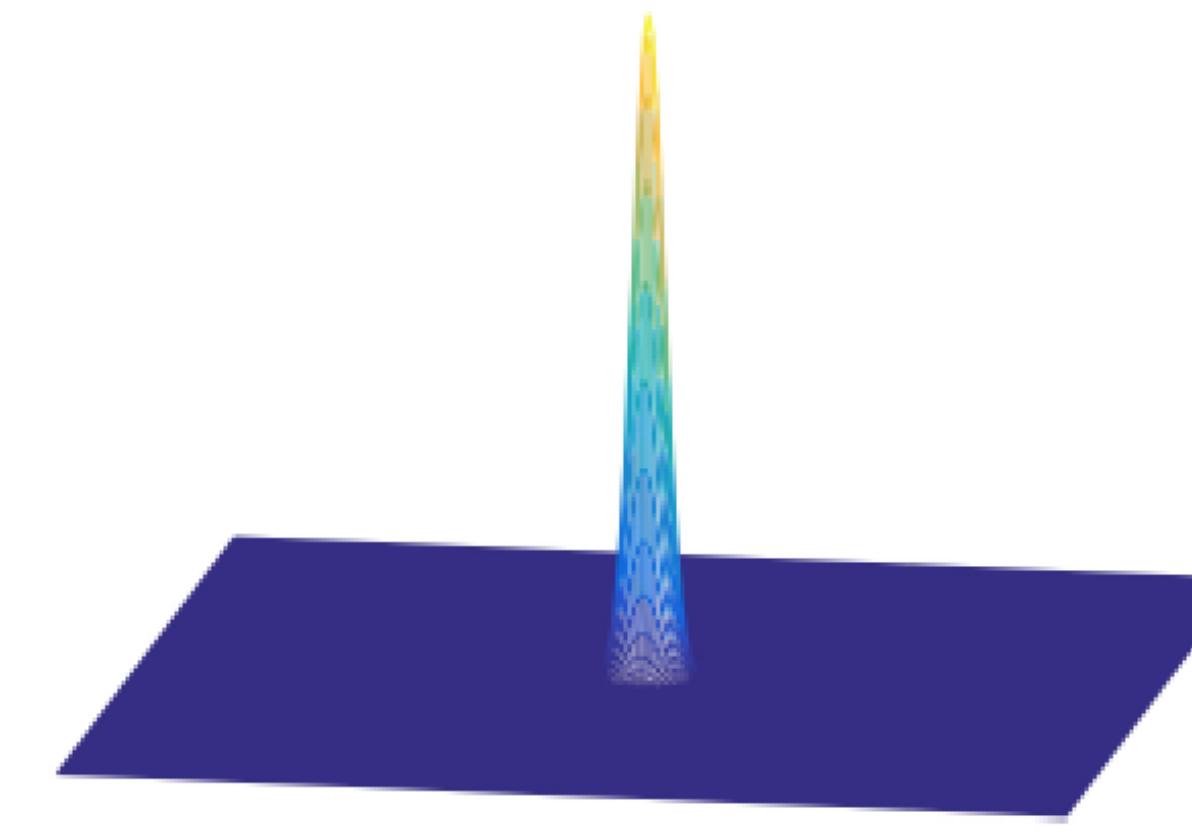
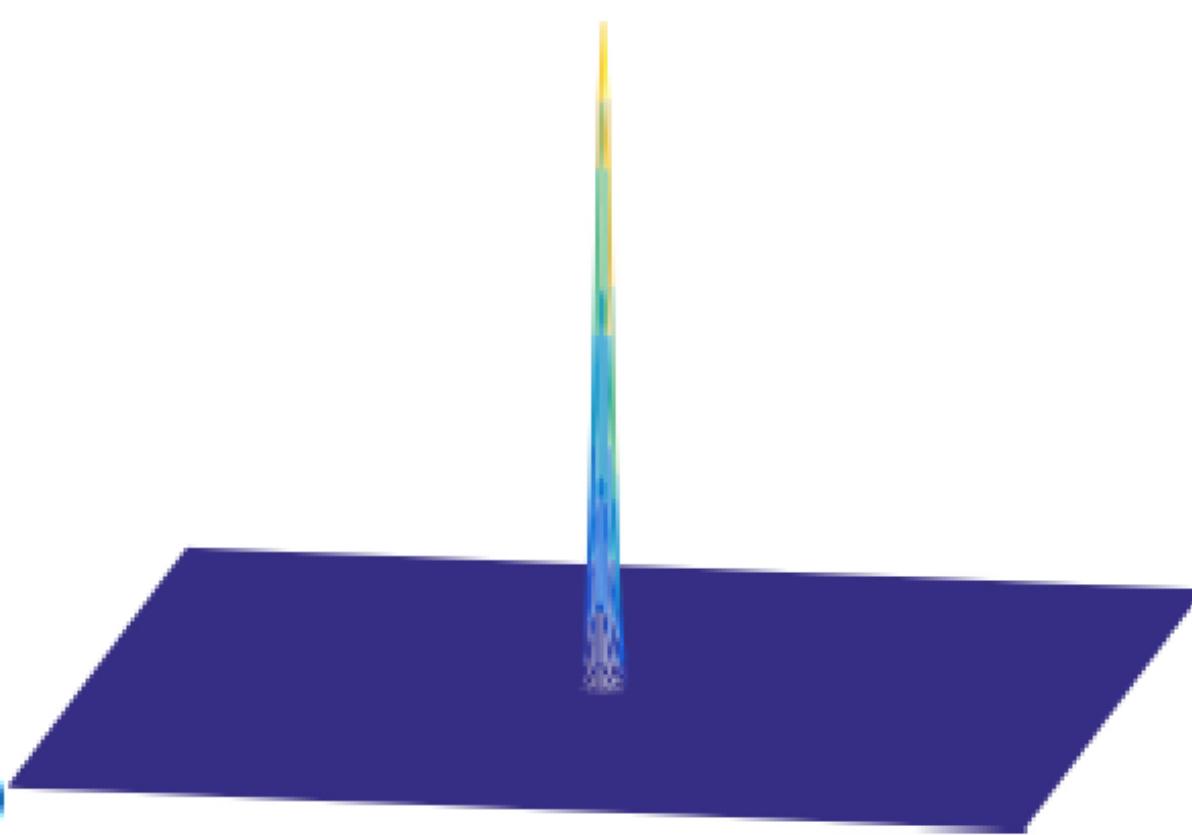
Discretization of the Gaussian:

At 3σ the amplitude of the Gaussian is around 1% of its central value

$$g[m, n; \sigma] = \exp -\frac{m^2 + n^2}{2\sigma^2}$$

Scale

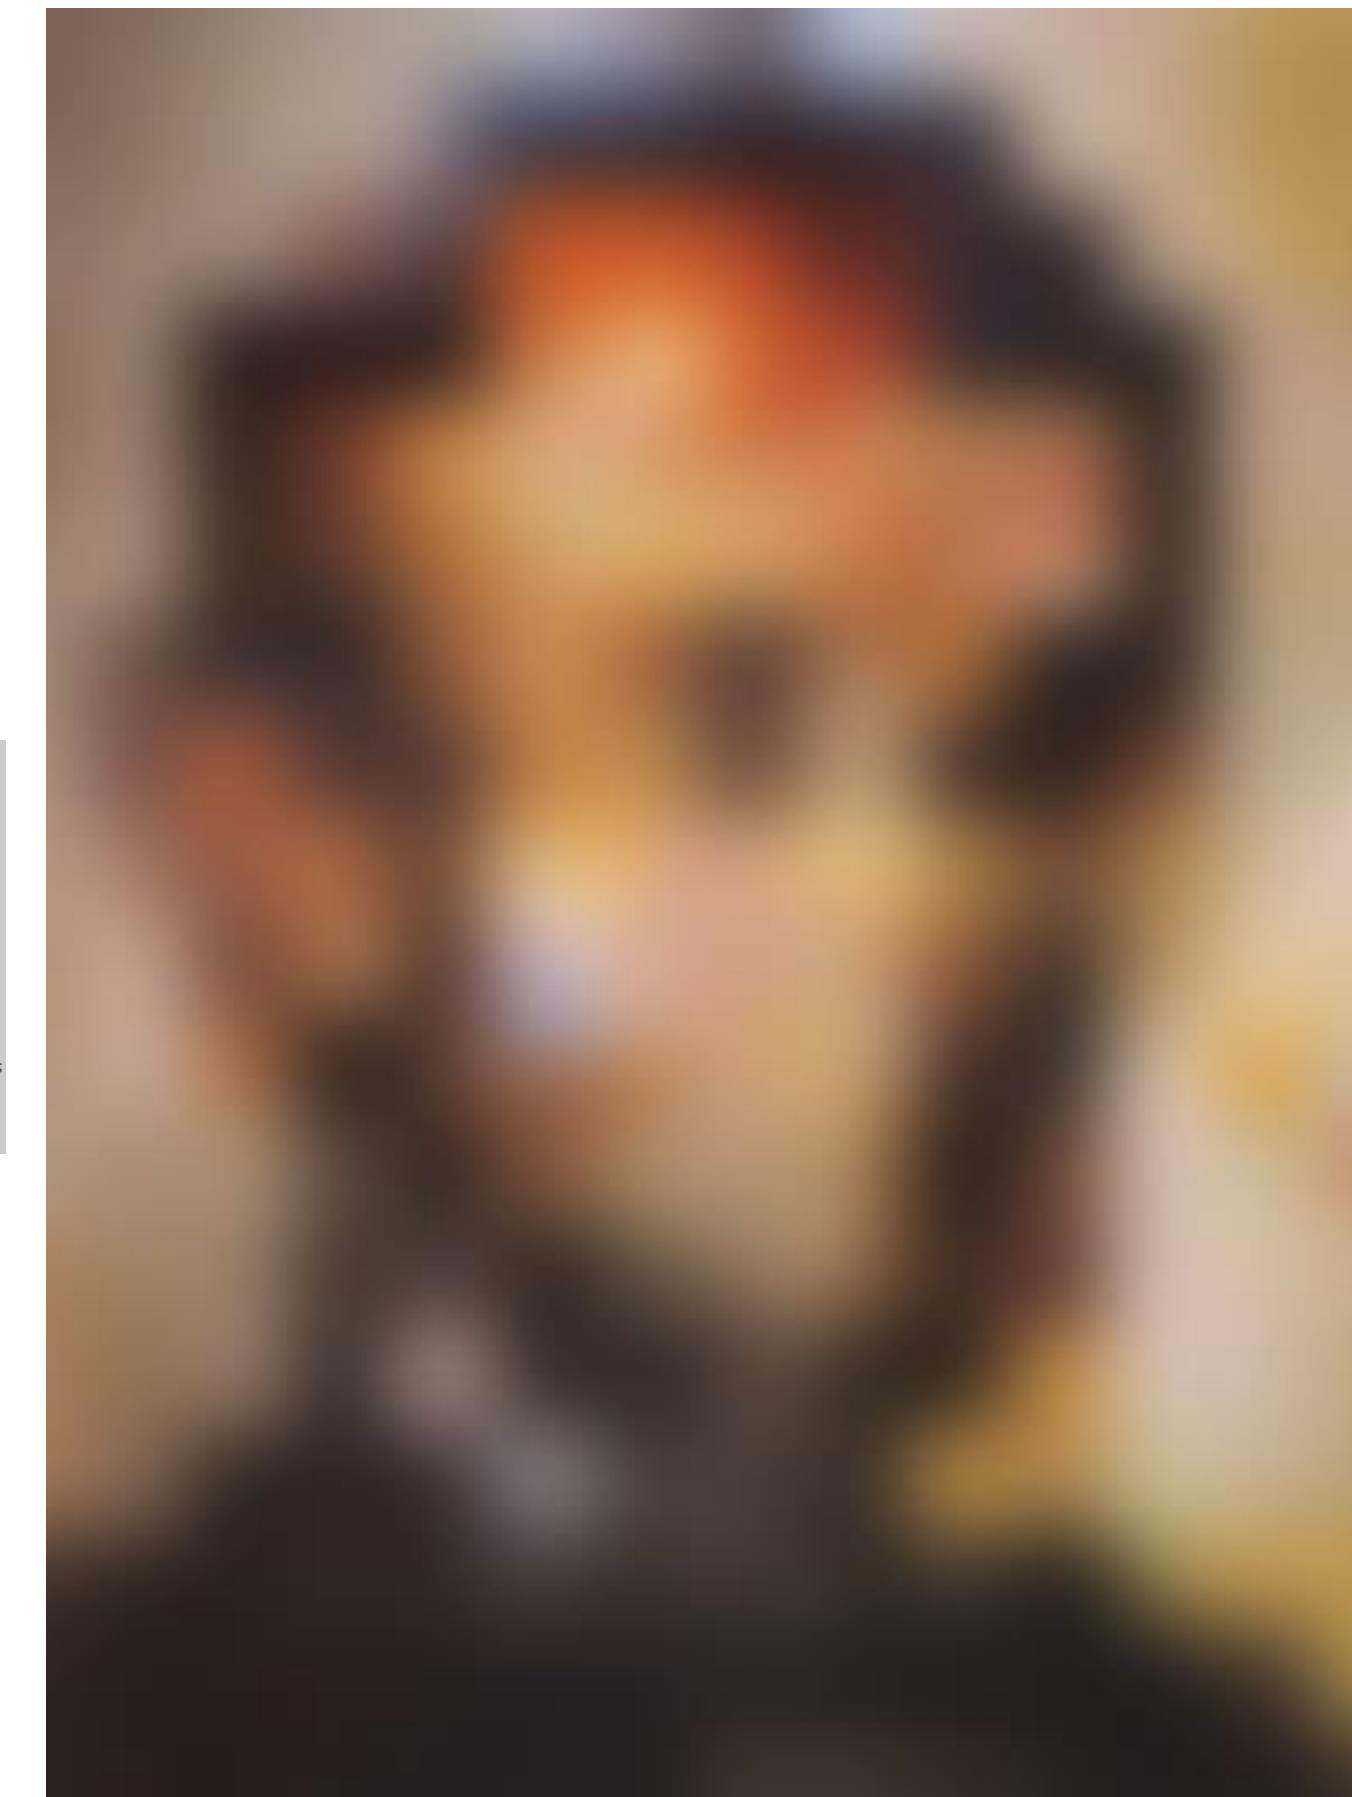
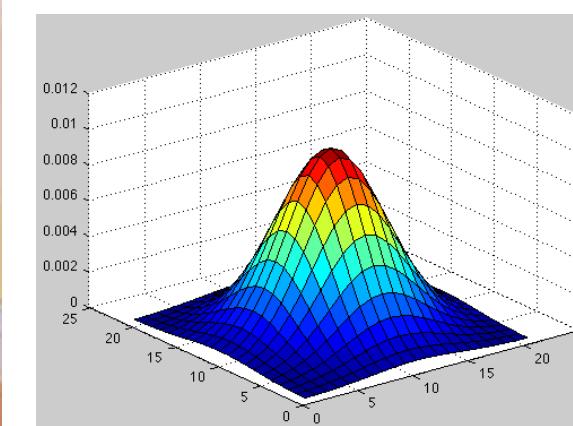
$$g[m, n; \sigma] = \exp -\frac{m^2 + n^2}{2\sigma^2}$$



Gaussian filter for low-pass filtering



Dali



Binomial filter

Binomial coefficients provide a compact approximation of the gaussian coefficients using only integers.

The simplest blur filter (low pass) is

$$[1 \ 1]$$

Binomial filters in the family of filters obtained as successive convolutions of $[1 \ 1]$

Binomial filter

$$\mathbf{b}_1 = [1 \ 1]$$

$$\mathbf{b}_2 = [1 \ 1] \odot [1 \ 1] = [1 \ 2 \ 1]$$

$$\mathbf{b}_3 = [1 \ 1] \odot [1 \ 1] \odot [1 \ 1] = [1 \ 3 \ 3 \ 1]$$

Binomial filter

b_1		1	1							$\sigma_1^2 = 1/4$
b_2		1	2	1						$\sigma_2^2 = 1/2$
b_3		1	3	3	1					$\sigma_3^2 = 3/4$
b_4		1	4	6	4	1				$\sigma_4^2 = 1$
b_5		1	5	10	10	5	1			$\sigma_5^2 = 5/4$
b_6		1	6	15	20	15	6	1		$\sigma_6^2 = 3/2$
b_7		1	7	21	35	35	21	7	1	$\sigma_7^2 = 7/4$
b_8	1	8	28	56	70	56	28	8	1	$\sigma_8^2 = 2$

Properties of binomial filters

- Sum of the values is 2^n
- The variance of b_n is $\sigma^2 = n/4$
- The convolution of two binomial filters is also a binomial filter

$$b_n \circ b_m = b_{n+m}$$

With a variance:

$$\sigma_n^2 + \sigma_m^2 = \sigma_{n+m}^2$$

These properties are analogous to the gaussian property in the continuous domain (but the binomial filter is different than a discretization of a gaussian)

B2[n]

$$b_{2,2} = b_{2,0} \circ b_{0,2} = \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} \circ \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

DERIVATIVES

DERIVATIVES

High pass-filters

$$[-1 \ 1]$$

$$\frac{\partial \mathbf{I}}{\partial x} \simeq \mathbf{I}(x, y) - \mathbf{I}(x - 1, y)$$



$g[m,n]$

$$\otimes \quad [-1, 1] =$$

$h[m,n]$



$f[m,n]$

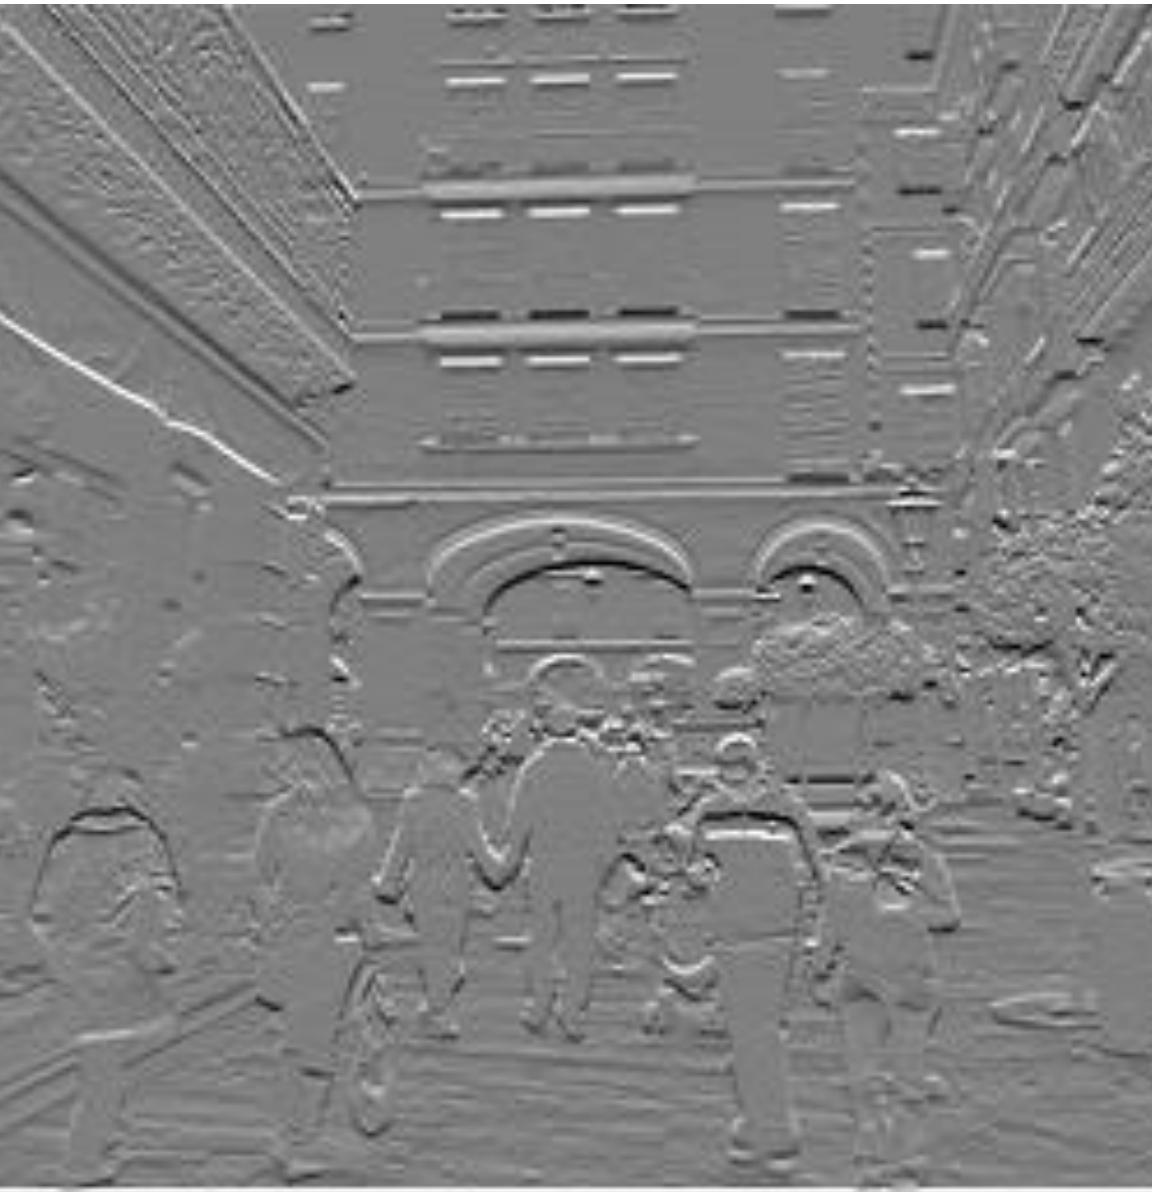
$$[-1 \ 1]^T$$



$g[m,n]$

$$\otimes \quad [-1, 1]^T =$$

$h[m,n]$



$f[m,n]$

Discrete derivatives

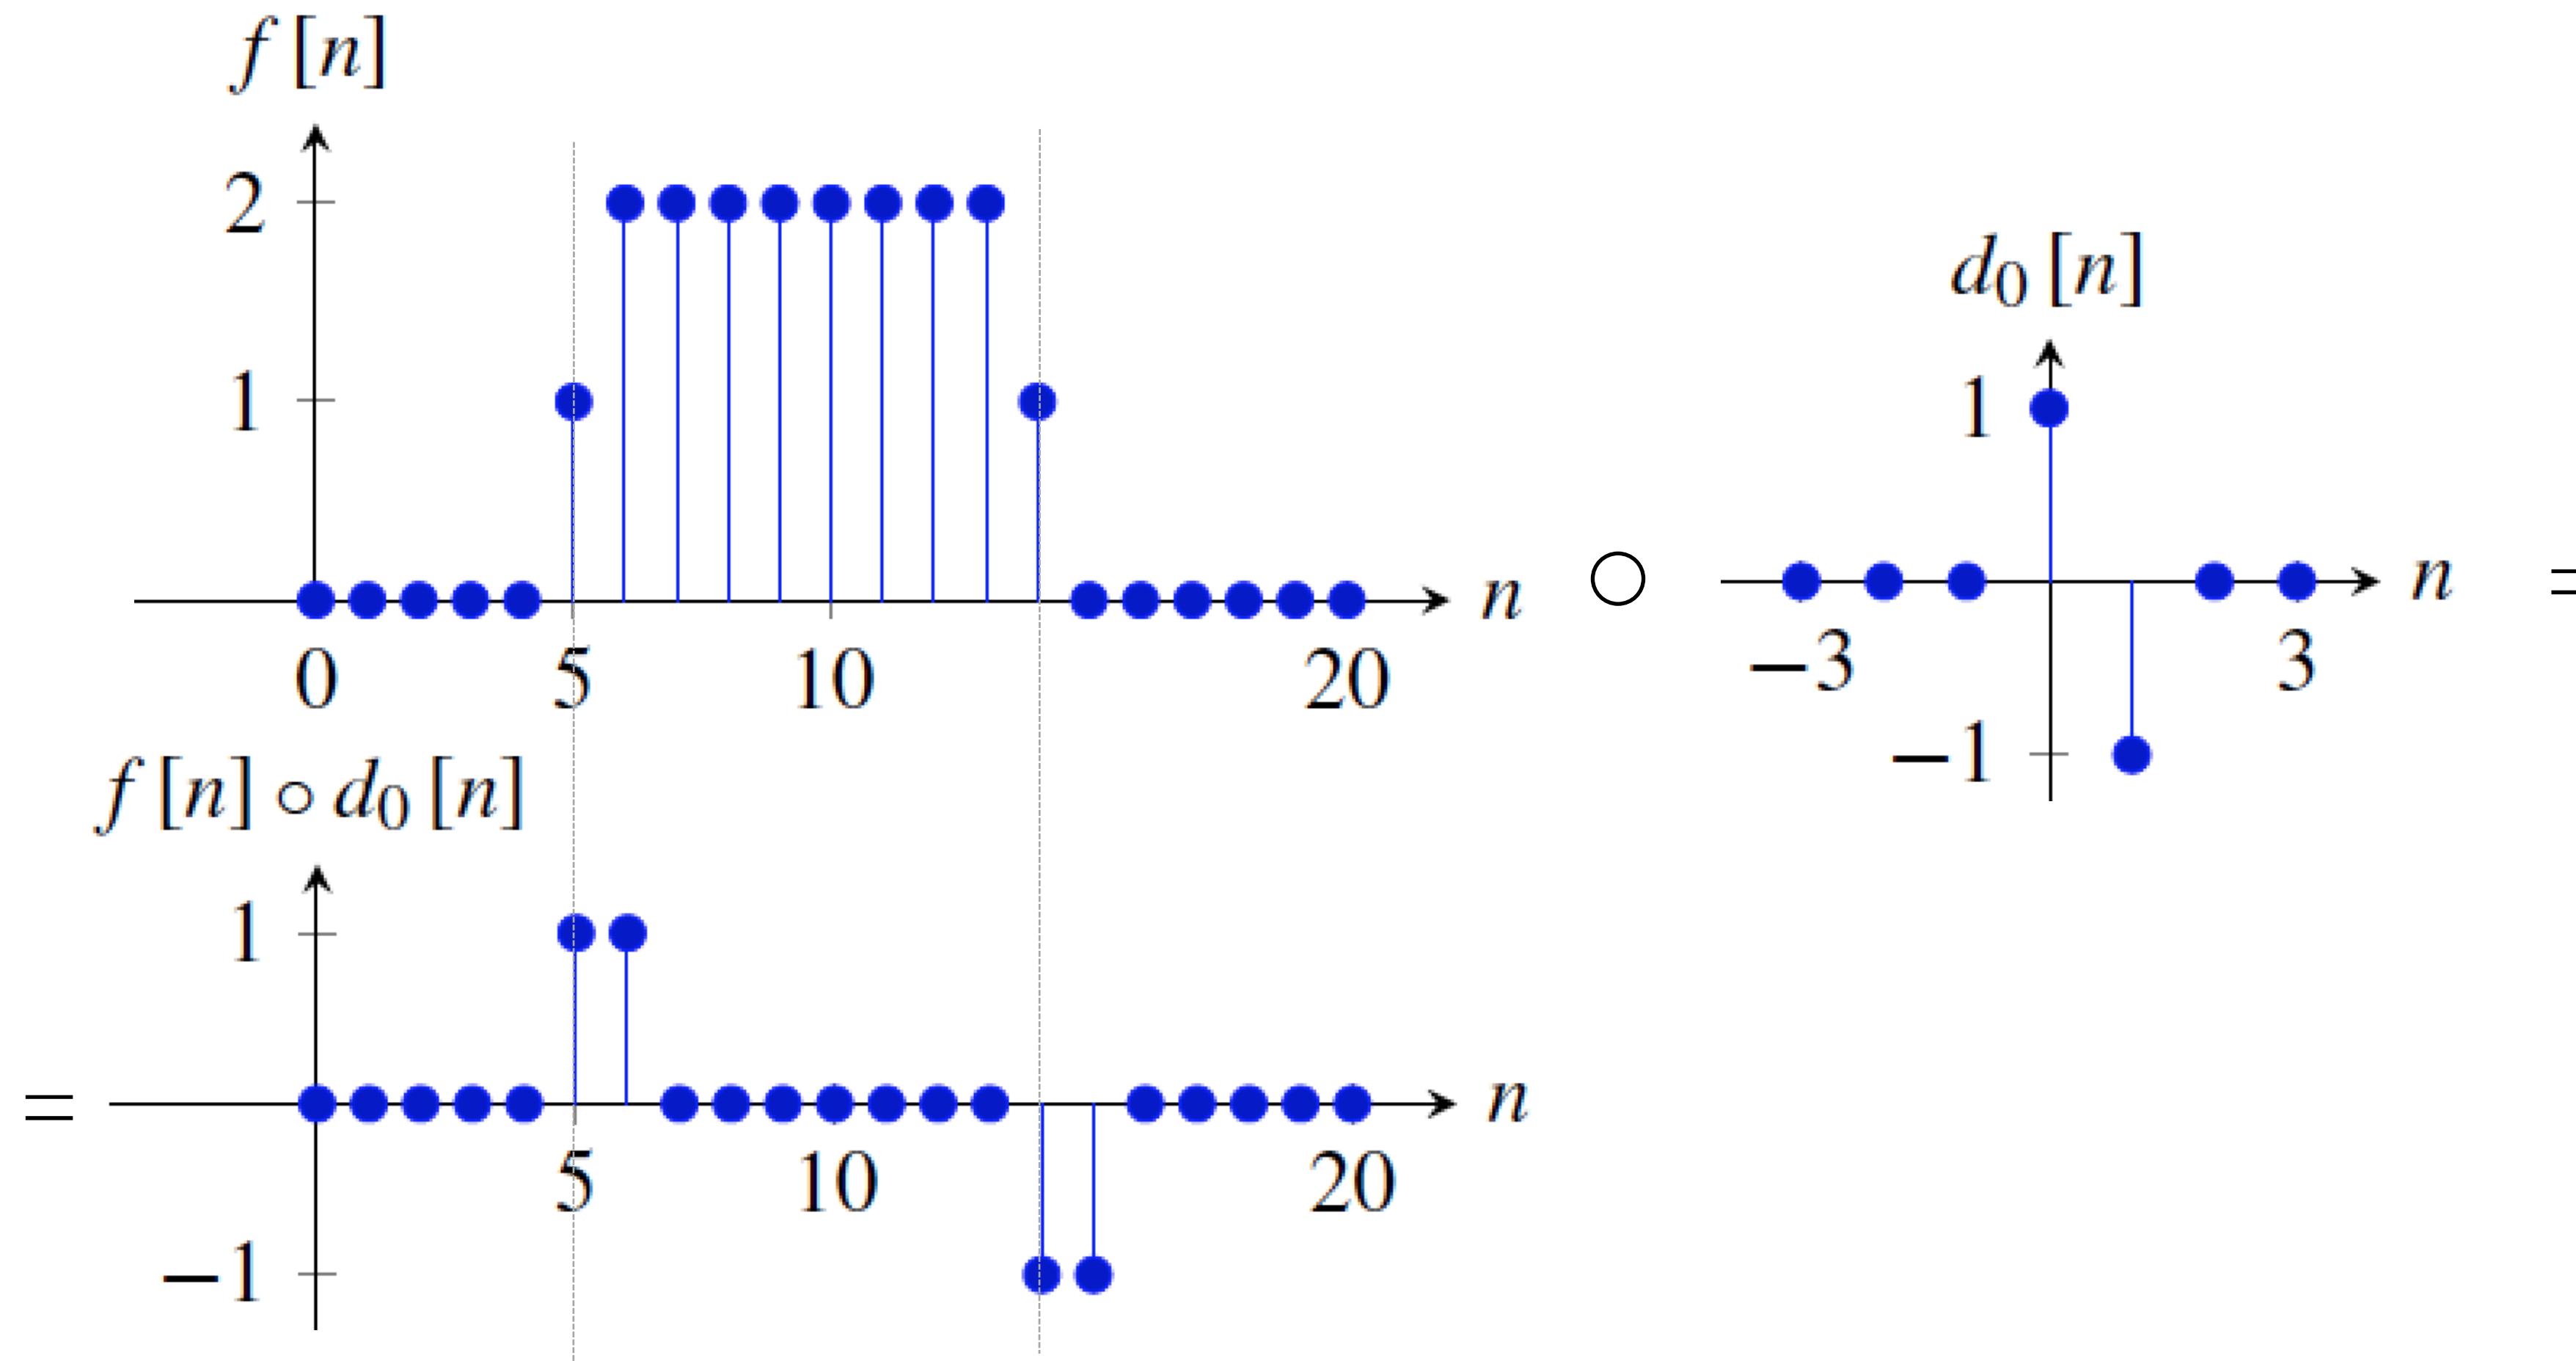
$$d_0 = [1, -1]$$

$$f \circ d_0 = f[n] - f[n-1]$$

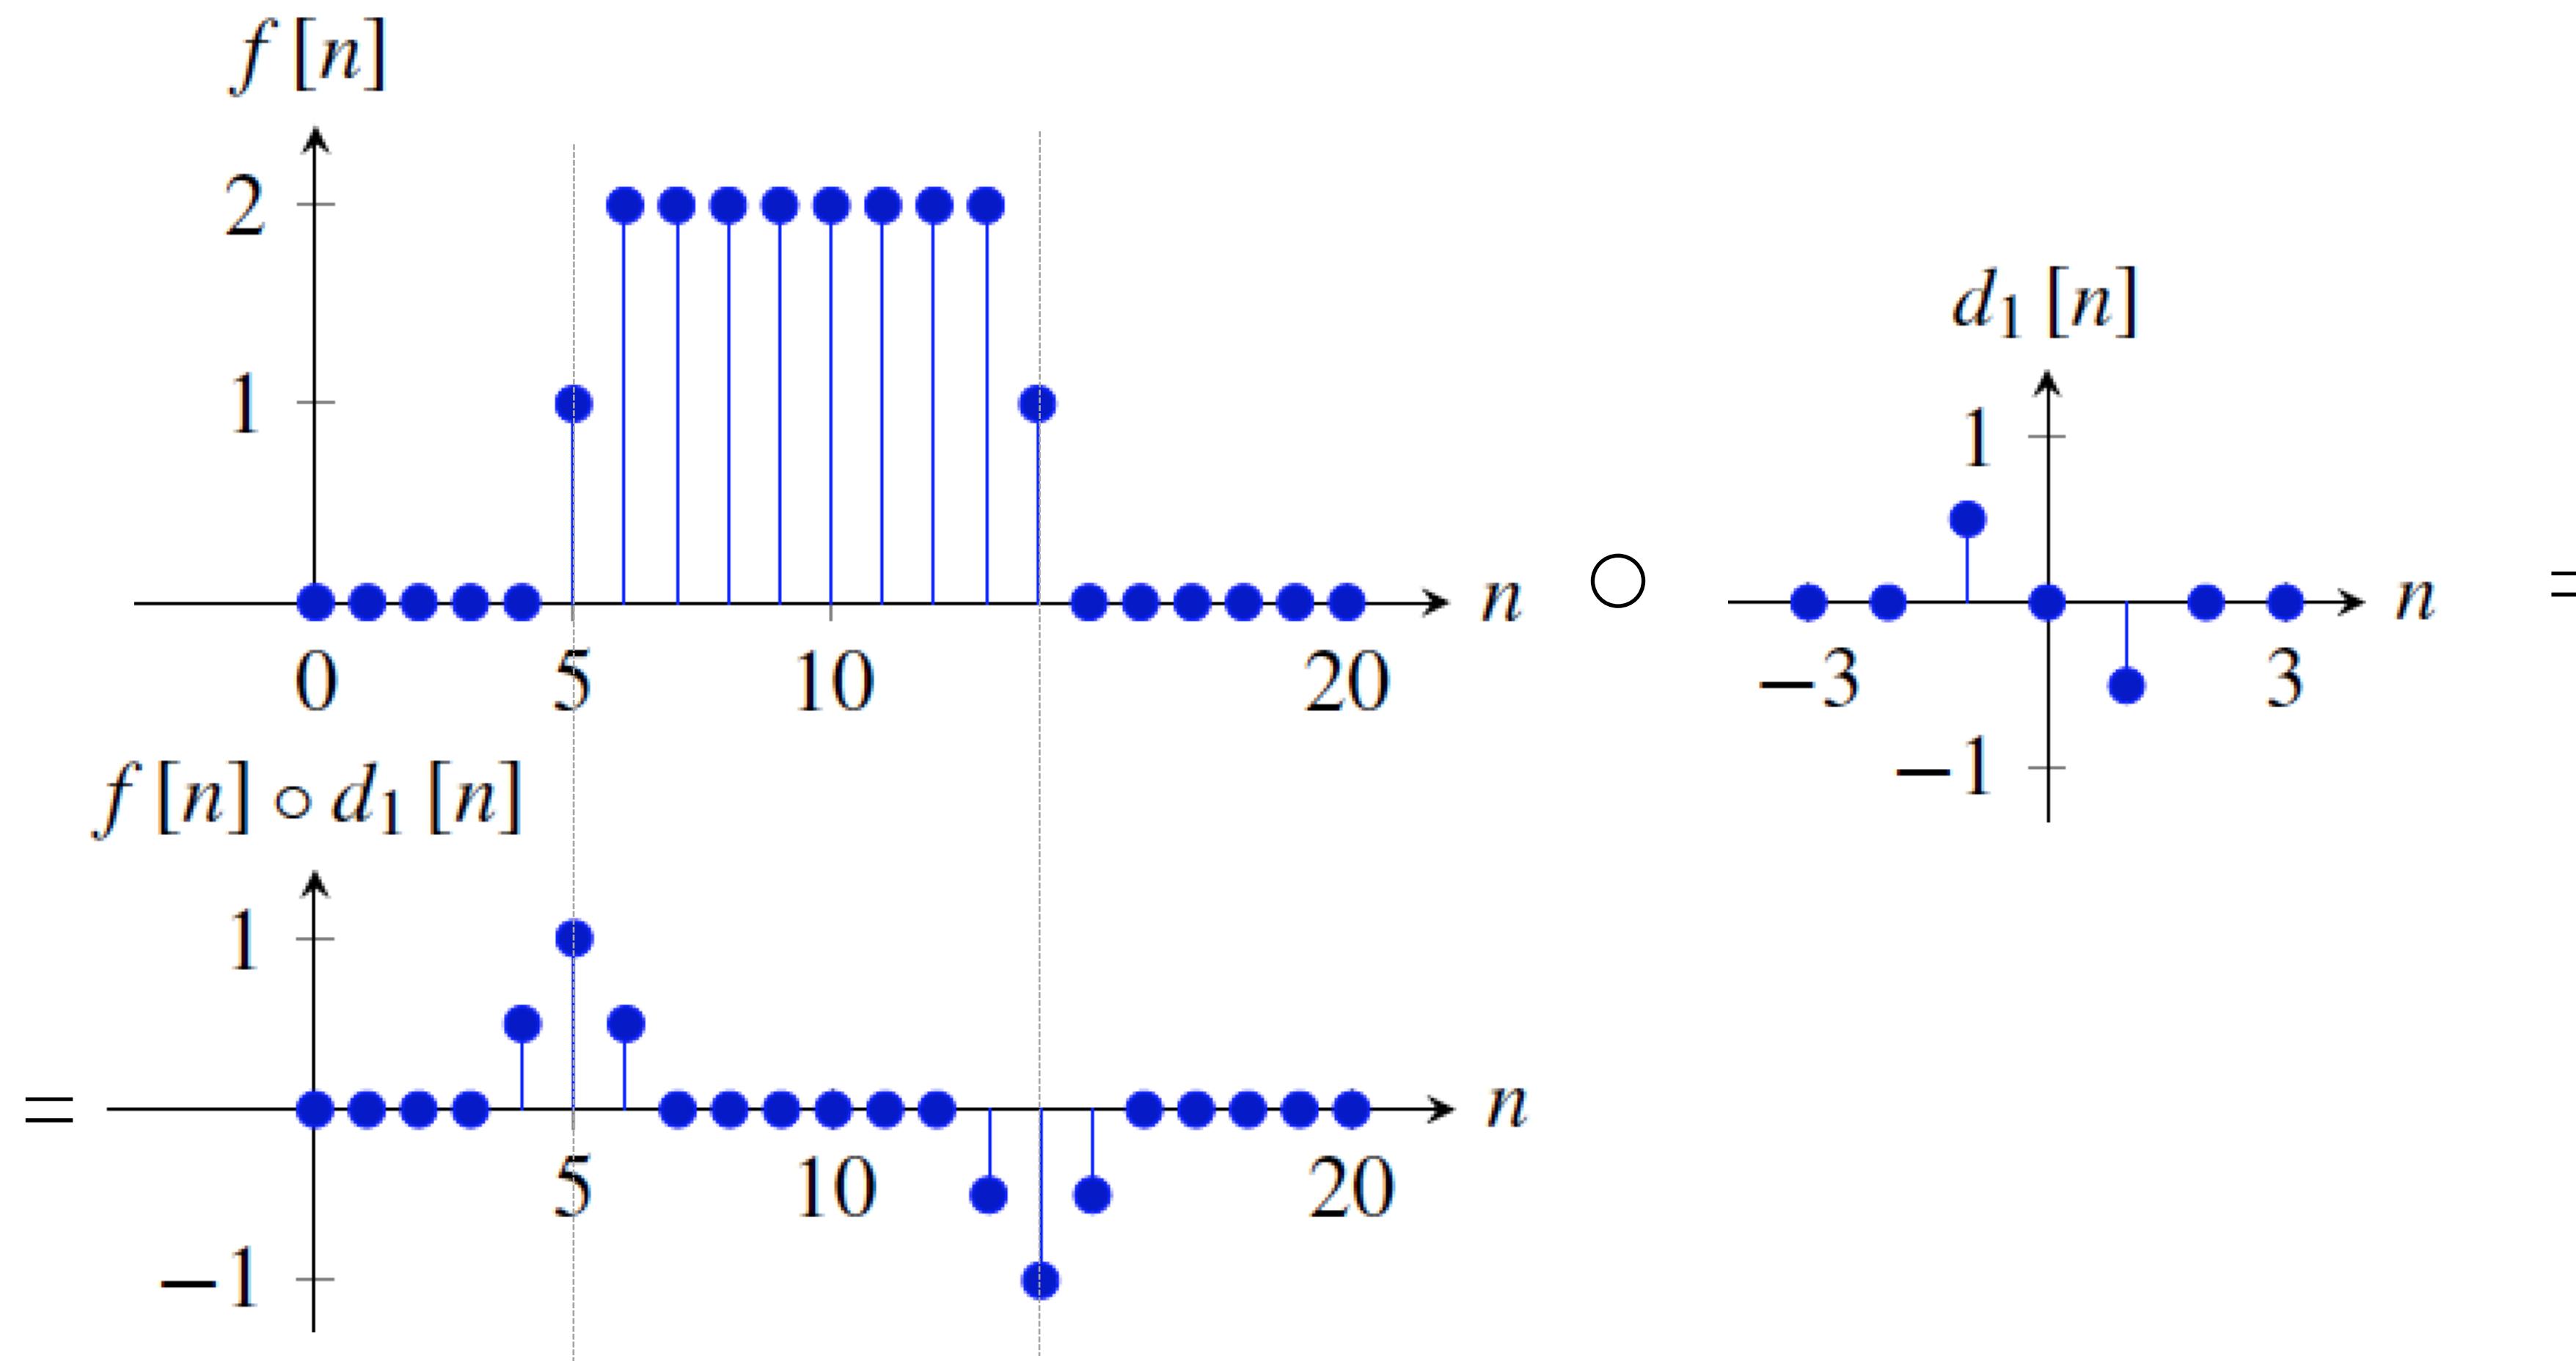
$$d_1 = [1, 0, -1]/2$$

$$f \circ d_1 = \frac{f[n+1] - f[n-1]}{2}$$

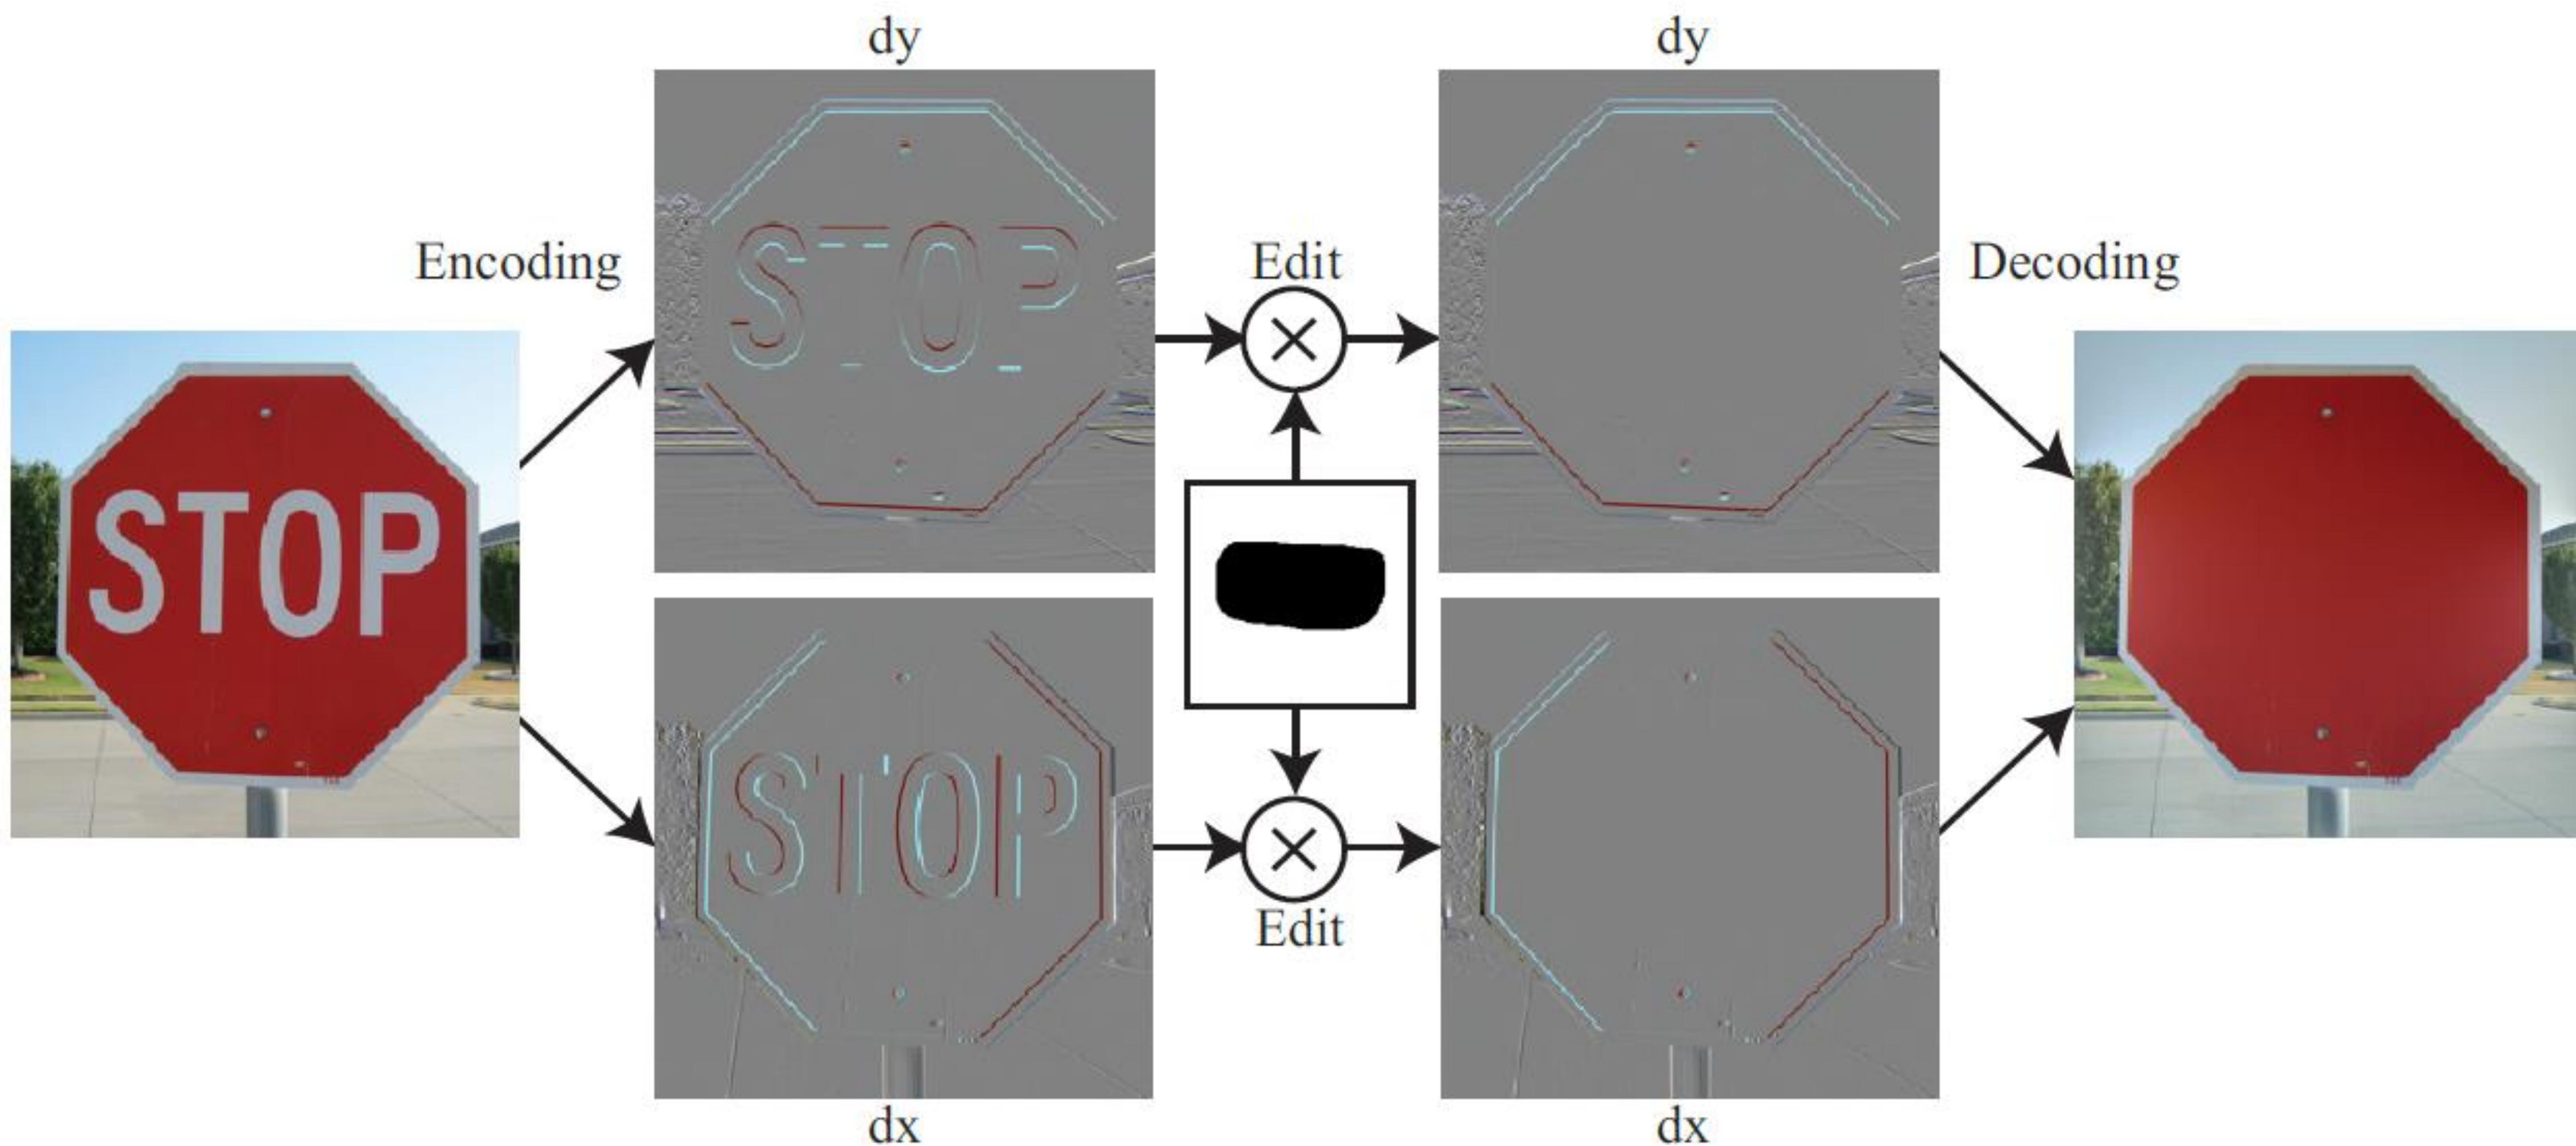
Discrete derivatives



Discrete derivatives



Gradient Domain



Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images

Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features

Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features



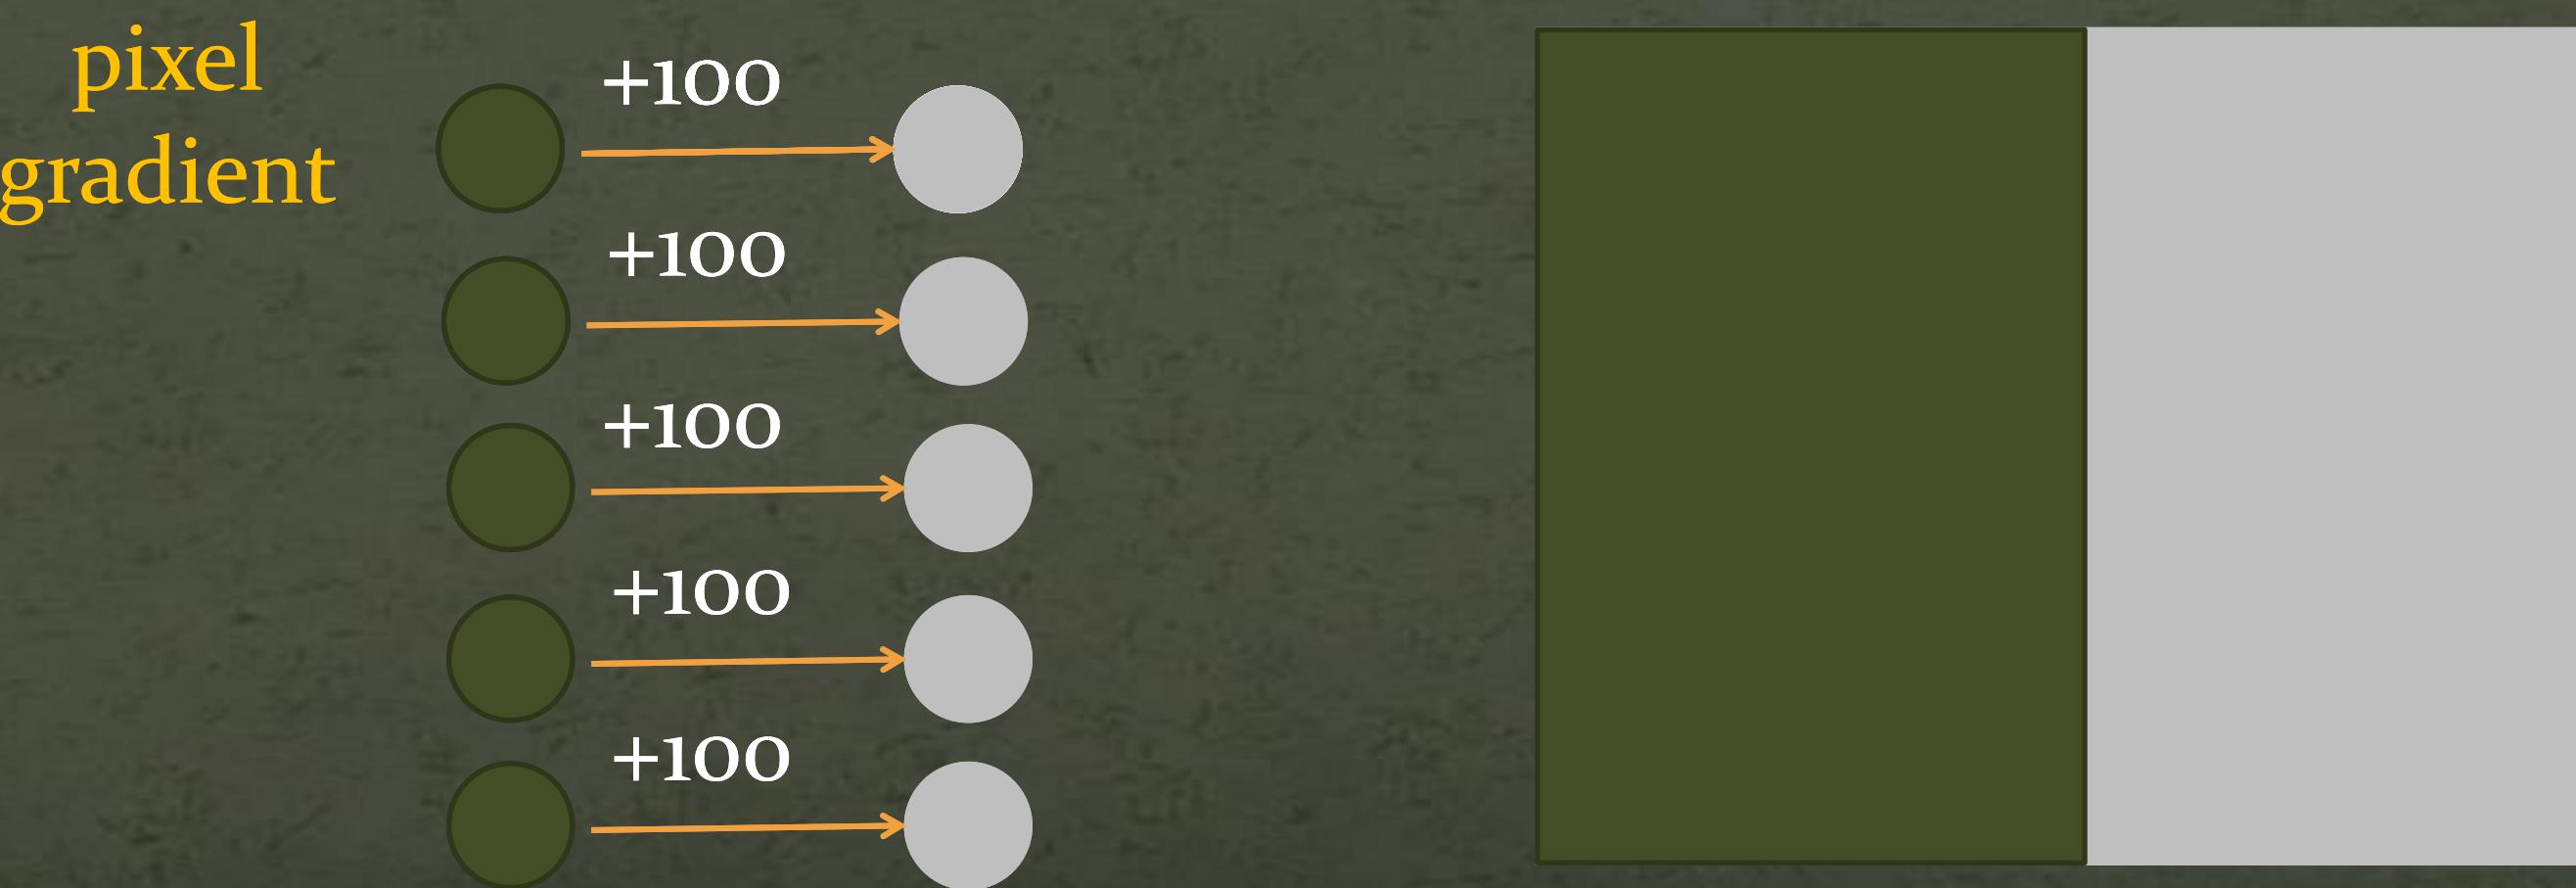
Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features
 - gradients – give rise to high level image-features



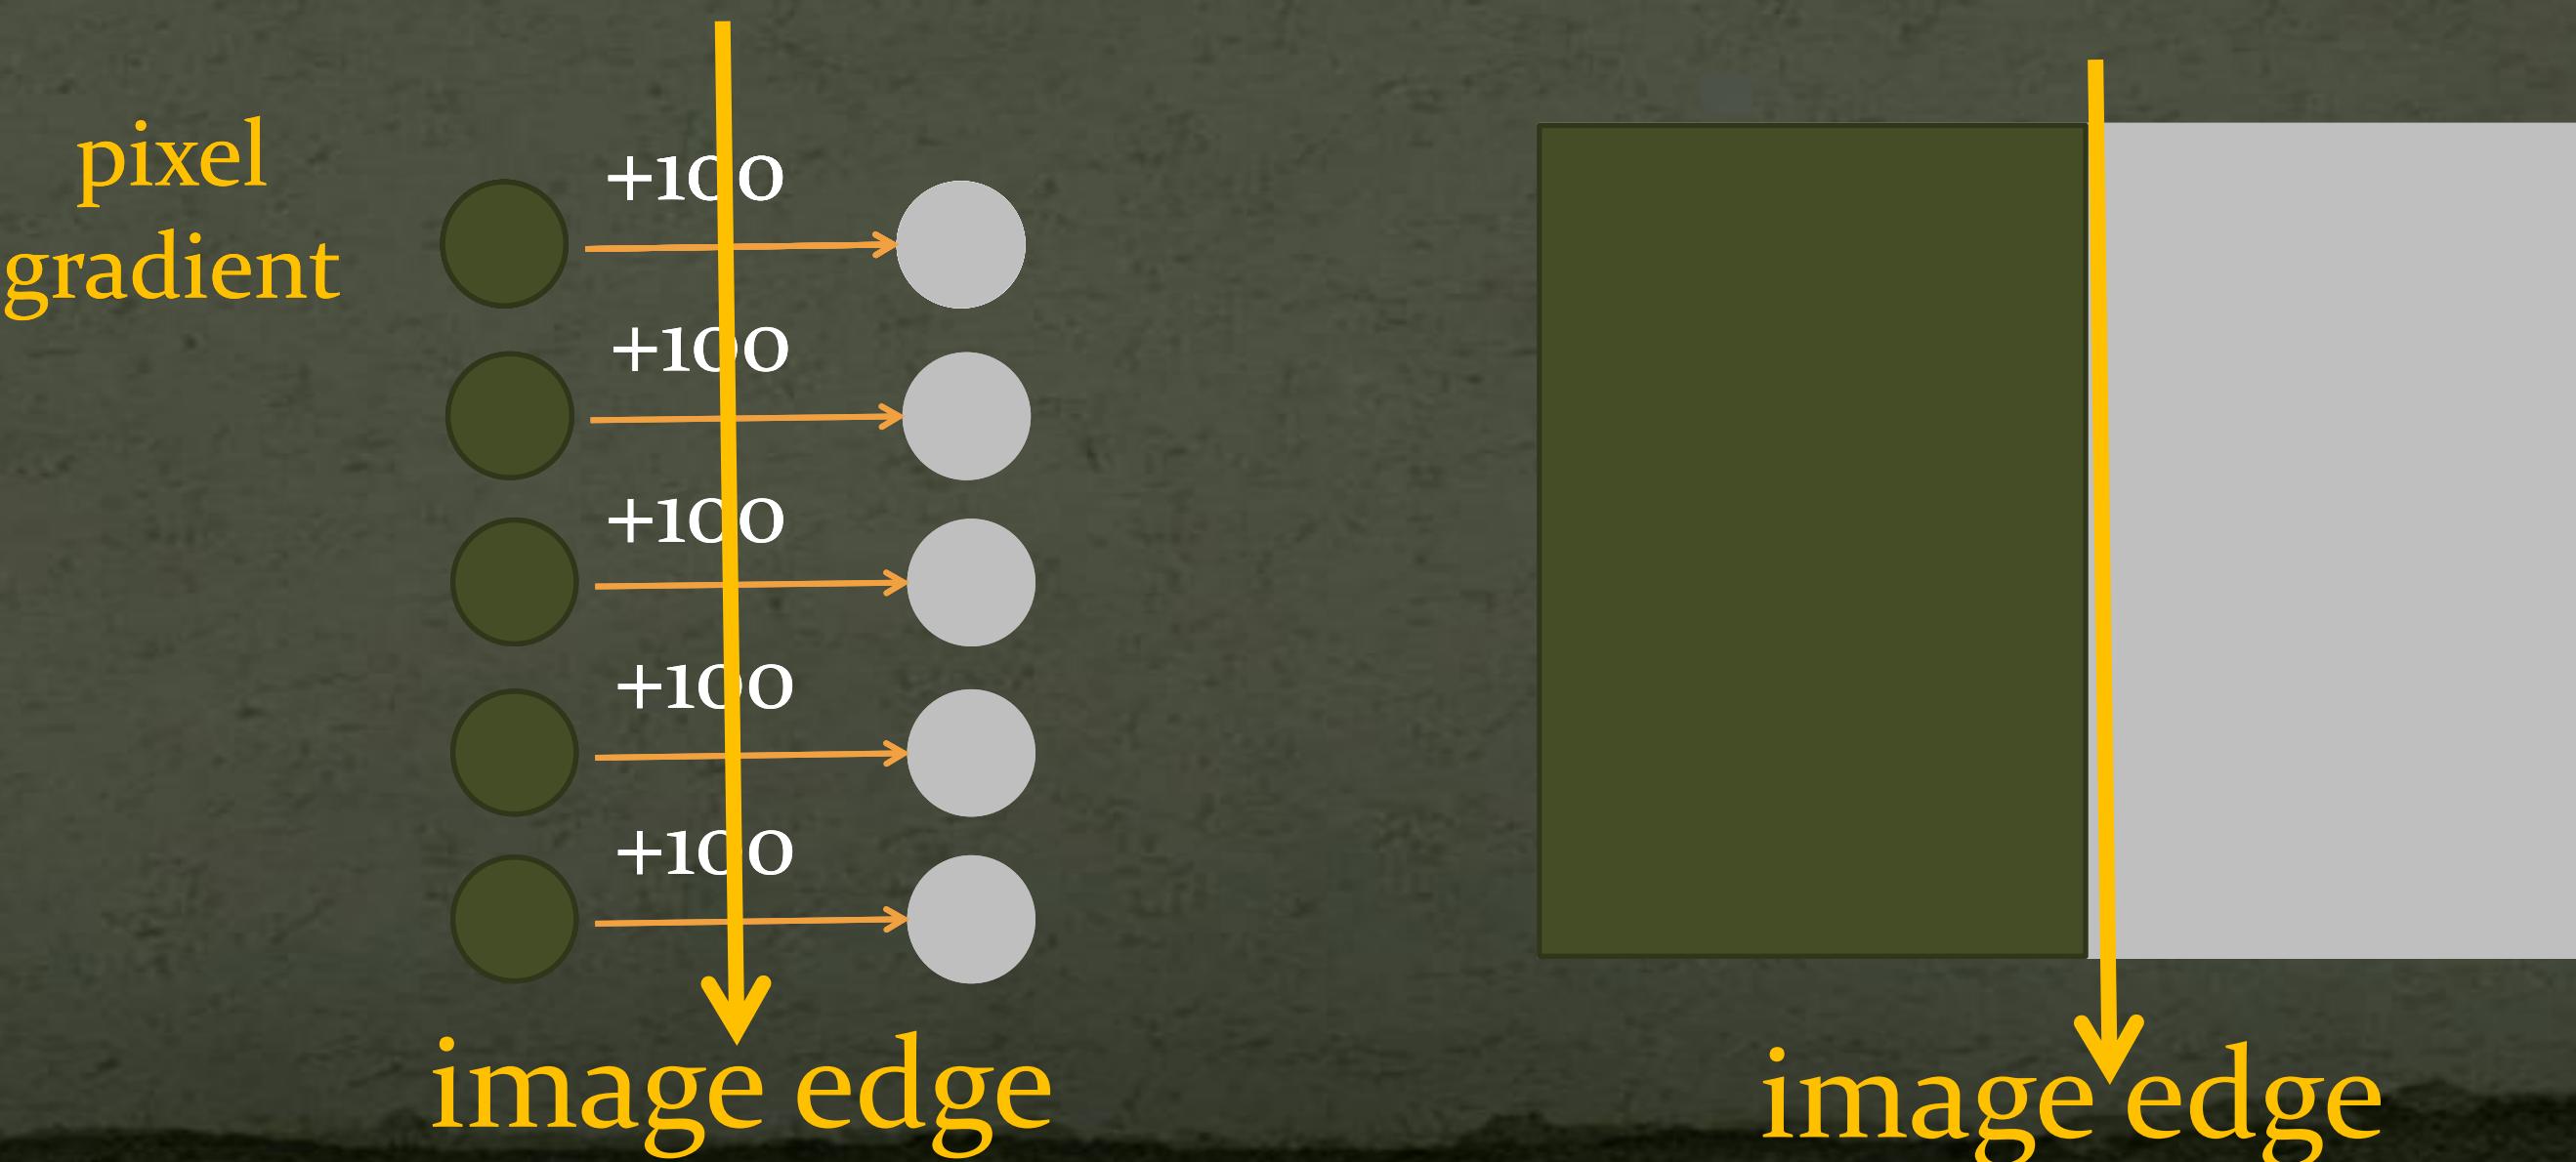
Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features
 - gradients – give rise to high level image-features



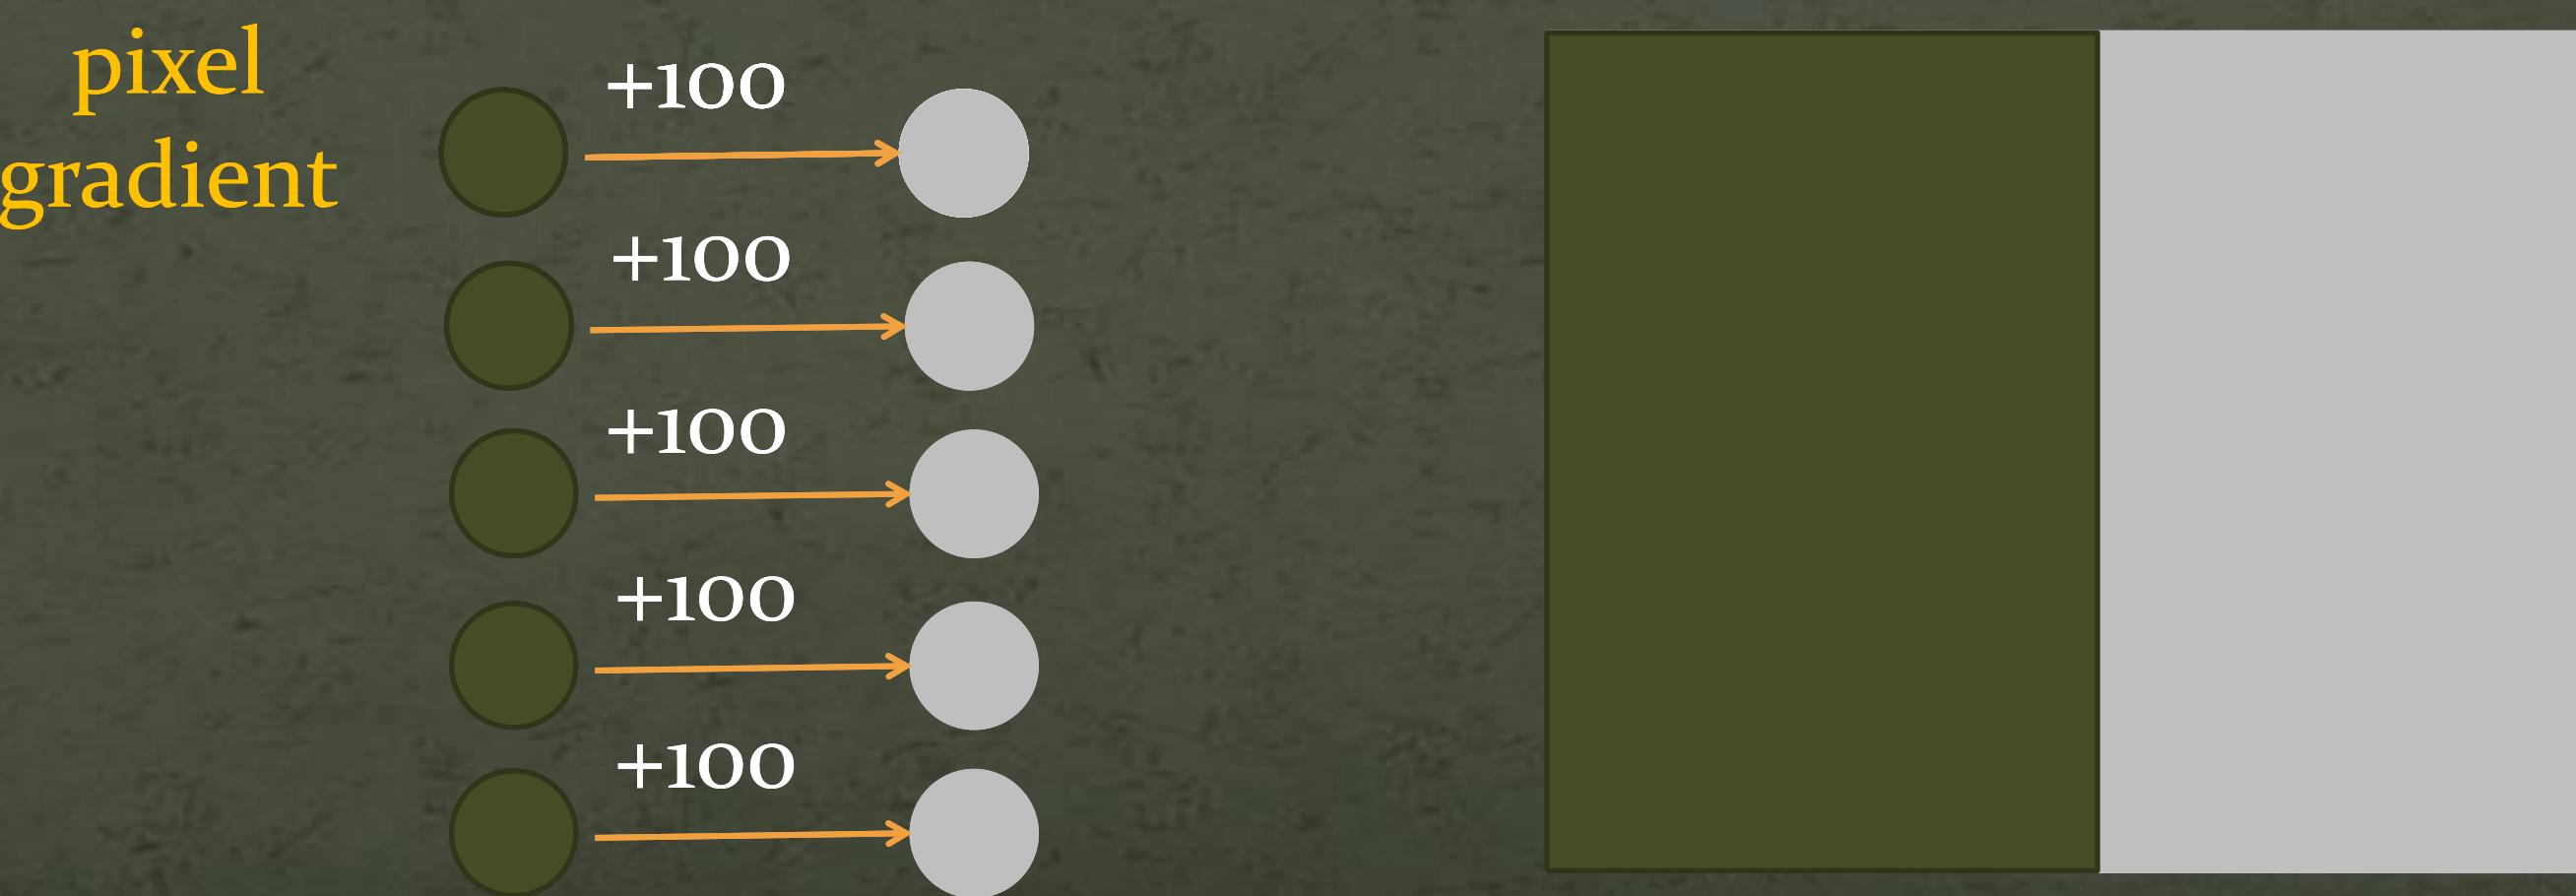
Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features
 - gradients – give rise to high level image-features



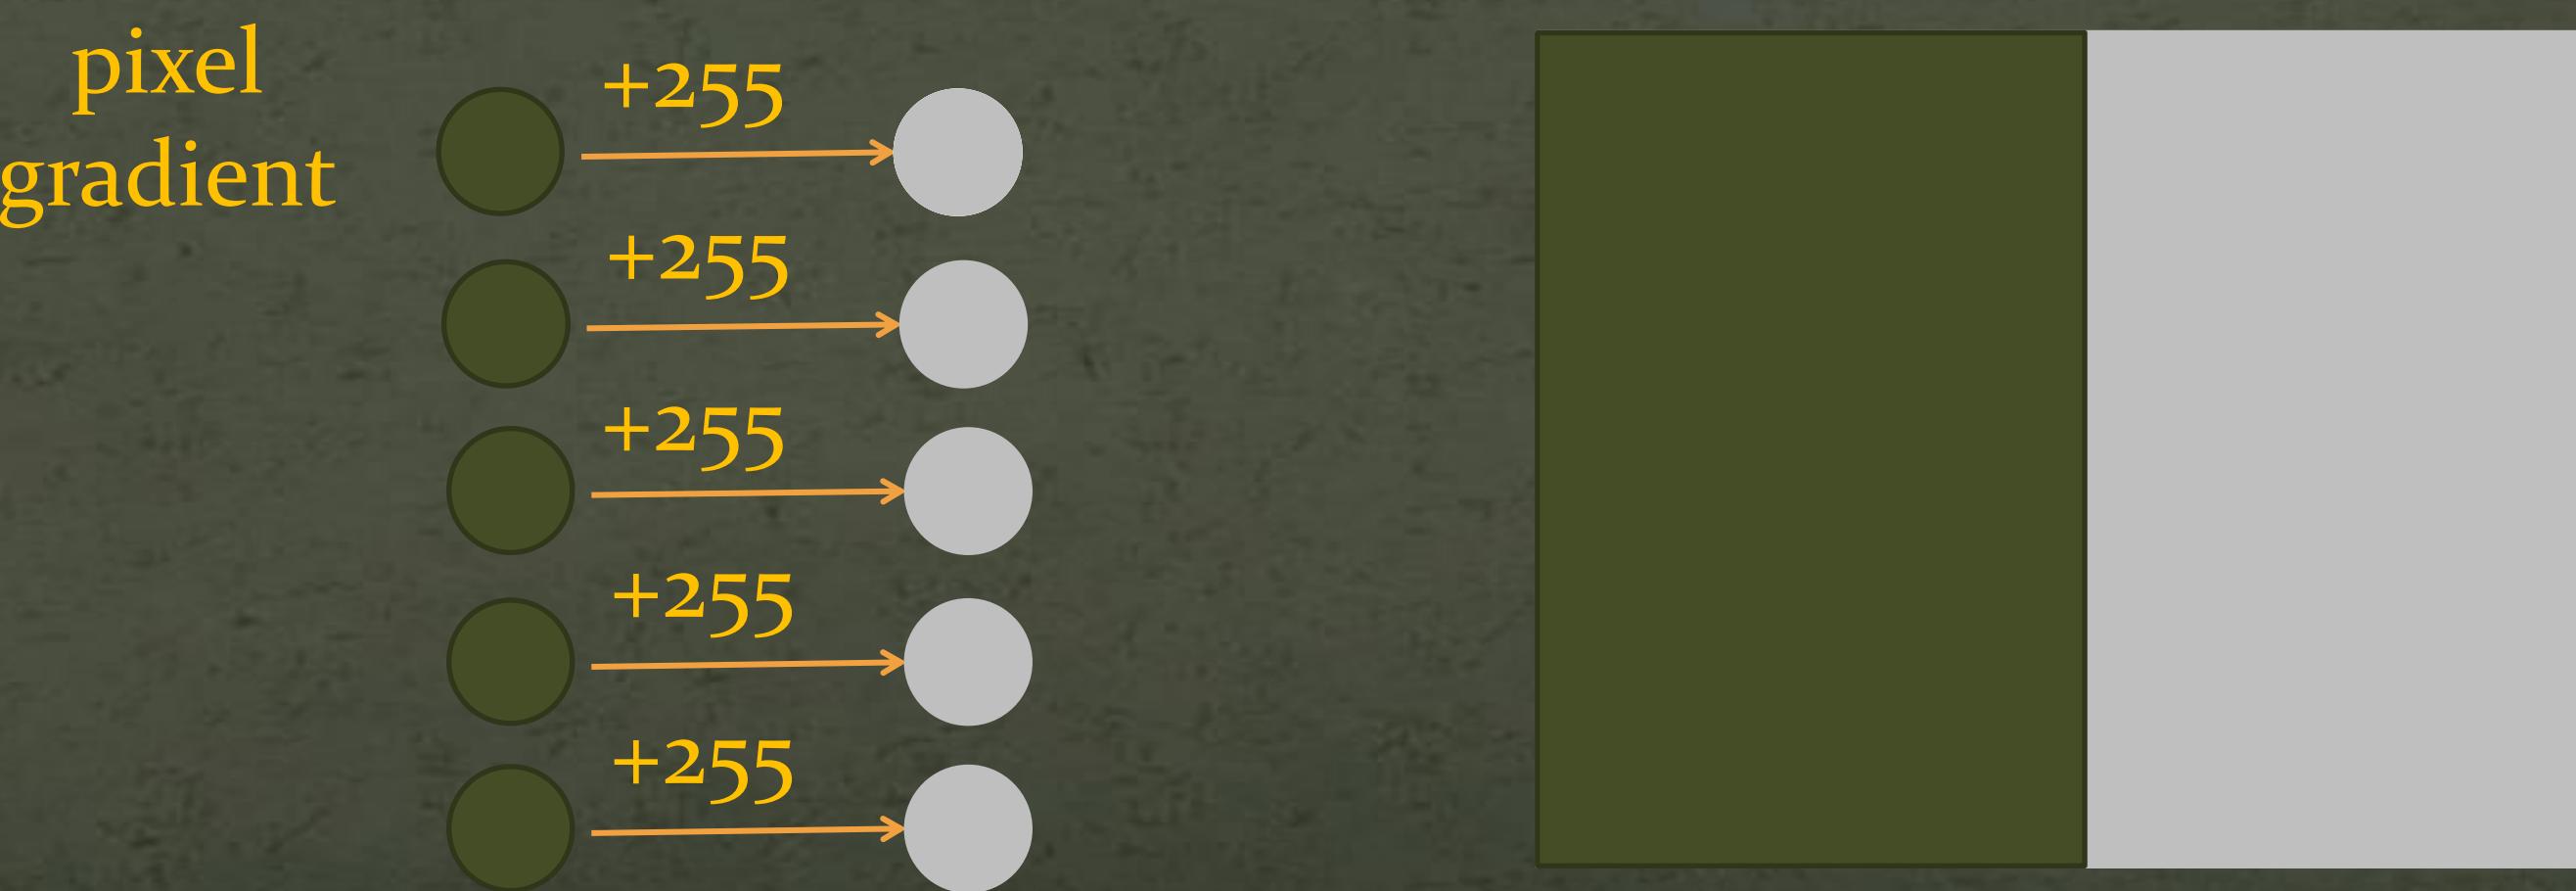
Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features
 - gradients – give rise to high level image-features
 - manipulate local gradients to manipulate global image interpretation



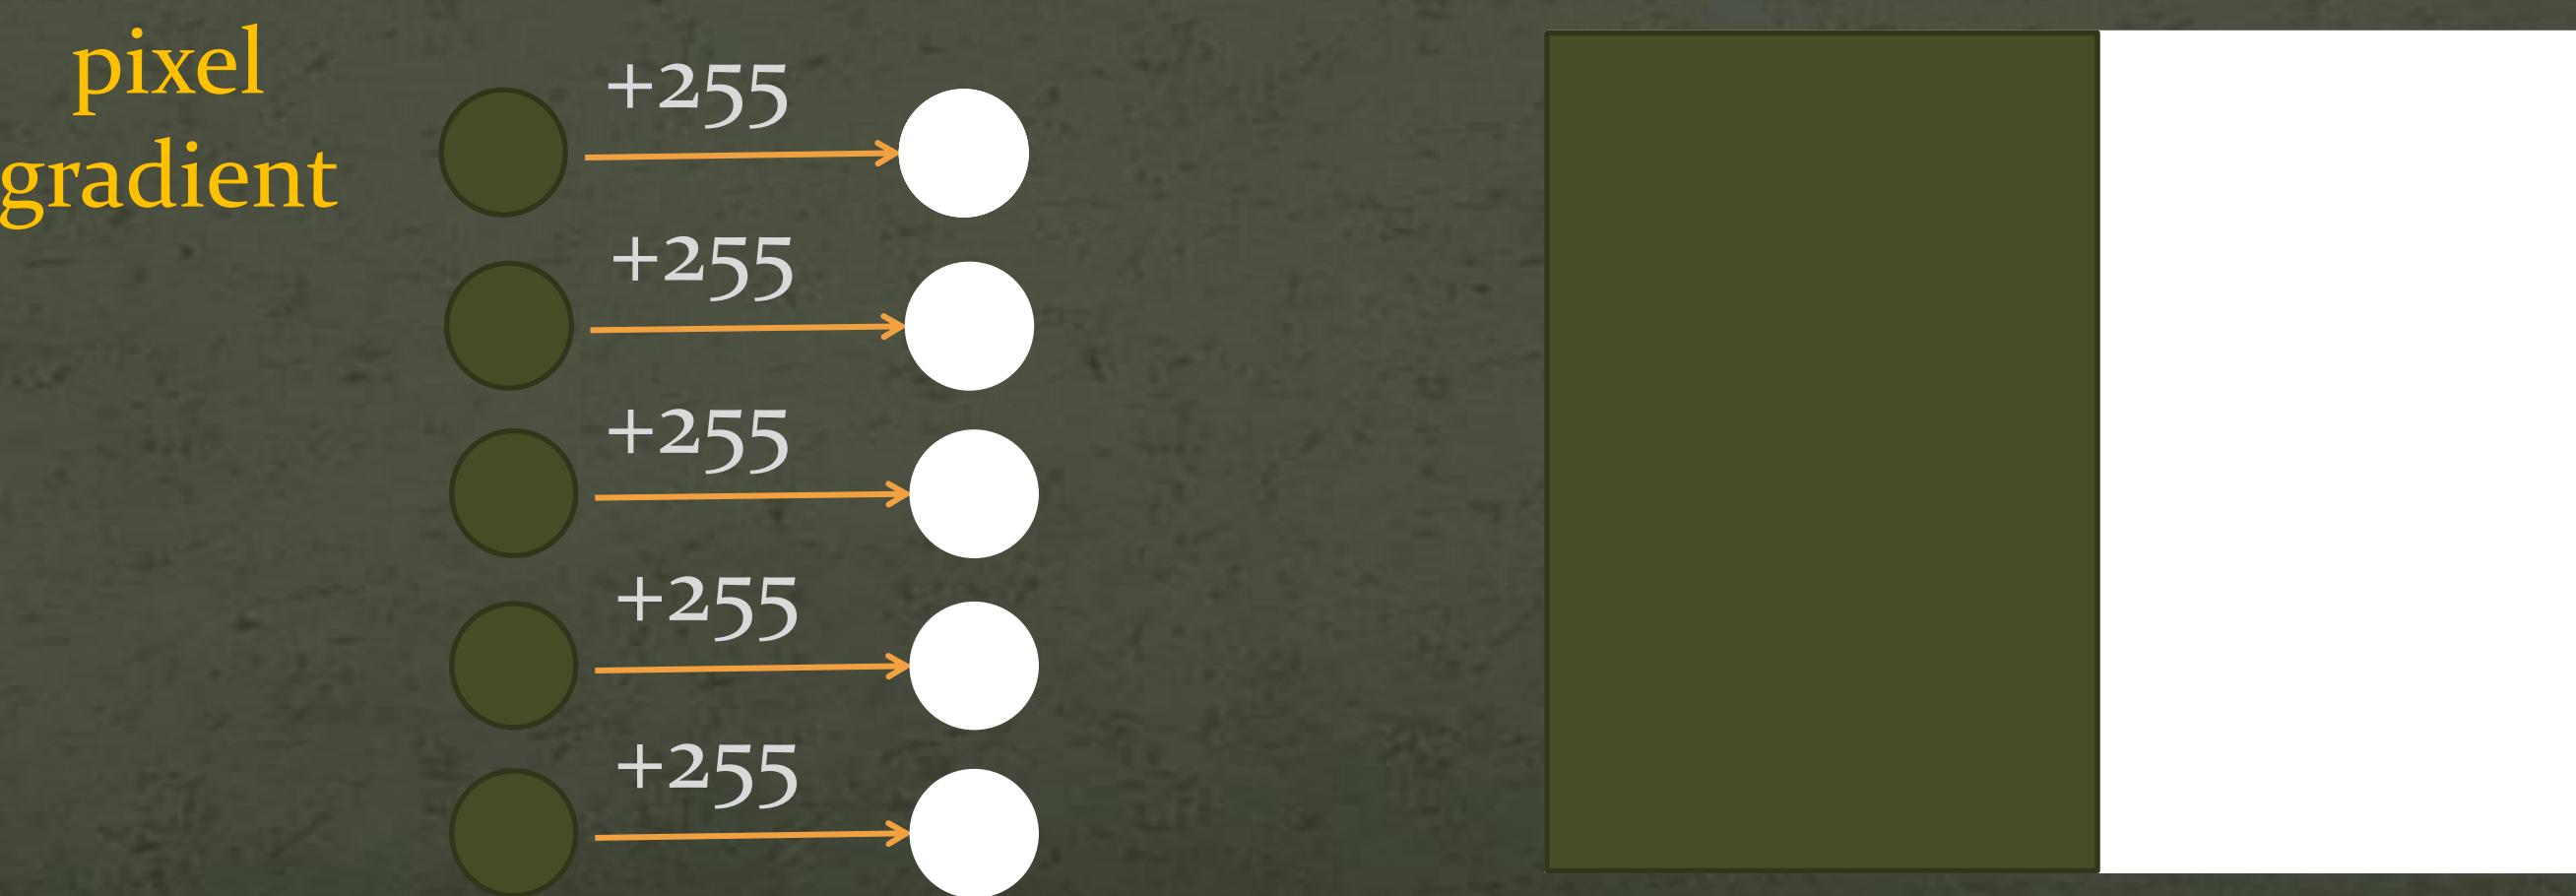
Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features
 - gradients – give rise to high level image-features
 - manipulate local gradients to manipulate global image interpretation



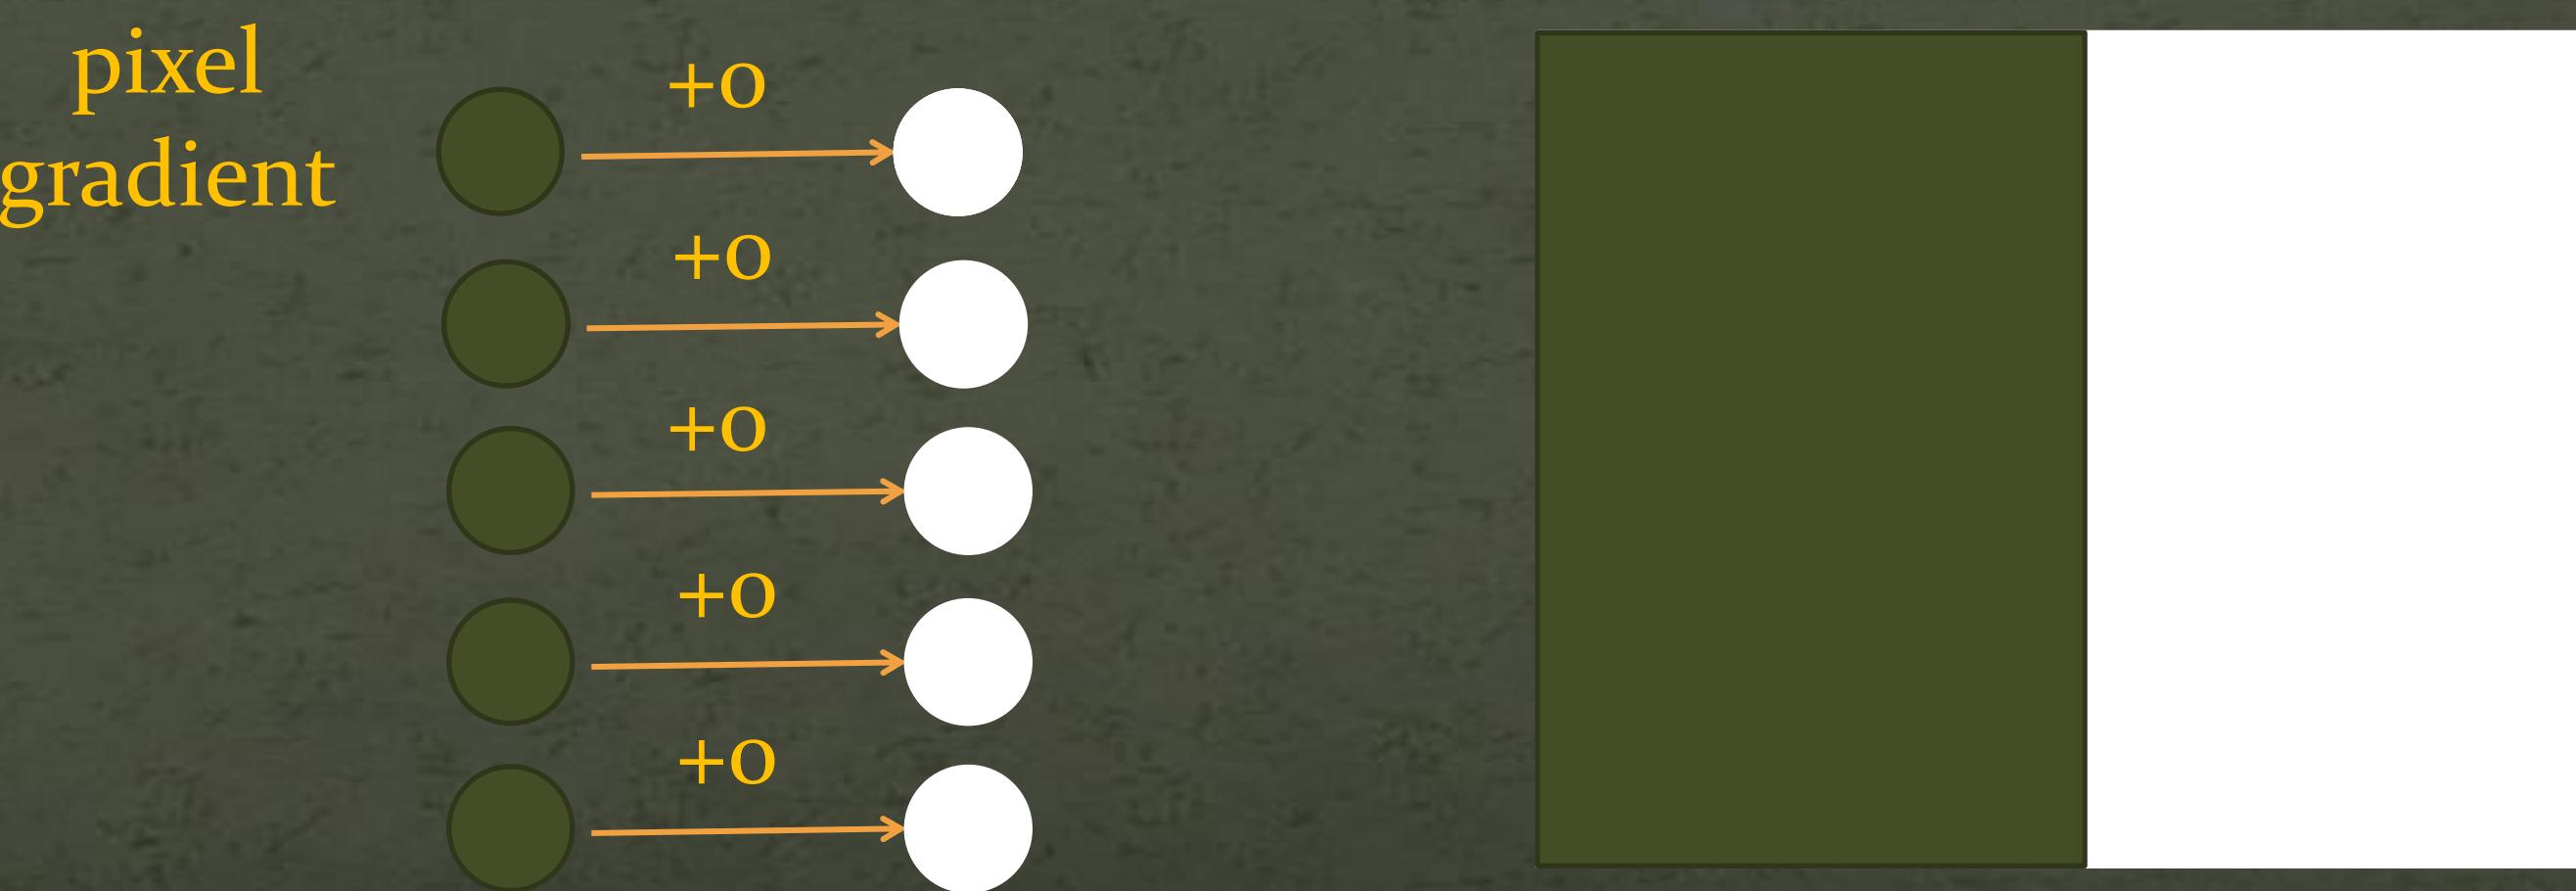
Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features
 - gradients – give rise to high level image-features
 - manipulate local gradients to manipulate global image interpretation



Motivation for gradient-domain filtering?

- Can be used to exert high-level control over images
 - gradients – low level image-features
 - gradients – give rise to high level image-features
 - manipulate local gradients to manipulate global image interpretation

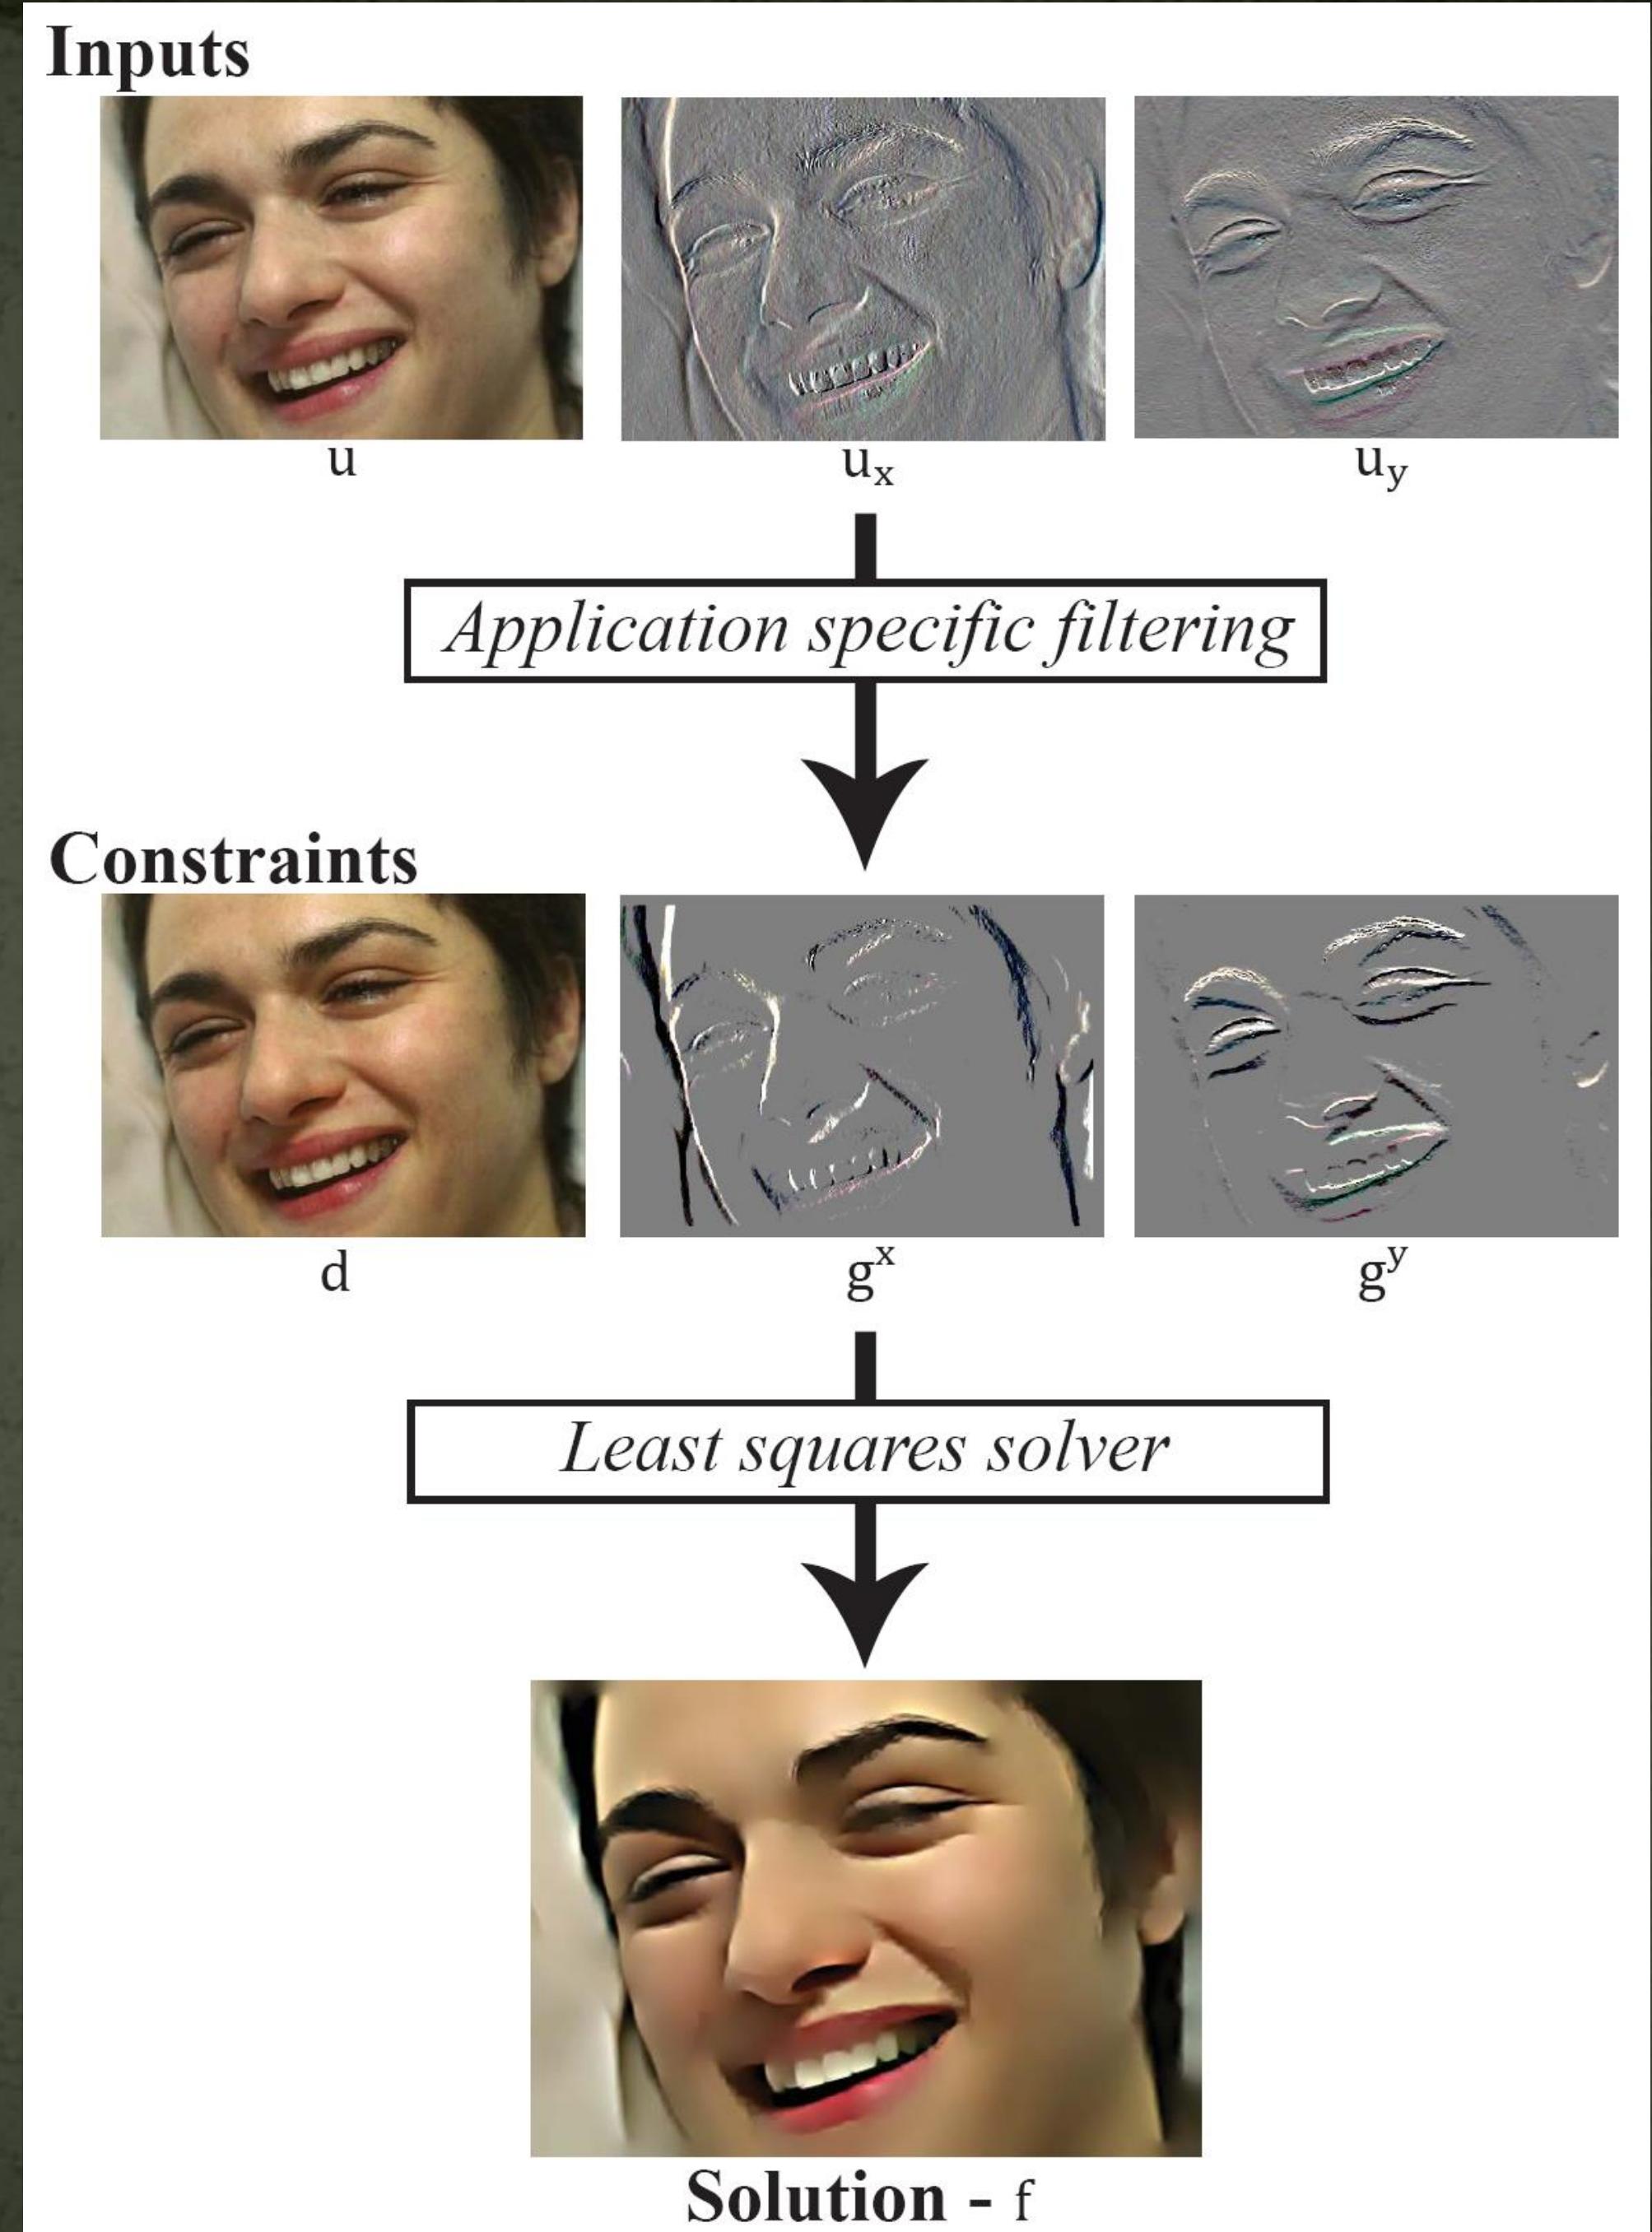


Motivation for gradient-domain filtering?

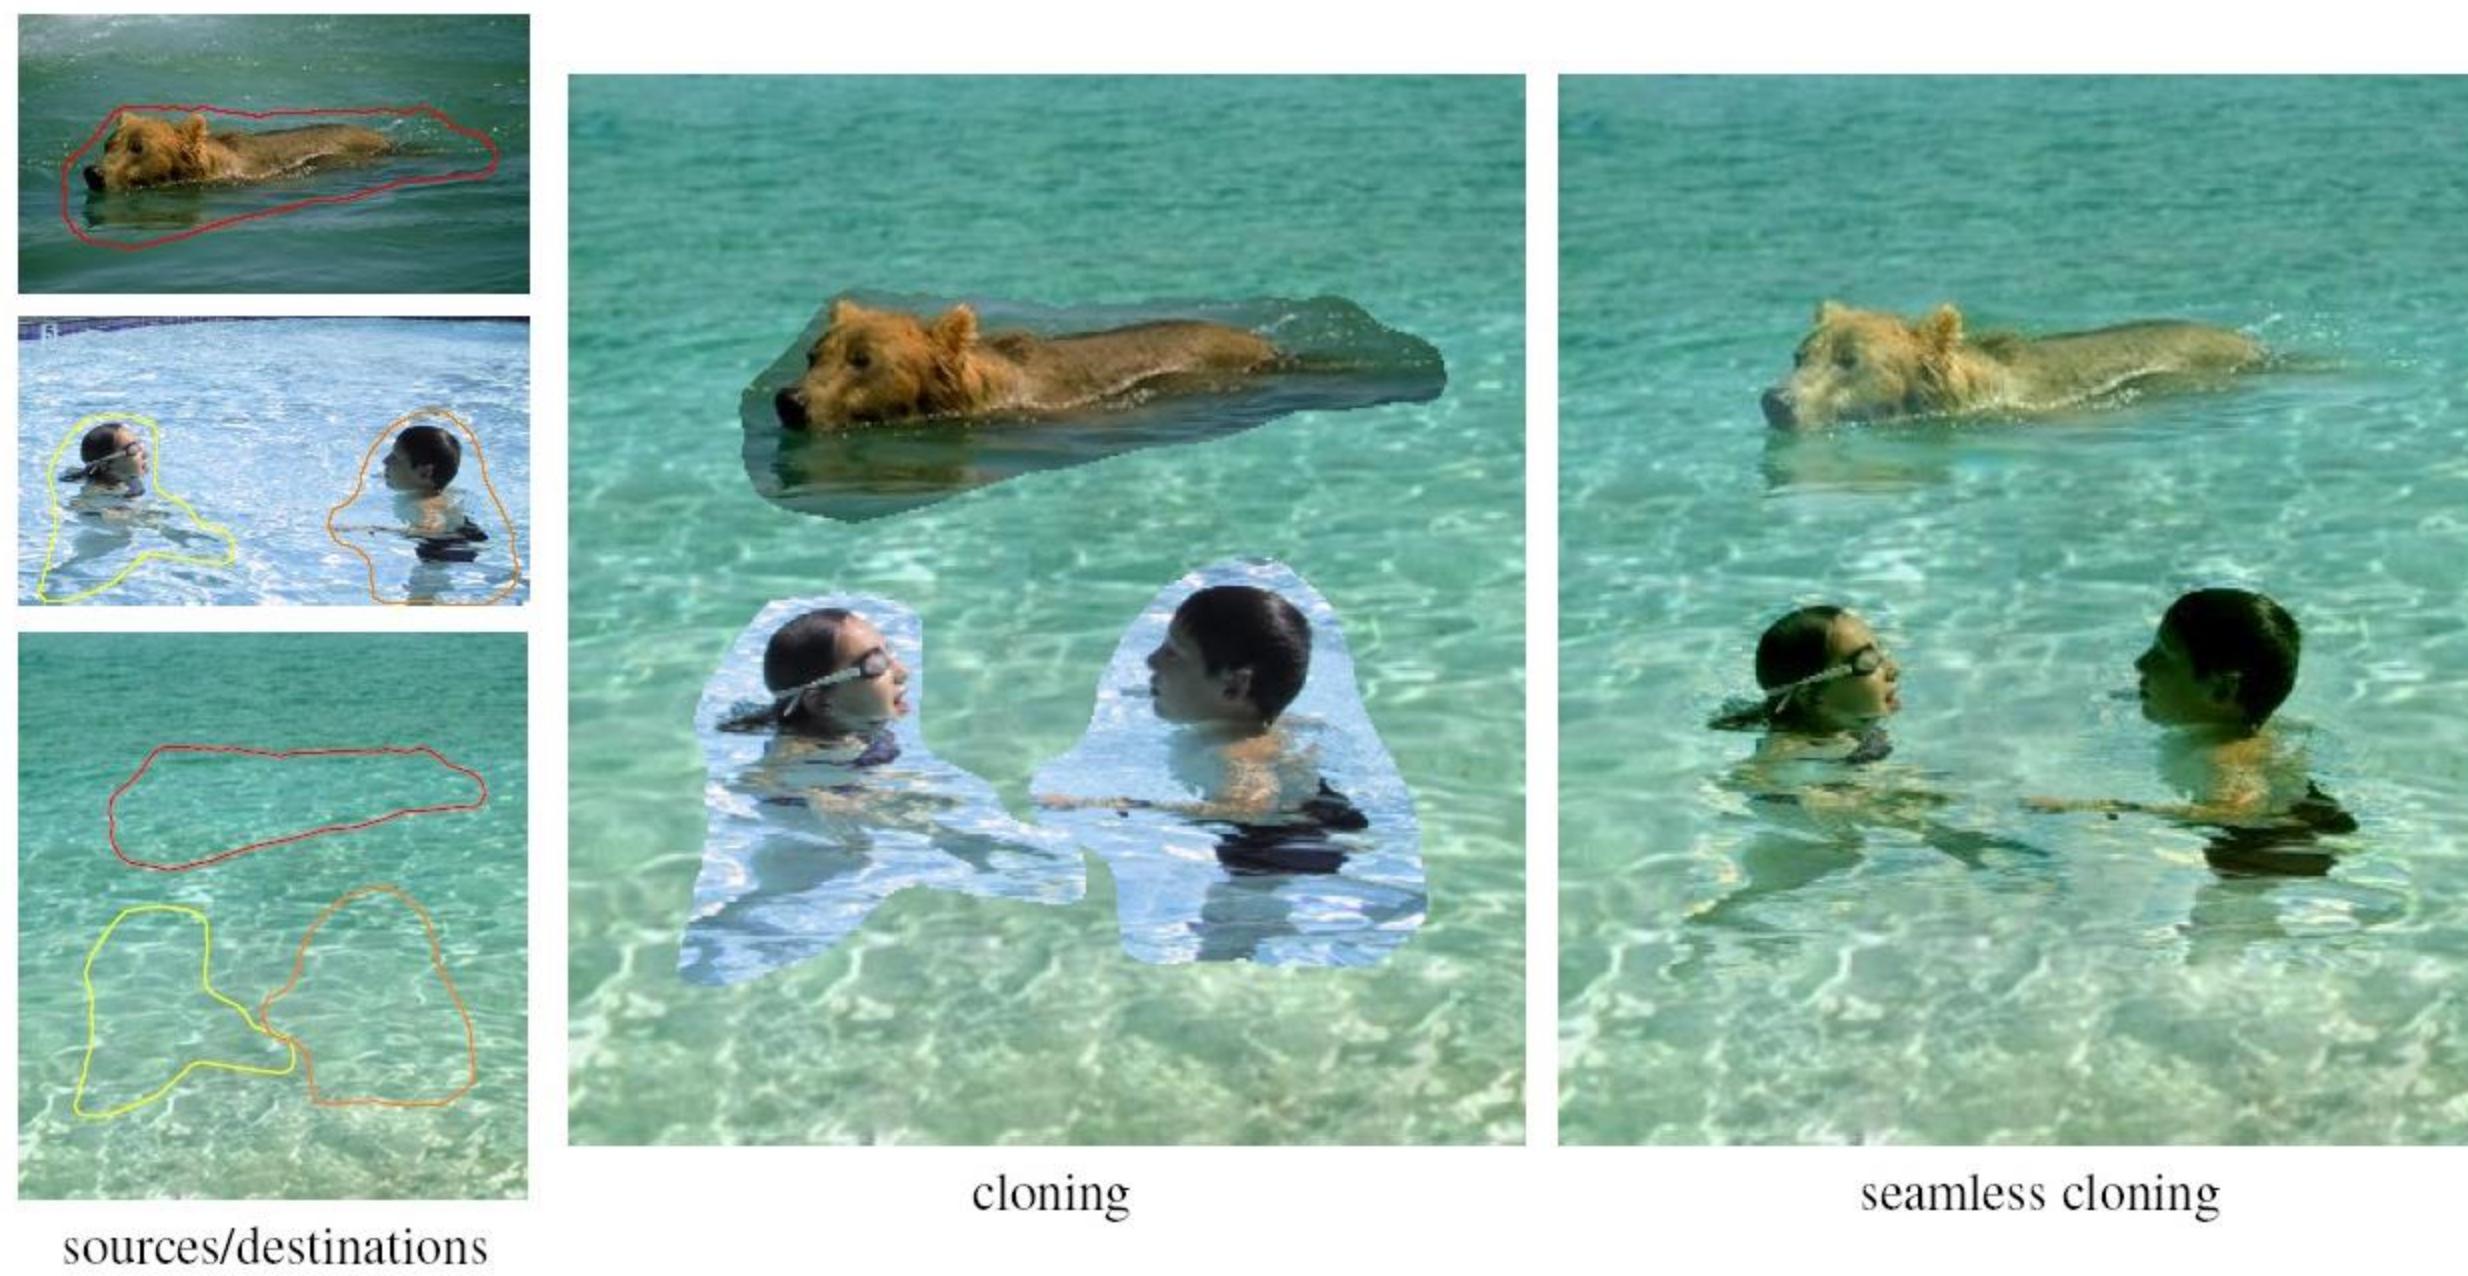
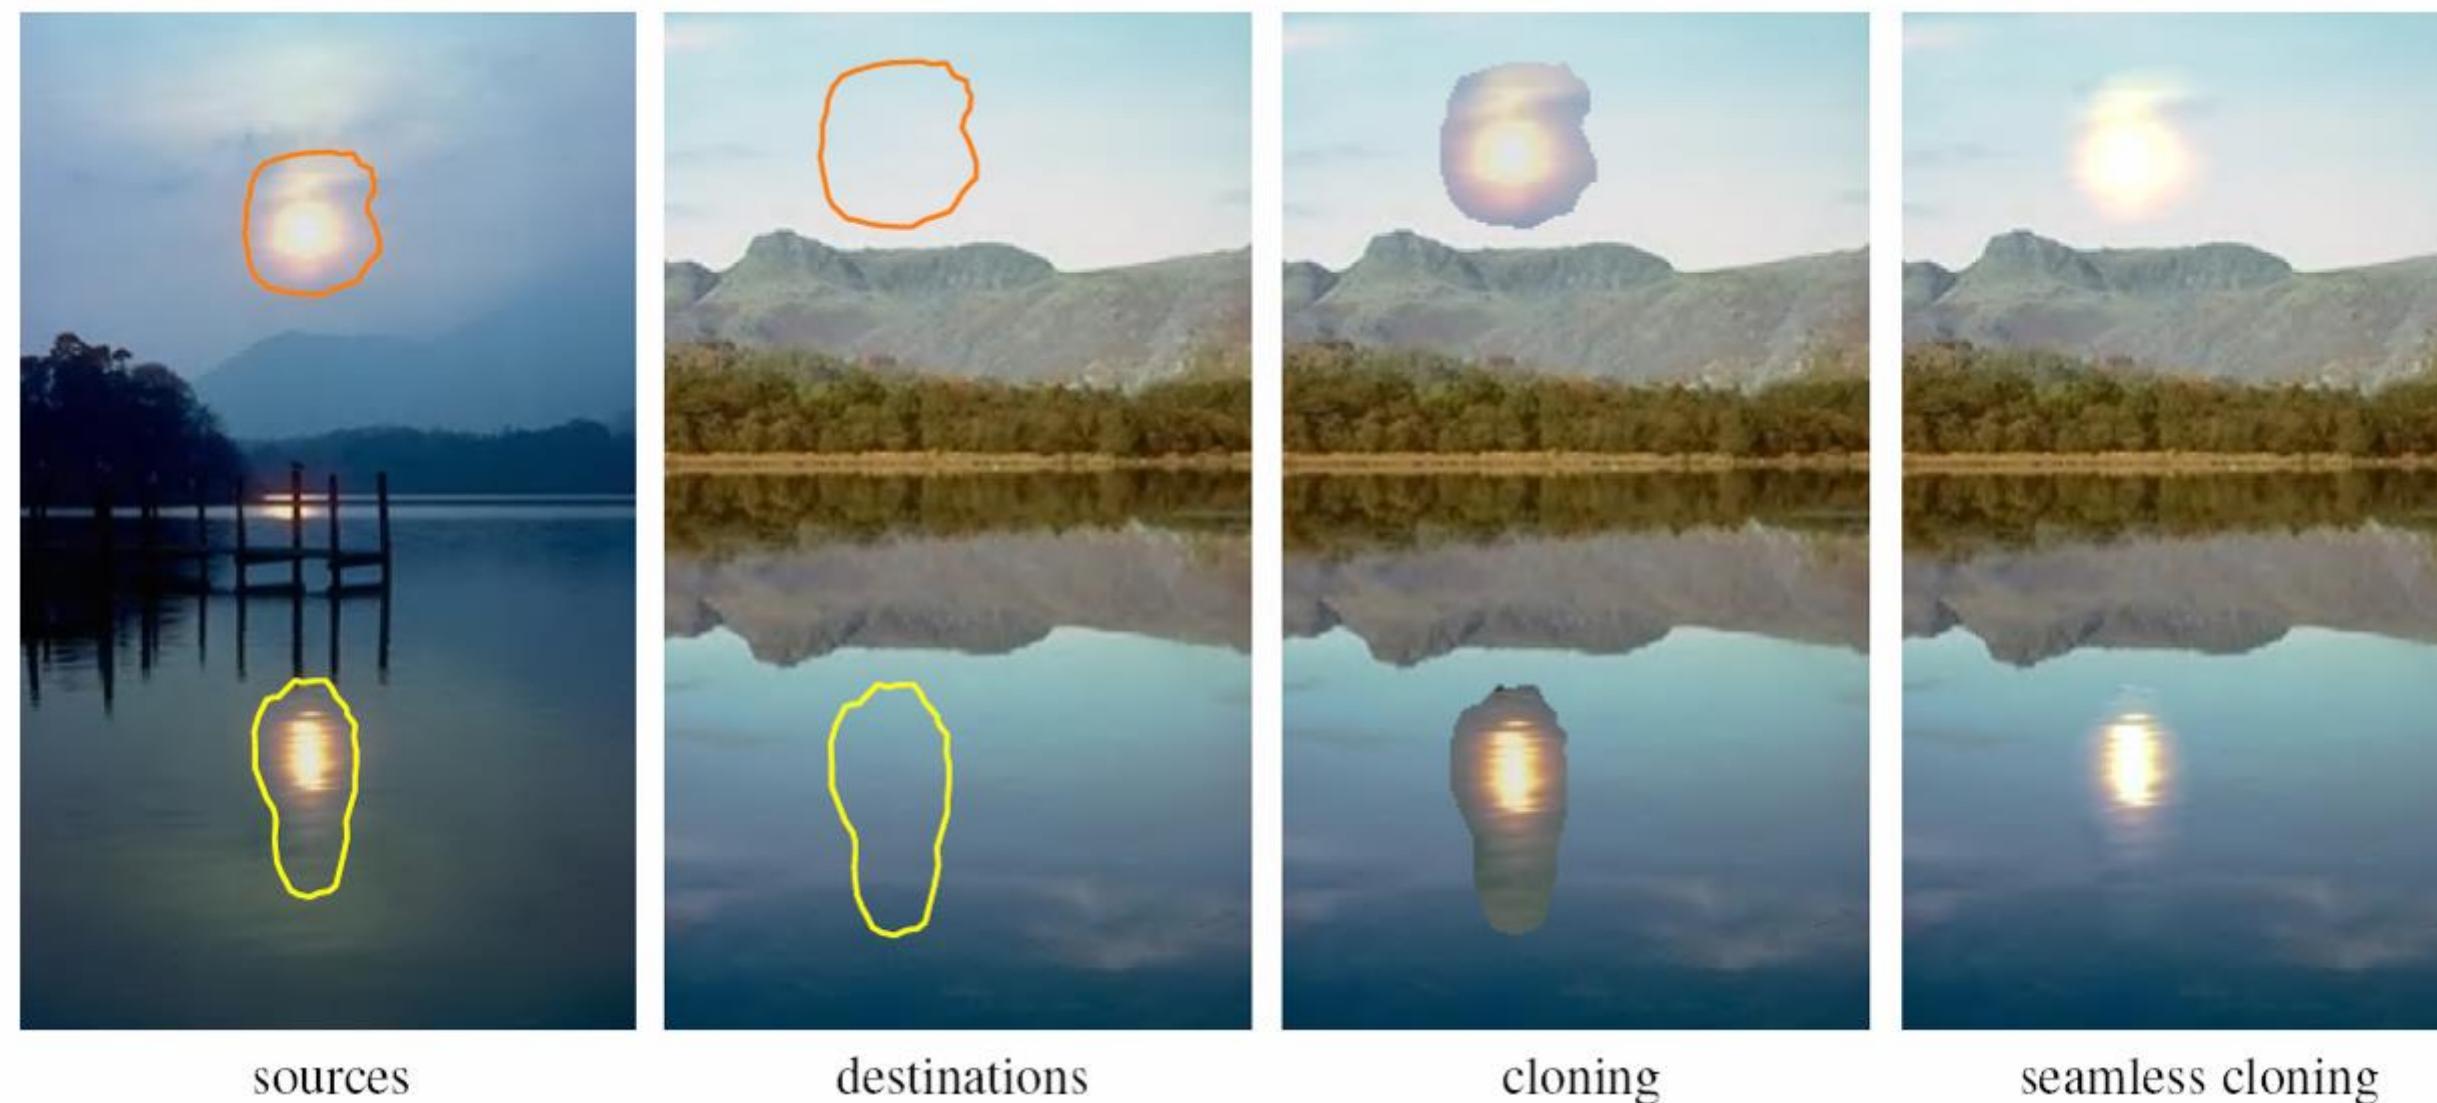
- Can be used to exert high-level control over images
 - gradients – low level image-features
 - gradients – give rise to high level image-features
 - manipulate local gradients to manipulate global image interpretation



Gradient Domain Editing



Perez et al., “Poisson Image Editing”, 2003



Drawing in Gradient Domain

Real-Time Gradient-Domain Painting

James McCann*
Carnegie Mellon University

Nancy S. Pollard†
Carnegie Mellon University

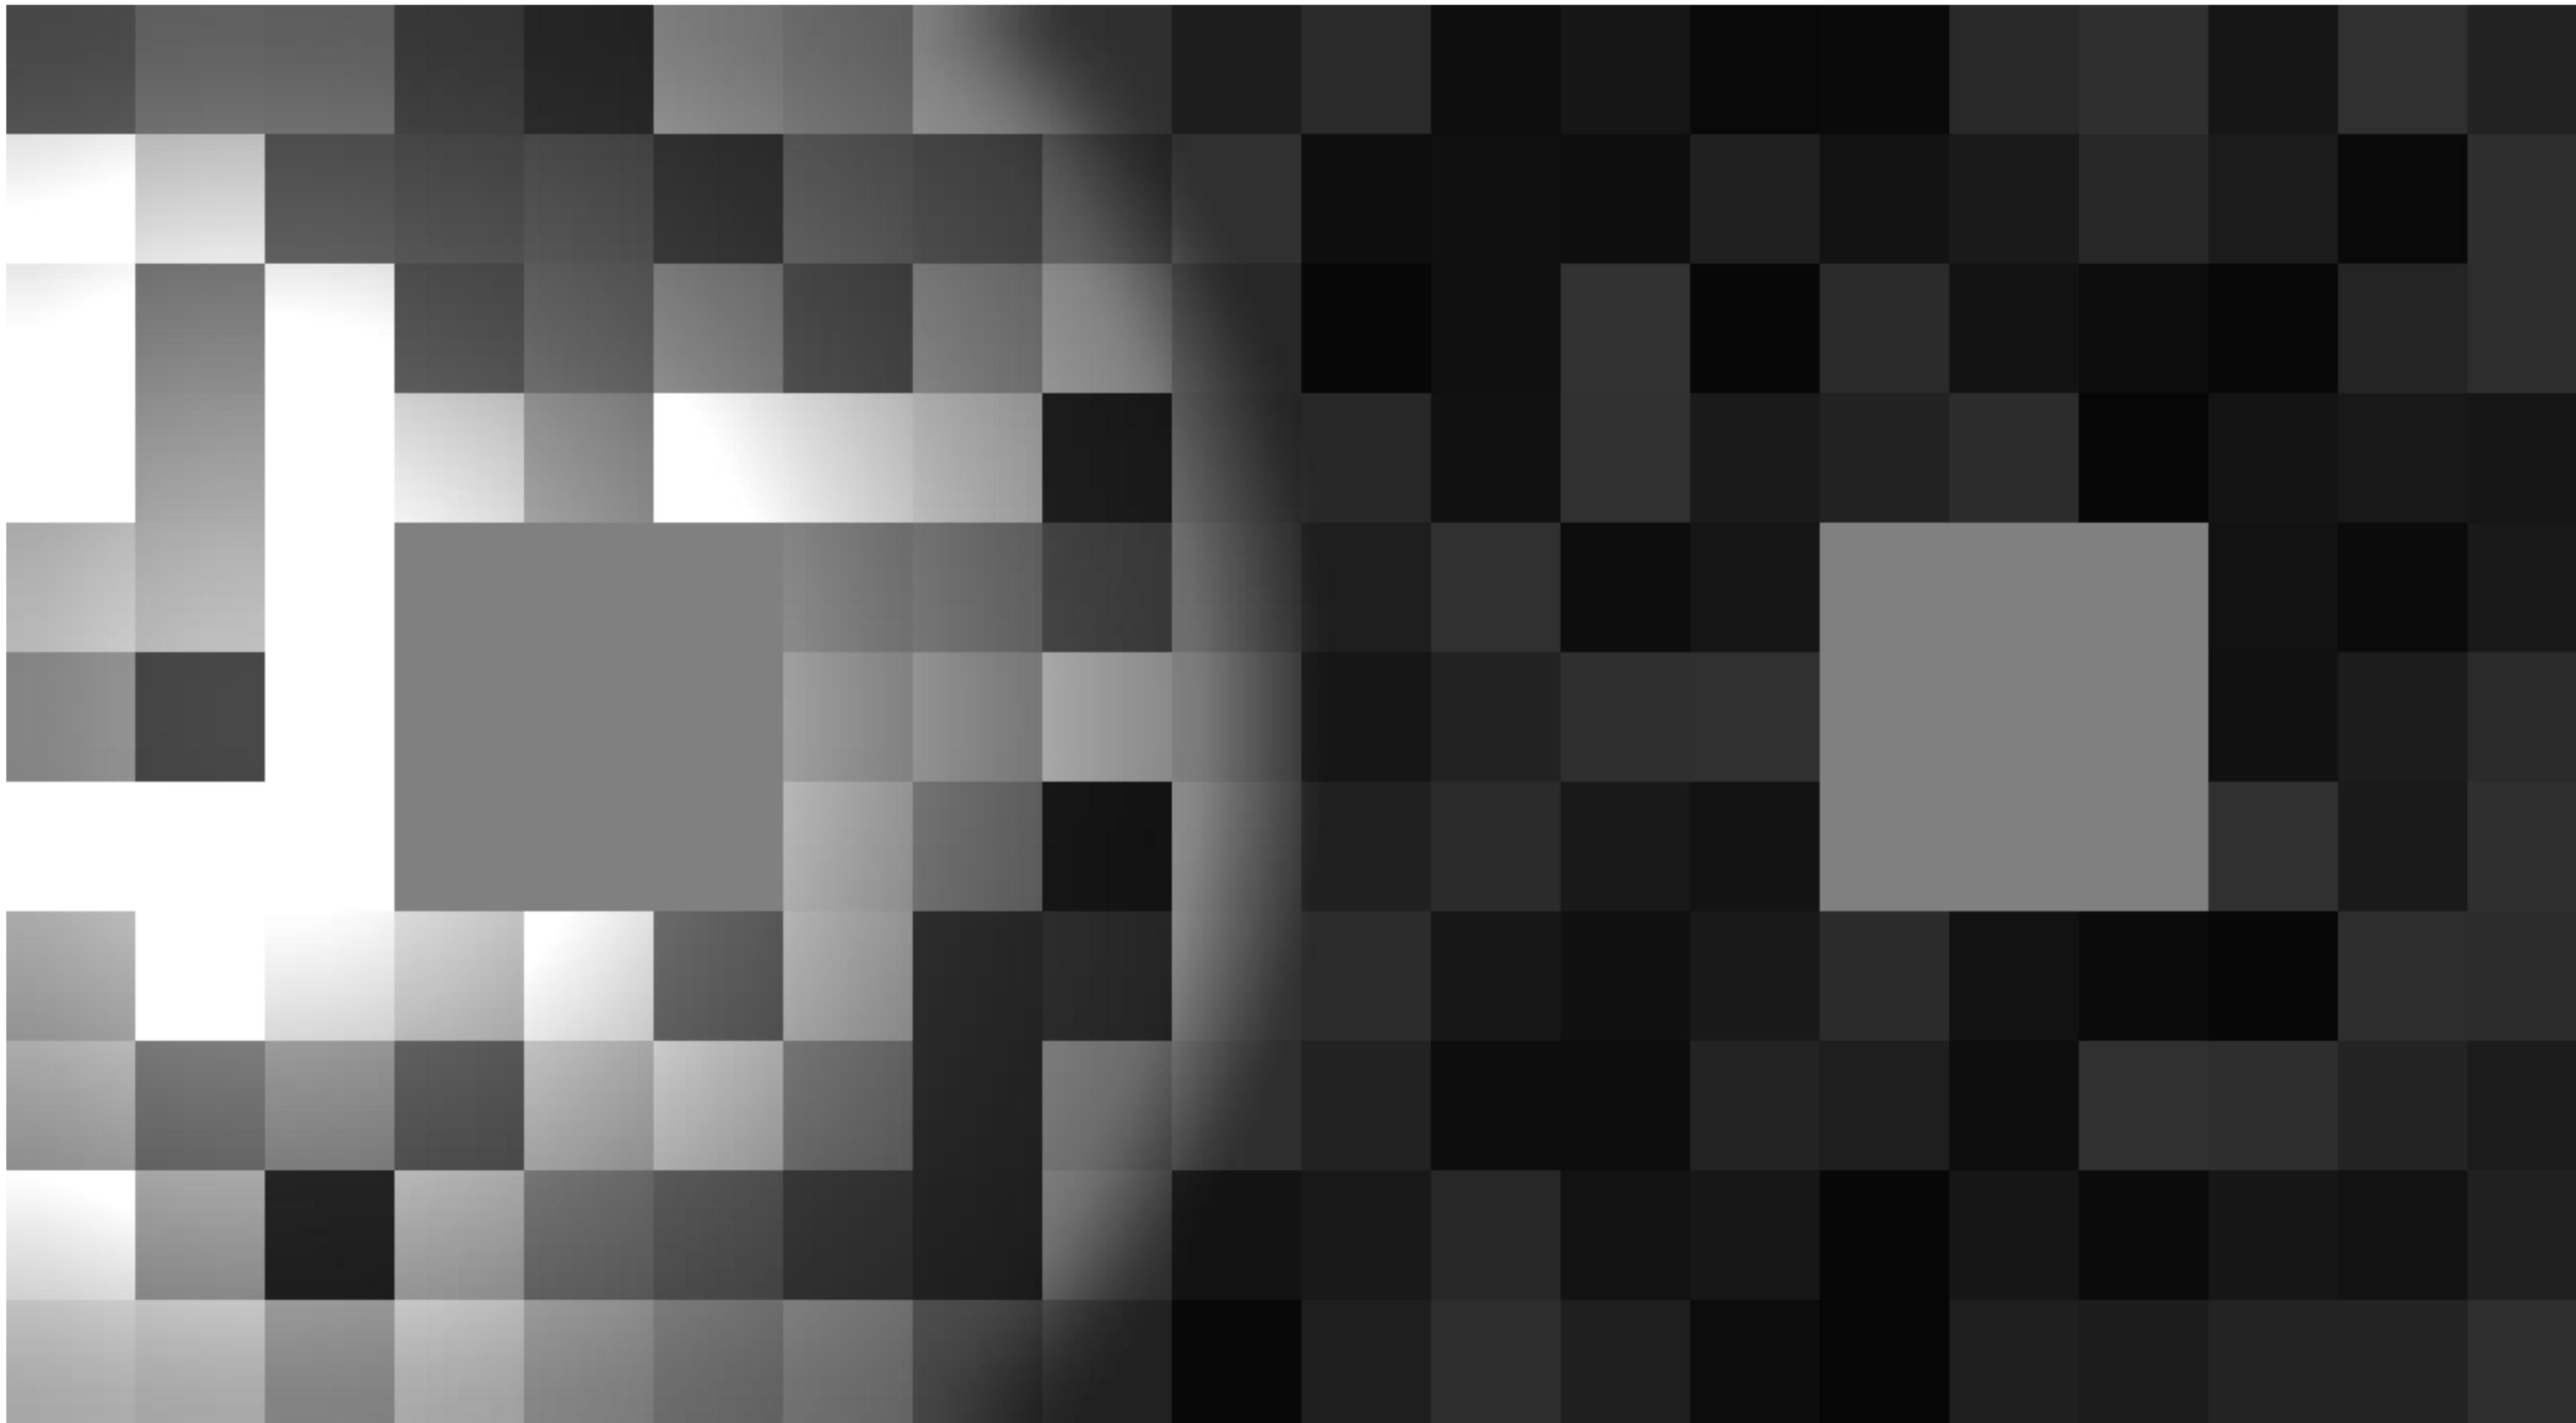


James McCann & Nancy Pollard
Real-Time Gradient-Domain Painting,
SIGGRAPH 2009
(paper came out of this class!)

<http://www.youtube.com/watch?v=RvhkAfrA0-w&feature=youtu.be>



Lightness Constancy



Simultaneous Contrast Illusion

Retinex (Land and McCann)

Intrinsic Images:

- *Image = reflectance image * illumination image*

$$\log \ell(x, y) = \log r(x, y) + \log l(x, y) \quad (18.38)$$

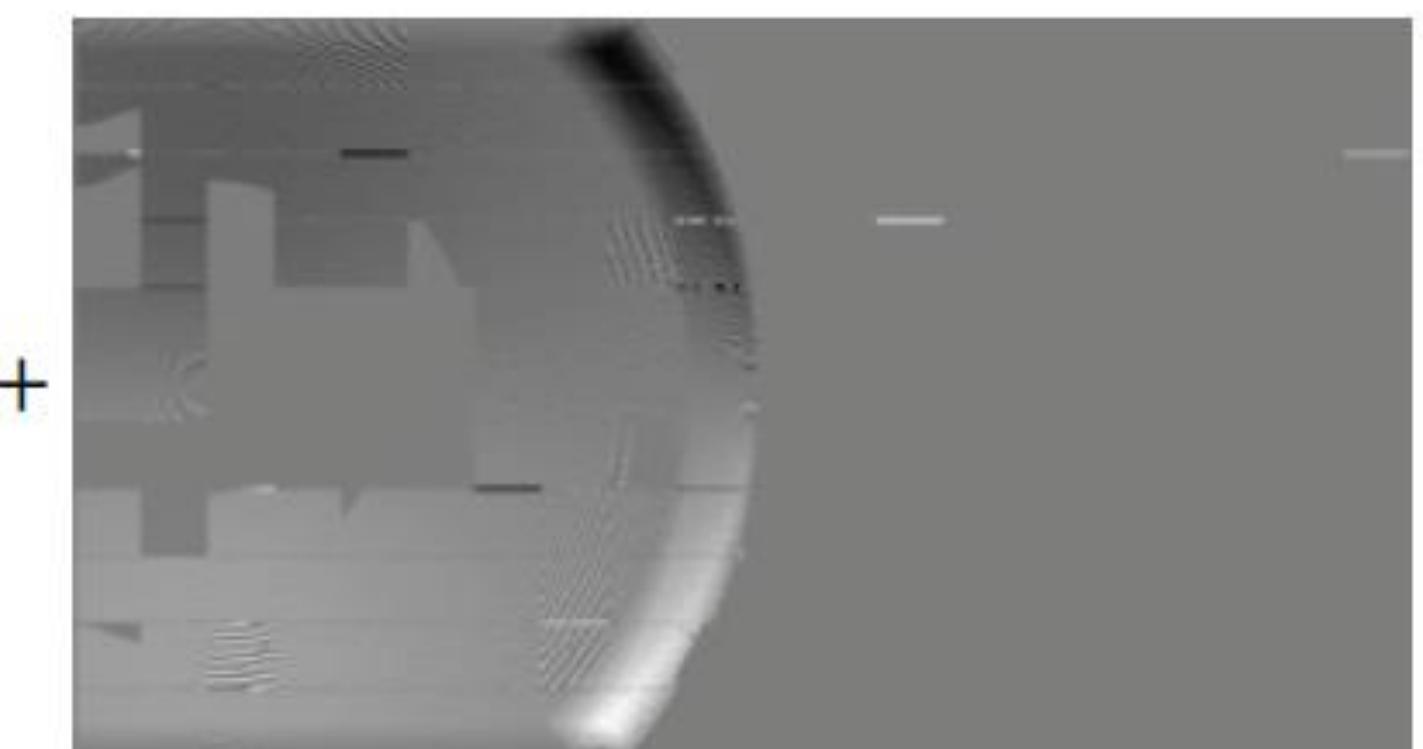
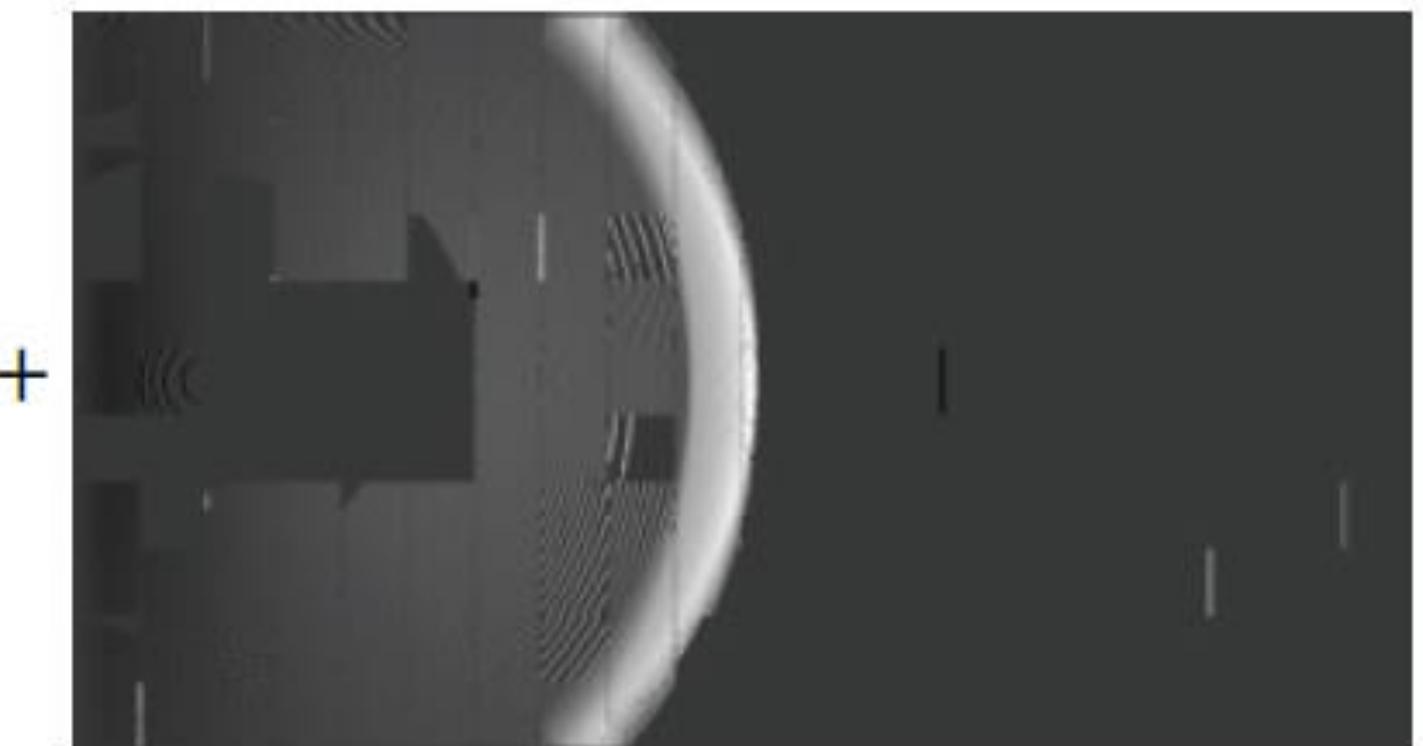
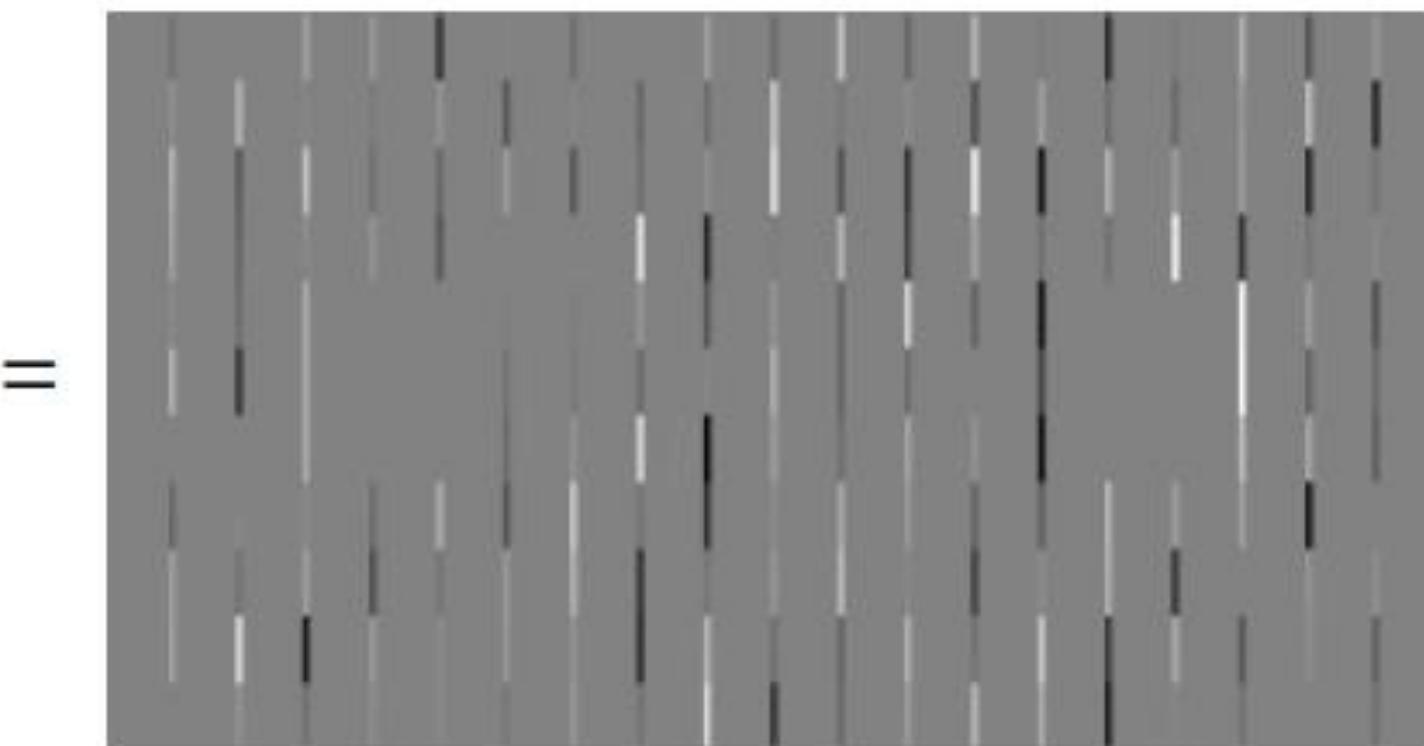
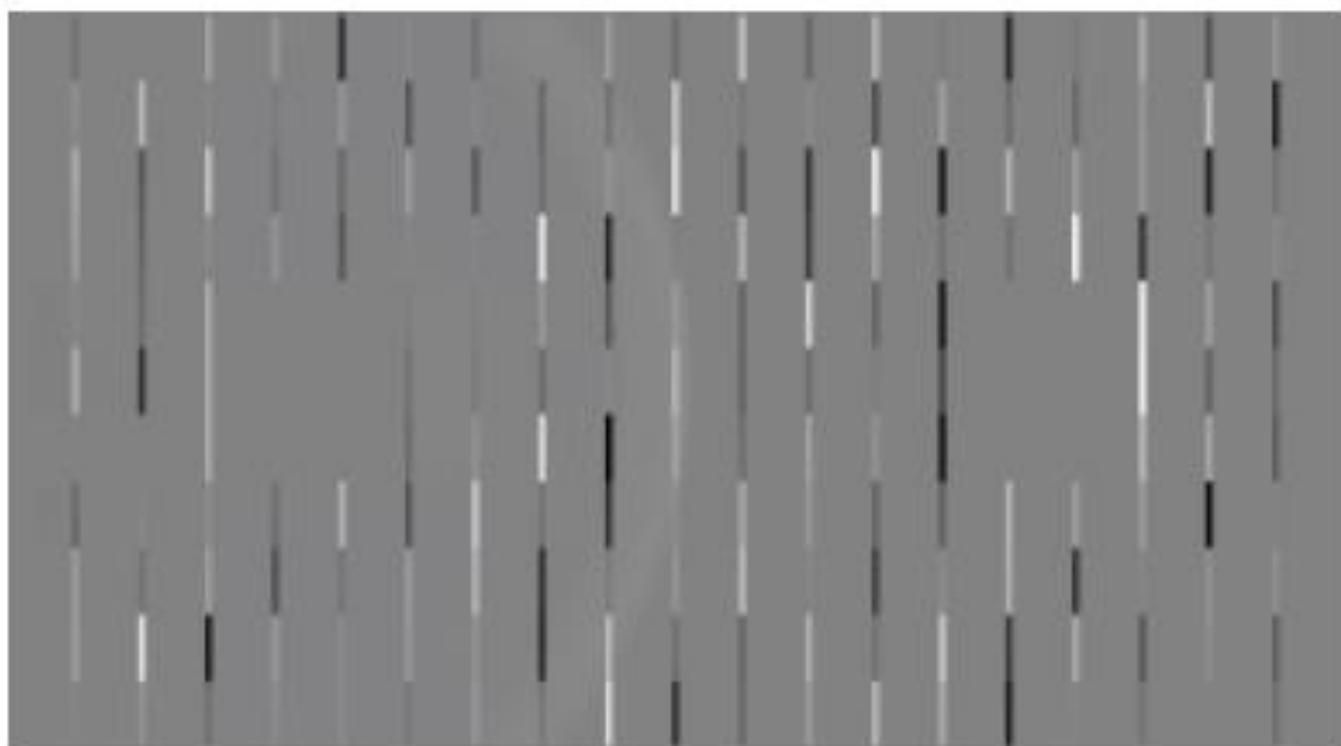
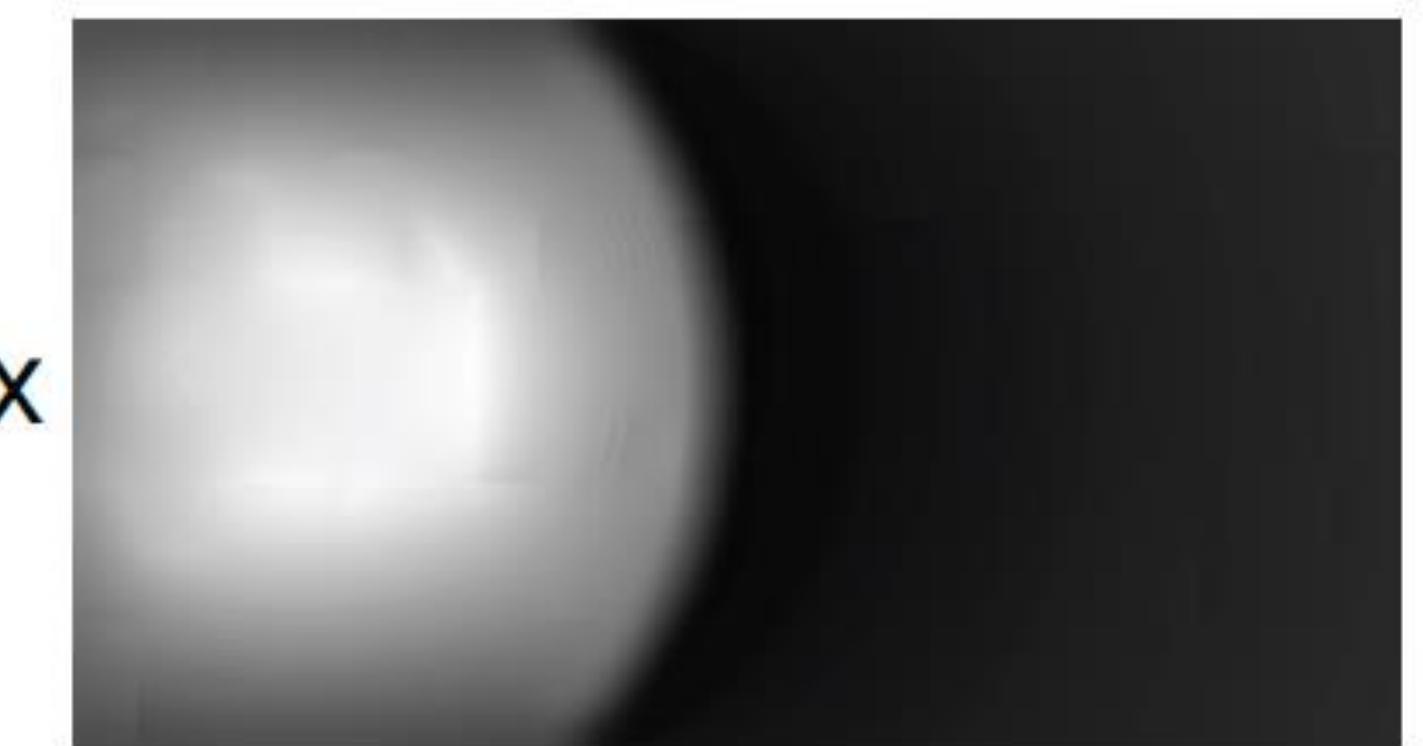
Taking derivatives along x and y is now simple:

$$\frac{\partial \log \ell(x, y)}{\partial x} = \frac{\partial \log r(x, y)}{\partial x} + \frac{\partial \log l(x, y)}{\partial x} \quad (18.39)$$

And the same thing is done for the derivative along y .

Any derivative larger than the threshold is assigned to the derivative of the reflectance image $r(x, y)$ and the ones smaller than a threshold are assigned to the illumination image $l(x, y)$:

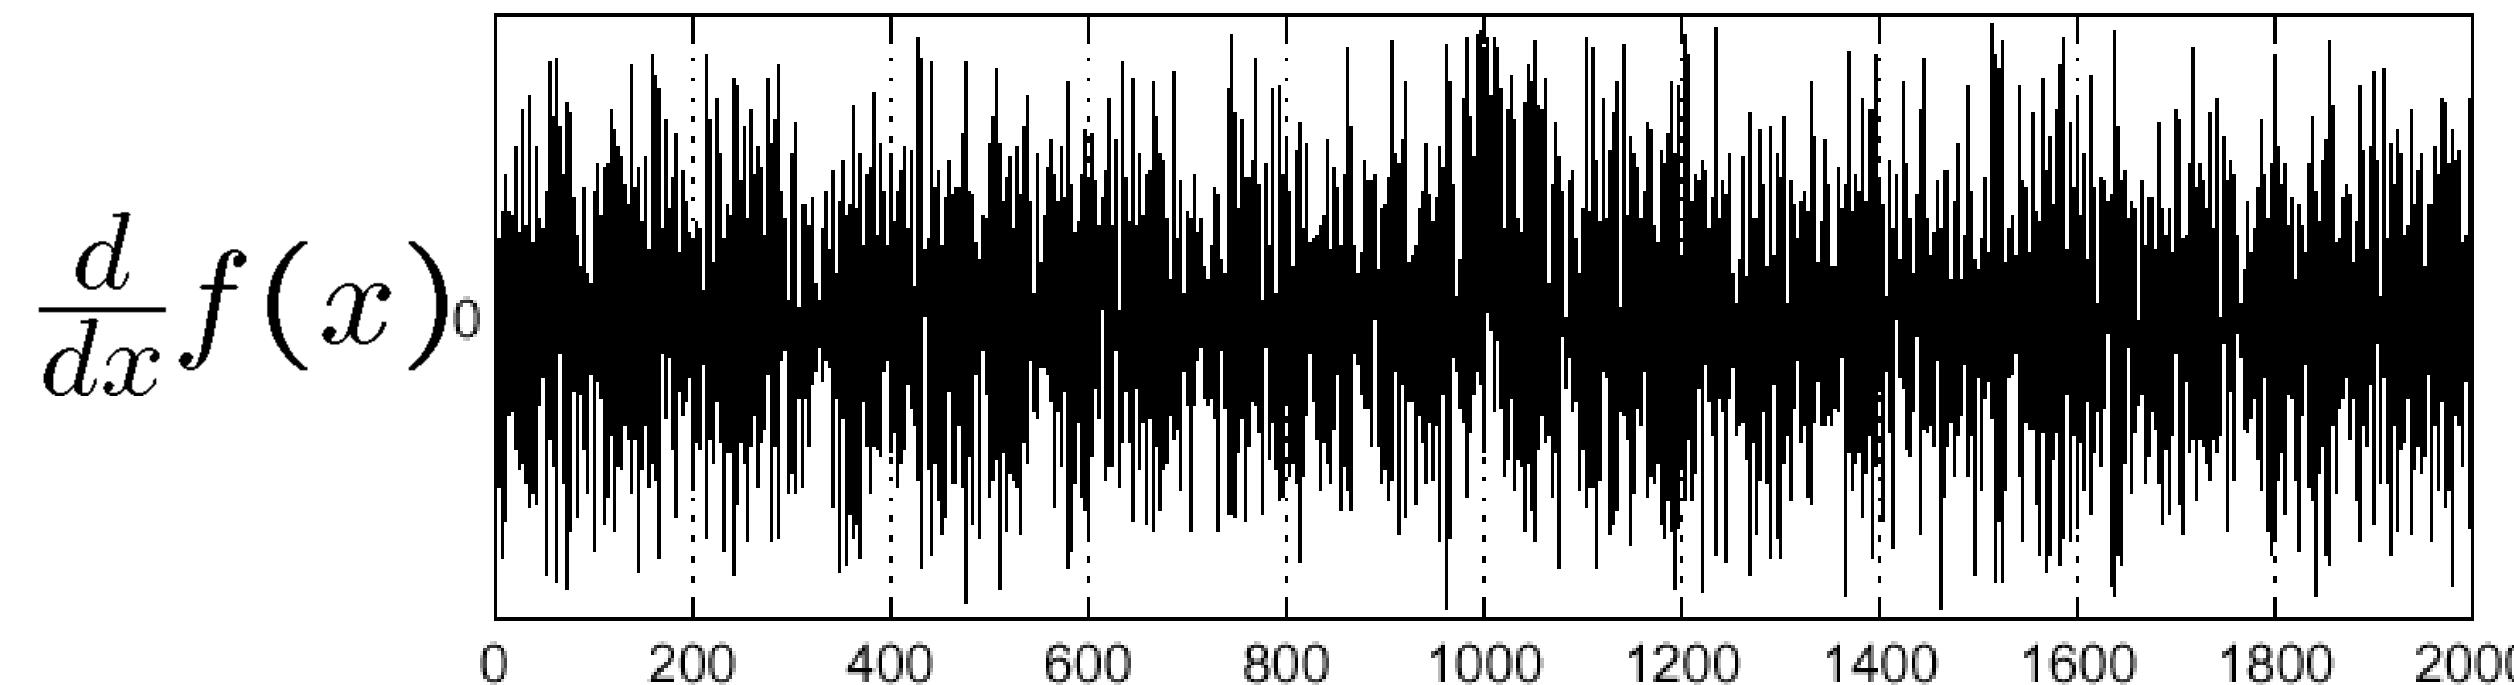
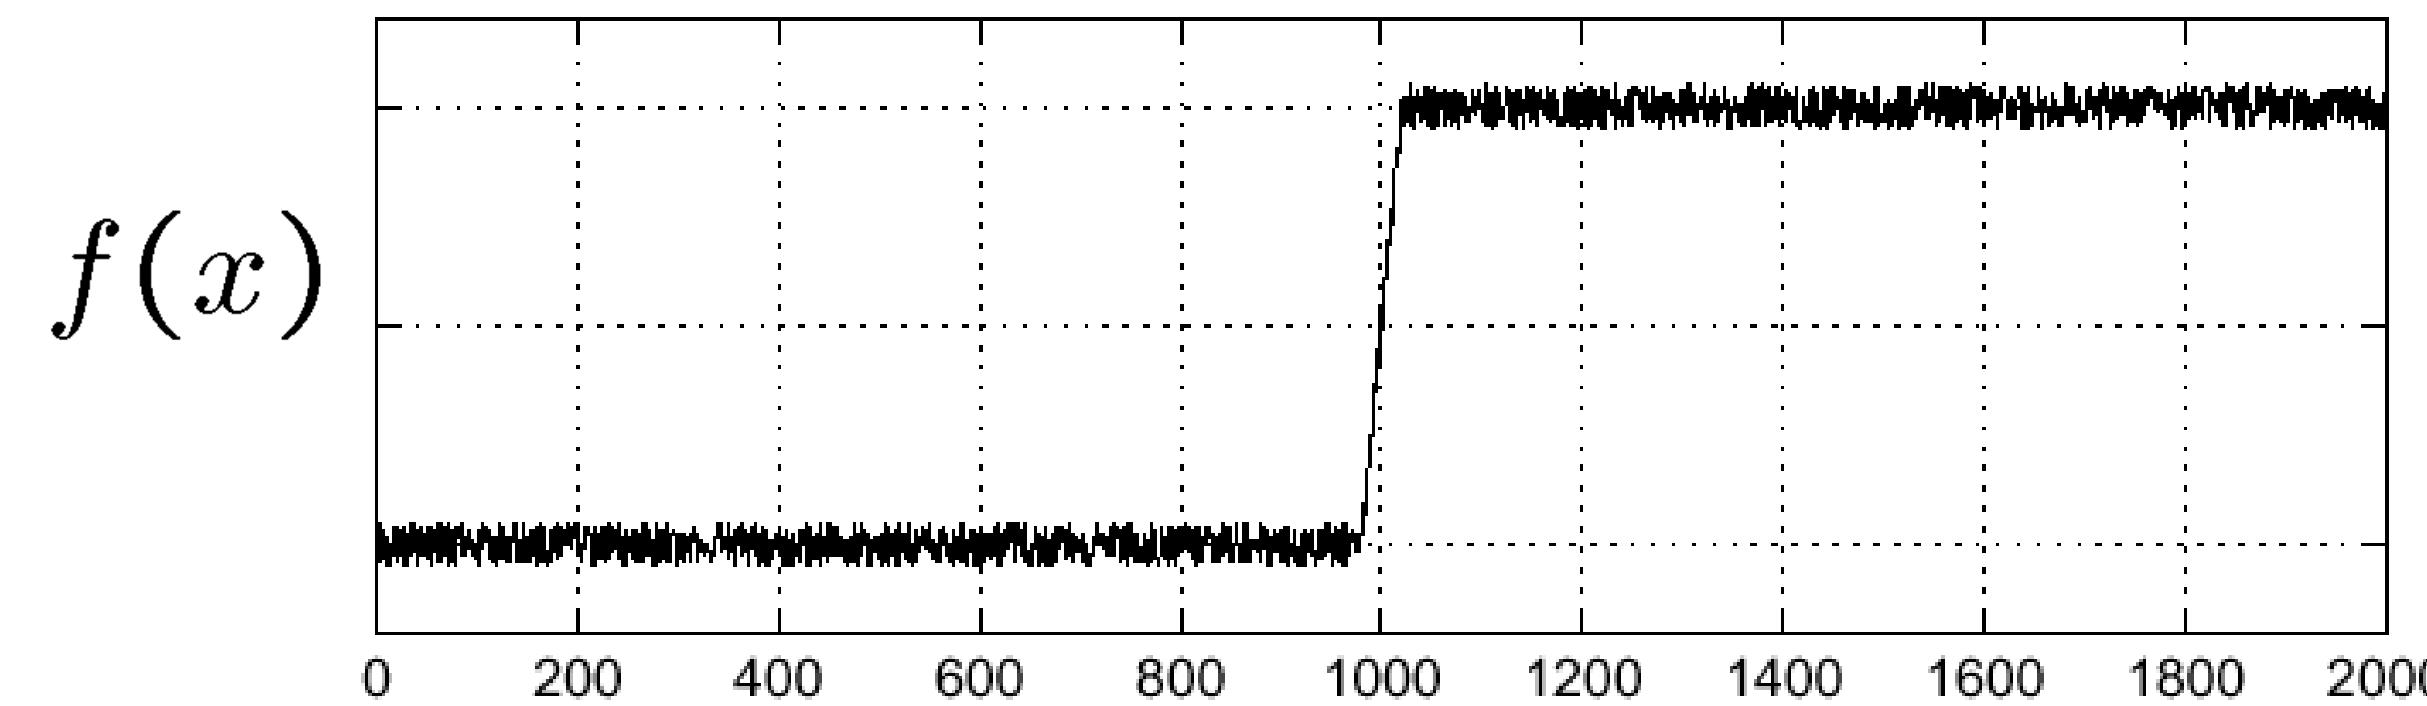
$$\frac{\partial \log r(x, y)}{\partial x} = \begin{cases} \frac{\partial \log \ell(x, y)}{\partial x} & \text{if } \left| \frac{\partial \log \ell(x, y)}{\partial x} \right| > T \\ 0 & \text{otherwise} \end{cases} \quad (18.40)$$

 $\ell(x, y)$  $r(x, y)$  $l(x, y)$ \times

Effects of noise

Consider a single row or column of the image

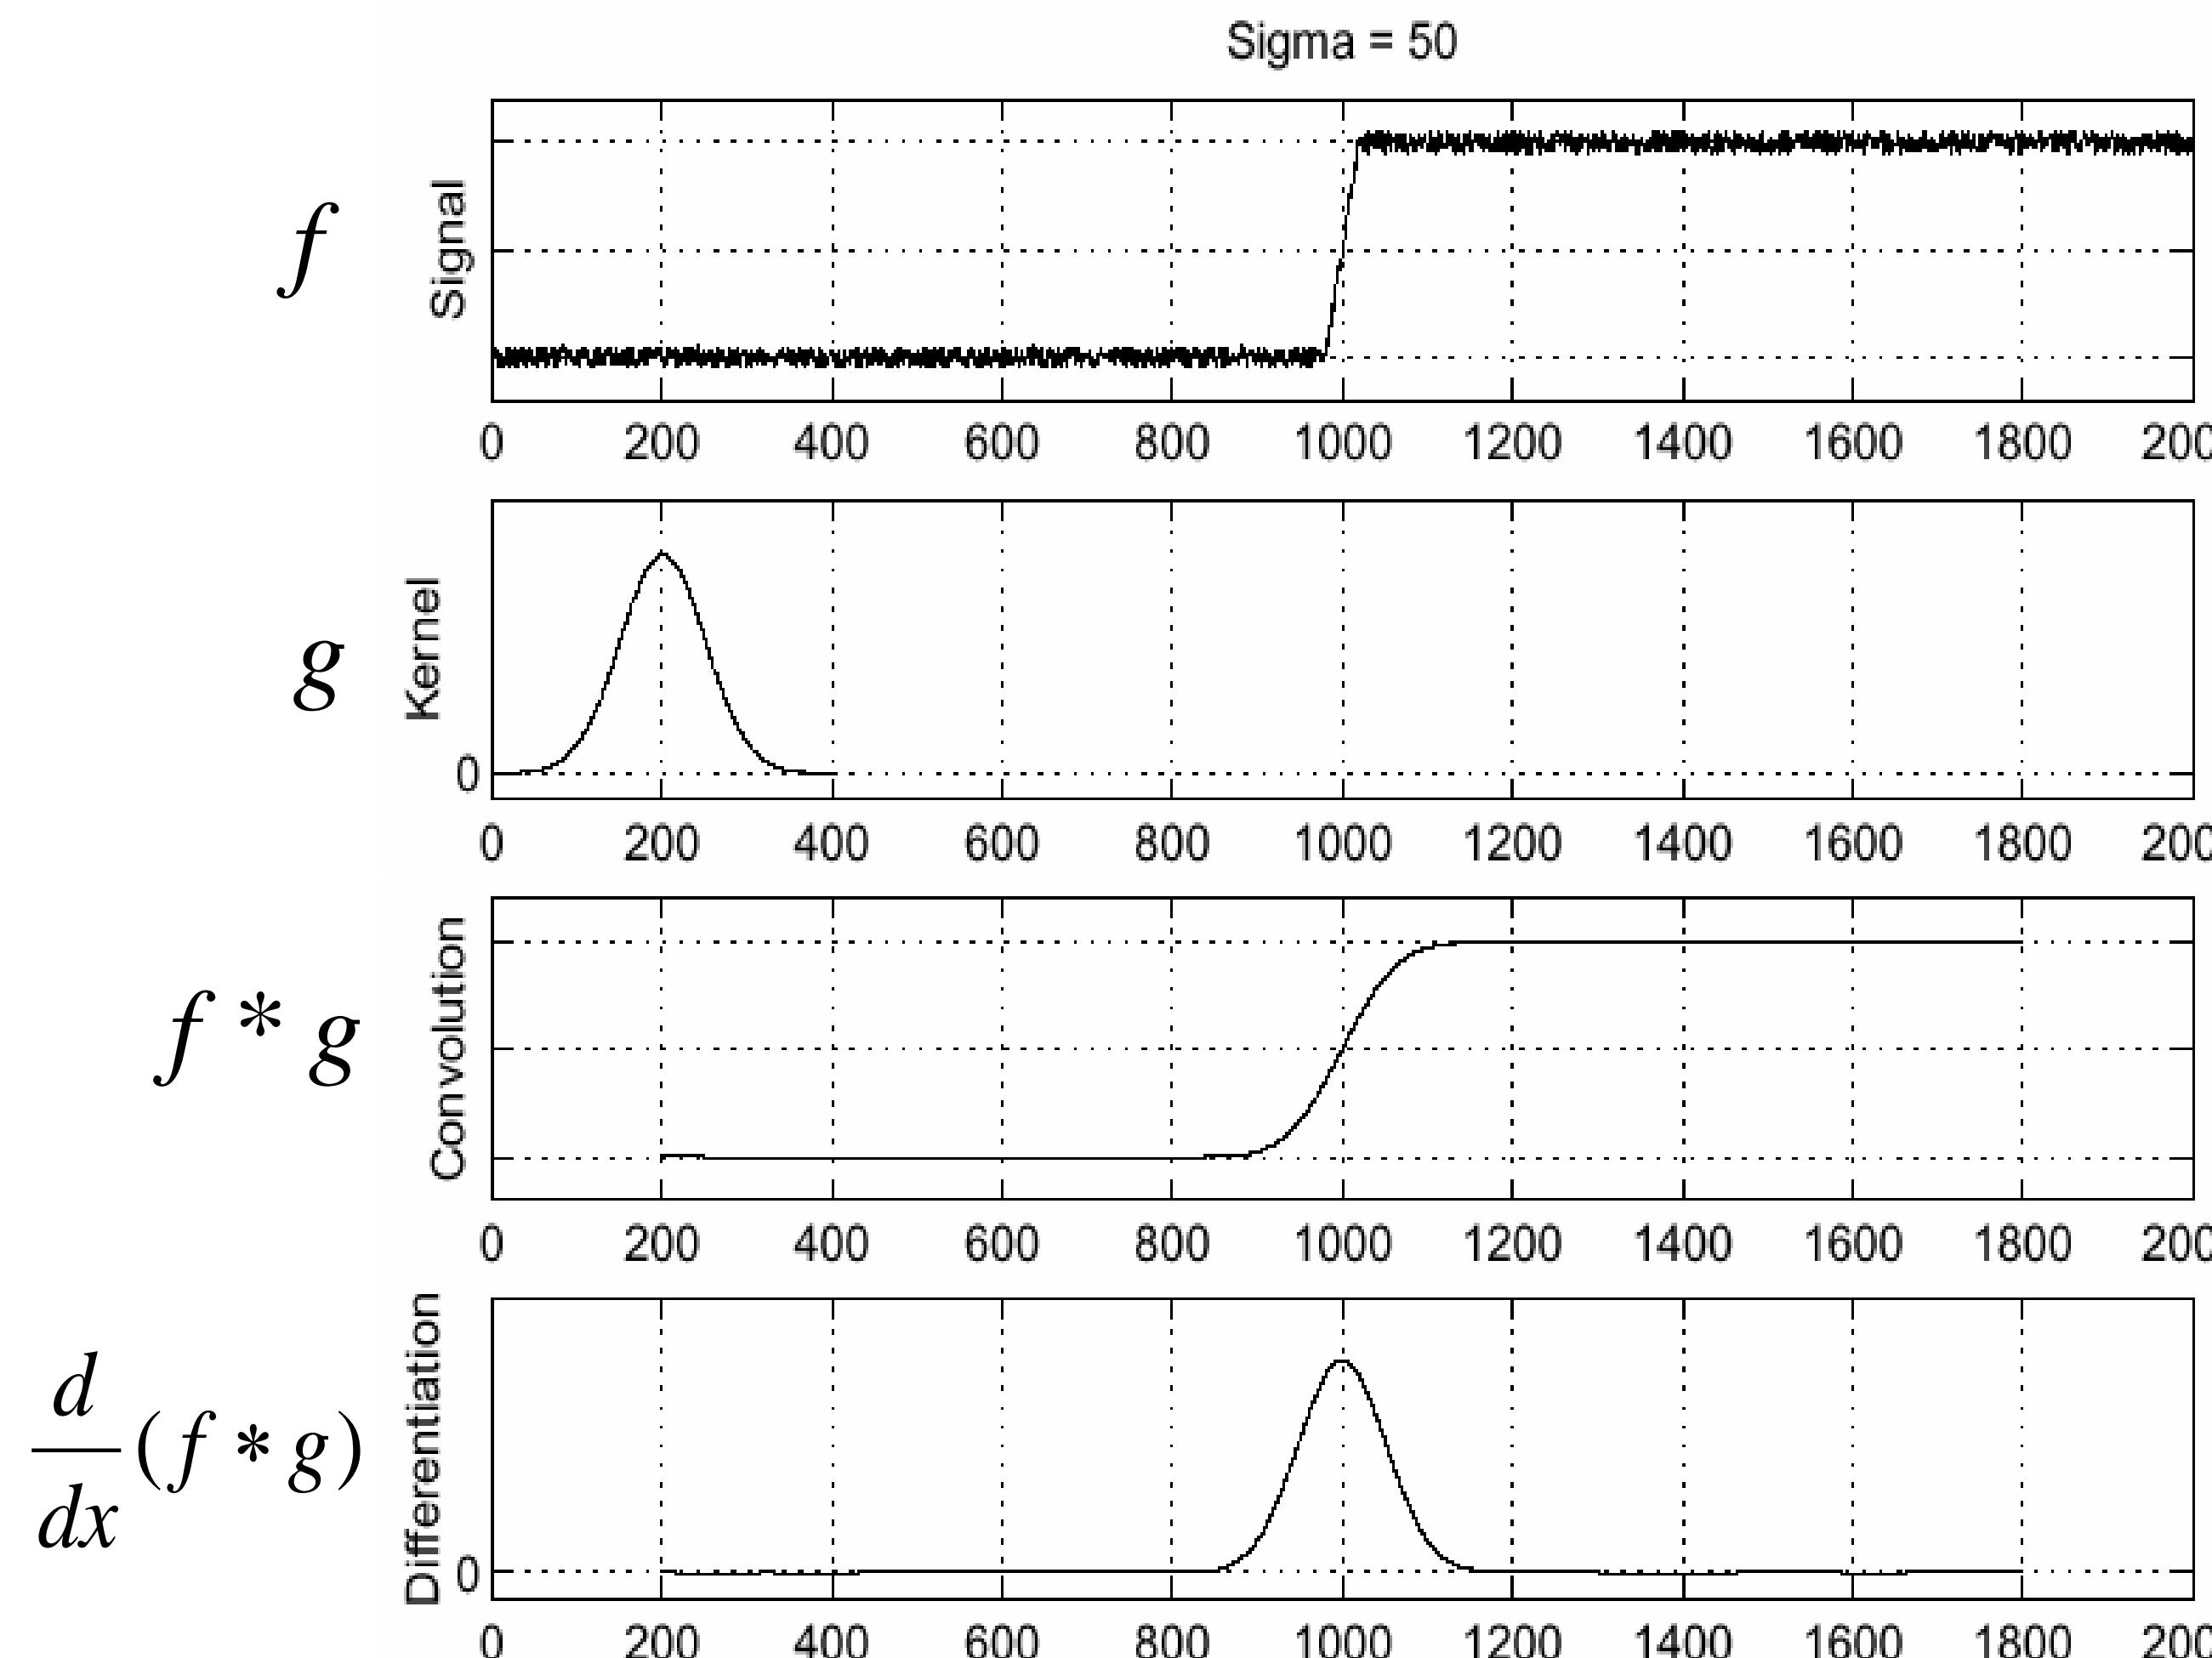
- Plotting intensity as a function of position gives a signal



Where is the edge?

Source: S. Seitz

Solution: smooth first

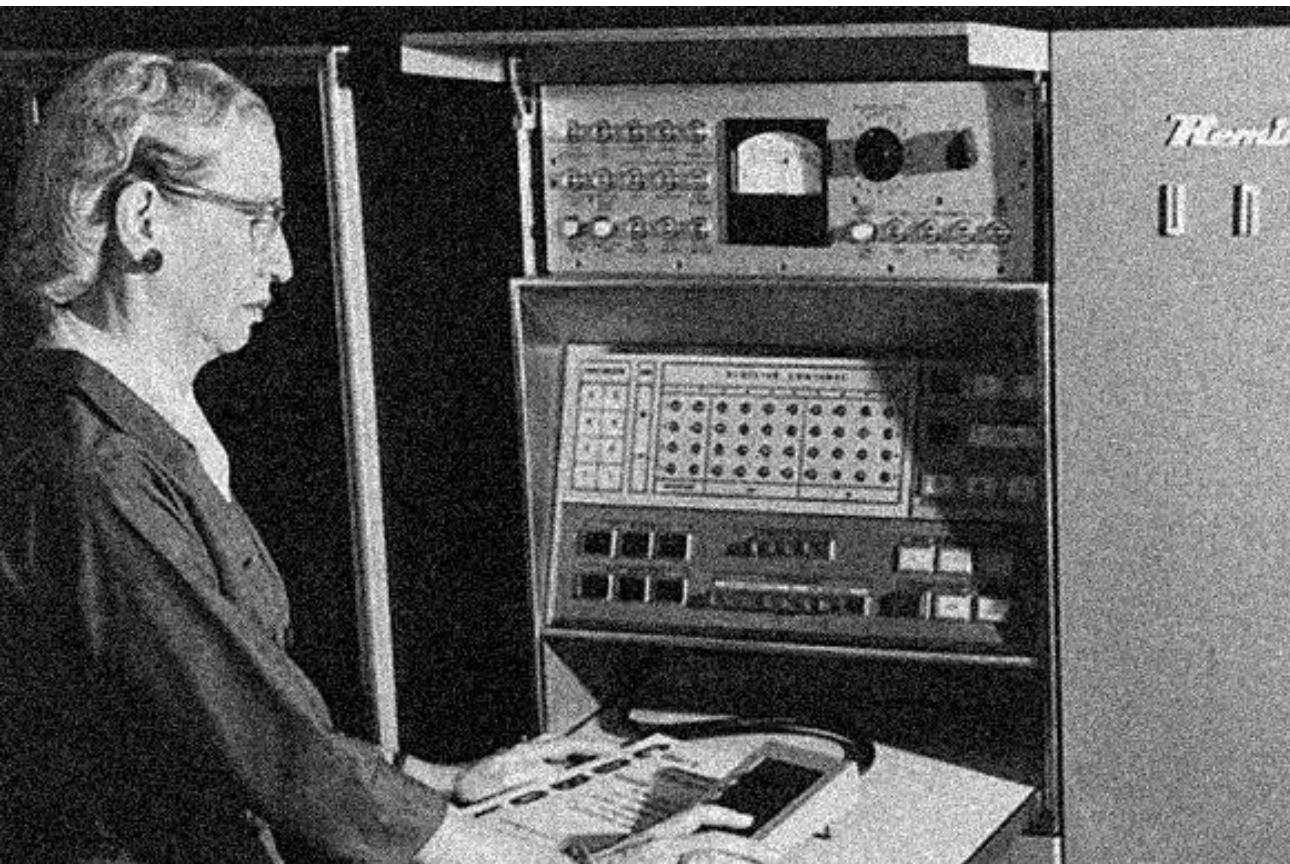


- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

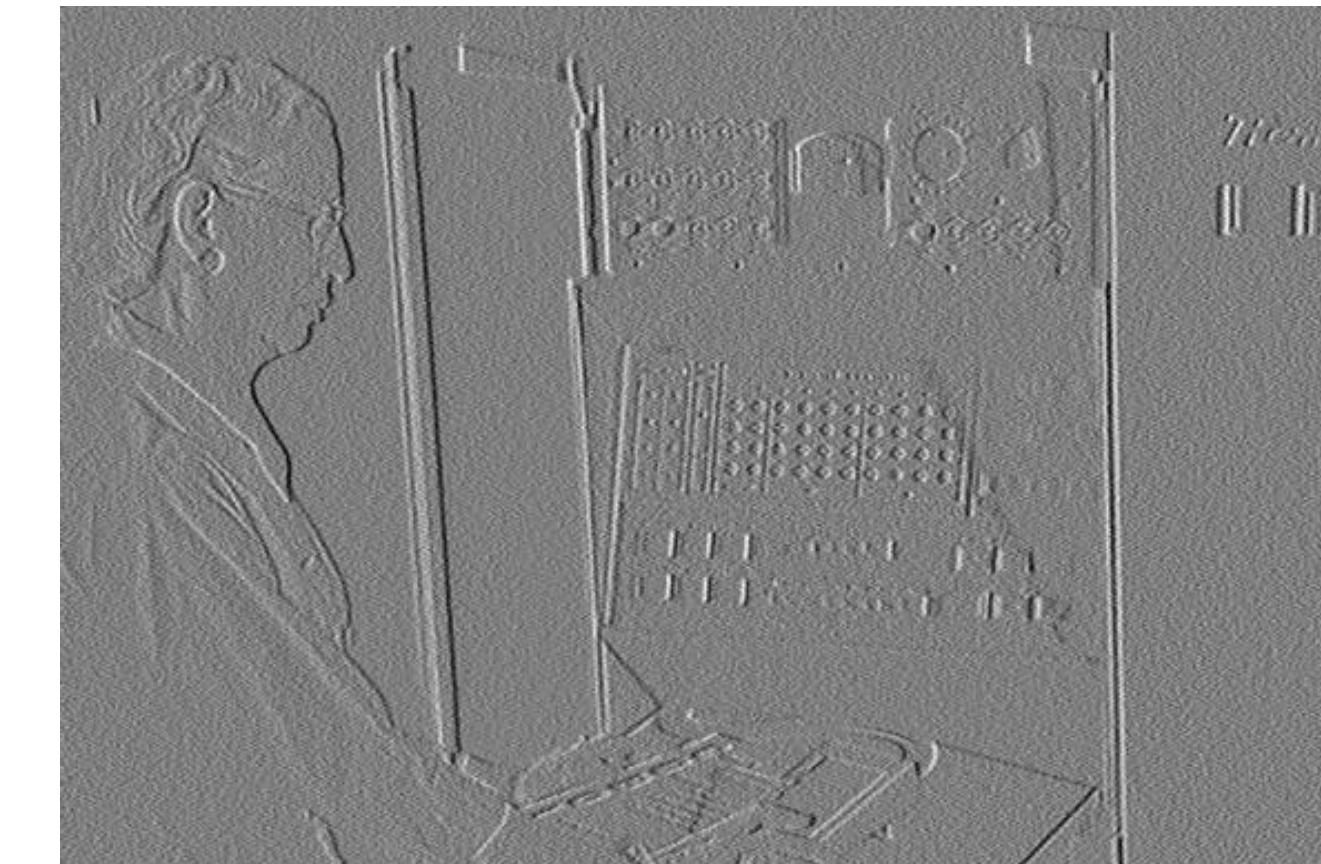
Source: S. Seitz

Noise in 2D

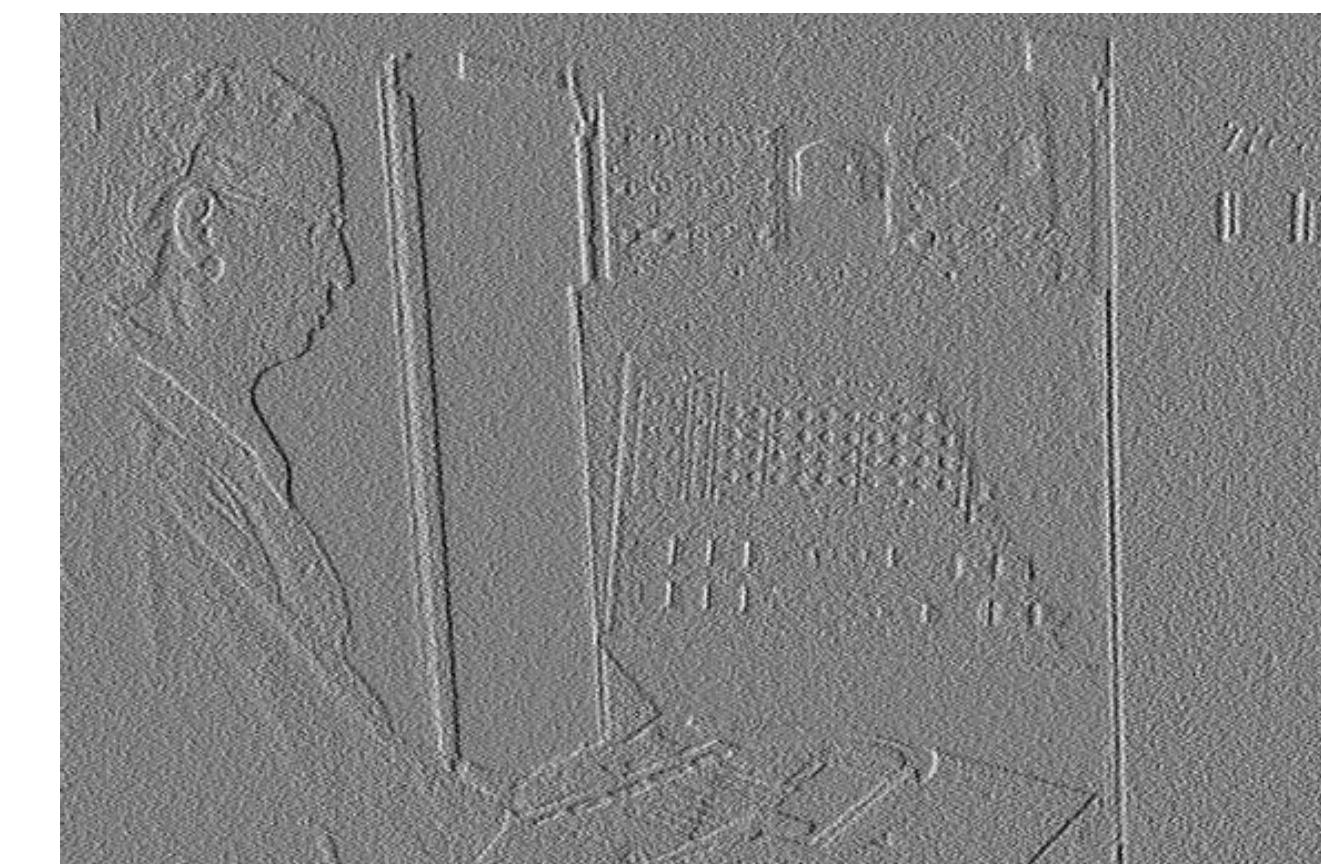
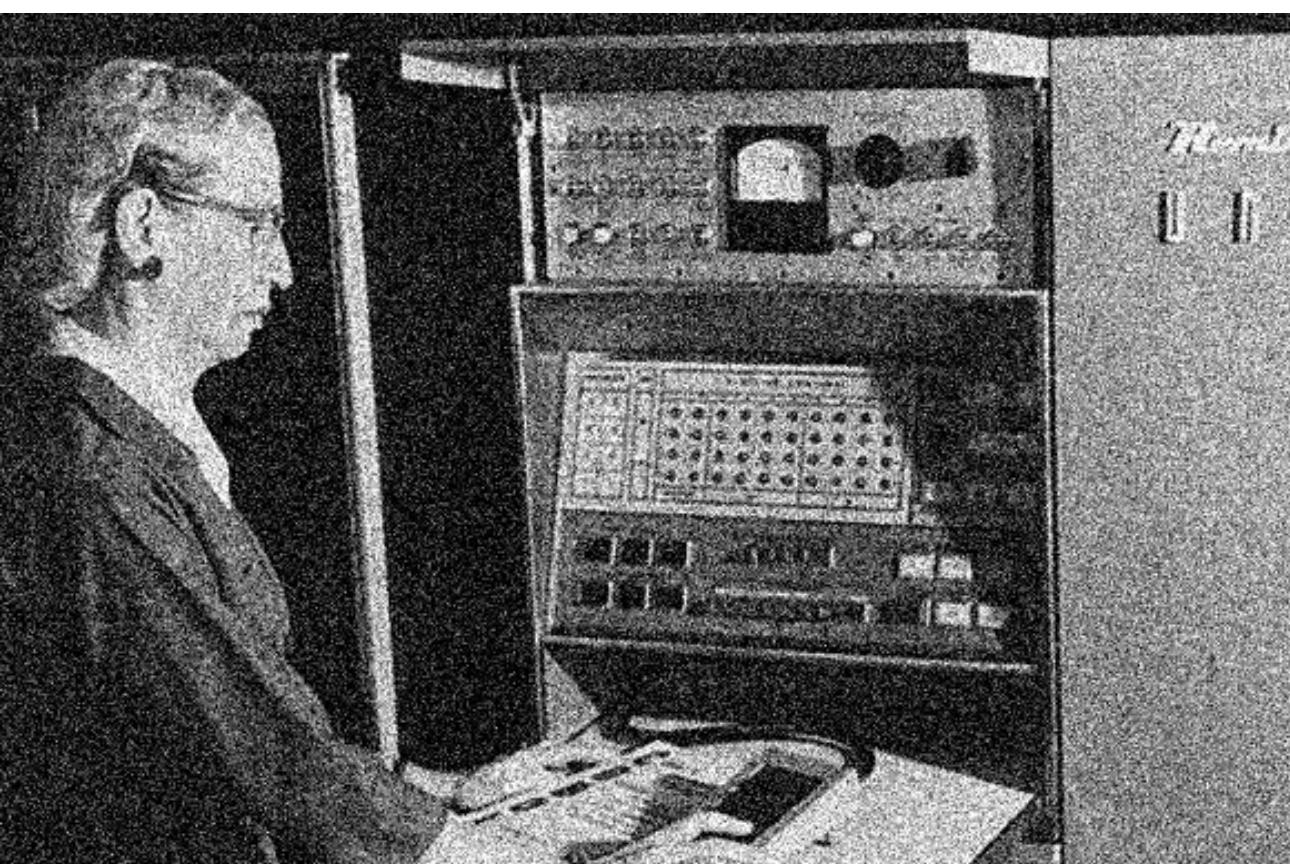
Noisy Input



I_x via $[-1,0,1]$



Zoom



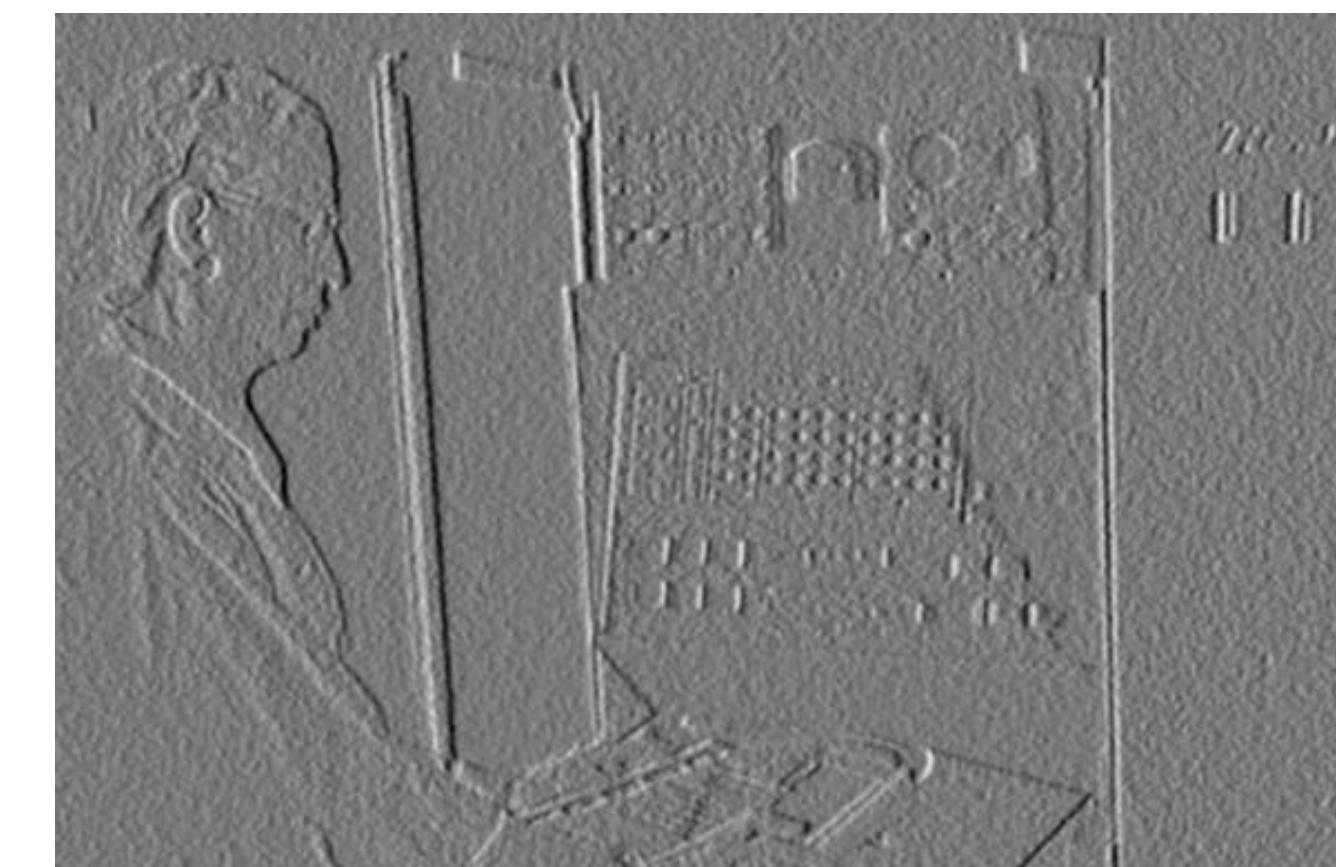
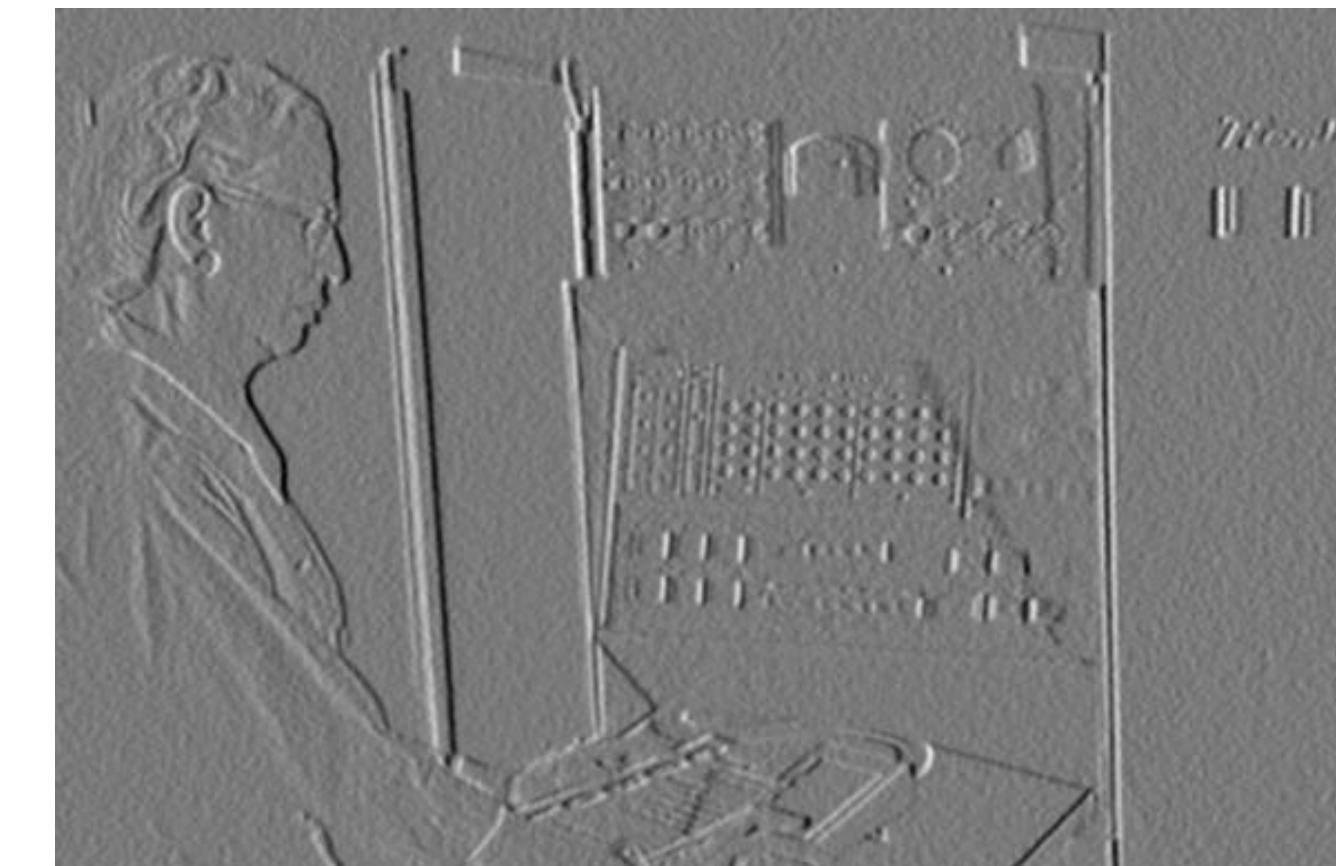
Source: D. Fouhey

Noise + Smoothing

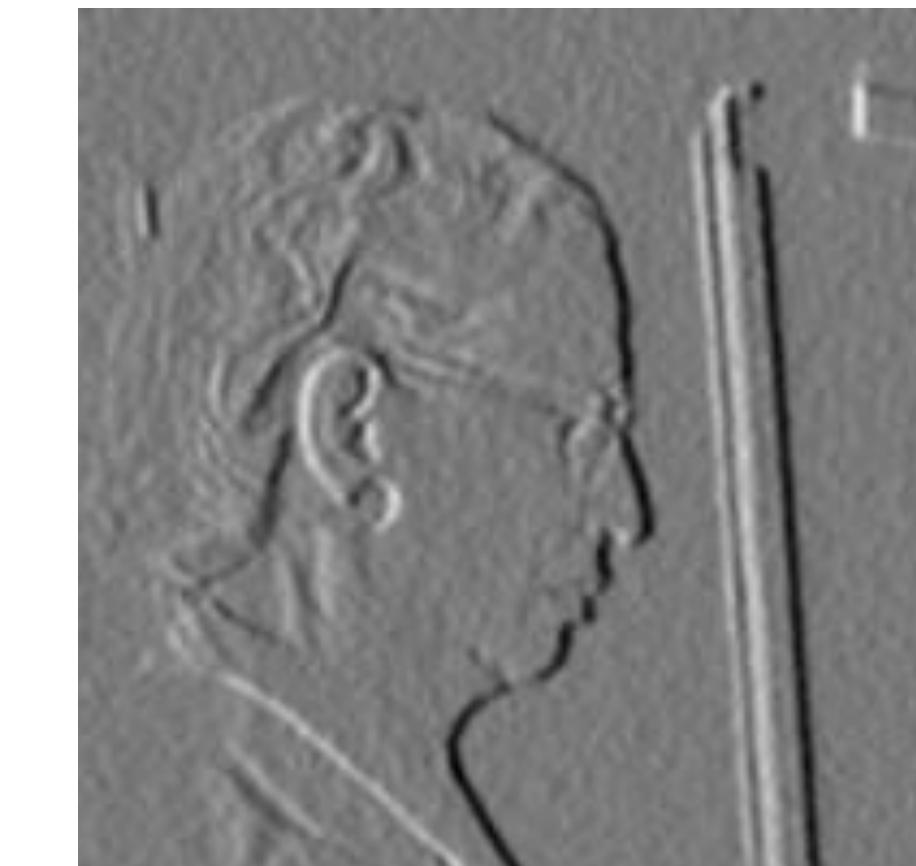
Smoothed Input



\mathbf{I}_x via $[-1,0,1]$

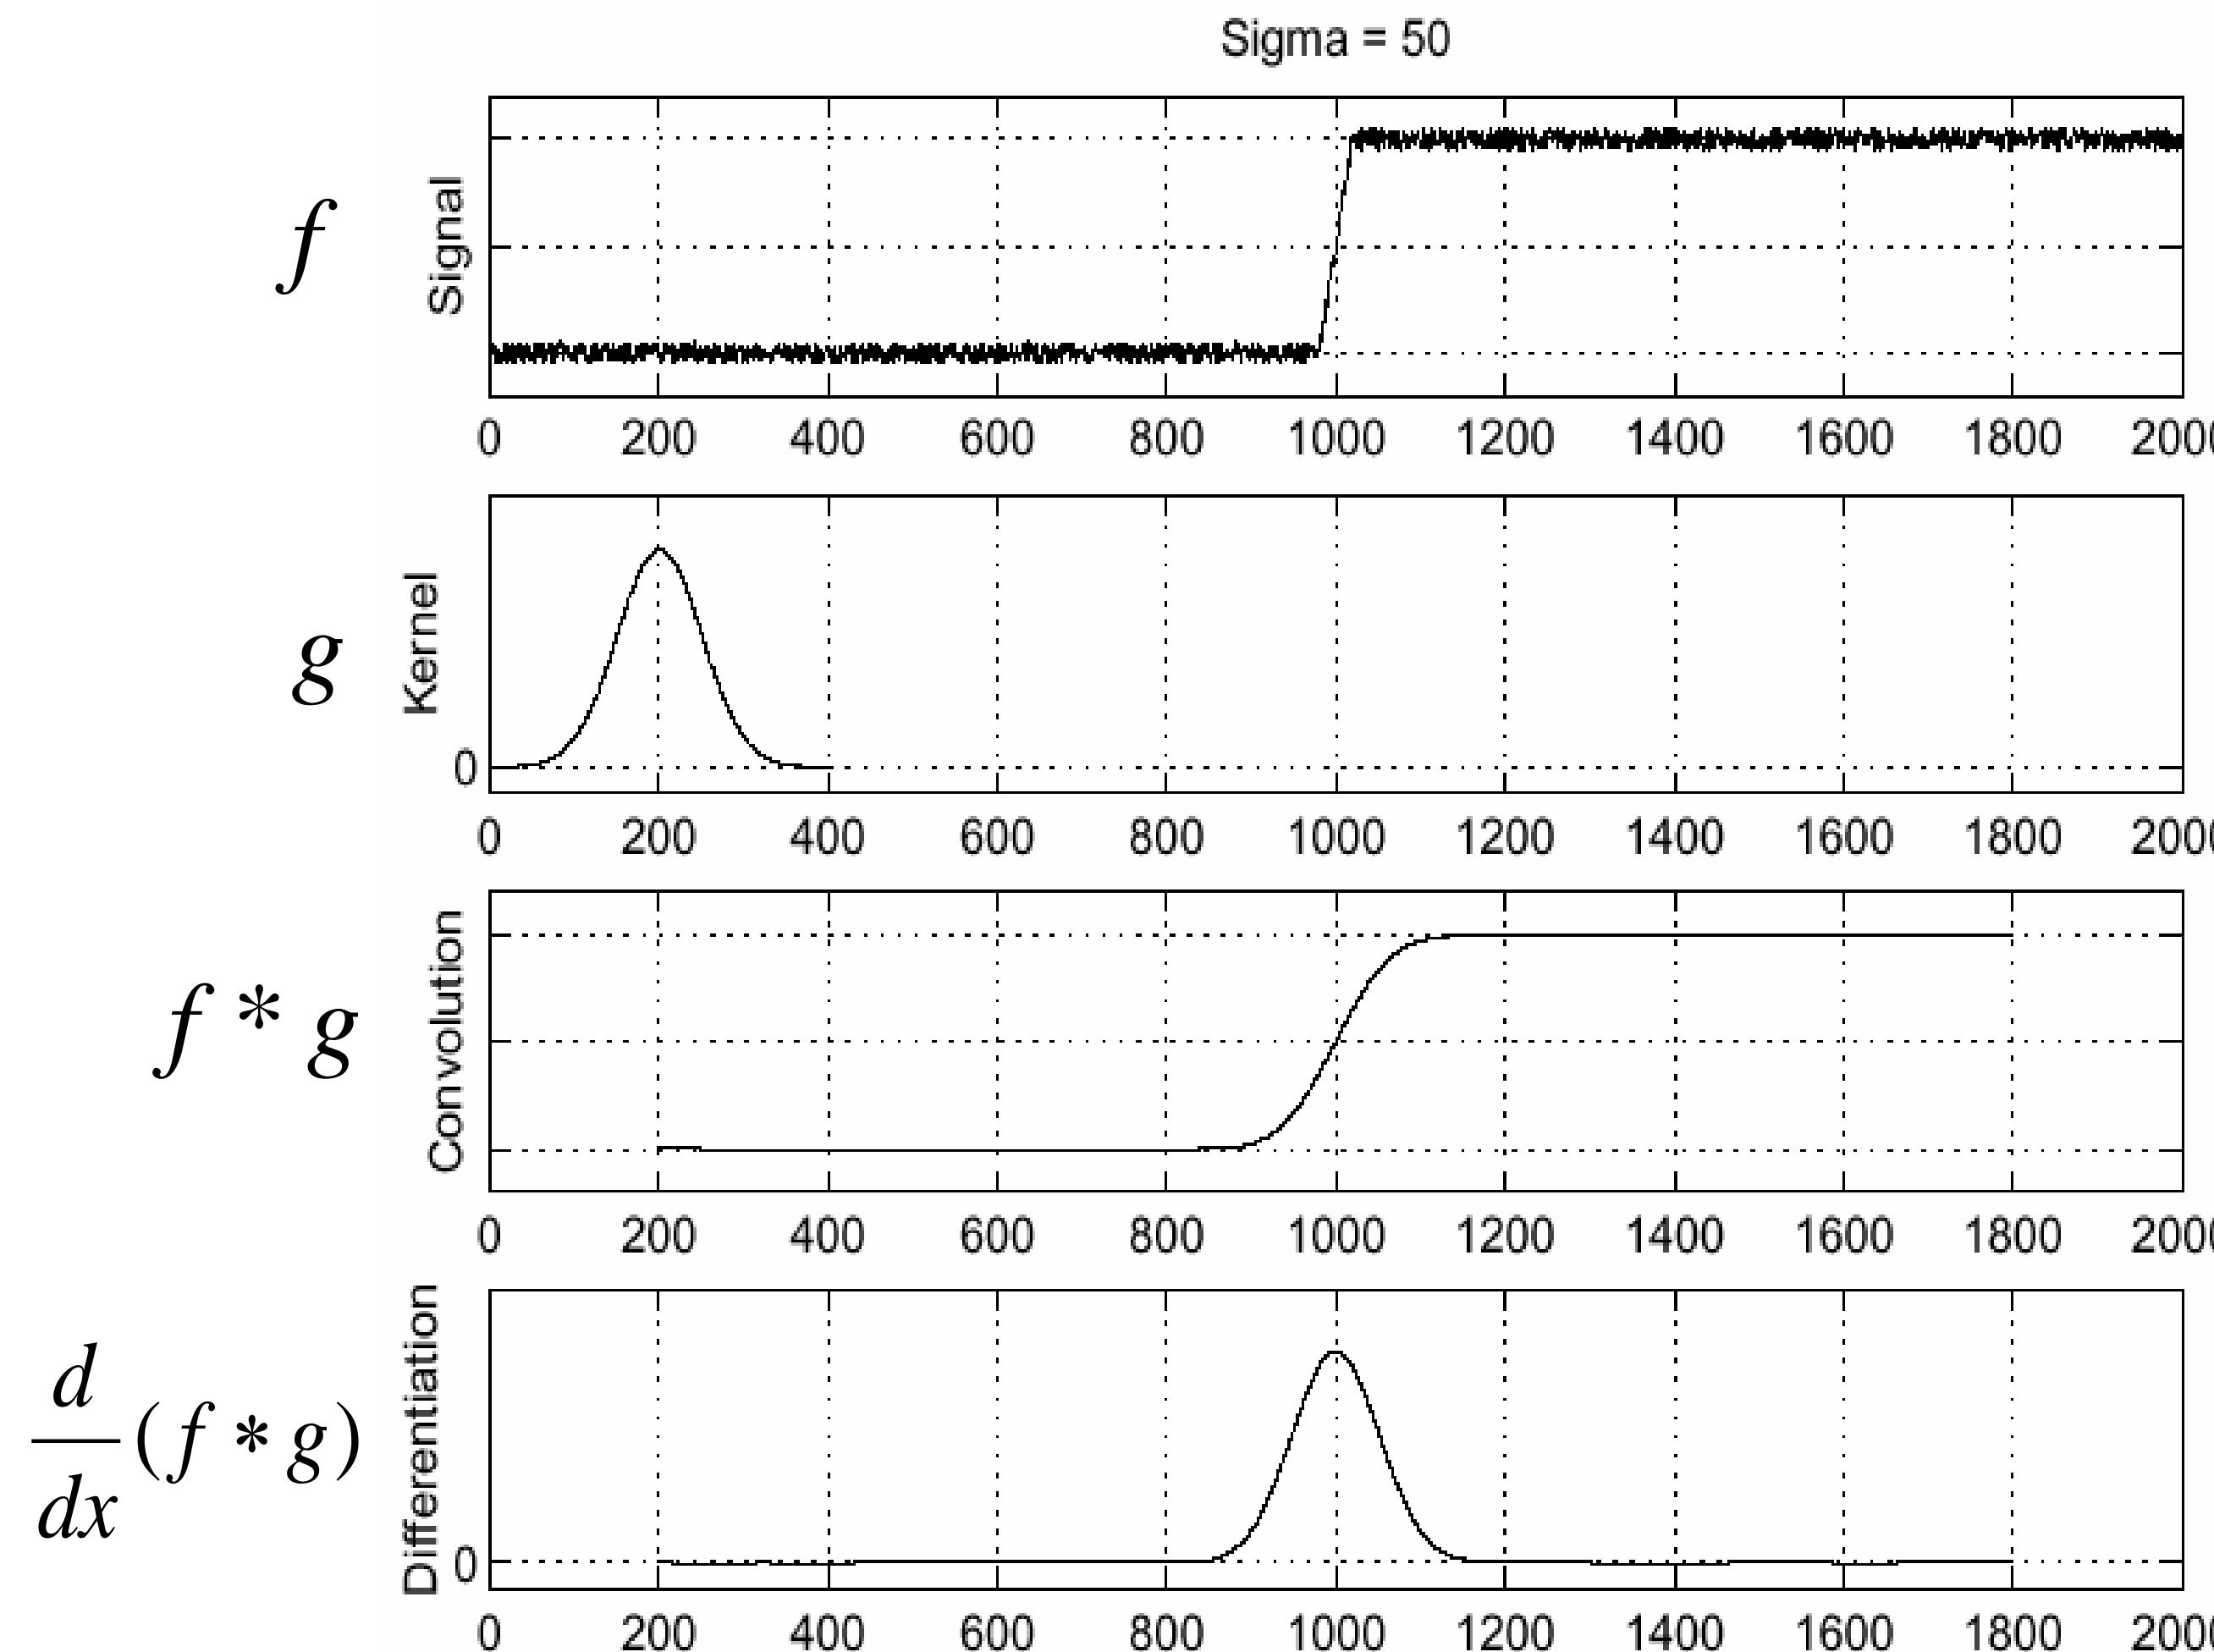


Zoom



Source: D. Fouhey

How many convolutions here?

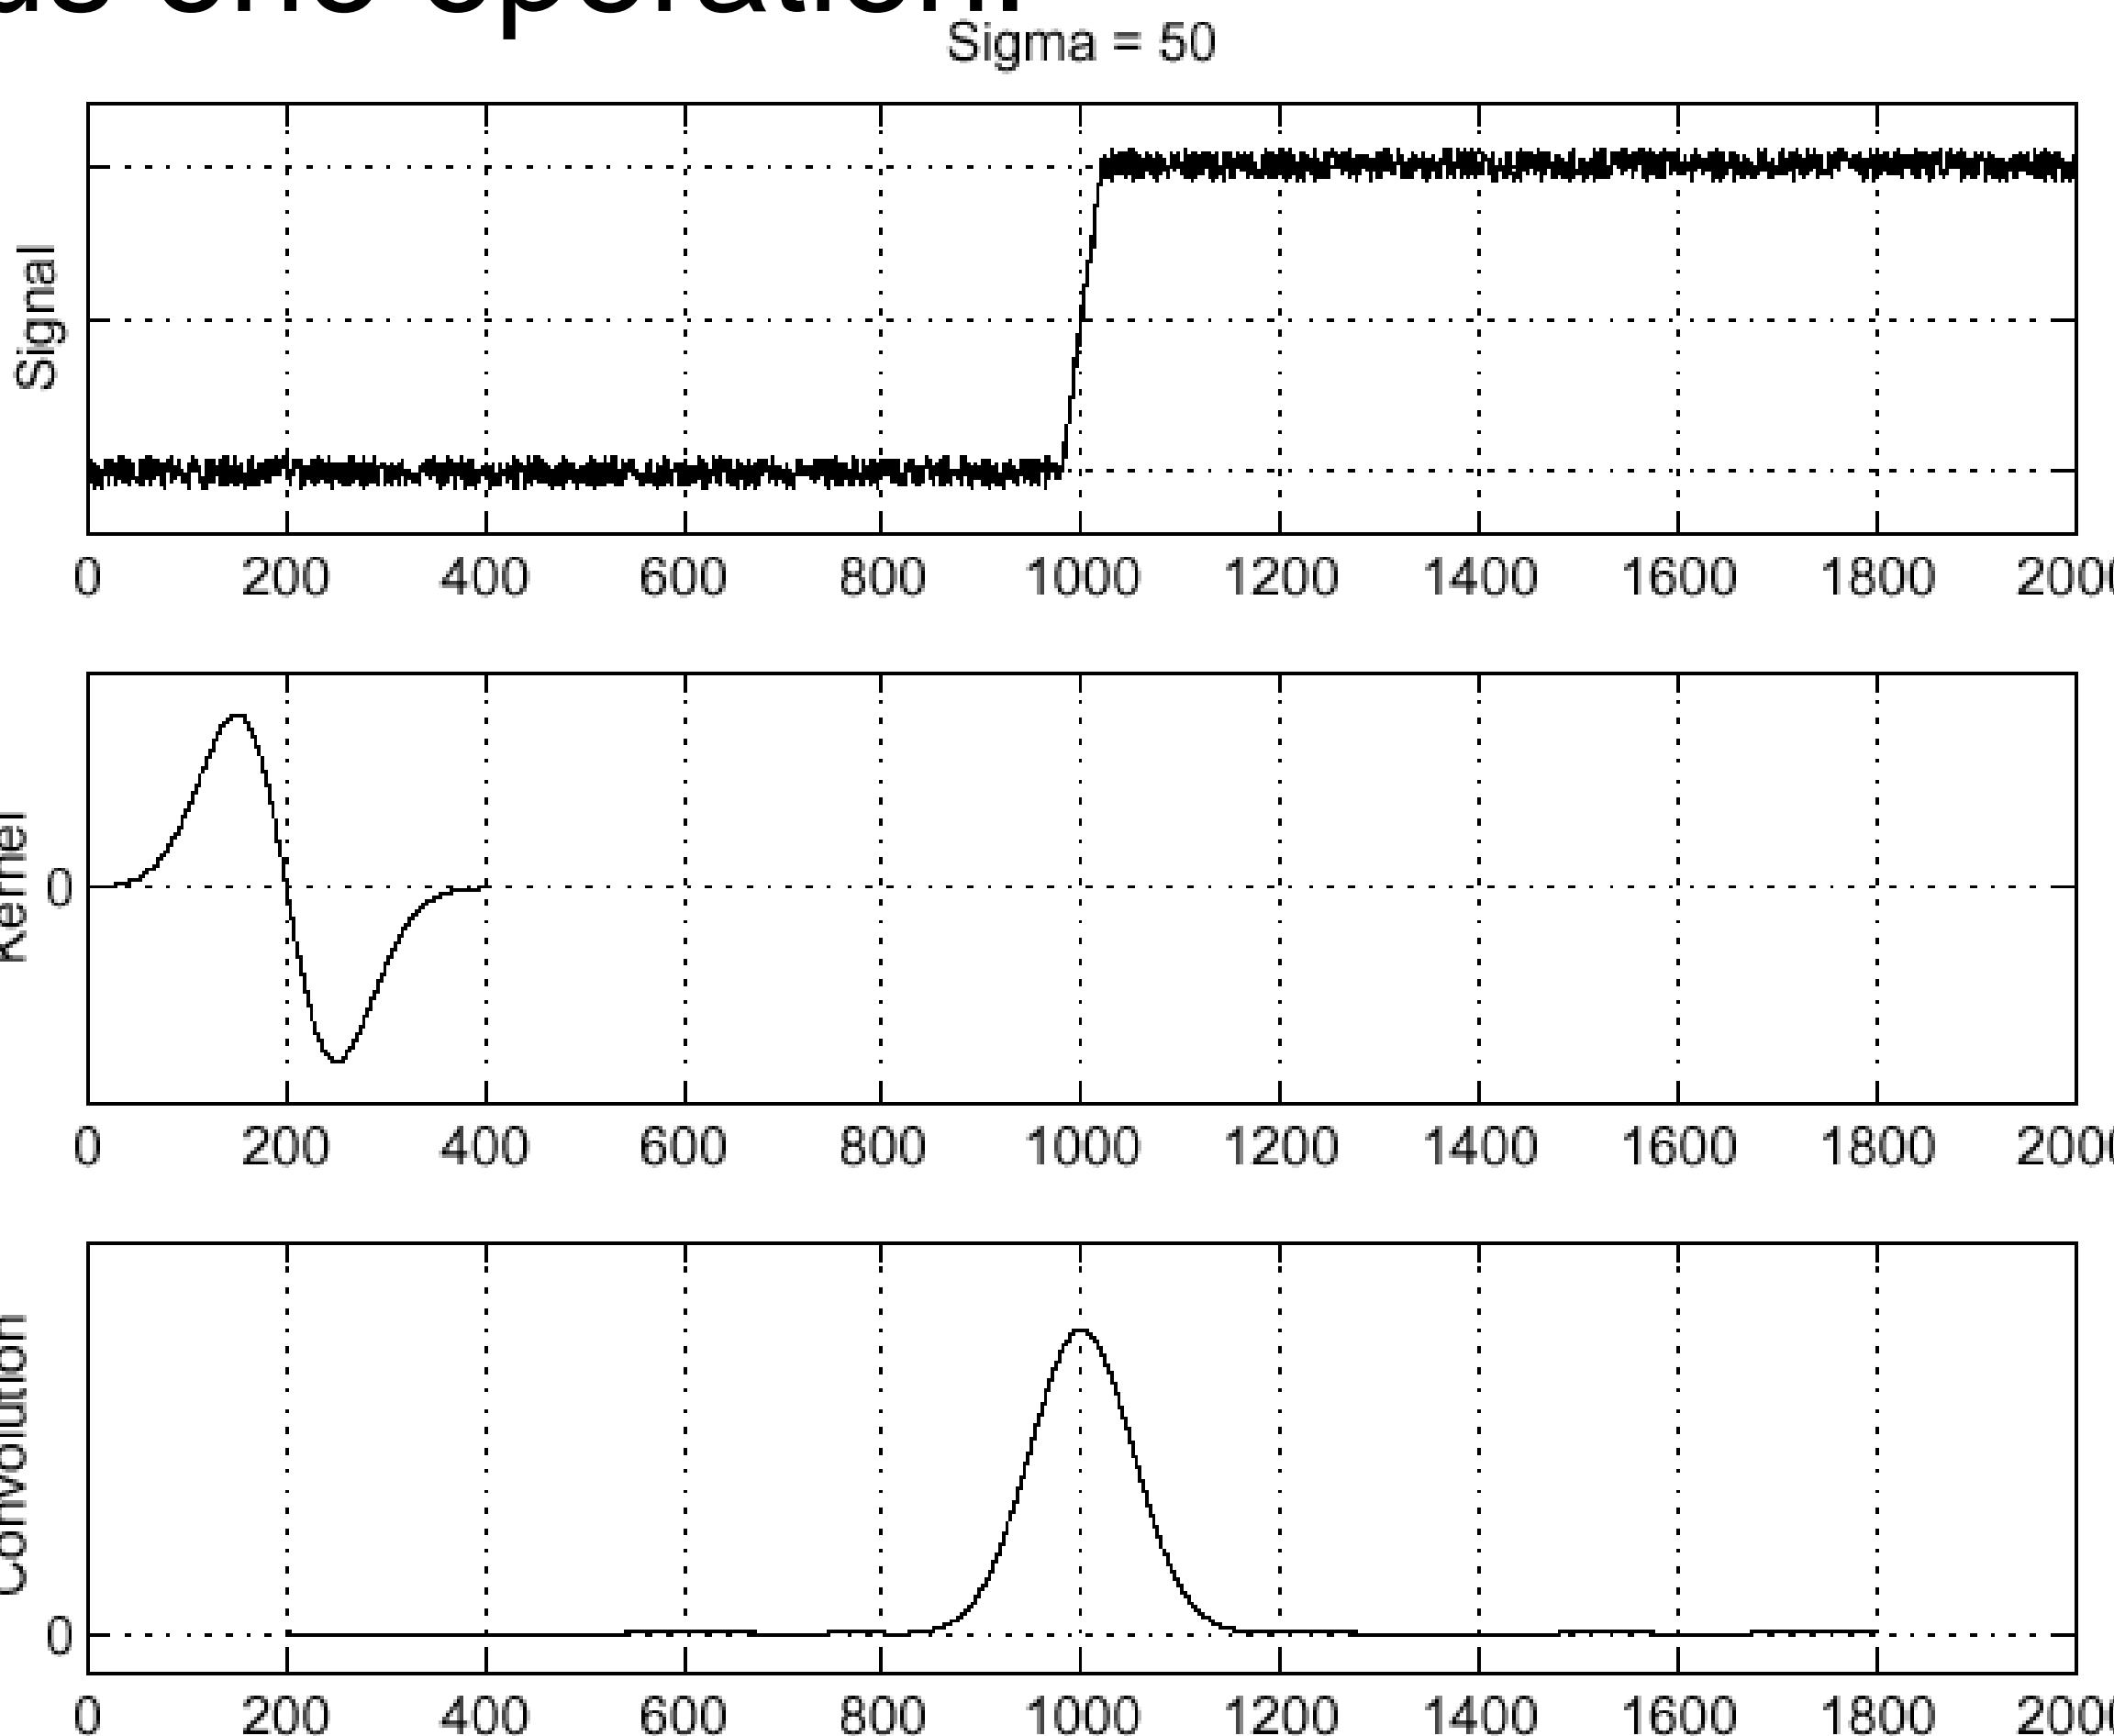


can we reduce this?

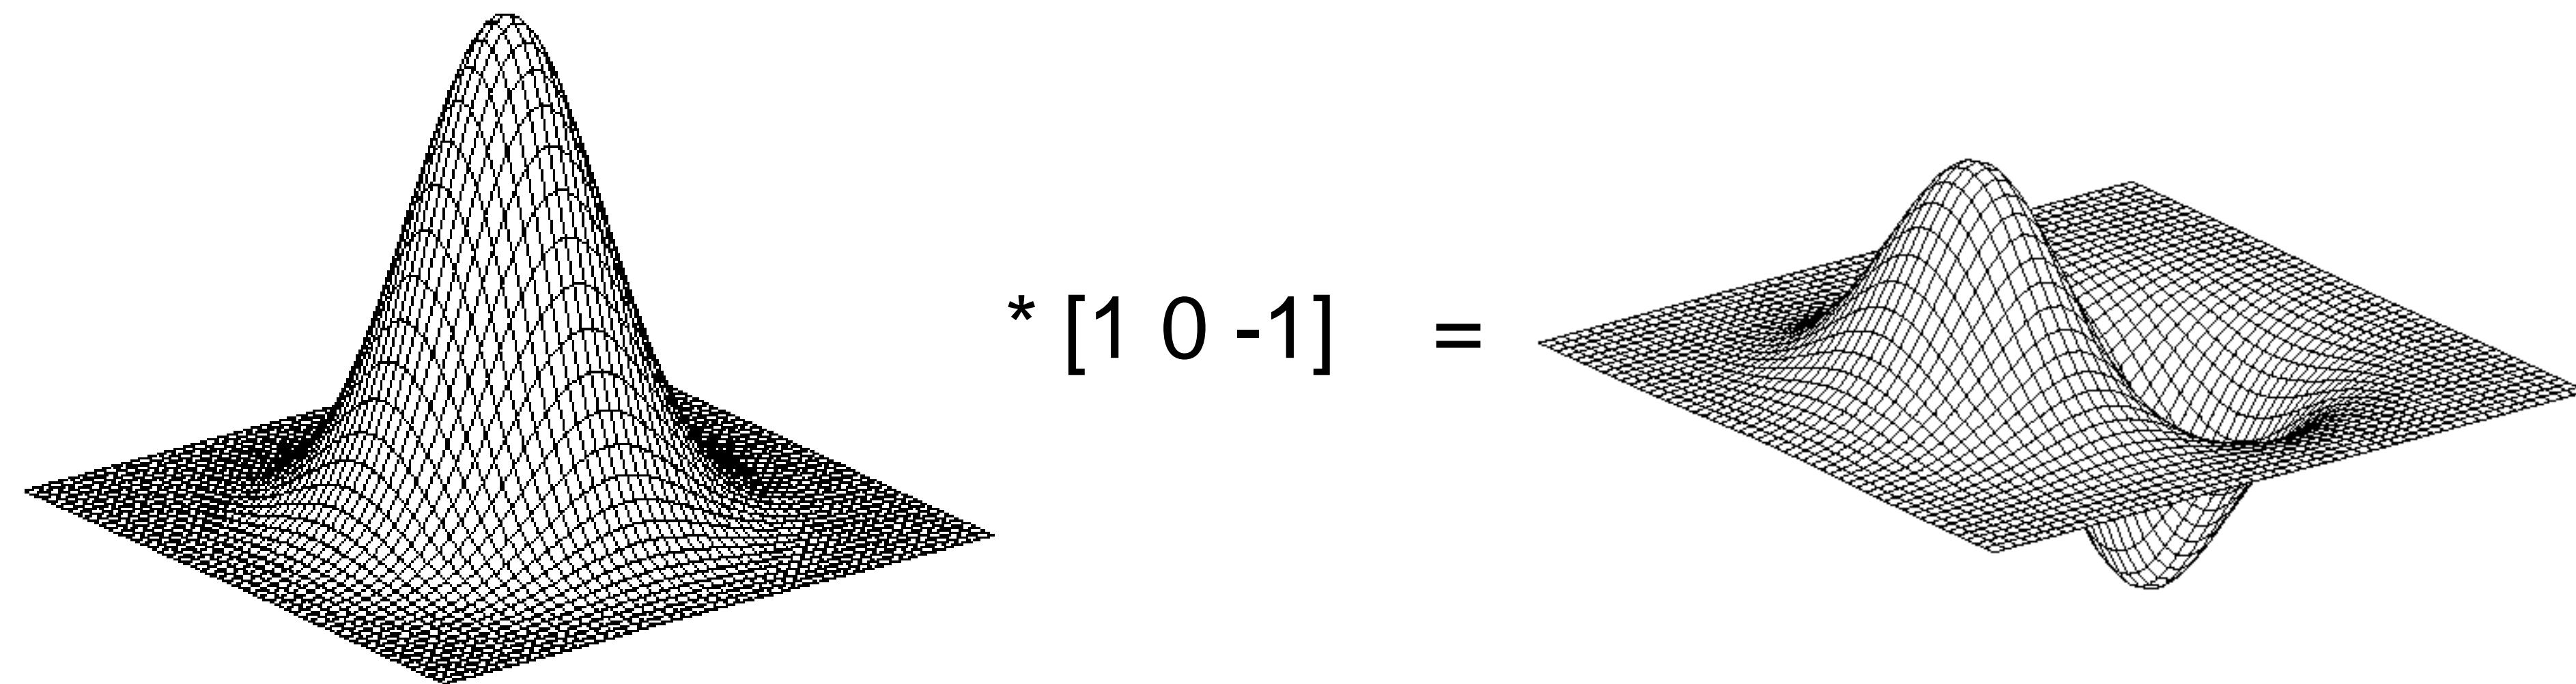
Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

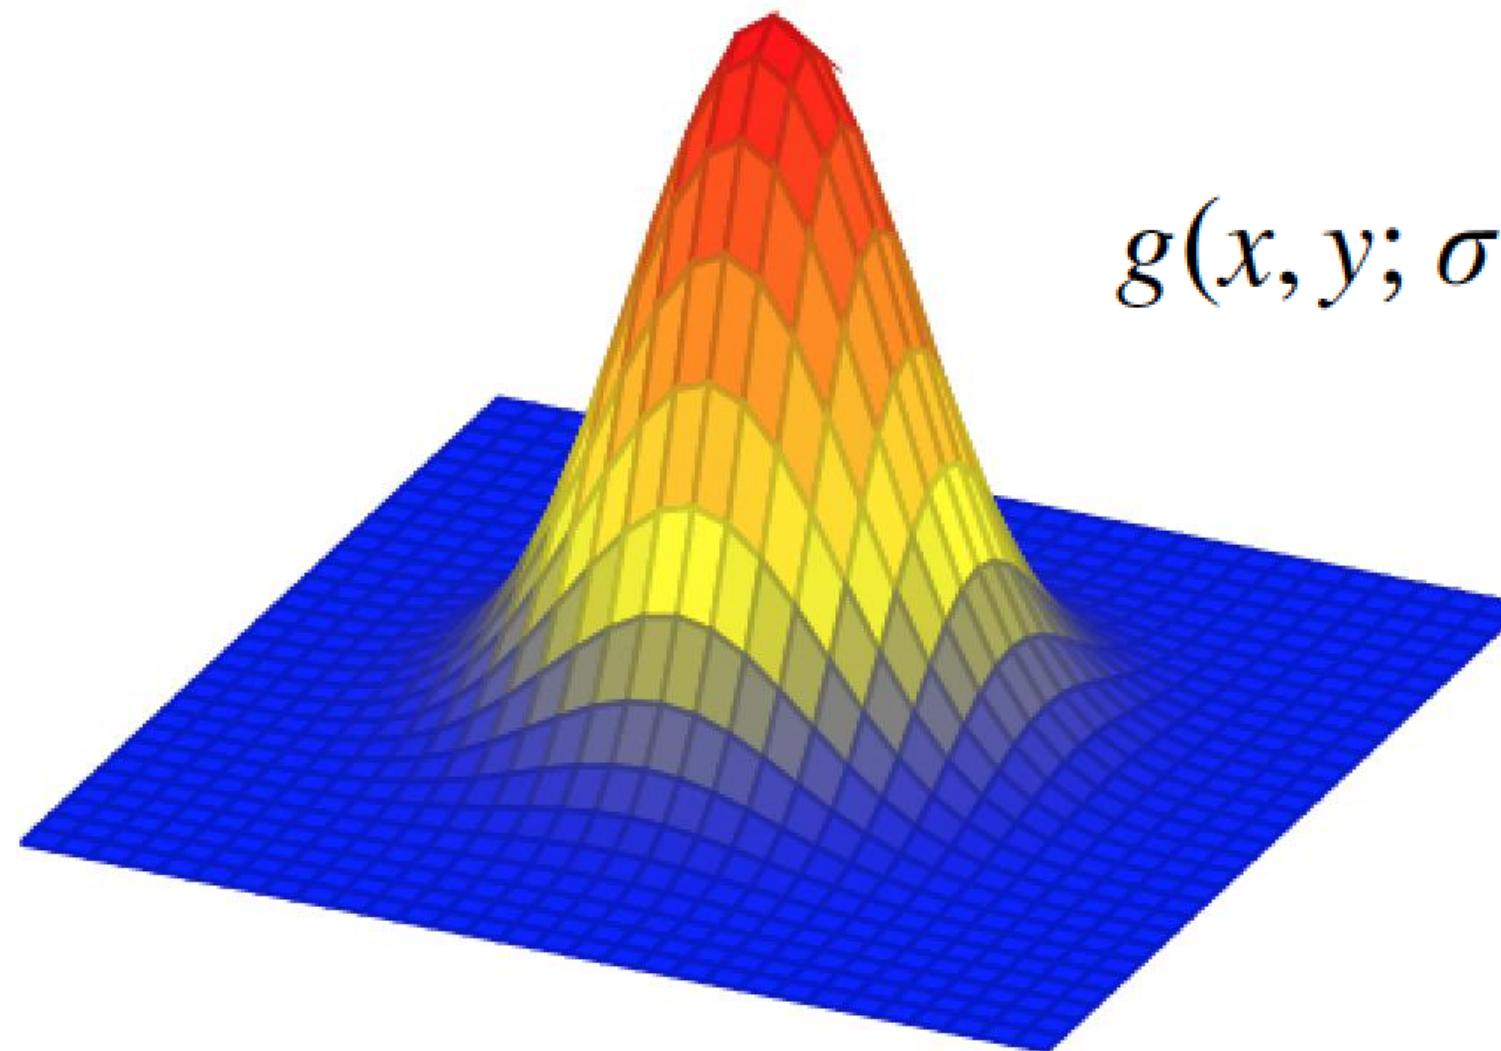
This saves us one operation:



Derivative of Gaussian filter



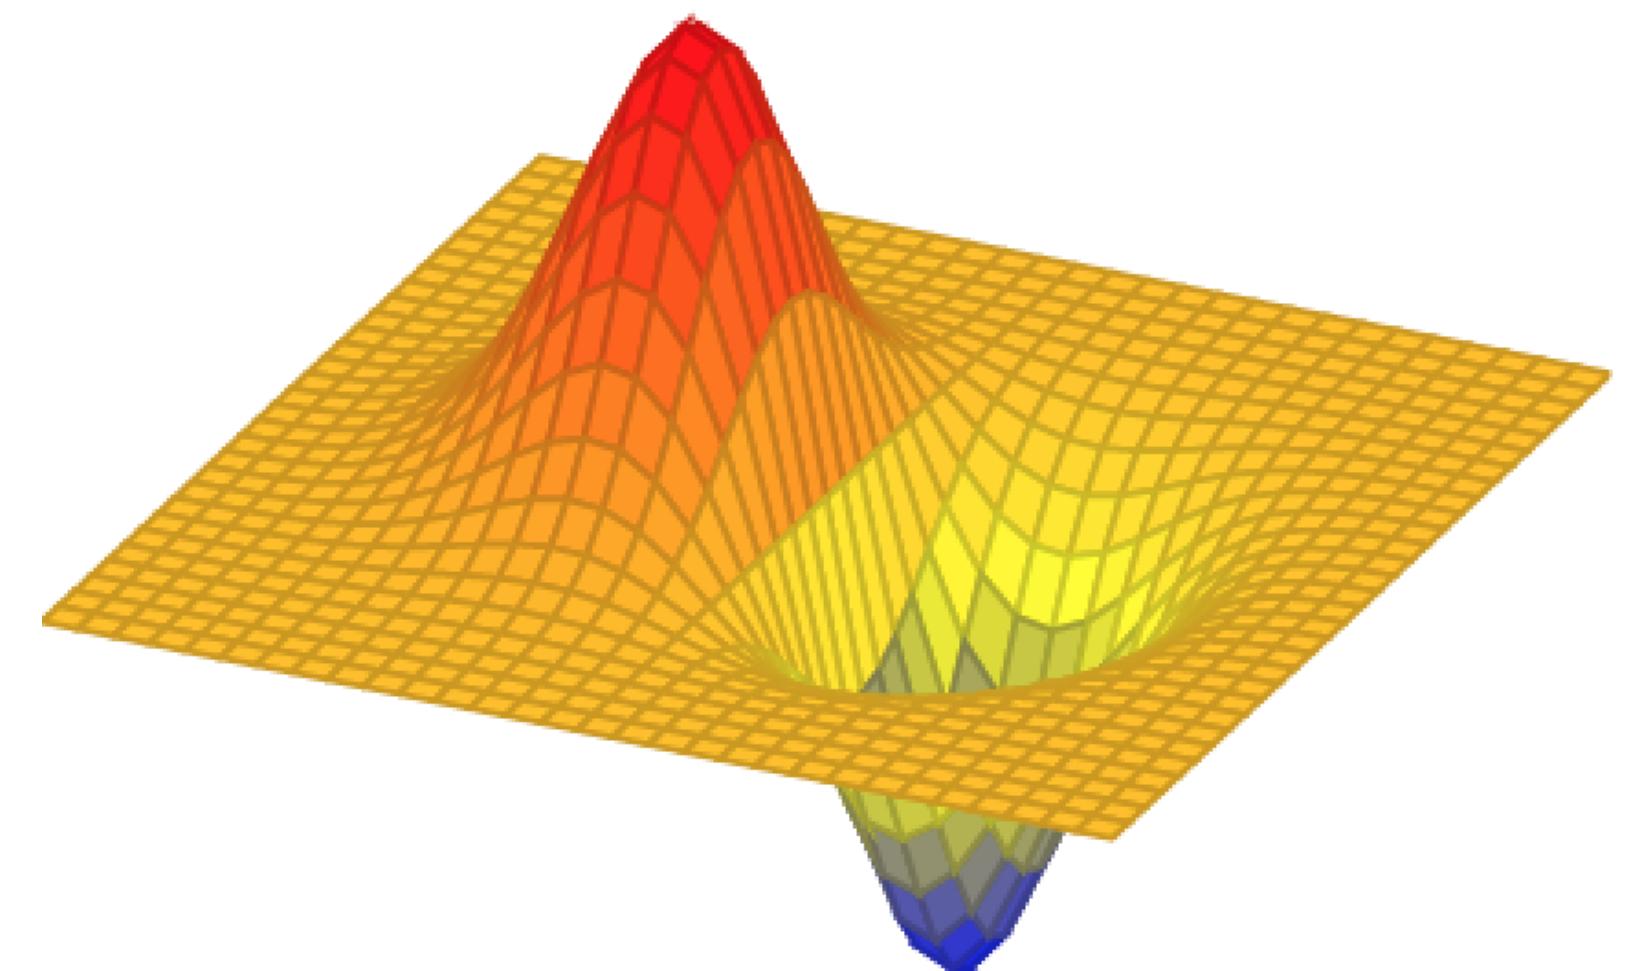
Gaussian derivatives



$$g(x, y; \sigma) = \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2}$$

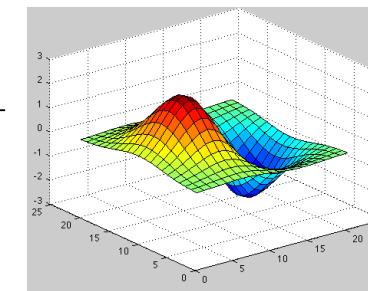
The continuous derivative is:

$$\begin{aligned} g_x(x, y; \sigma) &= \frac{\partial g(x, y; \sigma)}{\partial x} = \\ &= \frac{-x}{2\pi\sigma^4} \exp -\frac{x^2 + y^2}{2\sigma^2} \\ &= \frac{-x}{\sigma^2} g(x, y; \sigma) \end{aligned}$$

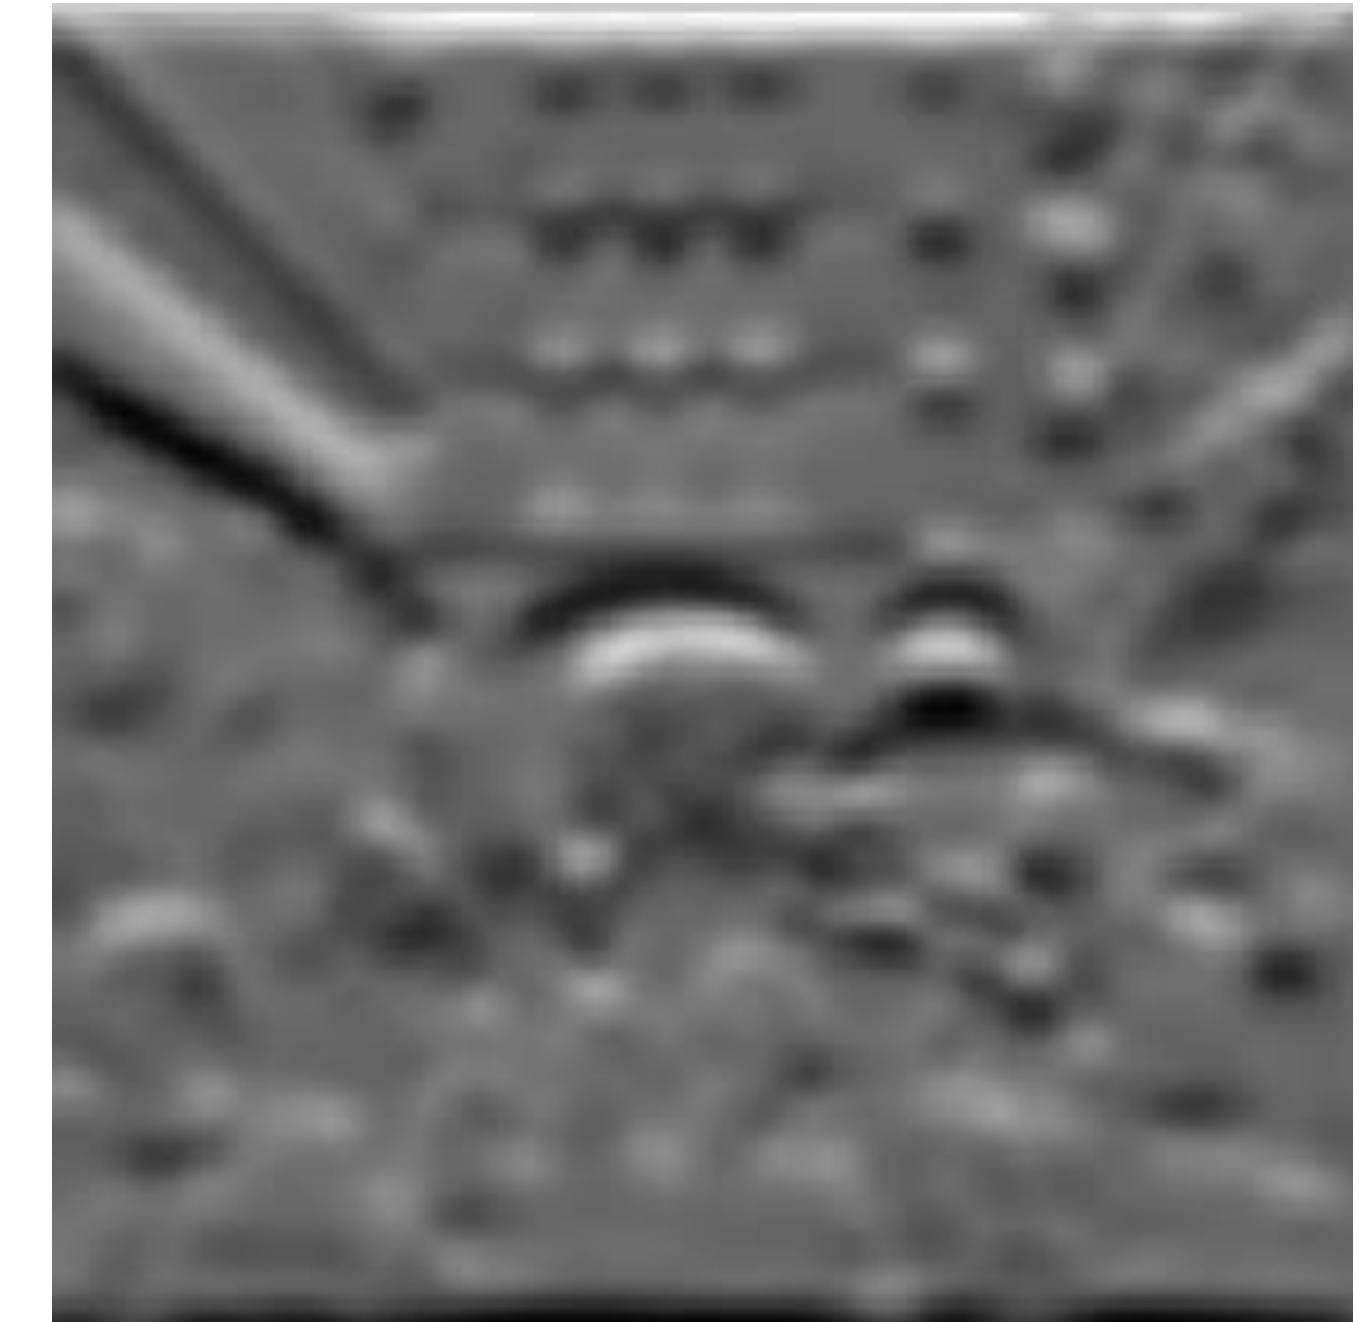
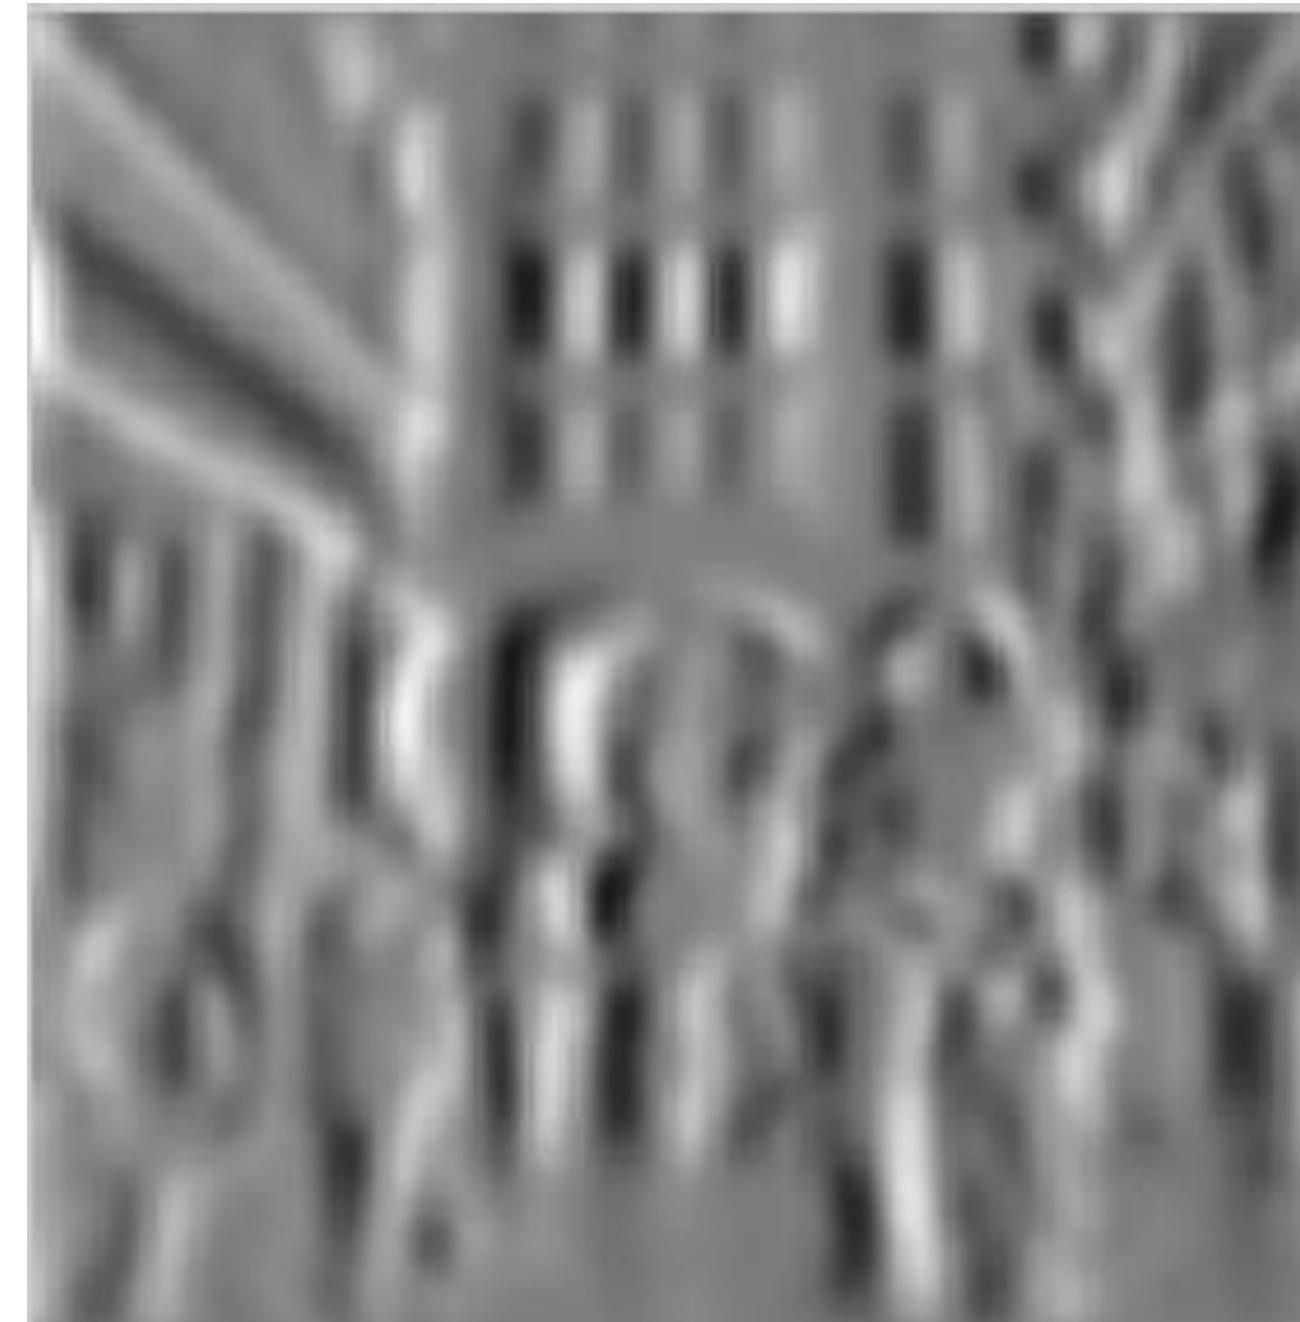
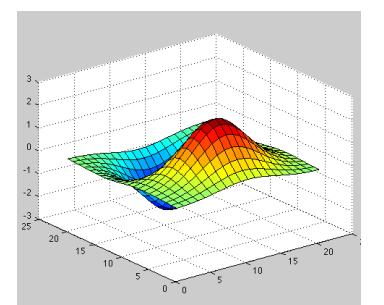


Orientation

$$g_x(x,y) = \frac{\partial g(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



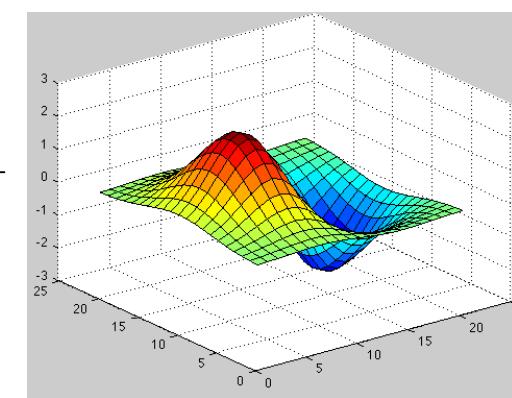
$$g_y(x,y) = \frac{\partial g(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



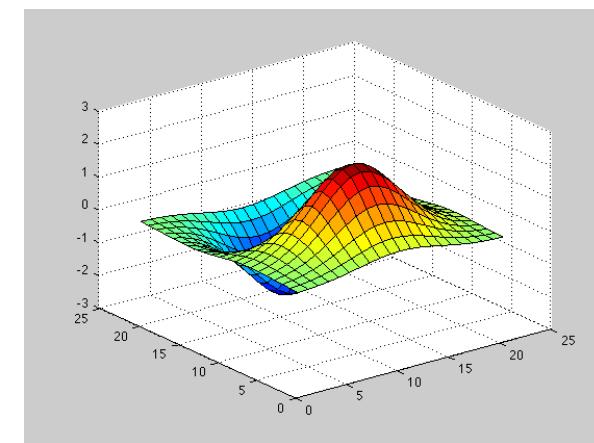
What about other orientations not axis aligned?

Orientation

$$g_x(x,y) = \frac{\partial g(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



$$g_y(x,y) = \frac{\partial g(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



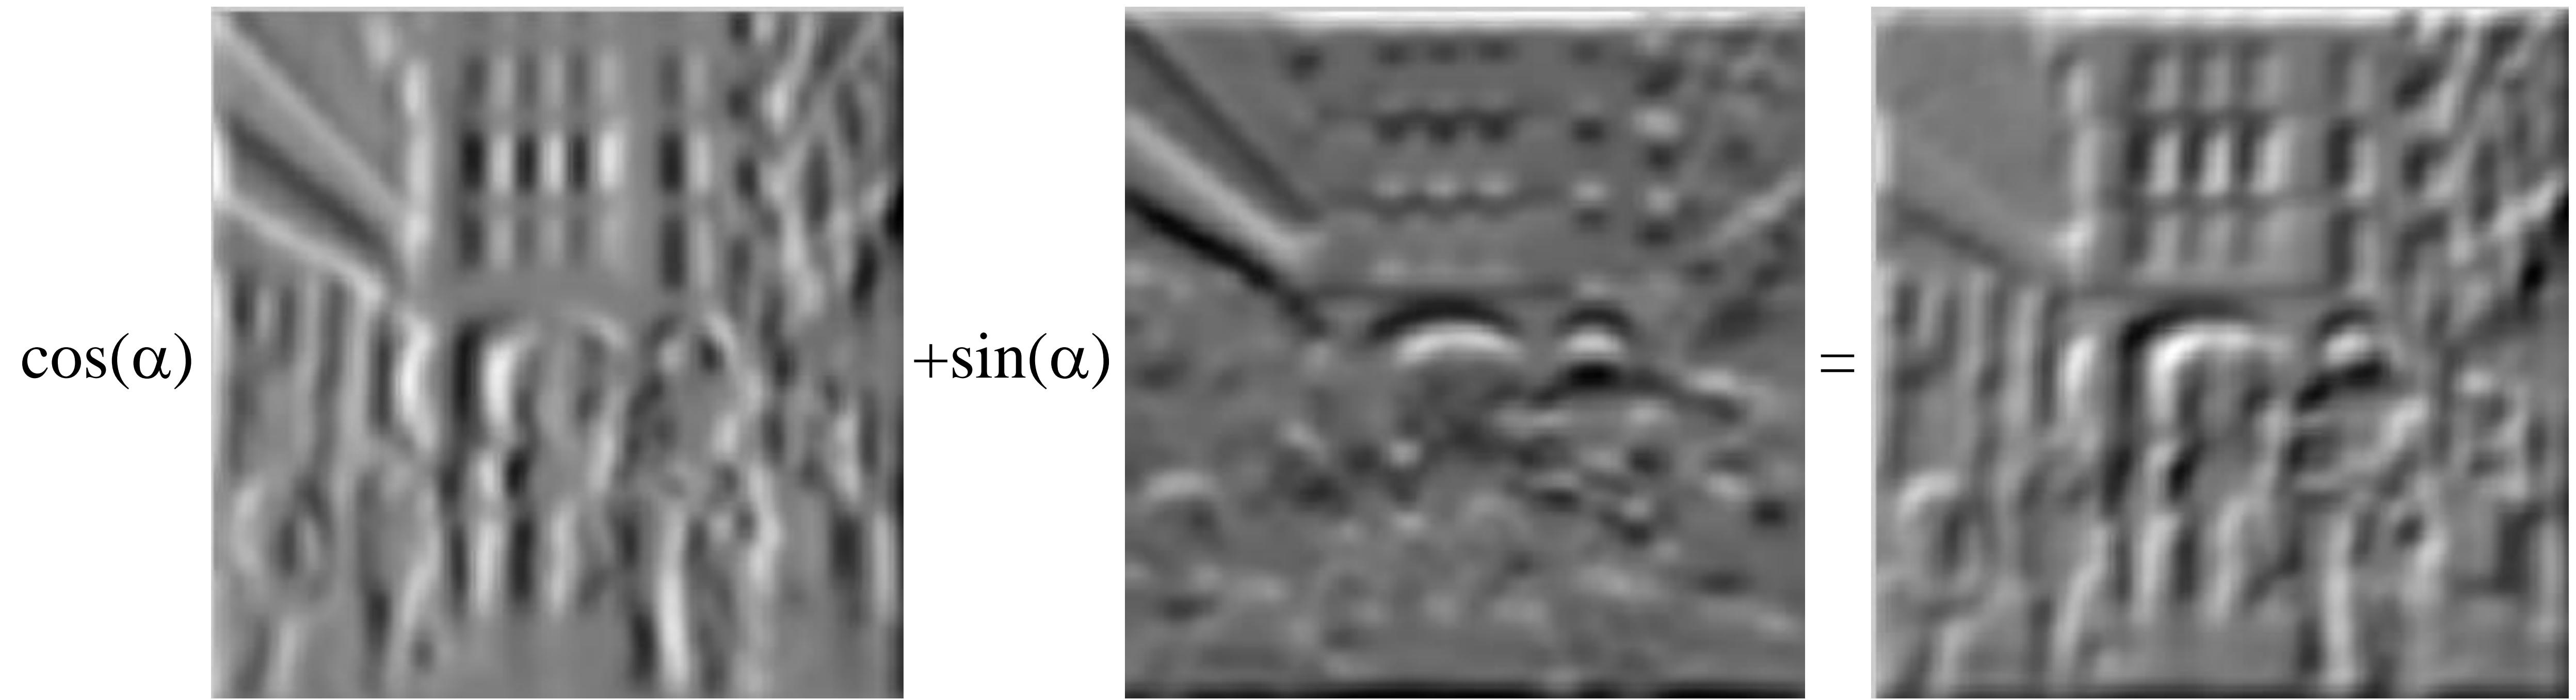
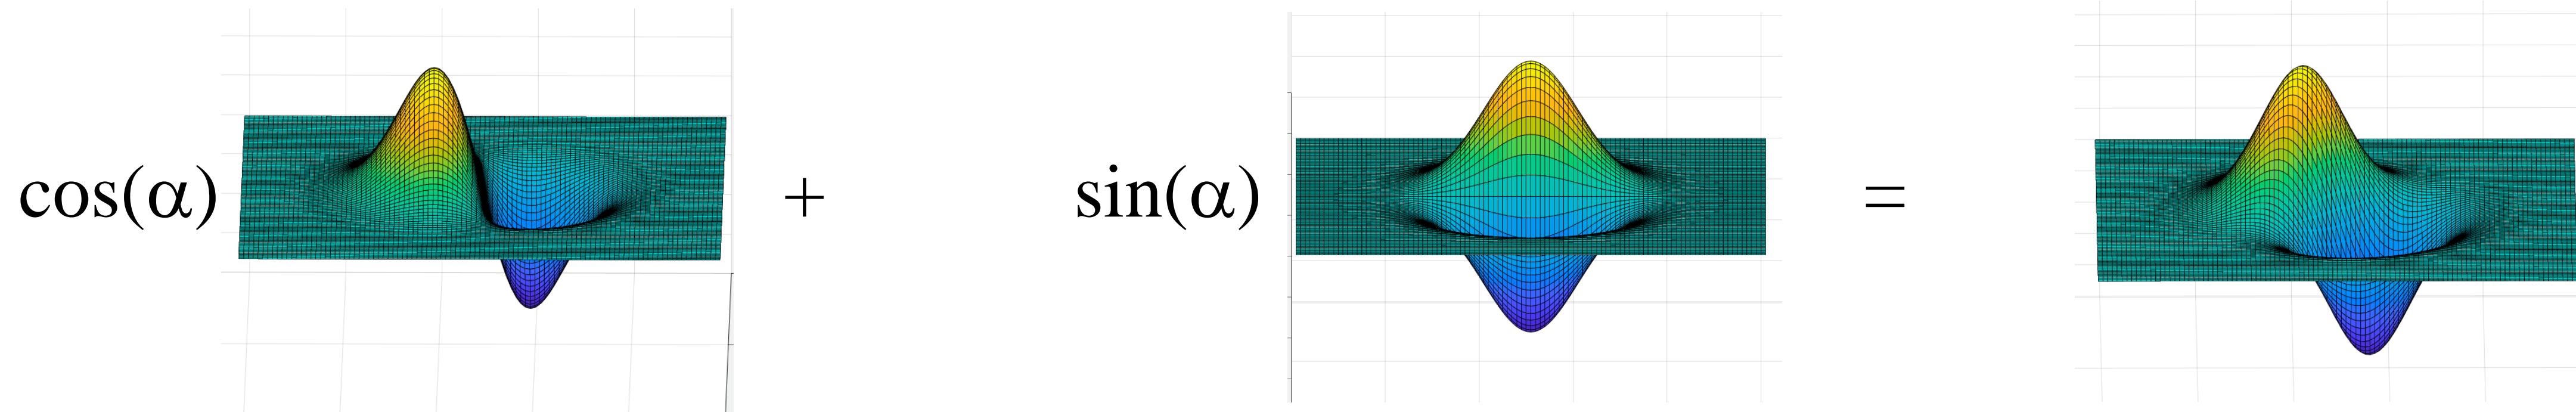
The smoothed directional gradient is a linear combination of two kernels

$$u^T \nabla g \otimes I = (\cos(\alpha)g_x(x,y) + \sin(\alpha)g_y(x,y)) \otimes I(x,y) =$$

Any orientation can be computed as a linear combination of two filtered images

$$= \cos(\alpha)g_x(x,y) \otimes I(x,y) + \sin(\alpha)g_y(x,y) \otimes I(x,y)$$

Example: “steering” to 45°

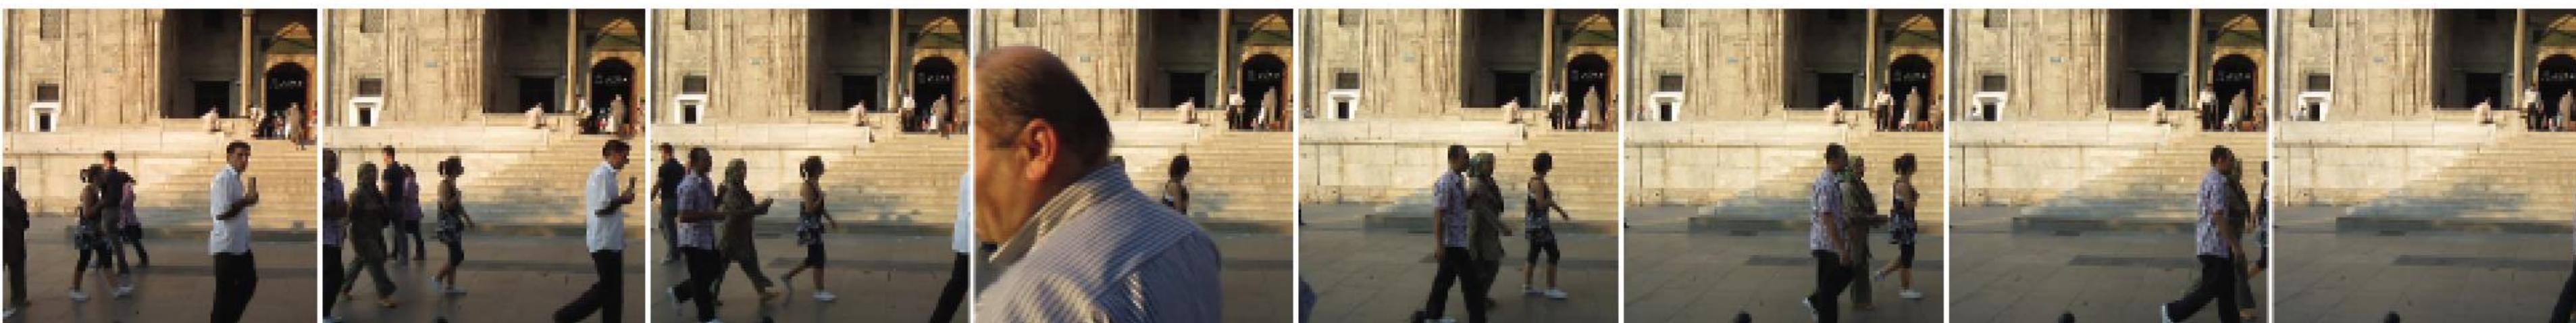


Temporal filtering

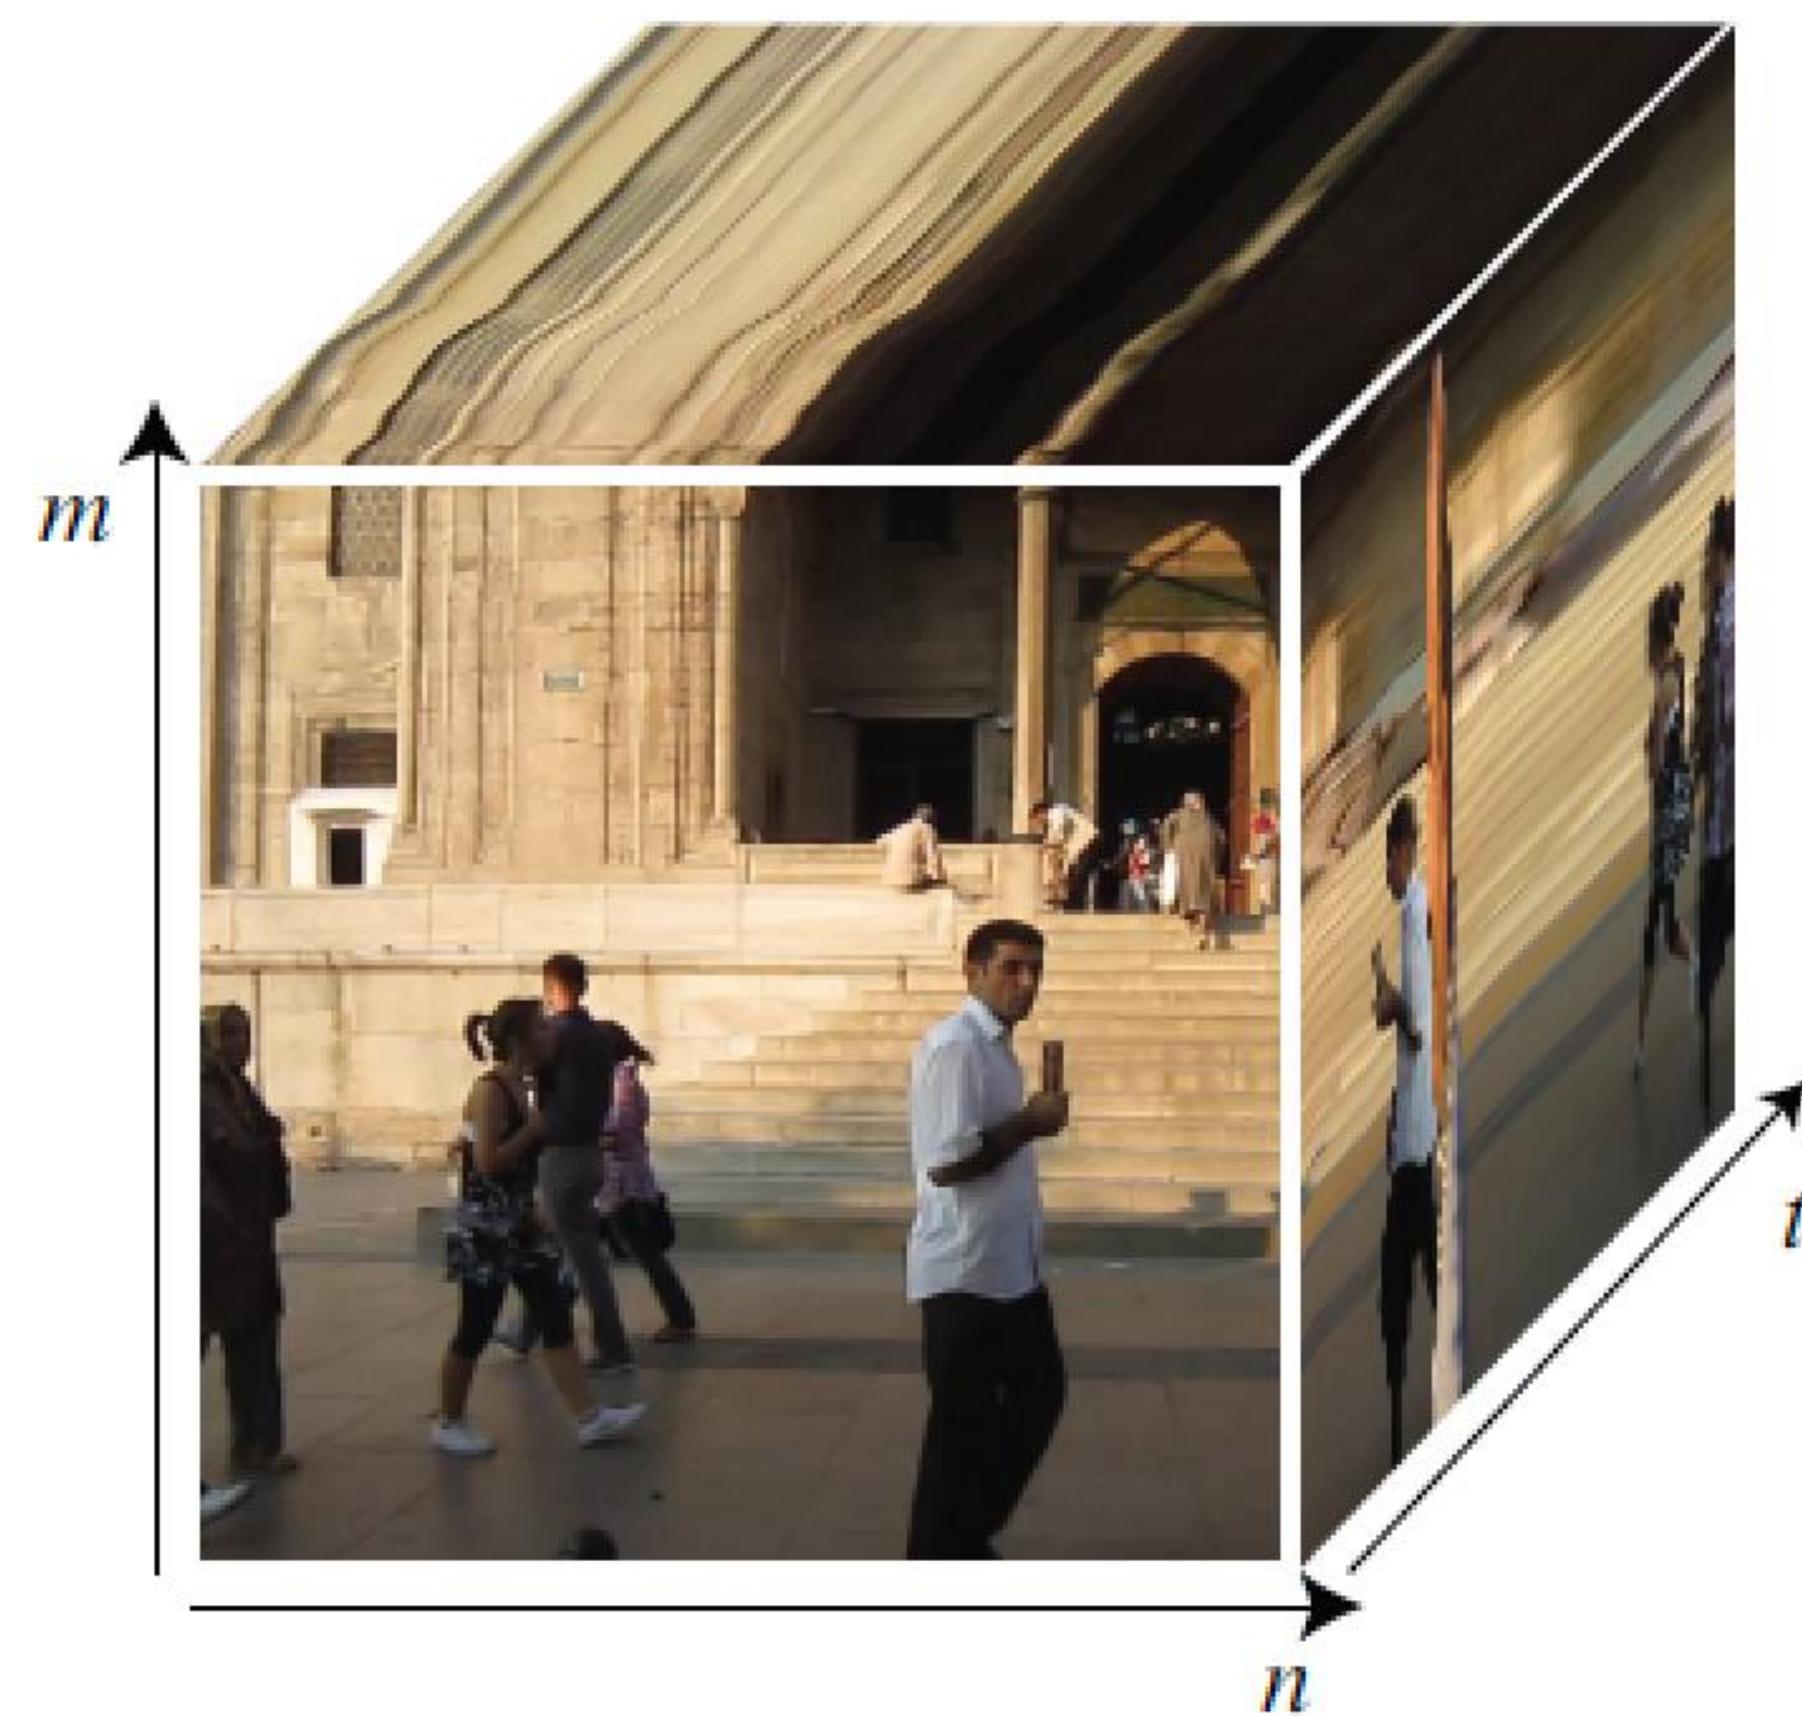


why filter videos over time?

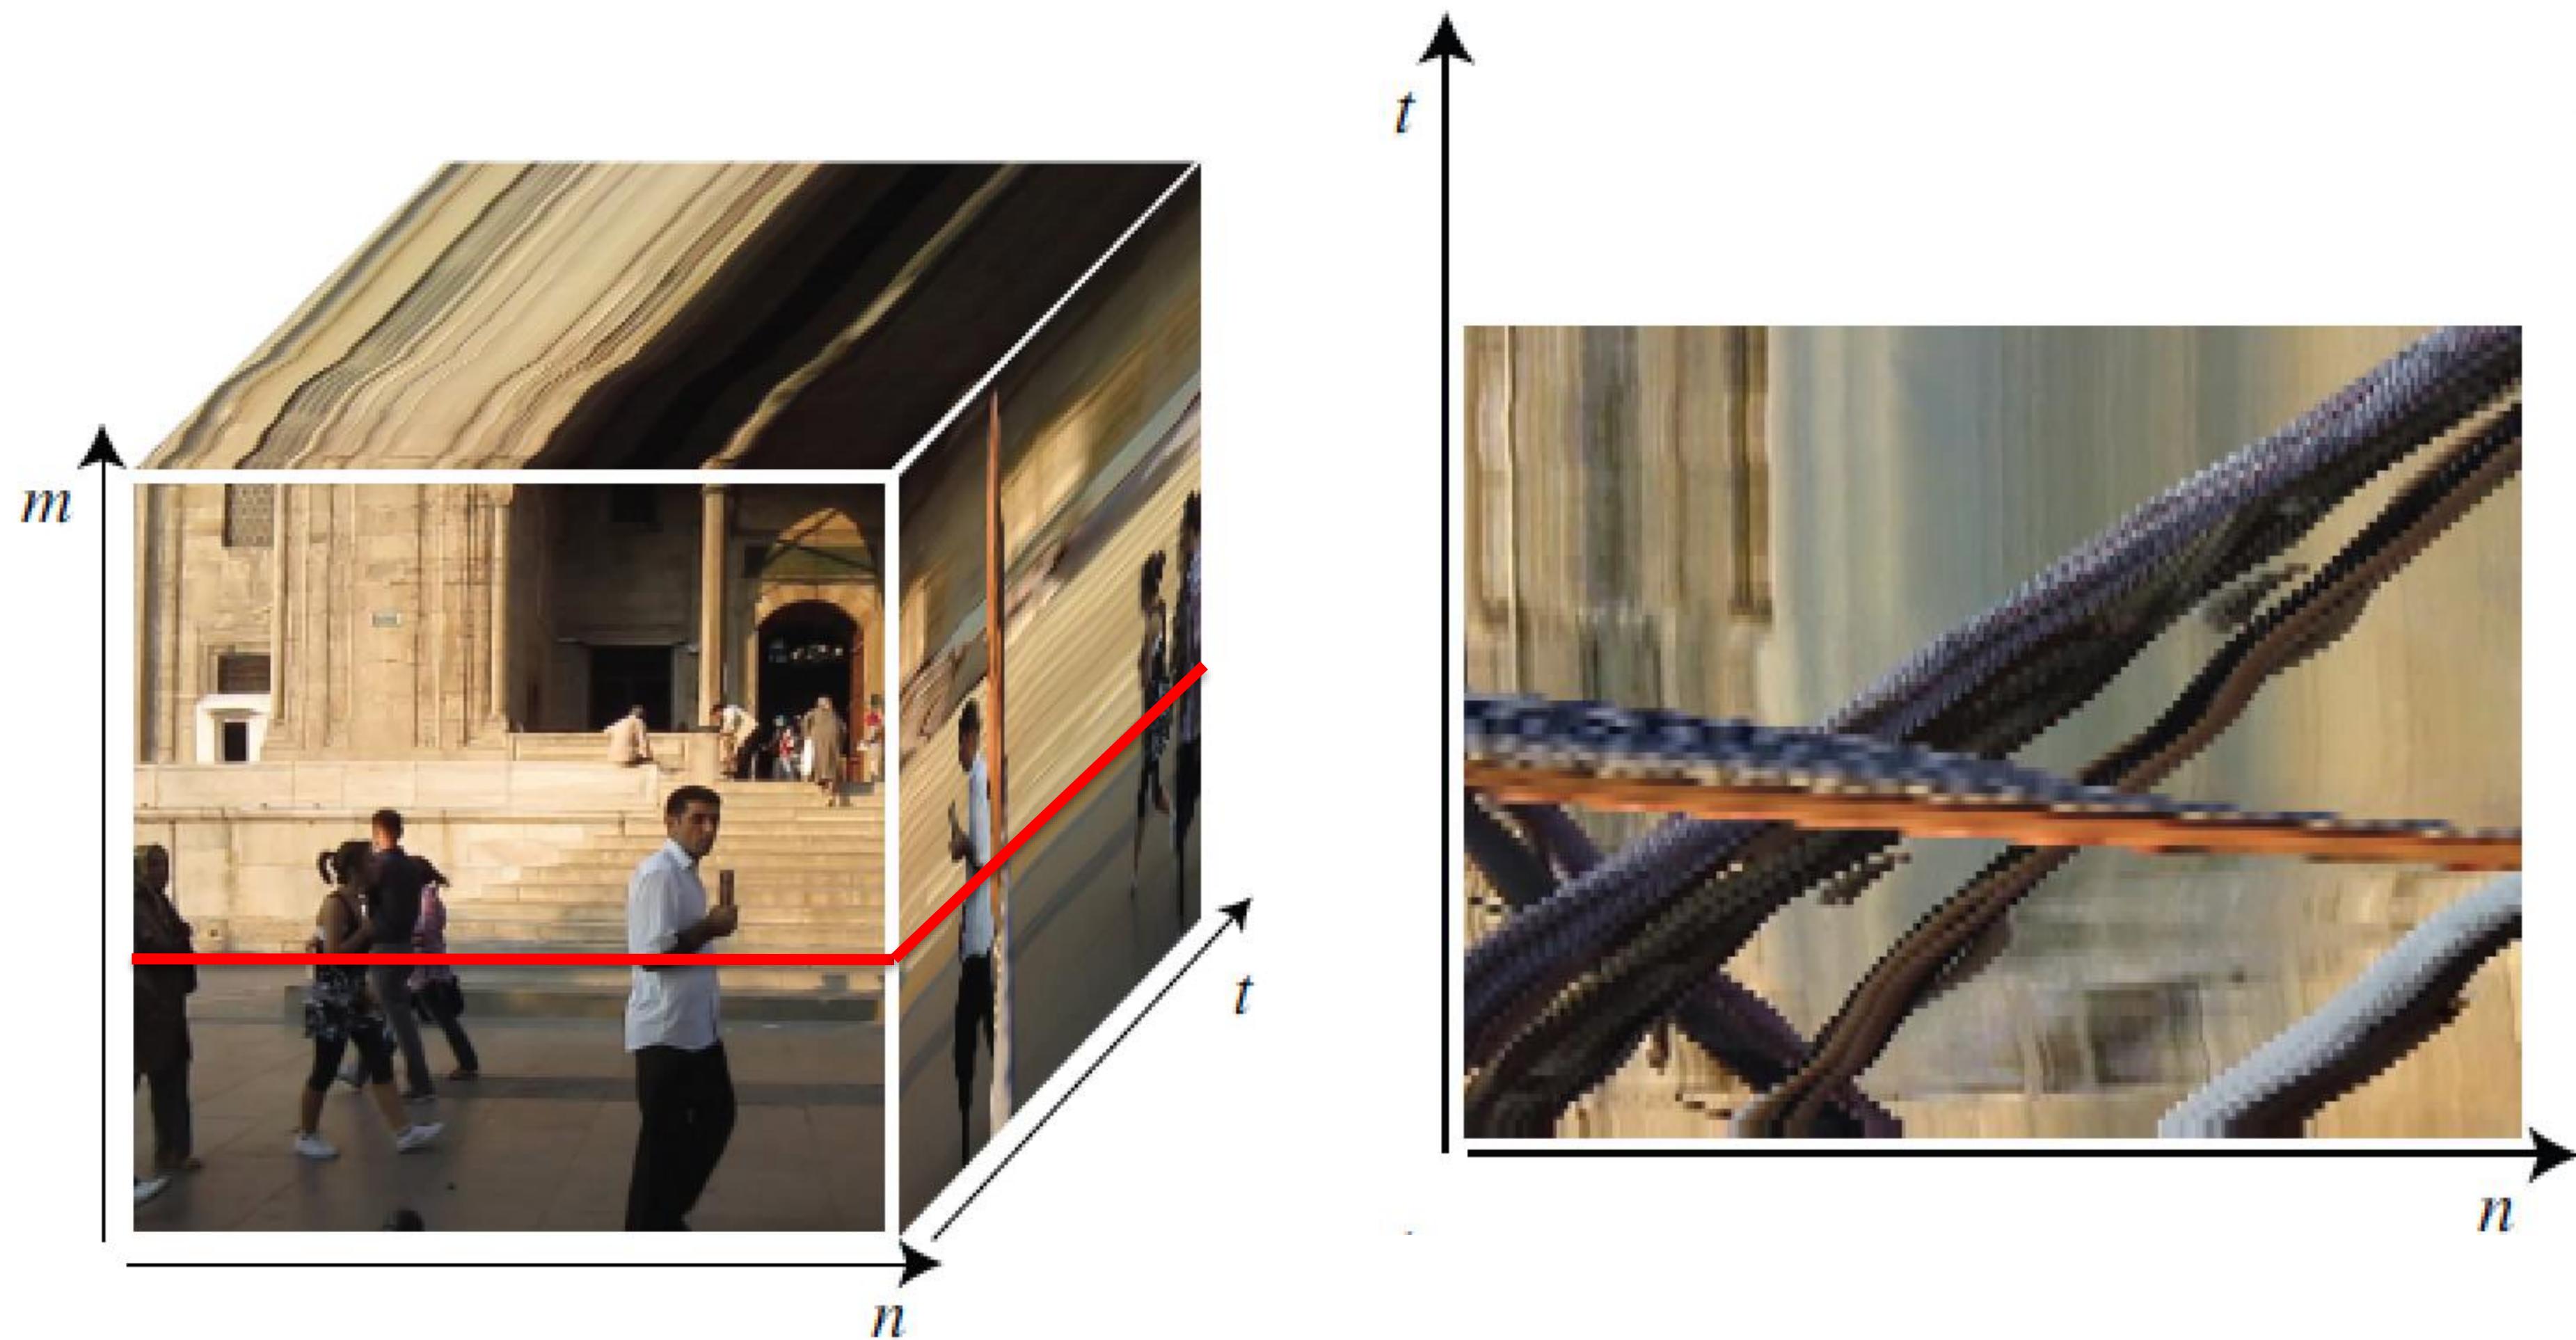
Sequences



time

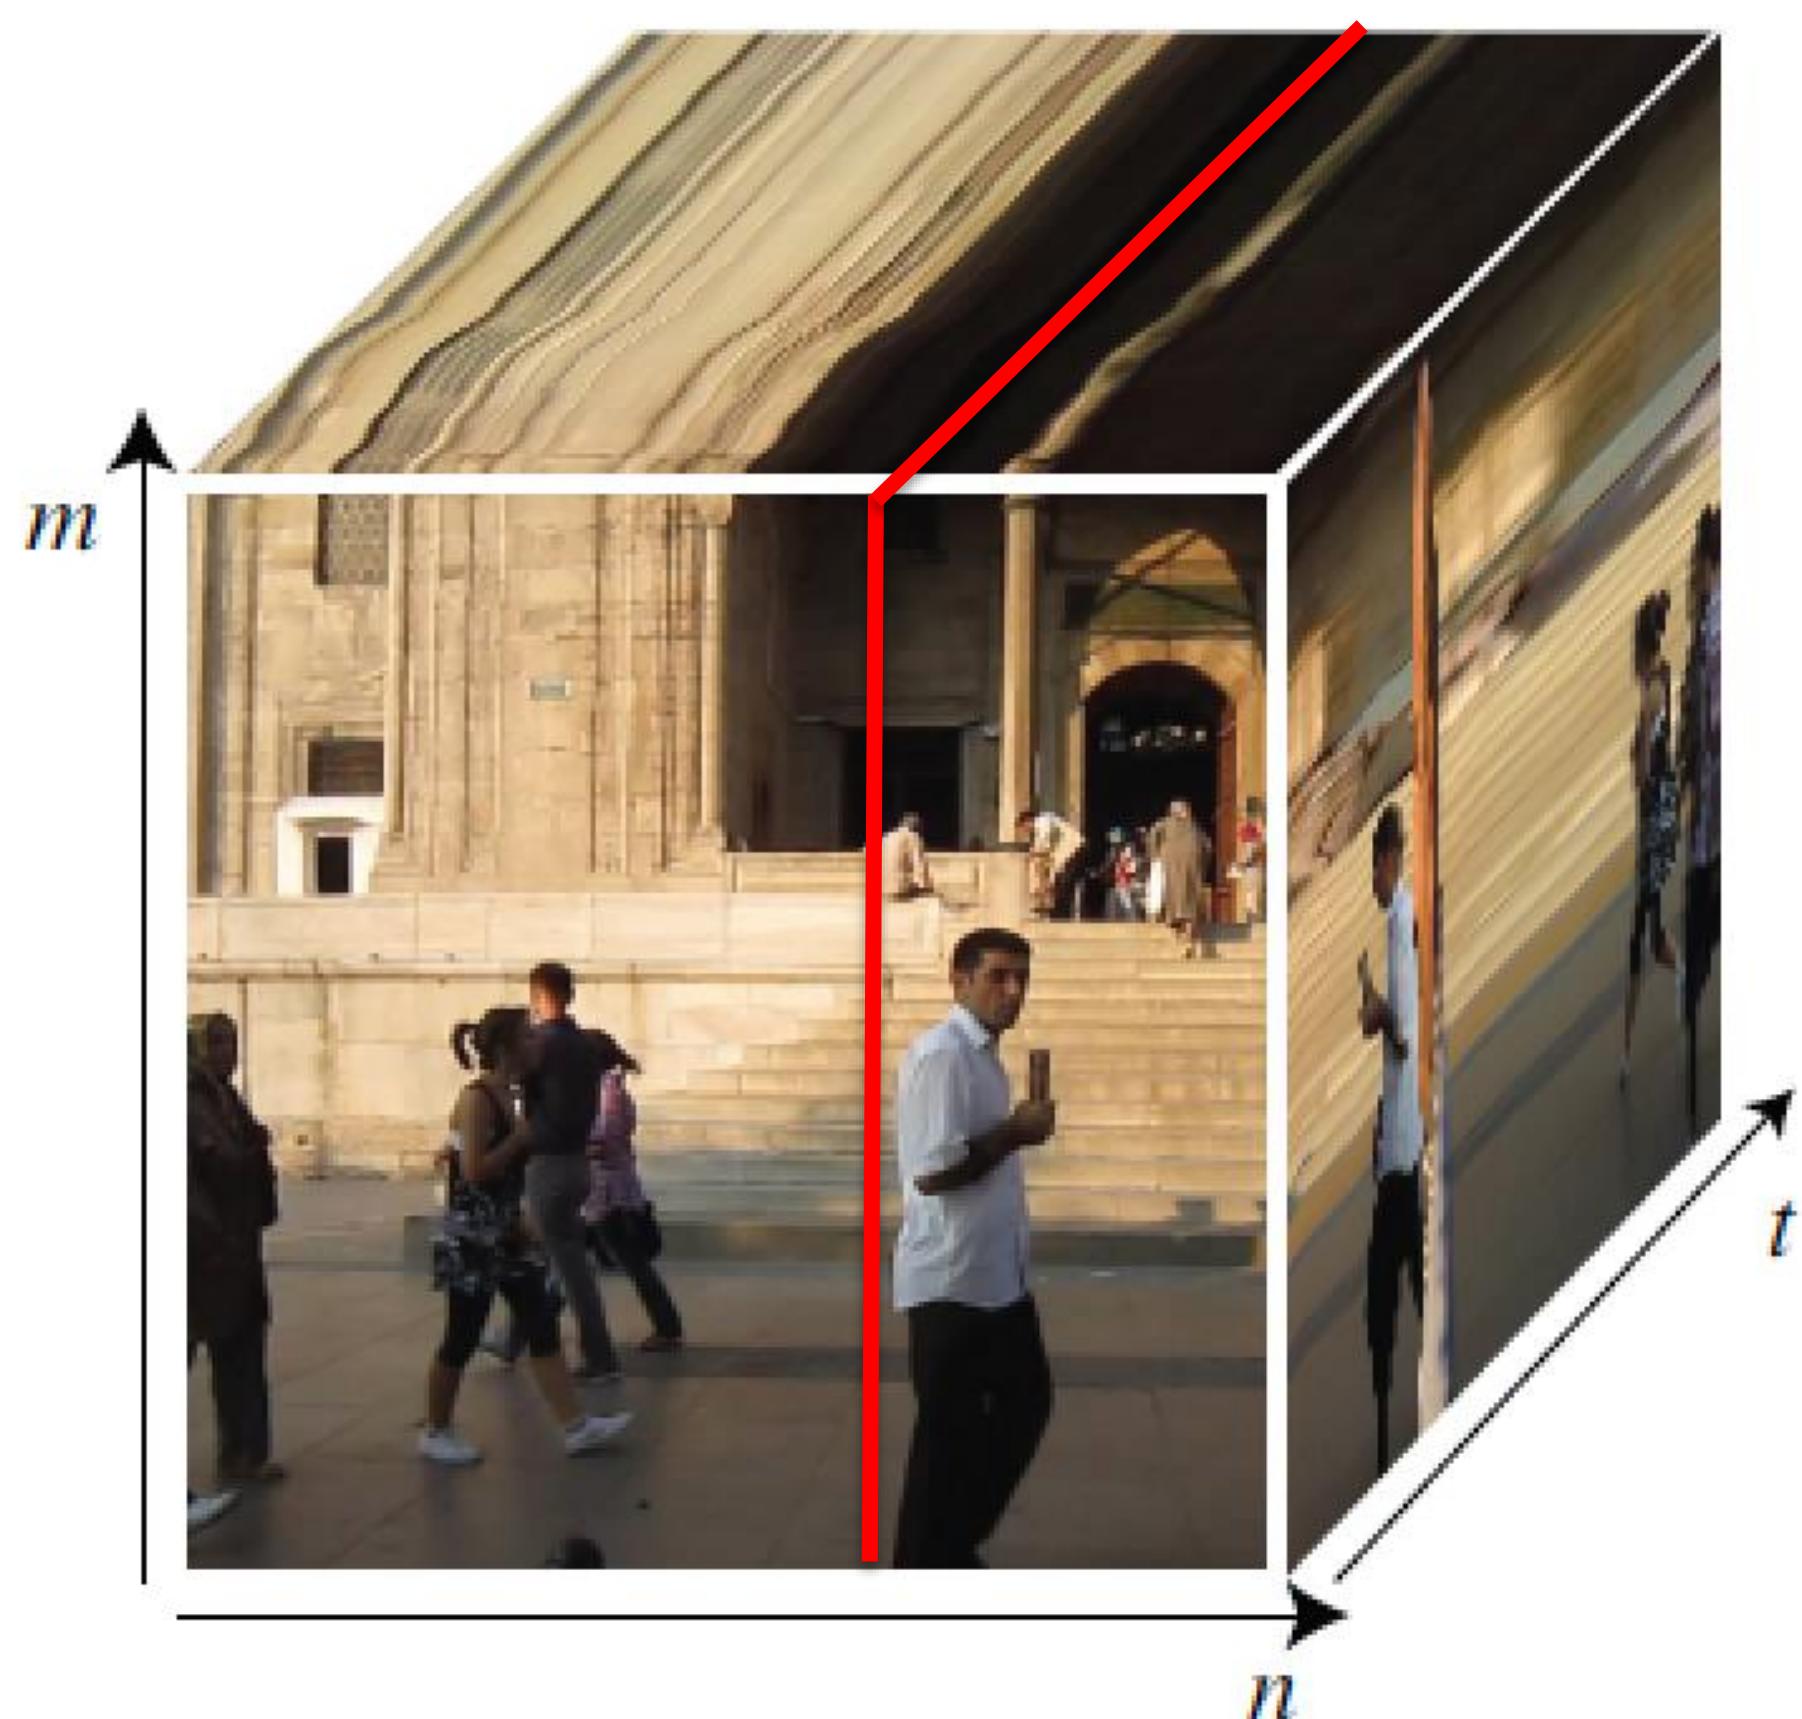


Sequences



Cube size = 128x128x90

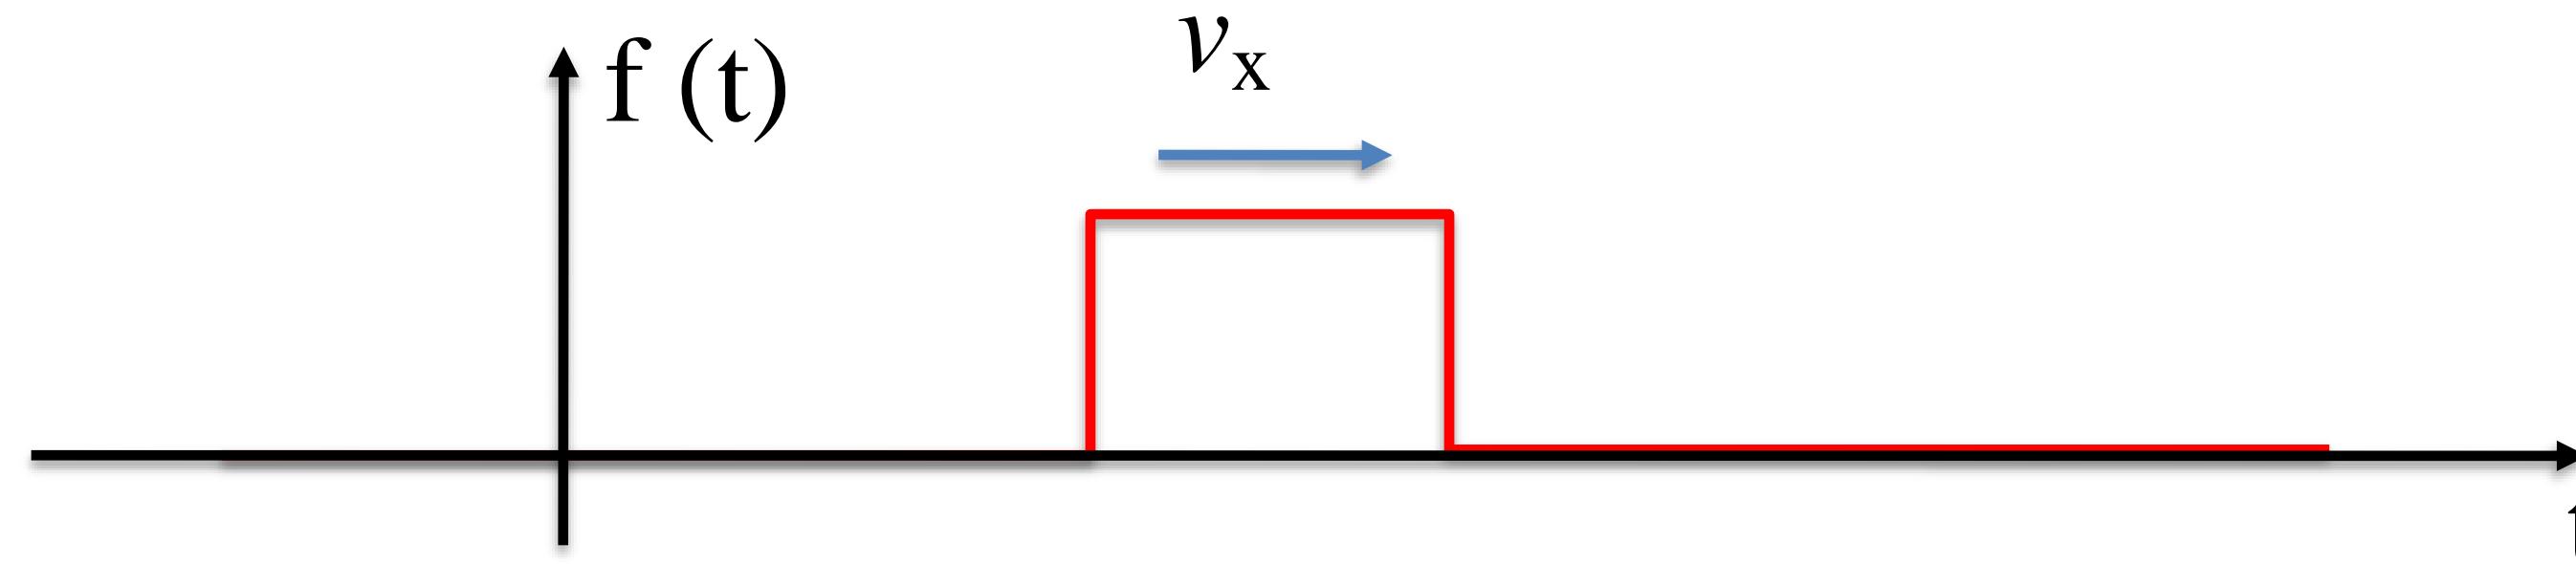
Sequences



Cube size = 128x128x90

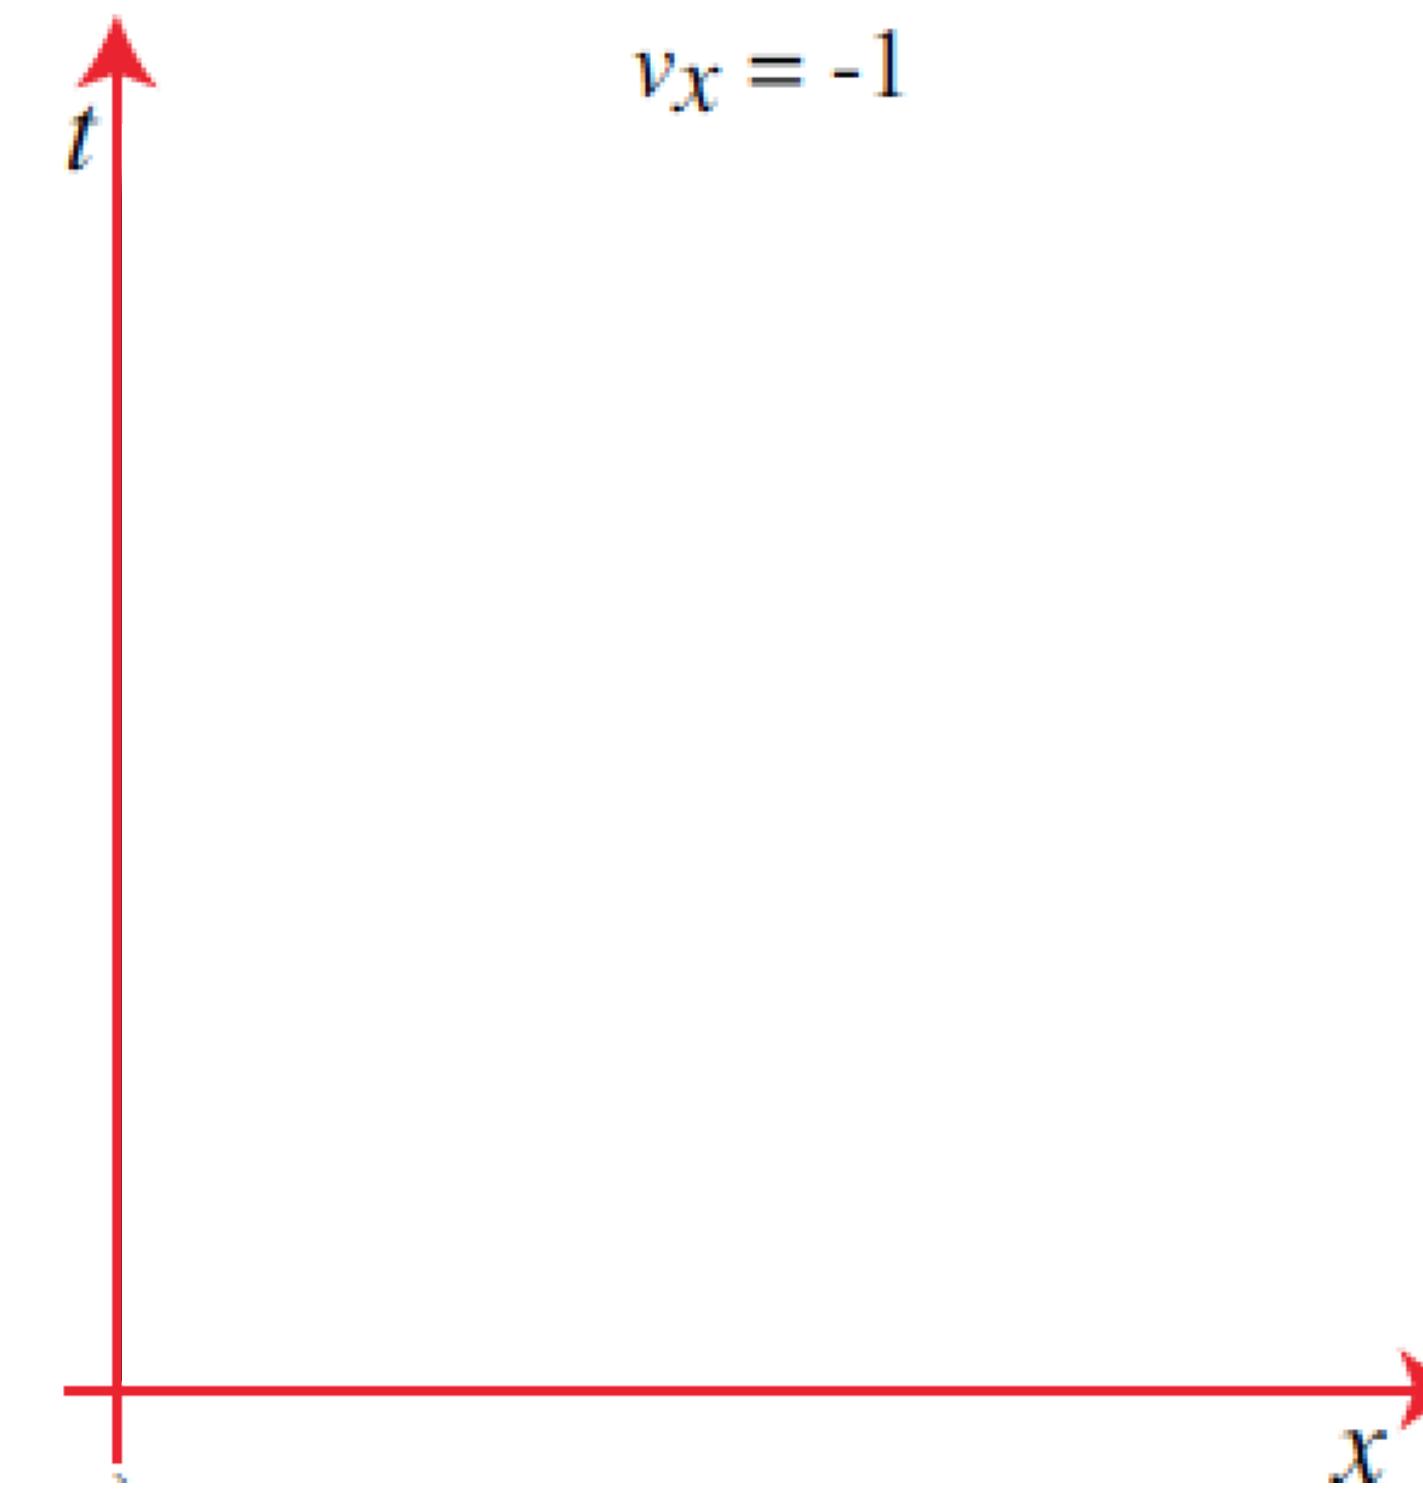
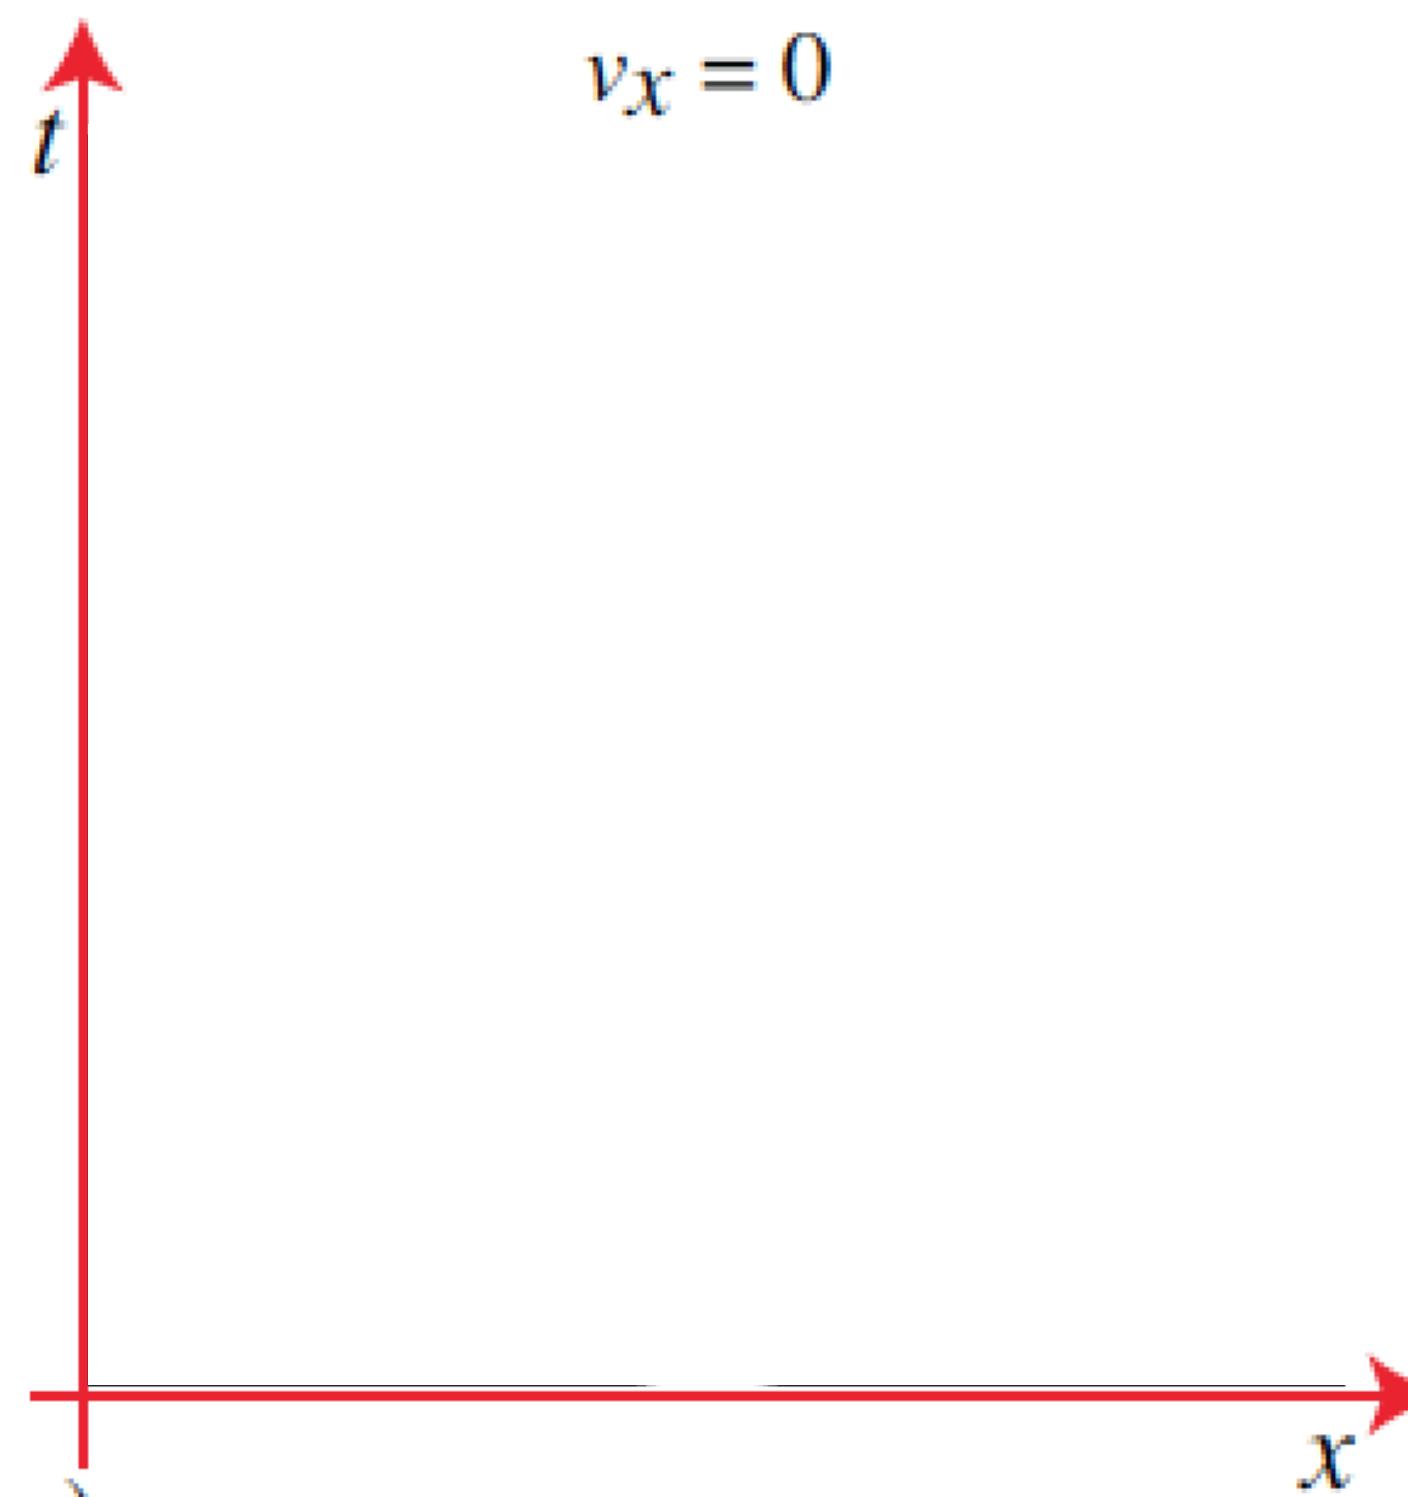


A box moving with speed v_x

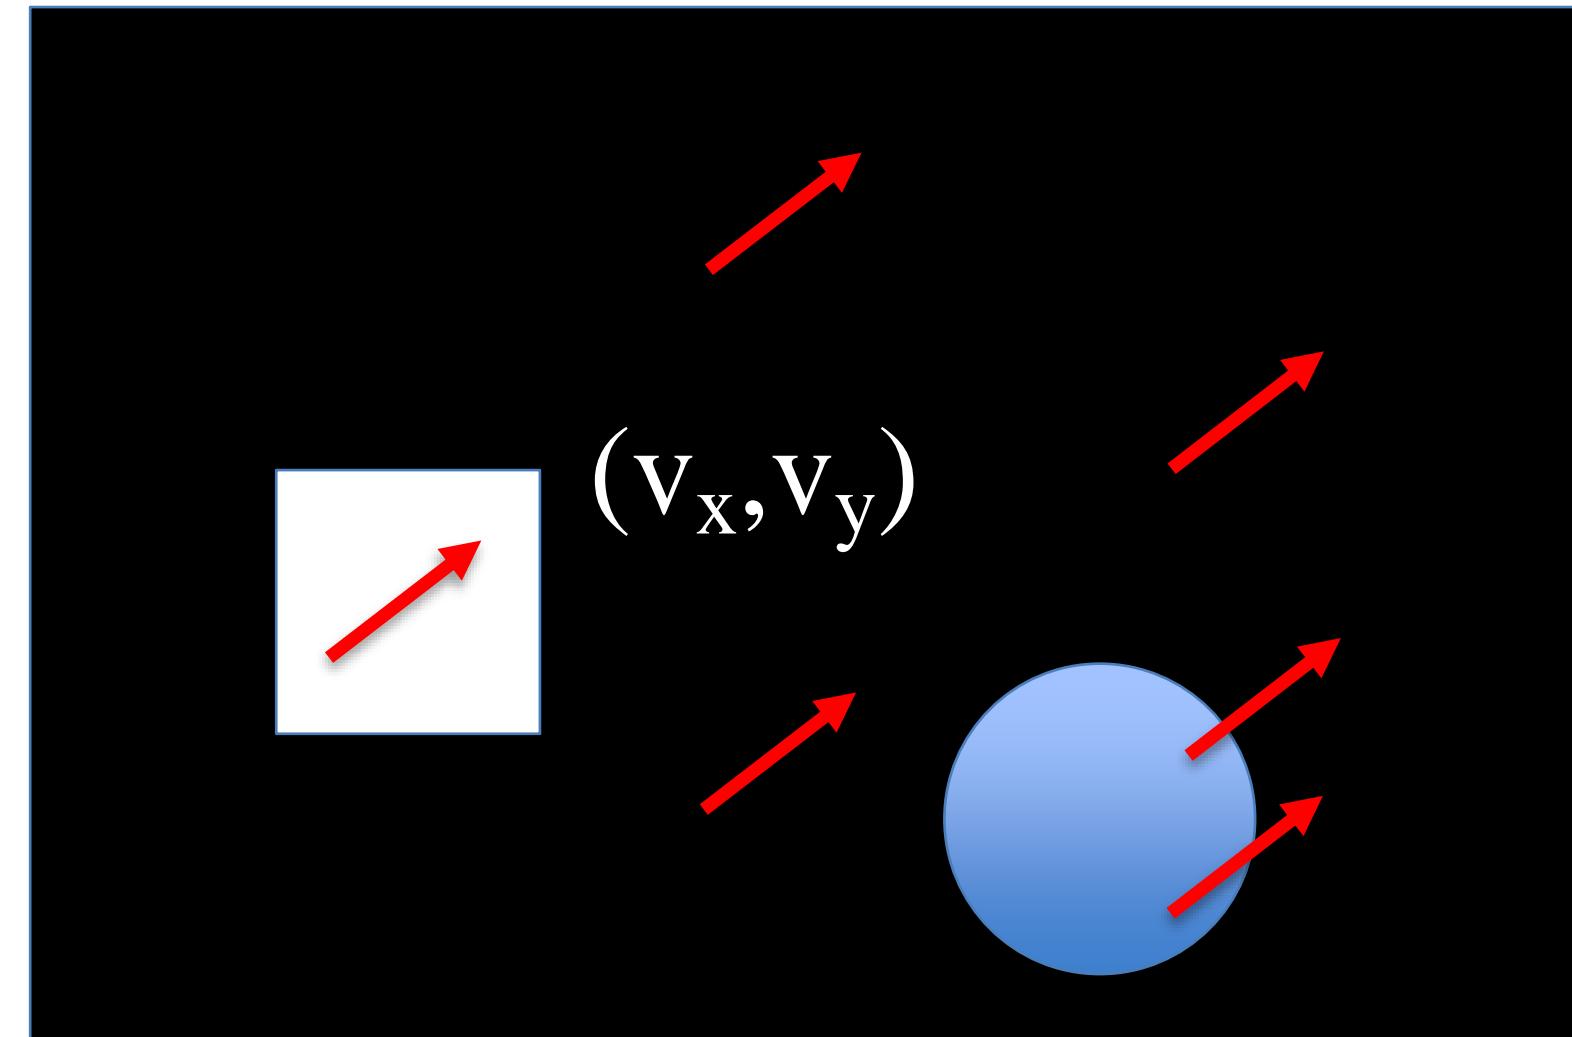


$$v_x = 0$$

$$v_x = -1$$



Global constant motion



A global motion of the image can be written as:

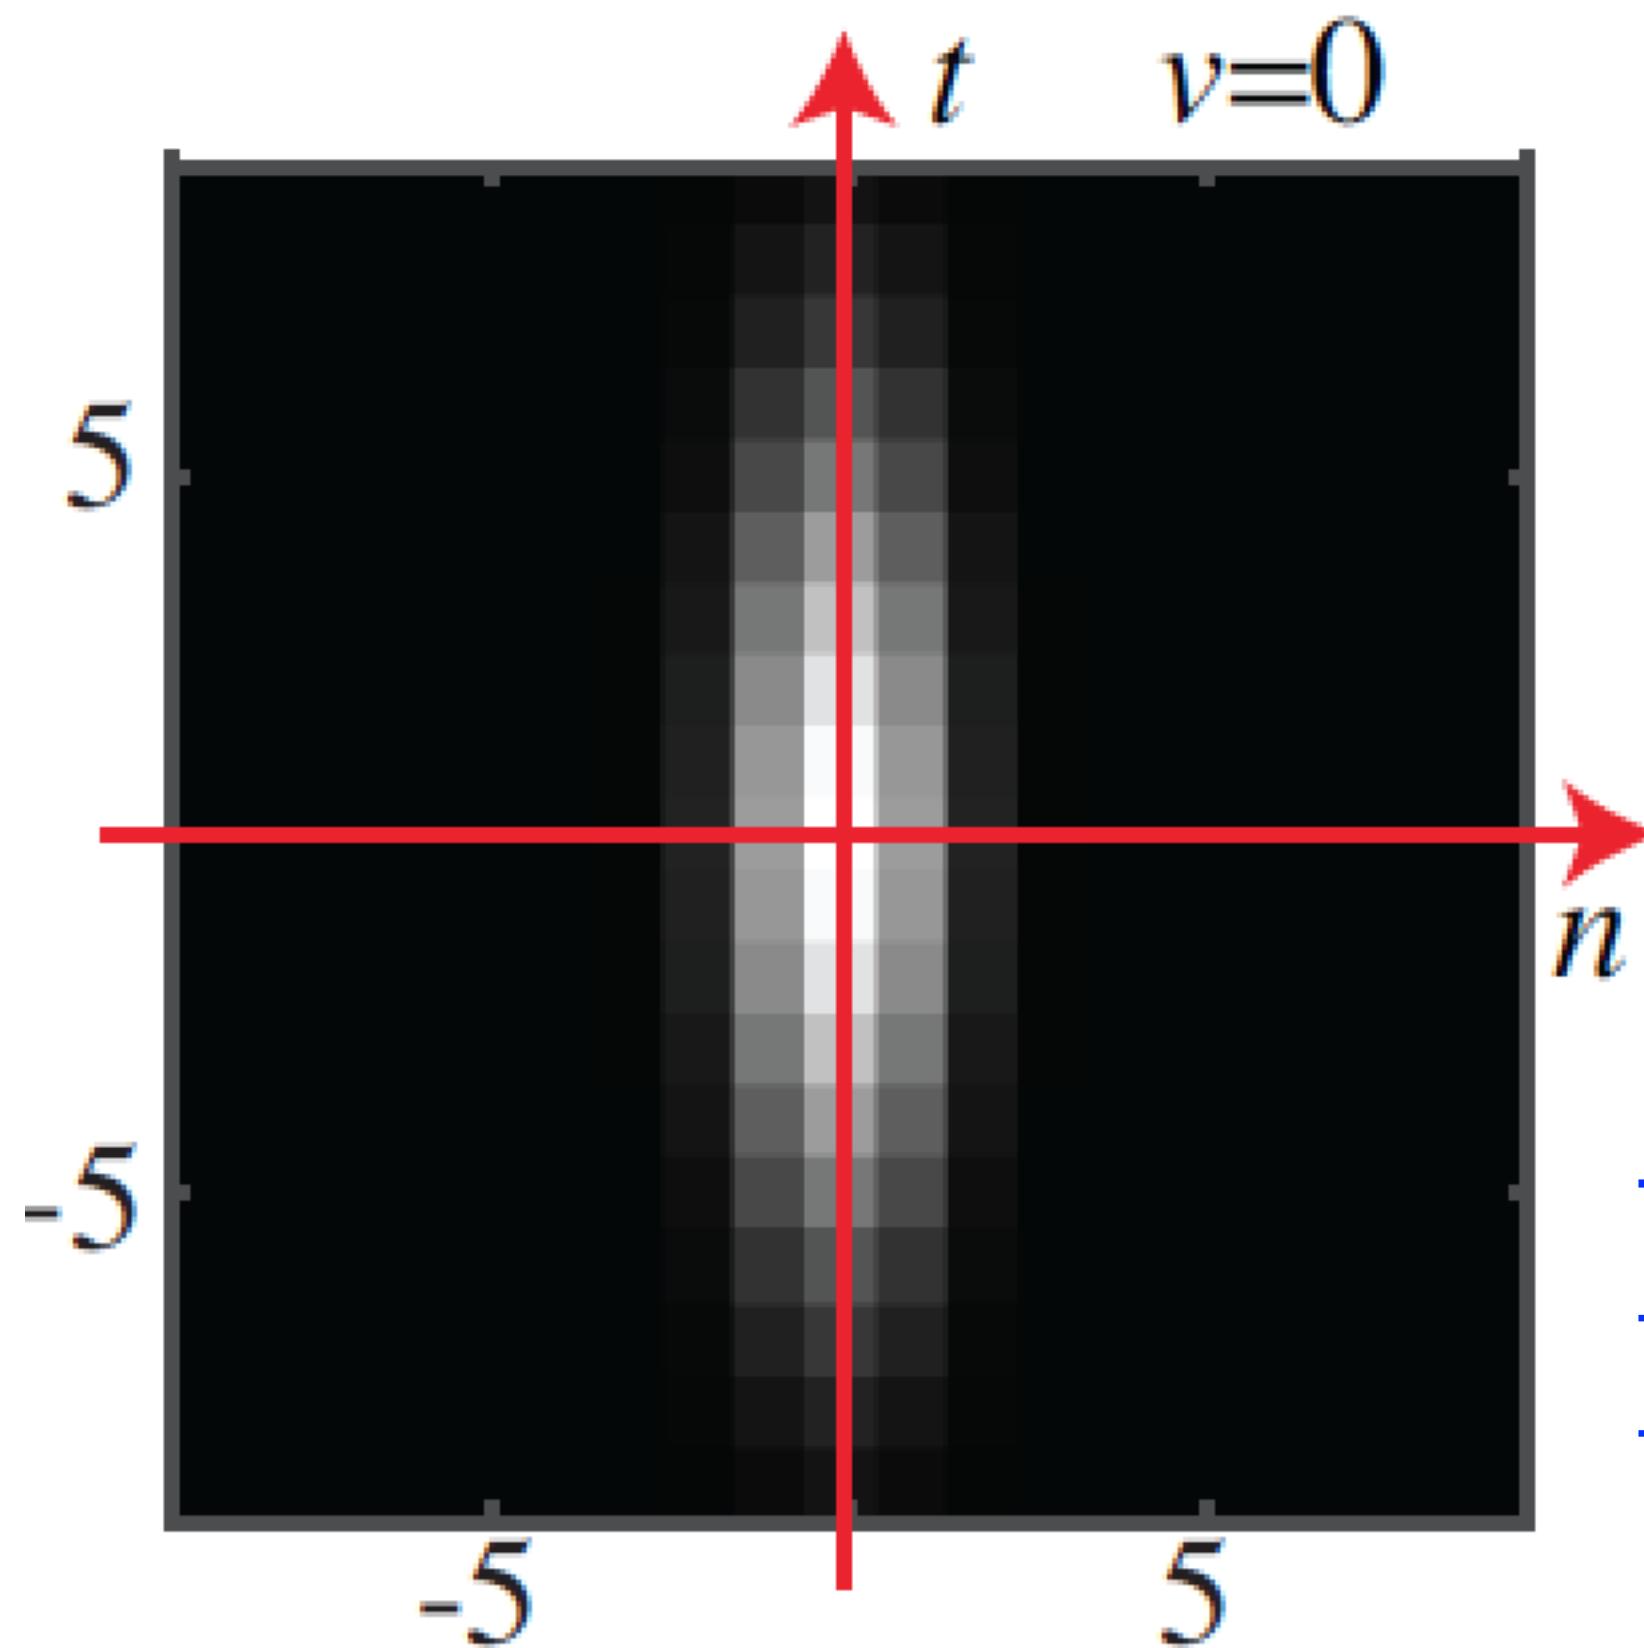
$$f(x, y, t) = f_0(x - v_x t, y - v_y t)$$

Where:

$$f_0(x, y) = f(x, y, 0)$$

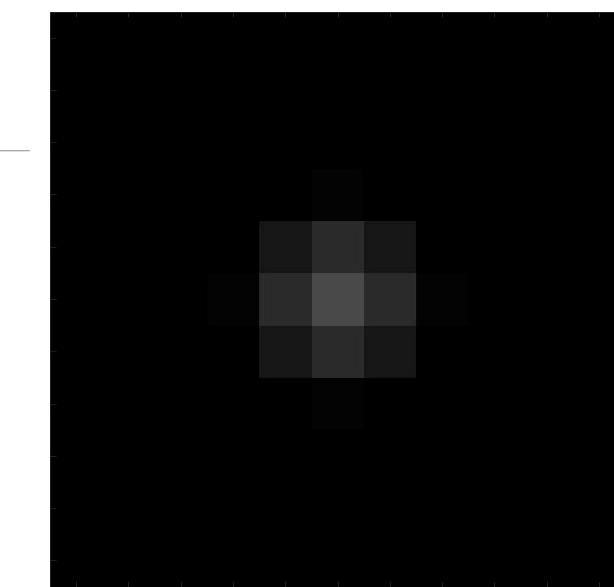
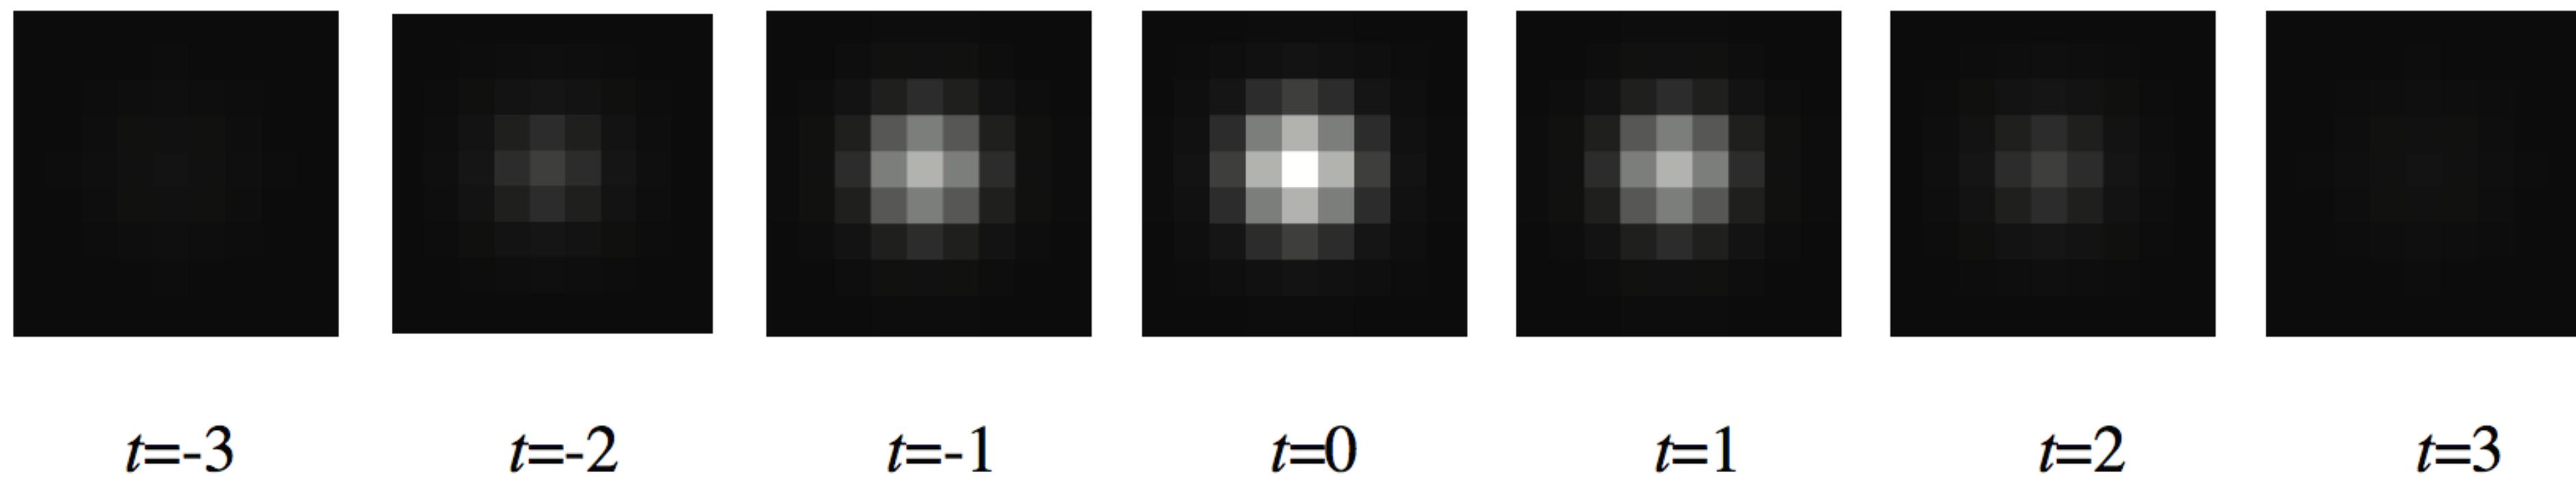
Temporal Gaussian

$$g(x, y, t; \sigma_x, \sigma_t) = \frac{1}{(2\pi)^{3/2} \sigma_x^2 \sigma_t} \exp -\frac{x^2 + y^2}{2\sigma_x^2} \exp -\frac{t^2}{2\sigma_t^2}$$



This filter keeps stationary things sharp, and blurs moving things.

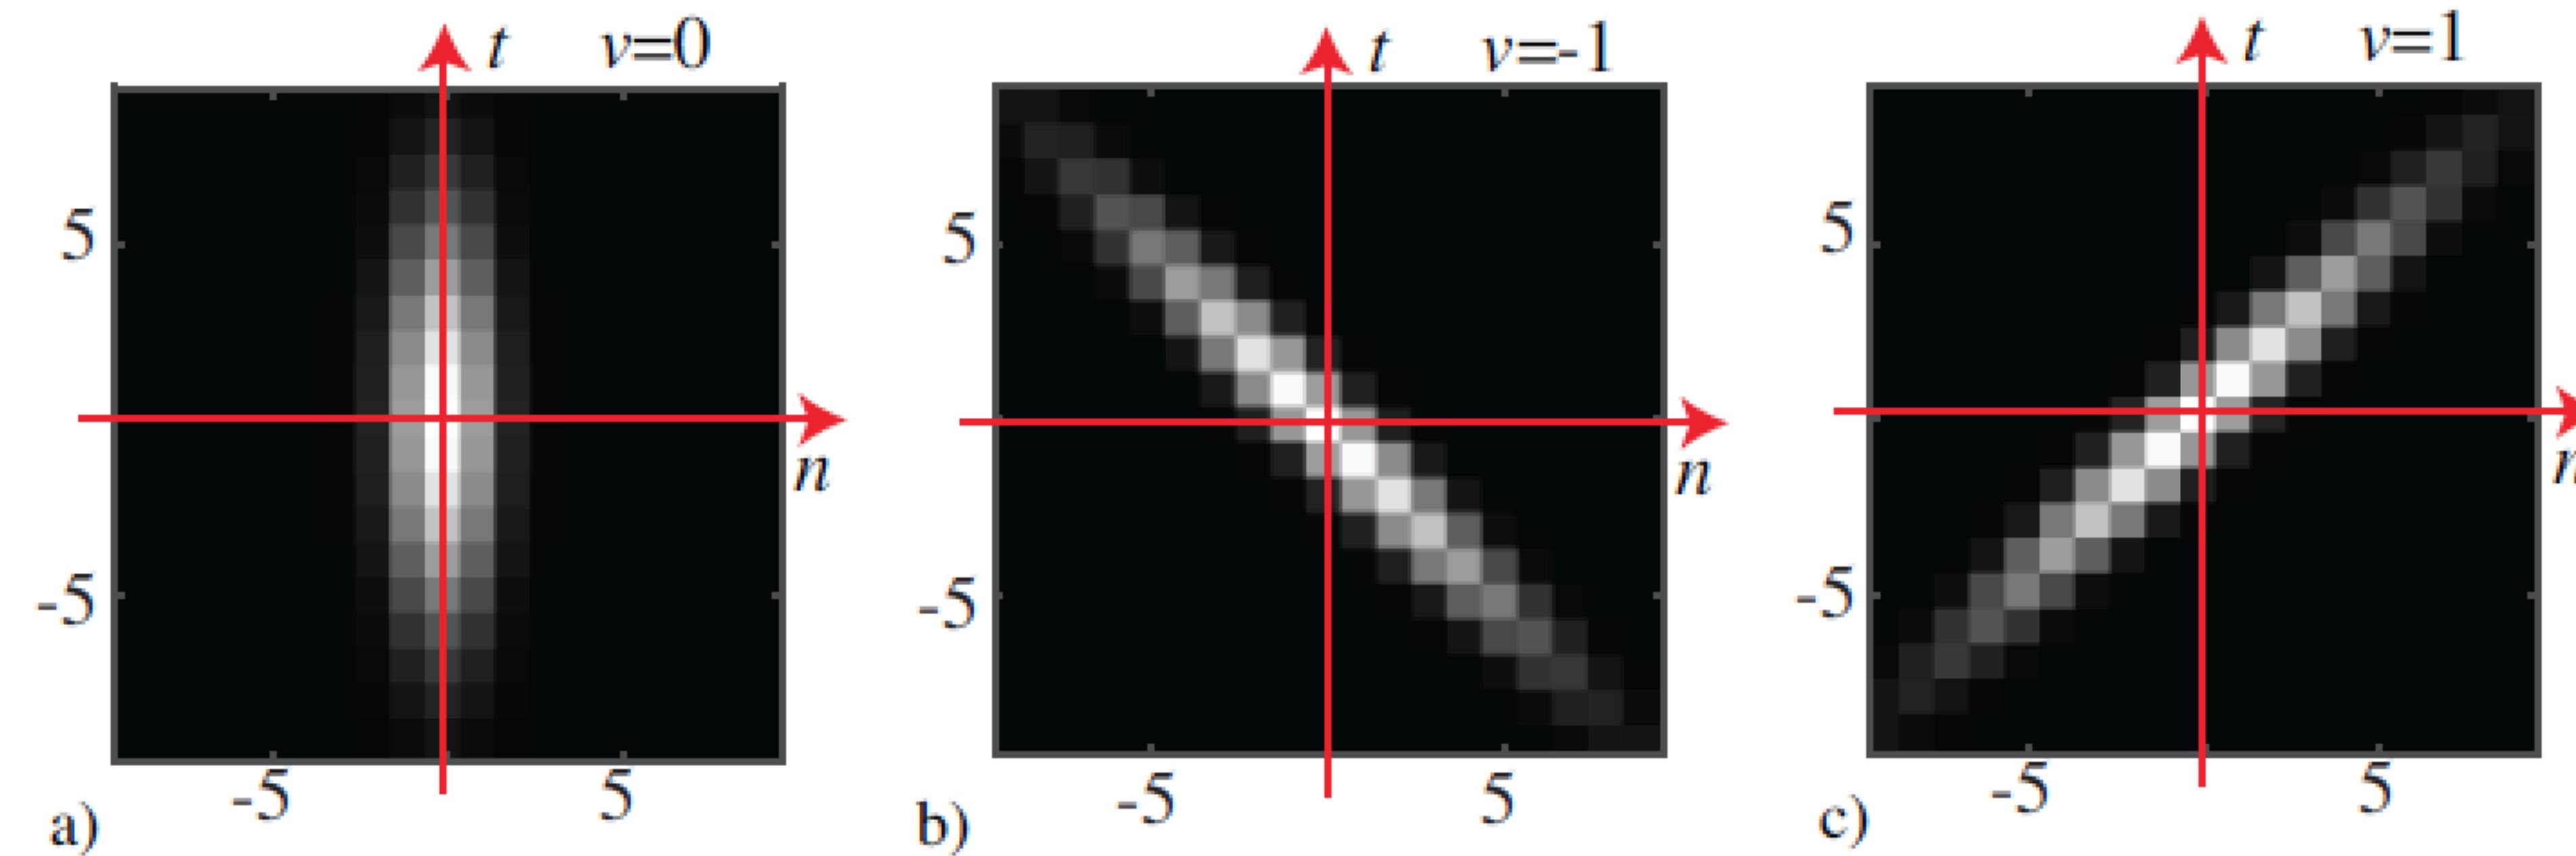
Spatio-temporal Gaussian

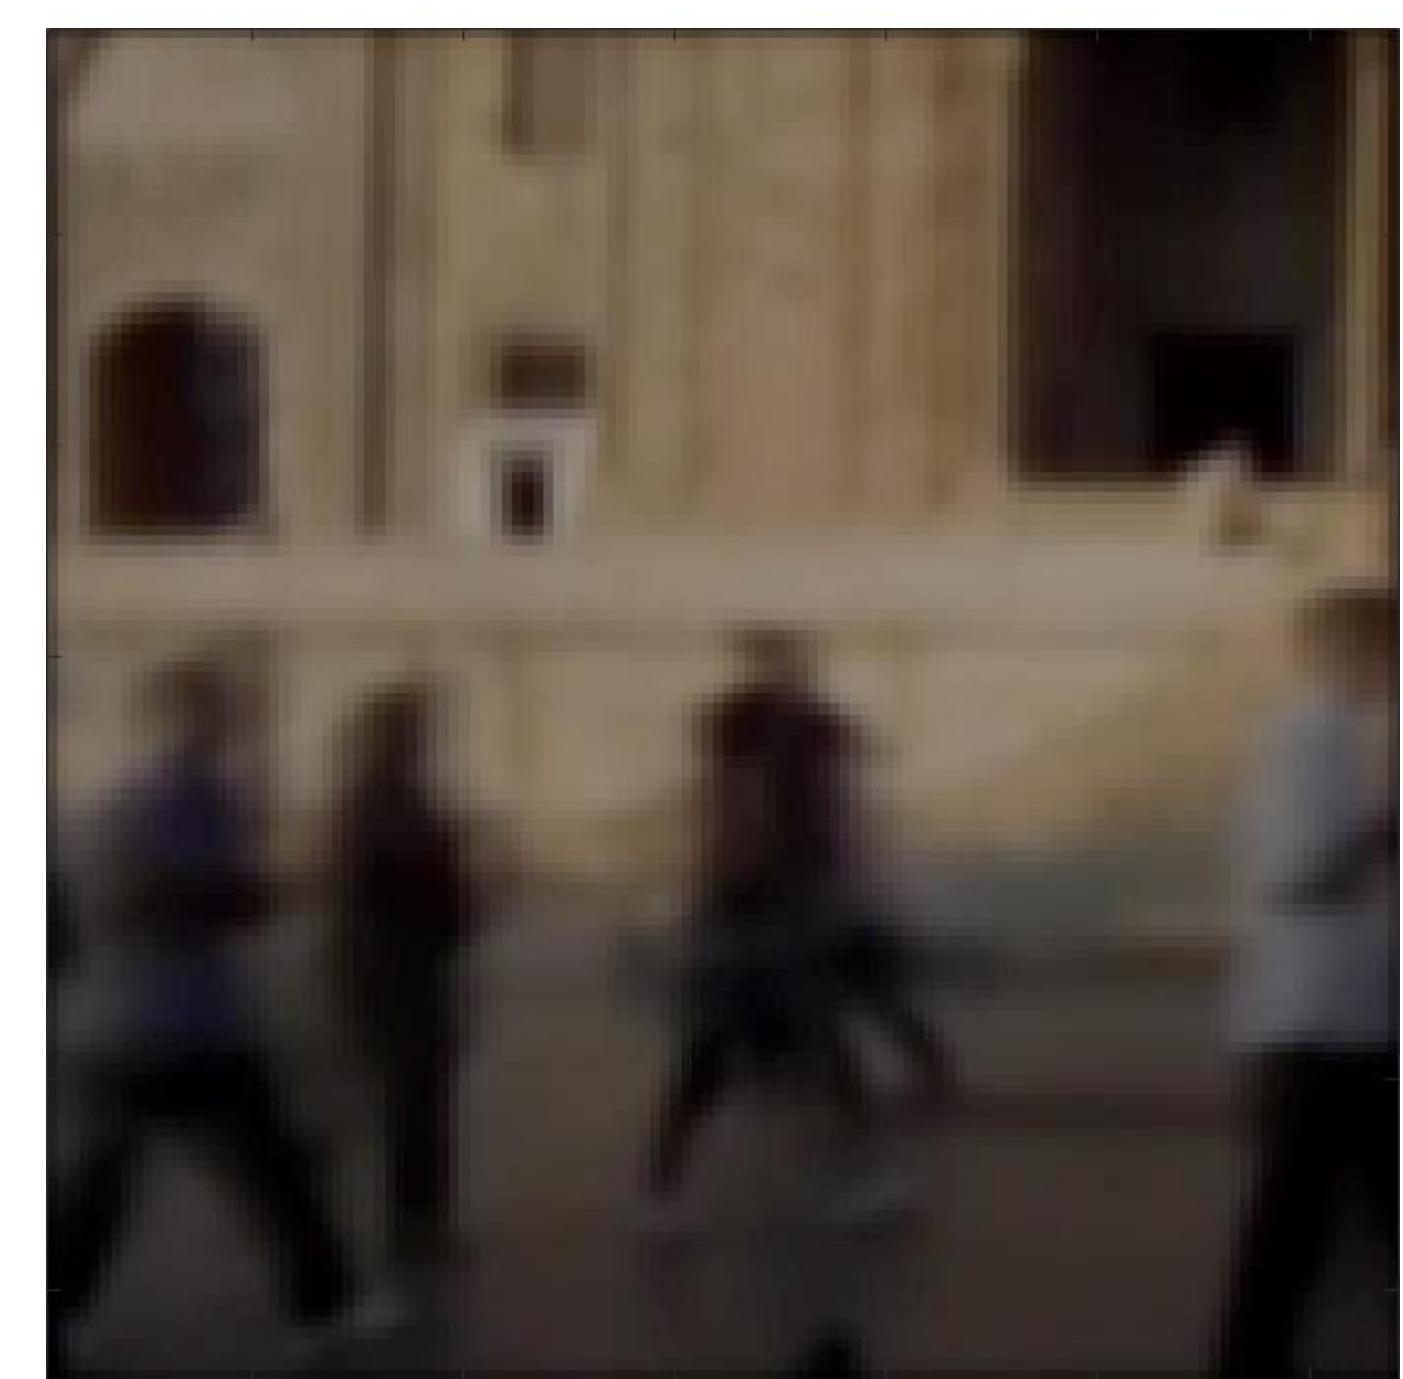
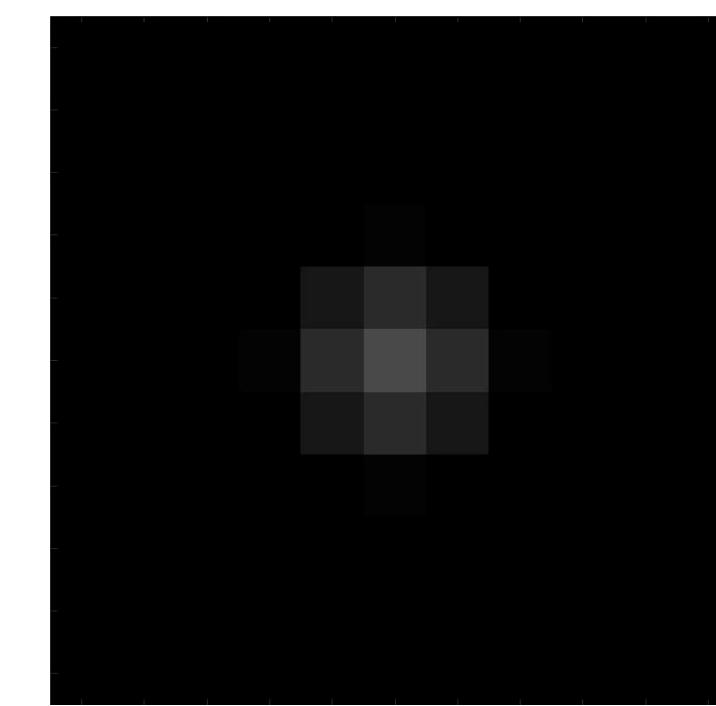
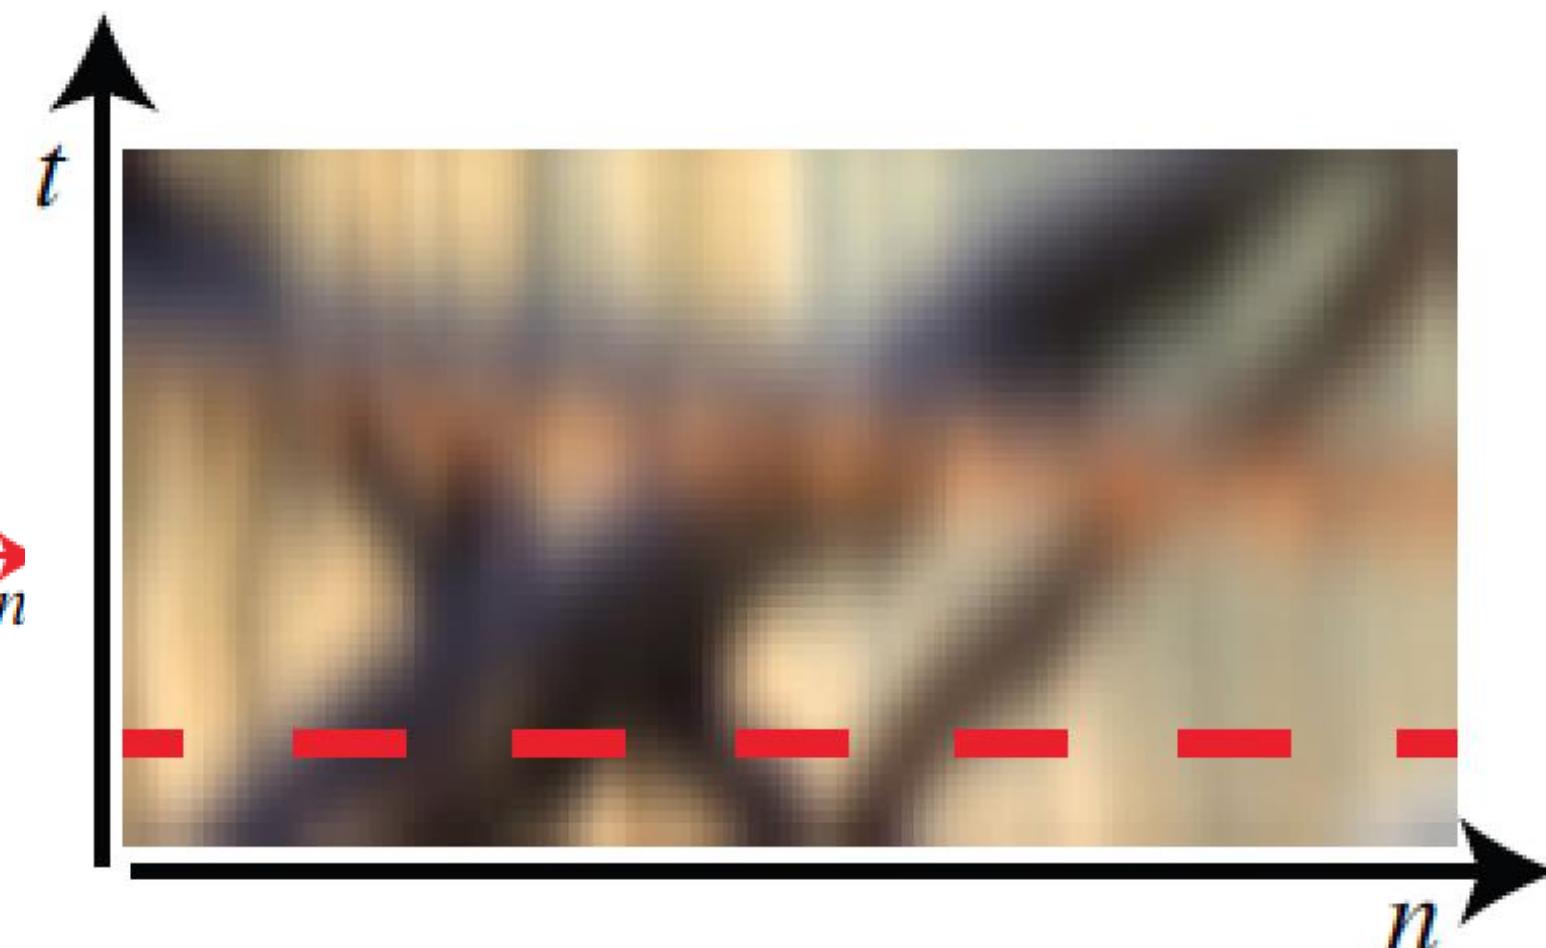
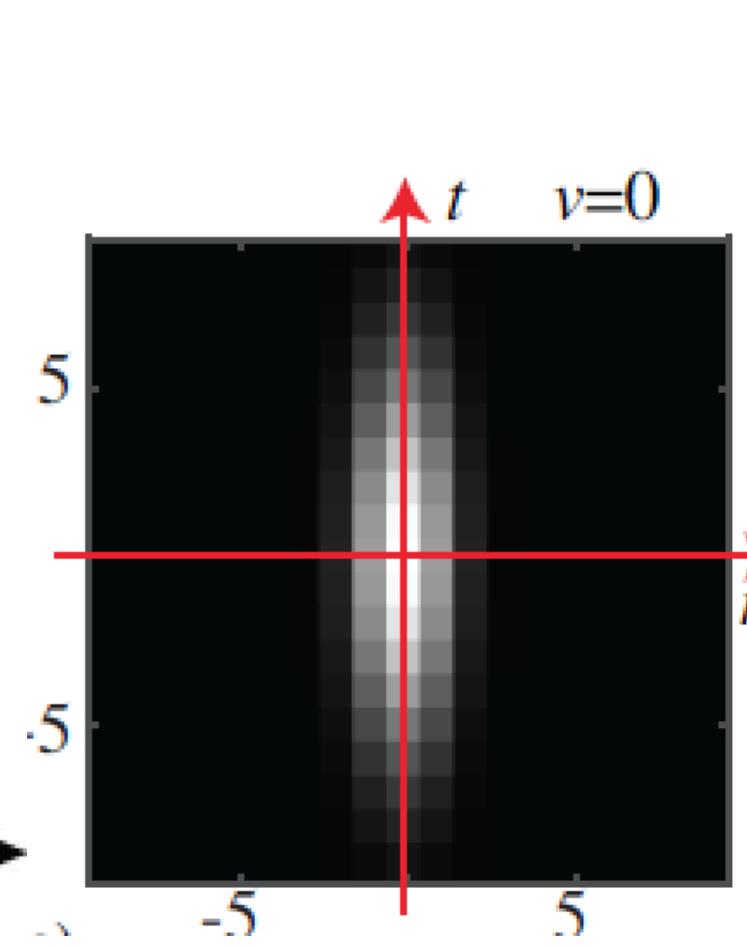
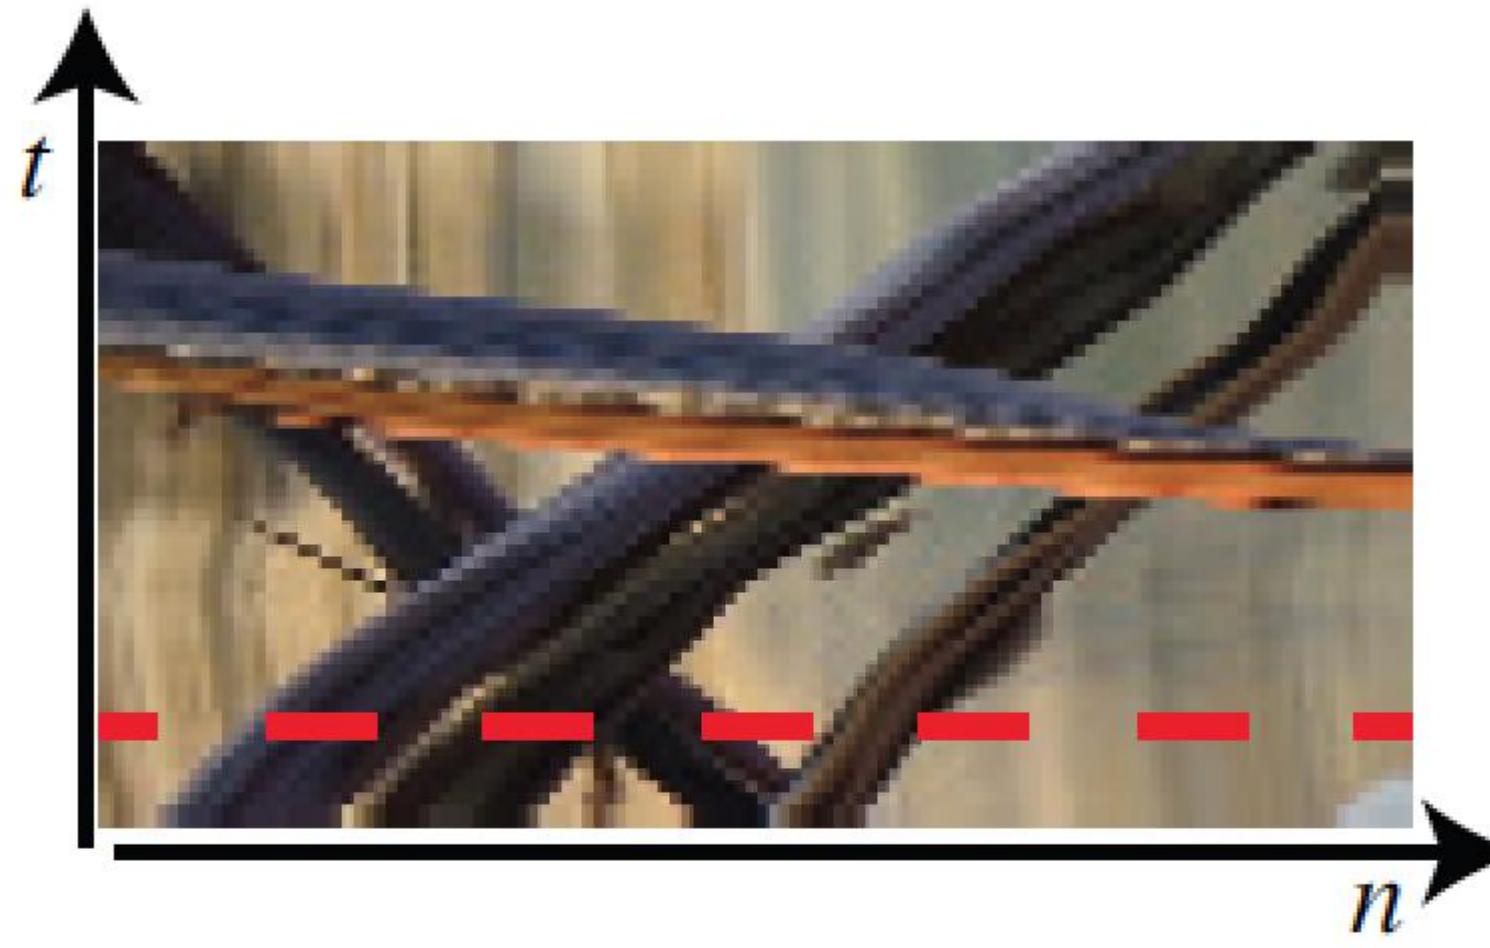


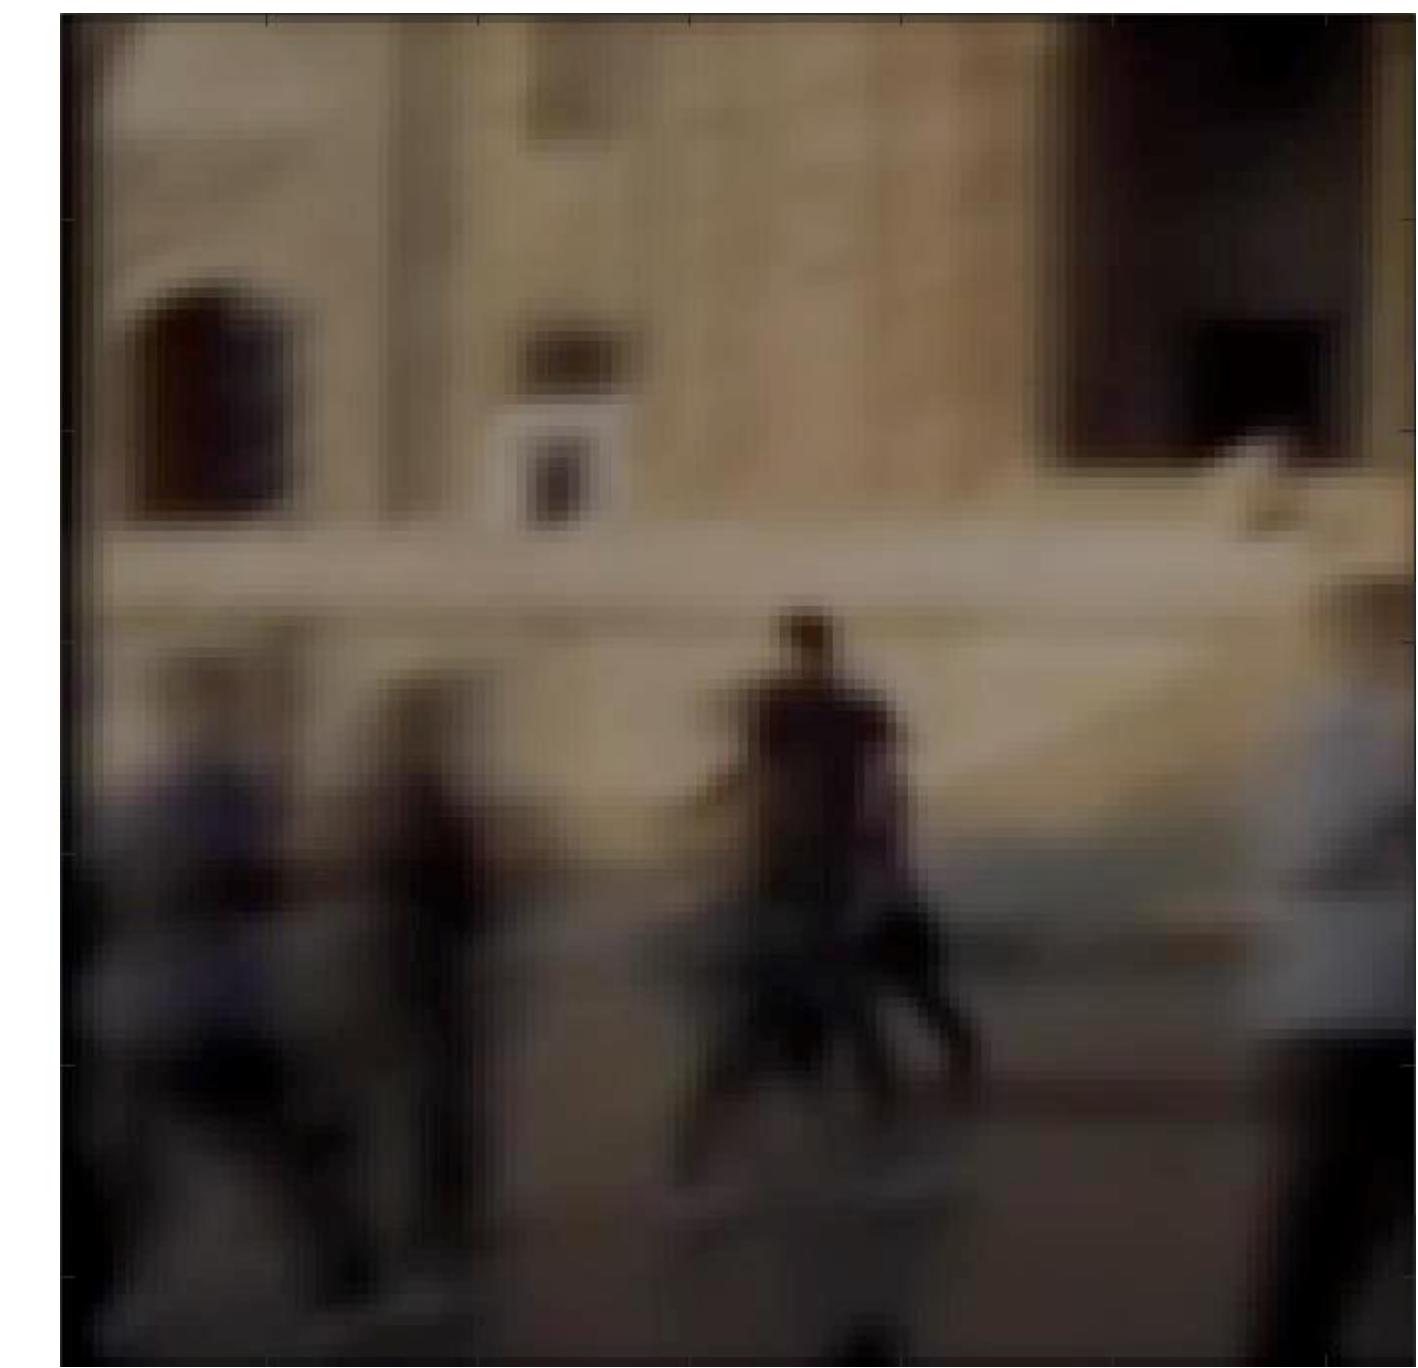
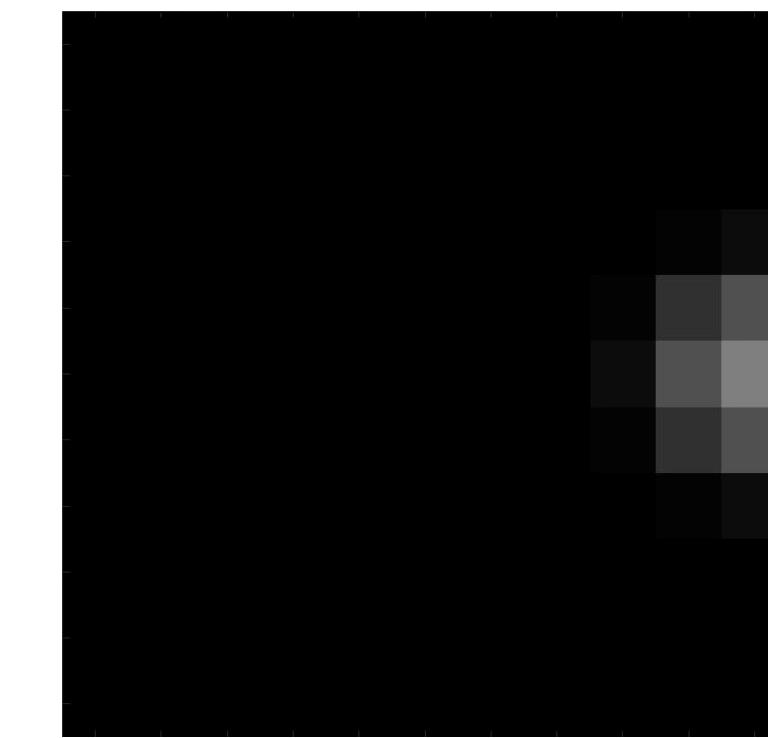
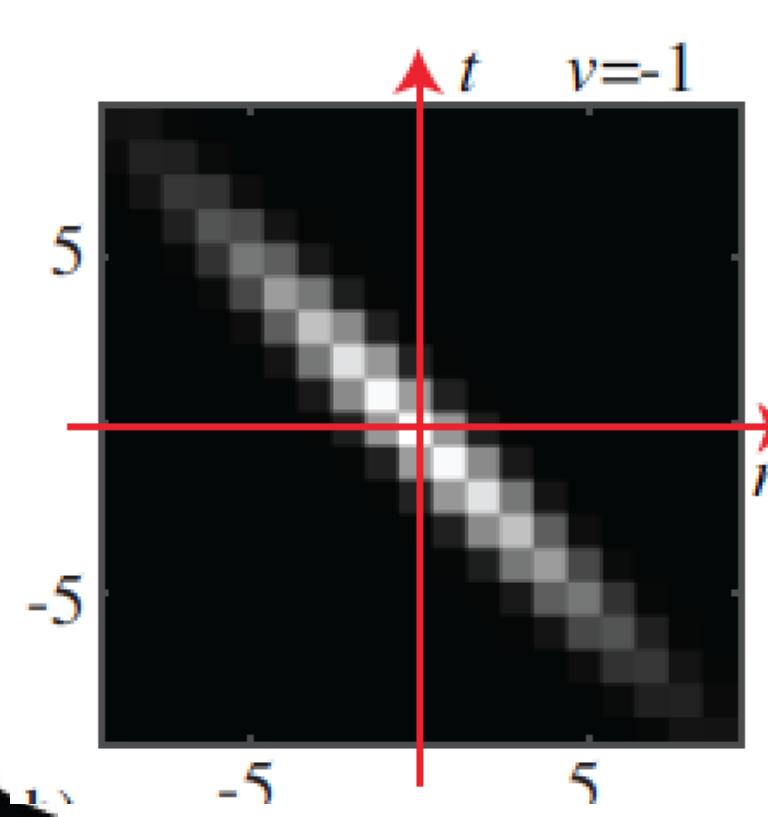
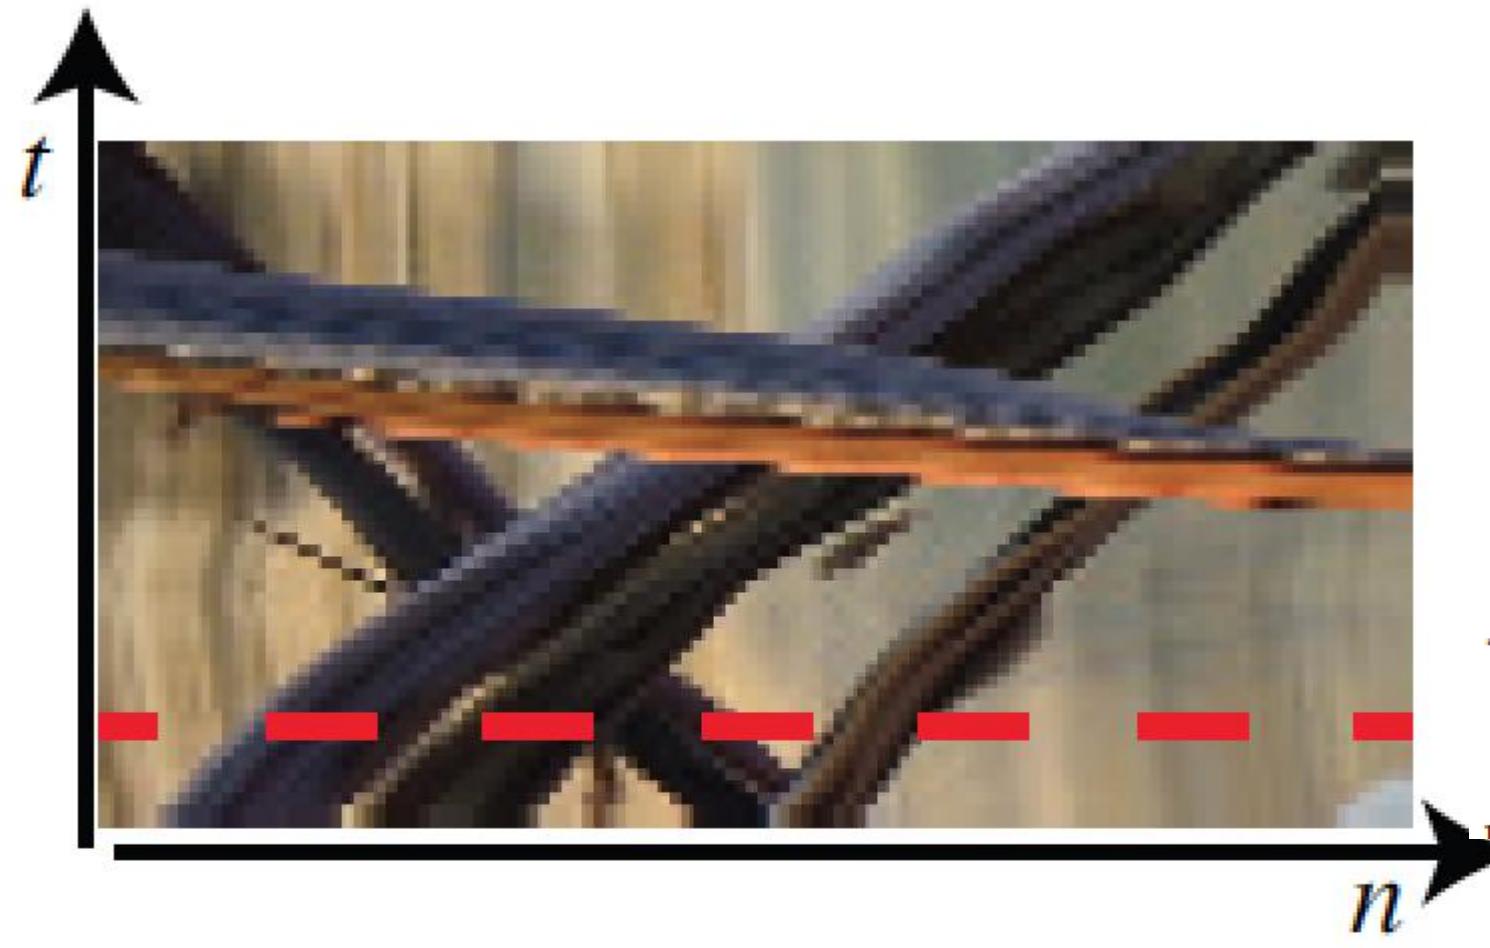
Spatio-temporal Gaussian

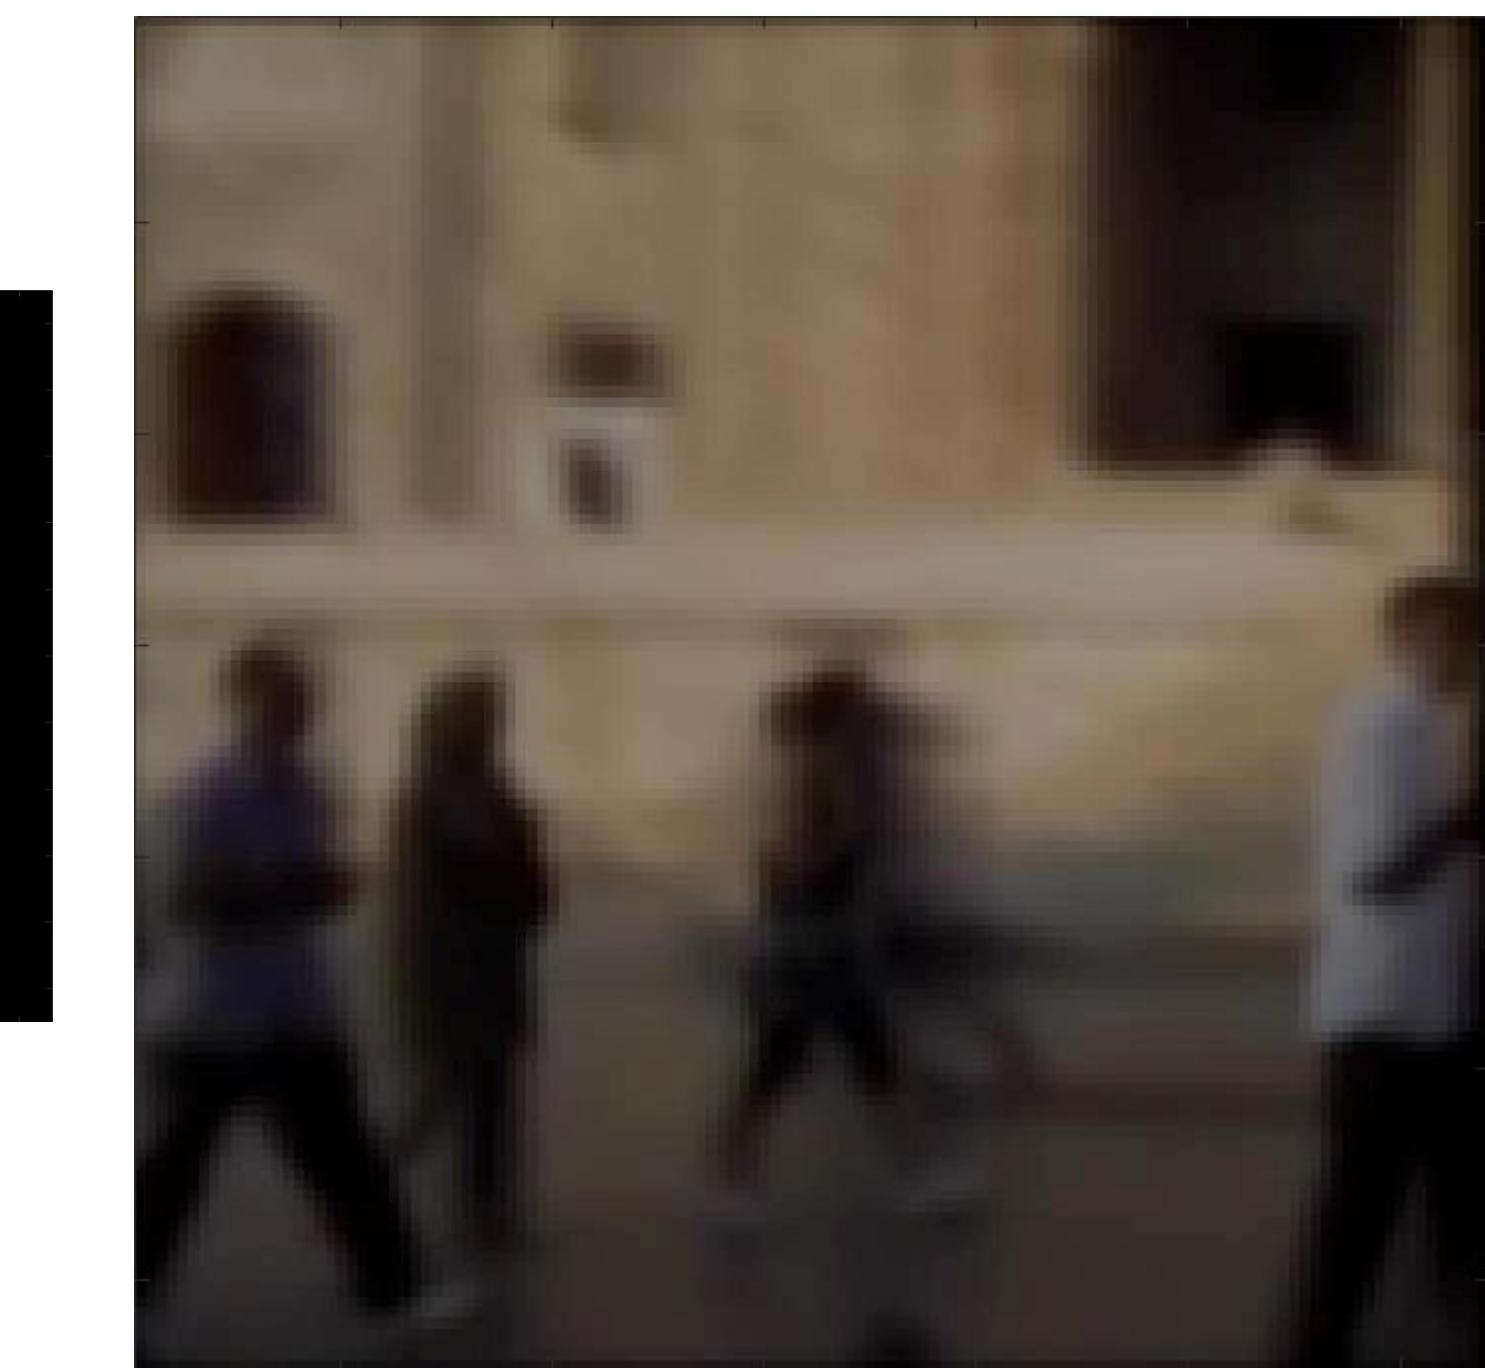
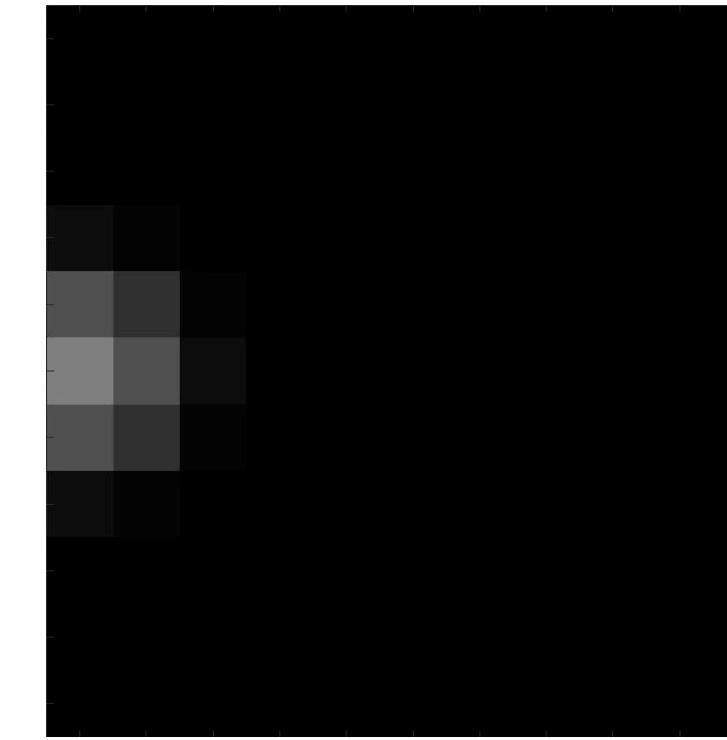
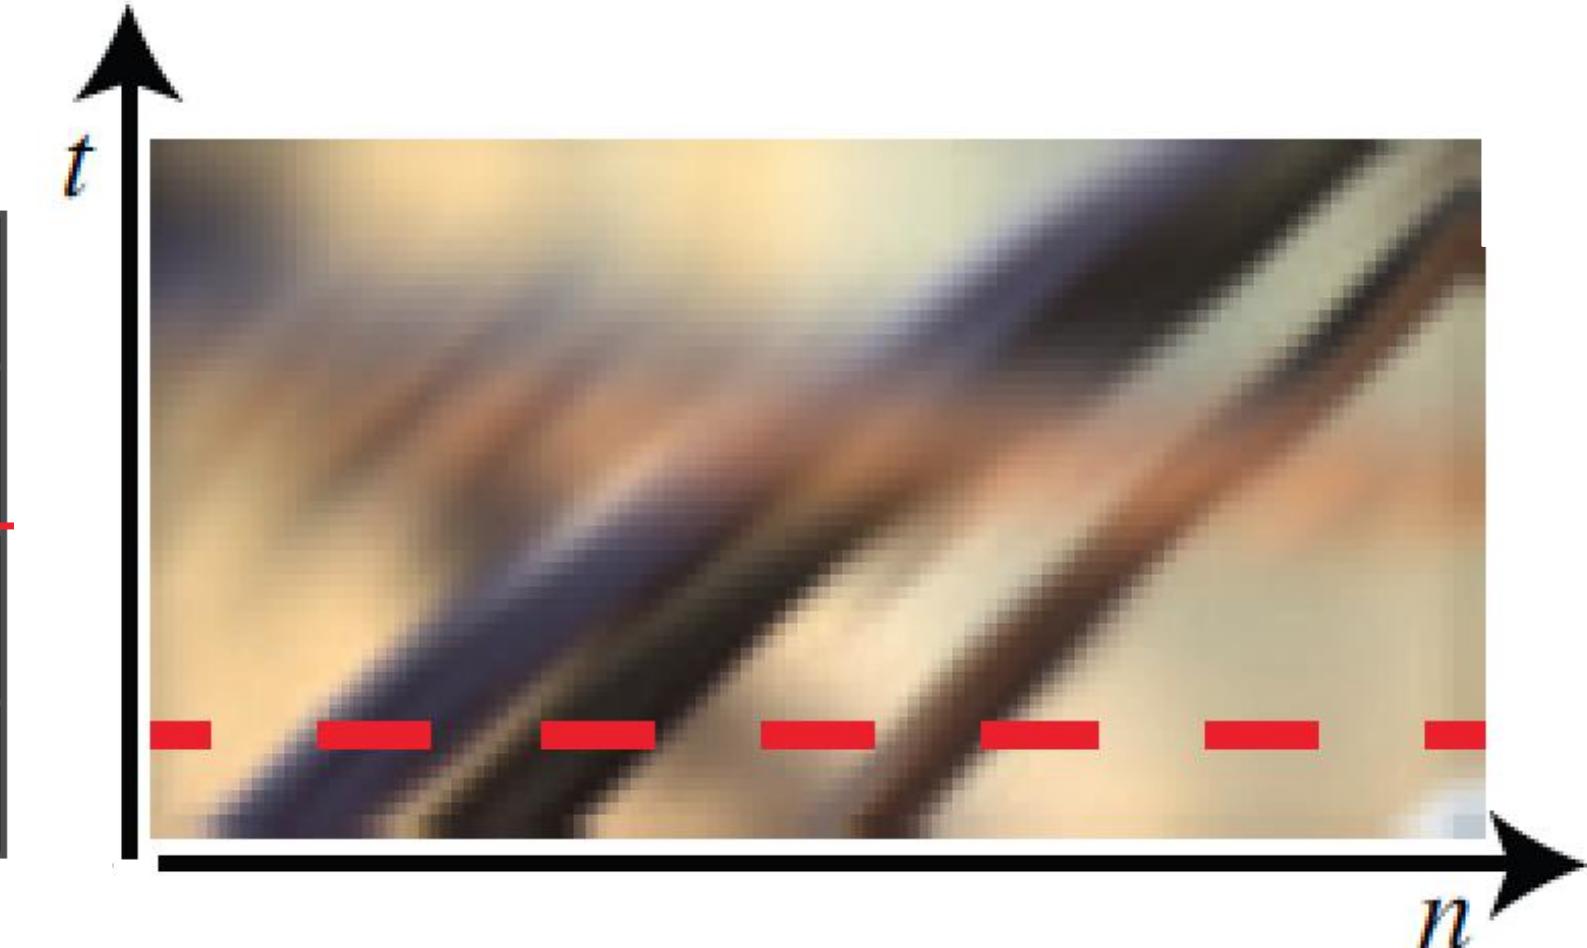
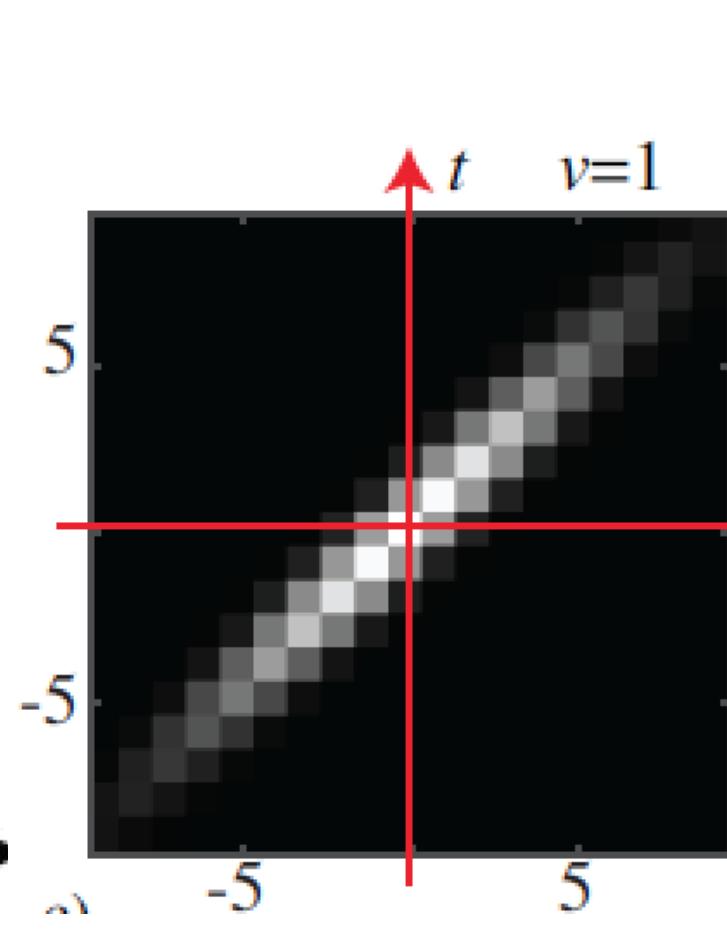
How could we create a filter that keeps sharp objects that move at some velocity (v_x , v_y) while blurring the rest?

$$g_{v_x, v_y}(x, y, t) = g(x - v_x t, y - v_y t, t)$$

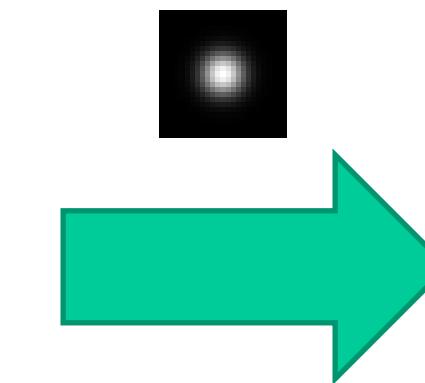
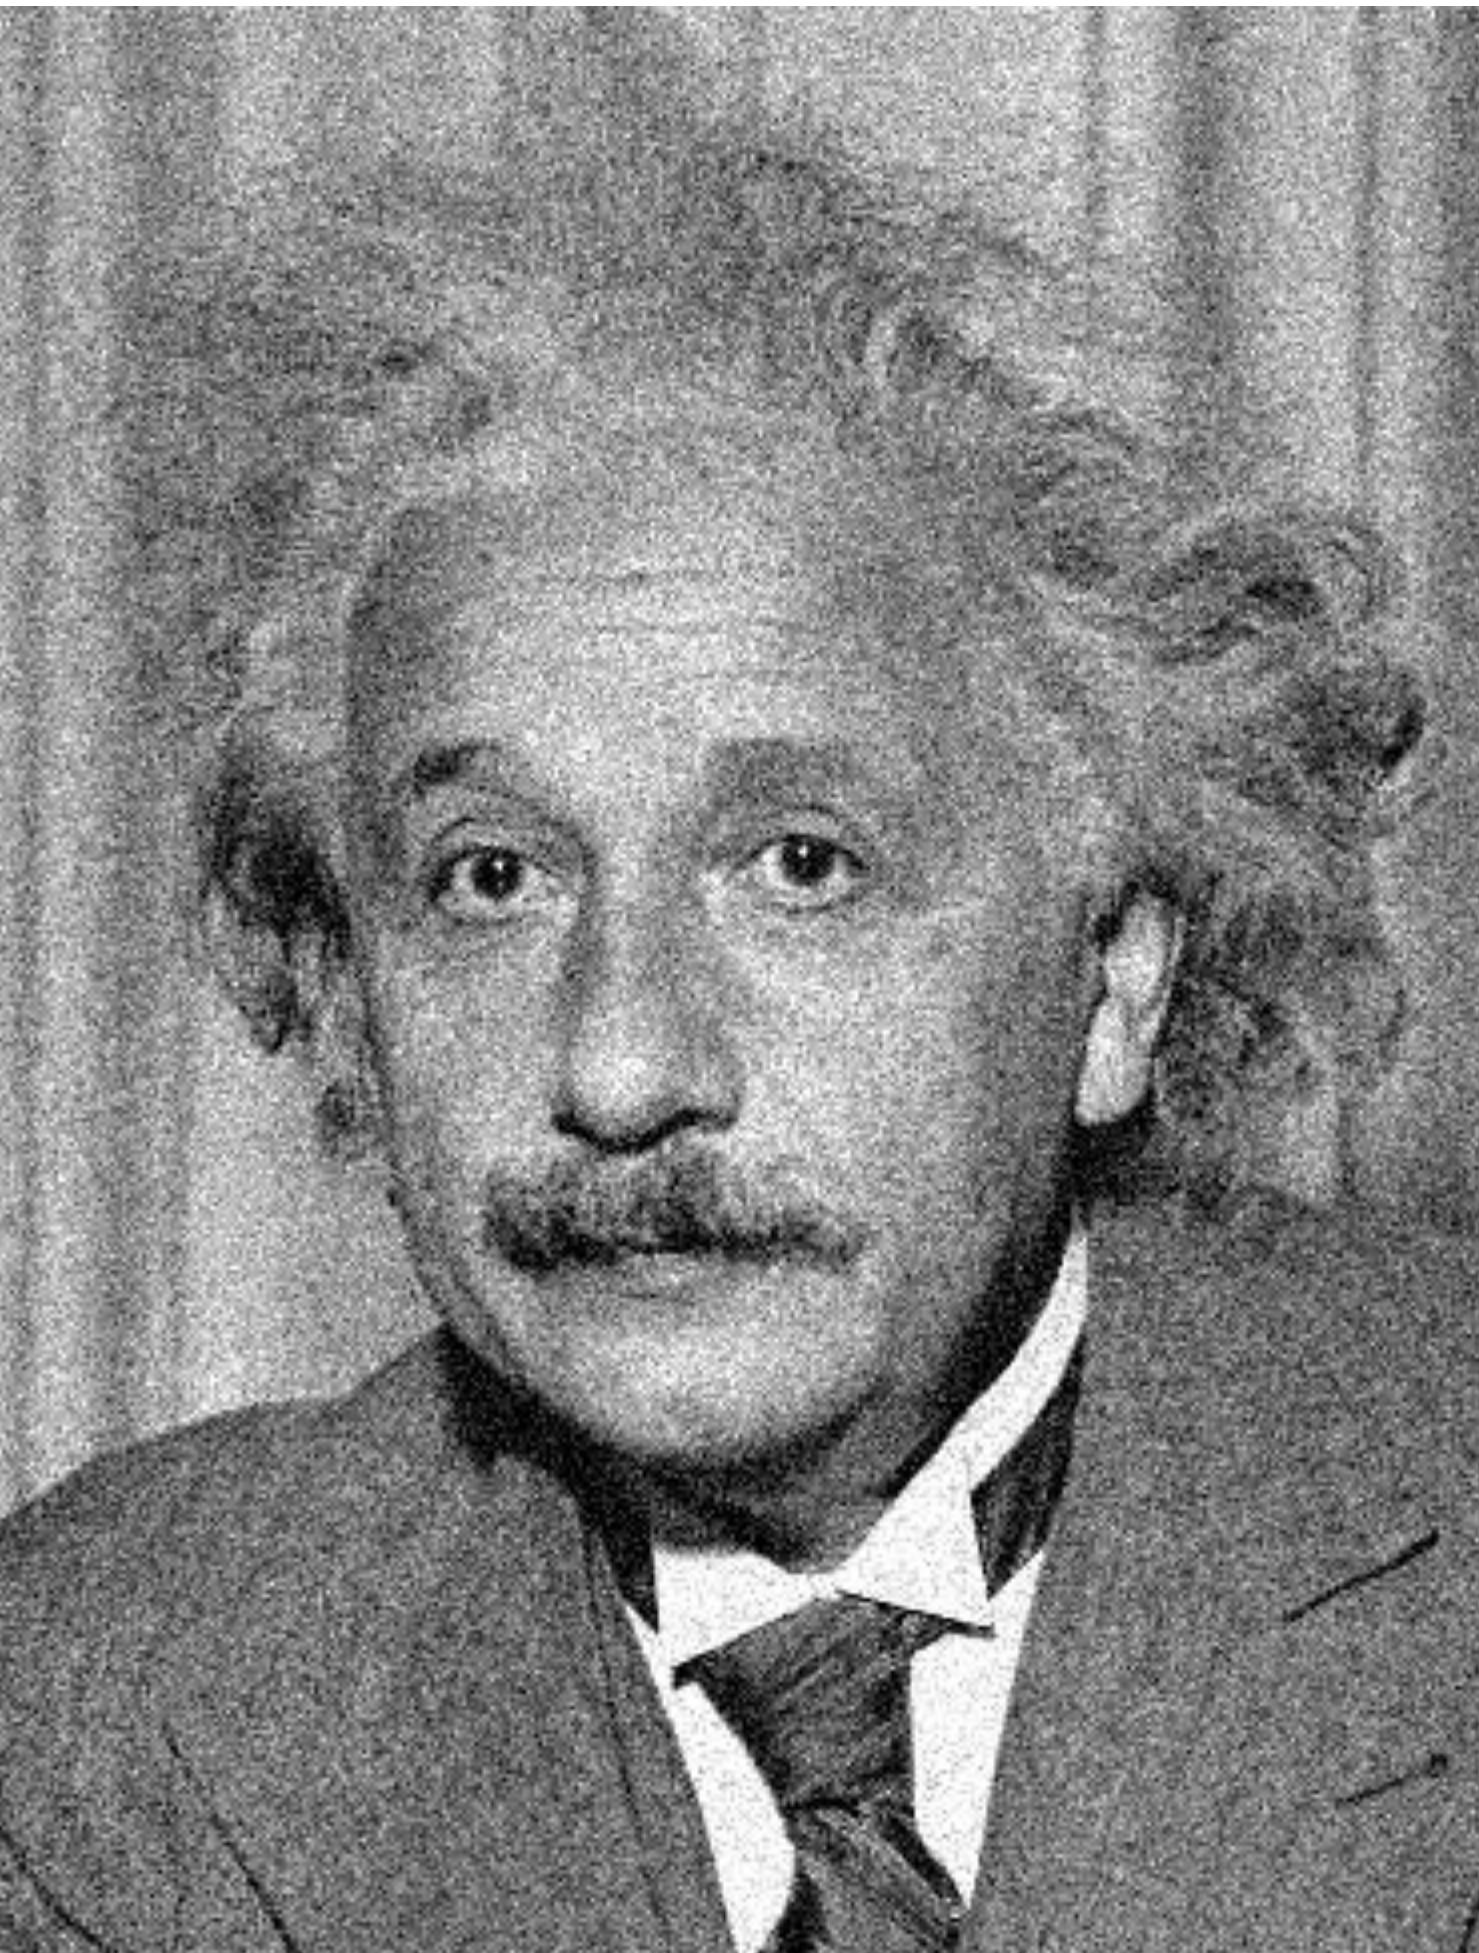




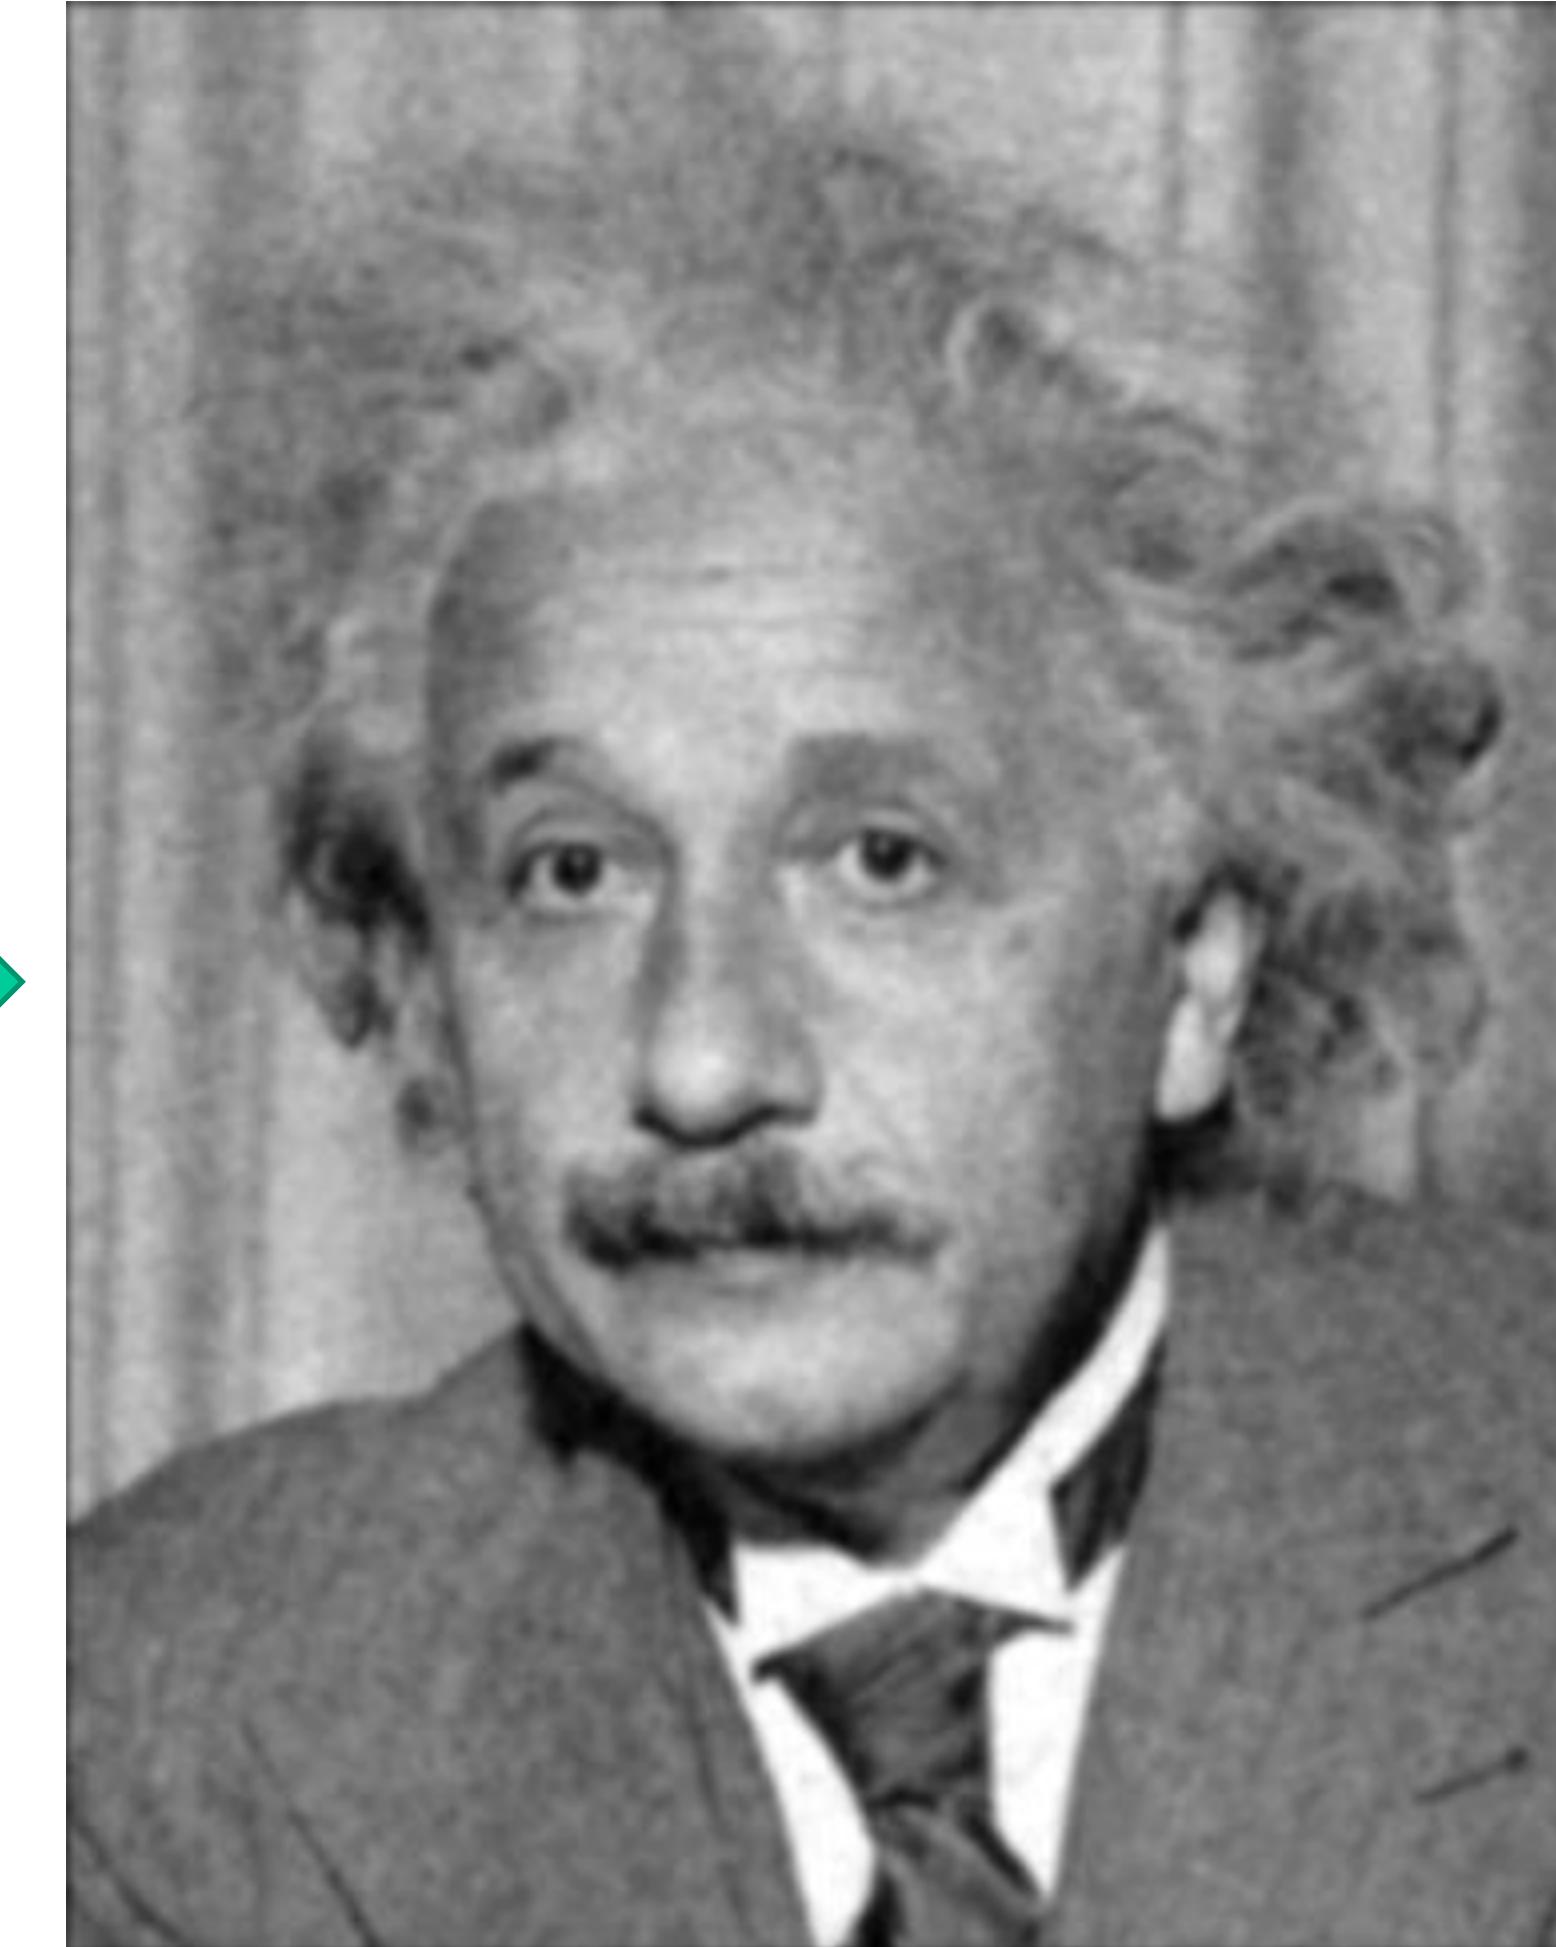




Non-linear Filters



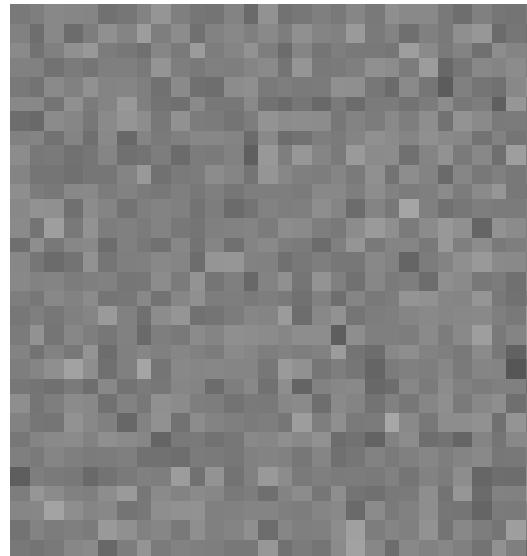
Gaussian
Filter



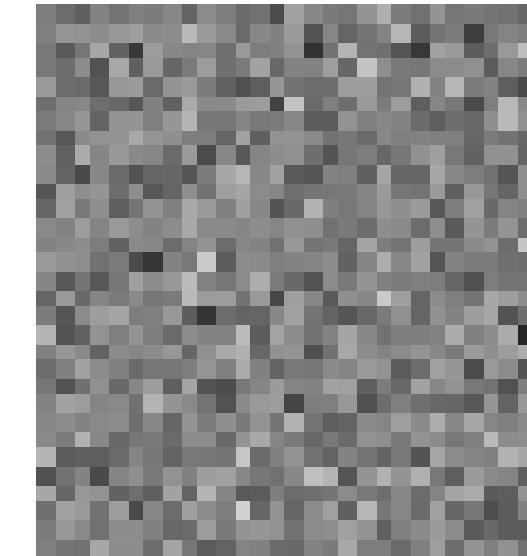
Additive Gaussian Noise

Reducing Gaussian noise

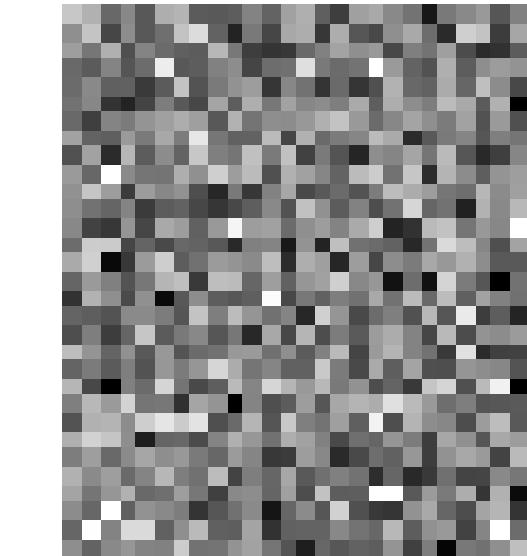
$\sigma=0.05$



$\sigma=0.1$



$\sigma=0.2$



no
smoothing



$\sigma=1$ pixel



$\sigma=2$ pixels



Smoothing with larger standard deviations suppresses noise,
but also blurs the image

Source: S. Lazebnik

Reducing salt-and-pepper noise by Gaussian smoothing

3x3



5x5

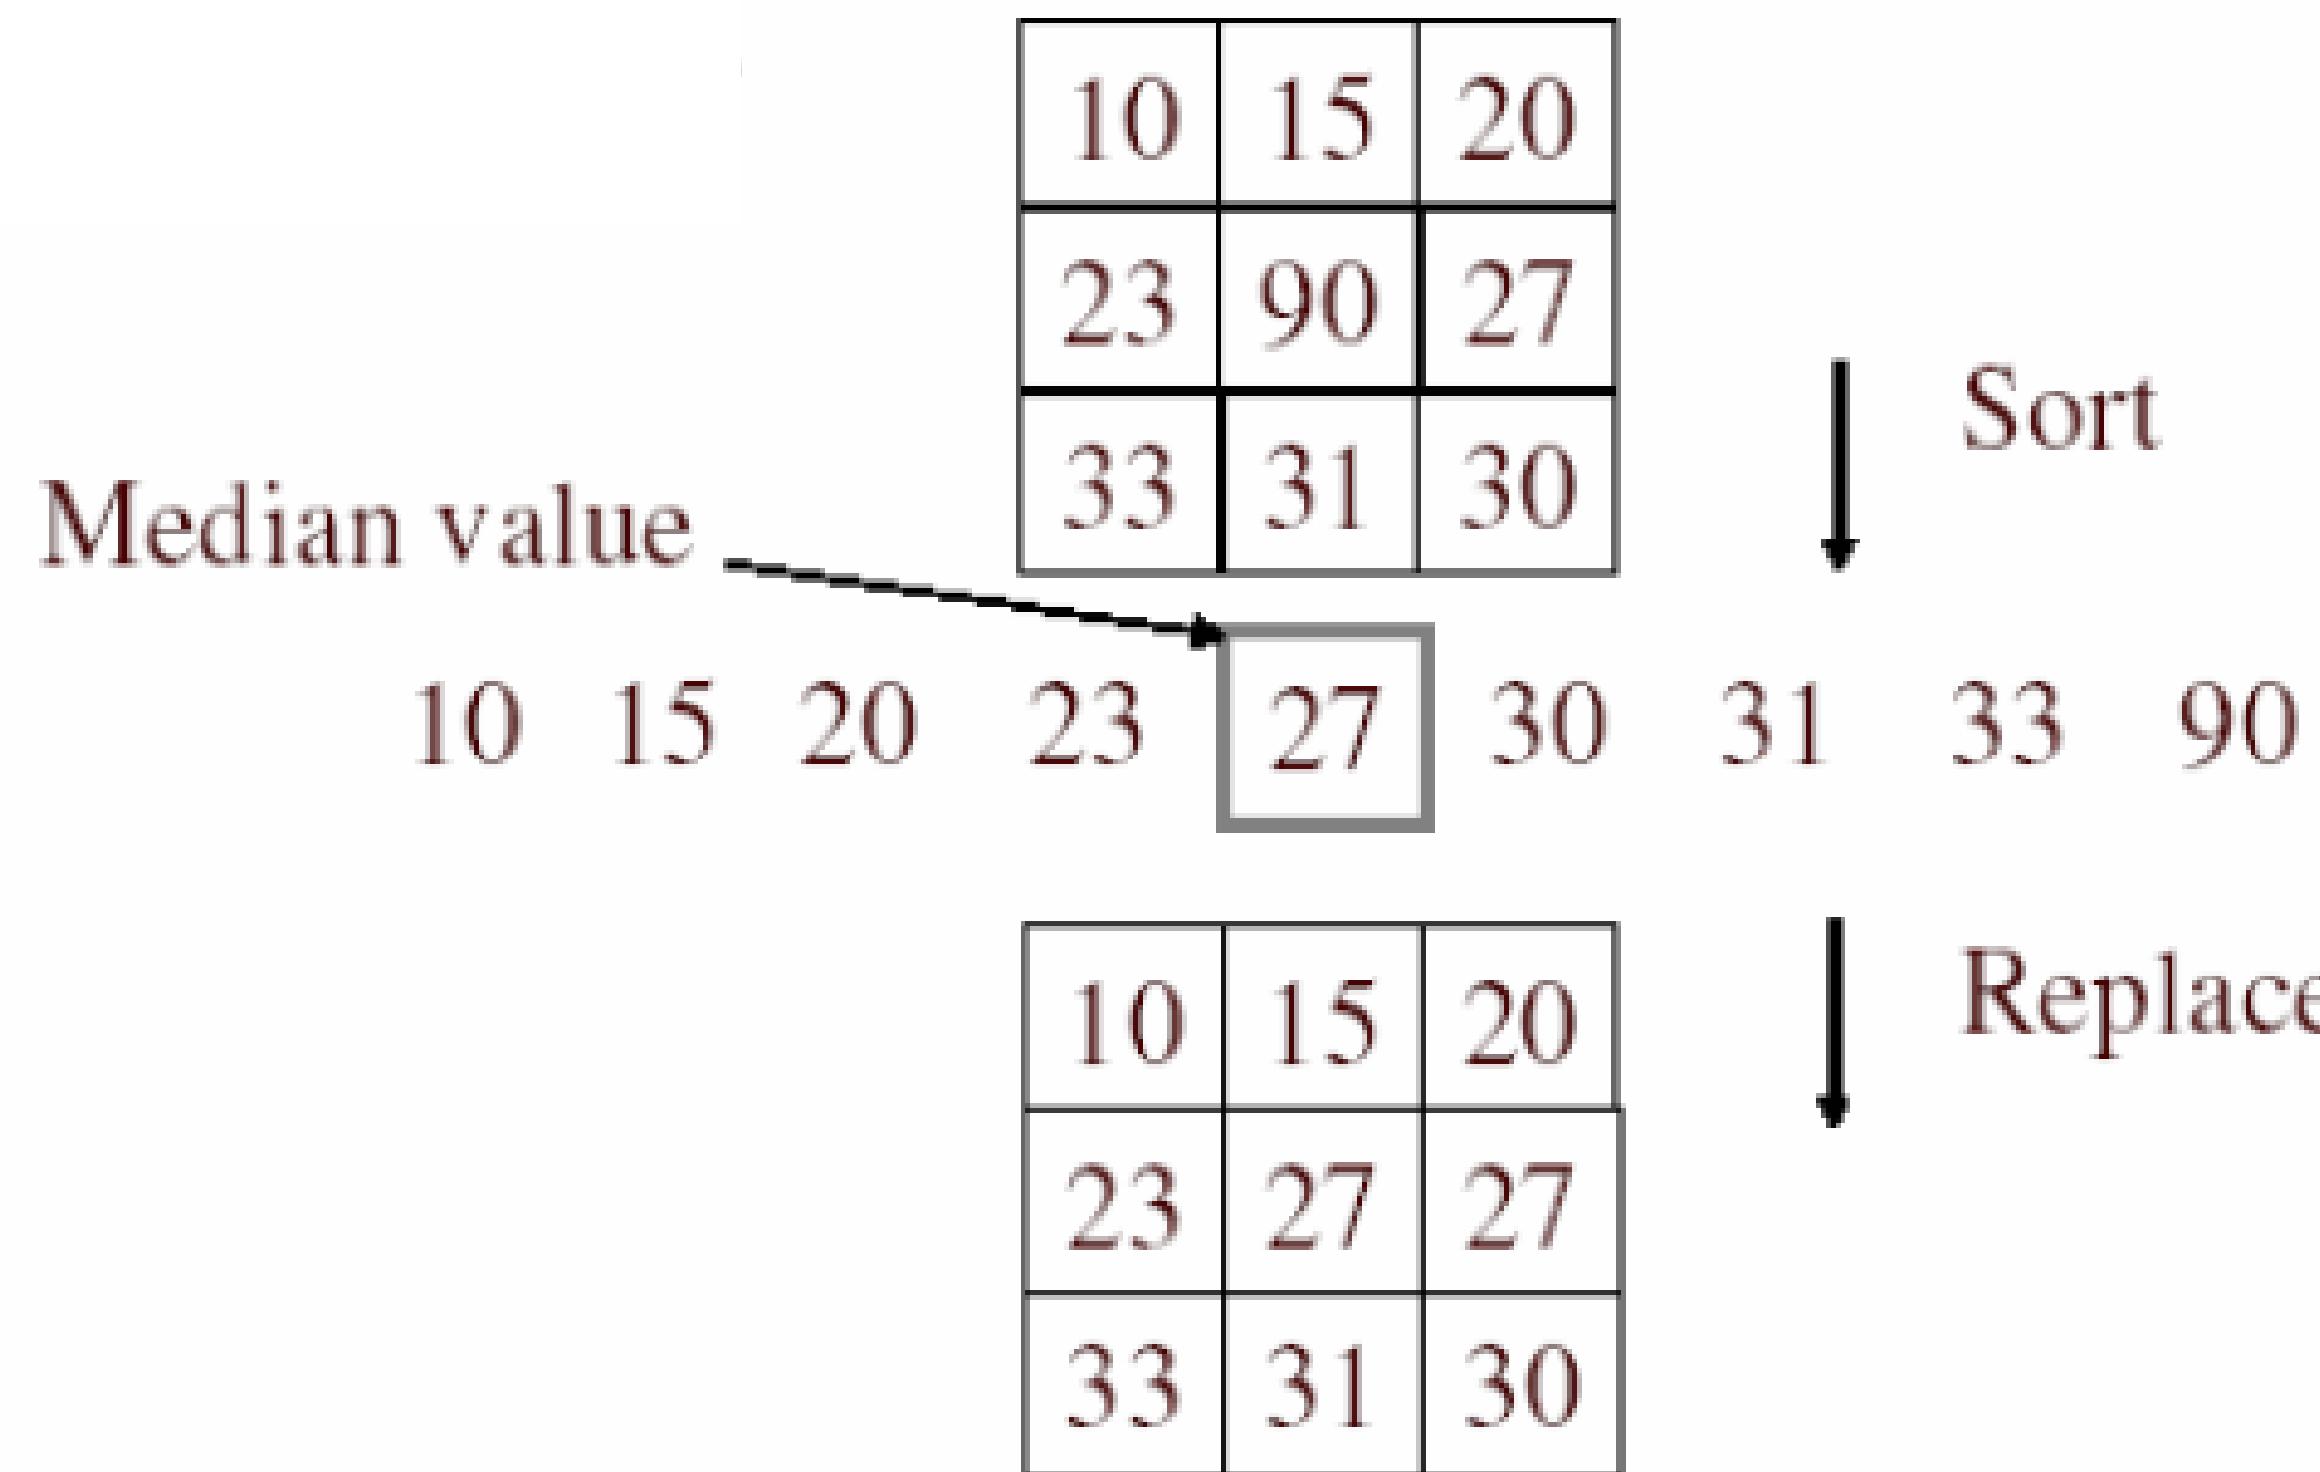


7x7



Alternative idea: Median filtering

A **median filter** operates over a window by selecting the median intensity in the window



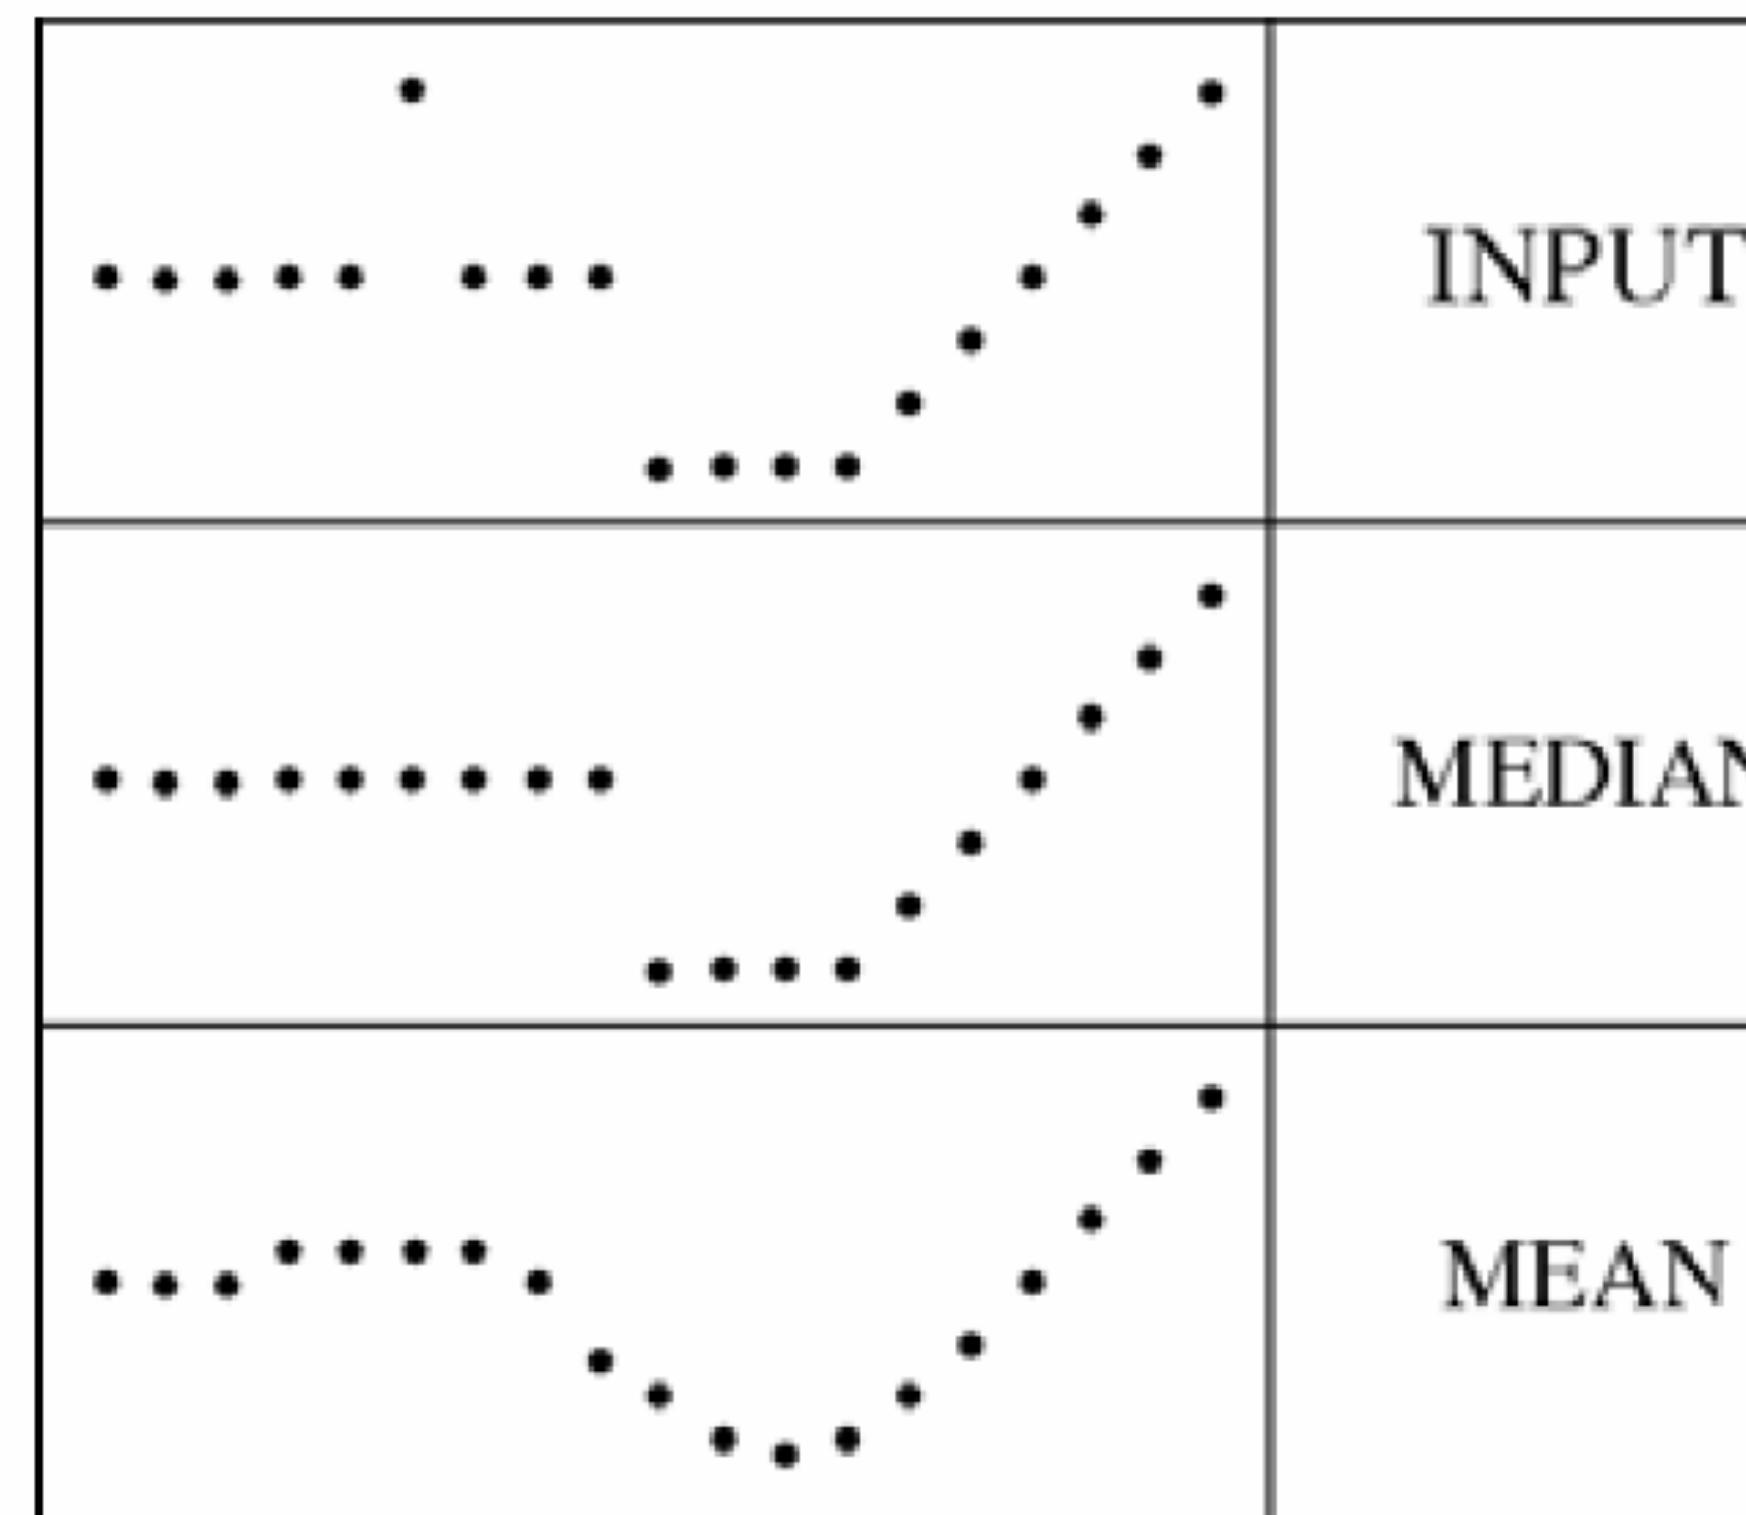
- Is median filtering linear?

Median filter

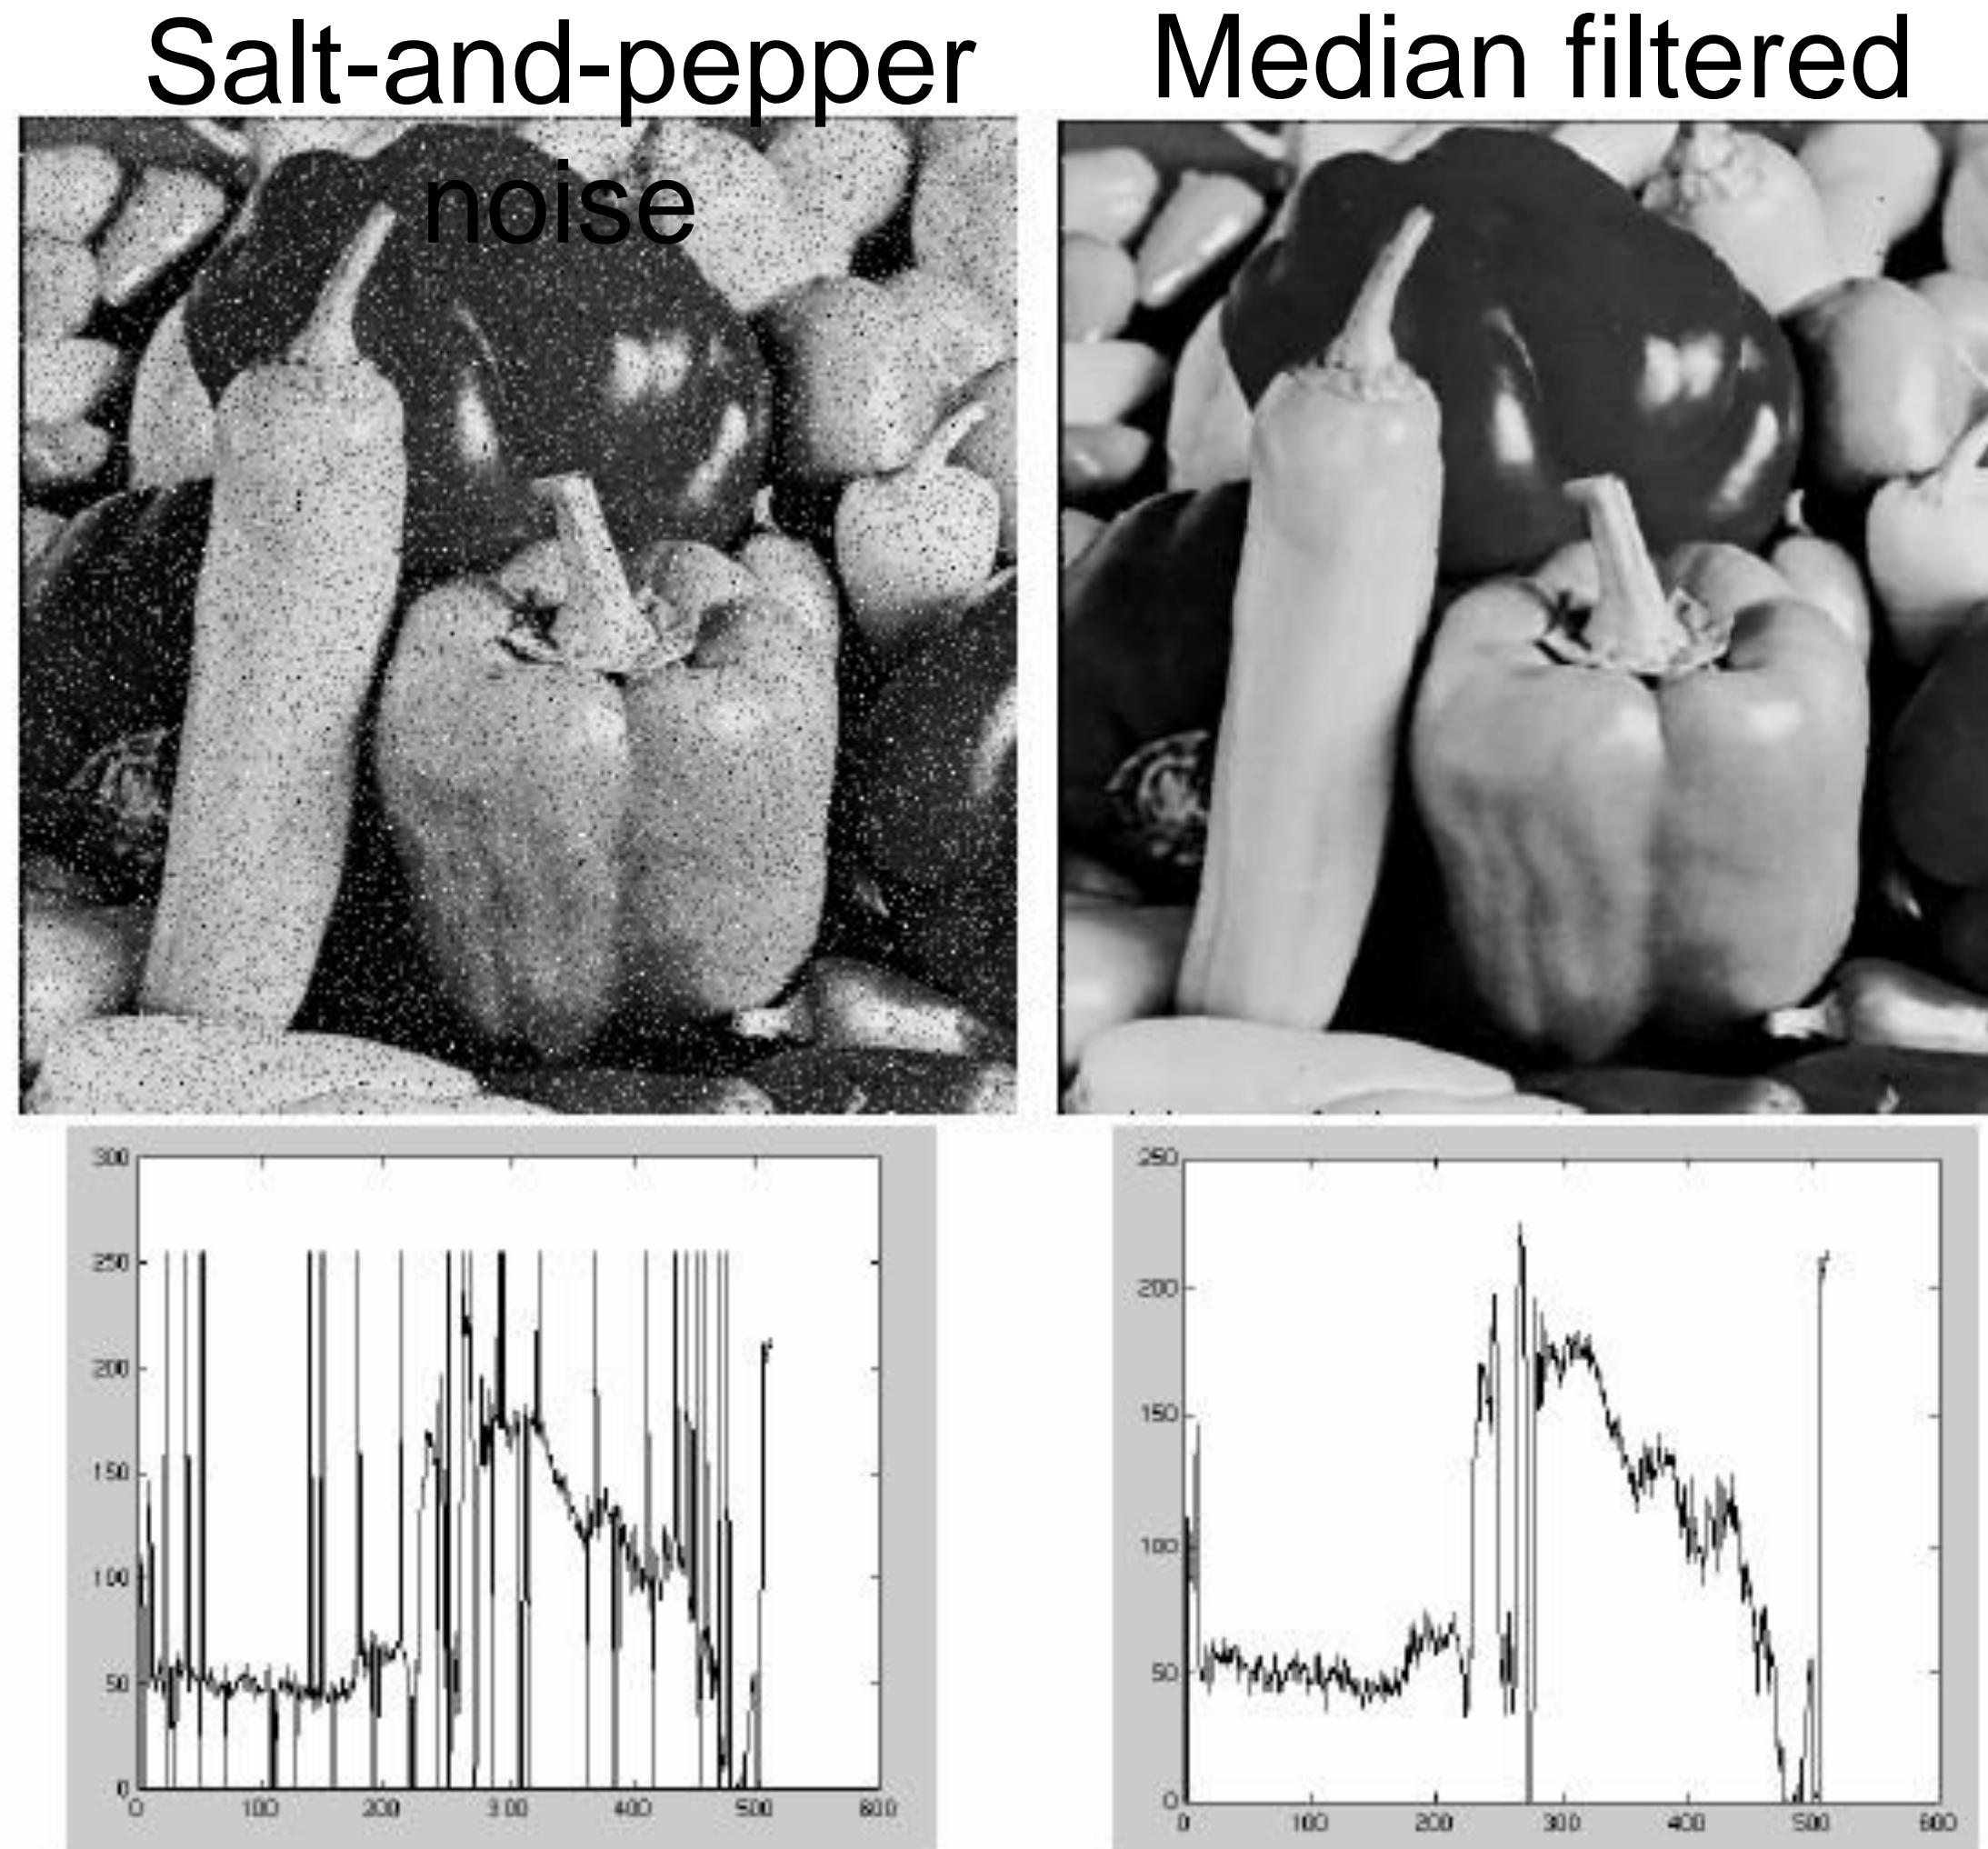
What advantage does median filtering have over Gaussian filtering?

- Robustness to outliers

filters have width 5 :



Median filter



MATLAB: `medfilt2(image, [h w])`

Source: M. Hebert

Median vs. Gaussian filtering

3x3



5x5



7x7



Gaussian

Median



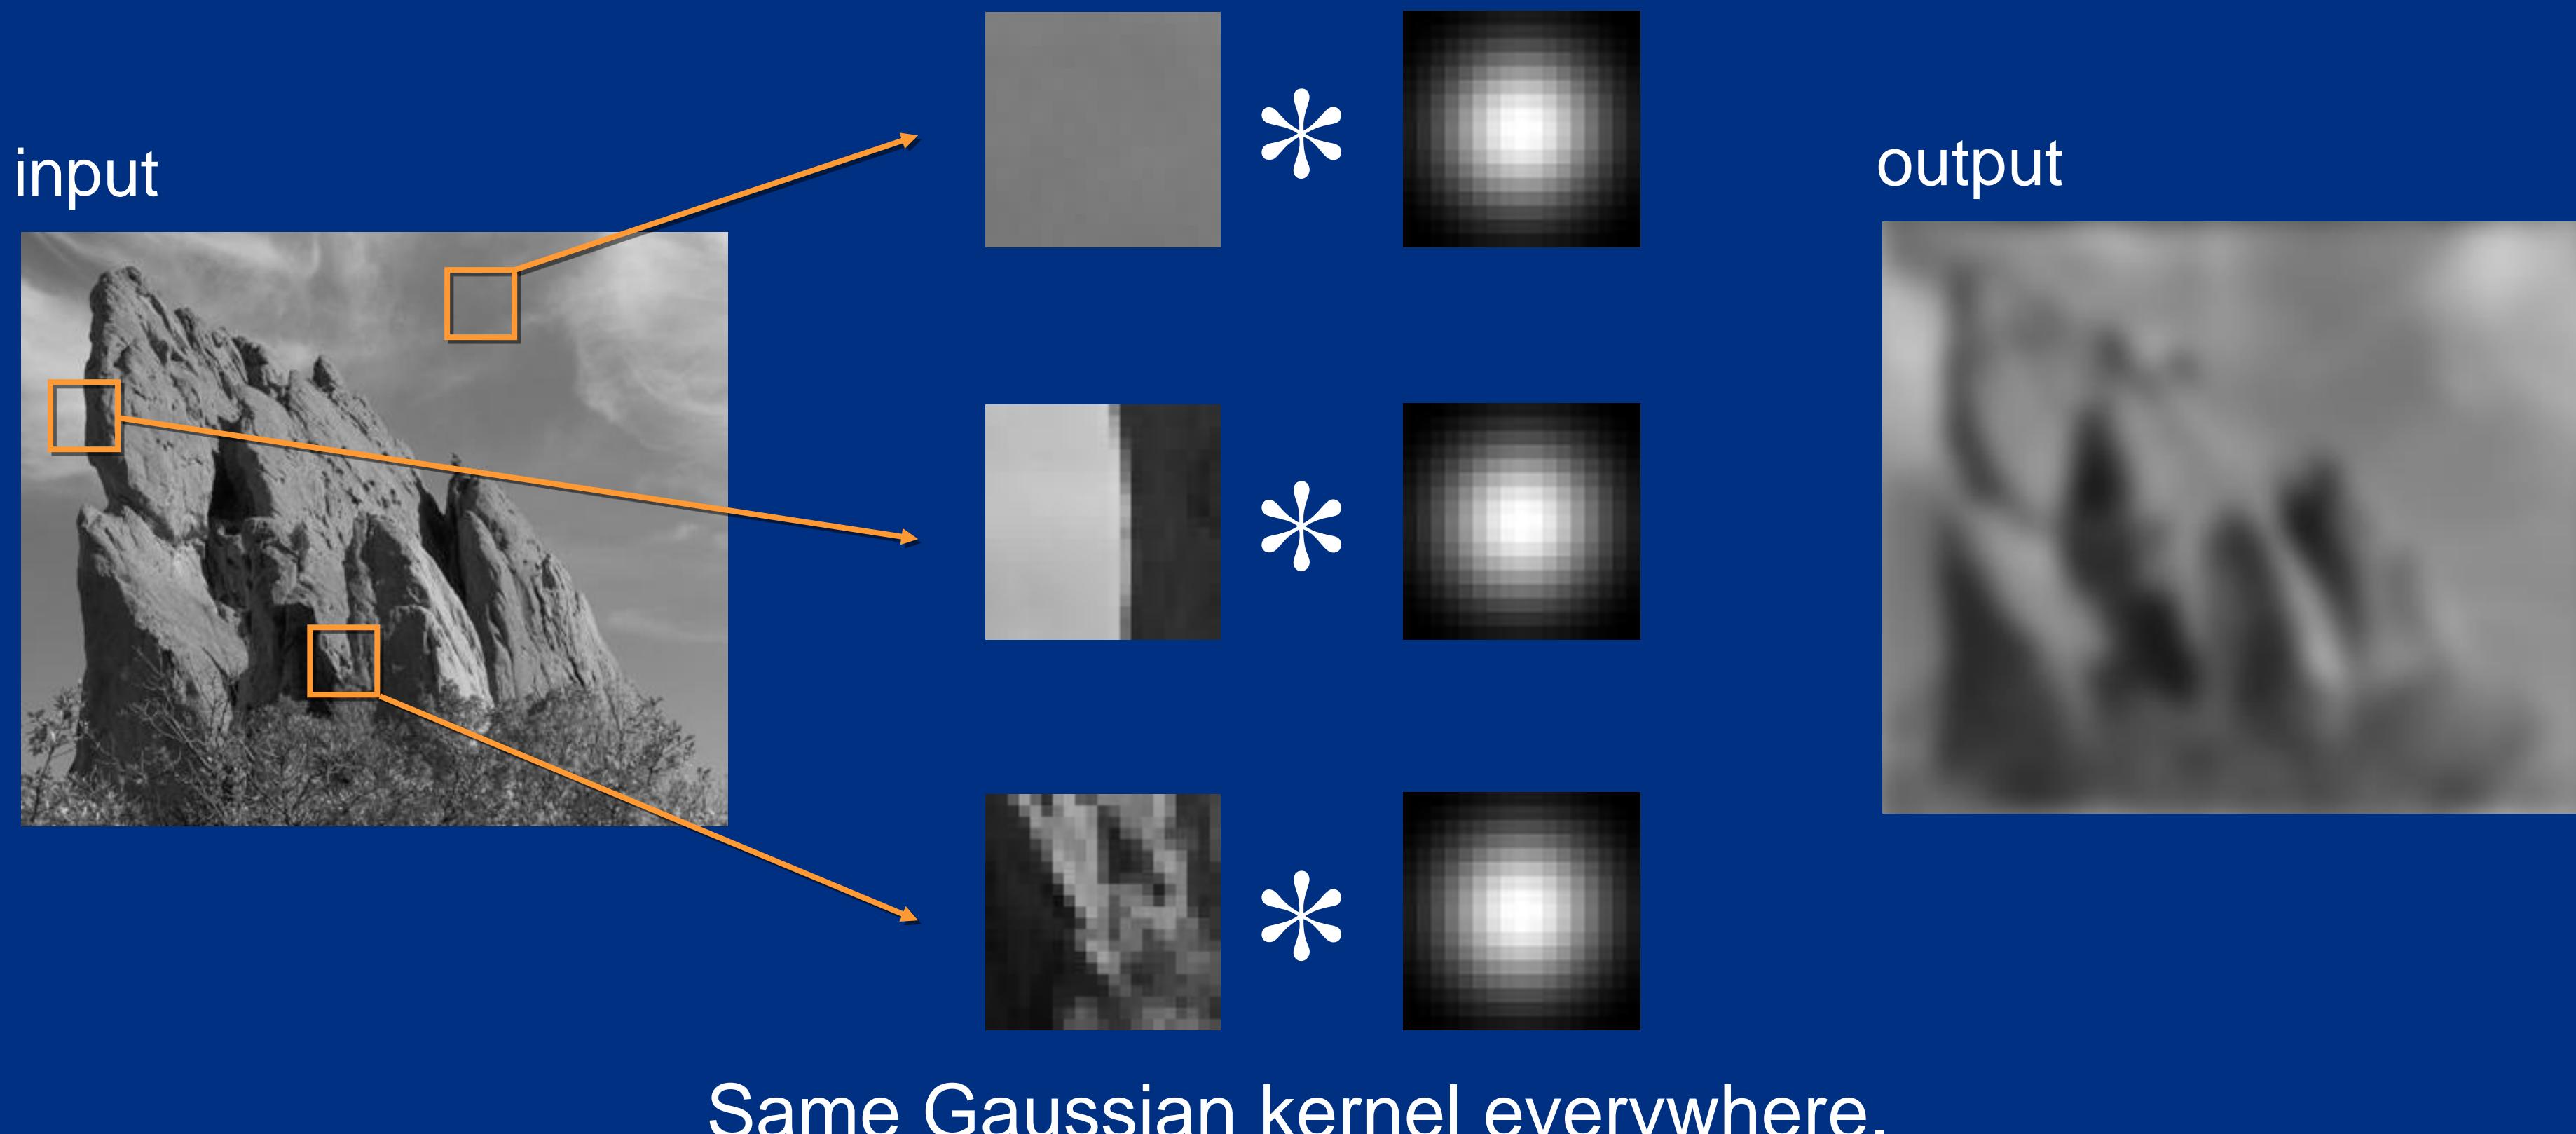
A Gentle Introduction to Bilateral Filtering and its Applications



“Fixing the Gaussian Blur”: the Bilateral Filter

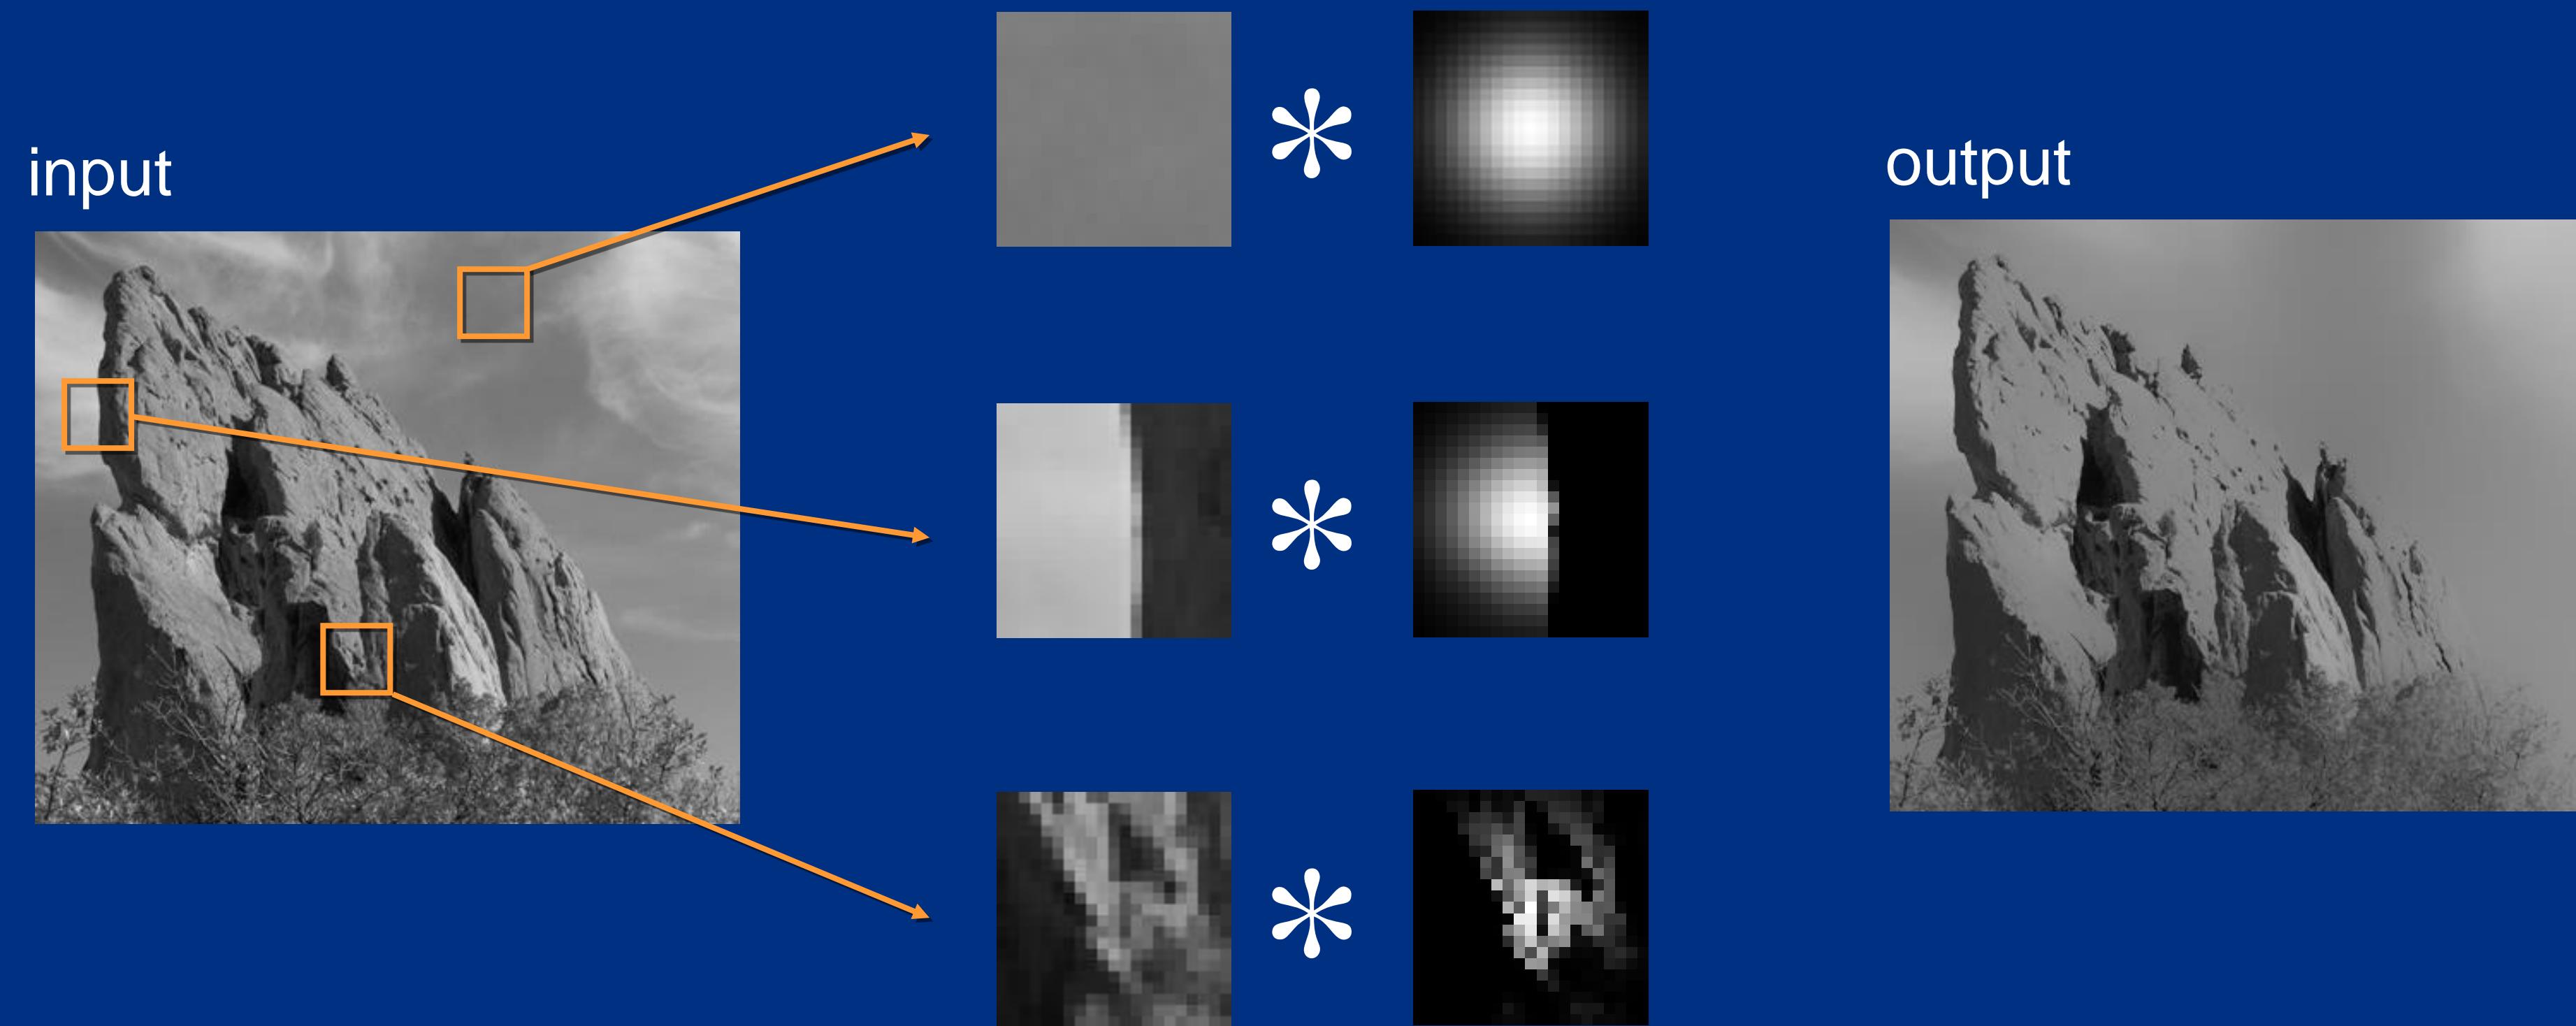
Sylvain Paris – MIT CSAIL

Blur Comes from Averaging across Edges



Bilateral Filter No Averaging across Edges

[Aurich 95, Smith 97, Tomasi 98]



The kernel shape depends on the image content.

Bilateral Filter Definition: an Additional Edge Term

Same idea: **weighted average of pixels.**

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

new
not new
new

normalization factor space weight range weight

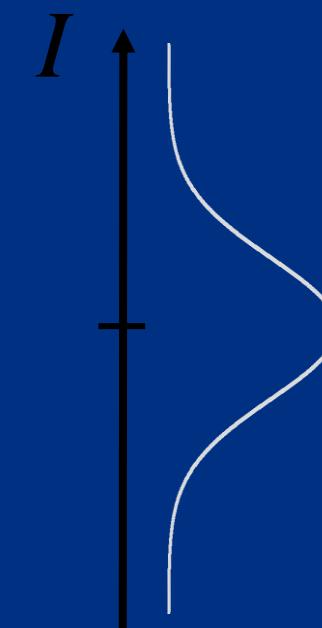
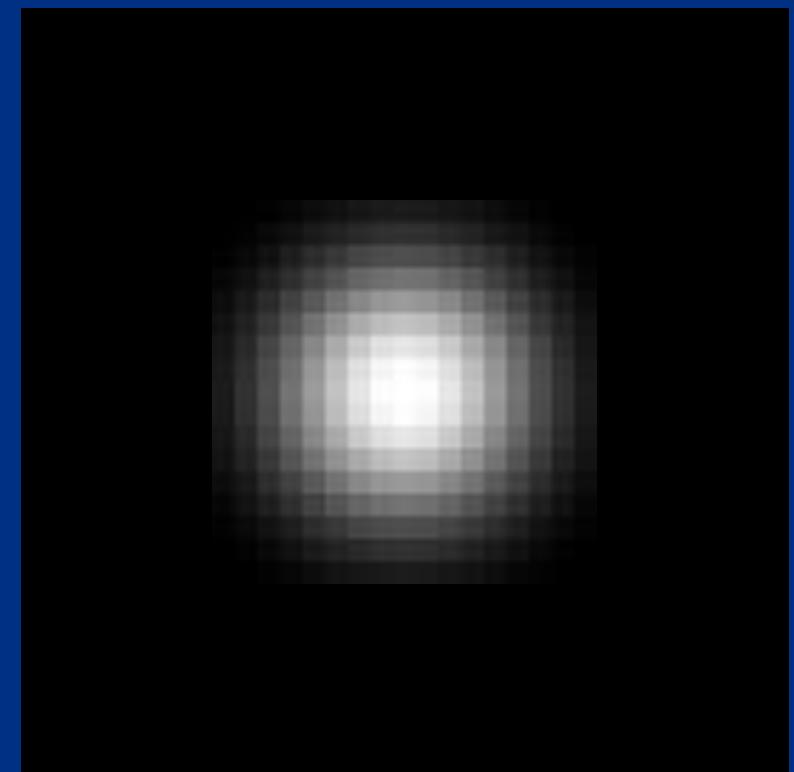
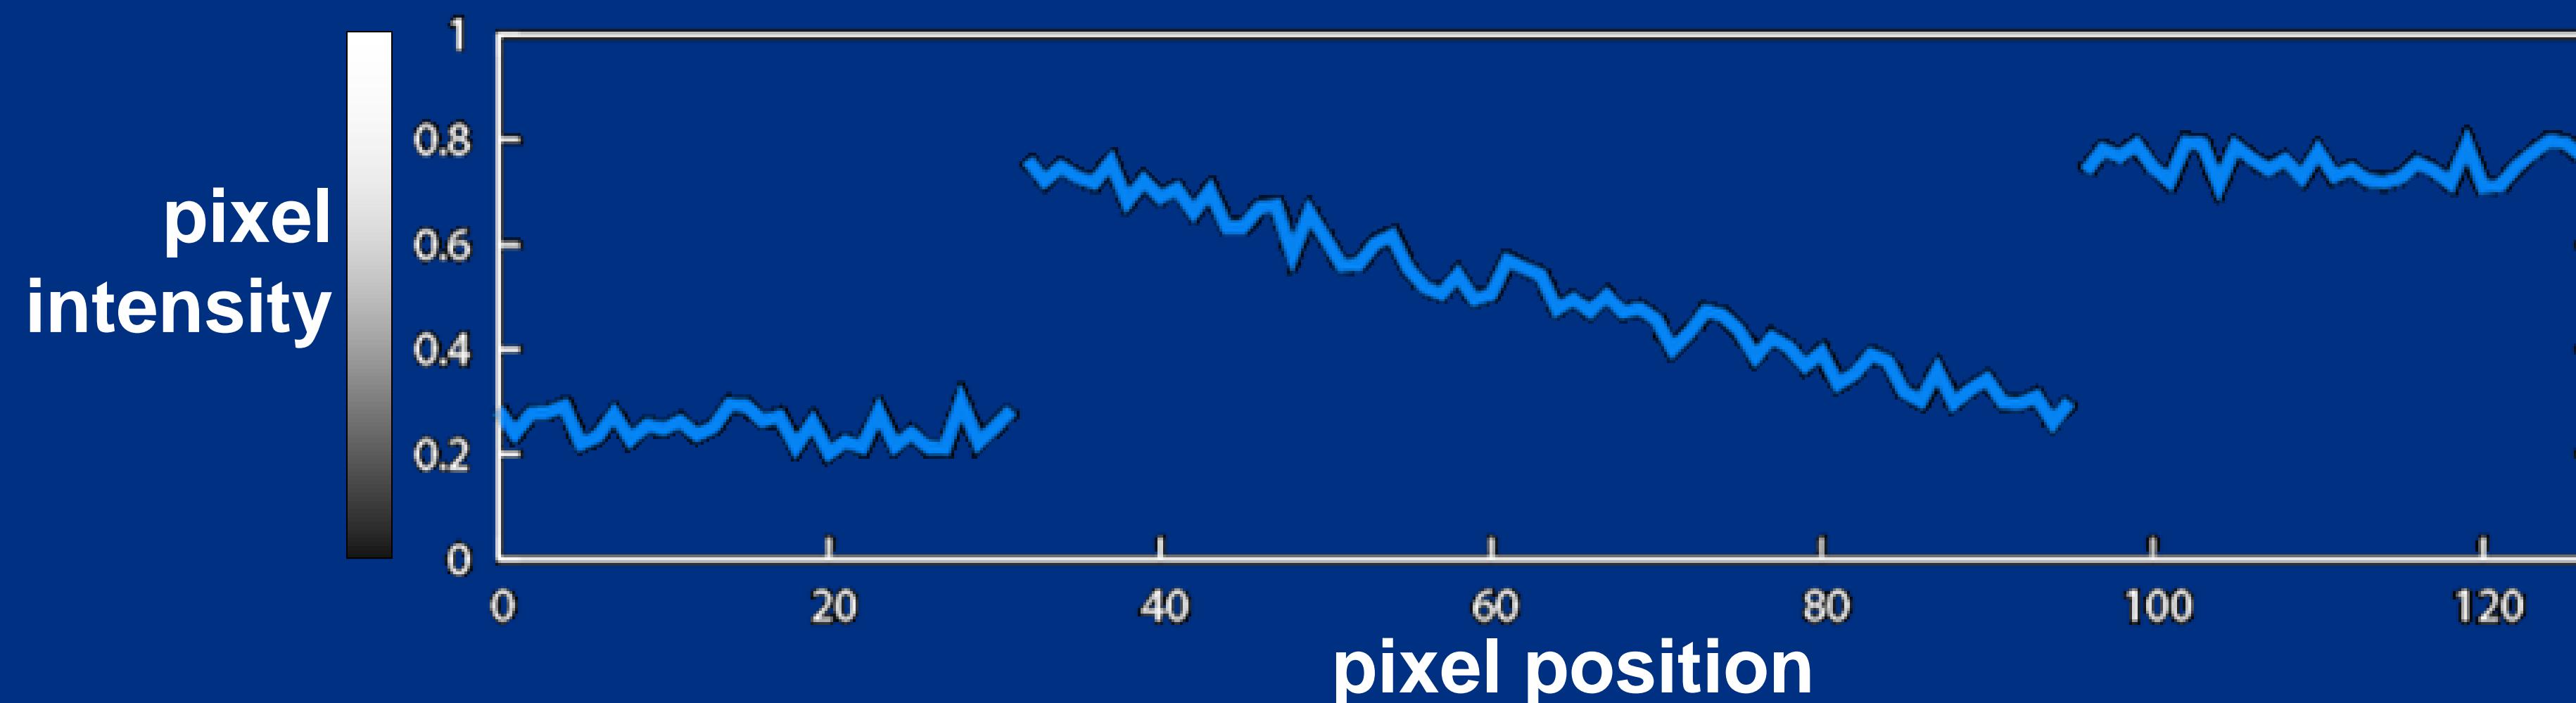


Illustration a 1D Image

- 1D image = line of pixels

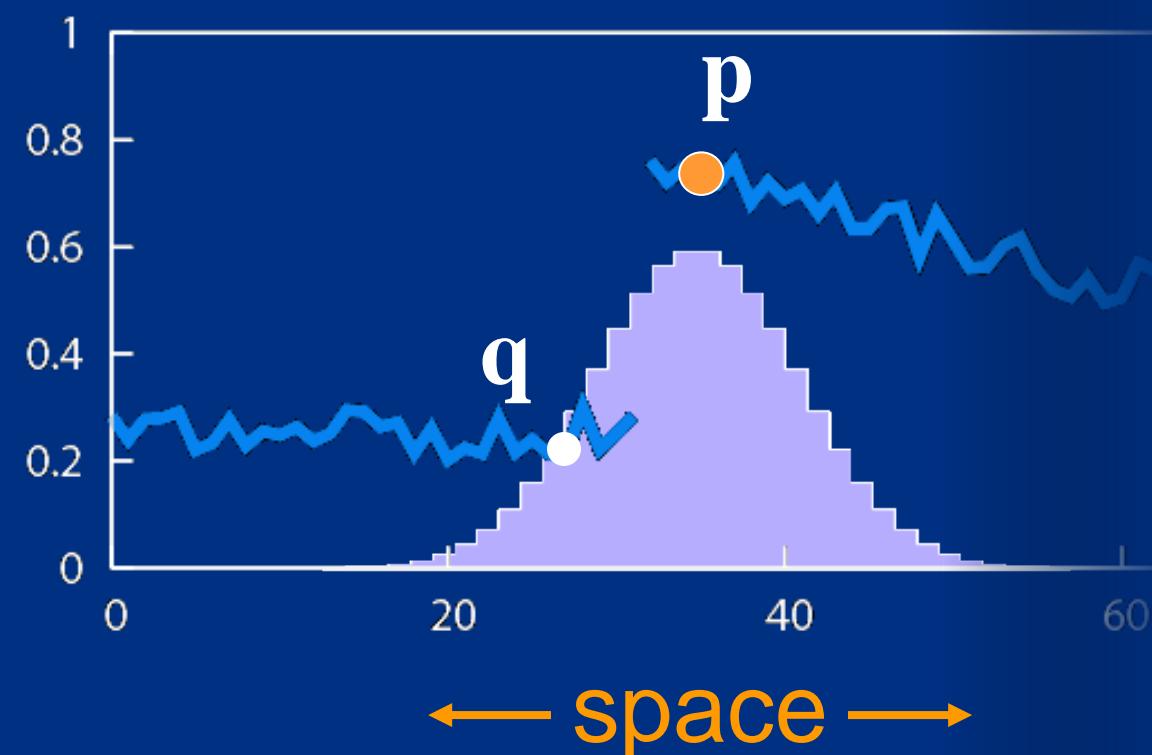


- Better visualized as a plot



Gaussian Blur and Bilateral Filter

Gaussian blur

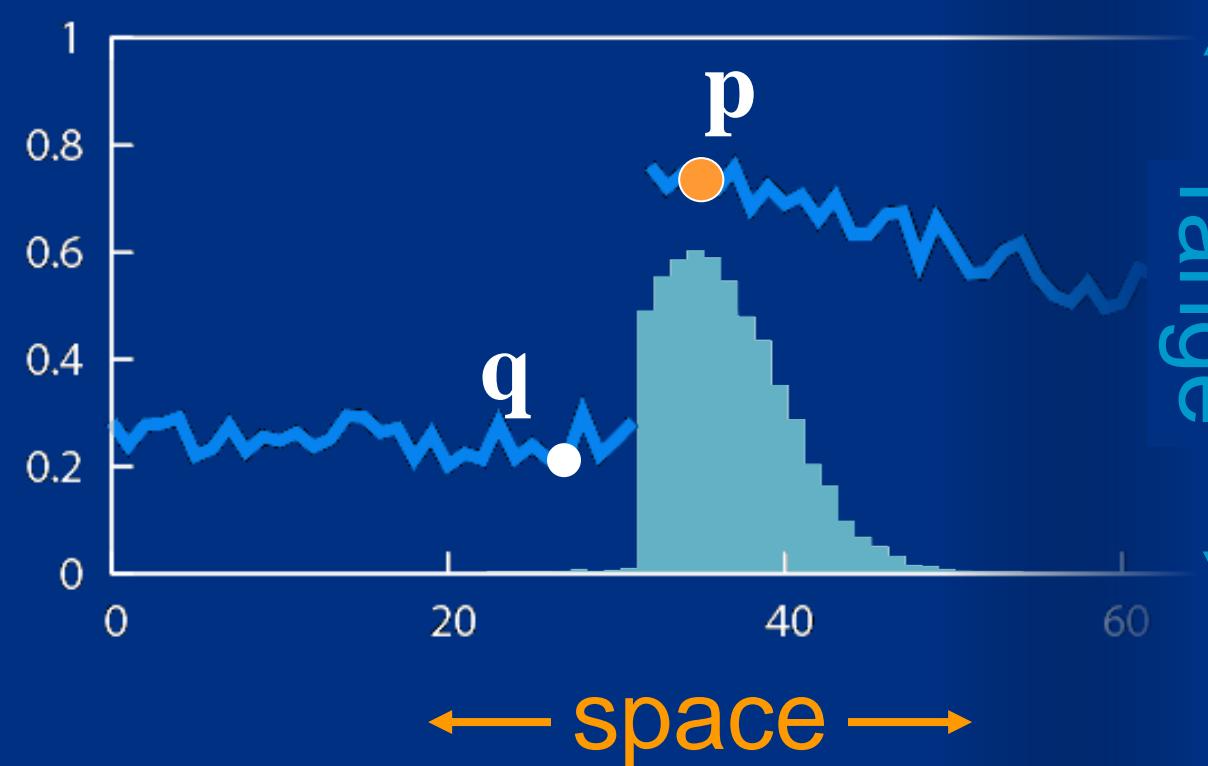


$$GB[I]_p = \sum_{q \in S} G_\sigma(\| p - q \|) I_q$$

space

Bilateral filter

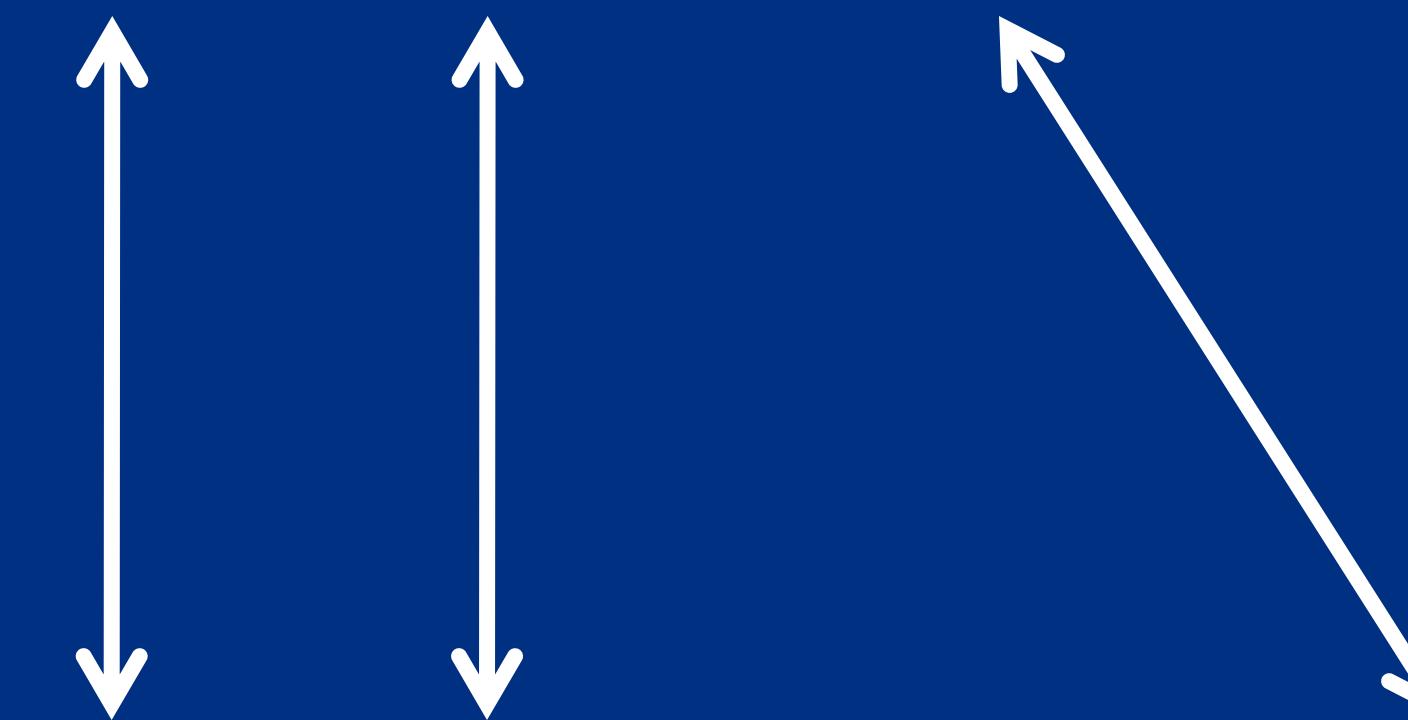
[Aurich 95, Smith 97, Tomasi 98]



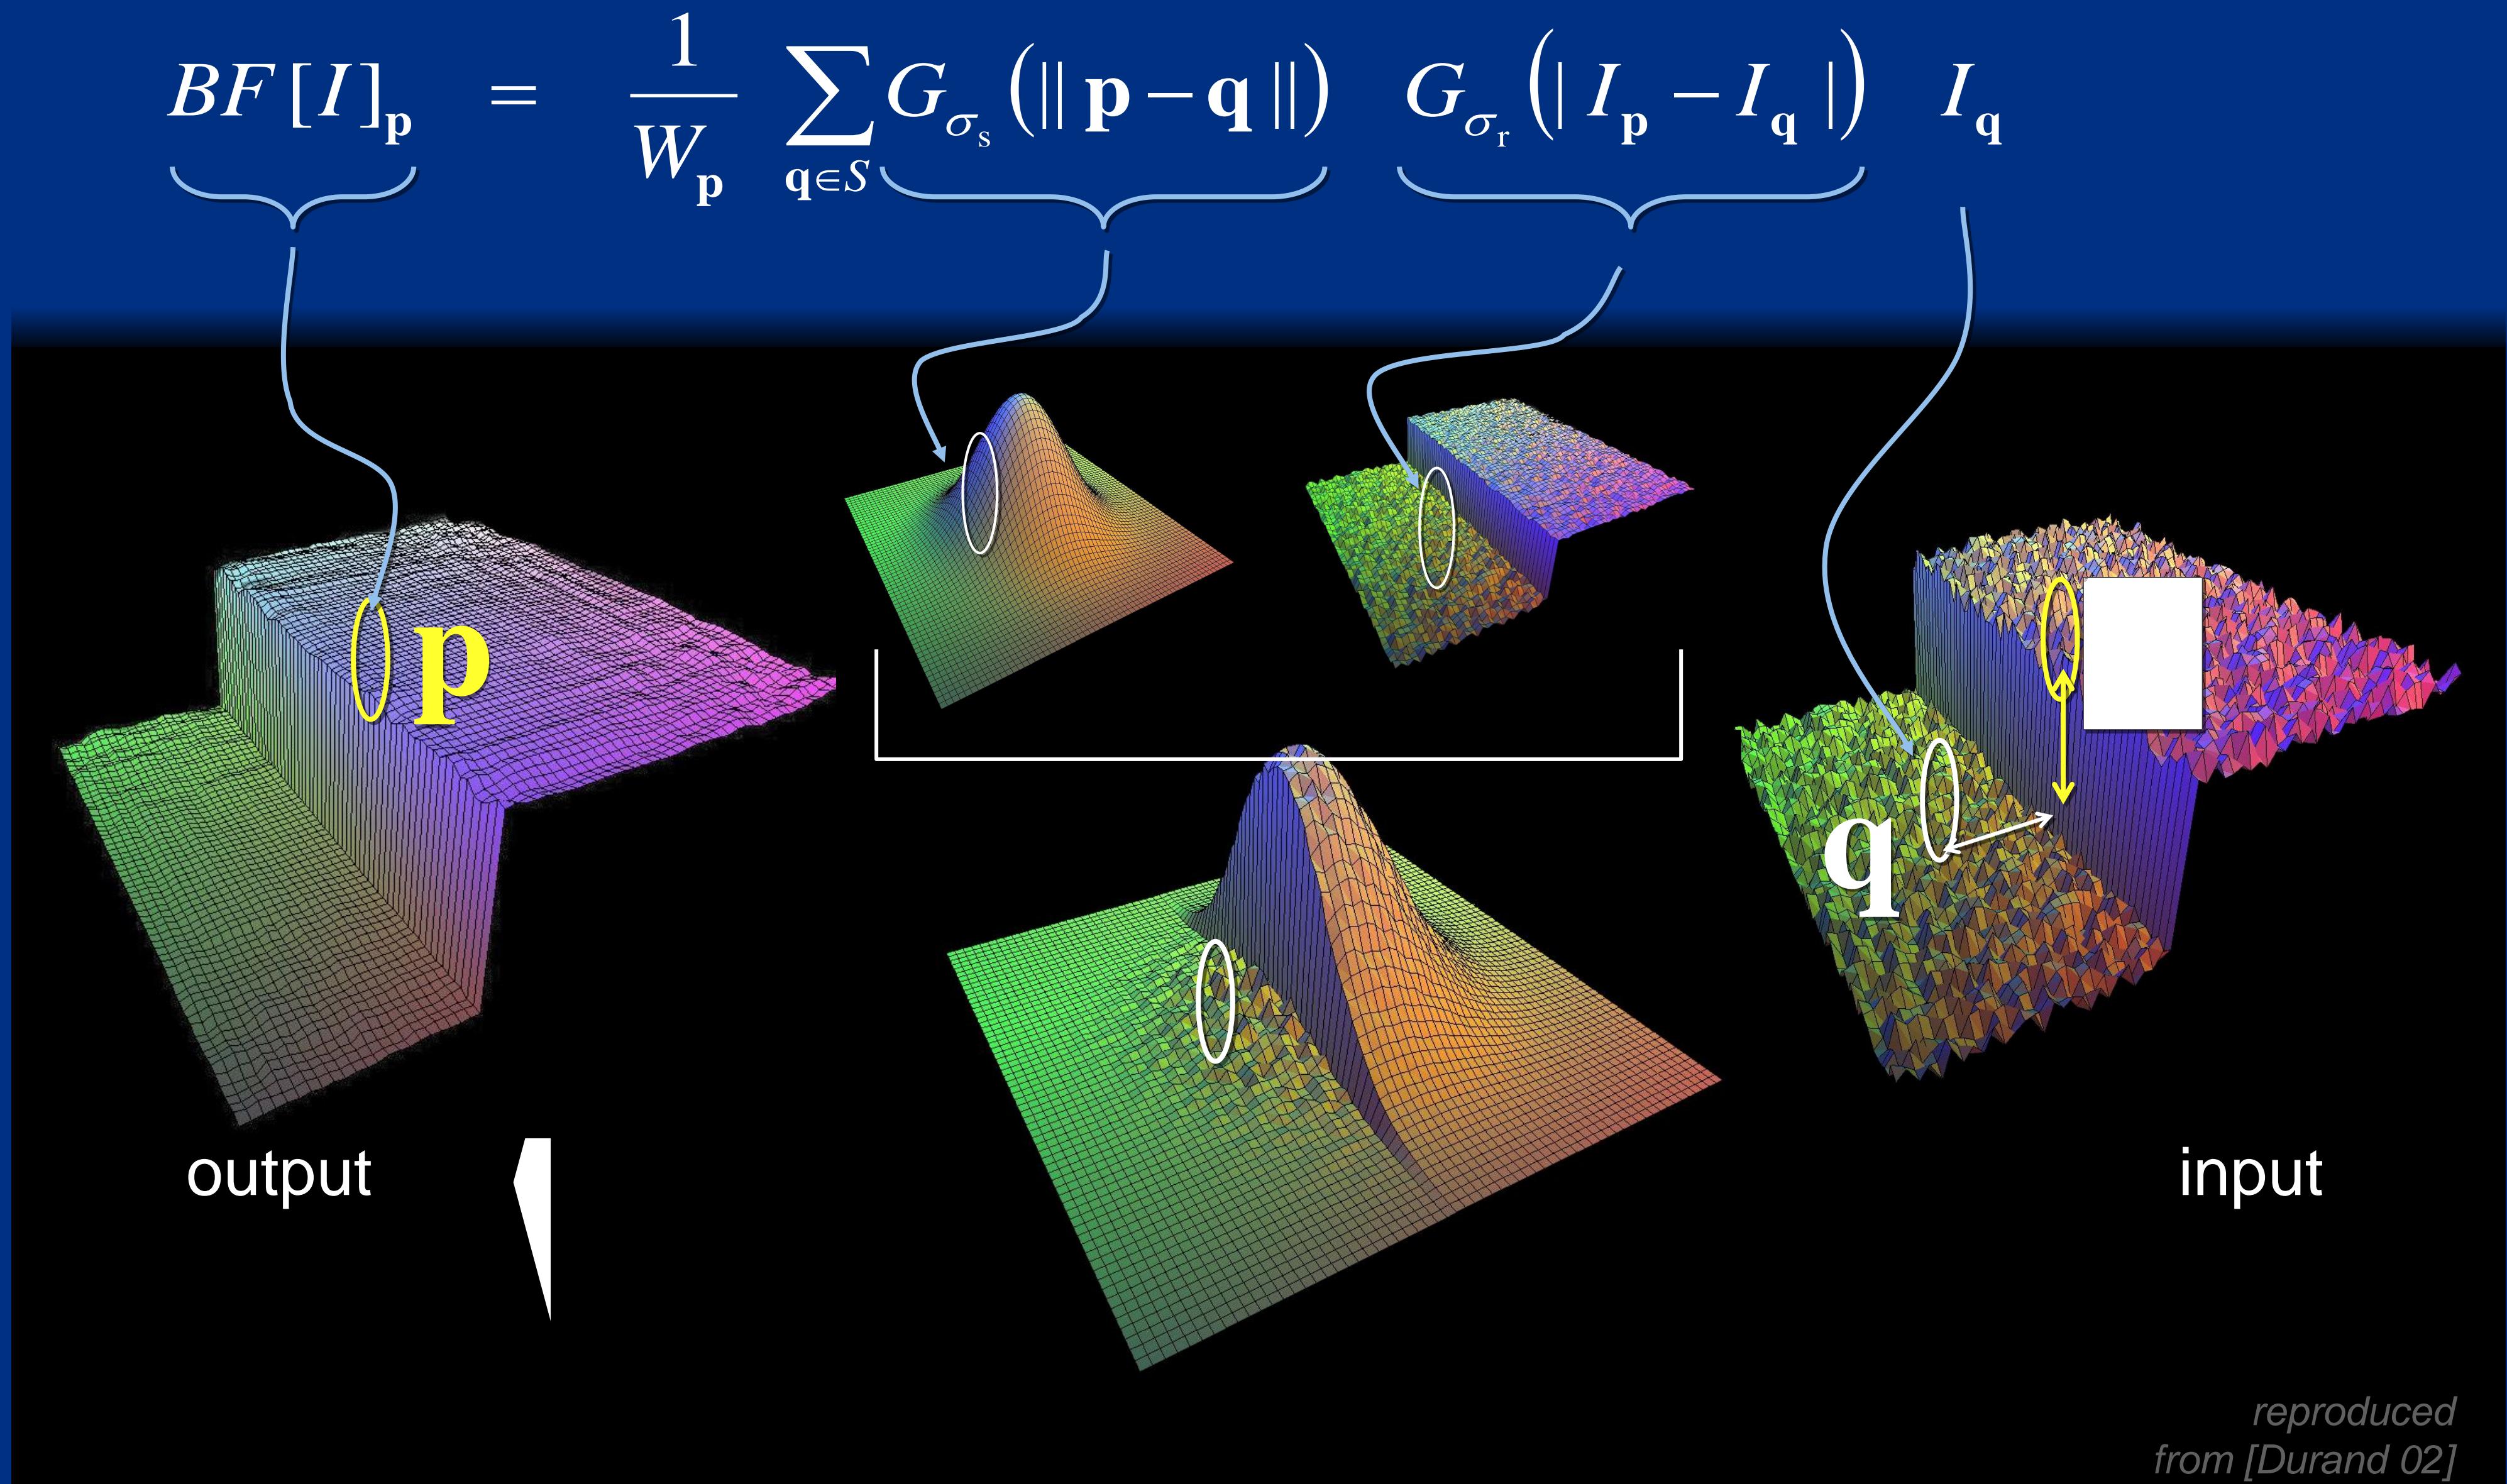
$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\| p - q \|) G_{\sigma_r}(|I_p - I_q|) I_q$$

space range

normalization



Bilateral Filter on a Height Field



reproduced
from [Durand 02]

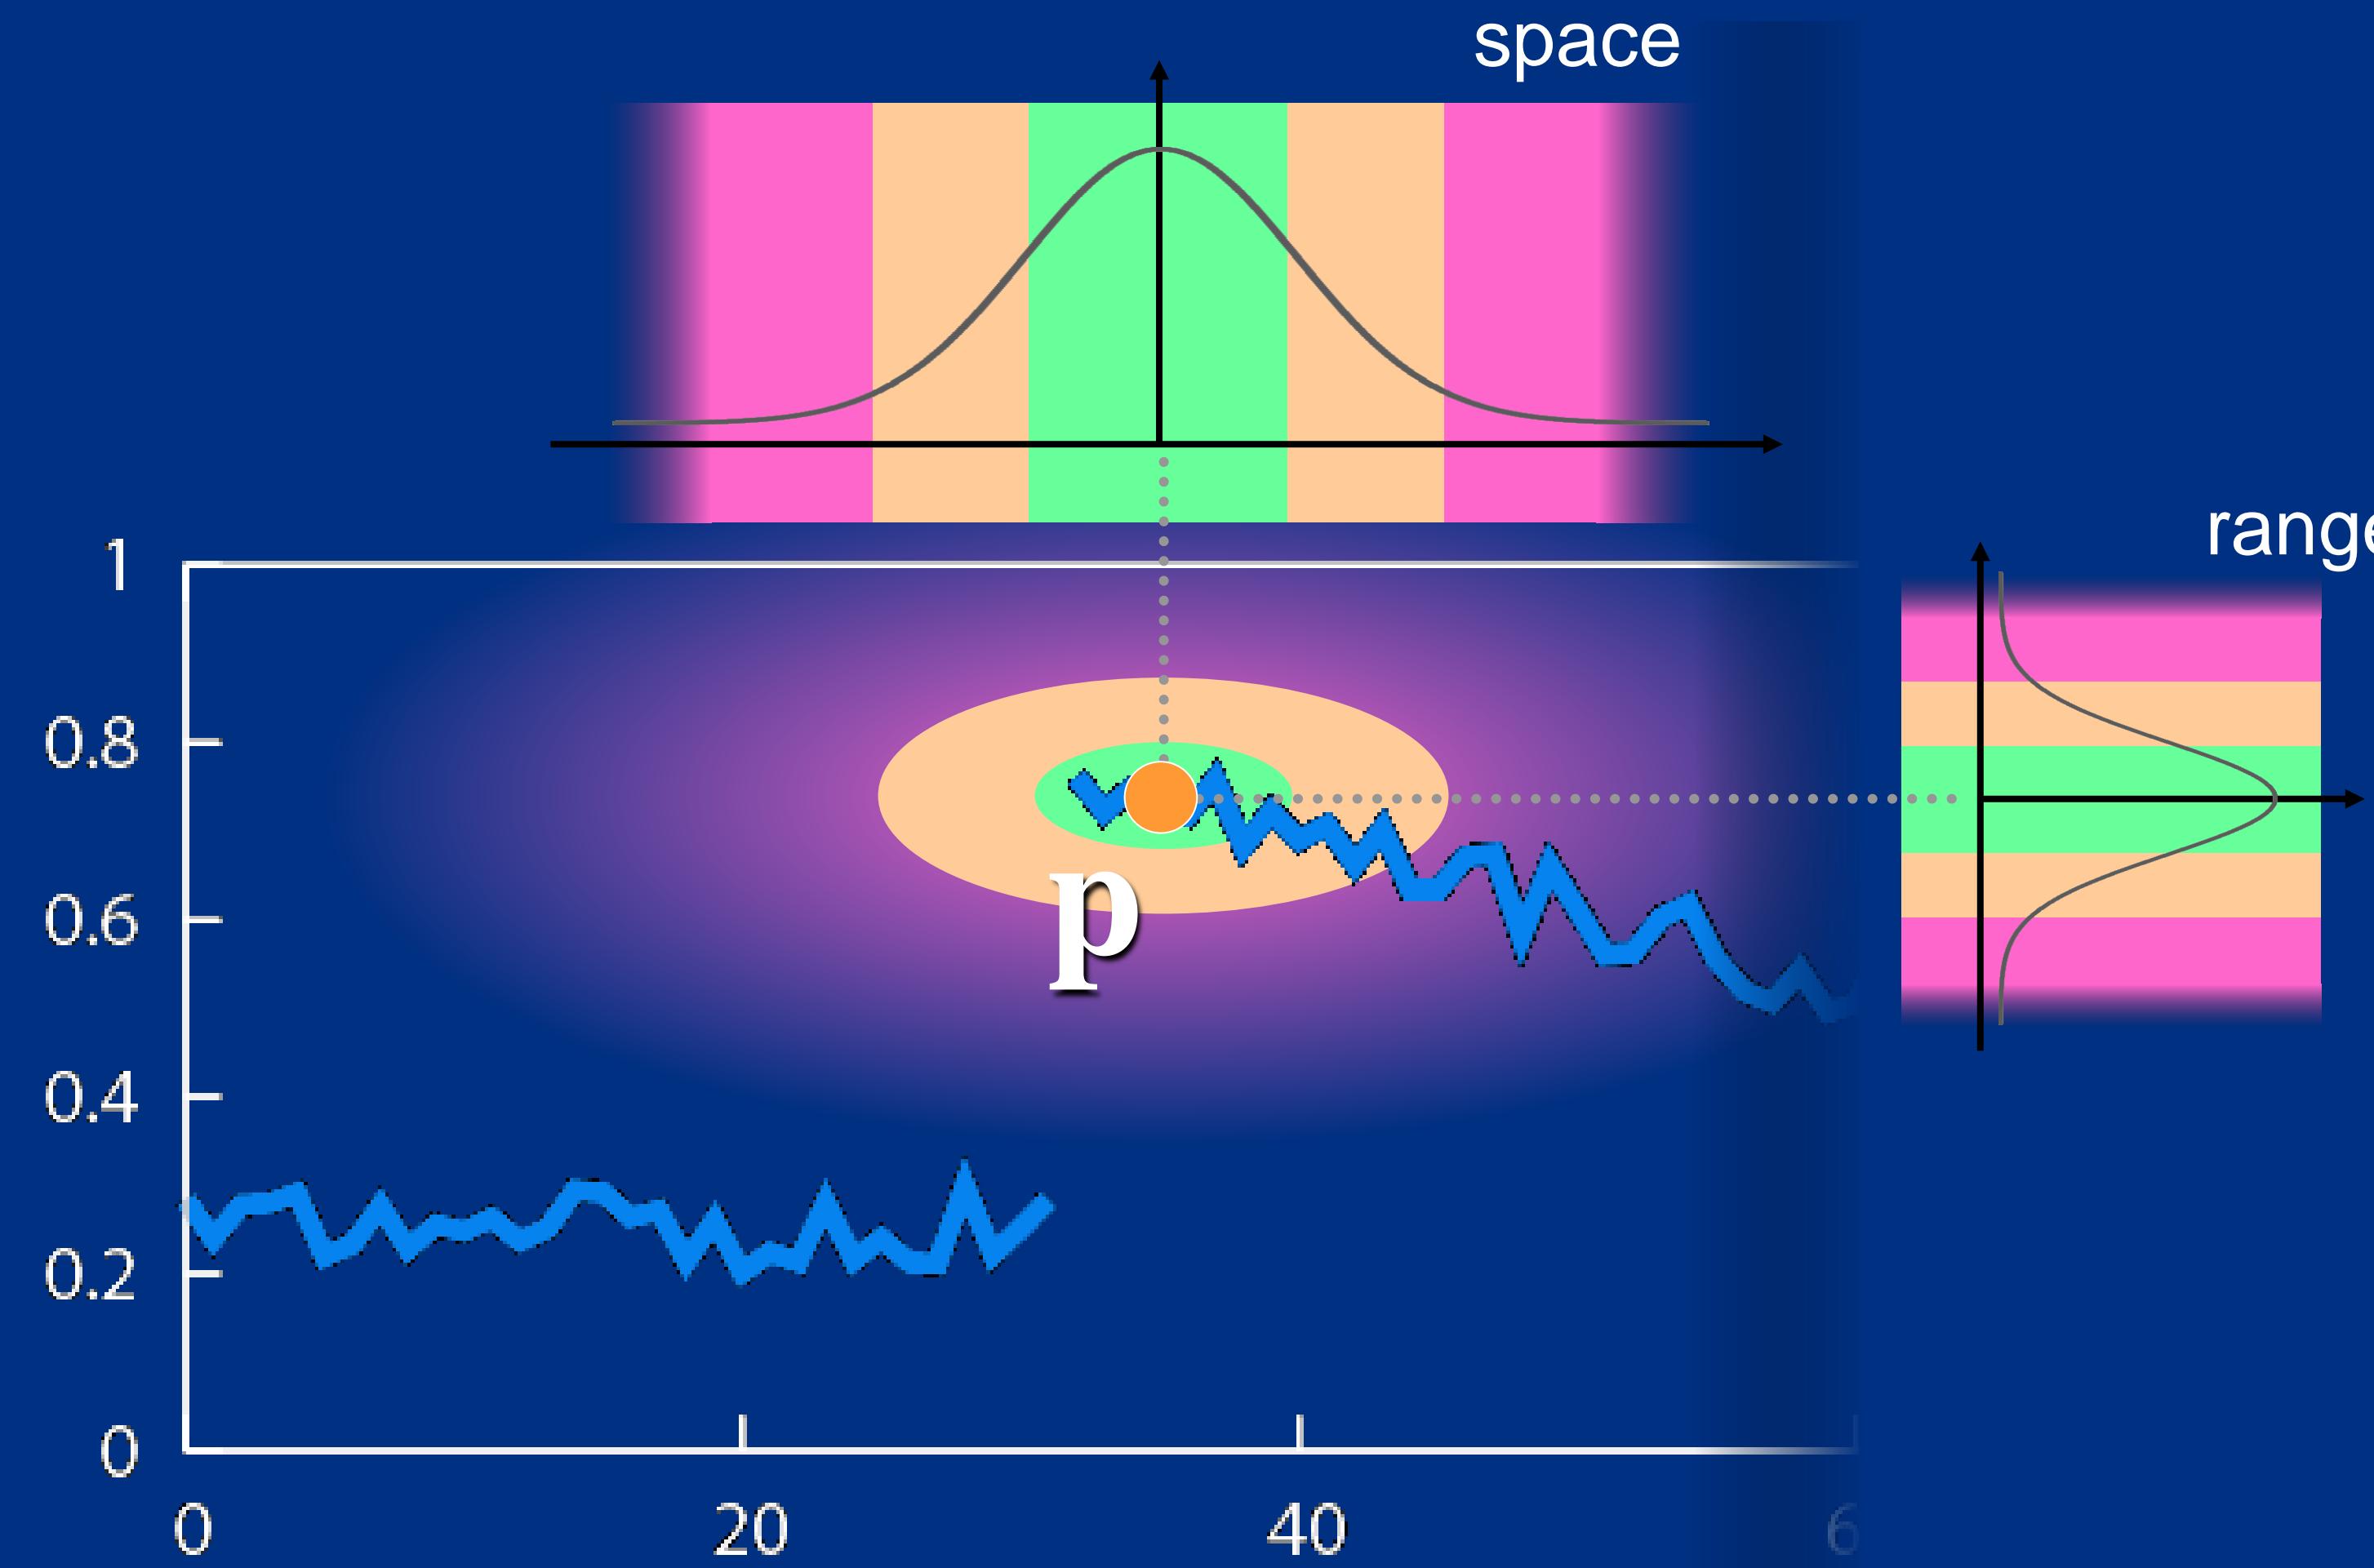
Space and Range Parameters

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

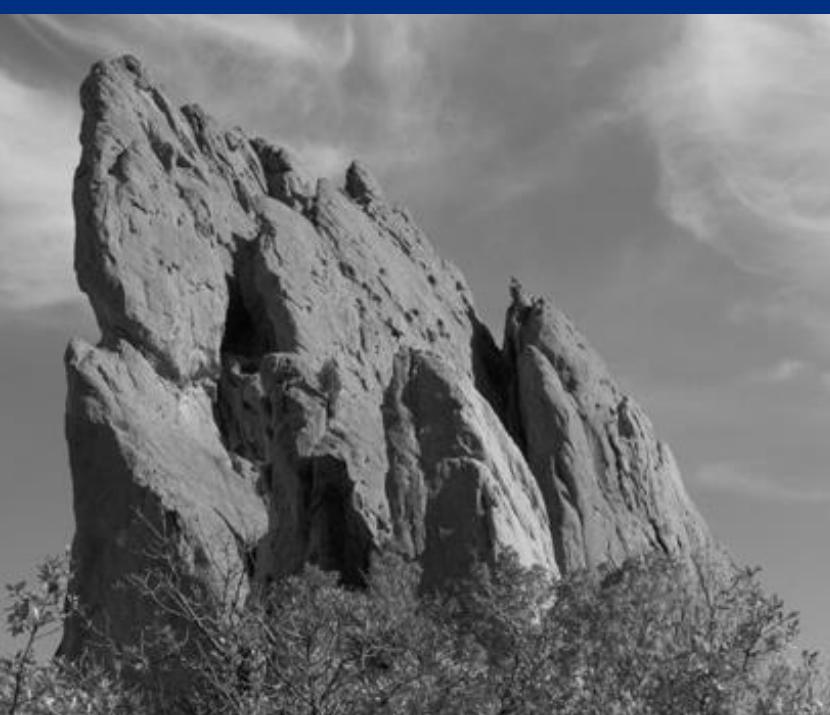

- space σ_s : spatial extent of the kernel, size of the considered neighborhood.
- range σ_r : “minimum” amplitude of an edge

Influence of Pixels

Only pixels close in space and in range are considered.



Exploring the Parameter Space



input



$\sigma_r = 0.1$



$\sigma_r = 0.25$



$\sigma_r = \infty$
(Gaussian blur)



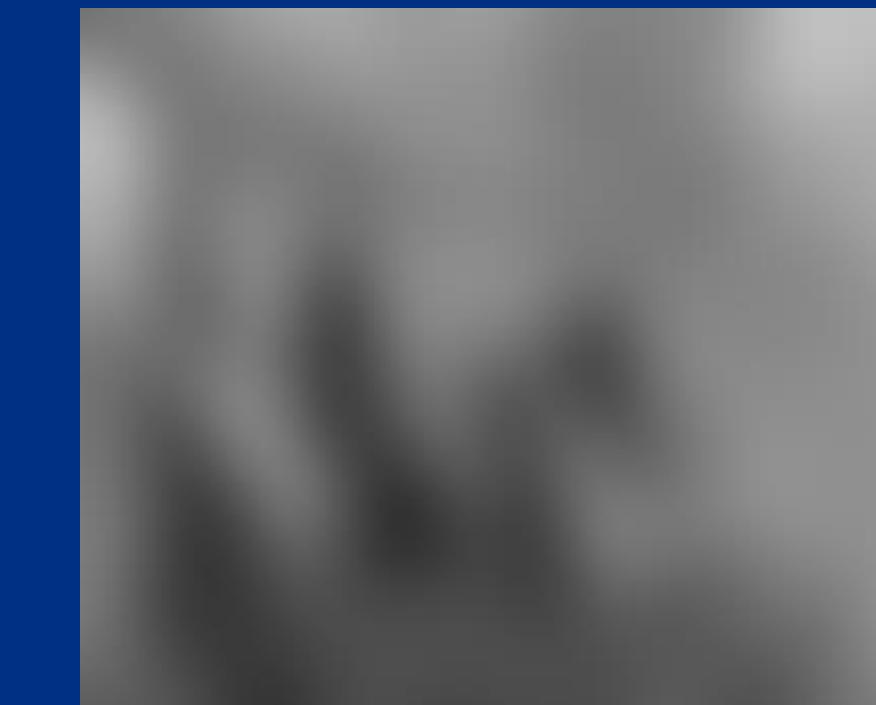
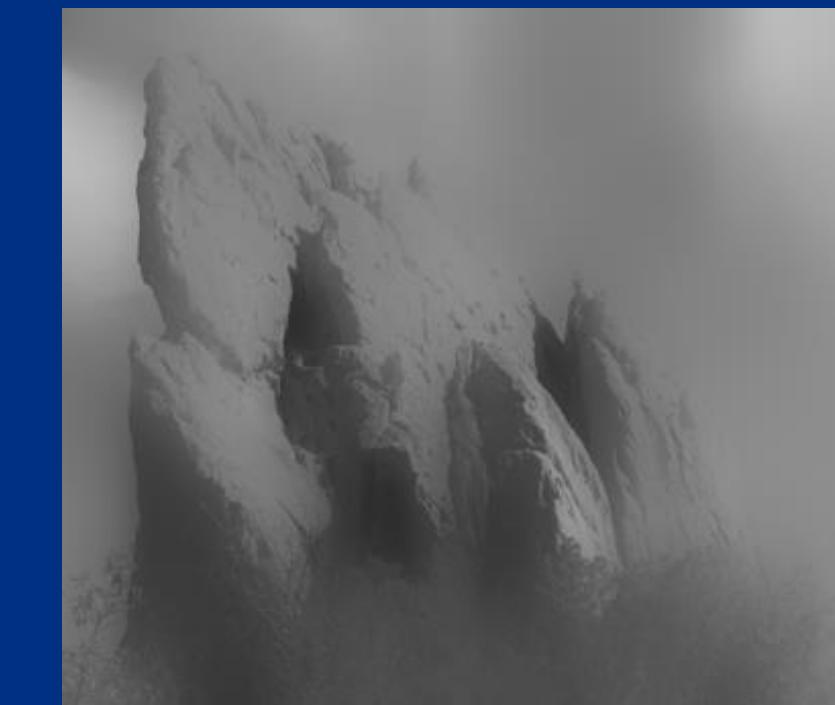
$\sigma_s = 2$



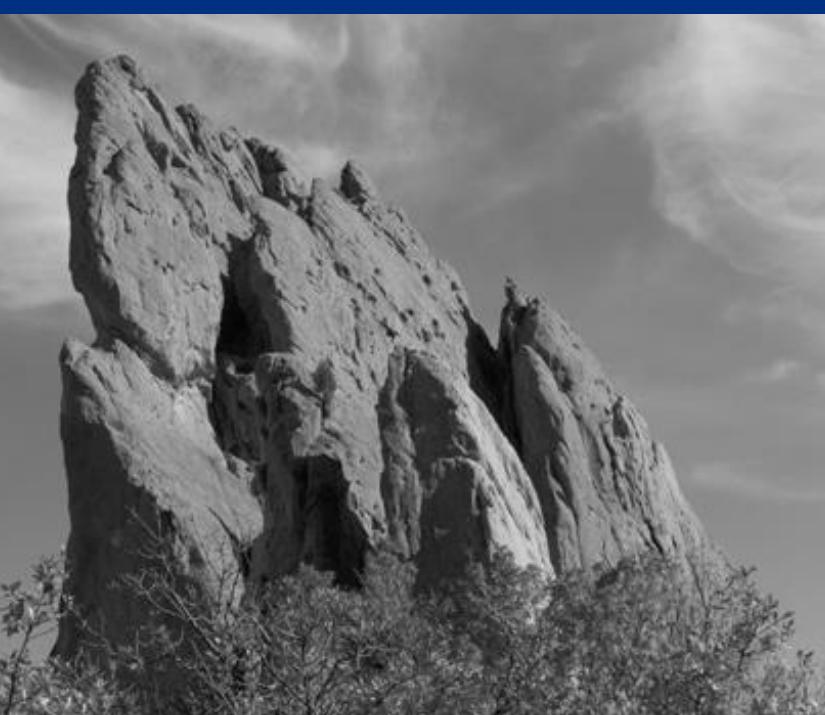
$\sigma_s = 6$



$\sigma_s = 18$

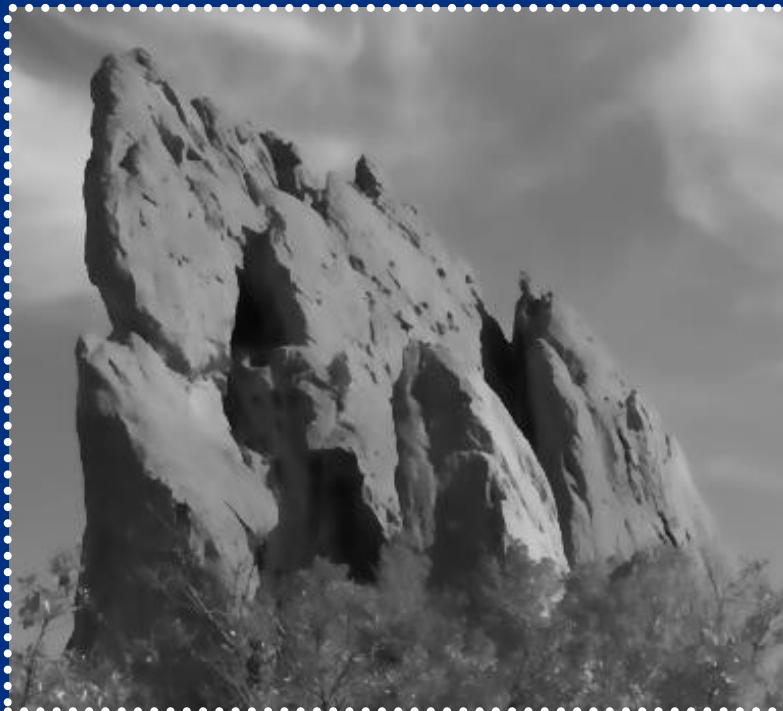


Varying the Range Parameter



input

$\sigma_r = 0.1$



$\sigma_r = 0.25$



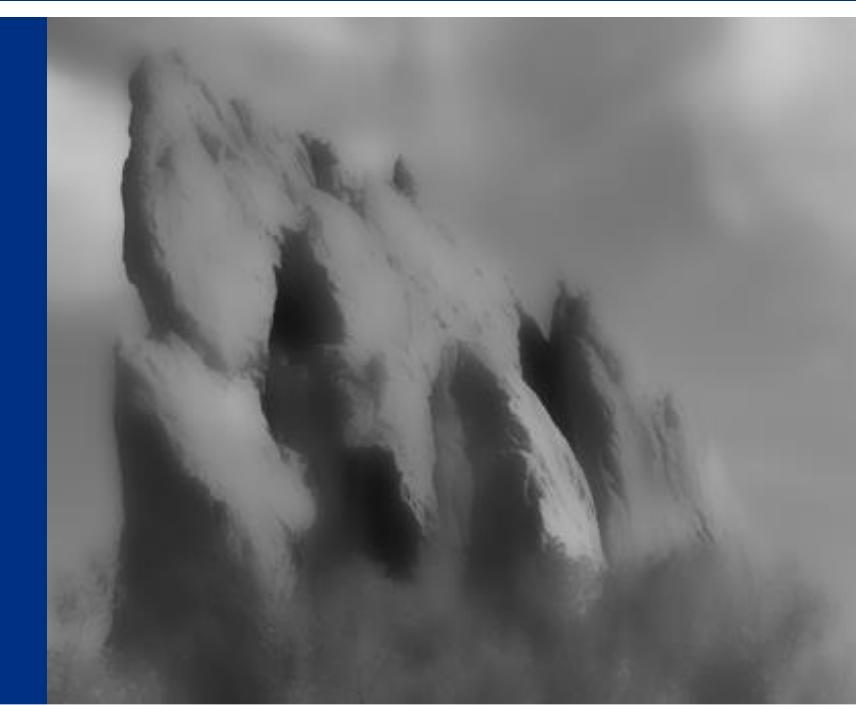
$\sigma_r = \infty$
(Gaussian blur)



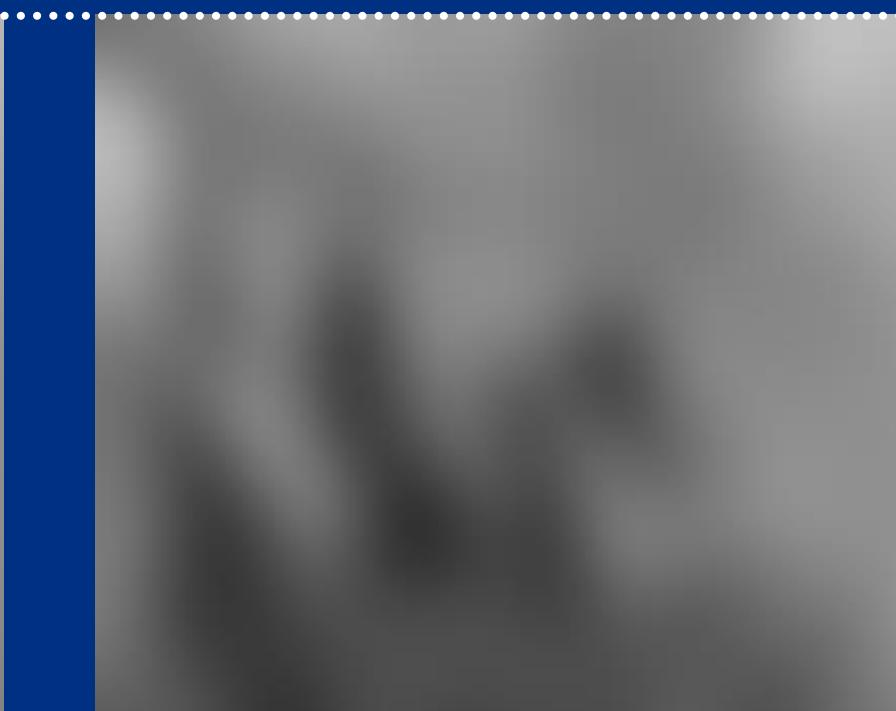
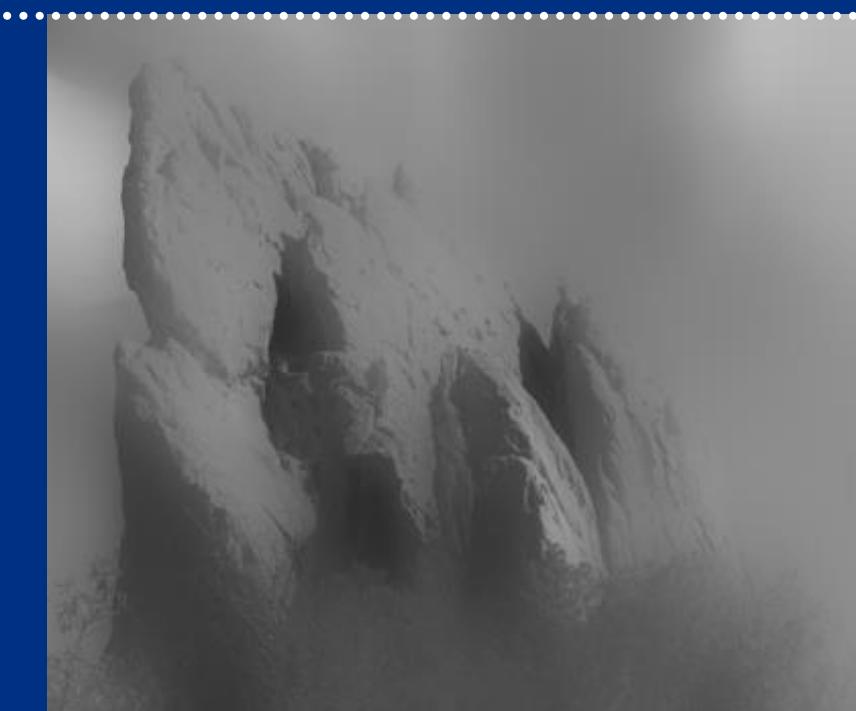
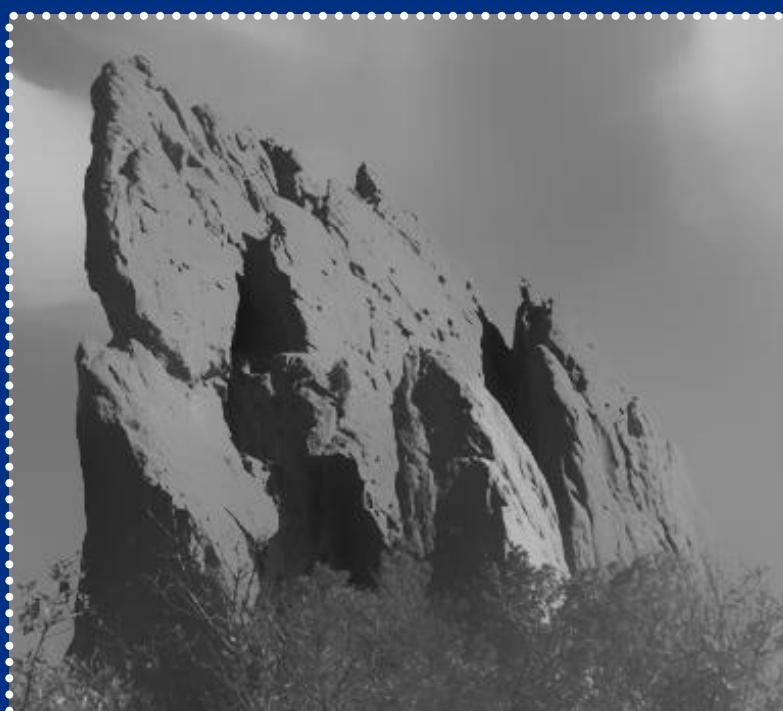
$\sigma_s = 2$



$\sigma_s = 6$



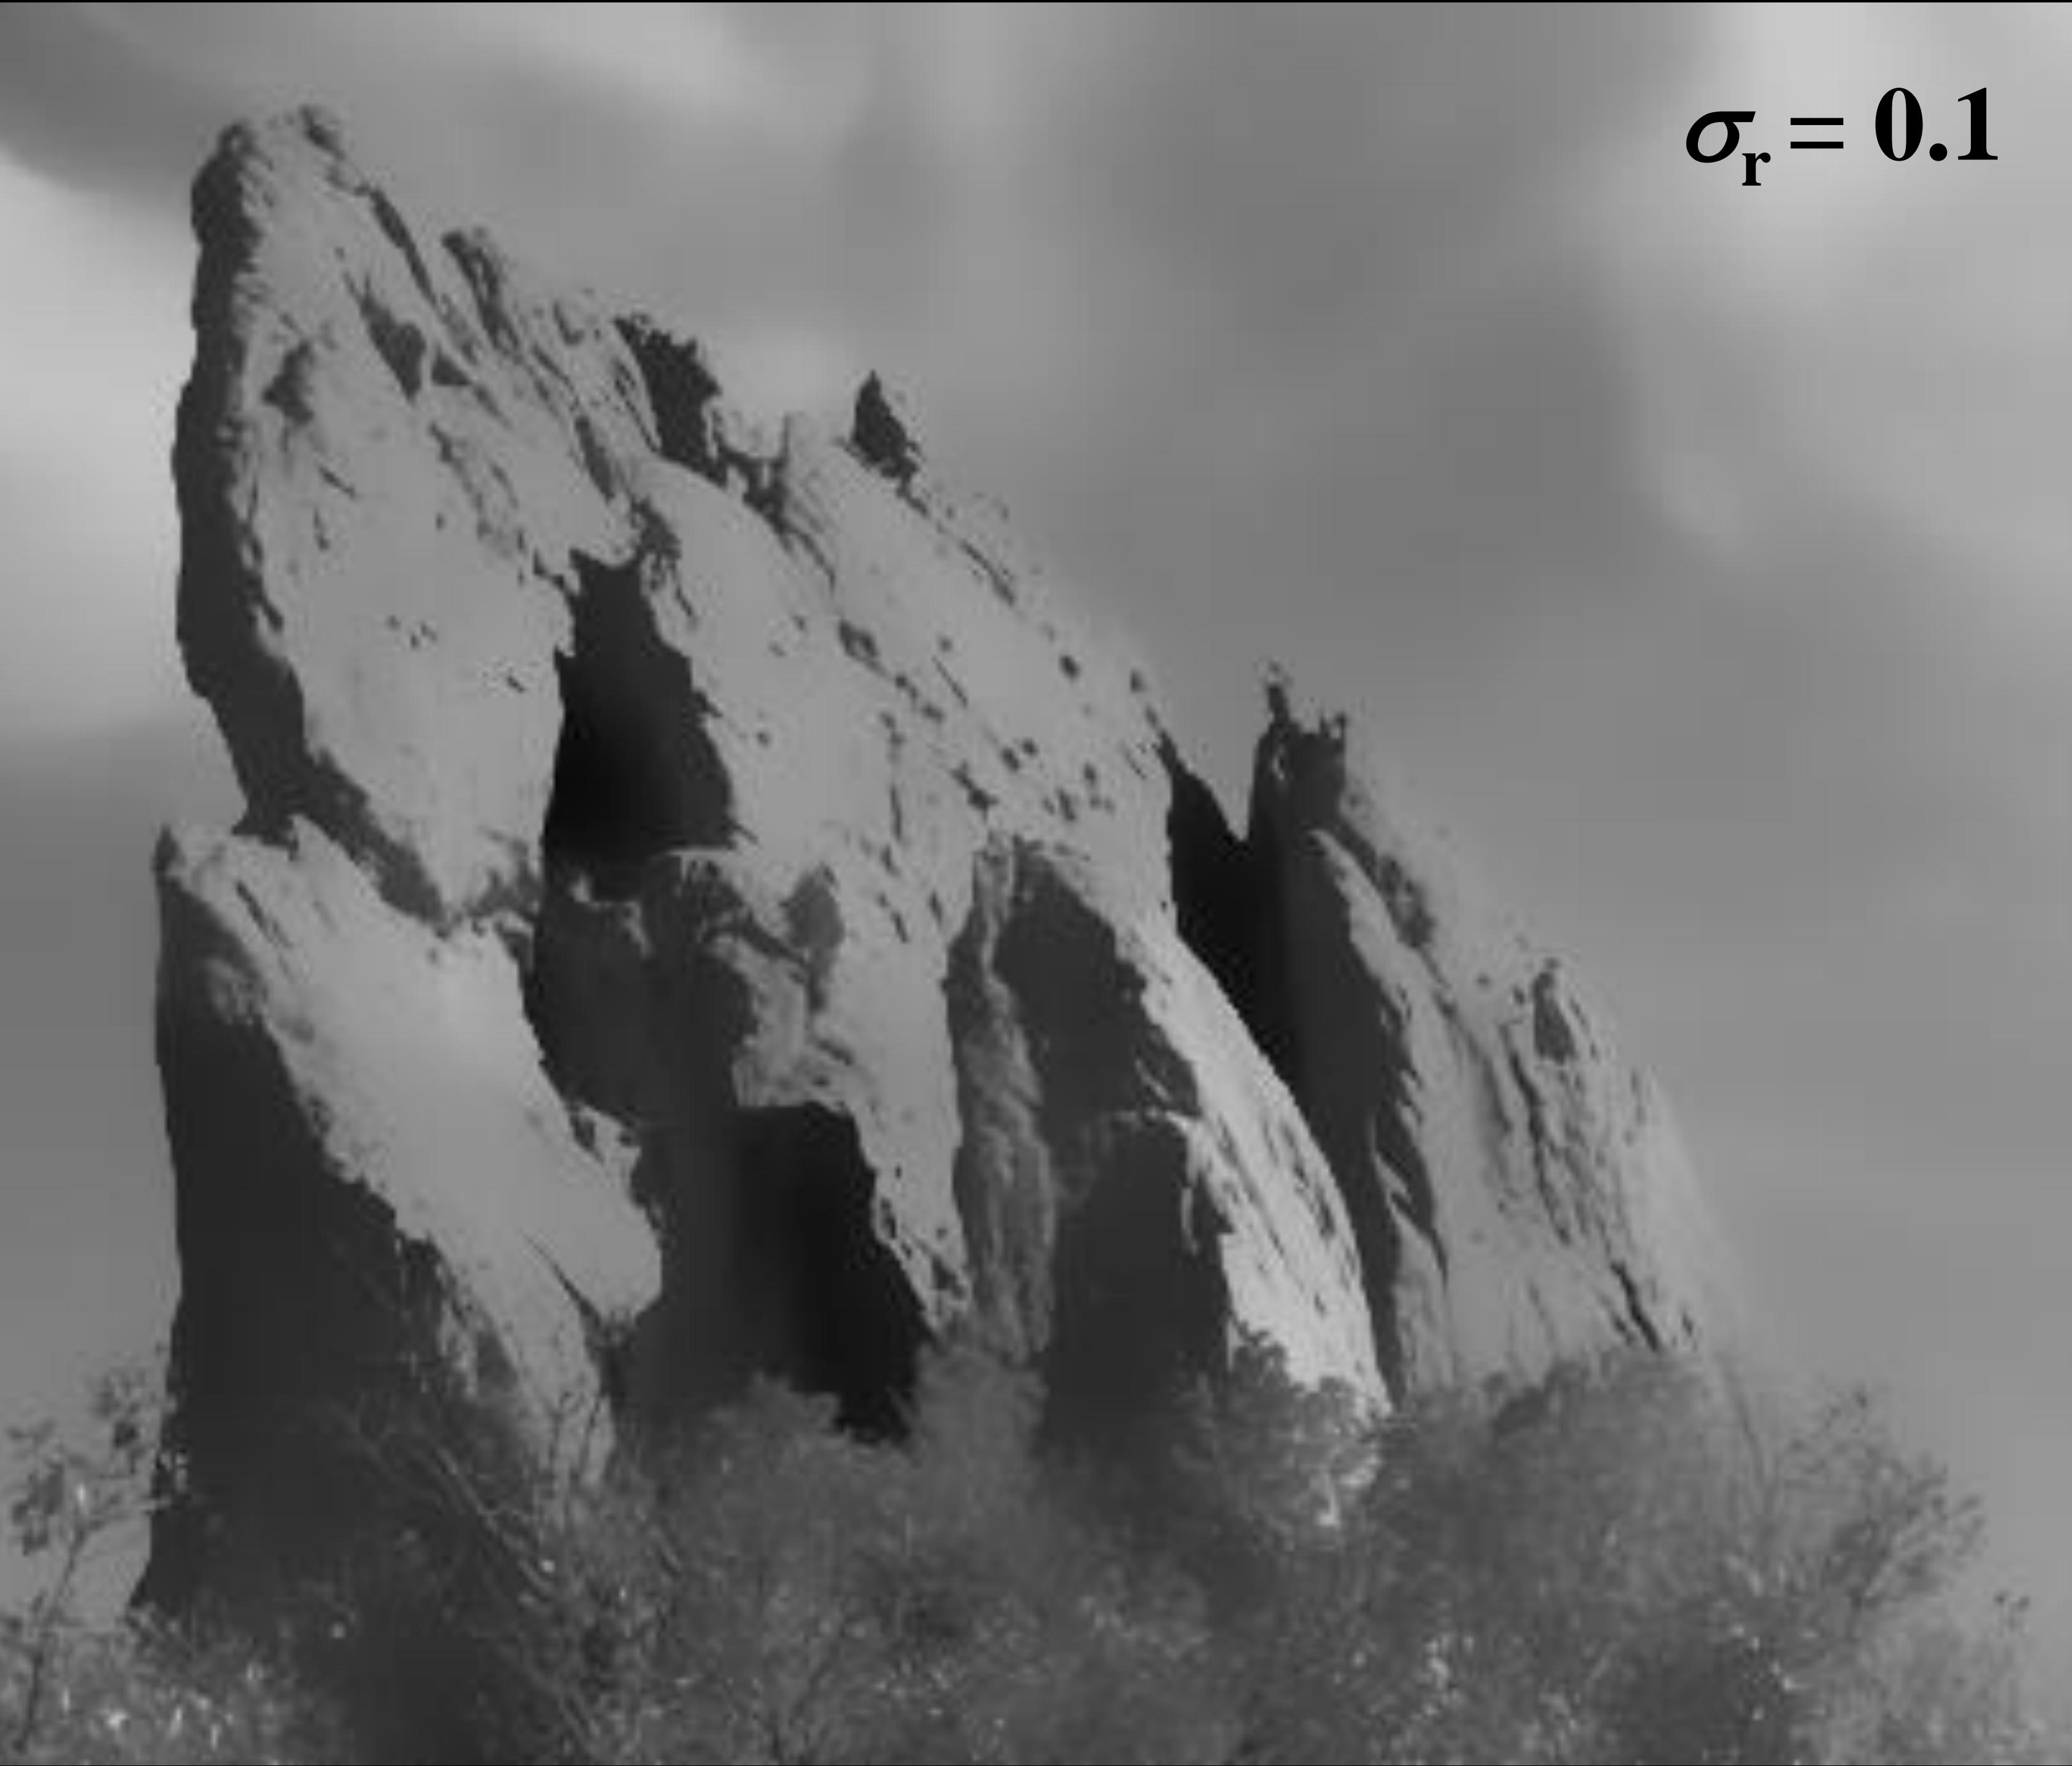
$\sigma_s = 18$

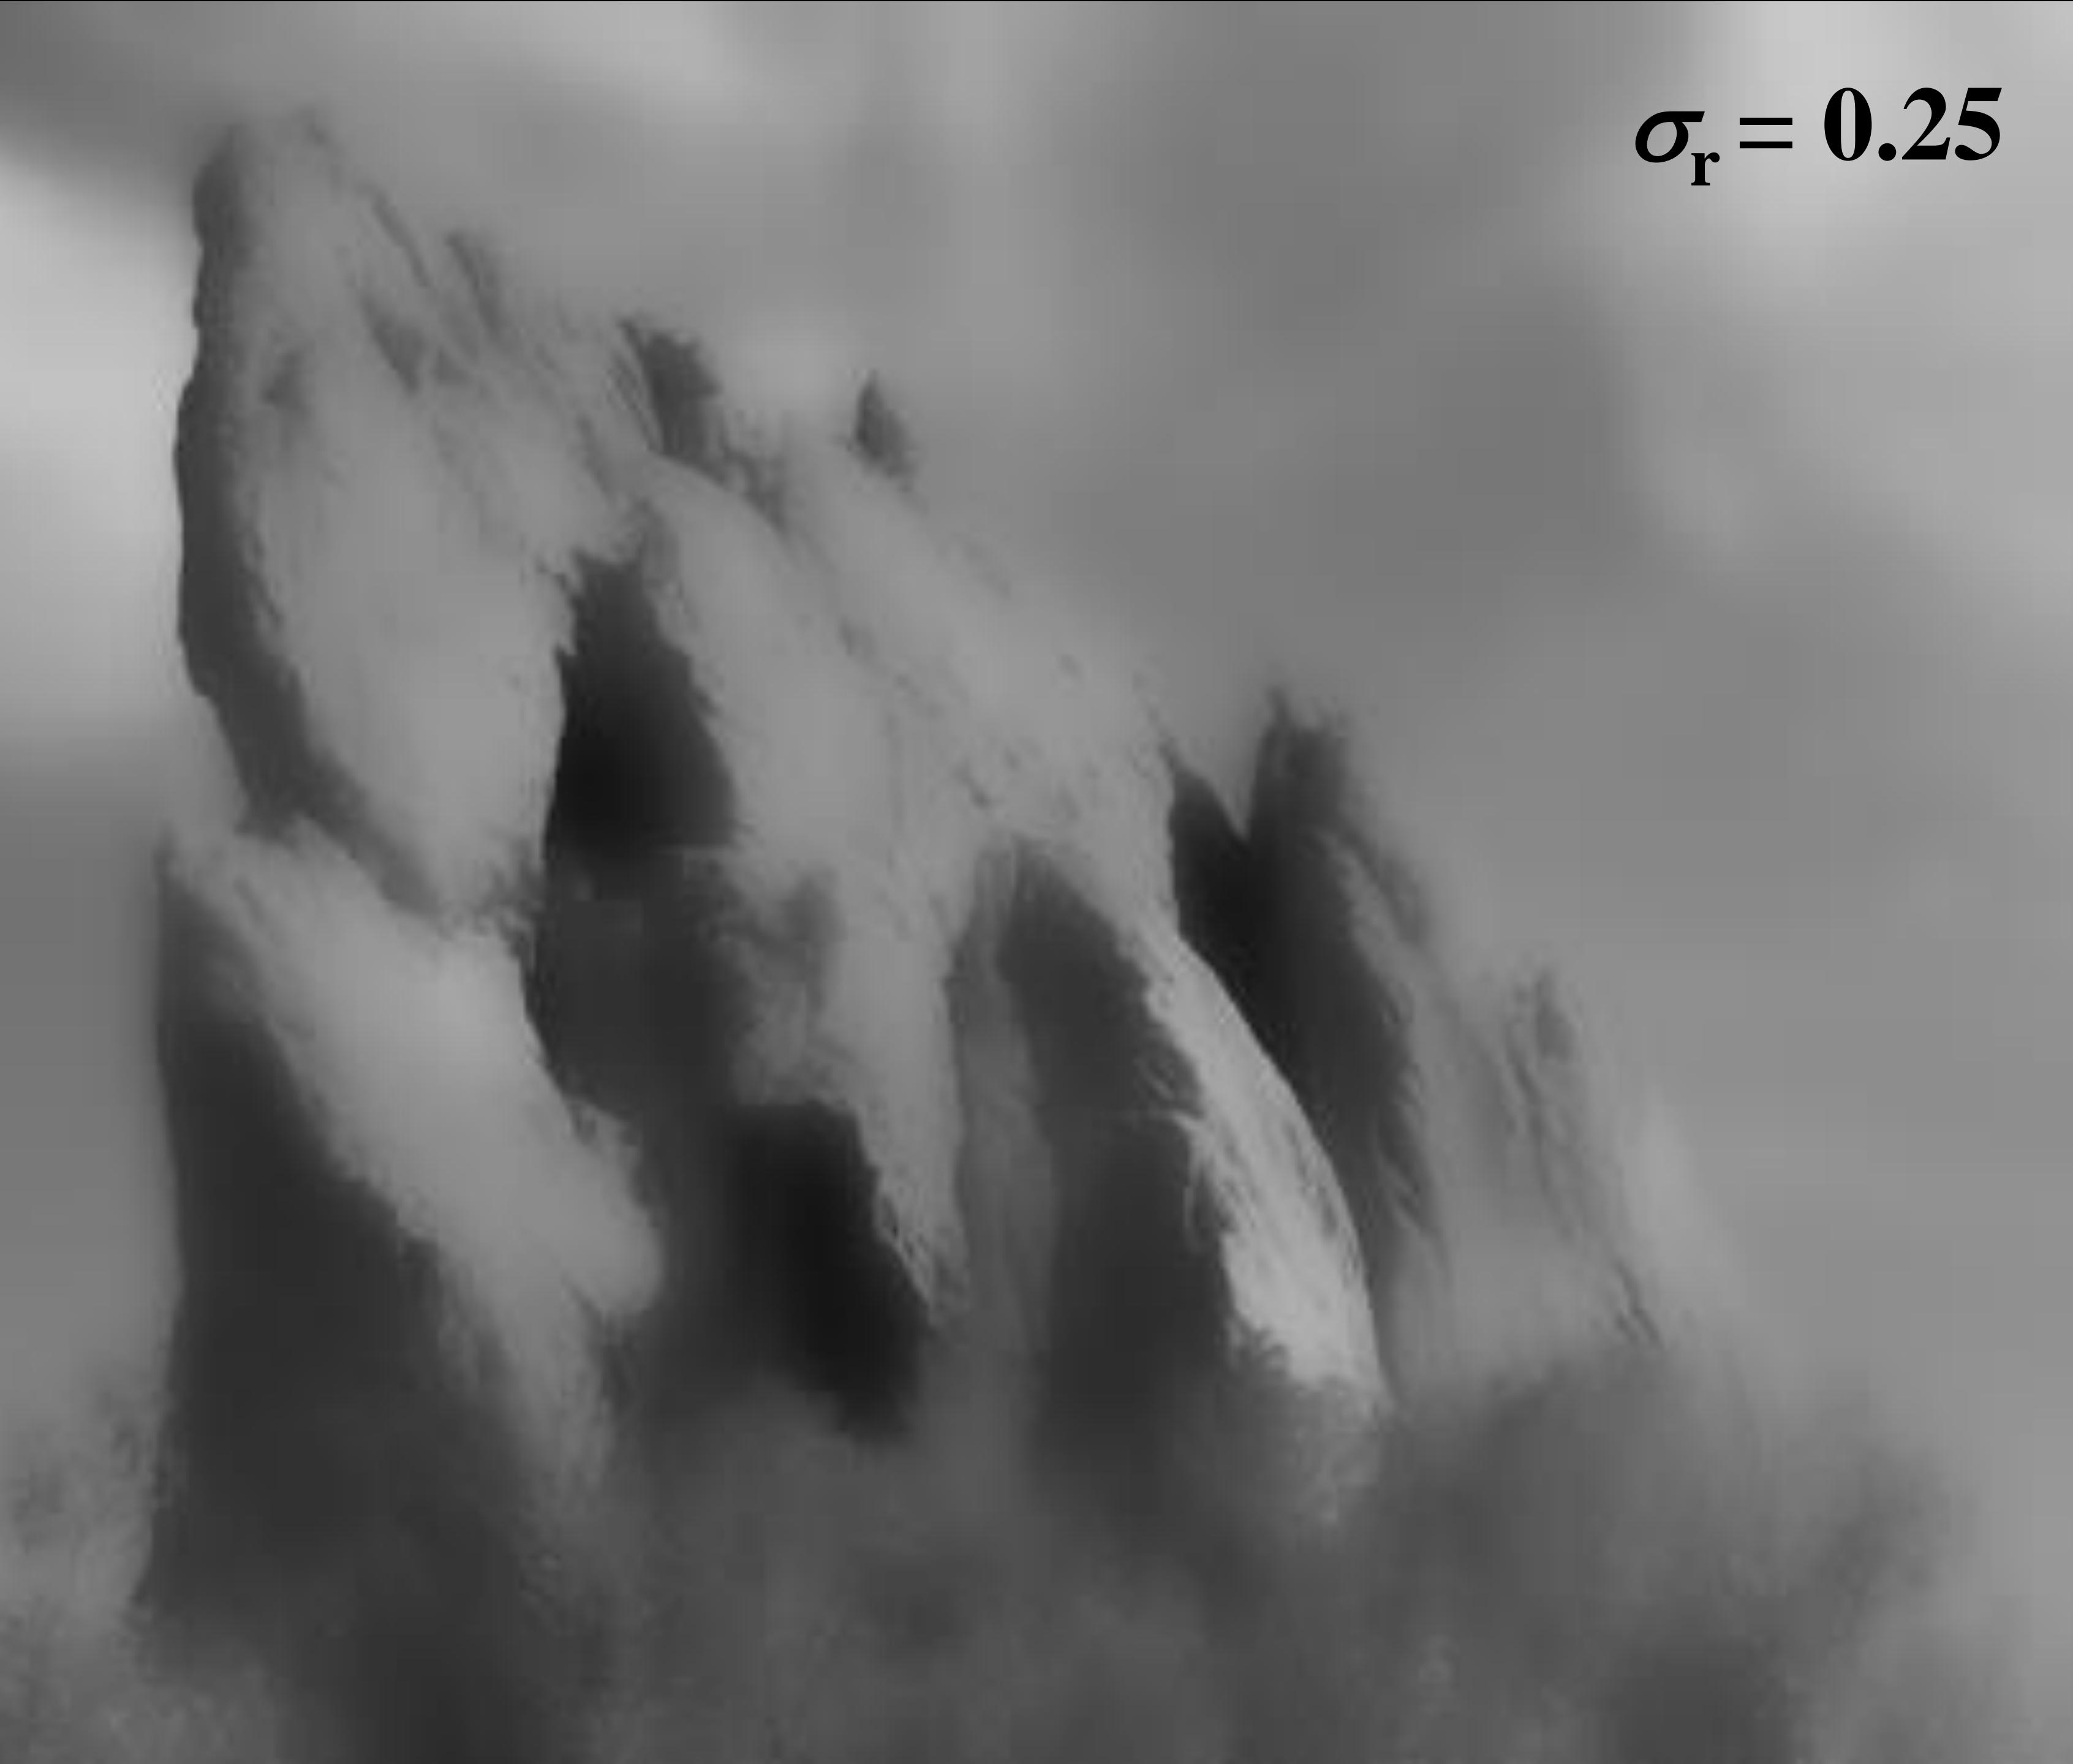


input



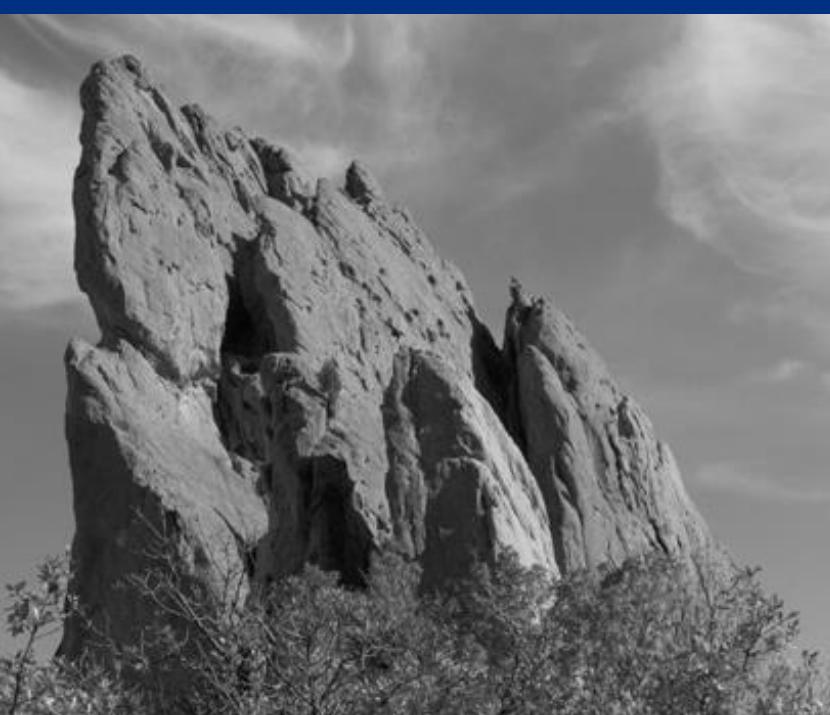
$\sigma_r = 0.1$



$\sigma_r = 0.25$ 

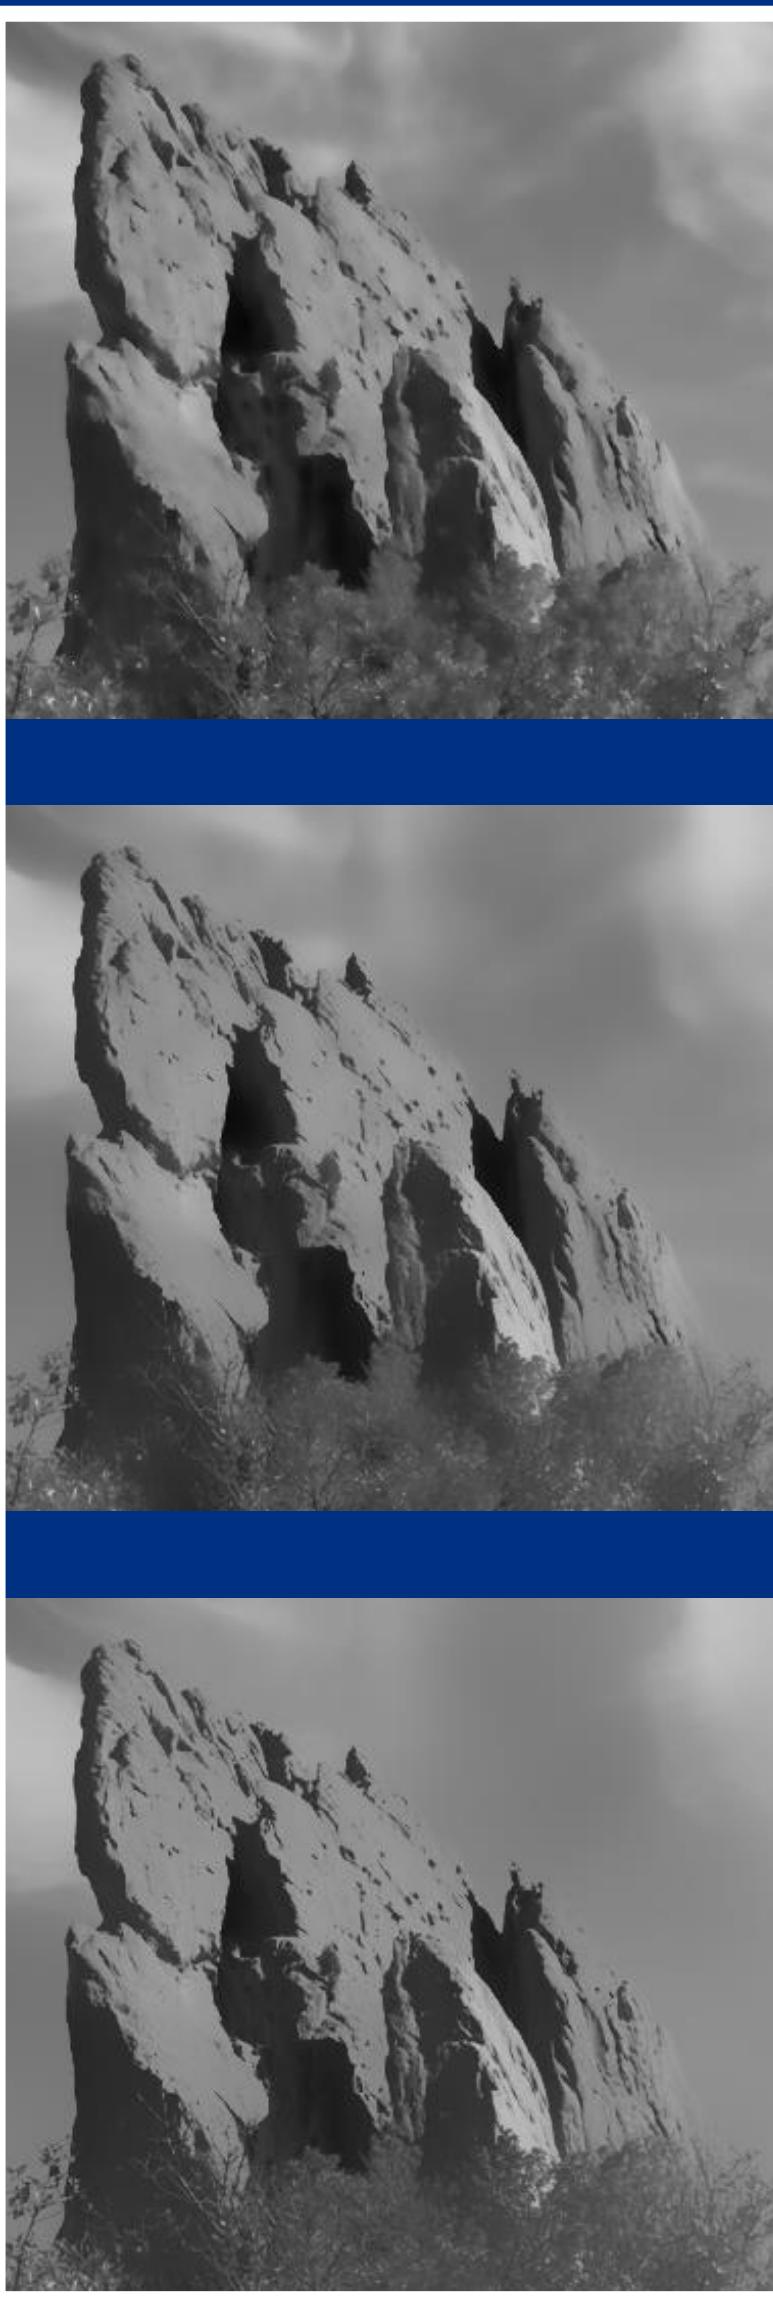
$\sigma_r = \infty$
(Gaussian blur)

Varying the Space Parameter



input

$\sigma_r = 0.1$



$\sigma_s = 2$



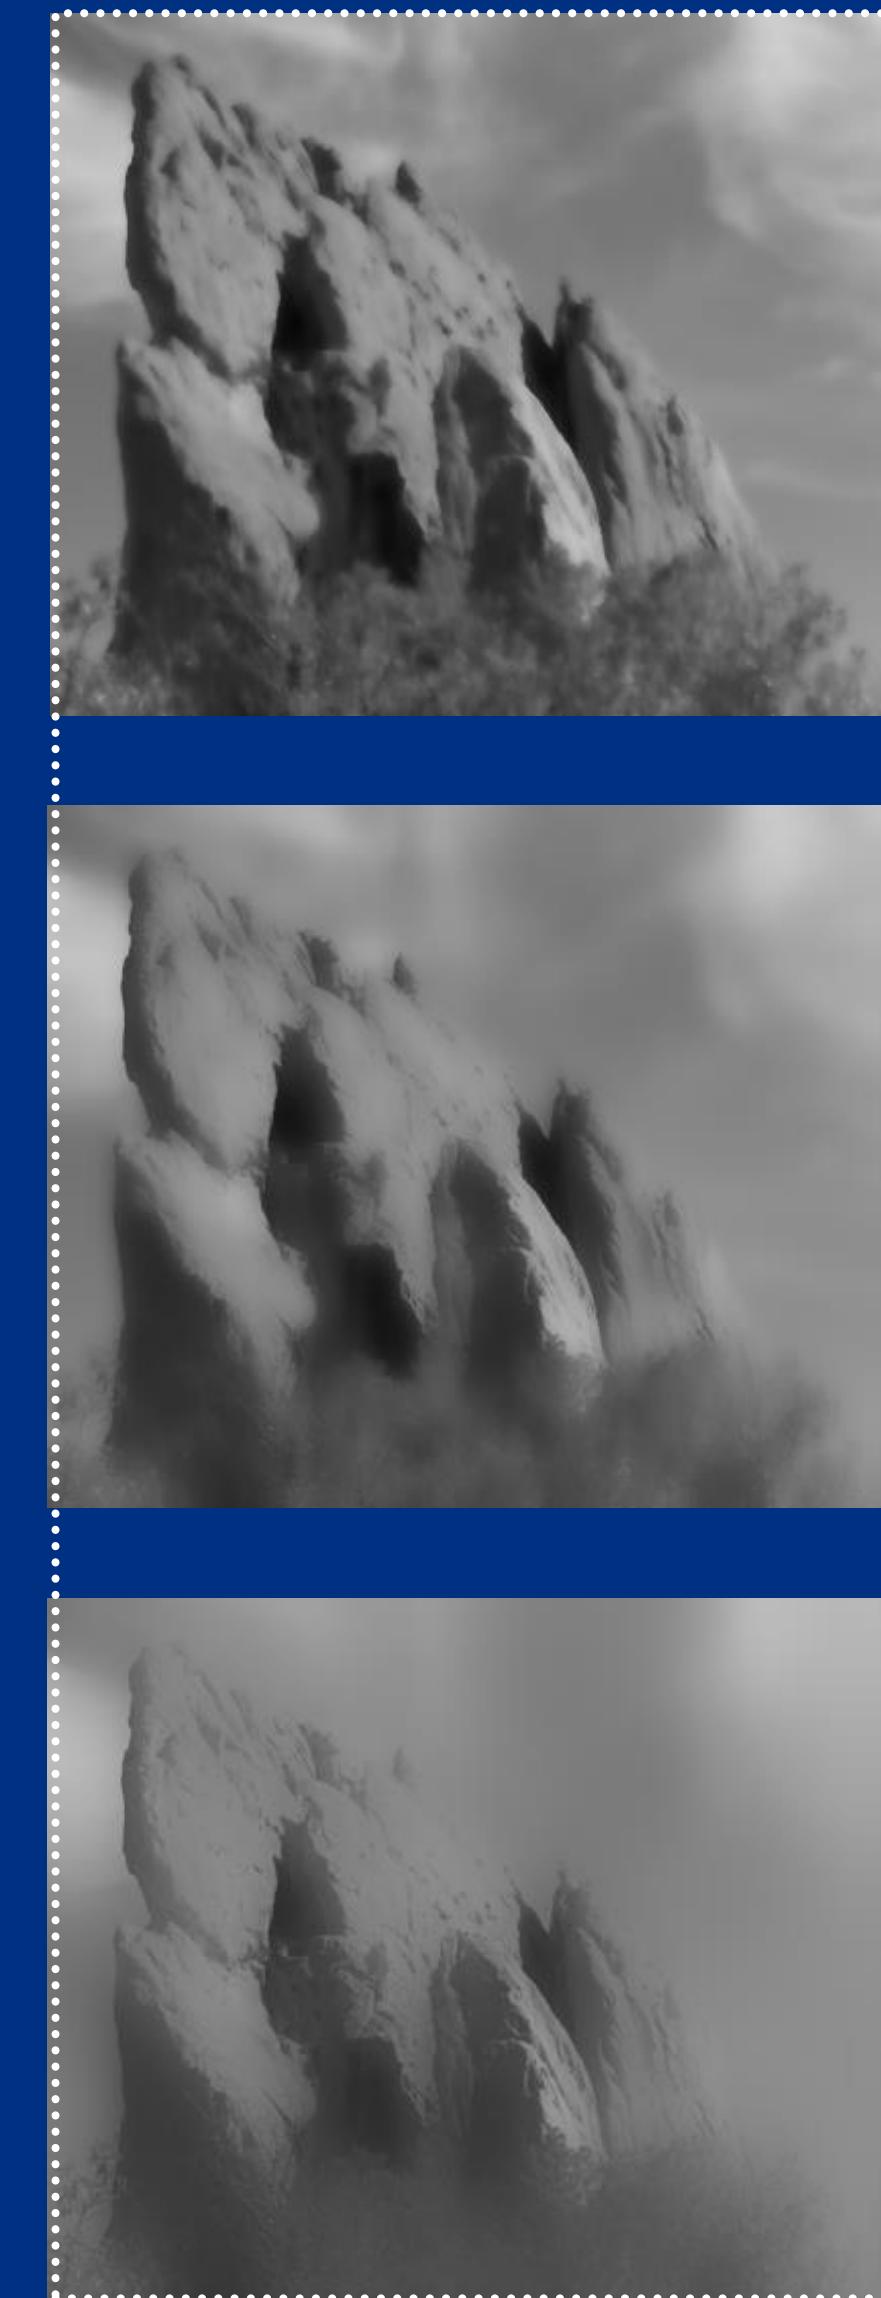
$\sigma_s = 6$



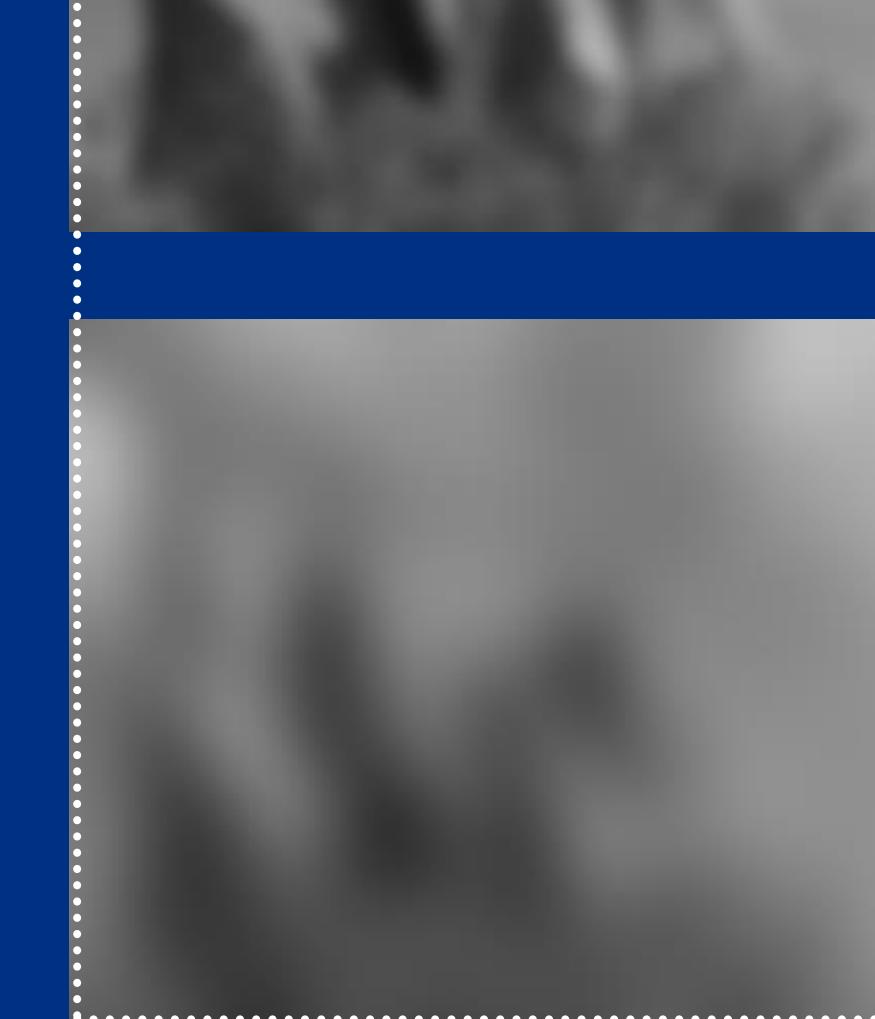
$\sigma_s = 18$



$\sigma_r = 0.25$



$\sigma_r = \infty$
(Gaussian blur)



input



$\sigma_s = 2$ 

$\sigma_s = 6$



$\sigma_s = 18$ 