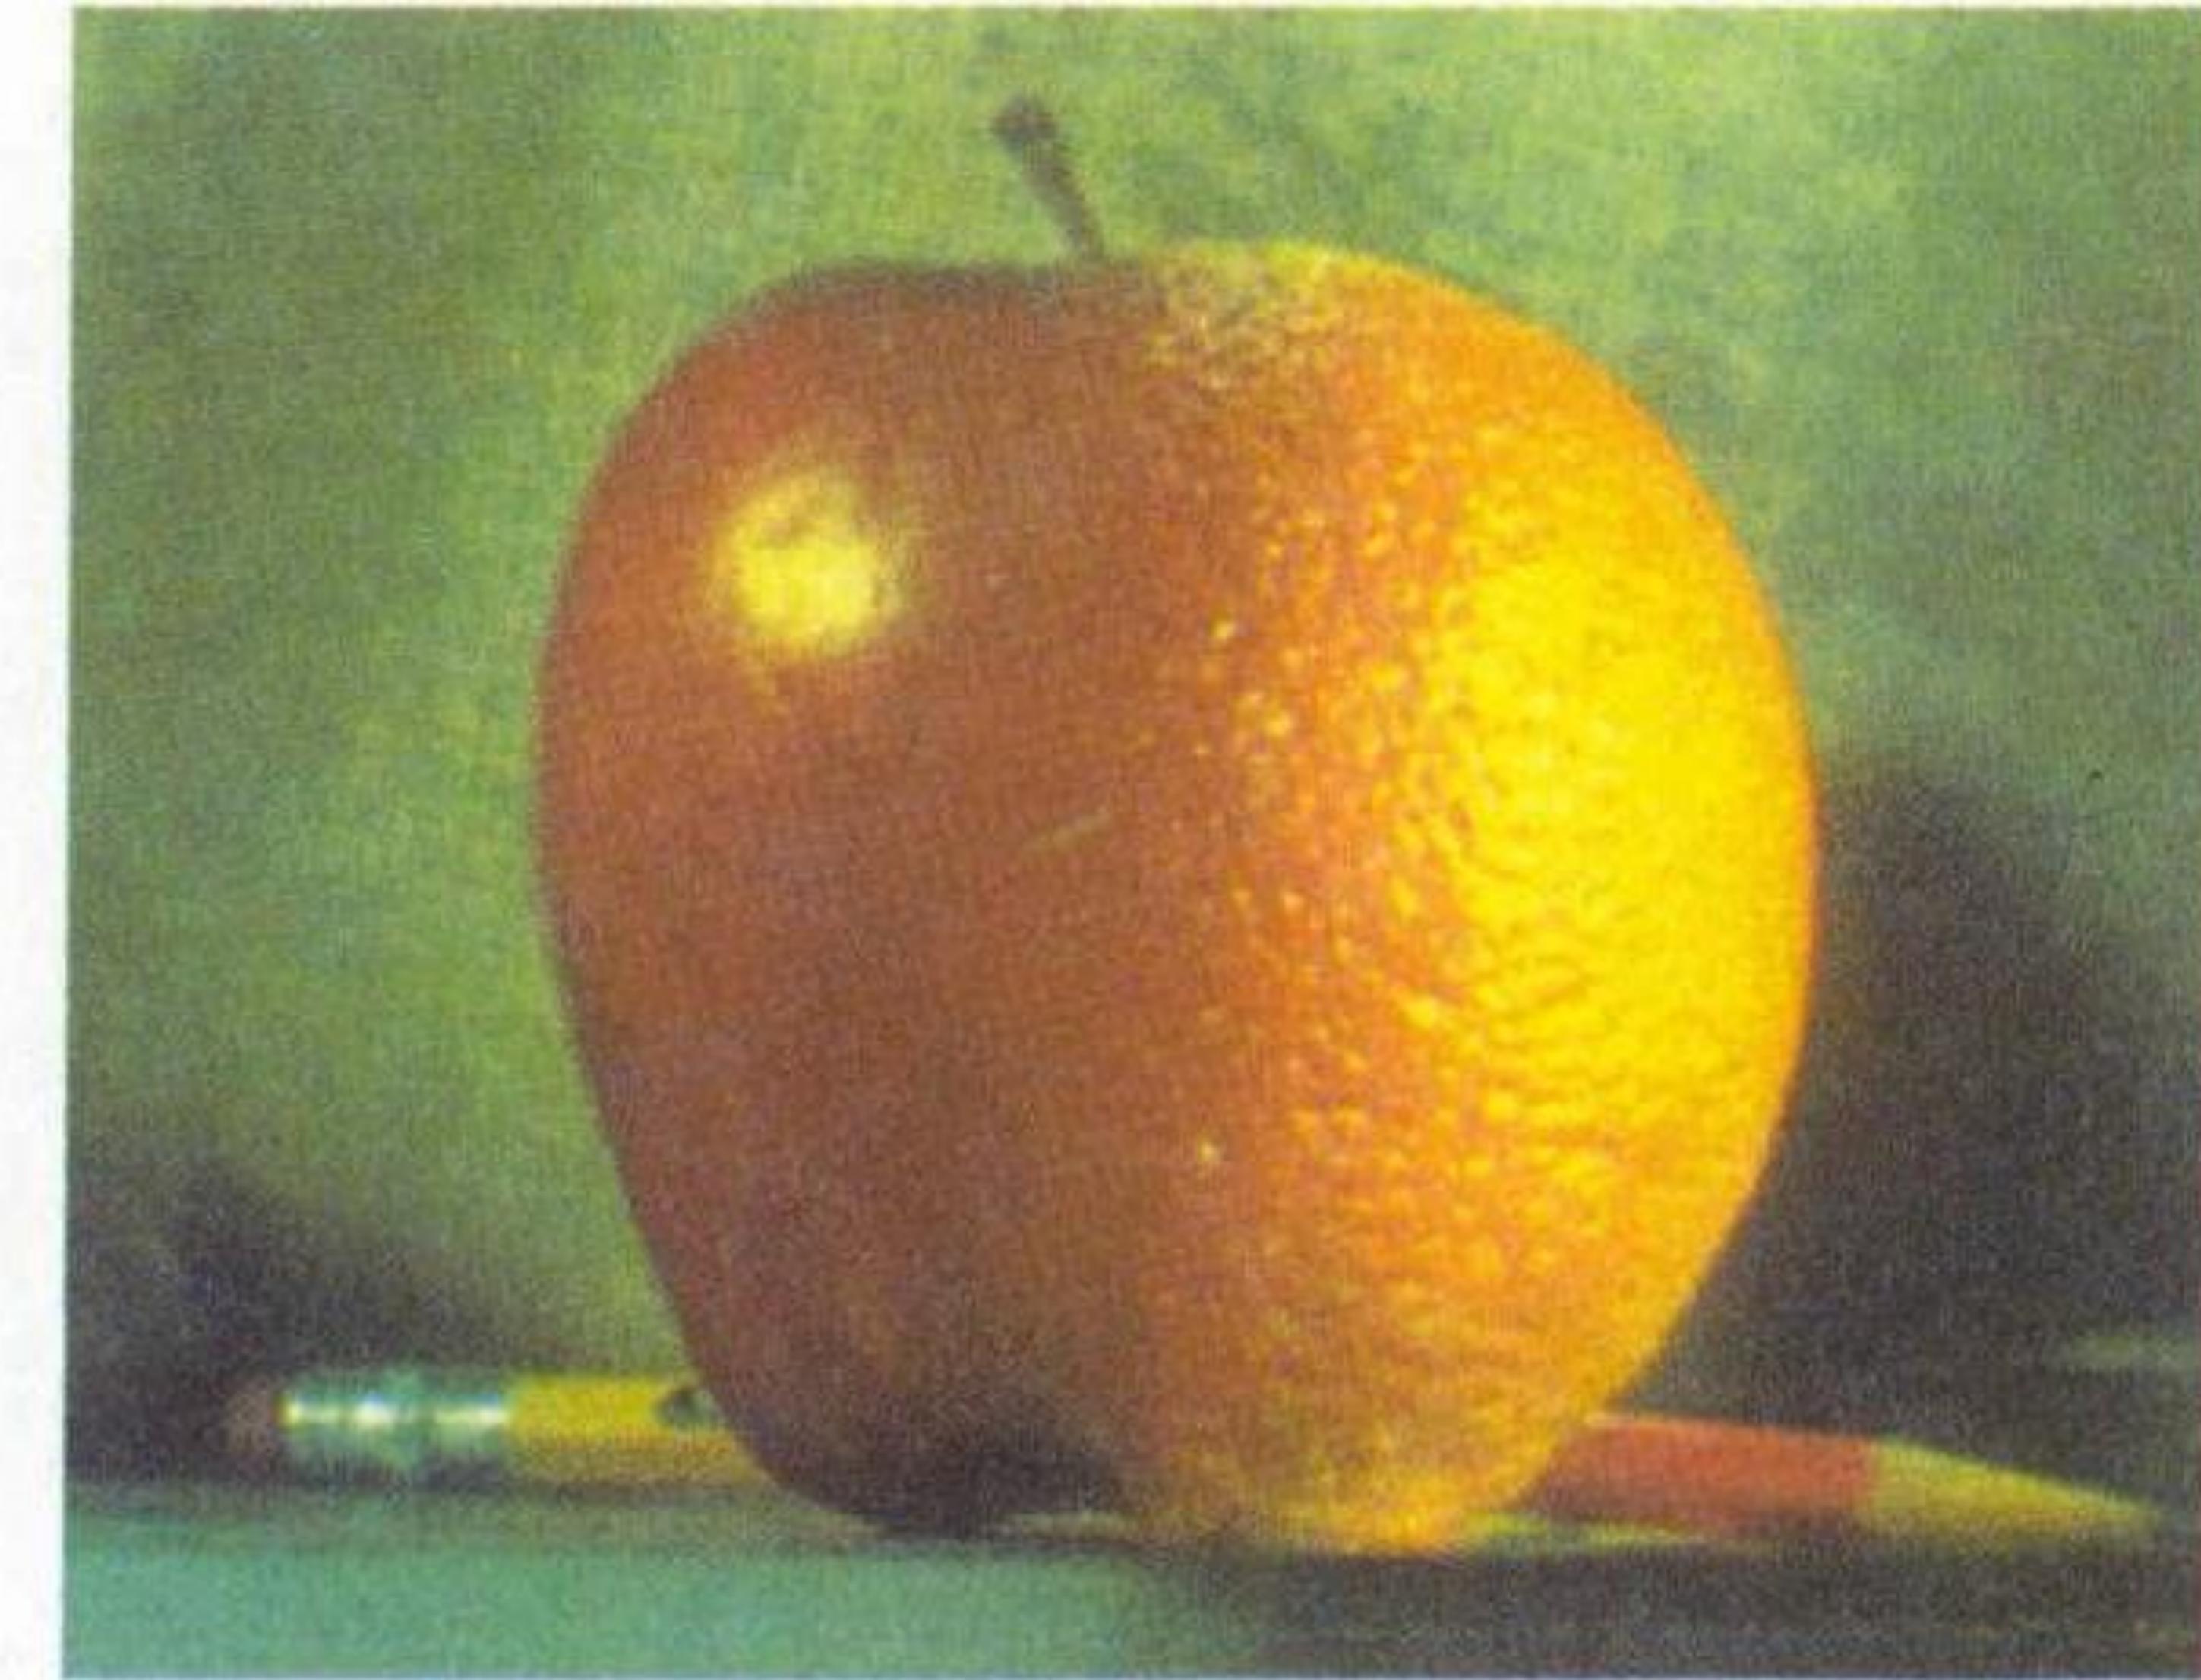


# Sampling, Reconstruction, Pyramids



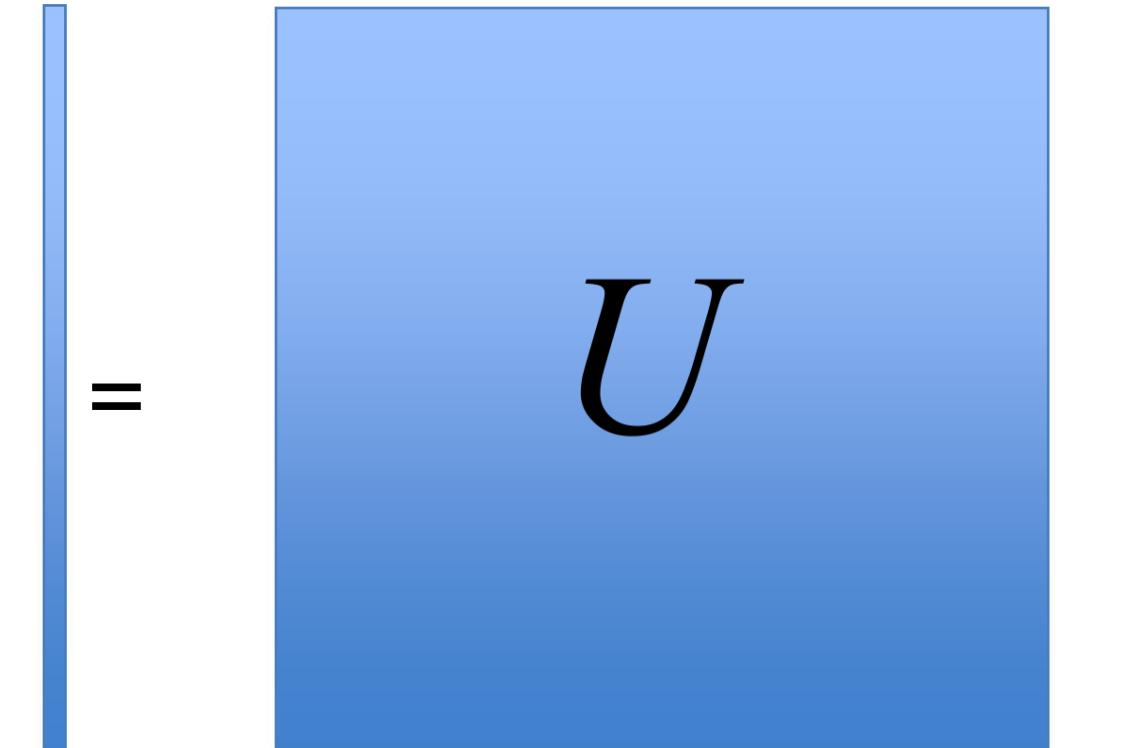
# Linear image transformations

---

In analyzing images, it's often useful to make a change of basis.

$$\vec{\bar{F}} = \vec{U} \vec{f}$$

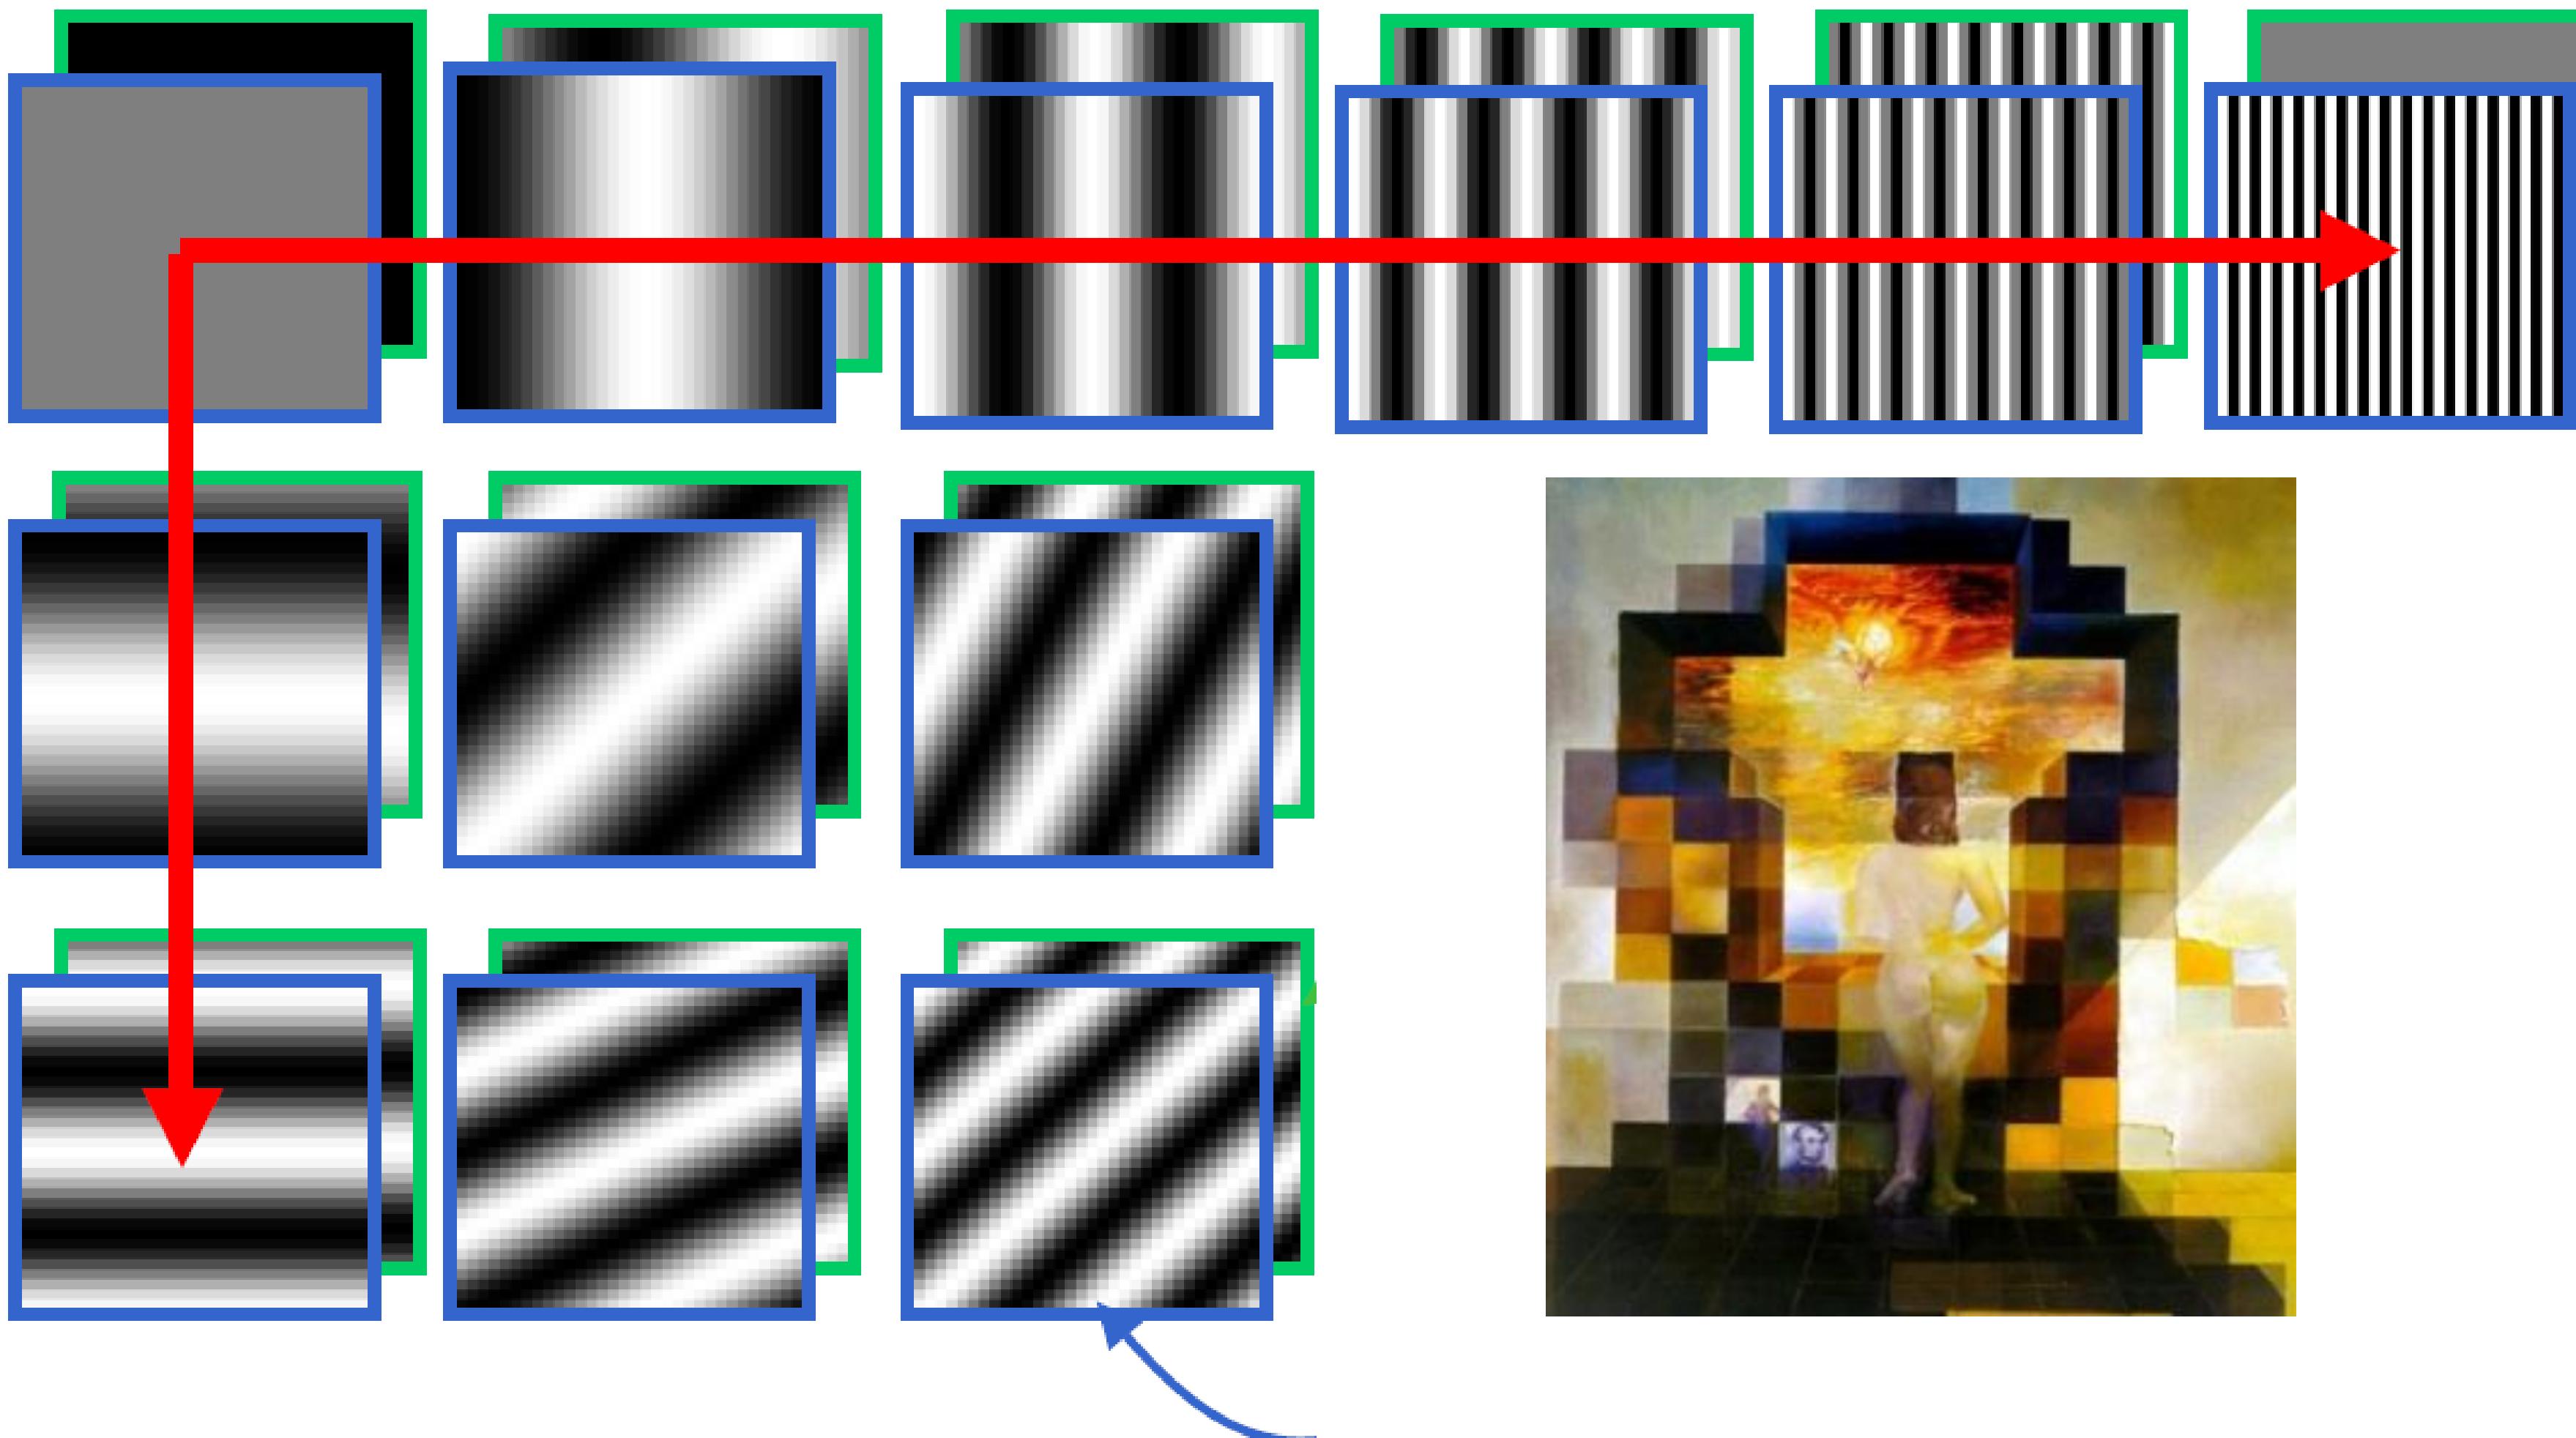
Transformed image →  $\vec{\bar{F}}$  =  $\vec{U} \vec{f}$  ← Vectorized image

=  Fourier transform, or  
Wavelet transform, or  
Steerable pyramid transform

# A nice set of basis

---

Teases away fast vs. slow changes in the image.



This change of basis has a special name...

# Jean Baptiste Joseph Fourier (1768-1830)

had crazy idea (1807)

*Any univariate function can be rewritten as a weighted sum of sines and cosines different frequencies.*

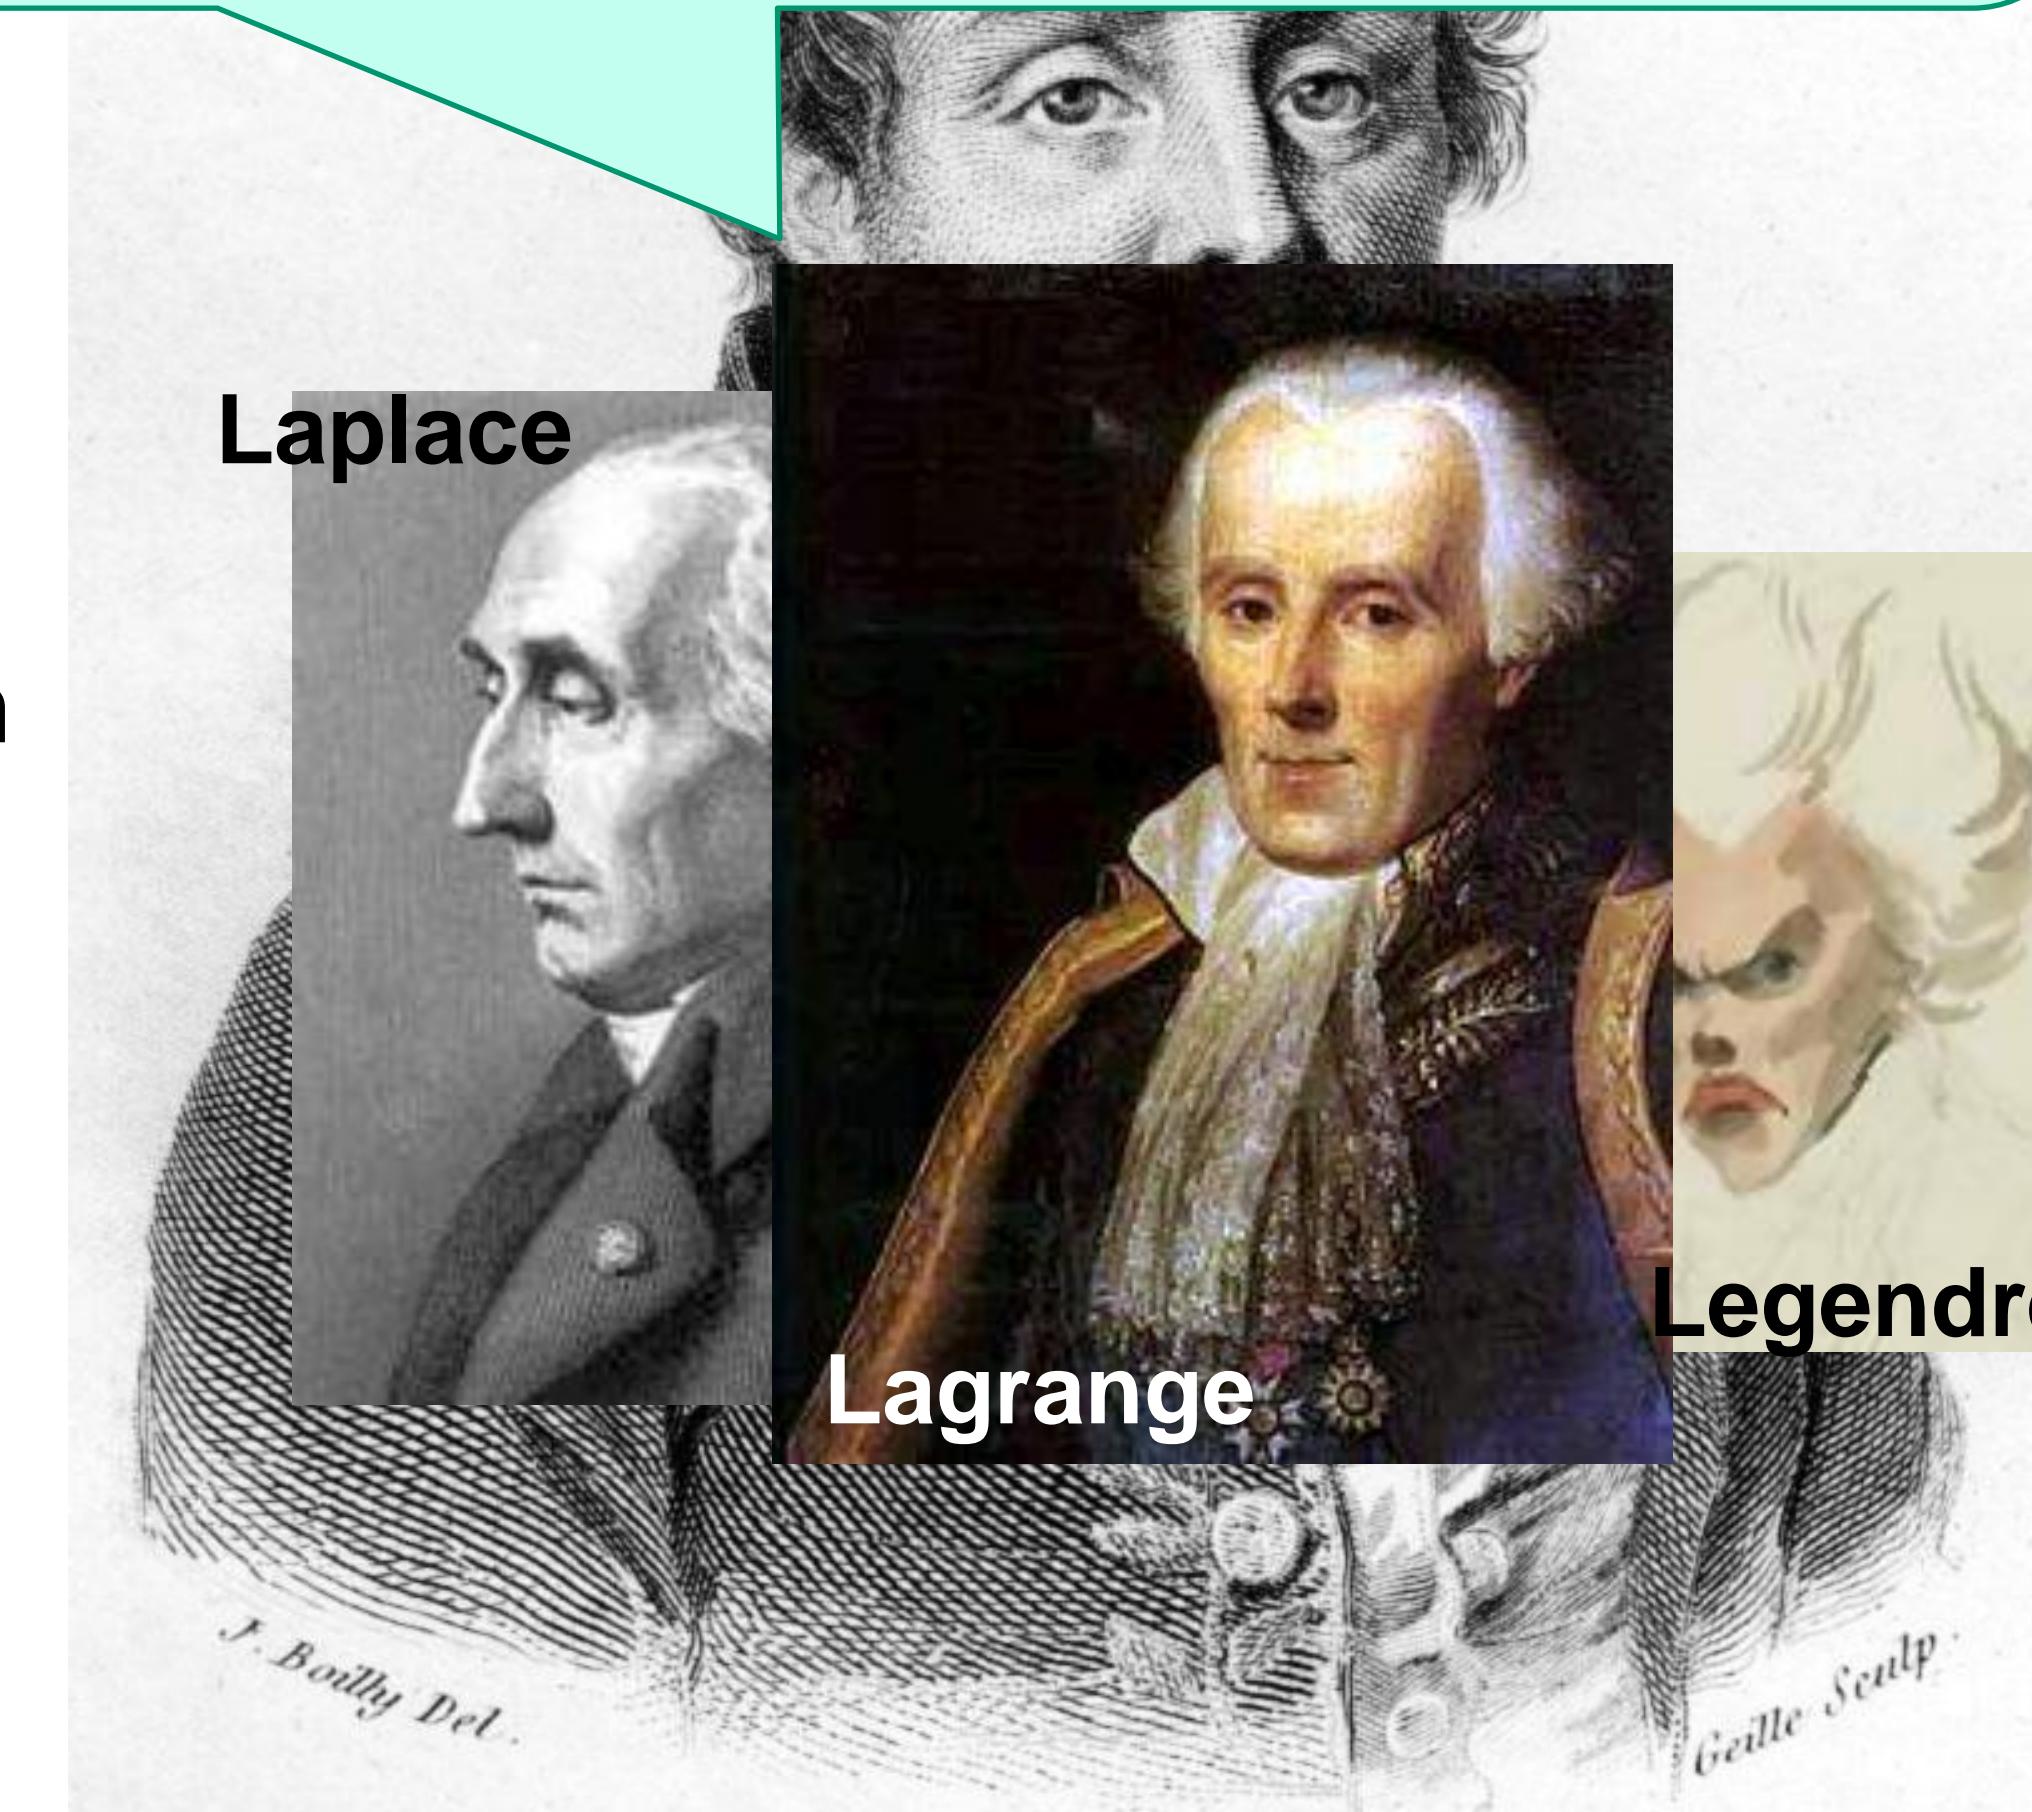
Don't believe it?

- Neither did Lagrange, Laplace, Poisson and other big wigs
- Not translated into English until 1878!

But it's (mostly) true!

- called Fourier Series

*...the manner in which the author arrives at these equations is not exempt of difficulties and... his analysis to integrate them still leaves something to be desired on the score of generality and even rigour.*



# A sum of sines

Our building block:

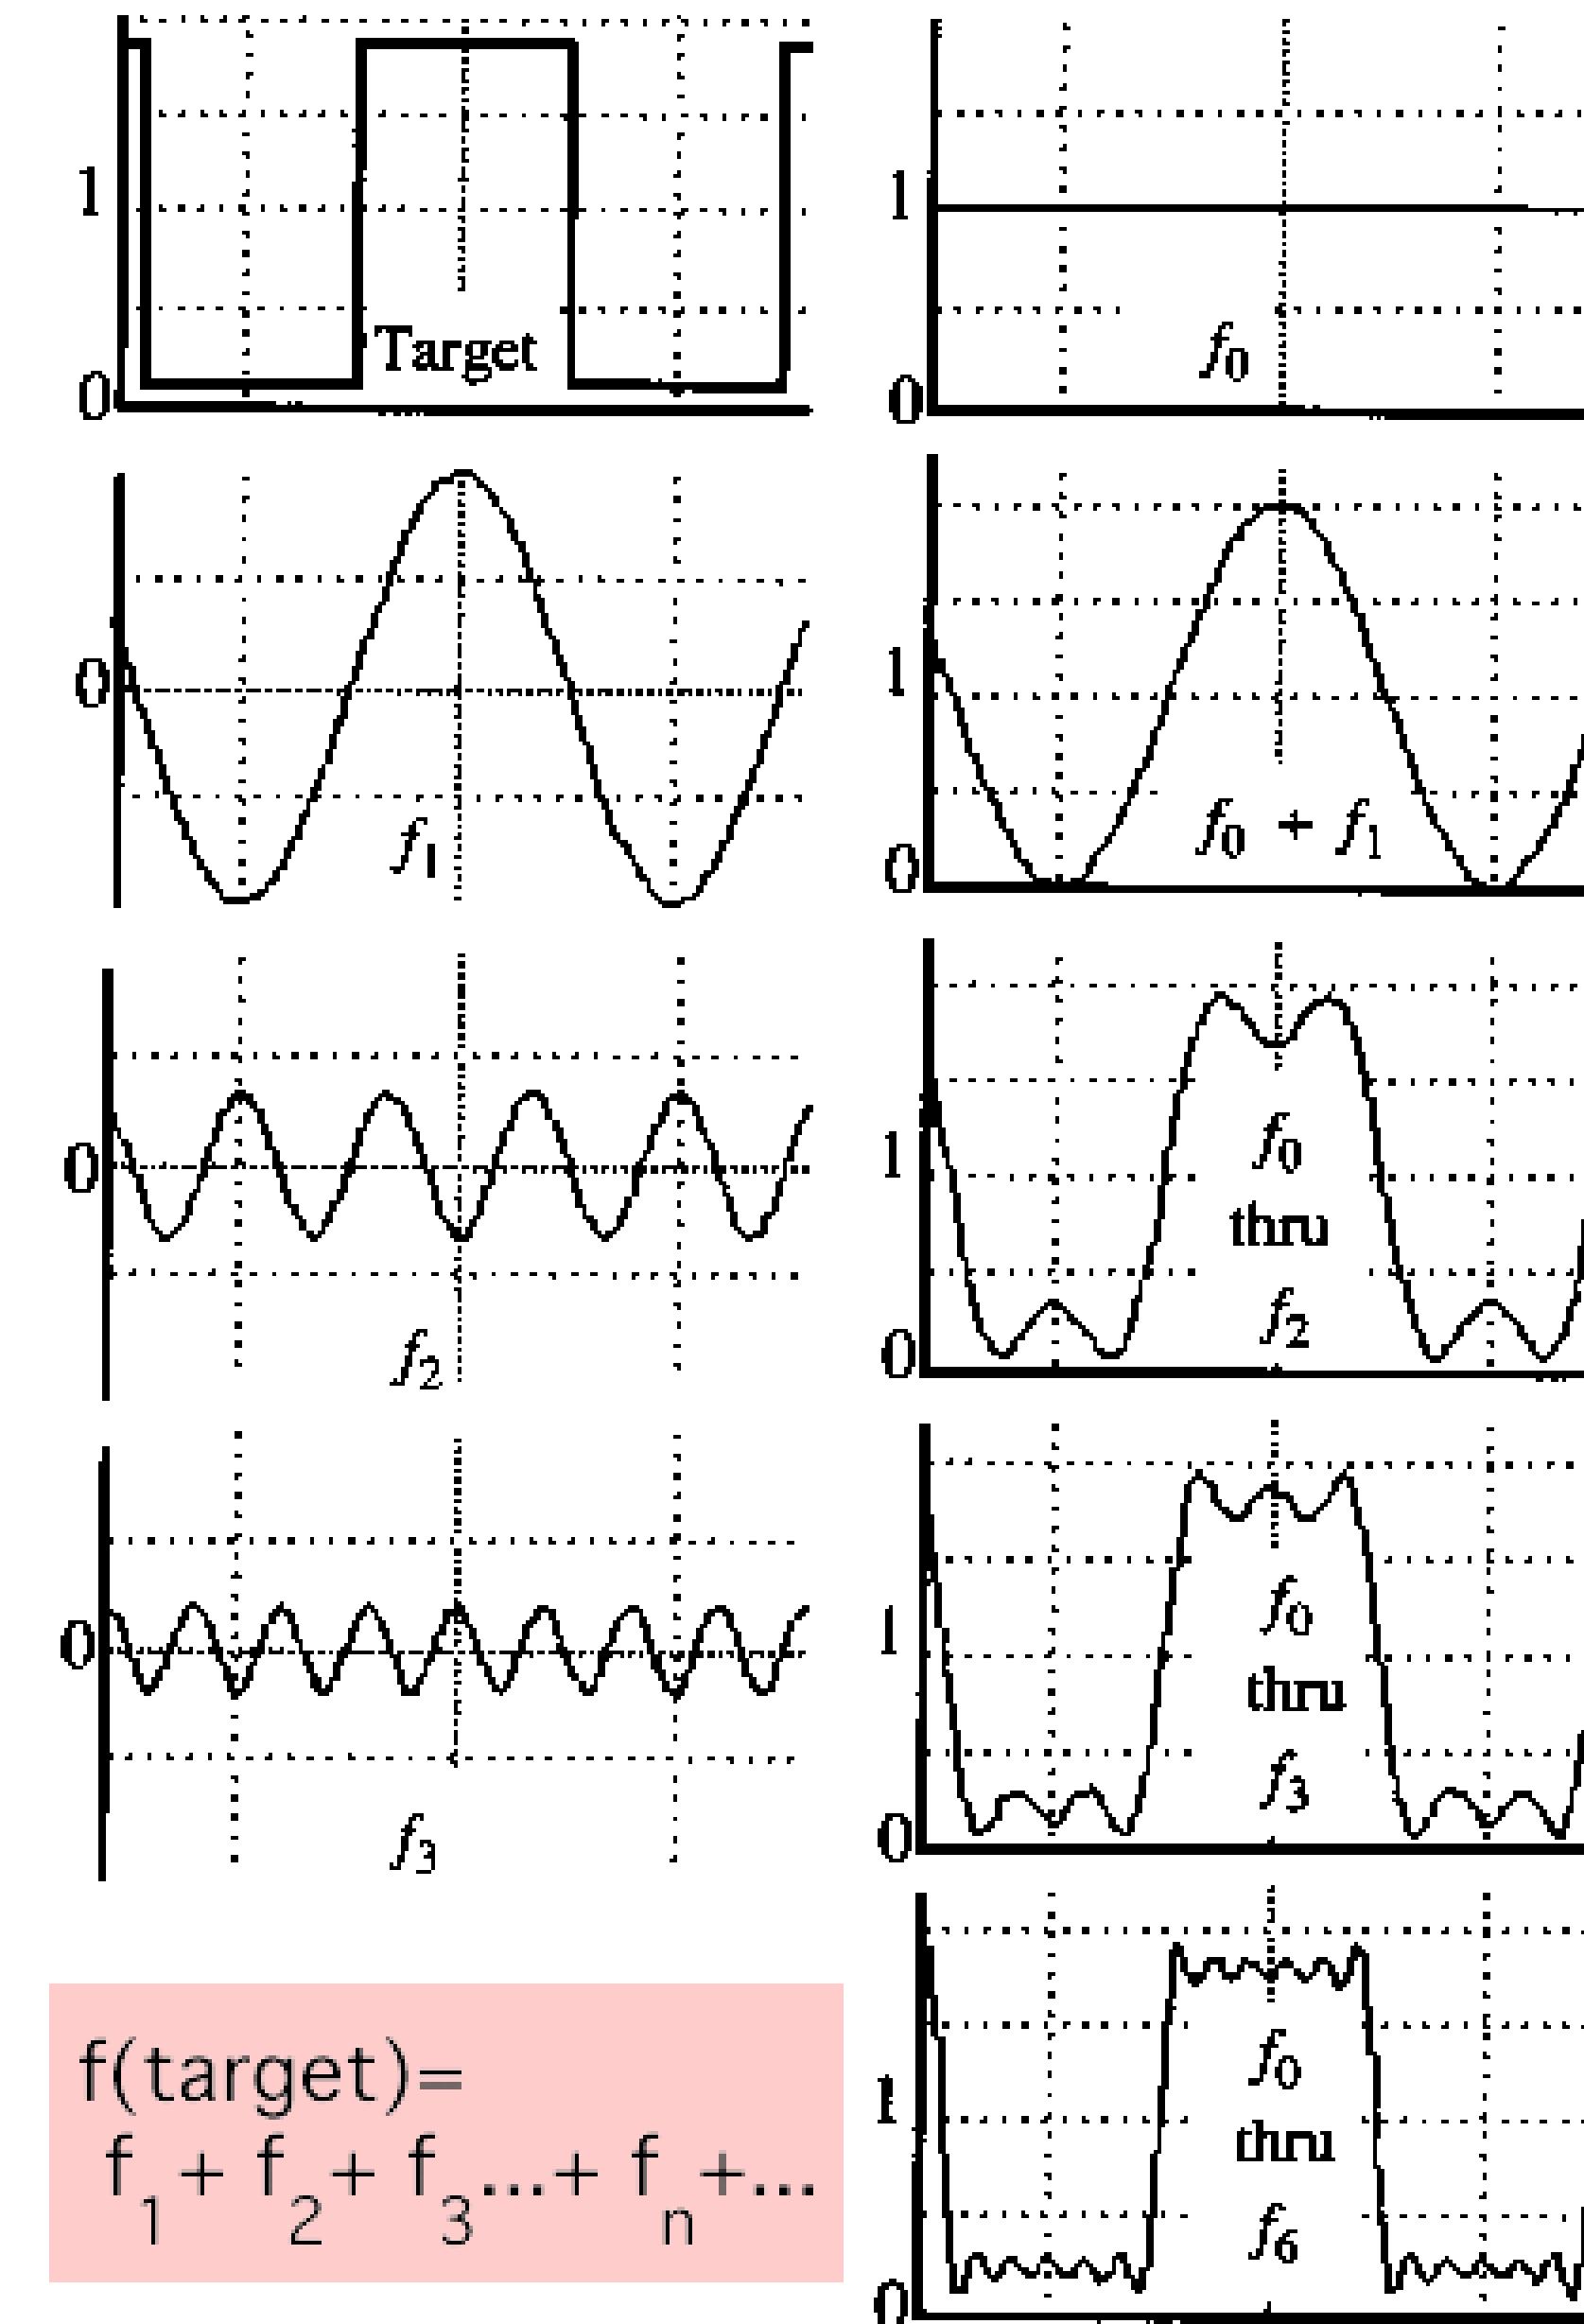
$$A \sin(\omega x + \phi)$$

Add enough of them to get any signal  $f(x)$  you want!

How many degrees of freedom?

What does each control?

Which one encodes the coarse vs. fine structure of the signal?



# Fourier Transform

---

We want to understand the frequency  $\omega$  of our signal. So, let's reparametrize the signal by  $\omega$  instead of  $x$ :



For every  $\omega$  from 0 to inf,  $F(\omega)$  holds the amplitude  $A$  and phase  $\phi$  of the corresponding sine  $A\sin(\omega x + \phi)$

- How does  $F$  hold both?

$$\begin{aligned} F(\omega) &= R(\omega) + iI(\omega) \\ A &= \pm \sqrt{R(\omega)^2 + I(\omega)^2} \quad \phi = \tan^{-1} \frac{I(\omega)}{R(\omega)} \end{aligned}$$

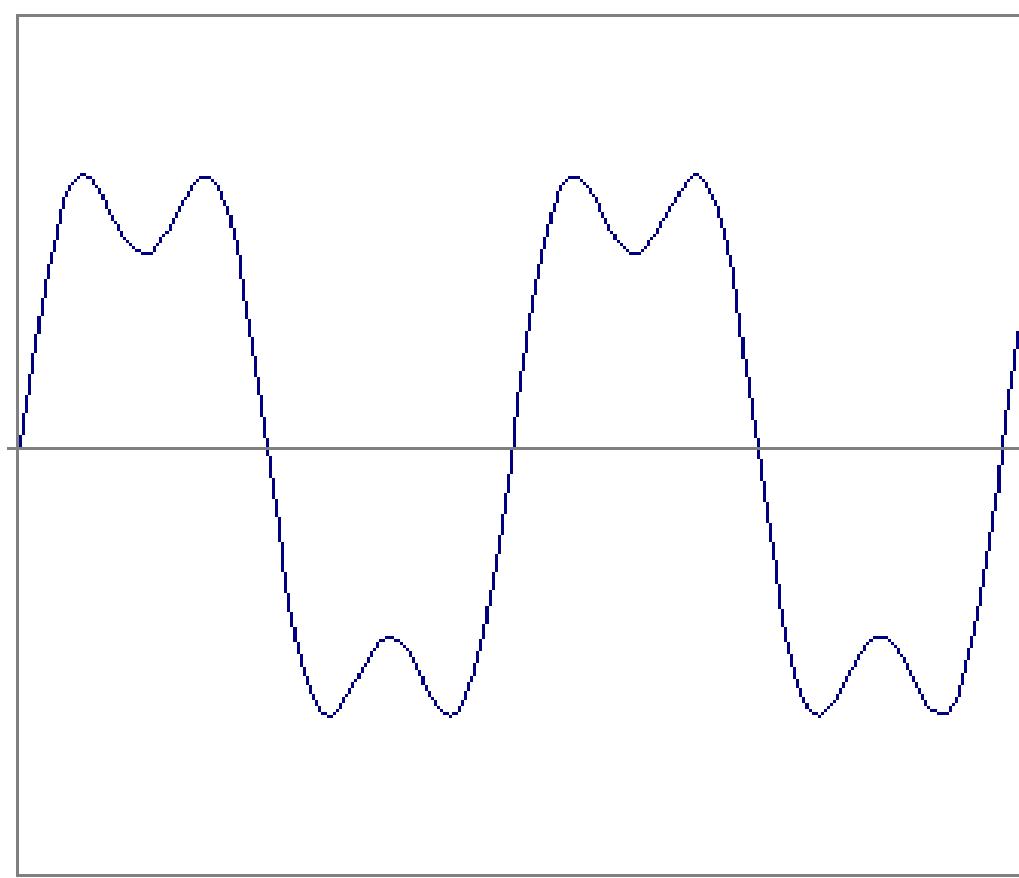
We can always go back:



# Time and Frequency

---

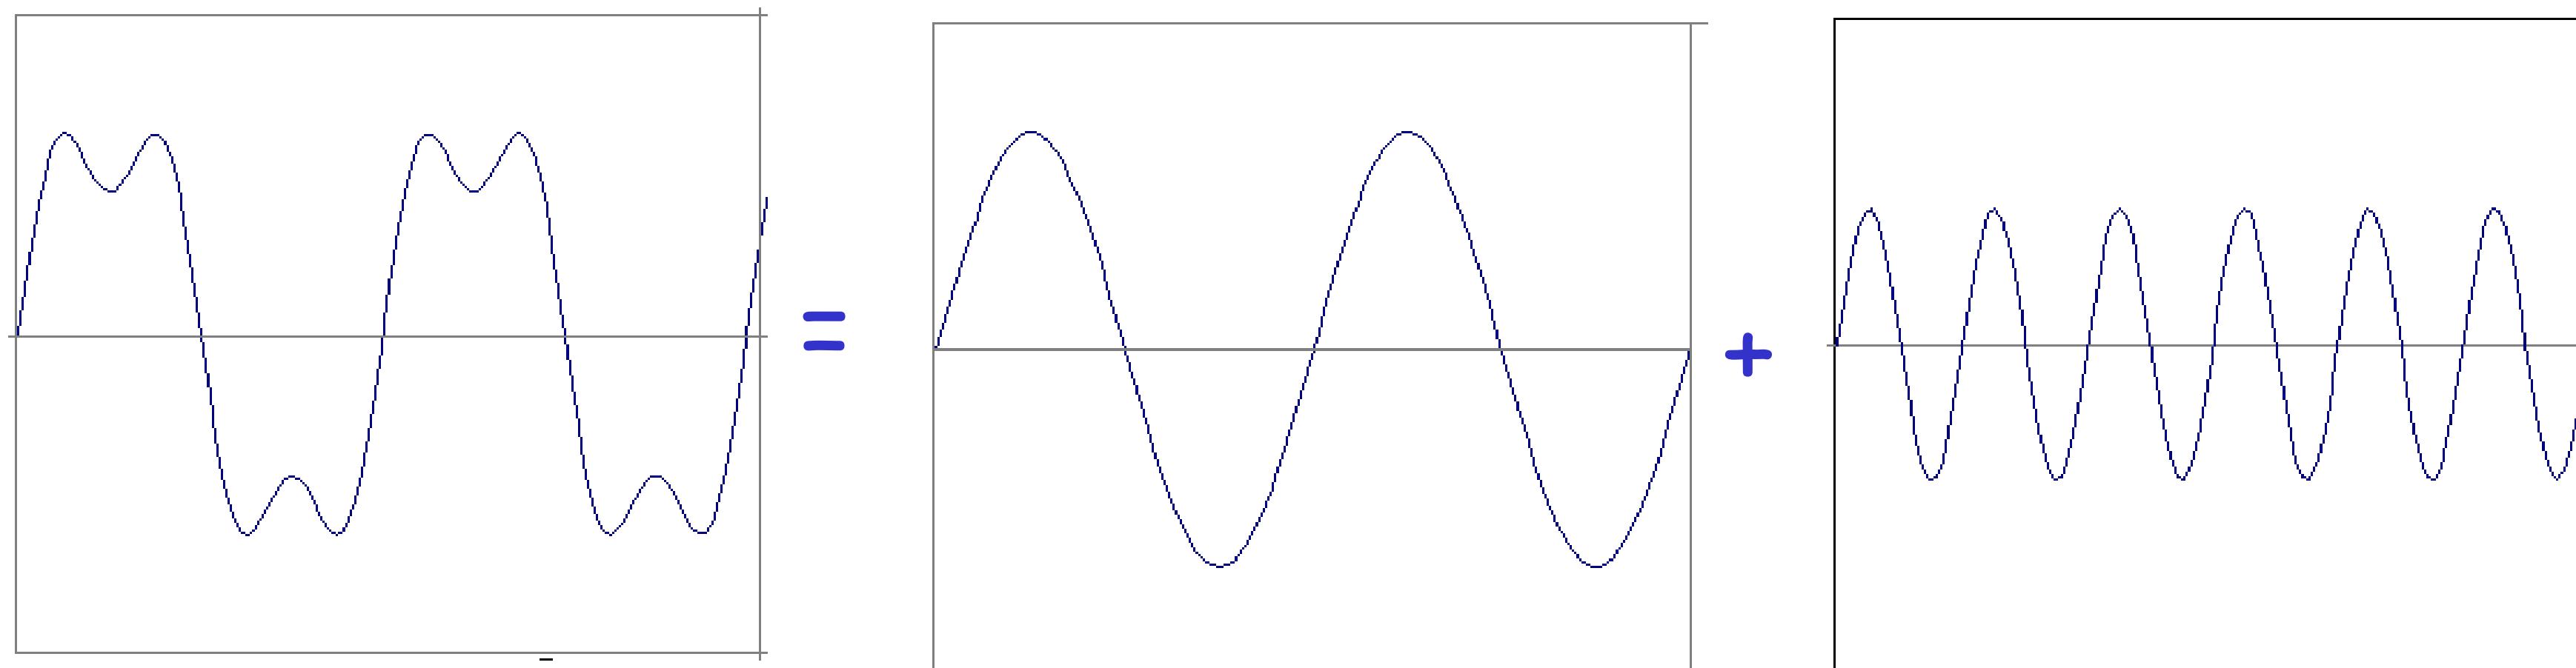
example :  $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f)t)$



# Time and Frequency

---

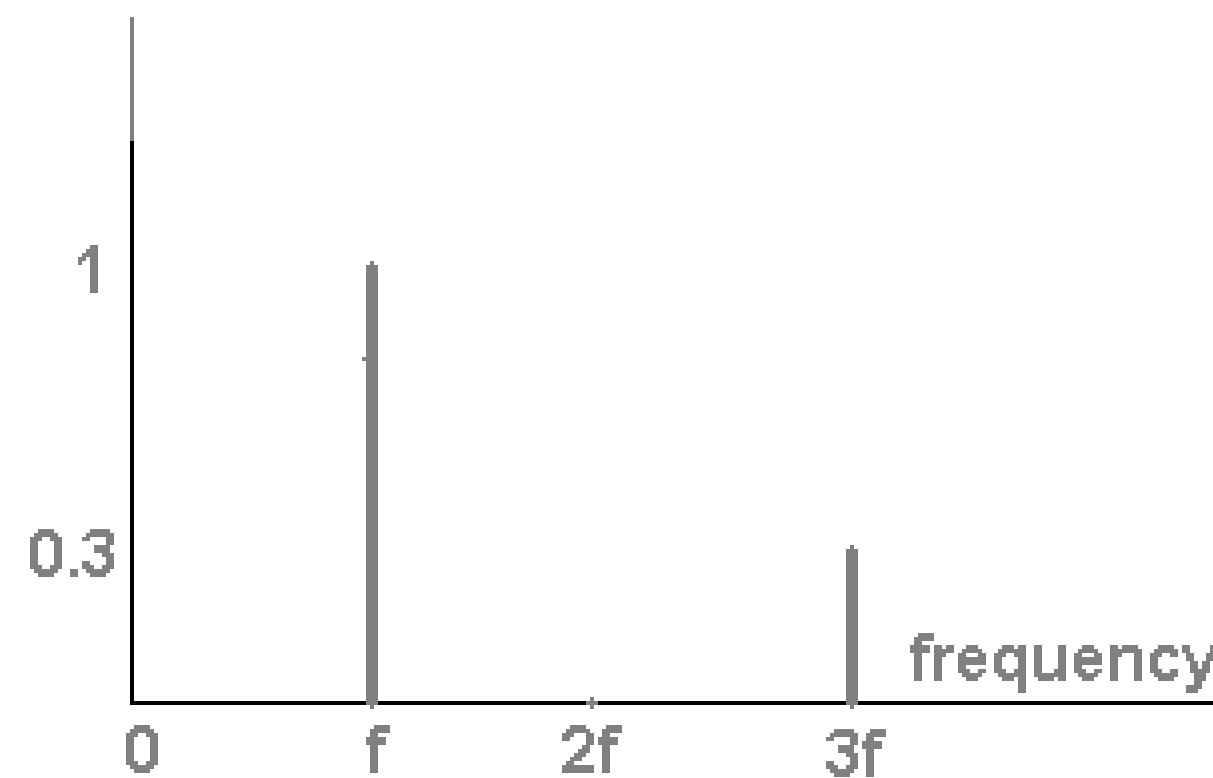
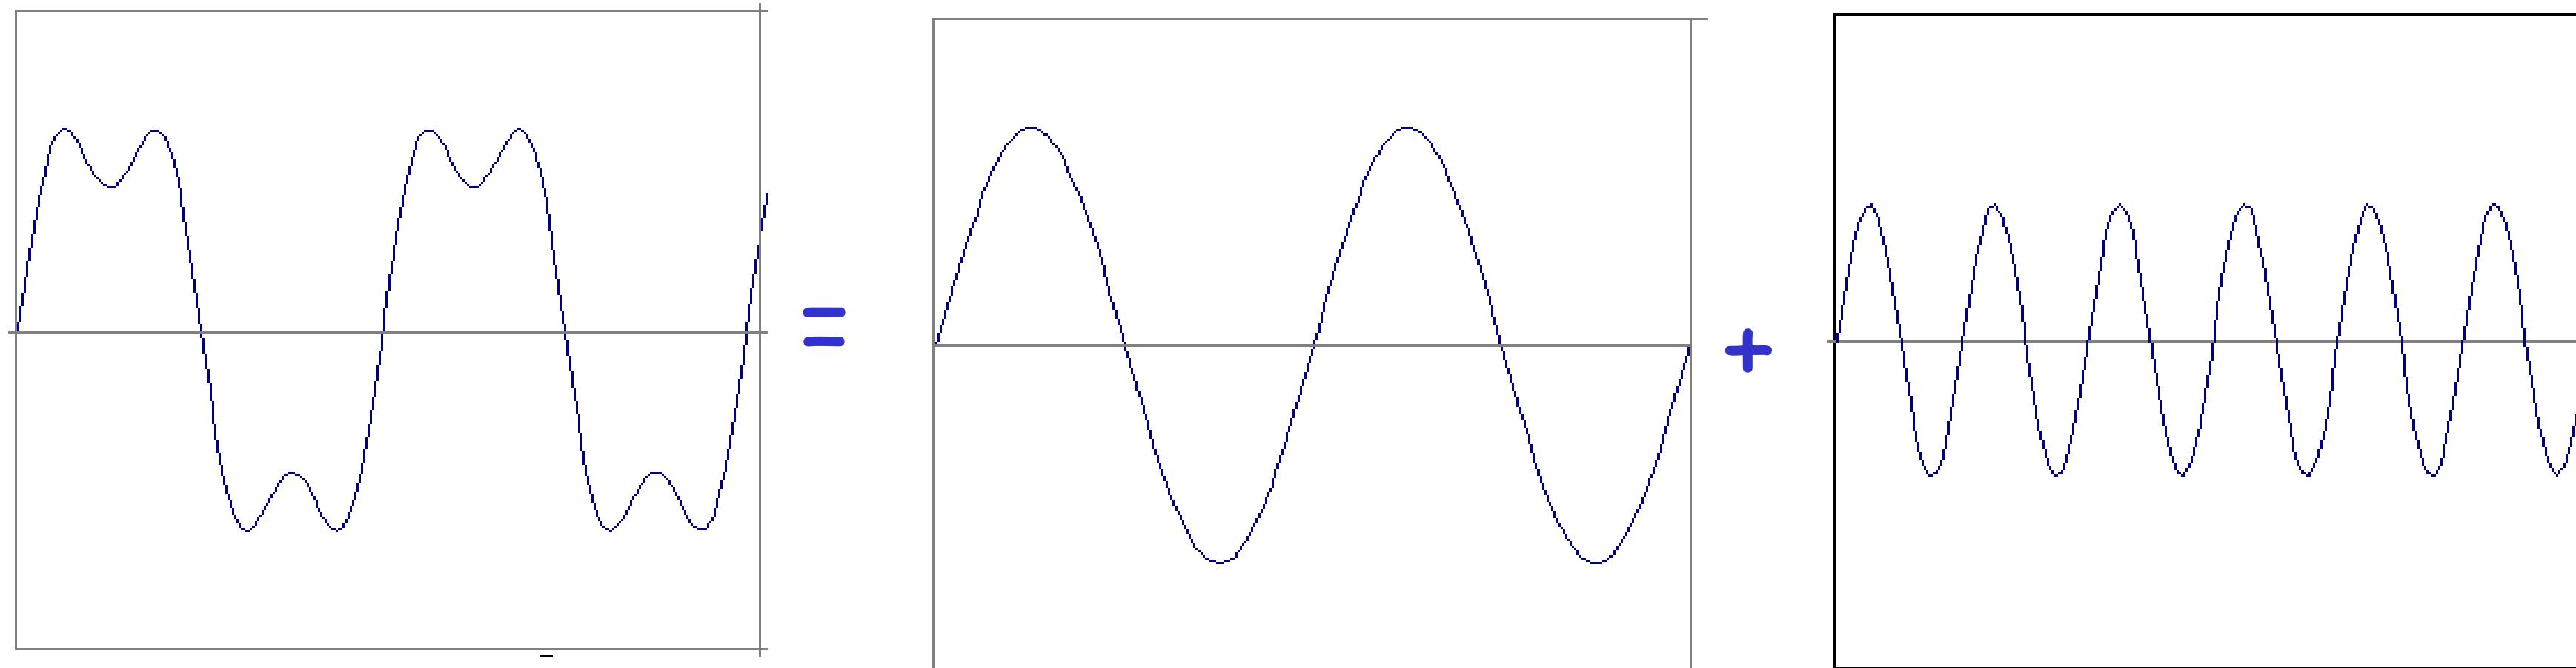
example :  $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f)t)$



# Frequency Spectra

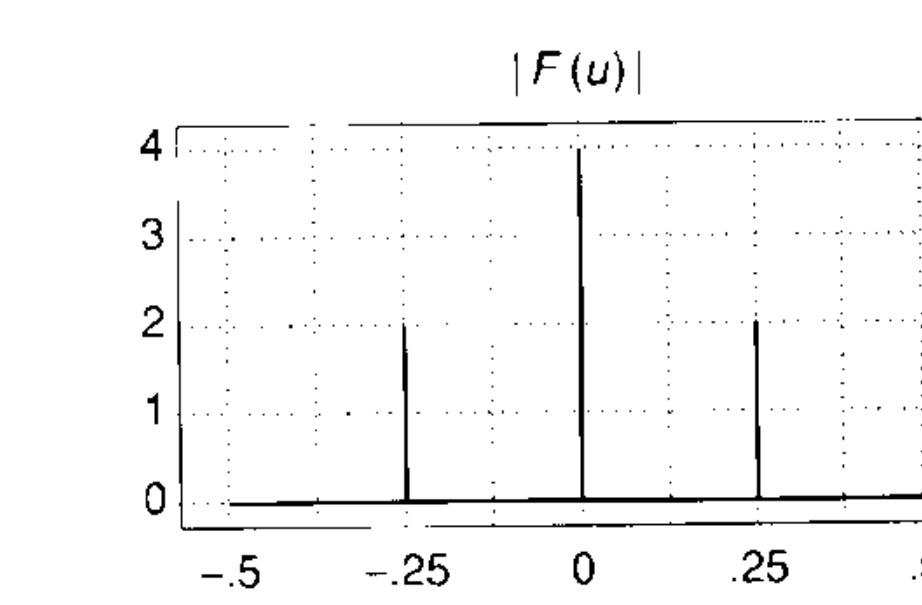
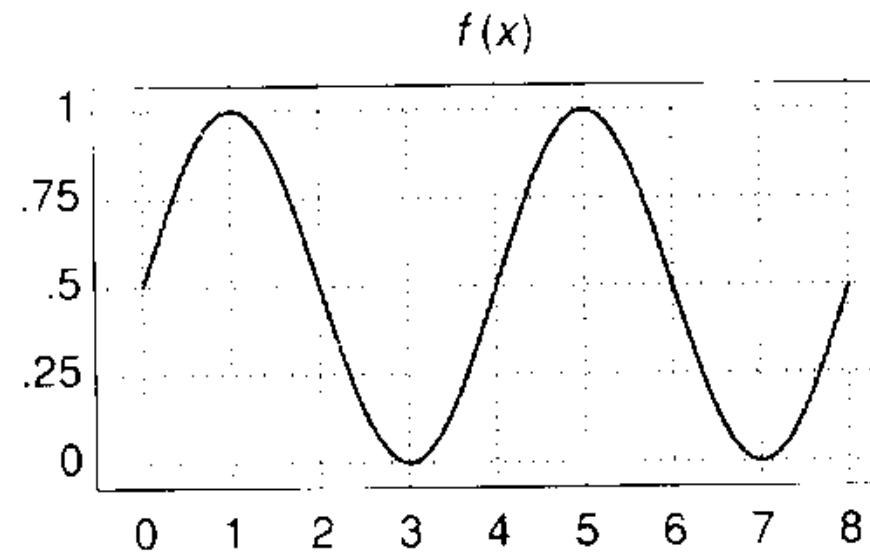
---

example :  $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f)t)$

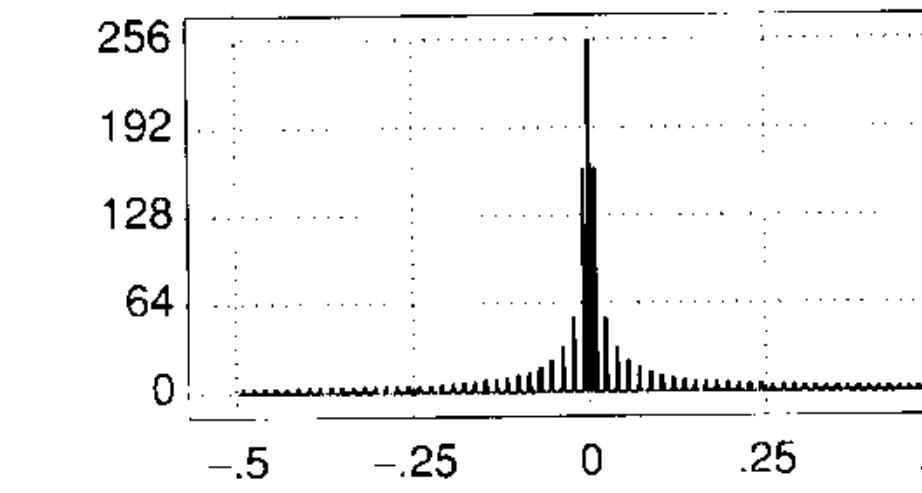
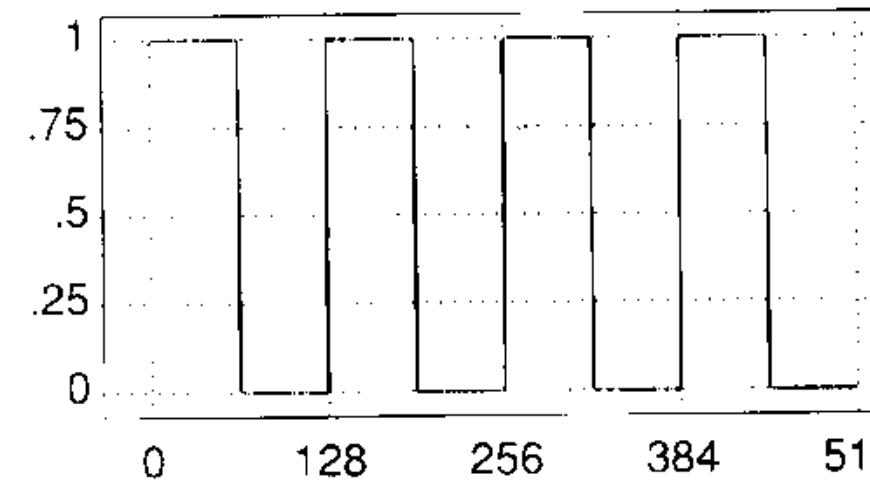


# Frequency Spectra

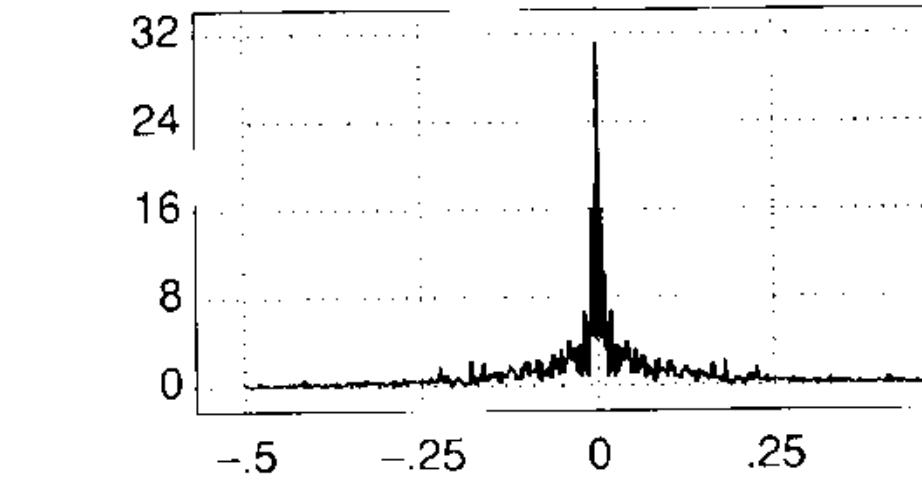
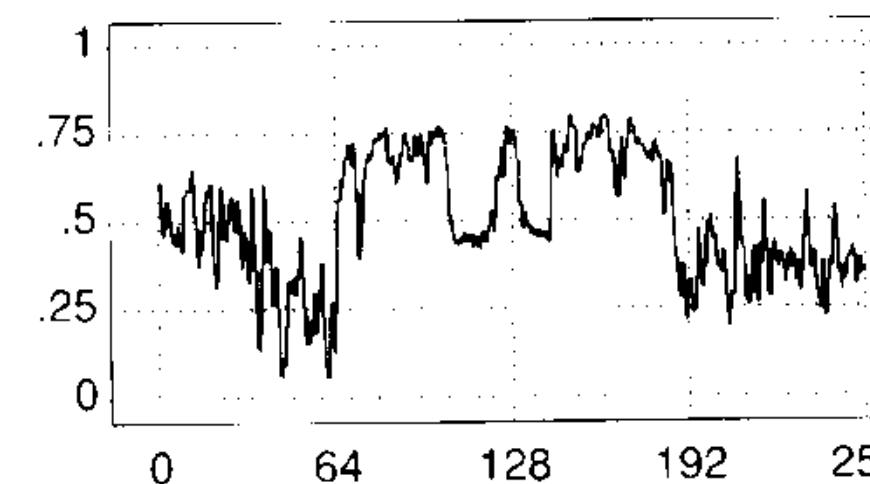
---



(a)



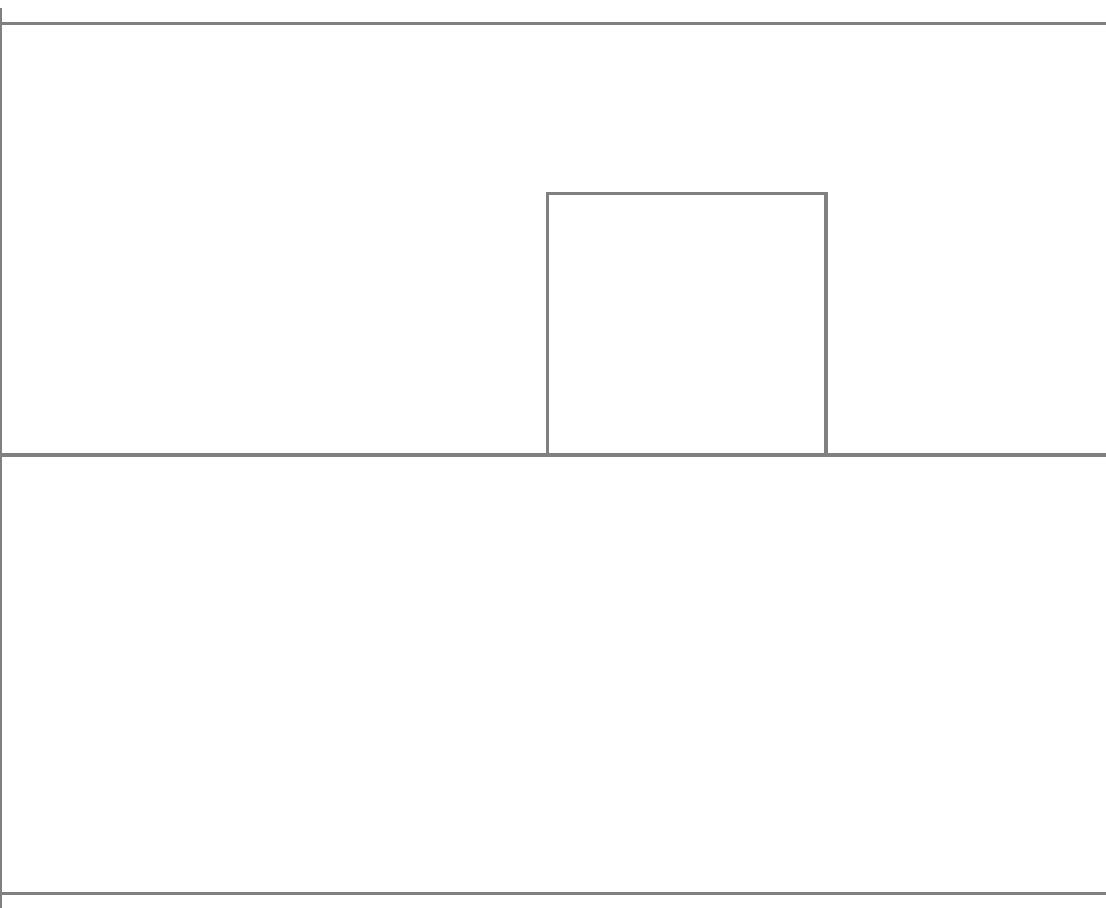
(b)



(c)

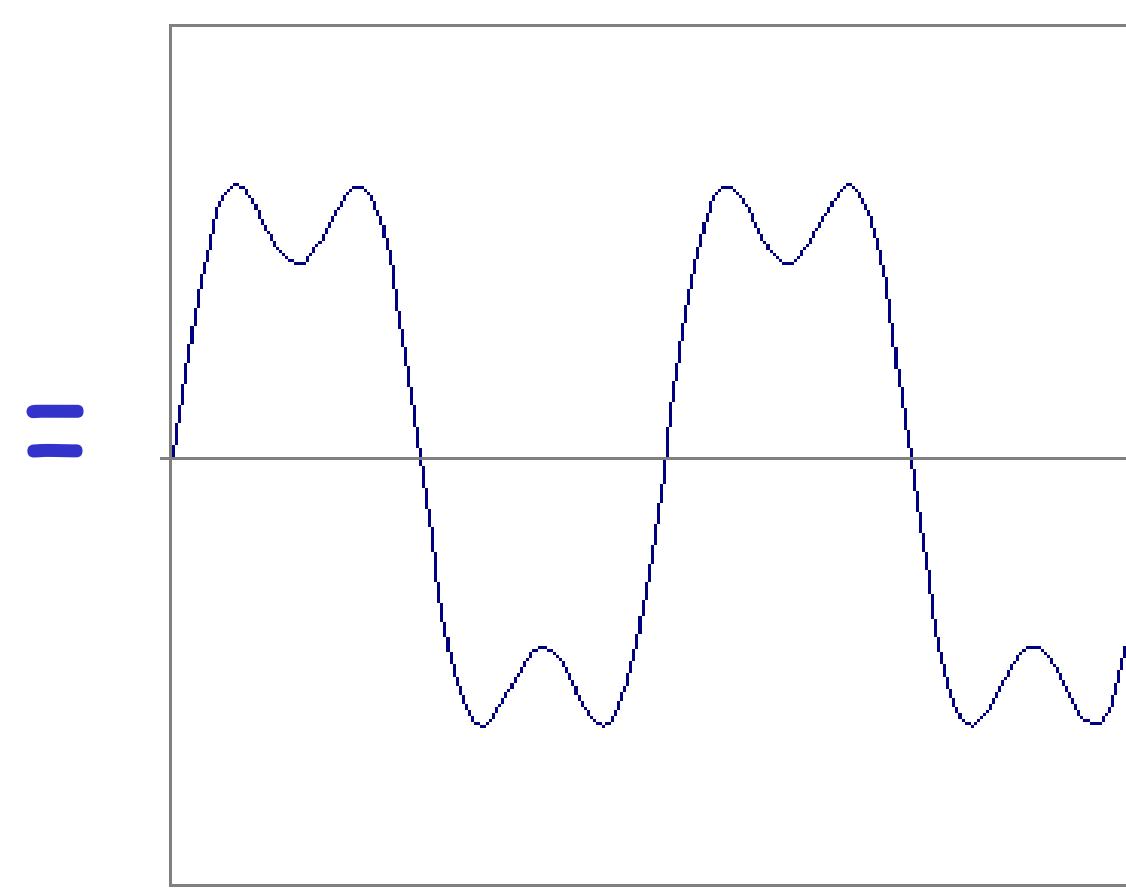
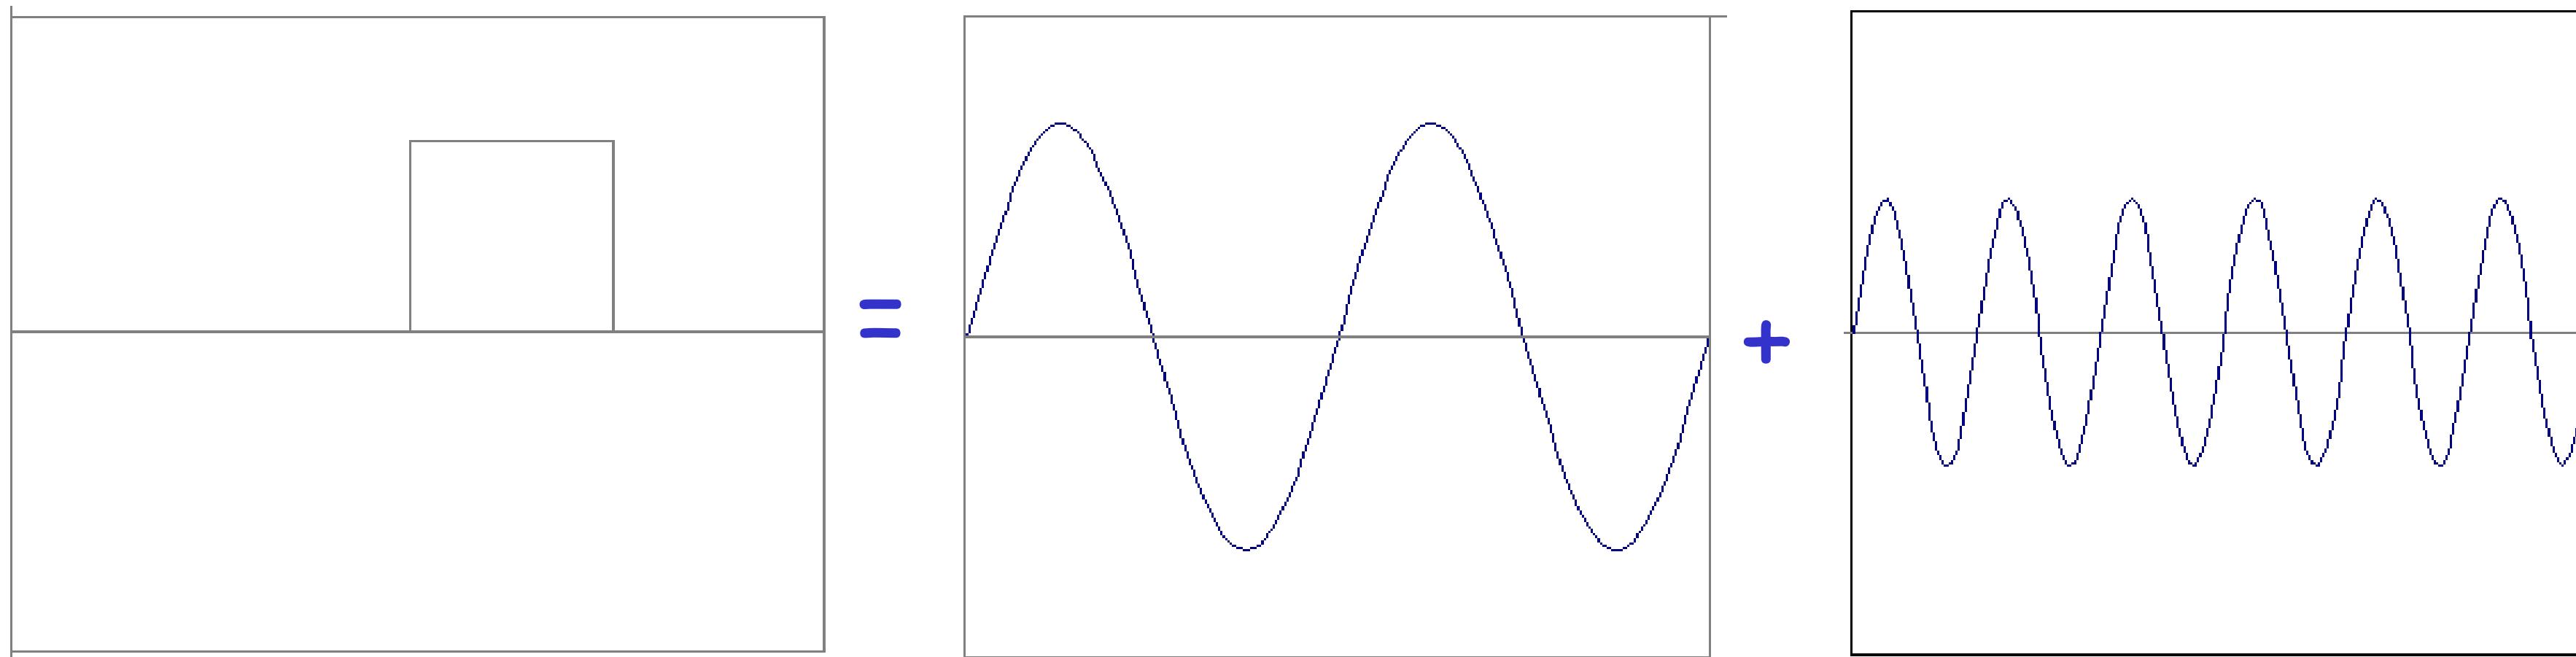
# Frequency Spectra

---



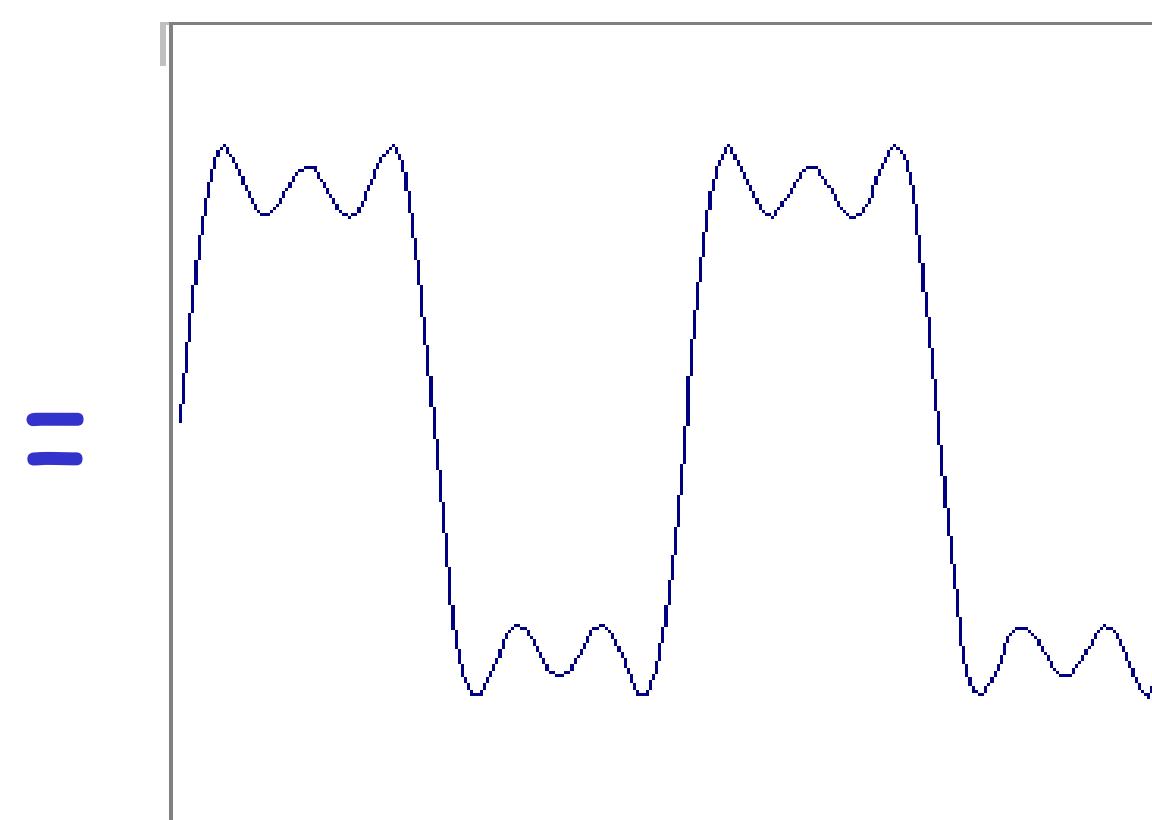
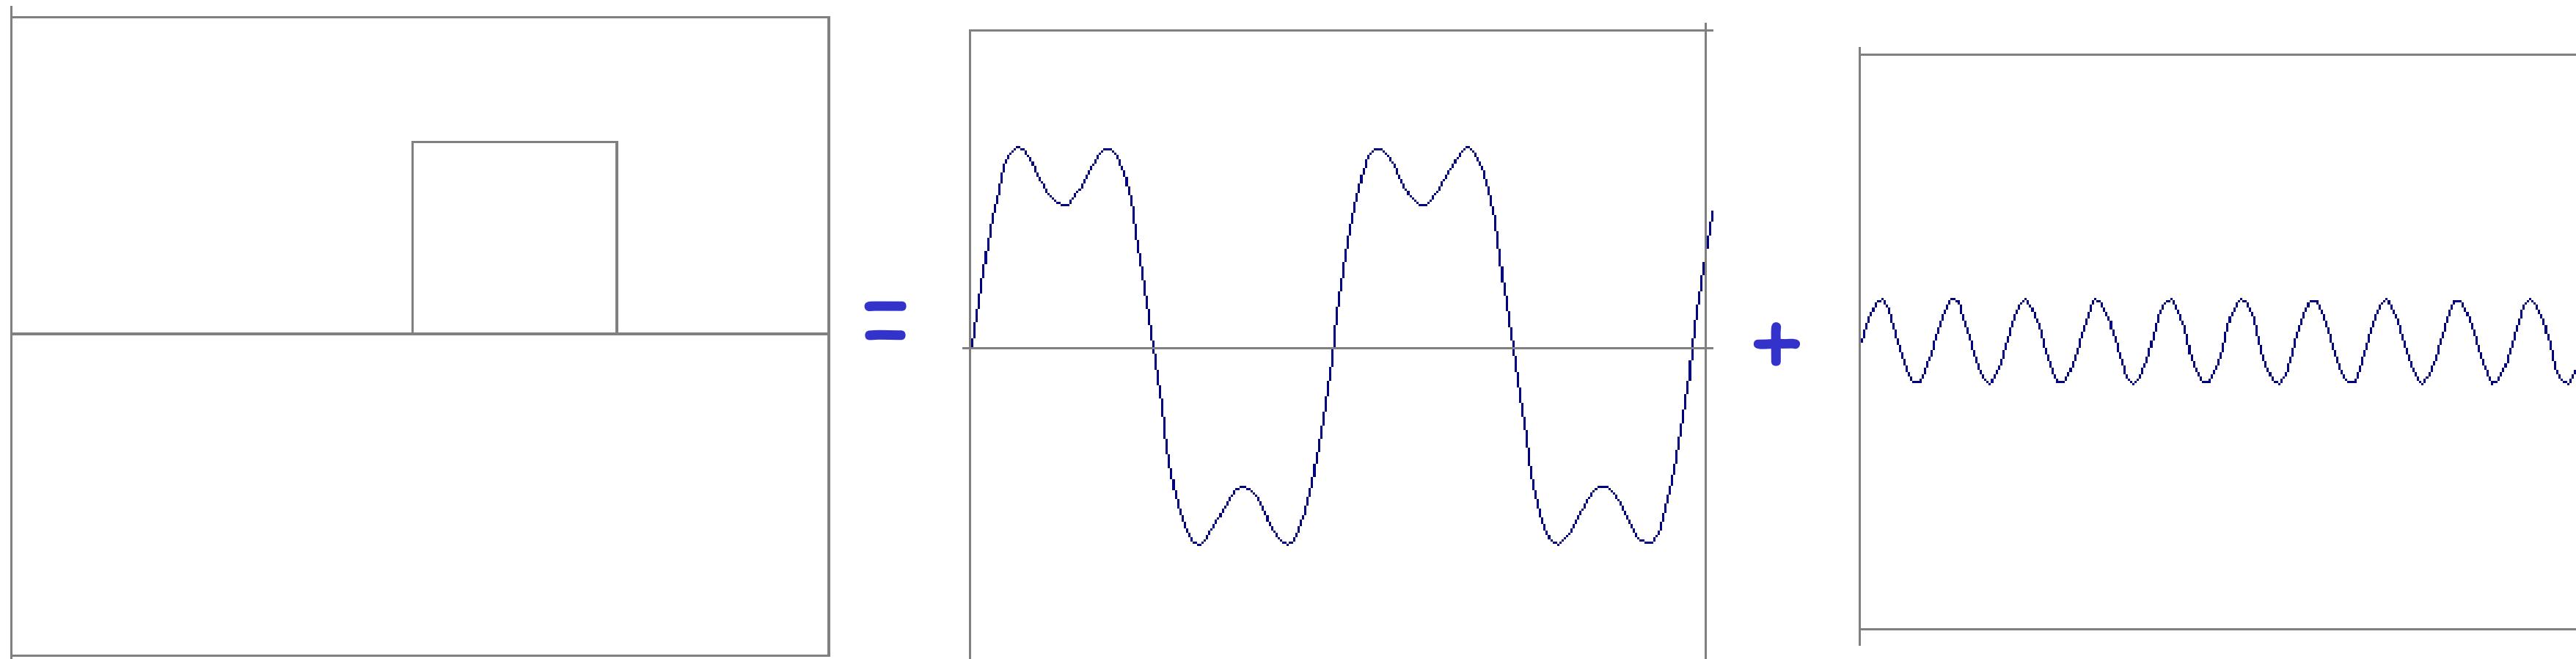
# Frequency Spectra

---



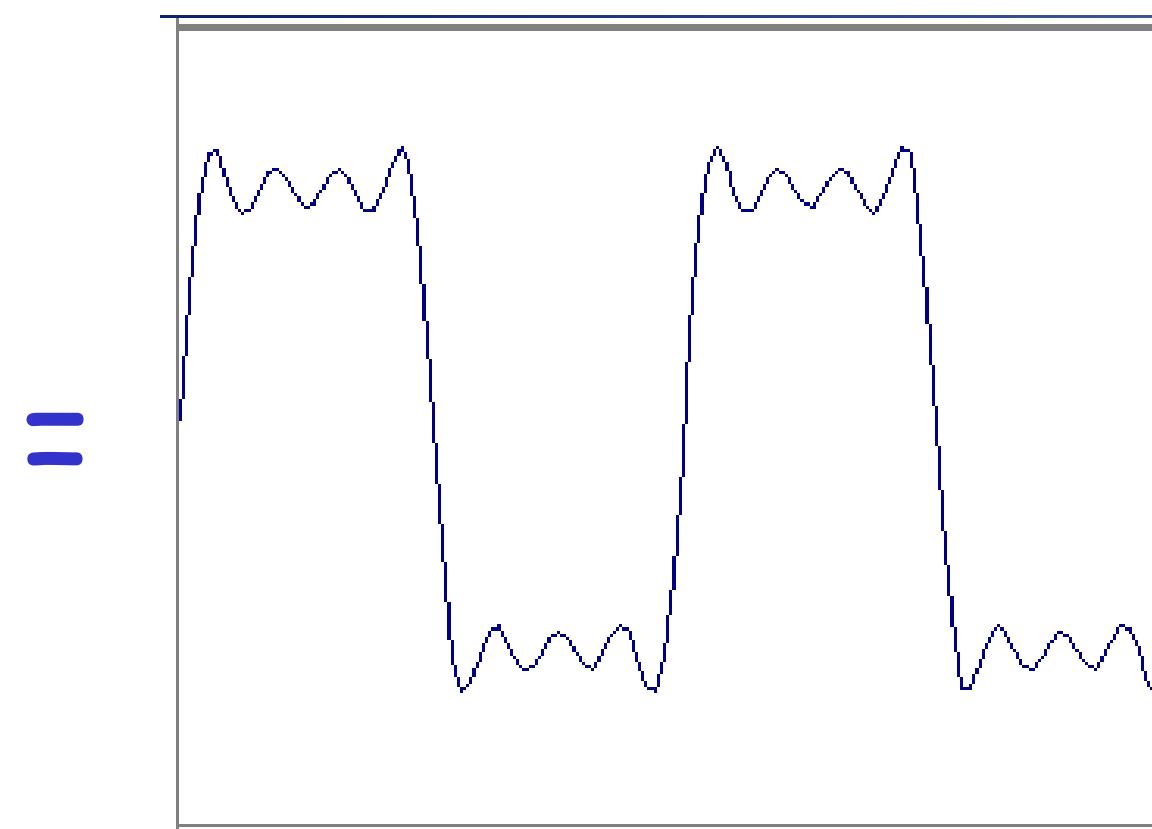
# Frequency Spectra

---



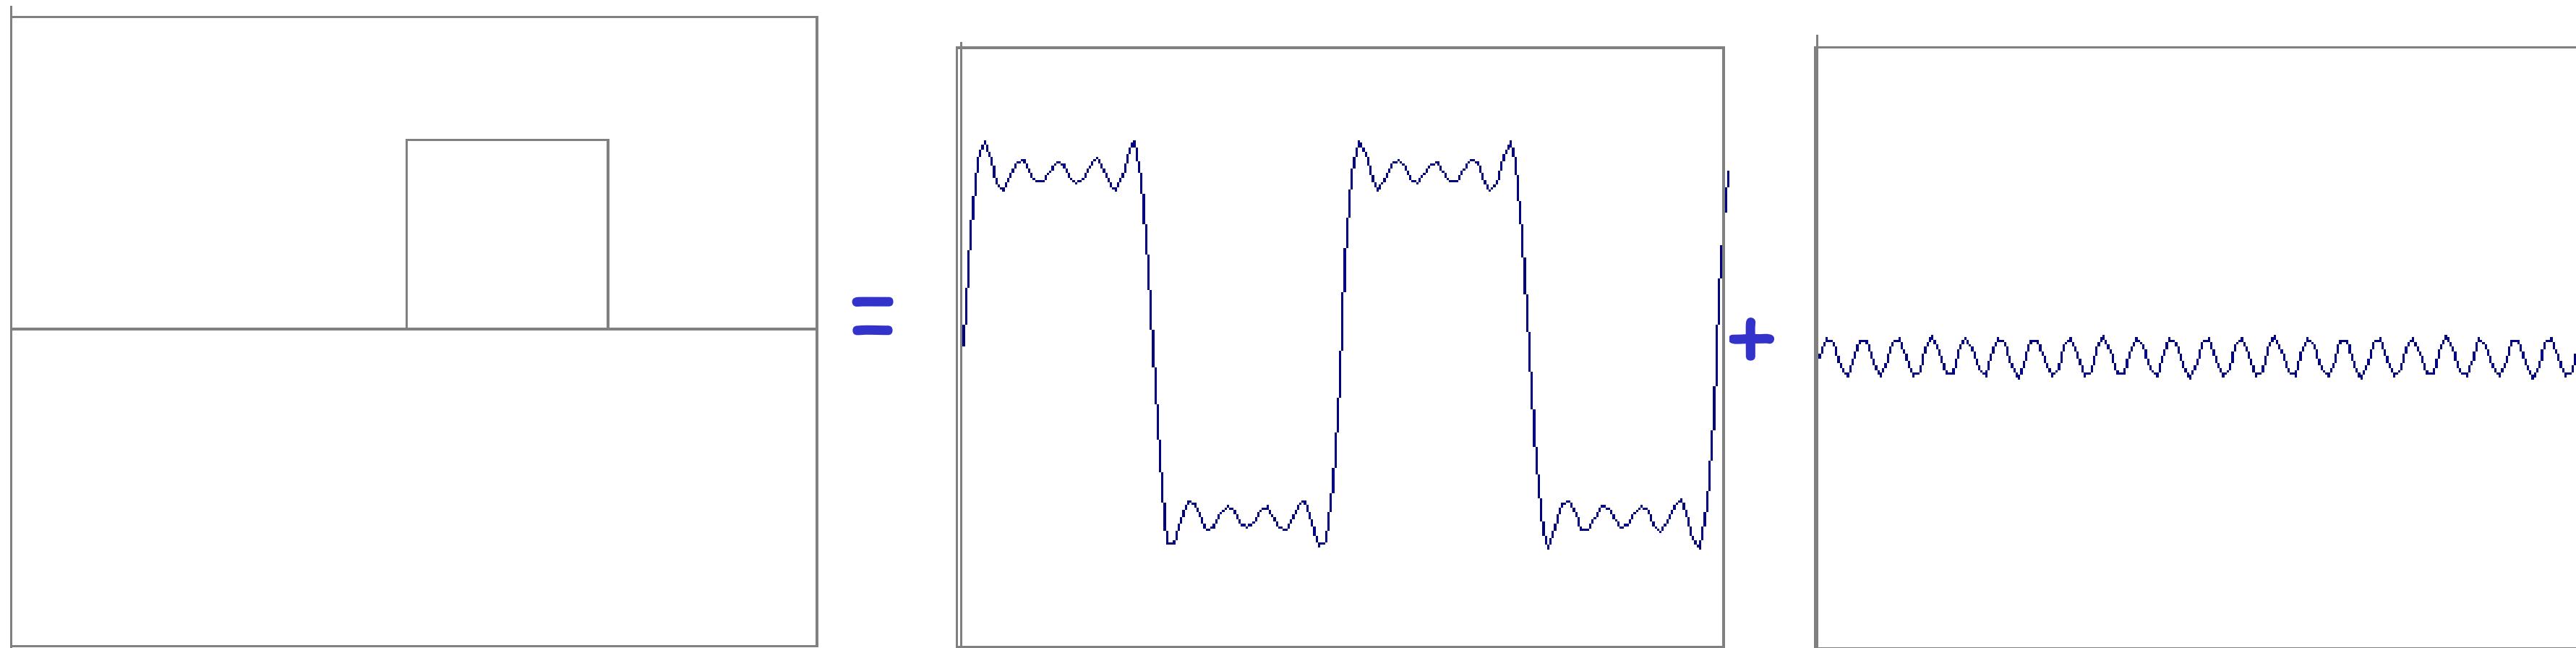
# Frequency Spectra

---



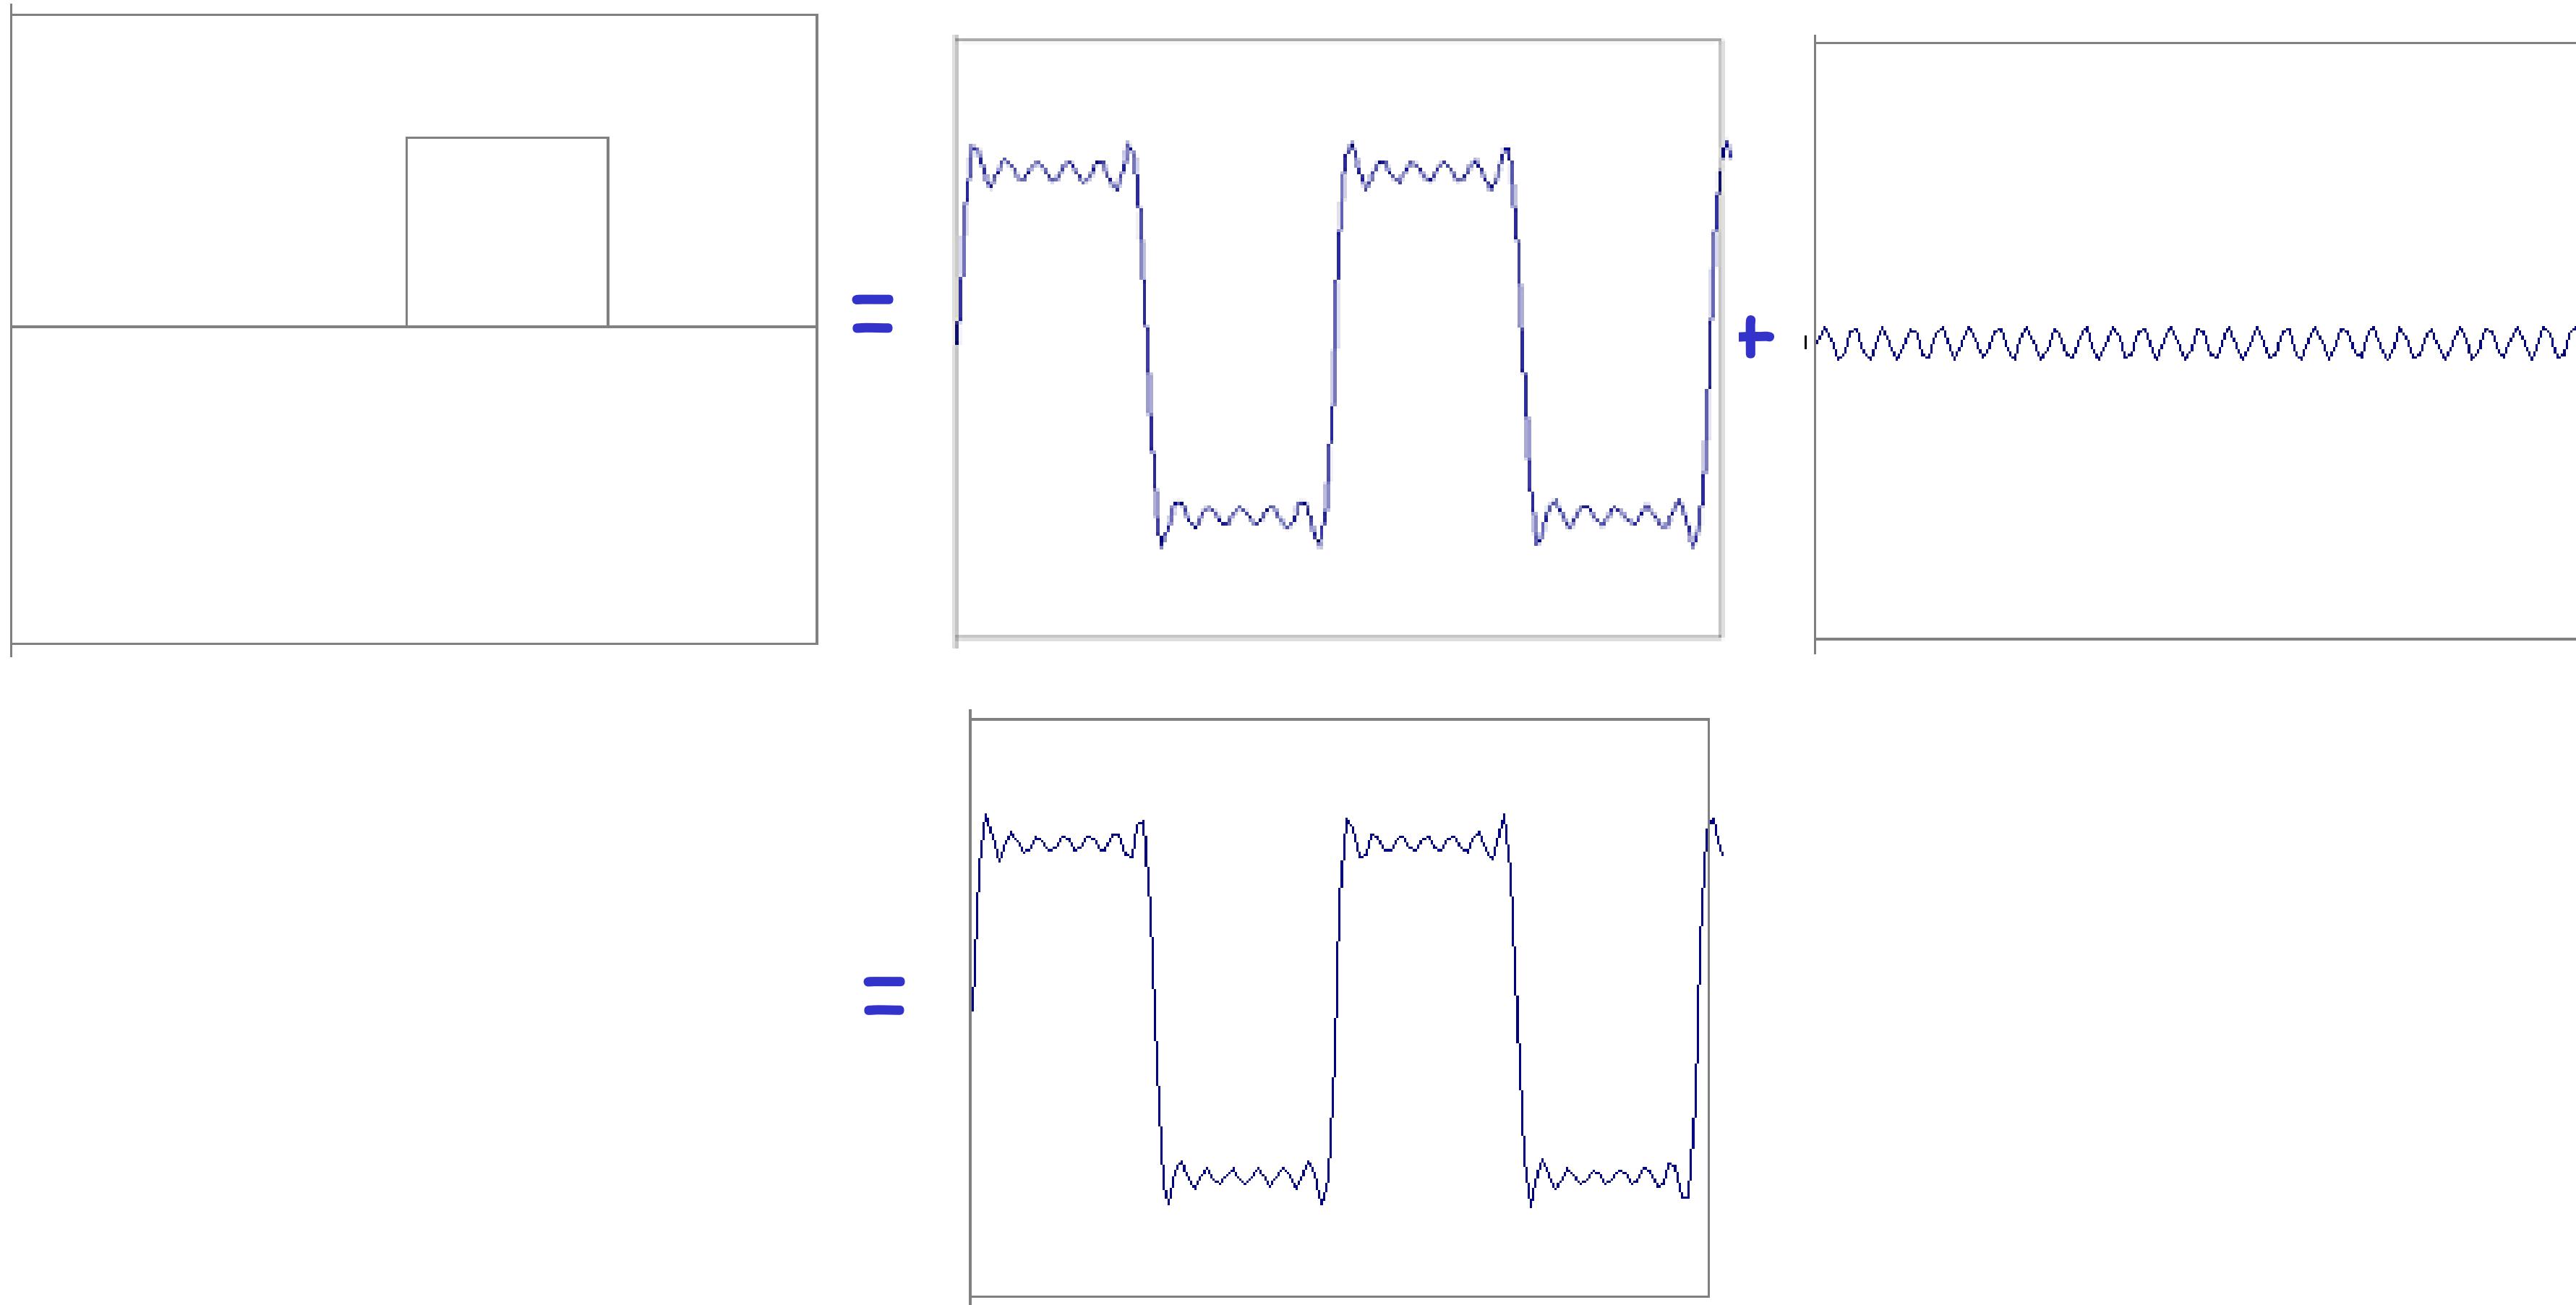
# Frequency Spectra

---



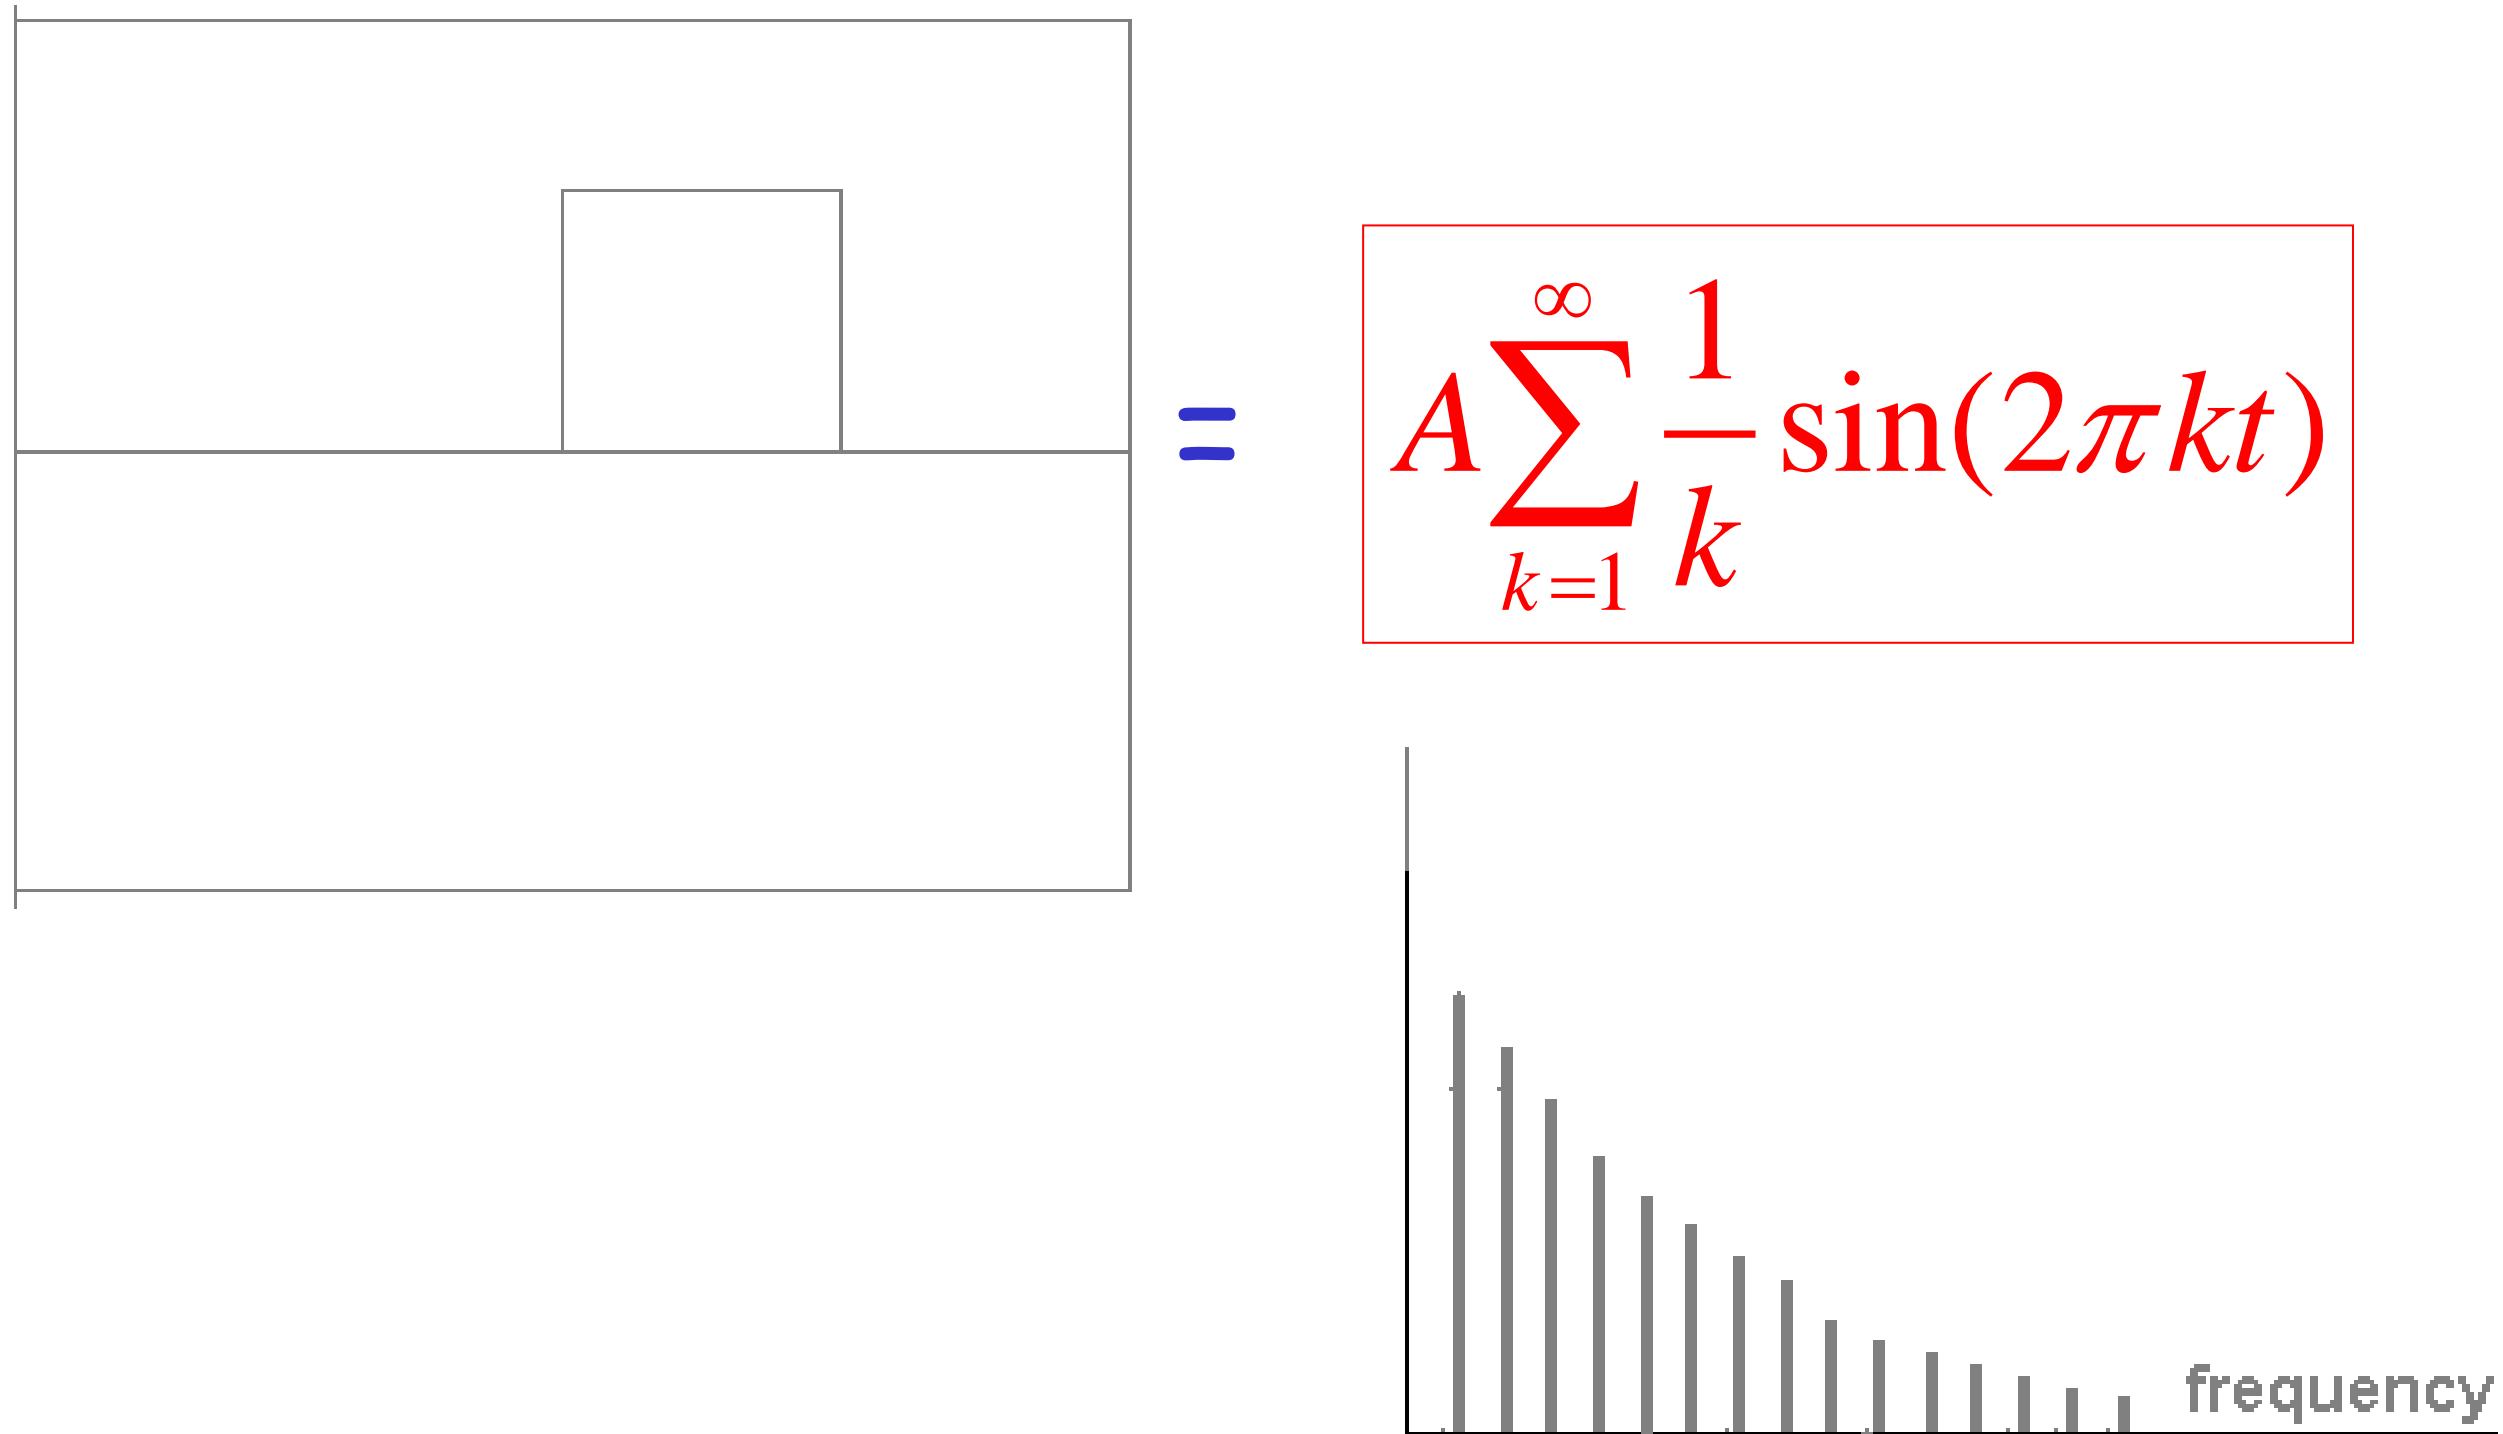
# Frequency Spectra

---



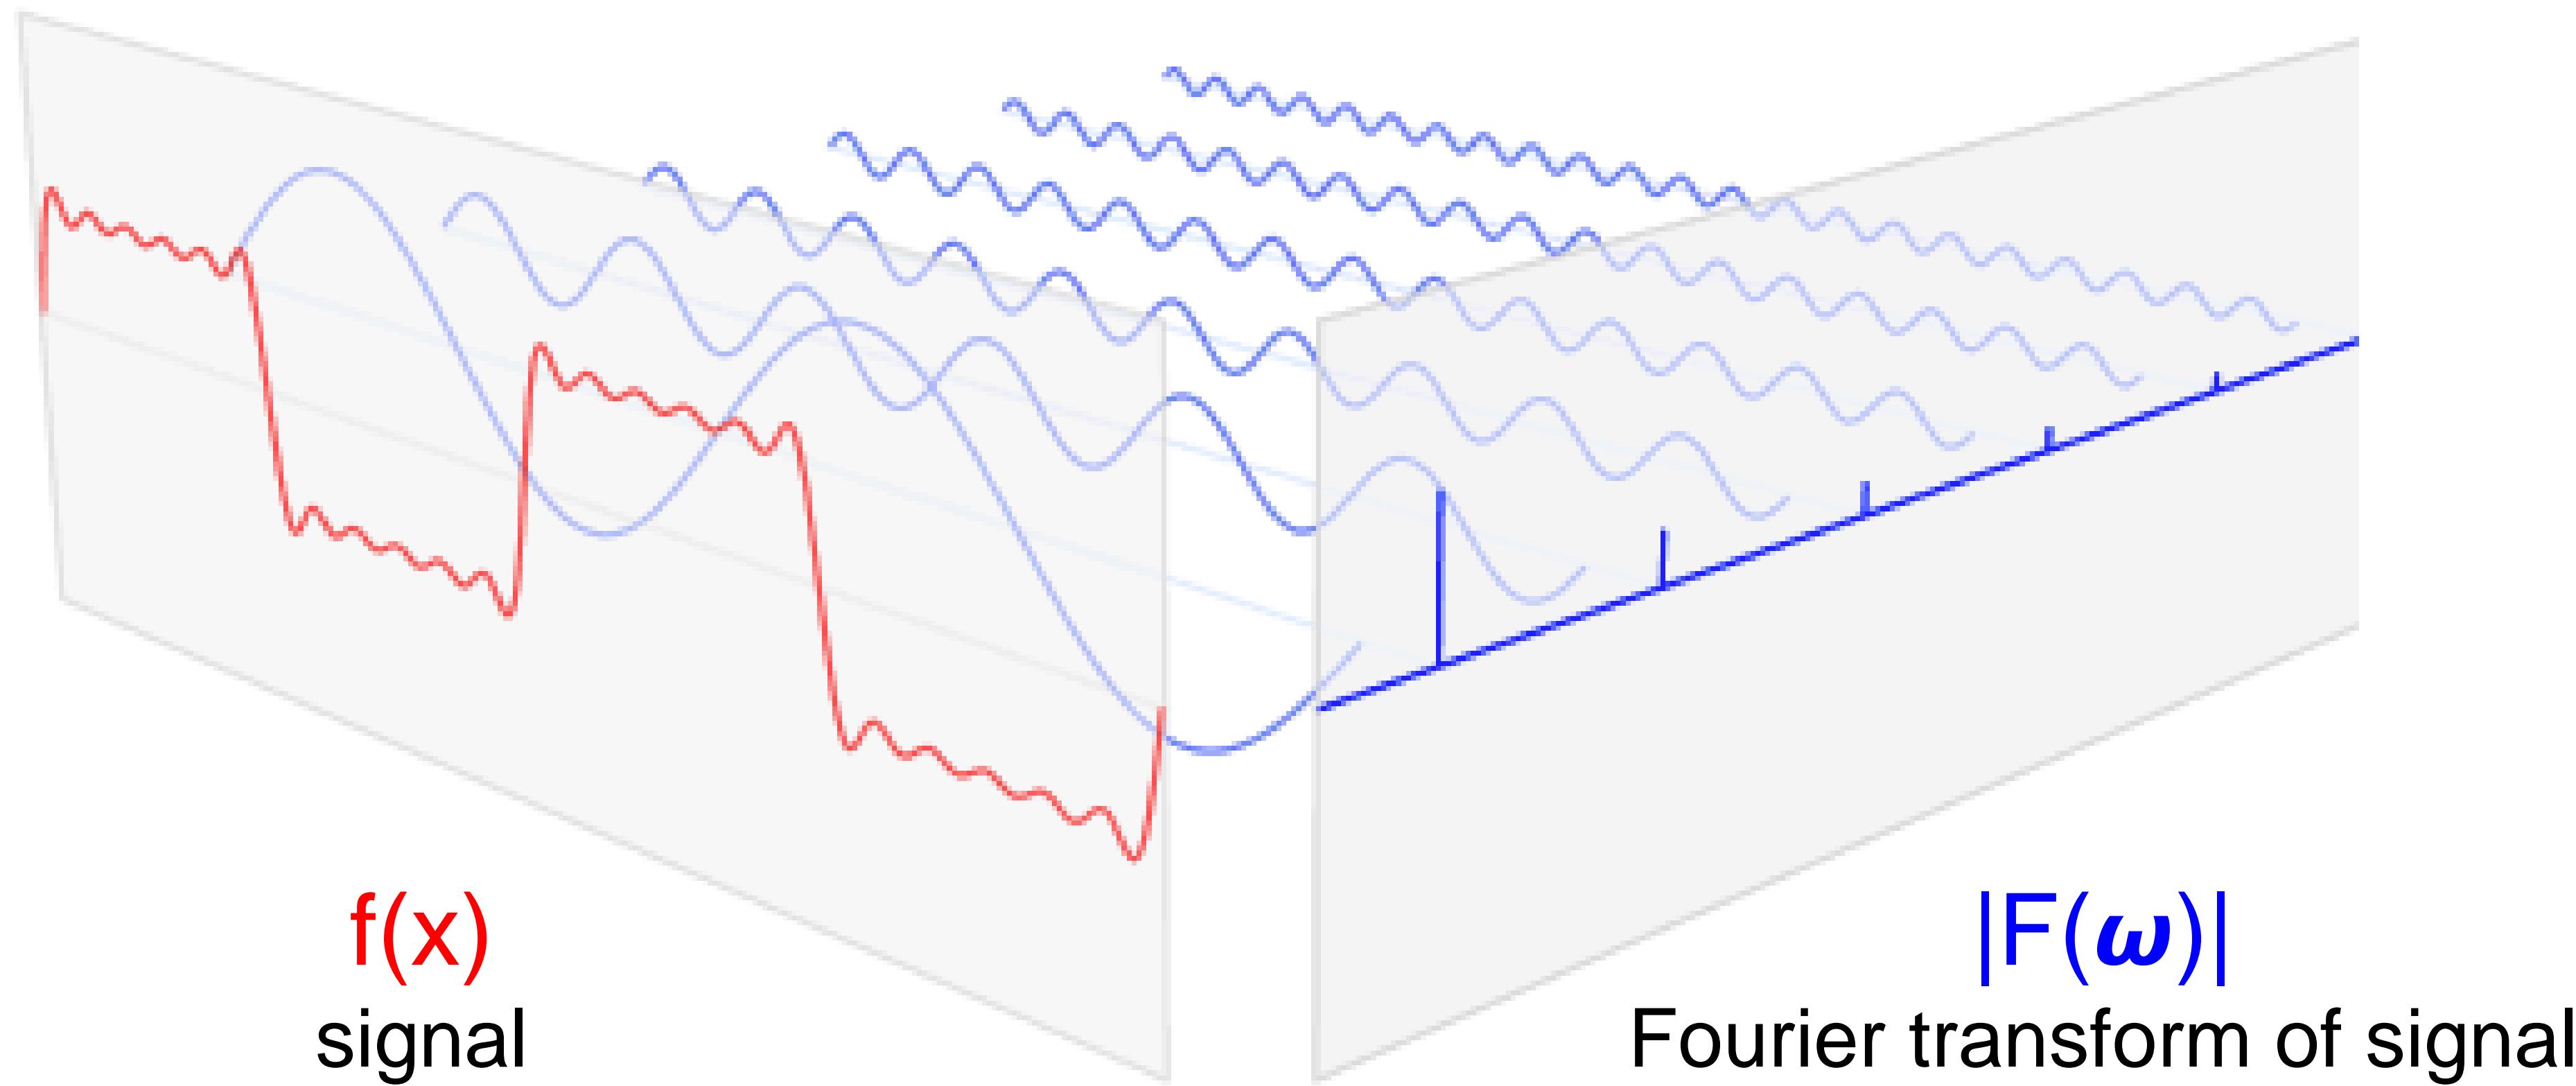
# Frequency Spectra

---



# Fourier Transform

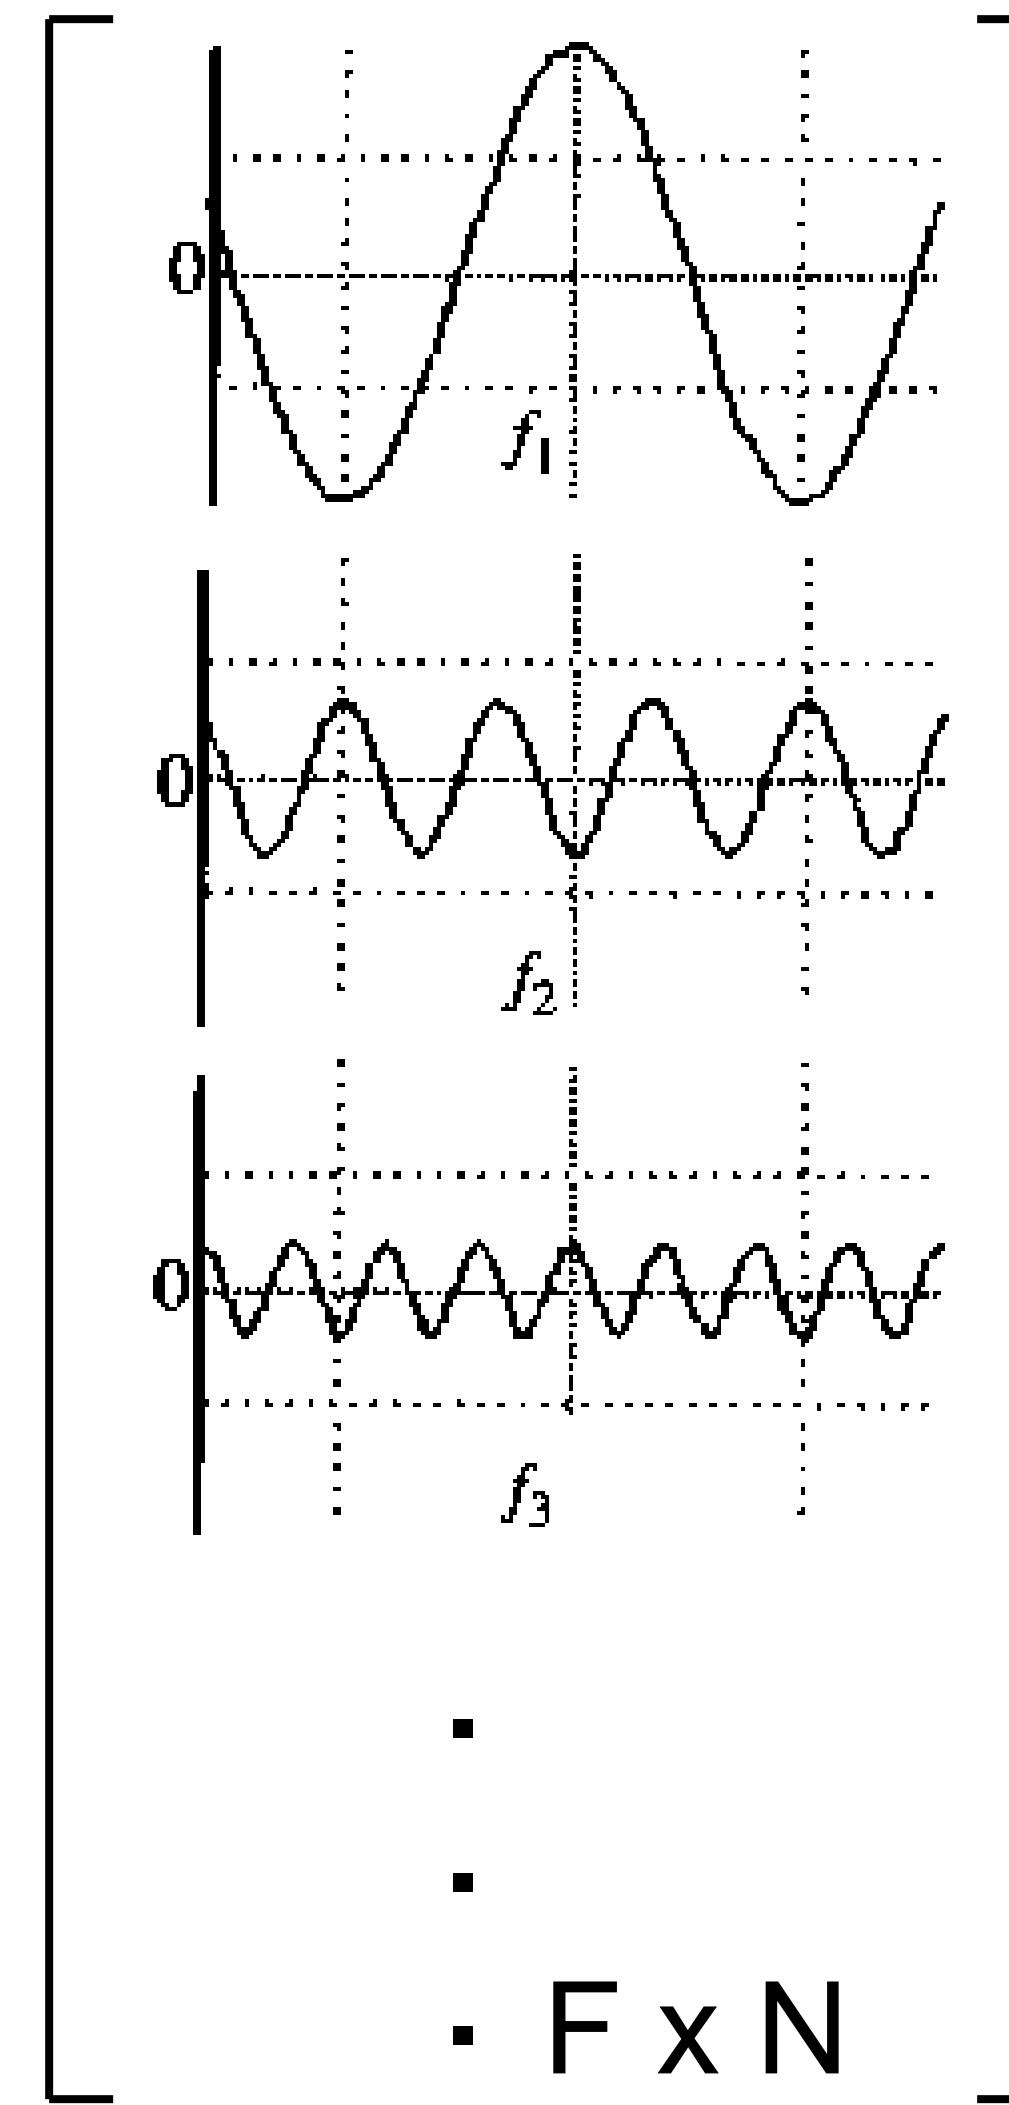
---



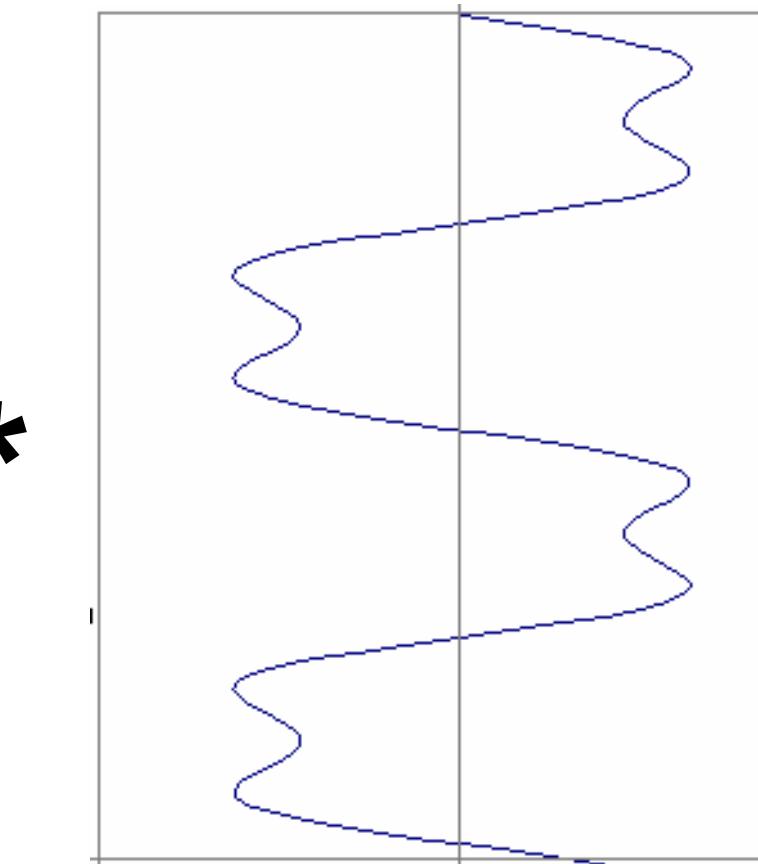
# FT: Just a change of basis

---

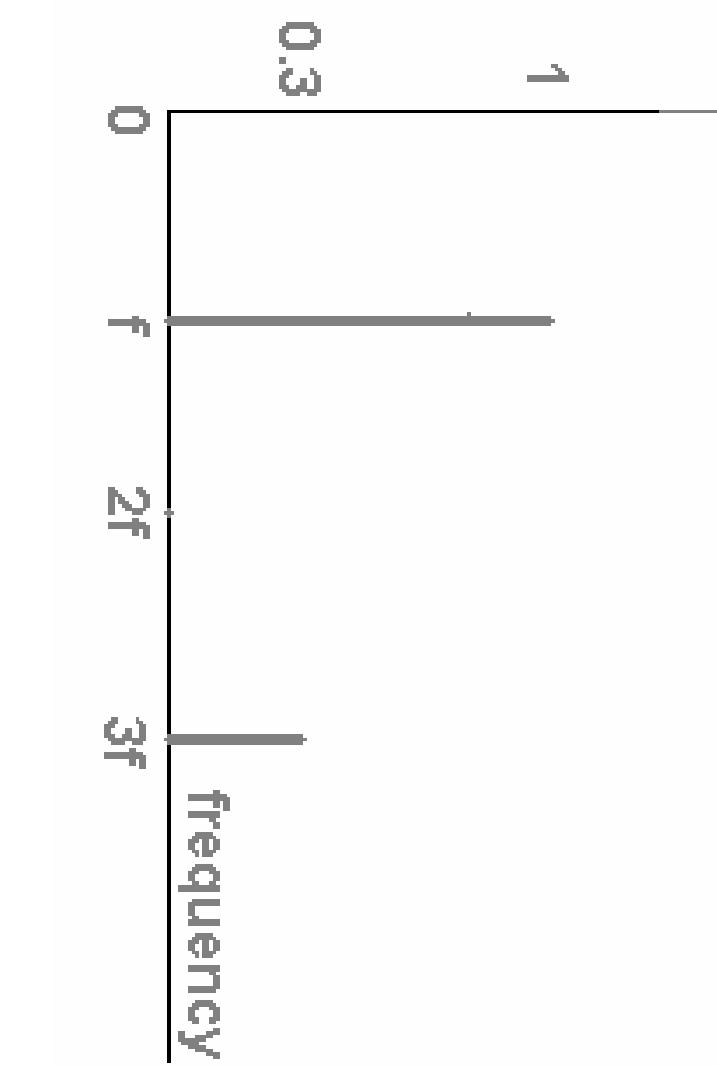
$$M * f(x) = F(\omega)$$



\*



=



$N \times 1$

$F \times 1$

# Self-inverting transforms

---

$$\vec{F} = U \vec{f} \longleftrightarrow \vec{f} = U^{-1} \vec{F}$$

Same basis functions are used for the inverse transform

$$\vec{f} = U^{-1} \vec{F}$$

$$= U^+ \vec{F}$$

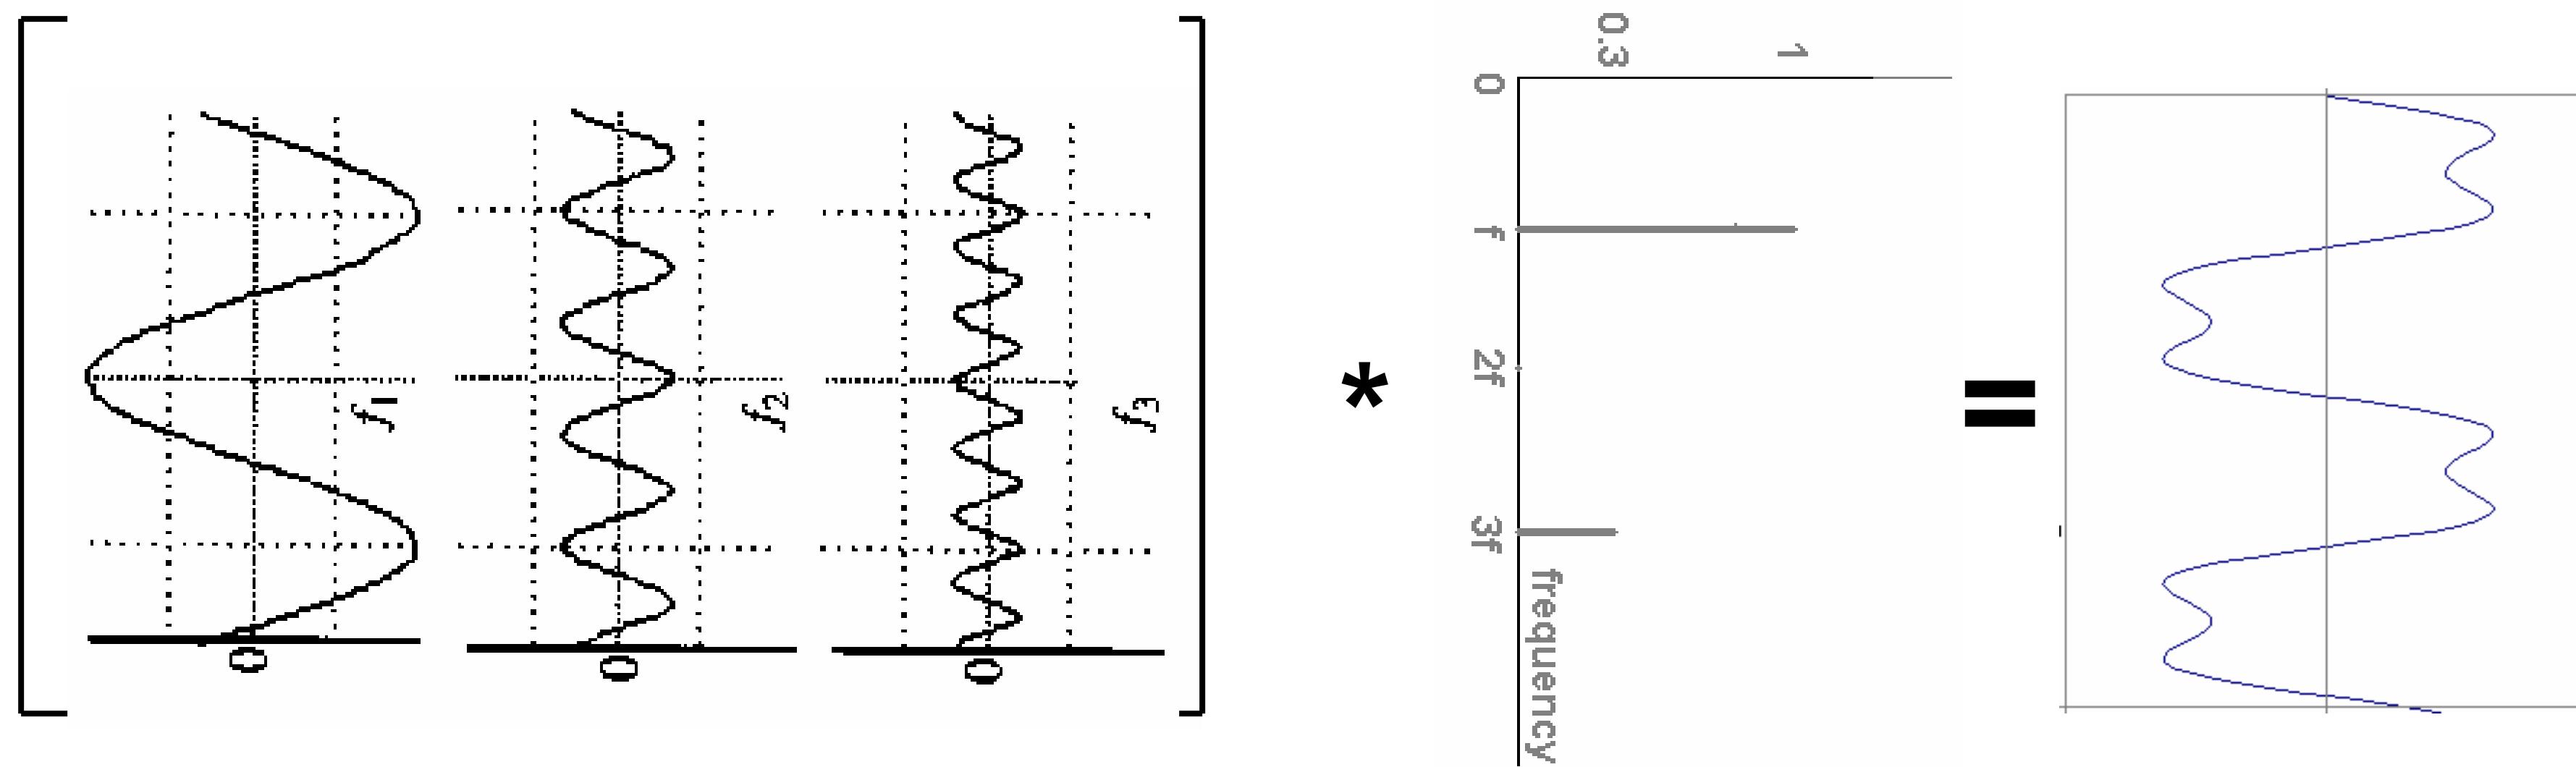


U transpose and complex conjugate

# IFT: Just a change of basis

---

$$M^{-1} * F(\omega) = f(x)$$



- $N \times F$

$F \times 1$

$N \times 1$

# Finally: Scary Math

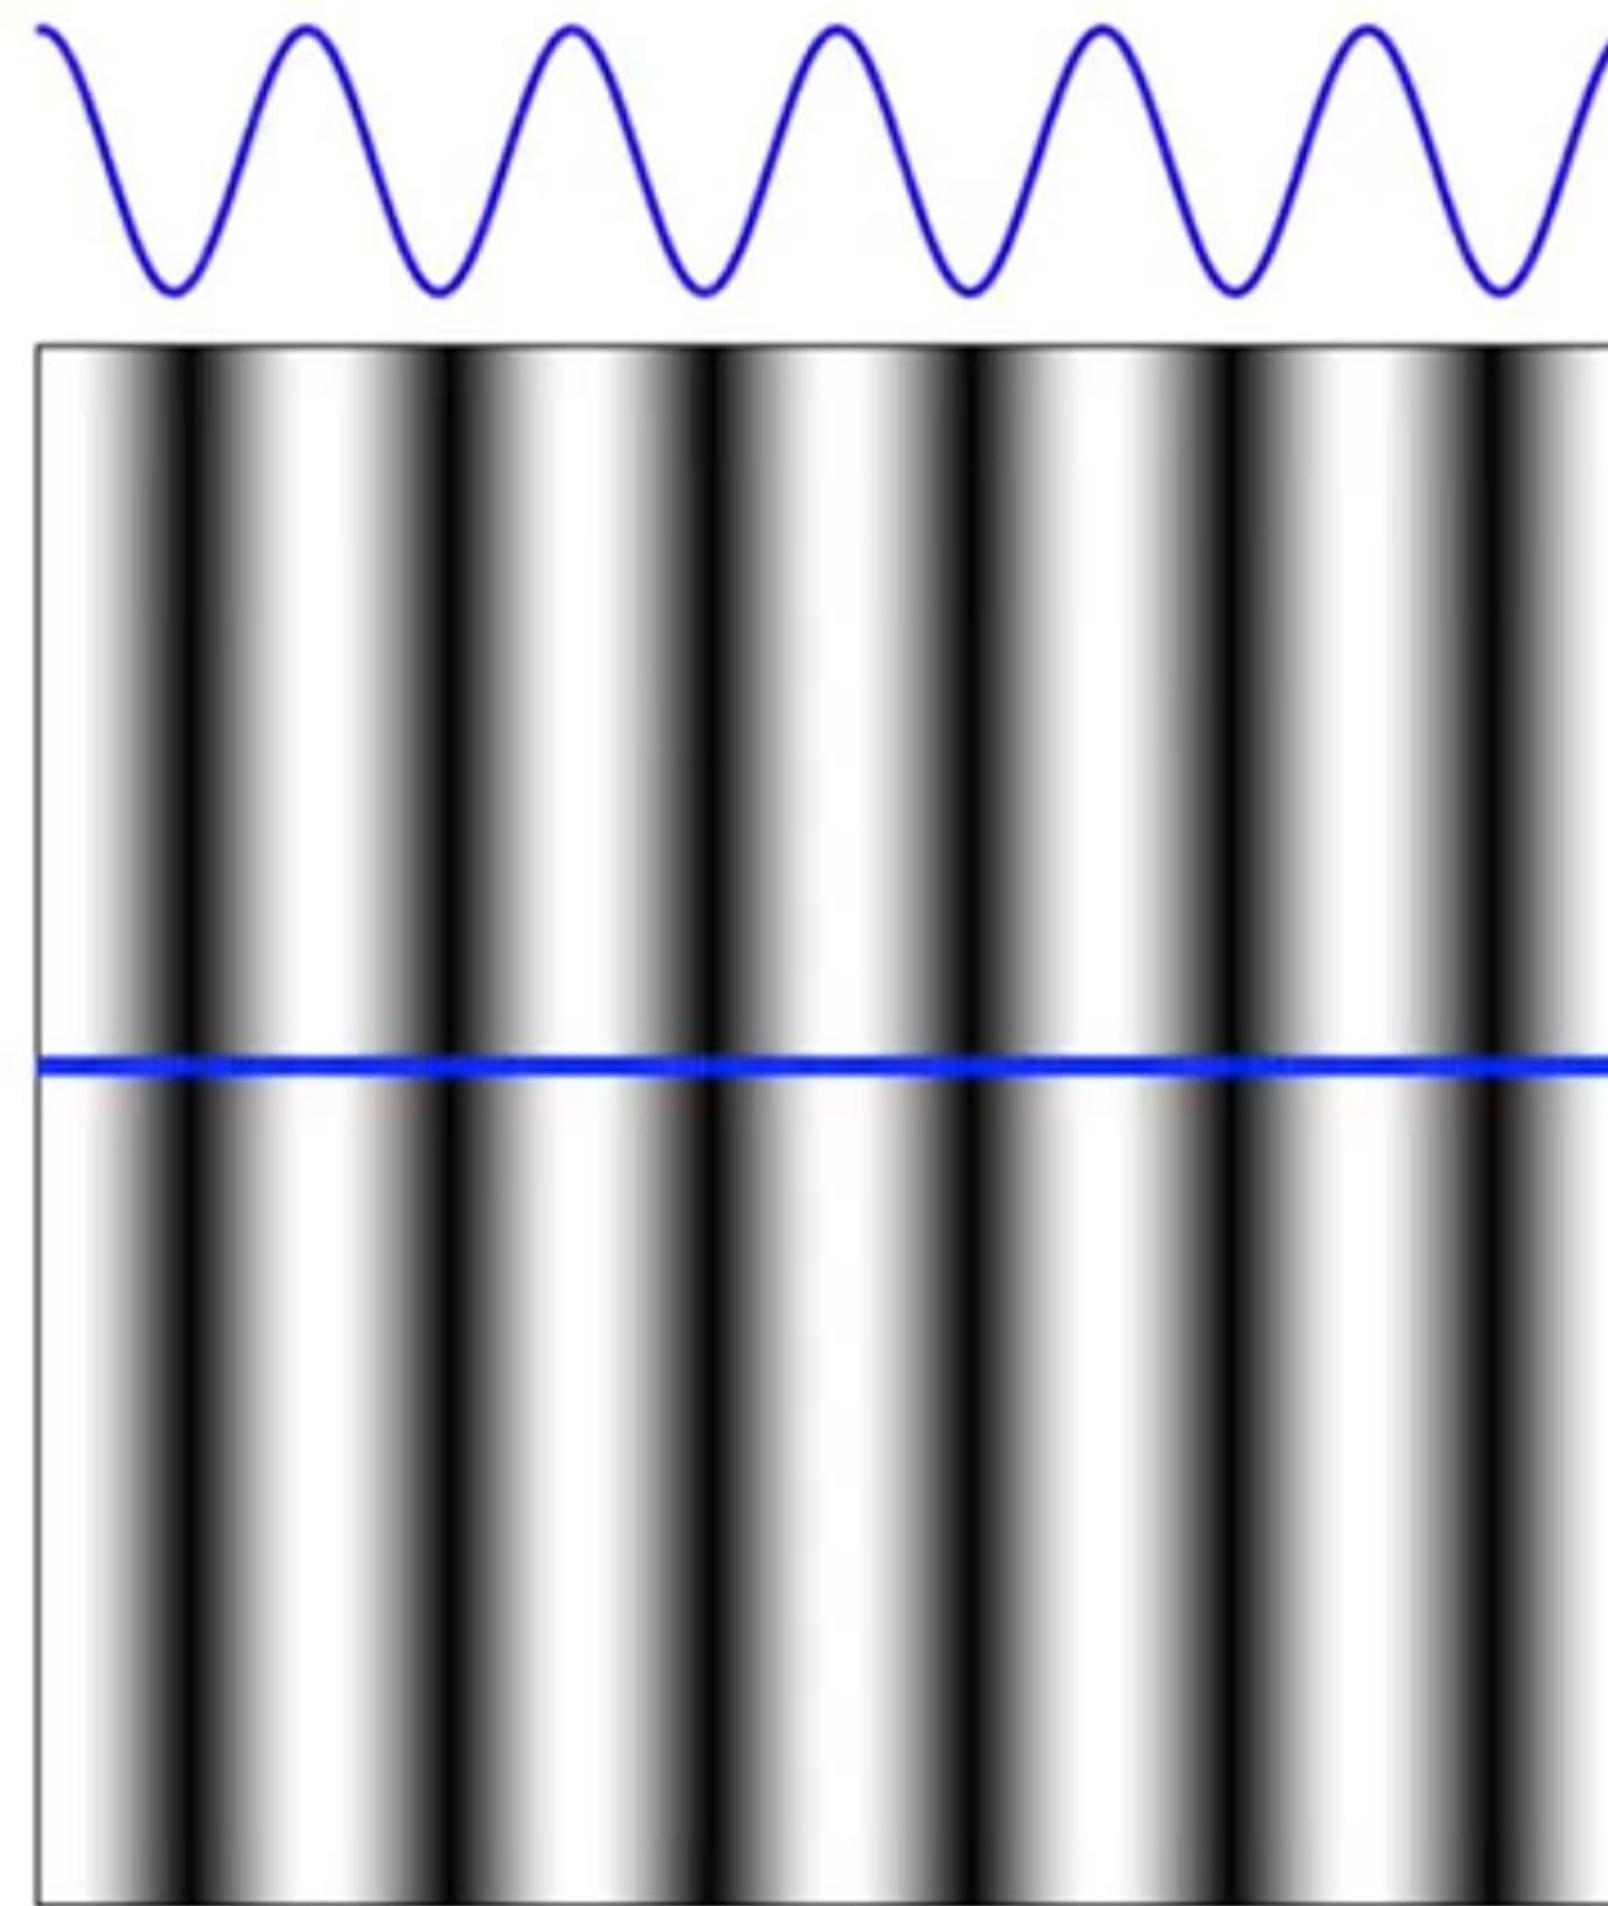
---

$$\text{Fourier Transform : } F(\omega) = \int_{-\infty}^{+\infty} f(x)e^{-i\omega x} dx$$

$$\text{Inverse Fourier Transform : } f(x) = \frac{1}{2\pi} \int_{-\infty}^{+\infty} F(\omega)e^{i\omega x} d\omega$$

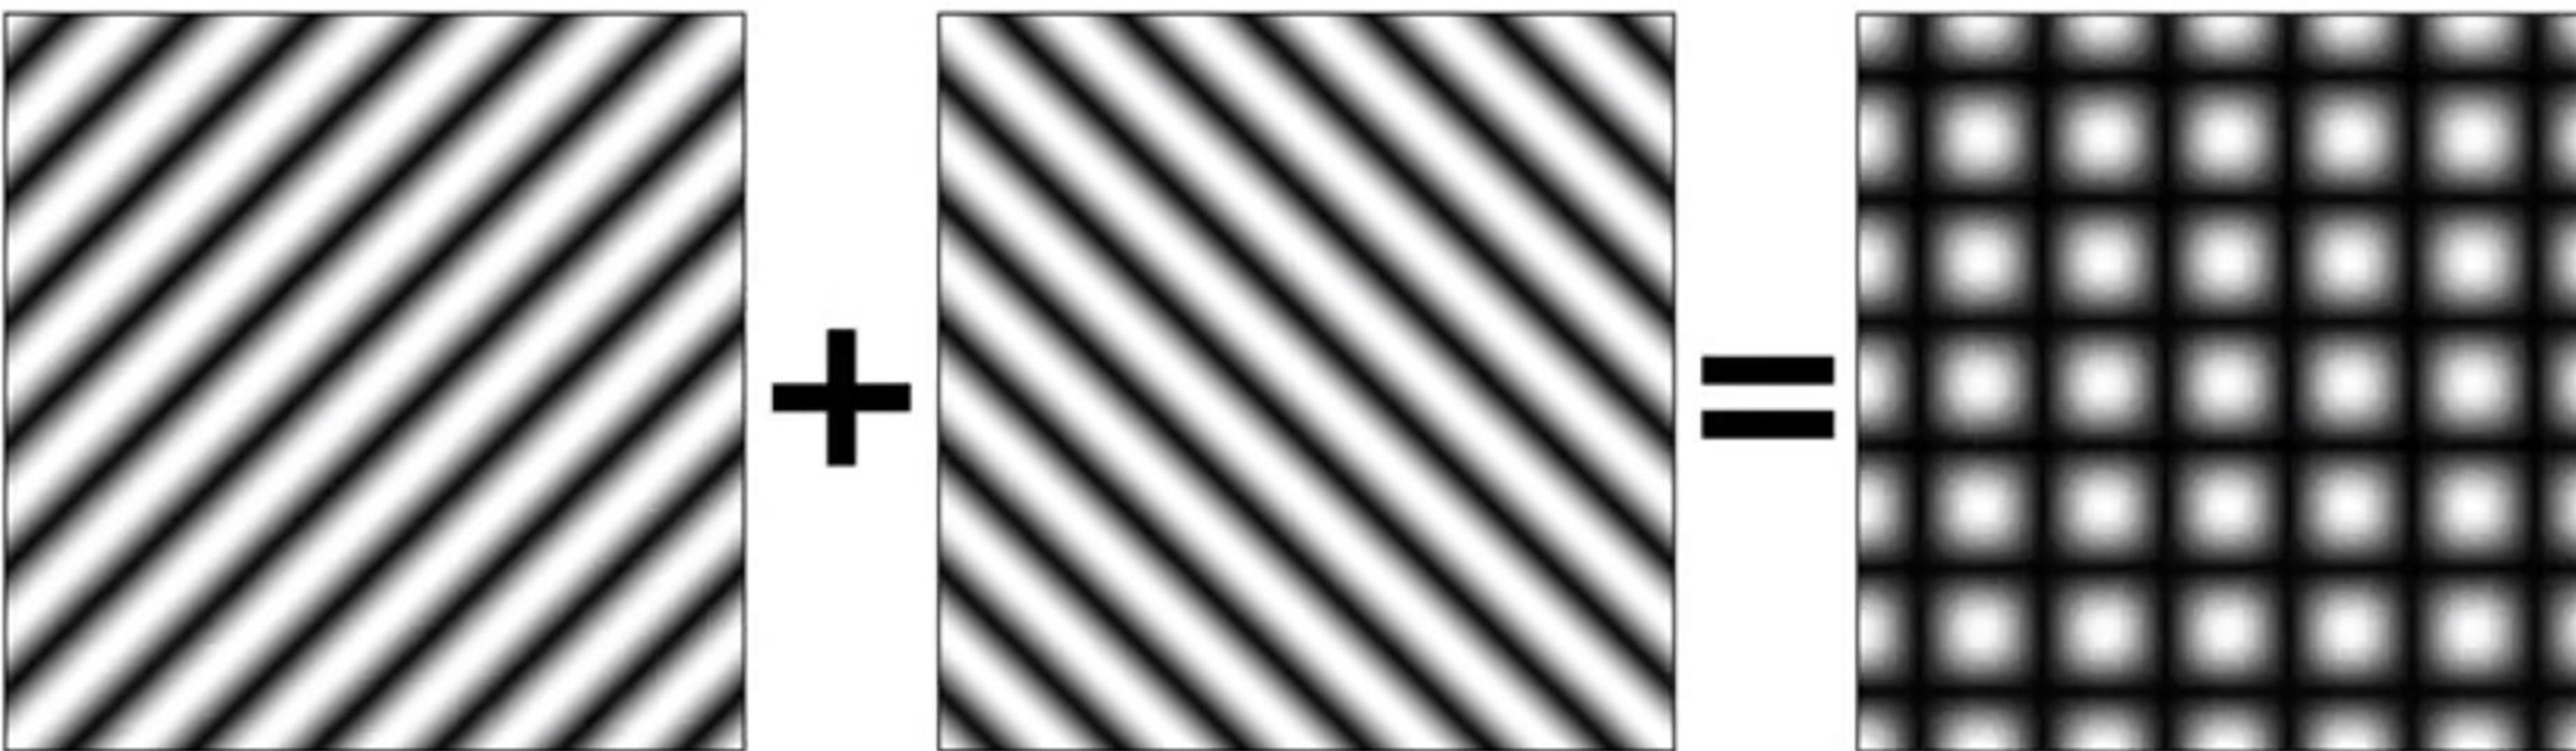
# Extending to 2D

---



# Addition still works in 2D

---



# Extension to 2D

---

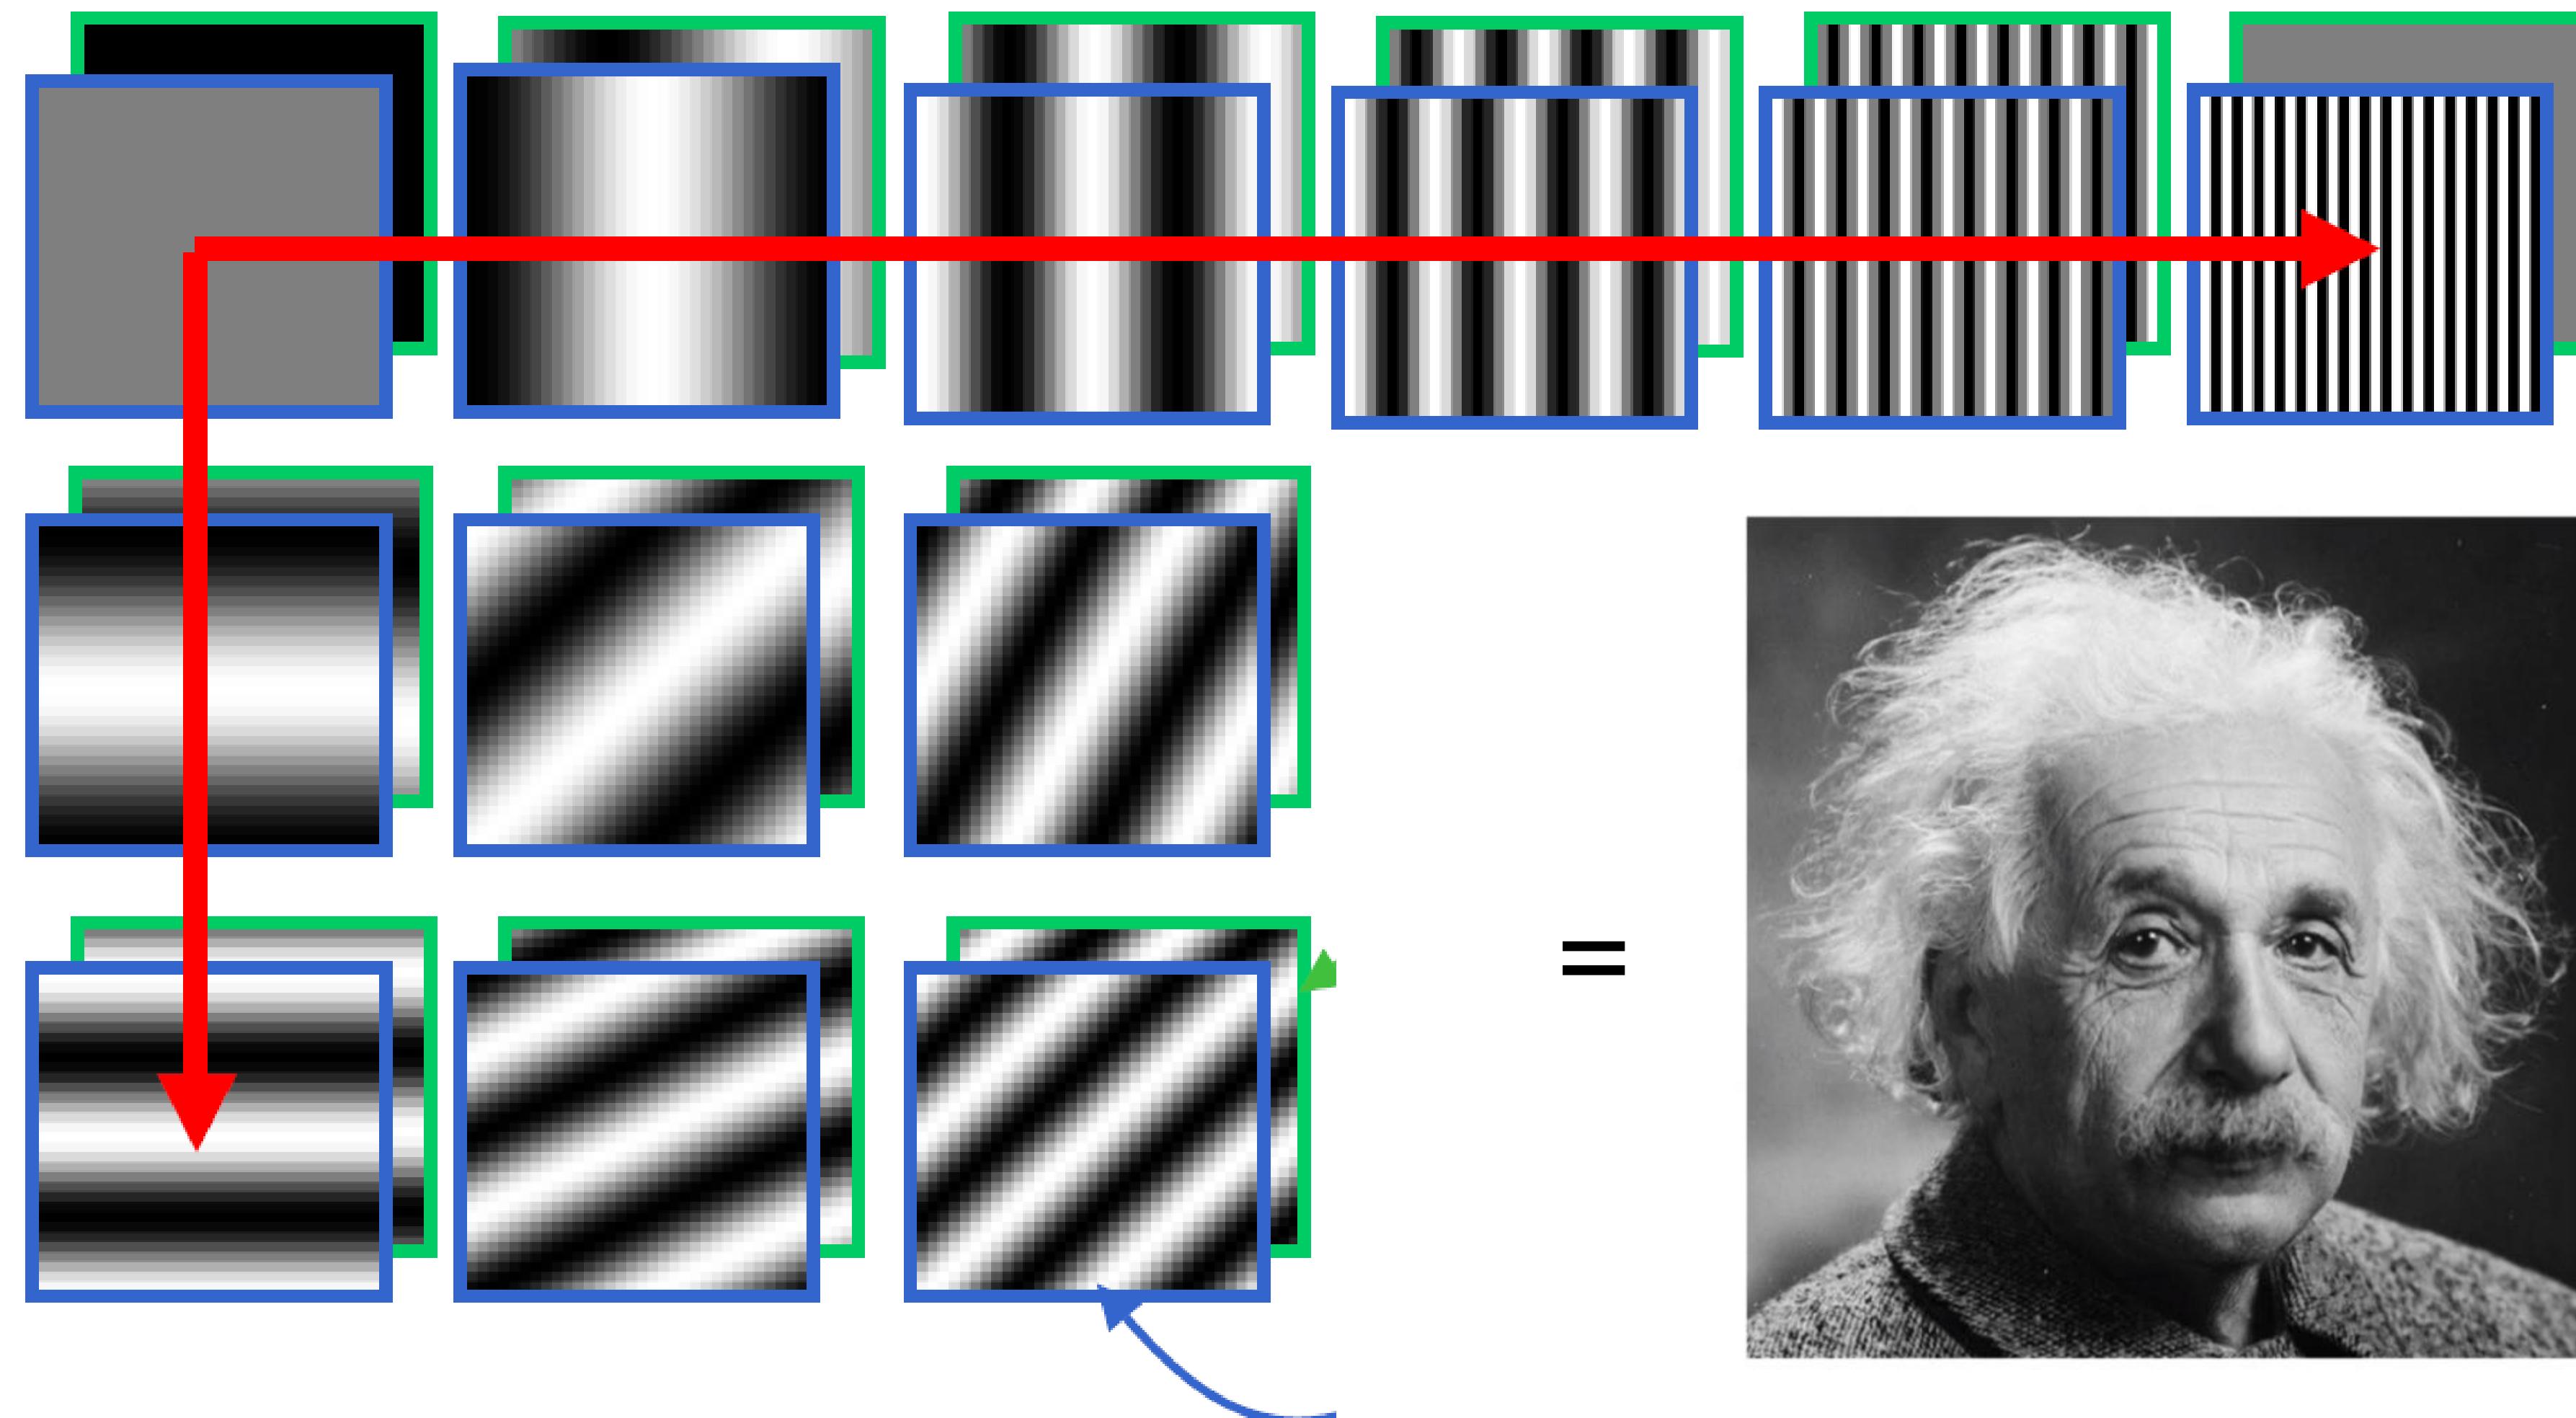
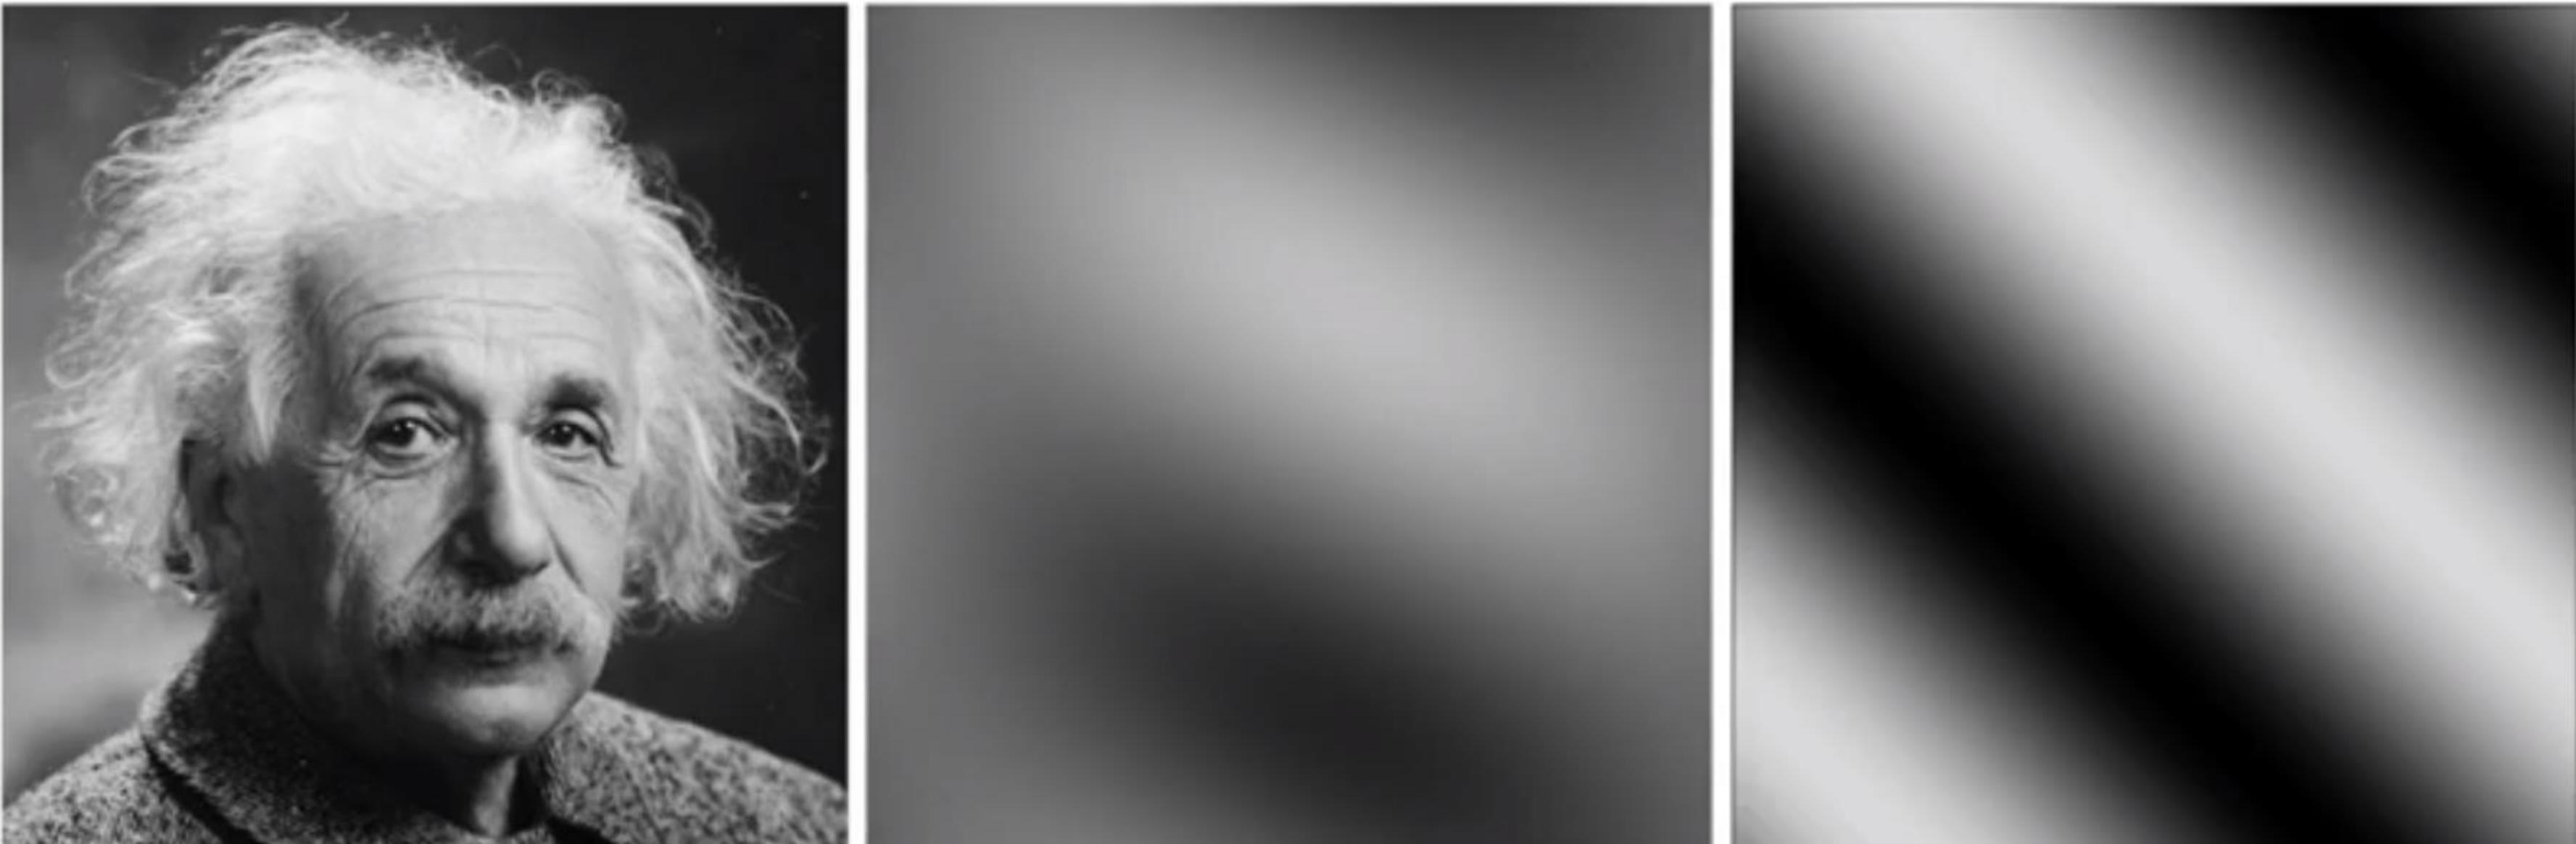


Image as a sum of basis images



Contrast x3



2



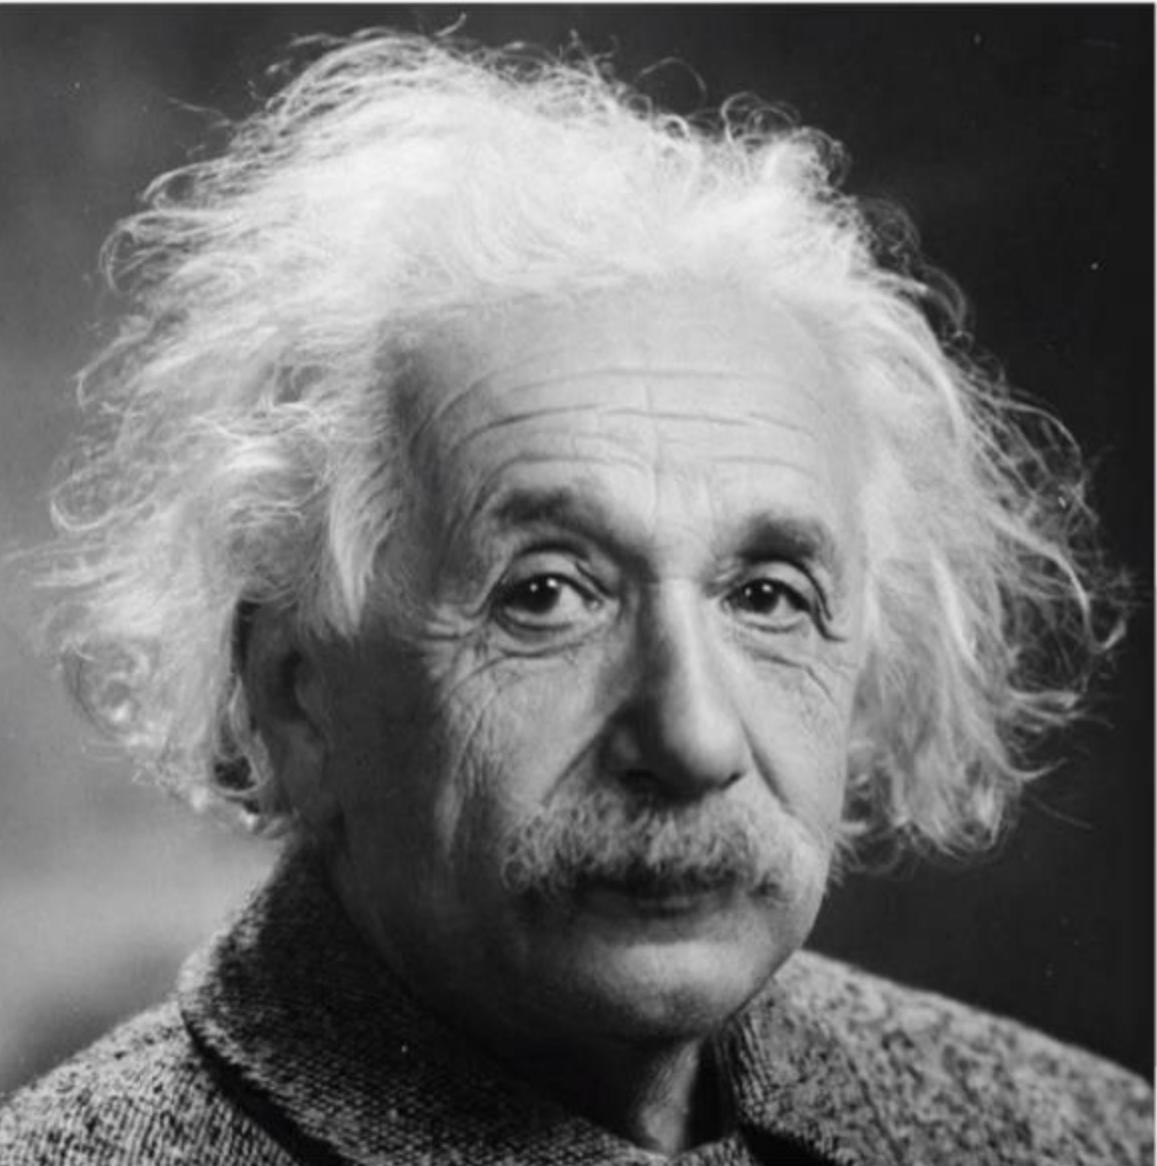
4



13



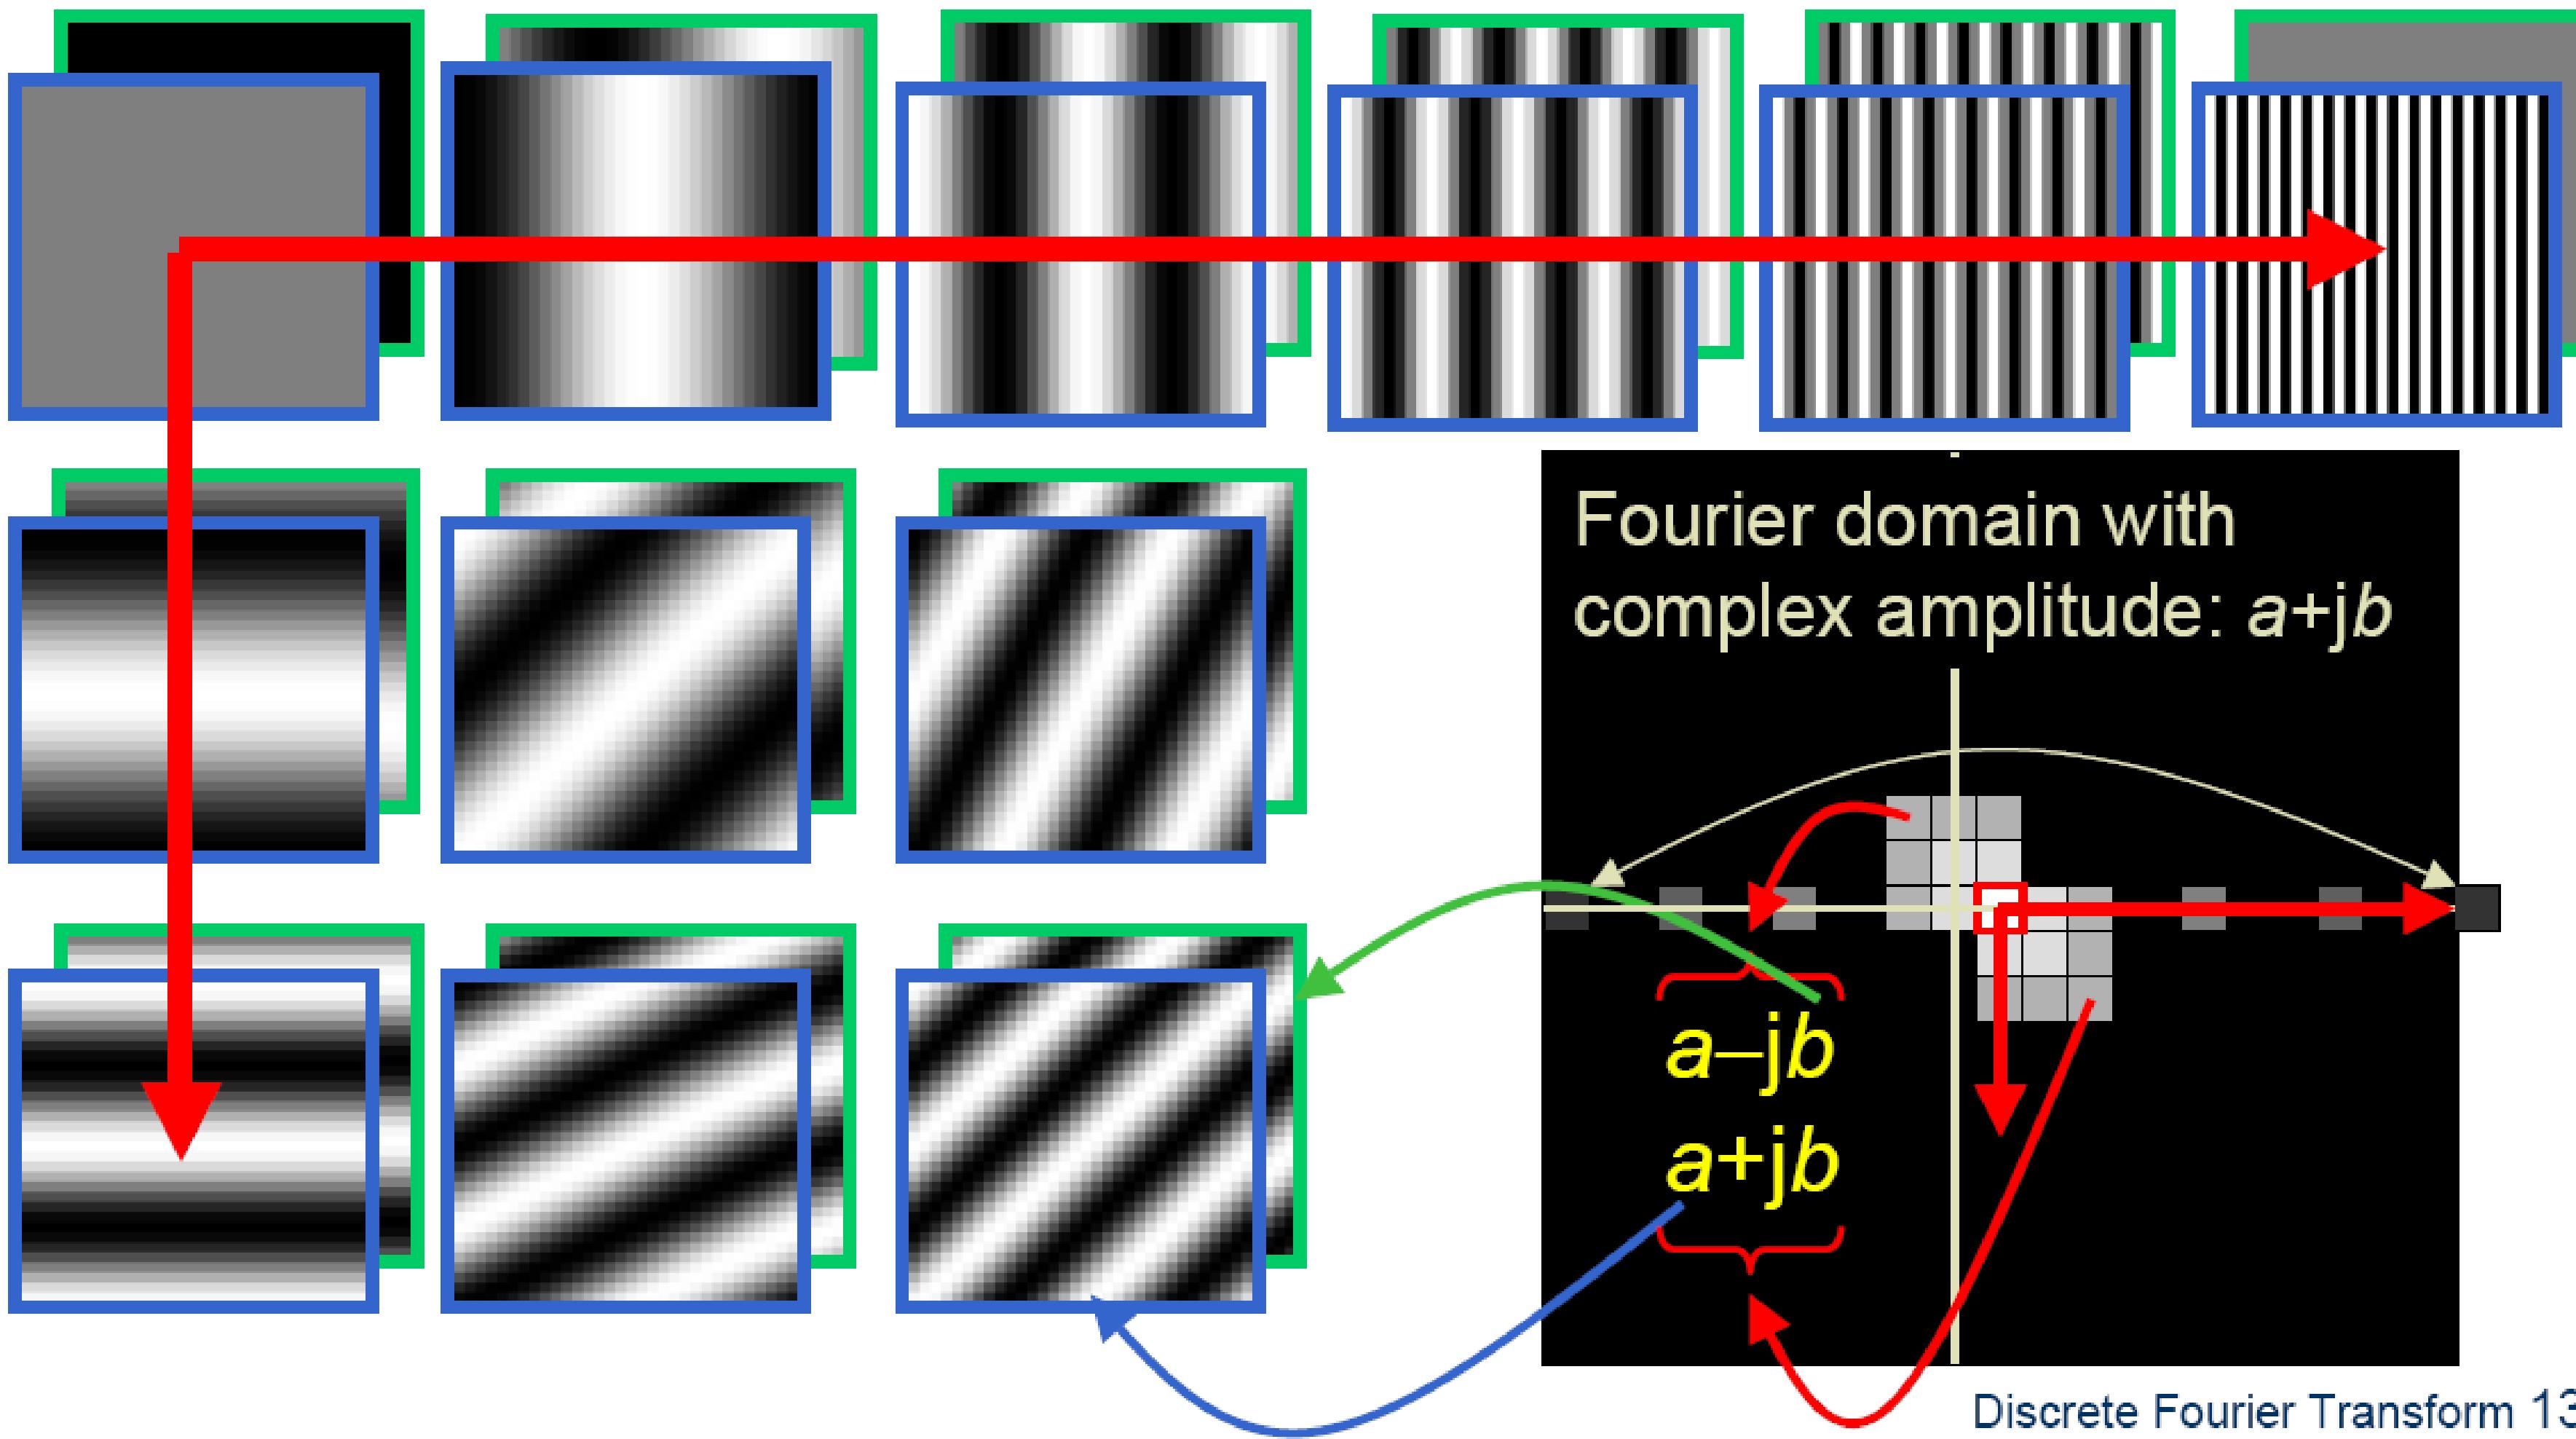
26



100

# Extension to 2D

---

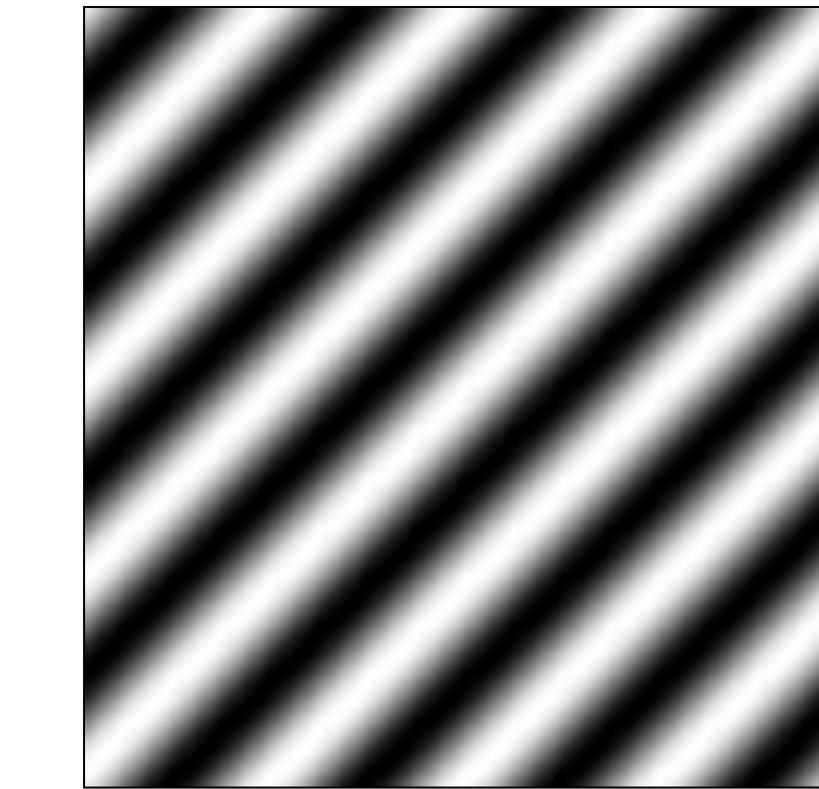
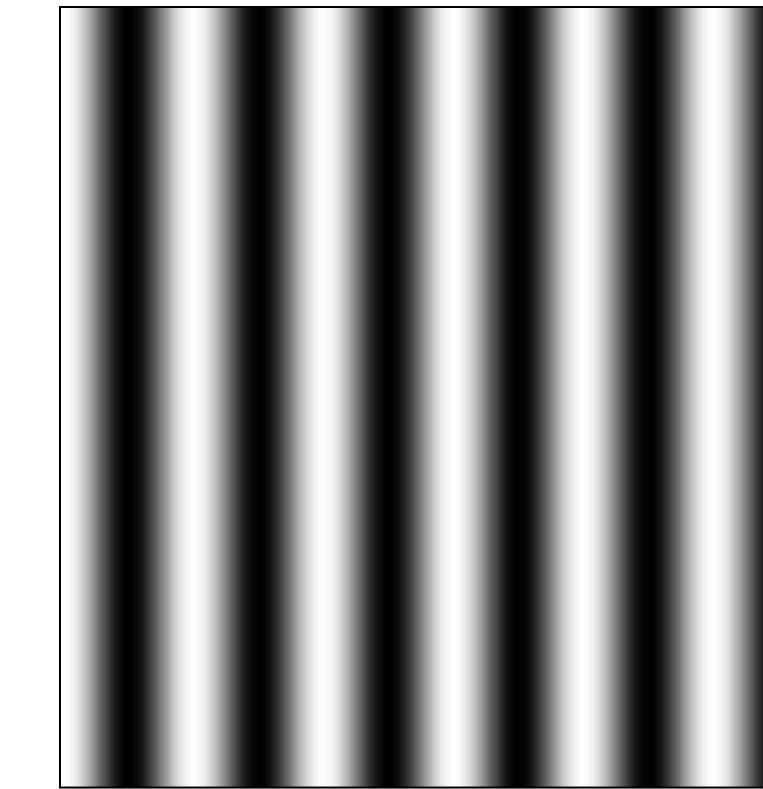
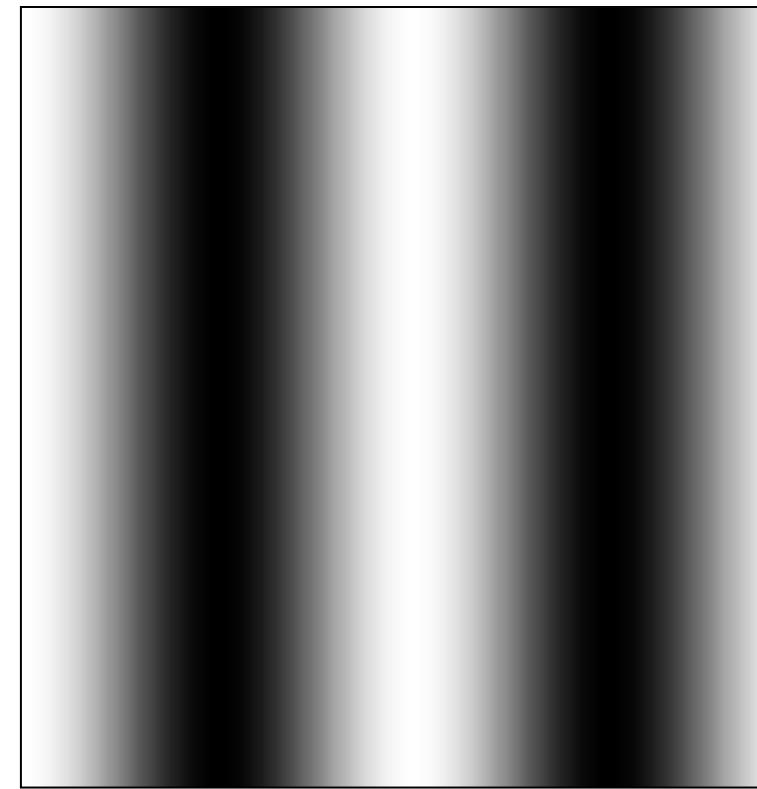


in Matlab, check out: `imagesc(log(abs(fftshift(fft2(im)))));`

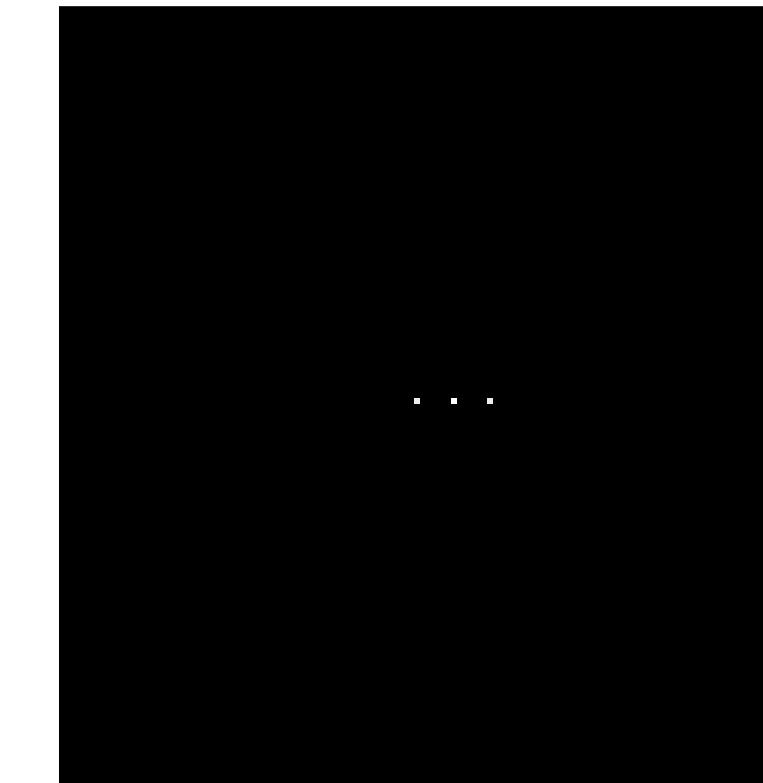
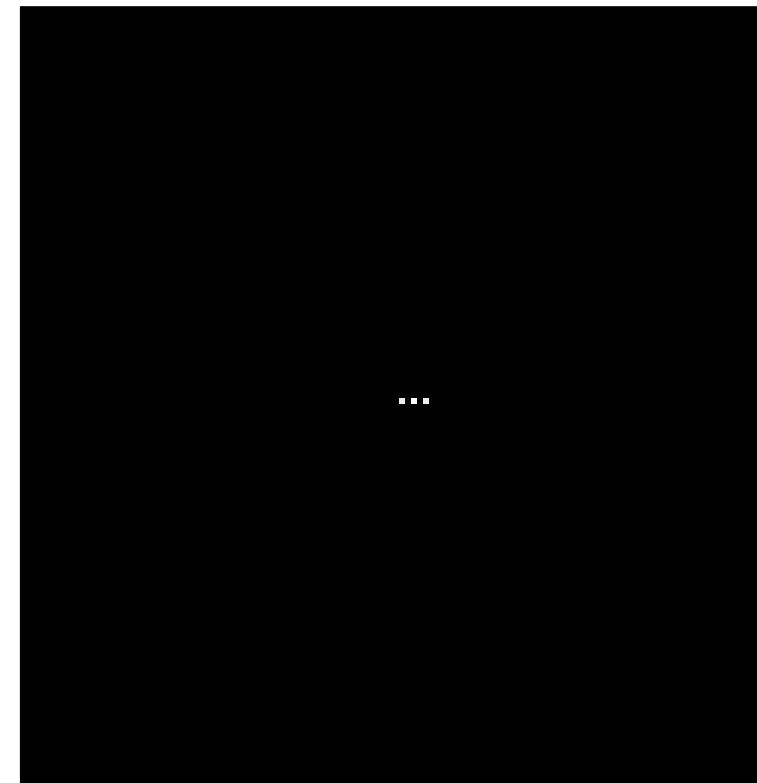
# Fourier analysis in images

---

Intensity Image

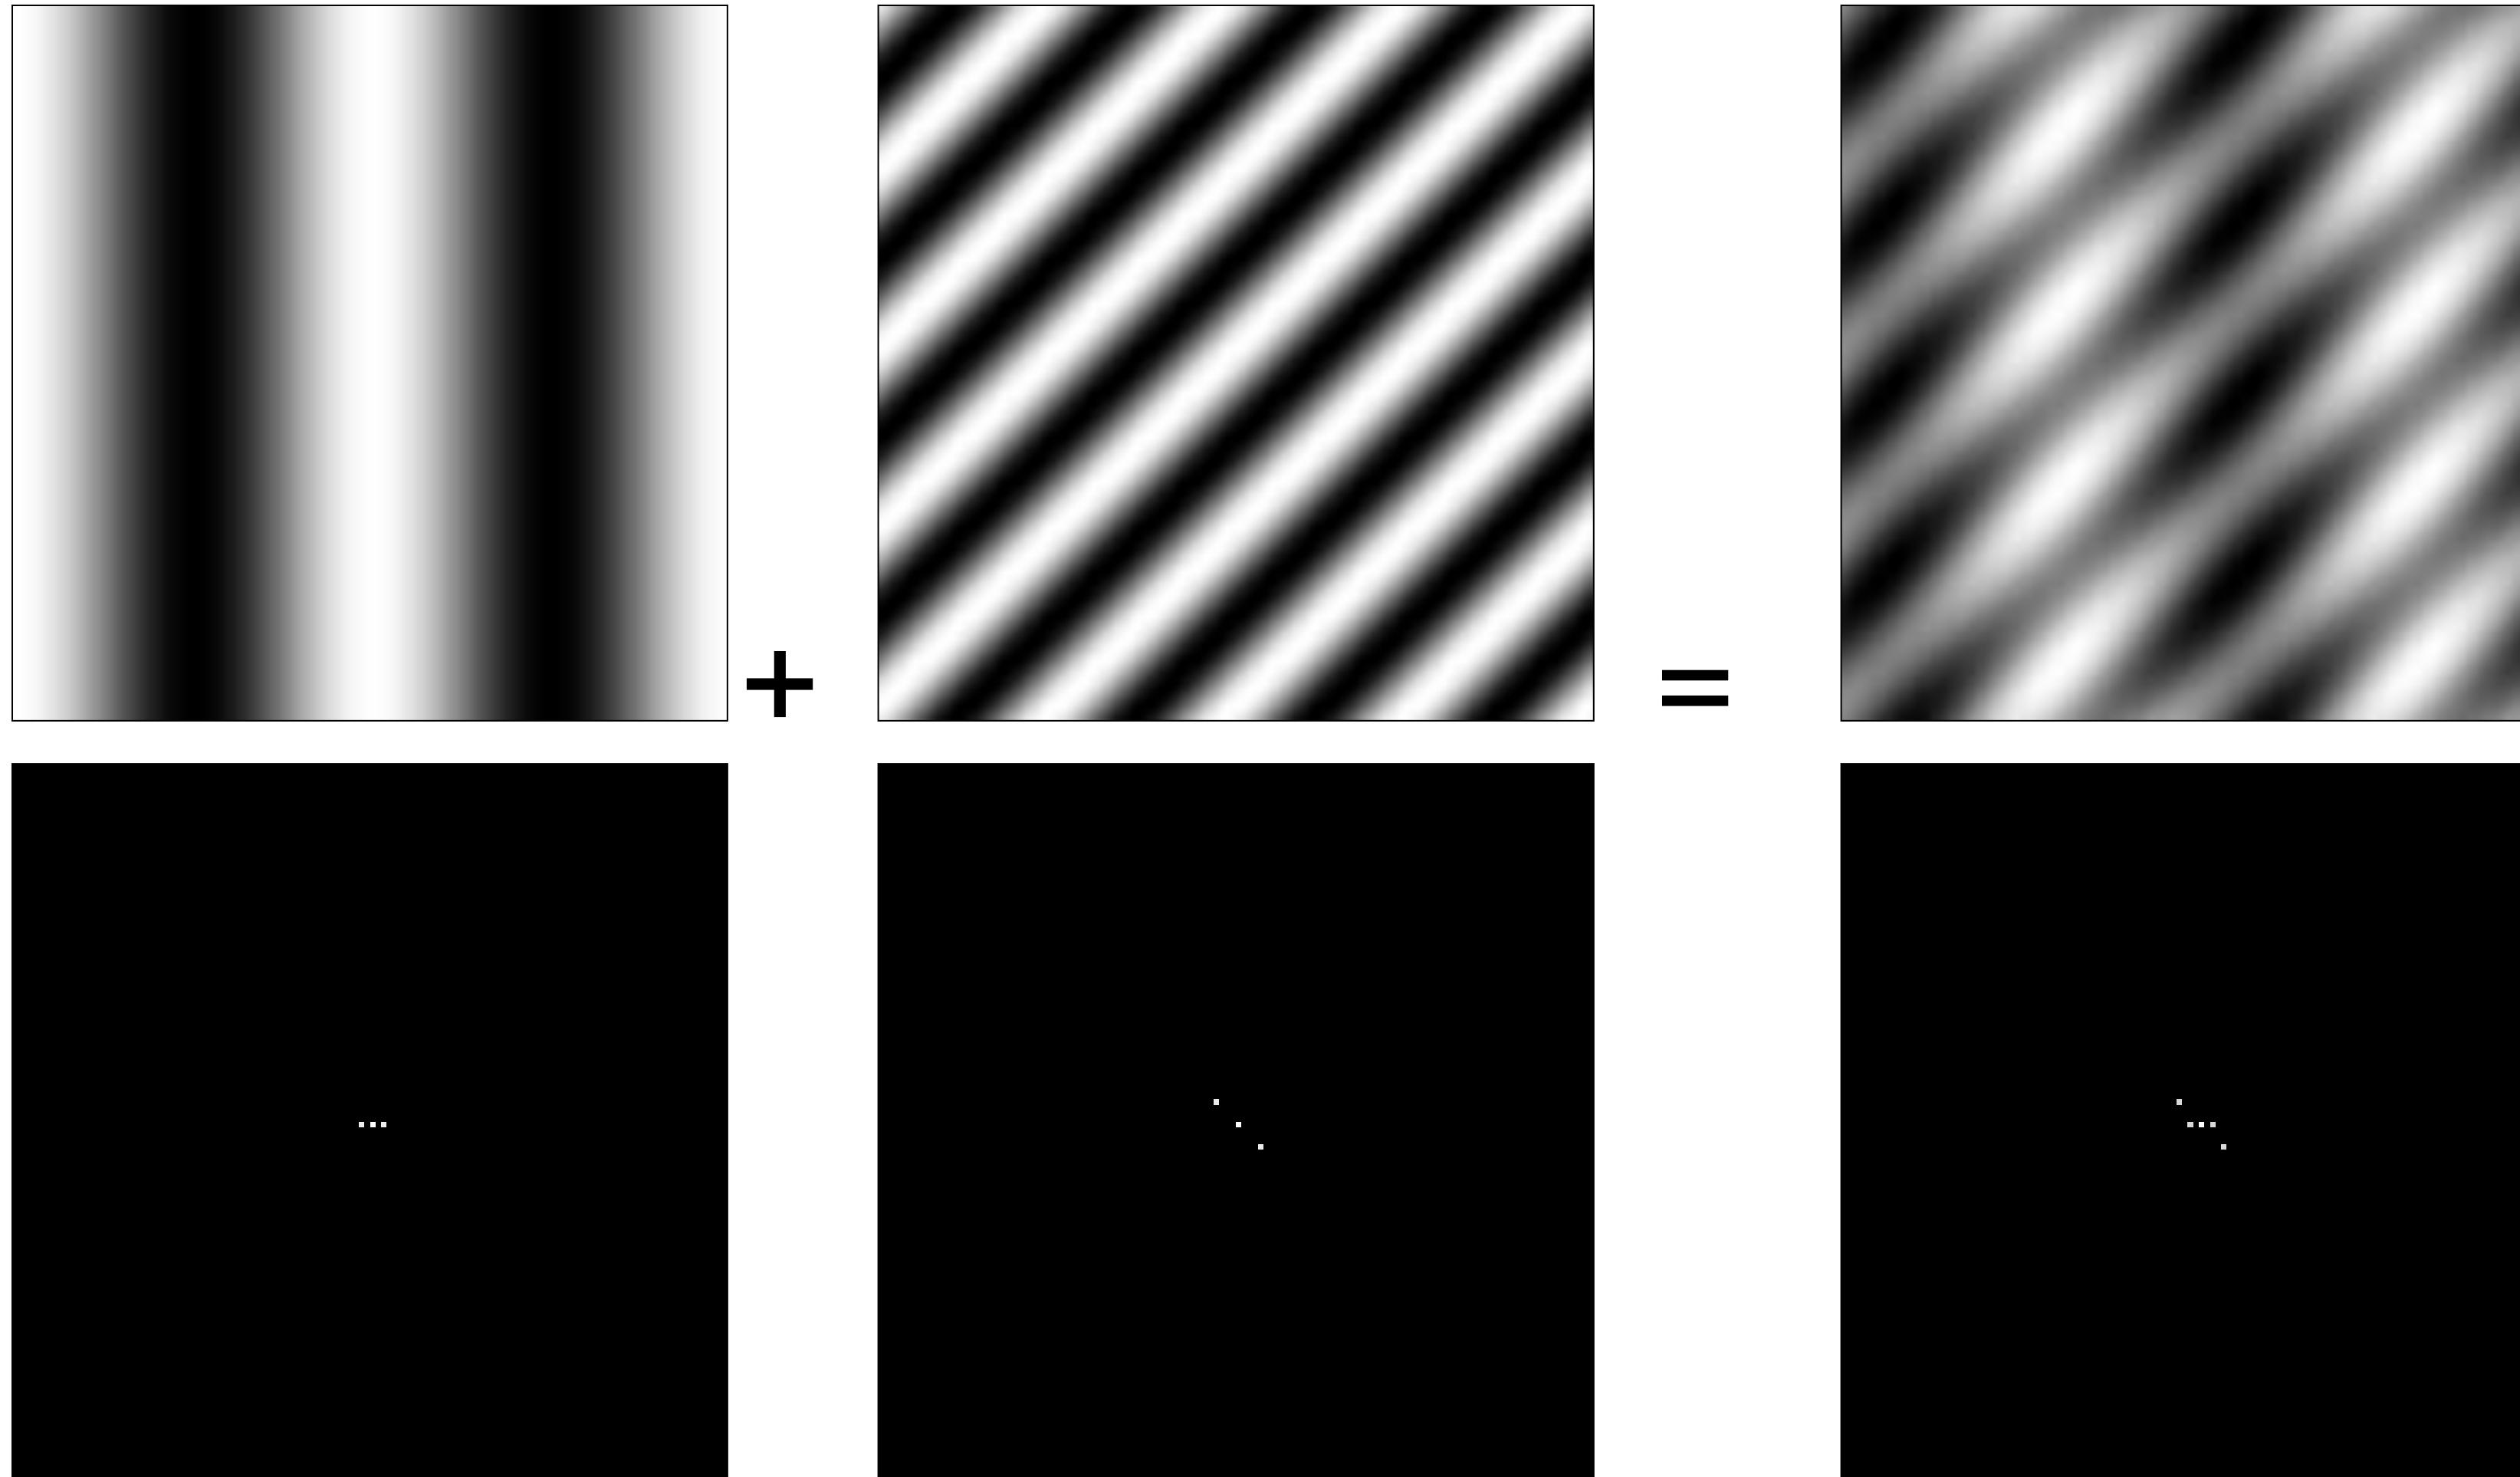


Fourier Image



# Signals can be composed

---

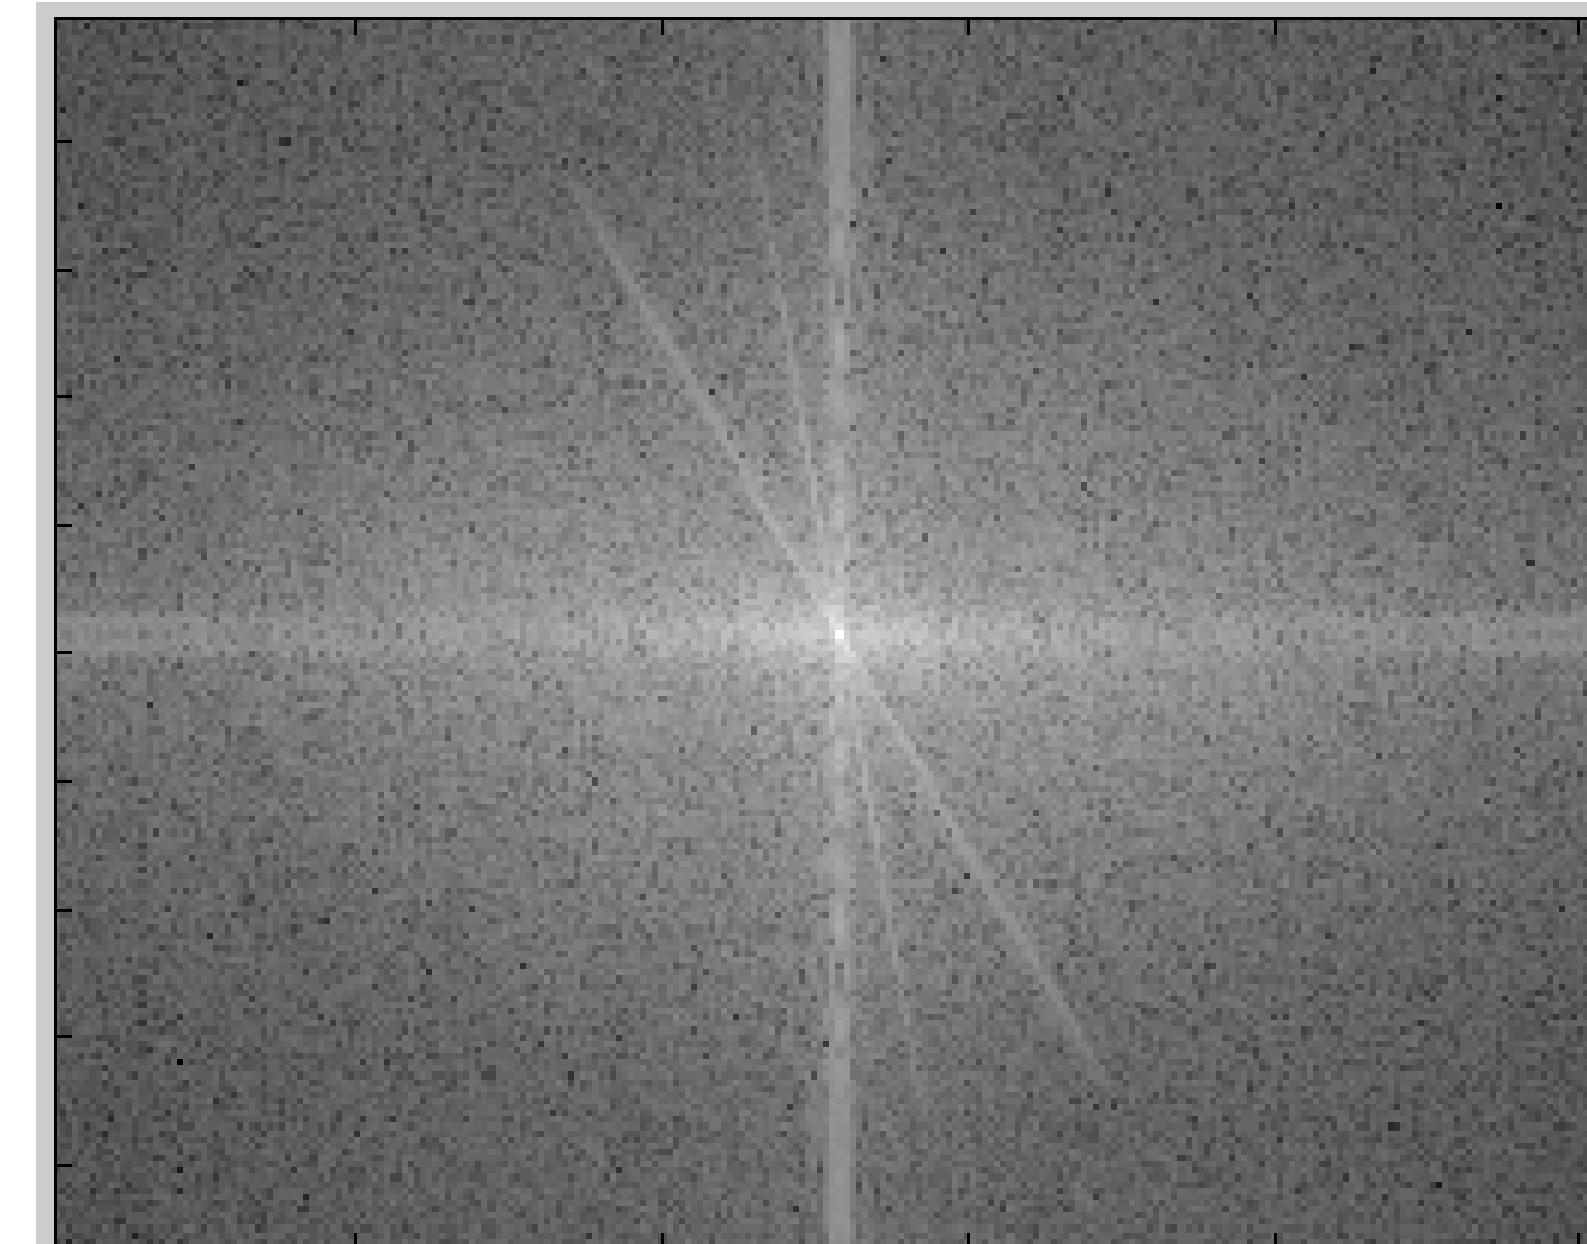


<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering>  
More: <http://www.cs.unm.edu/~brayer/vision/fourier.html>

# Man-made Scene

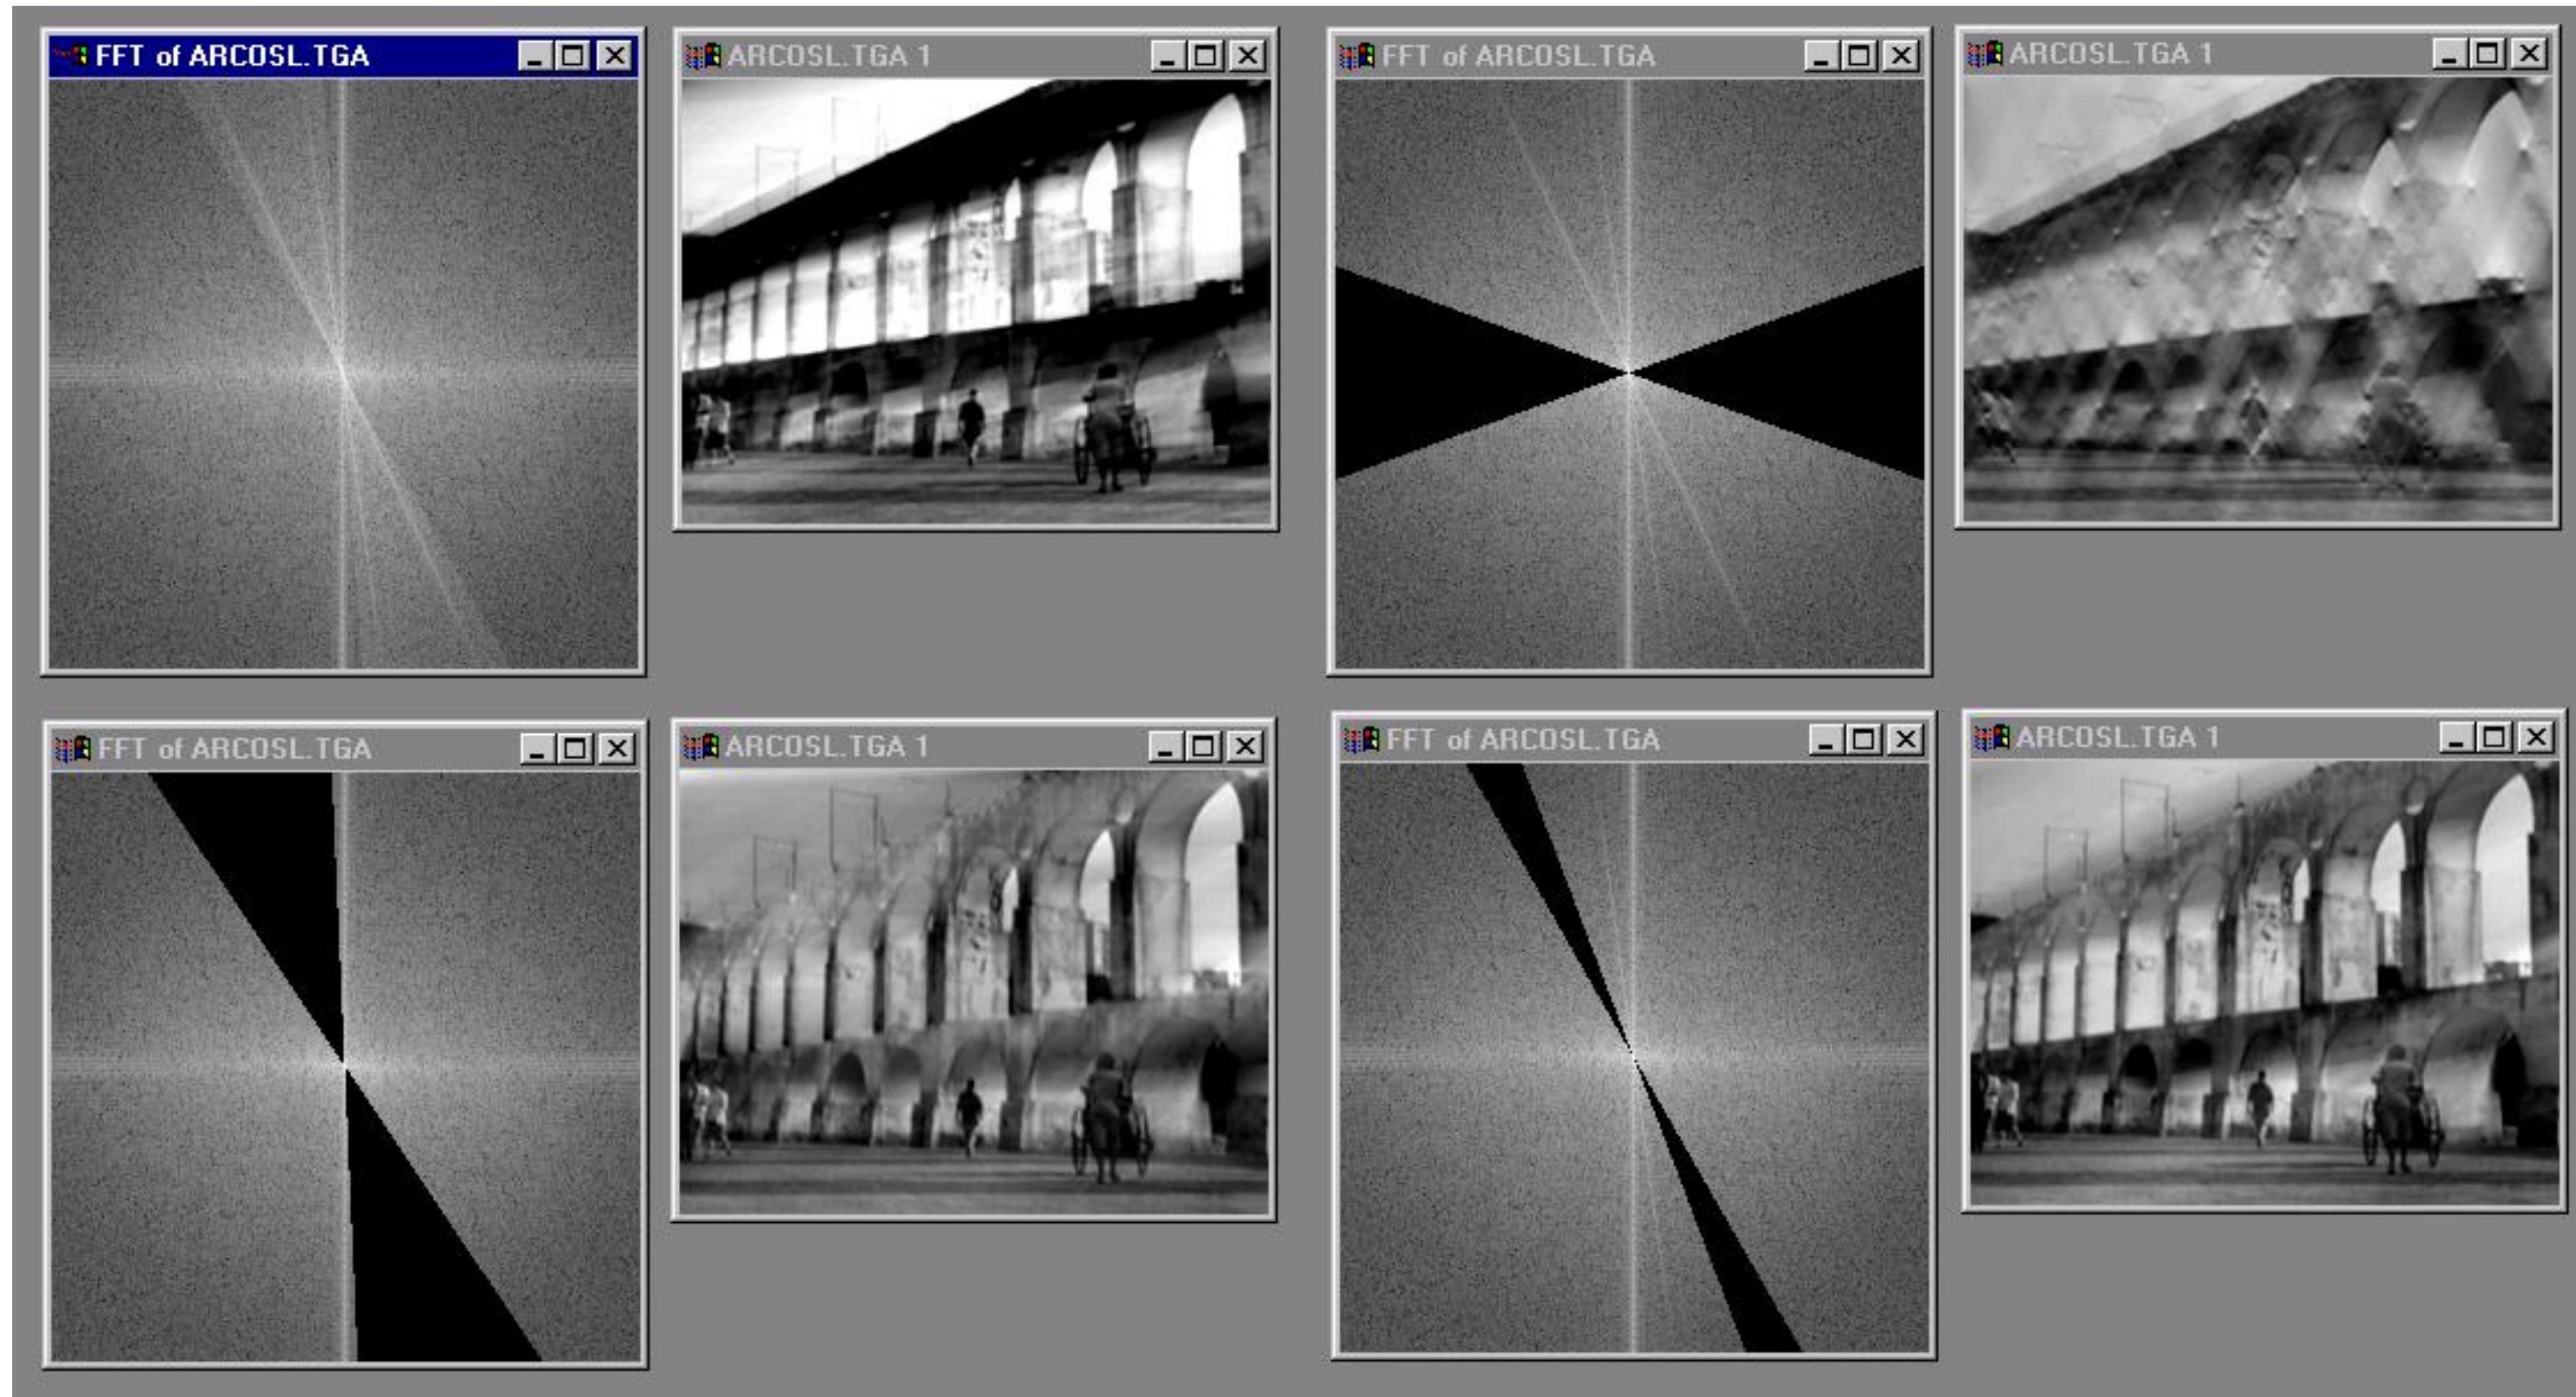
---

Amplitude Spectrum



# Can change spectrum, then reconstruct

---

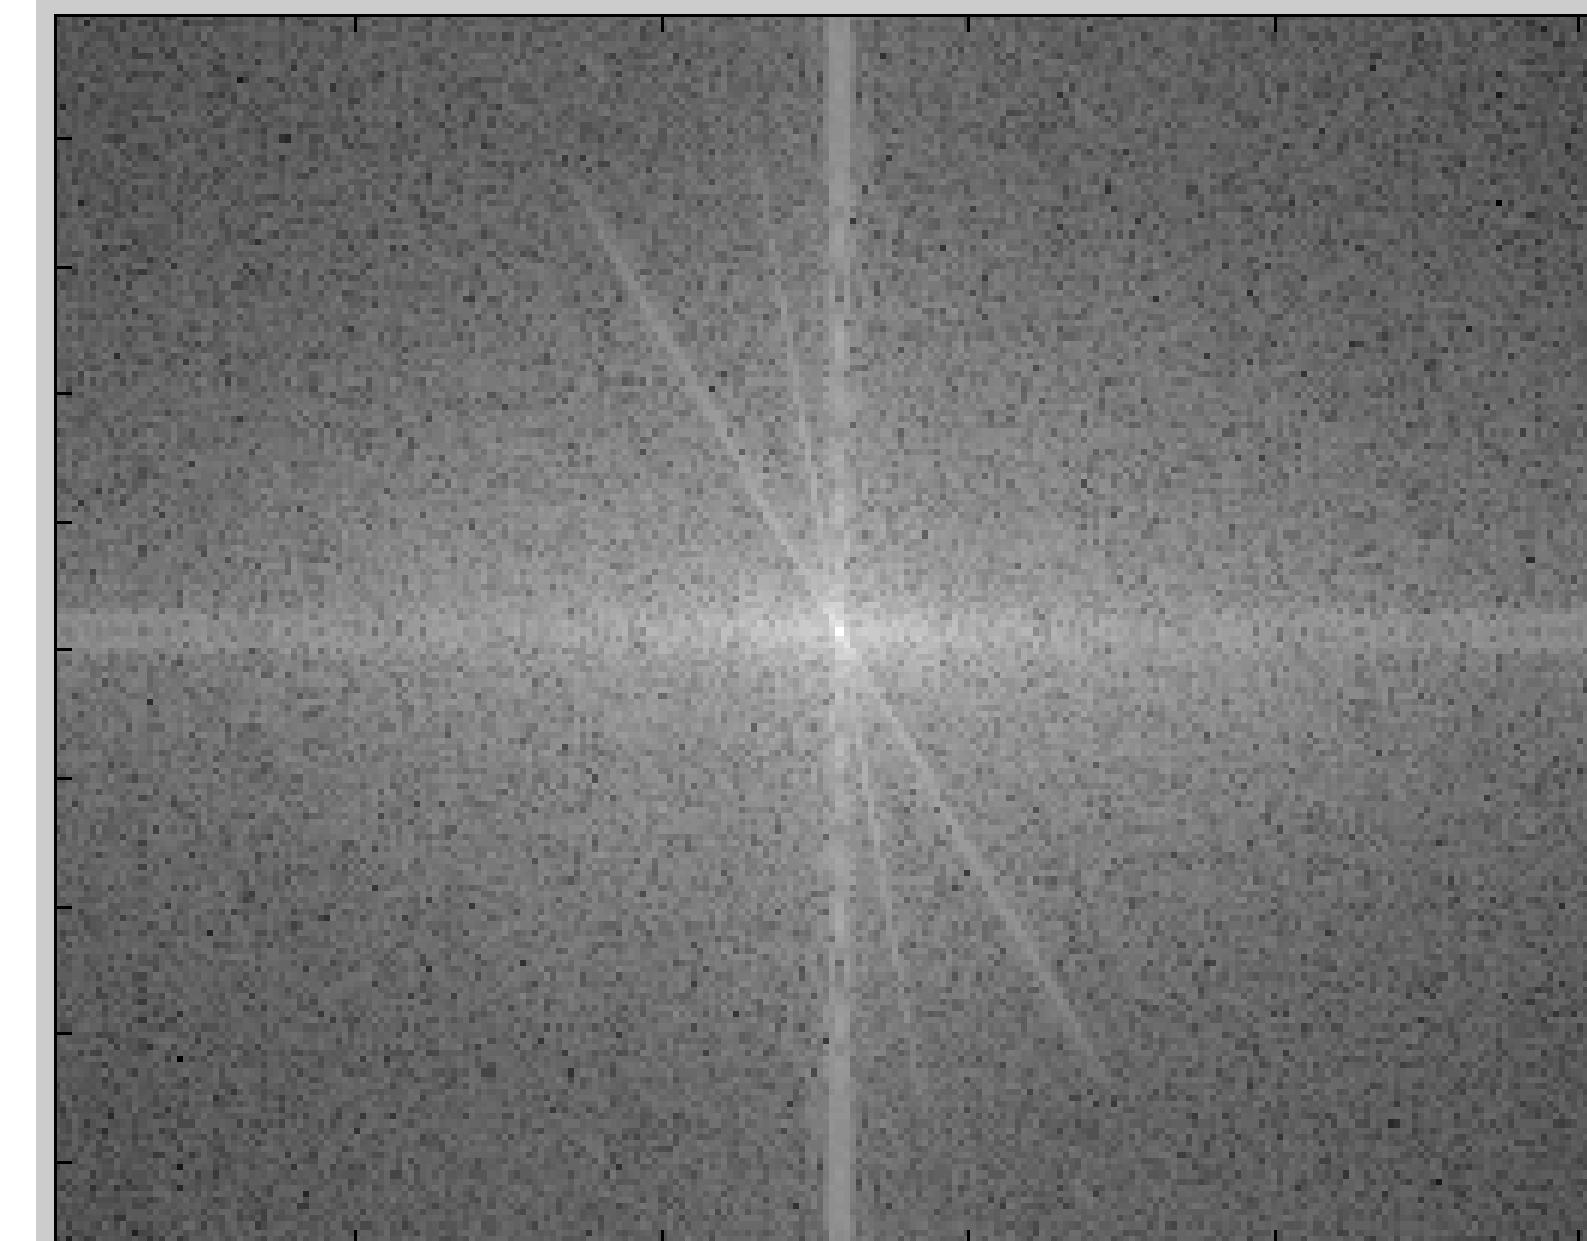


Local change in one domain, causes global change in the other

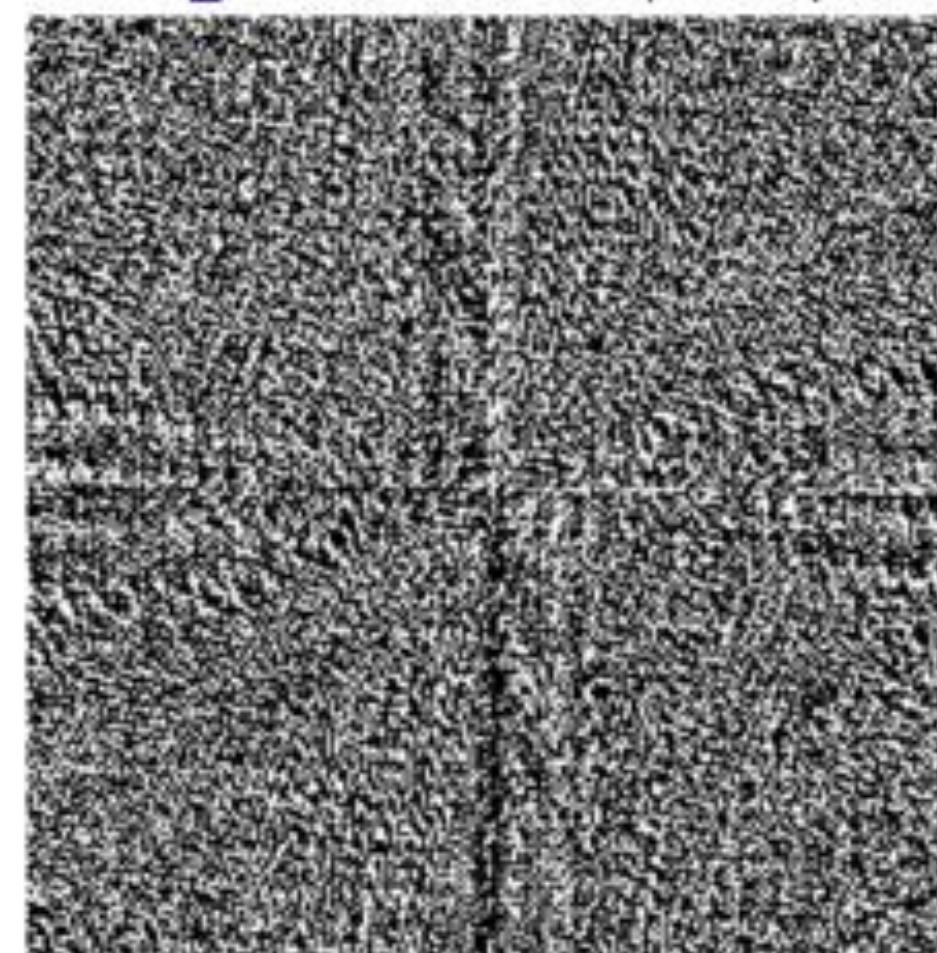
# What about phase?

---

Amplitude Spectrum

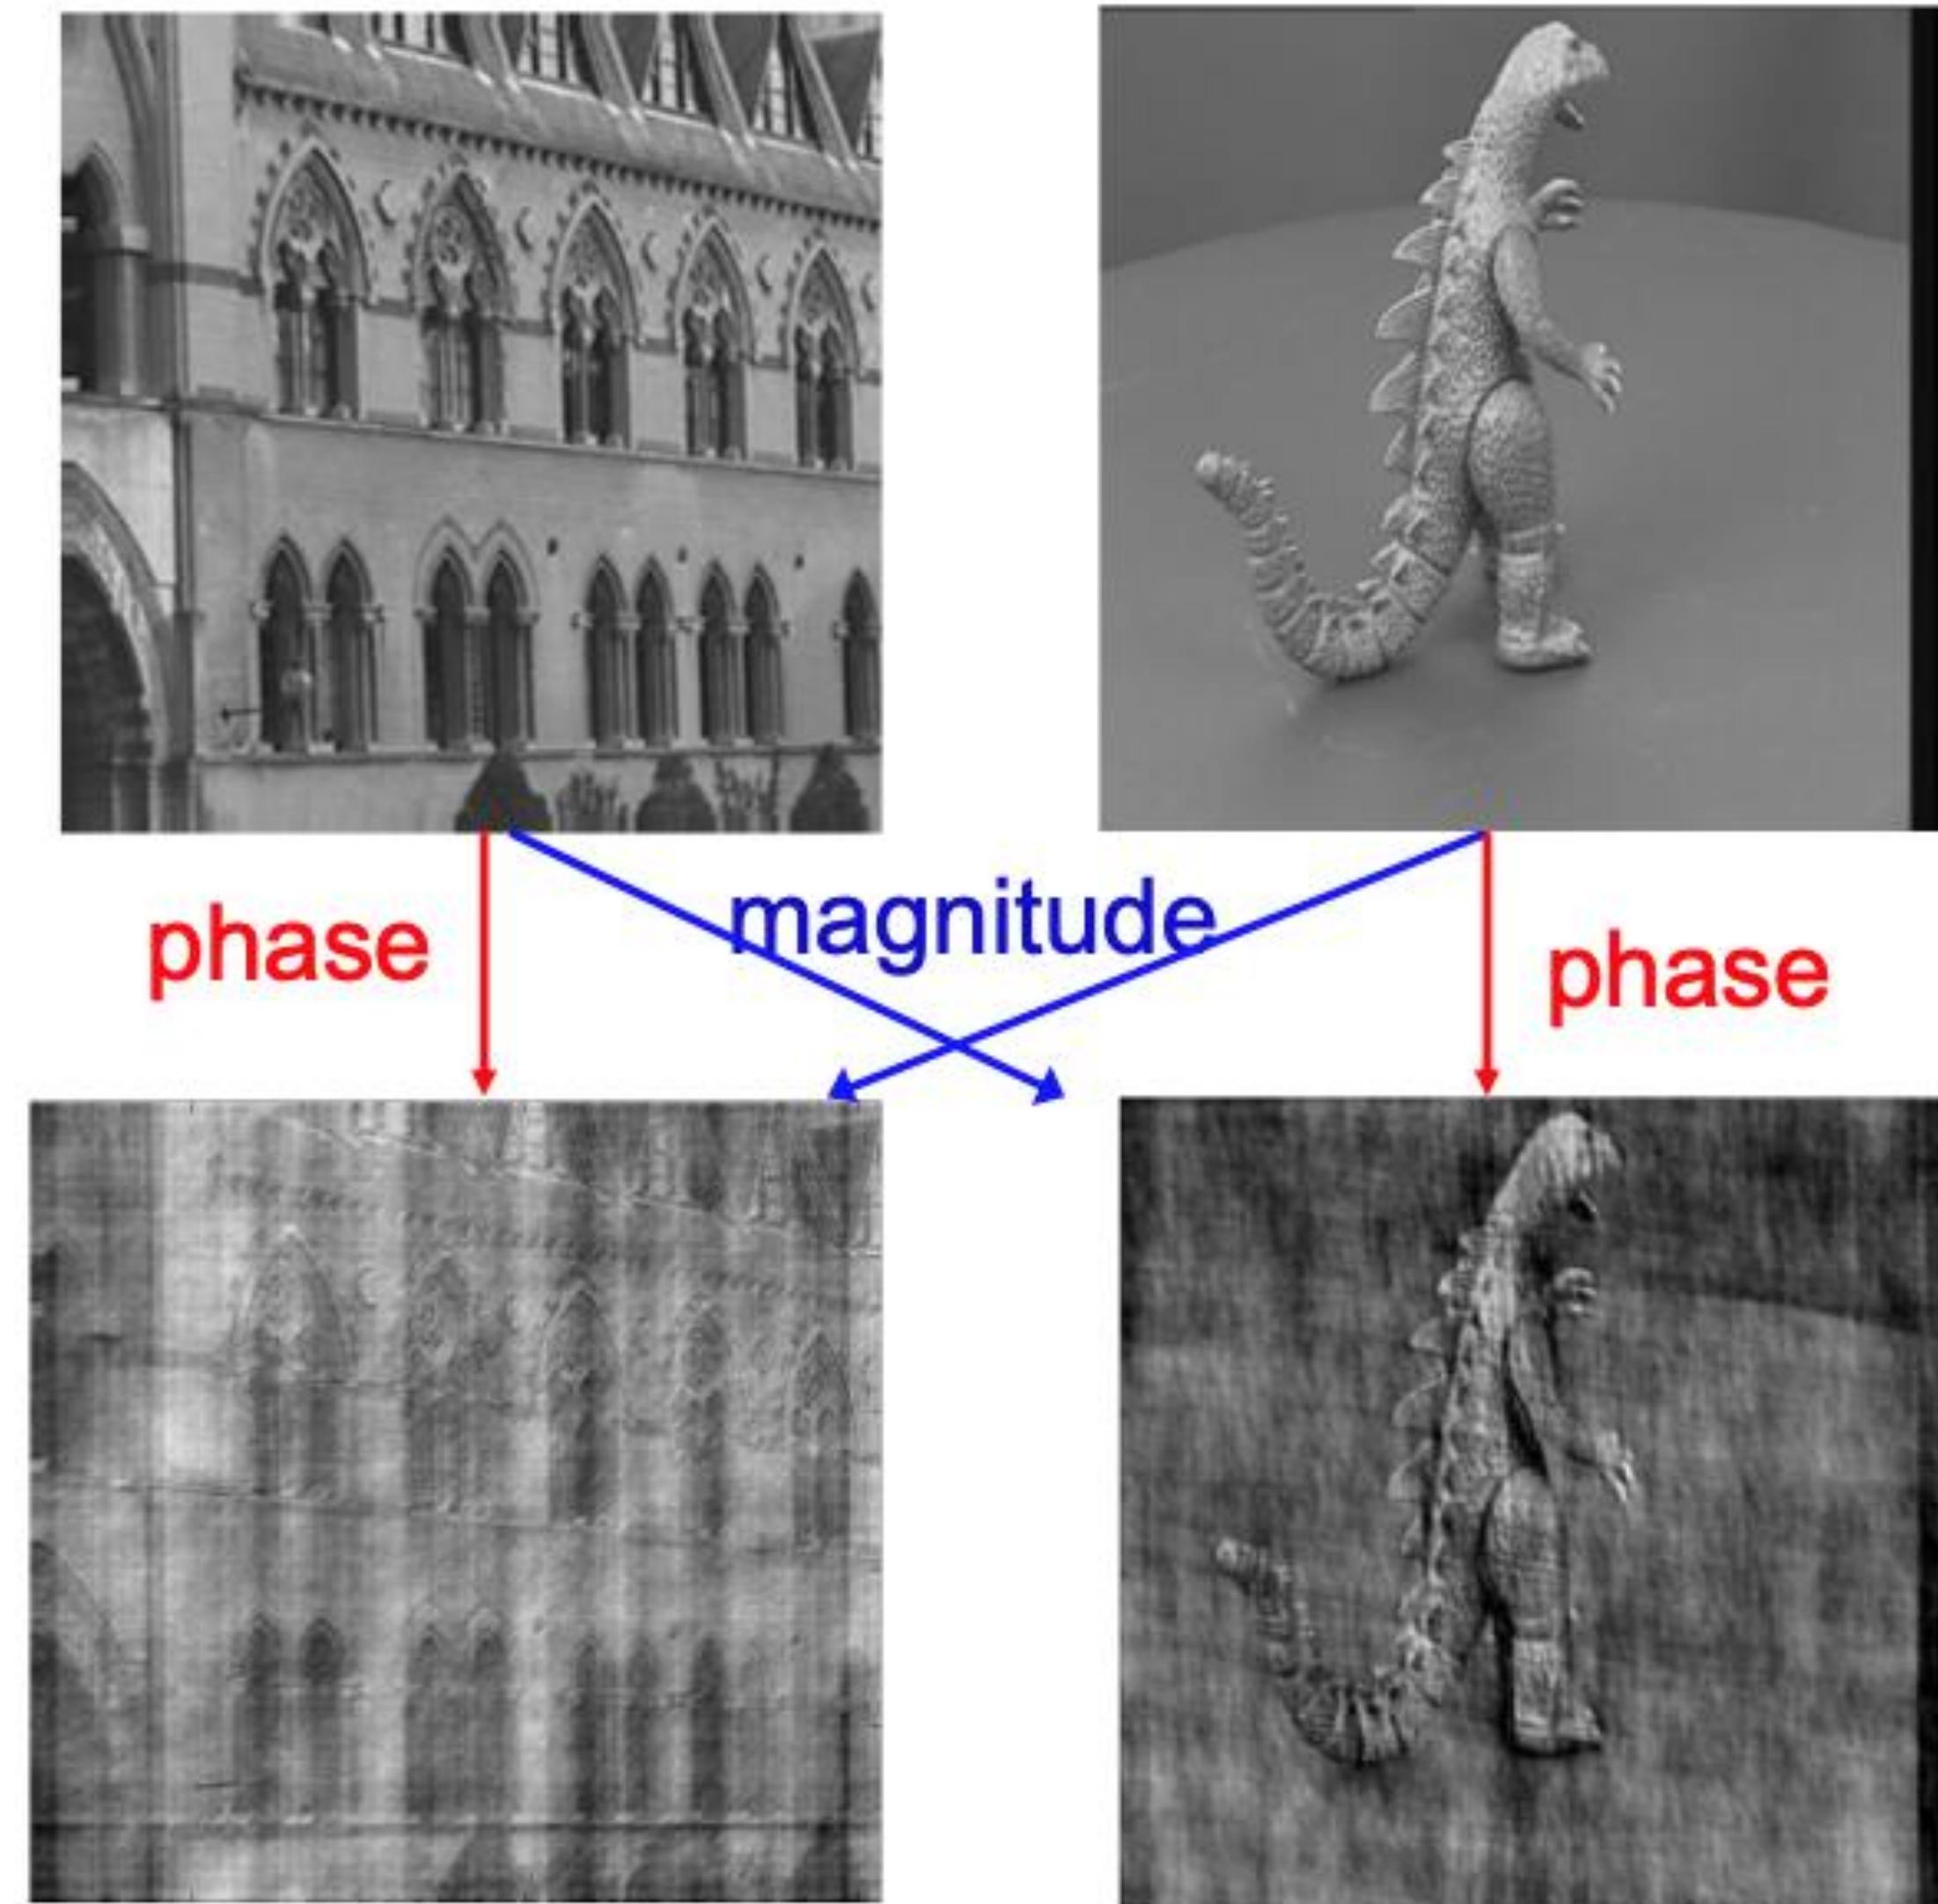


what does phase look like, you ask?  
(less visually informative)

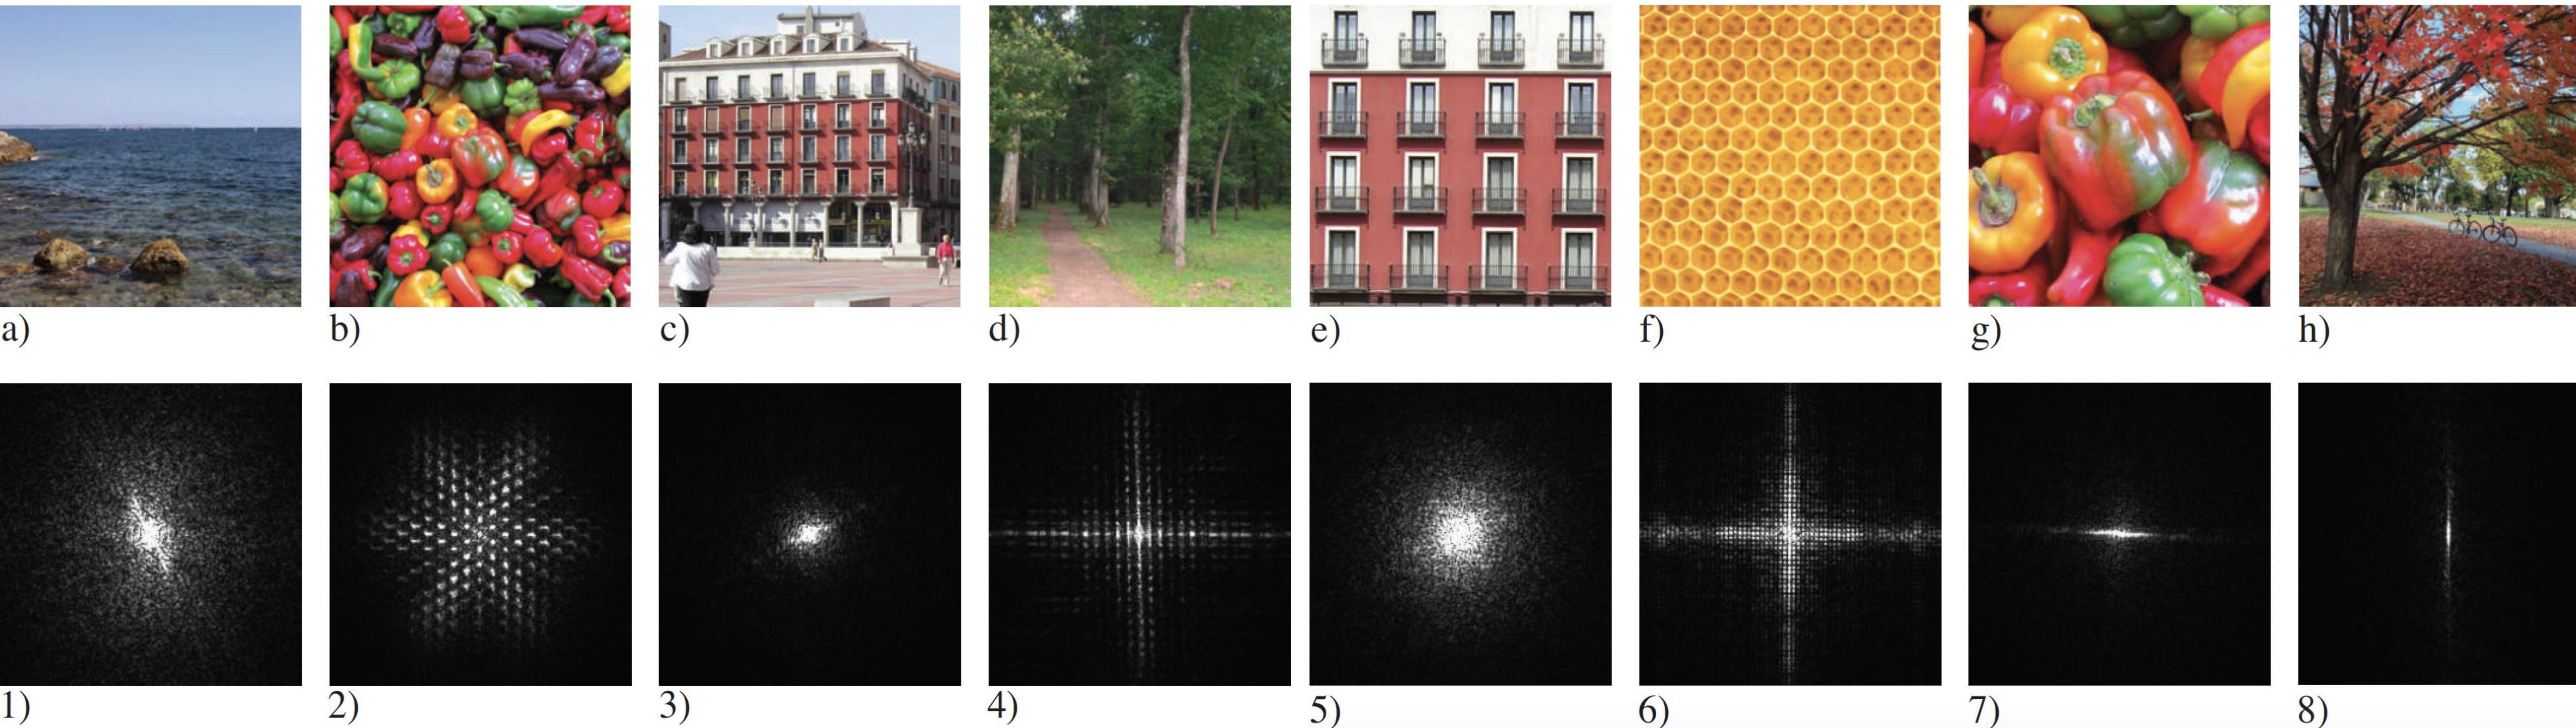


# The importance of Phase

---

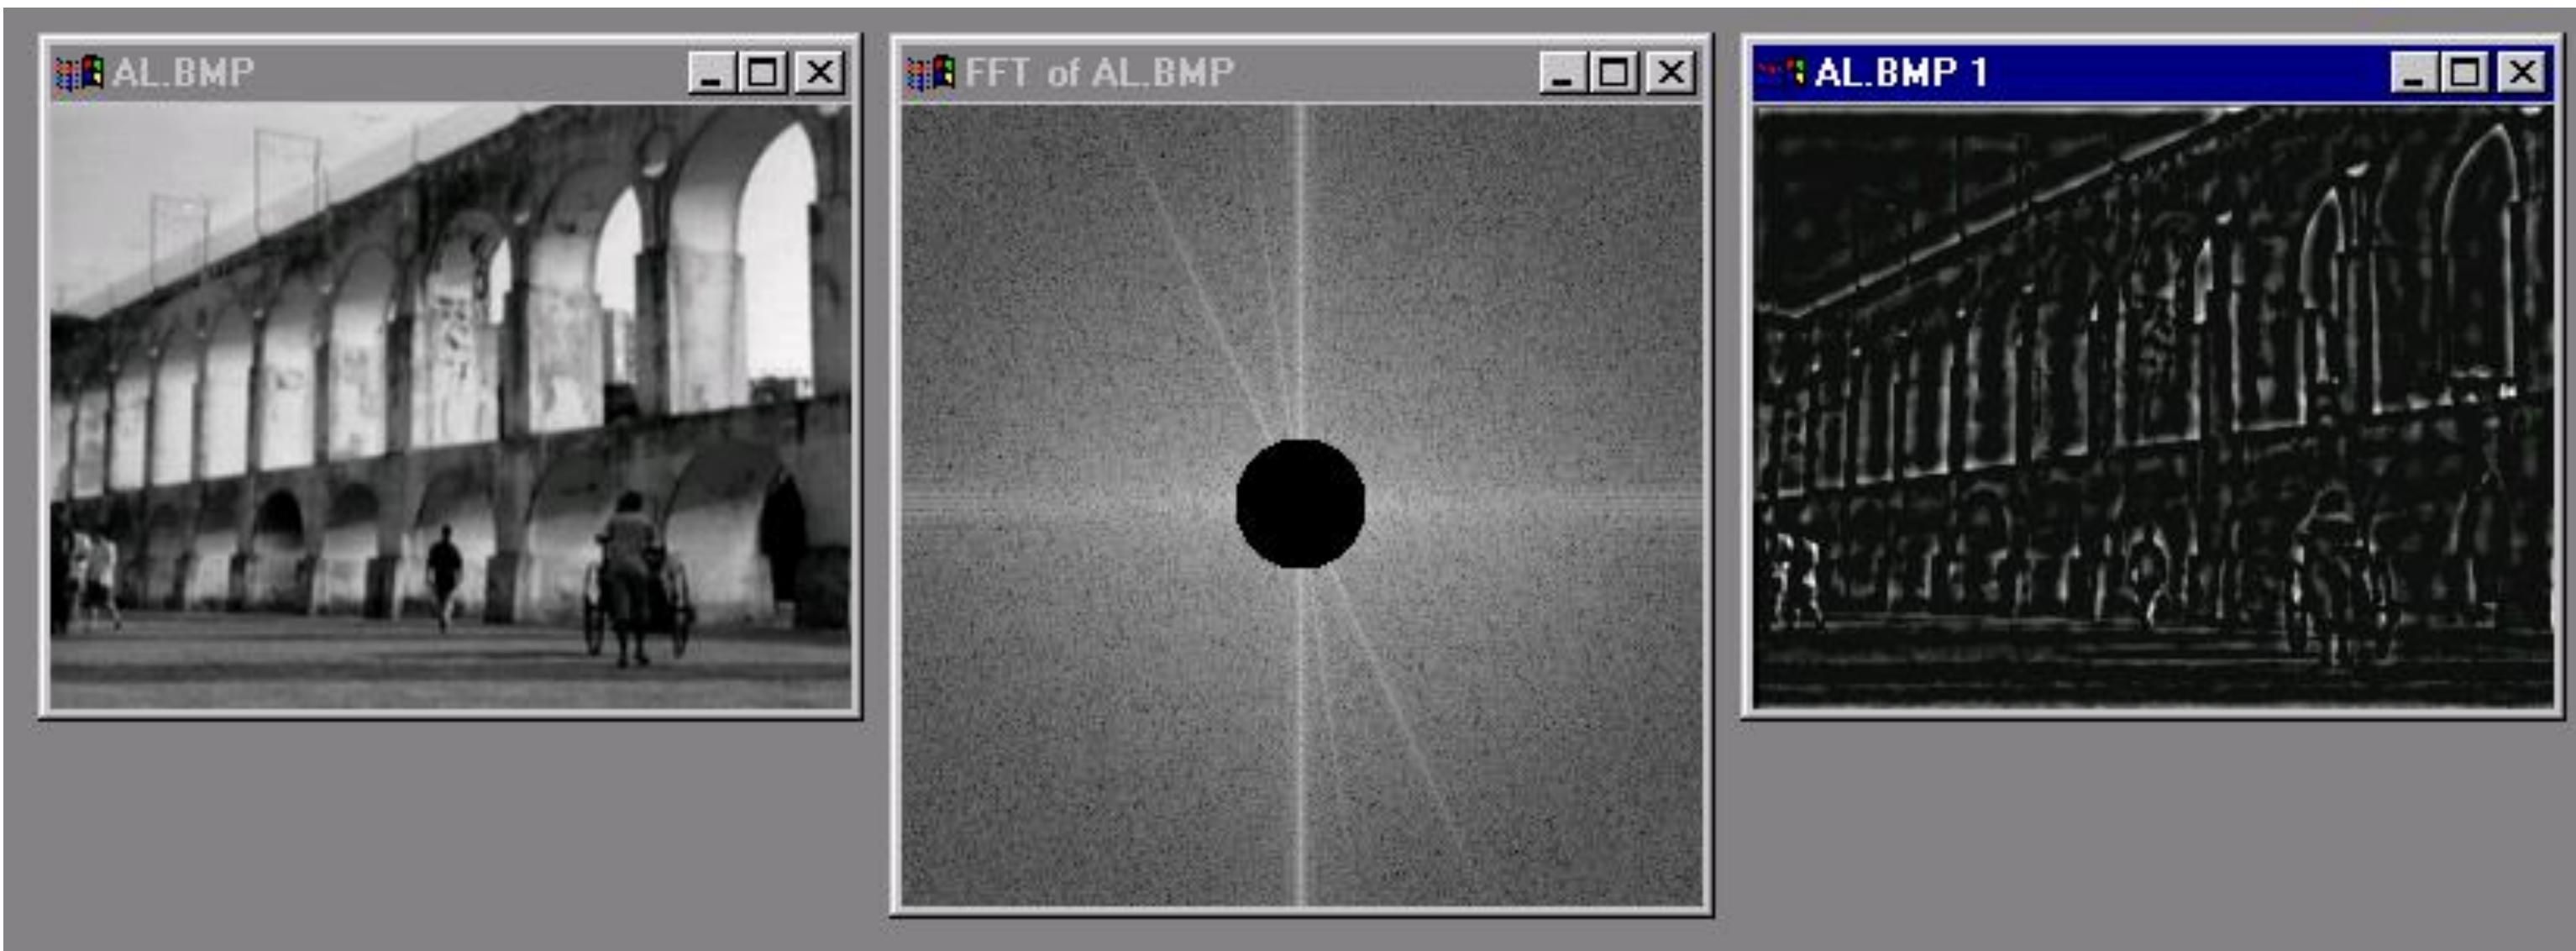


# The DFT Game: find the right pairs



# Low and High Pass filtering

---



# The Convolution Theorem

---

The greatest thing since sliced (banana) bread!

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

- The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms

$$\mathcal{F}^{-1}[gh] = \mathcal{F}^{-1}[g] * \mathcal{F}^{-1}[h]$$

- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

# 2D convolution theorem example

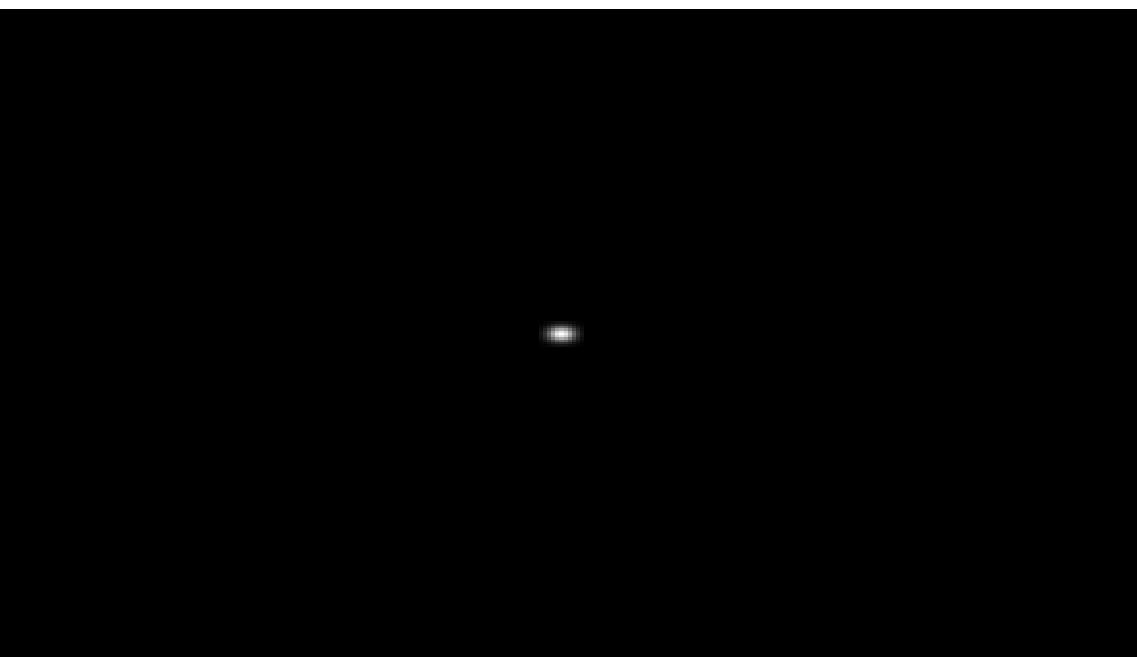
---

$f(x,y)$



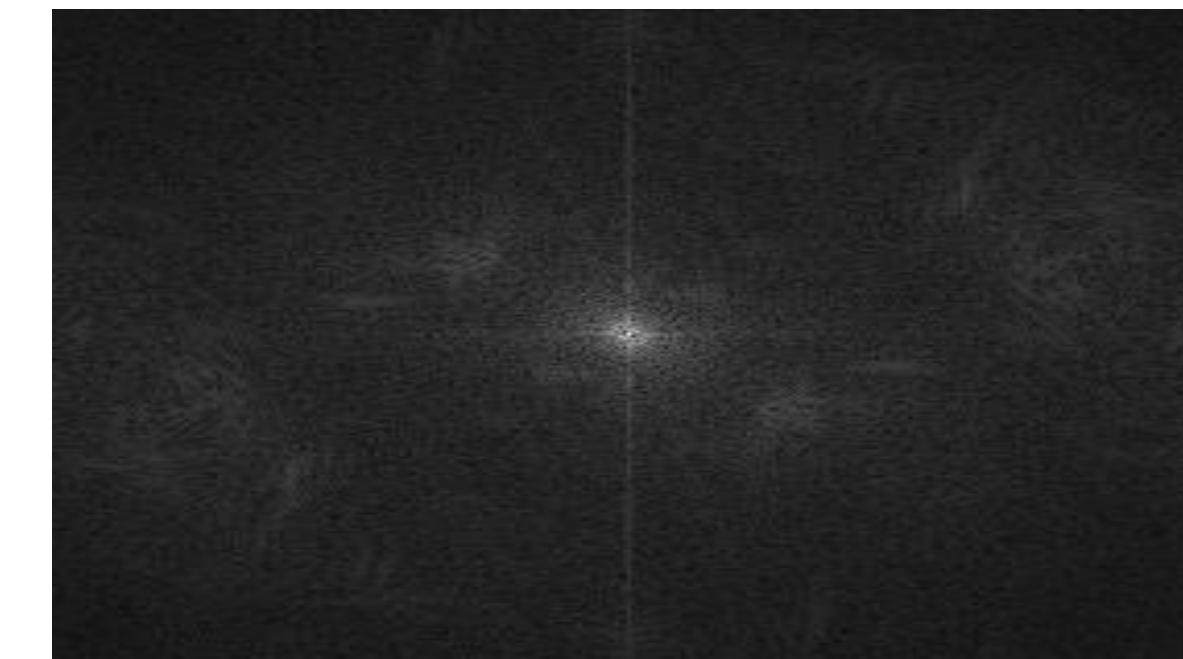
\*

$h(x,y)$



↓↓

$g(x,y)$



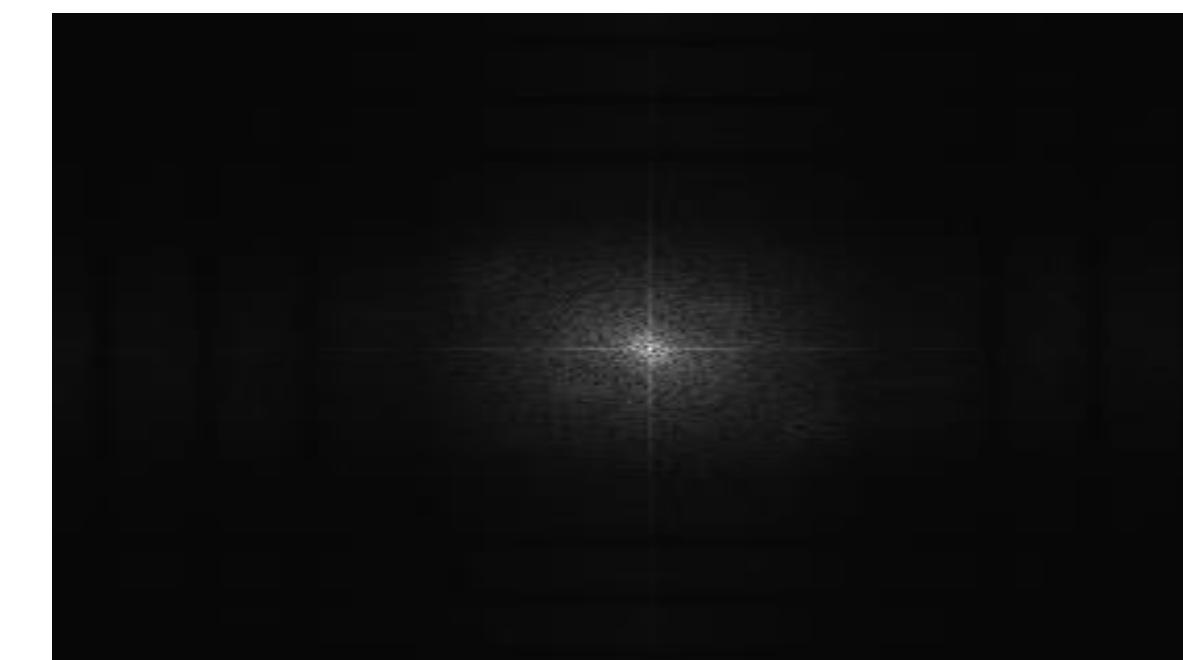
$|F(s_x,s_y)|$

×



$|H(s_x,s_y)|$

↓↓

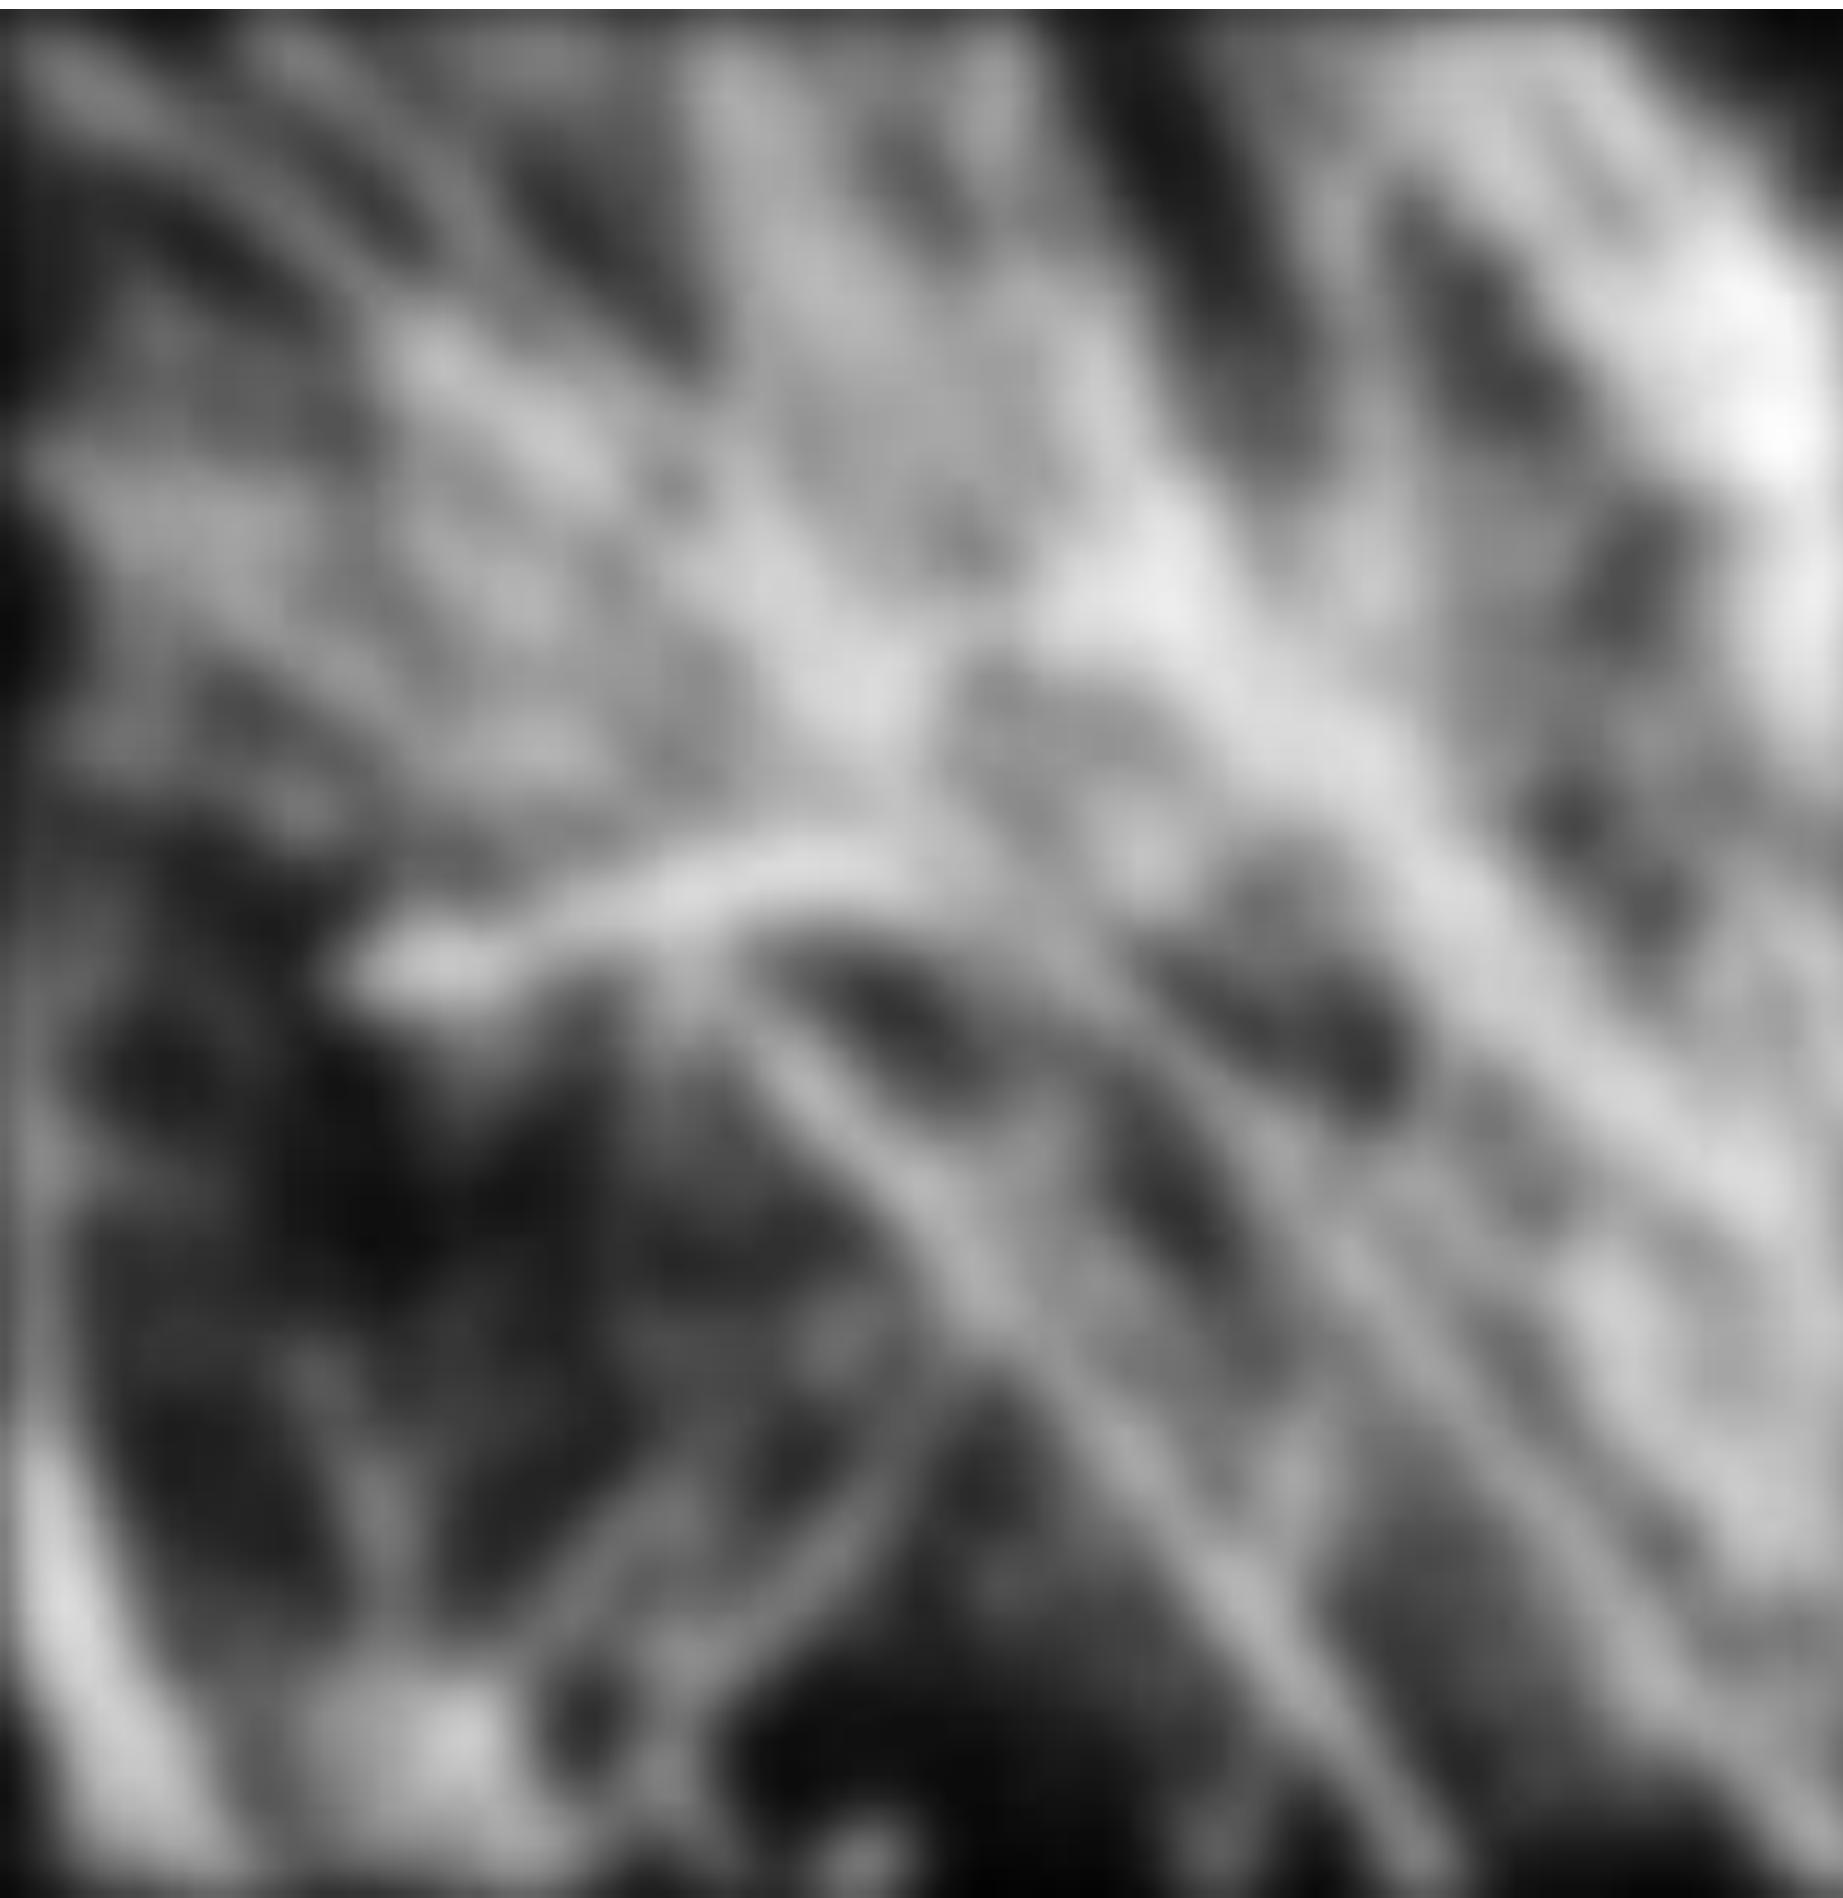
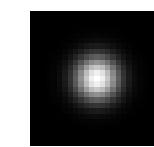


$|G(s_x,s_y)|$

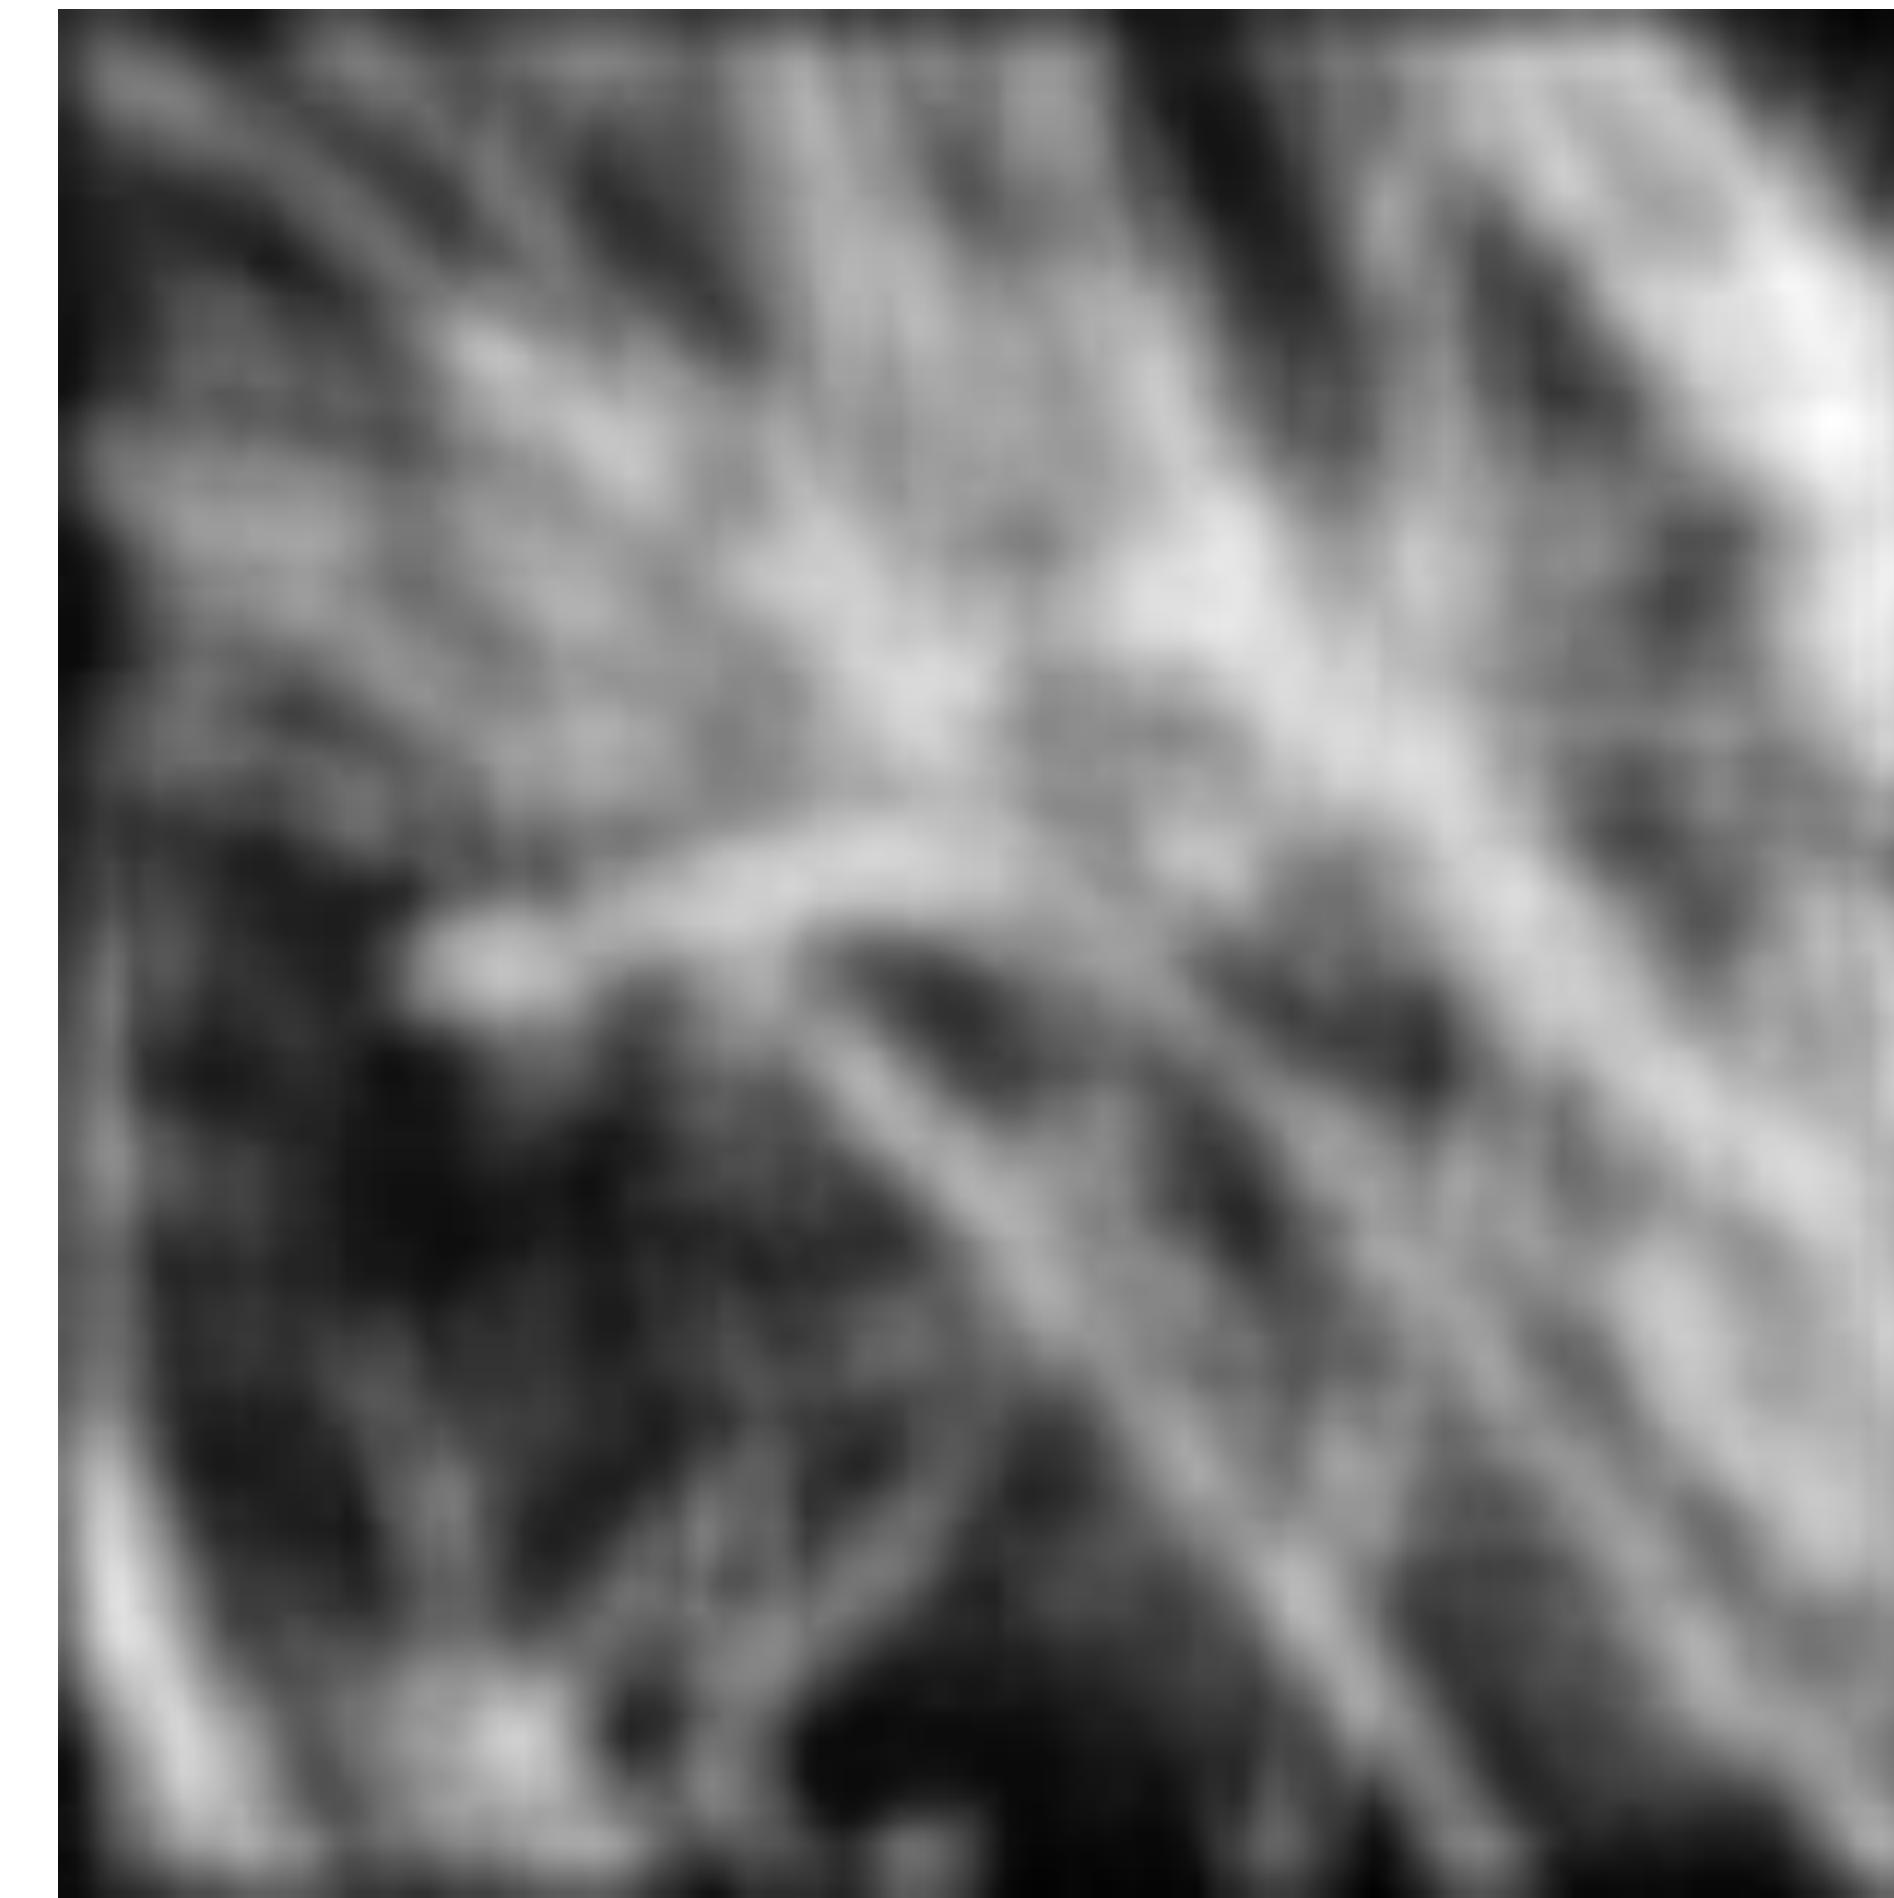
# Back to Filtering

**Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?**

Gaussian

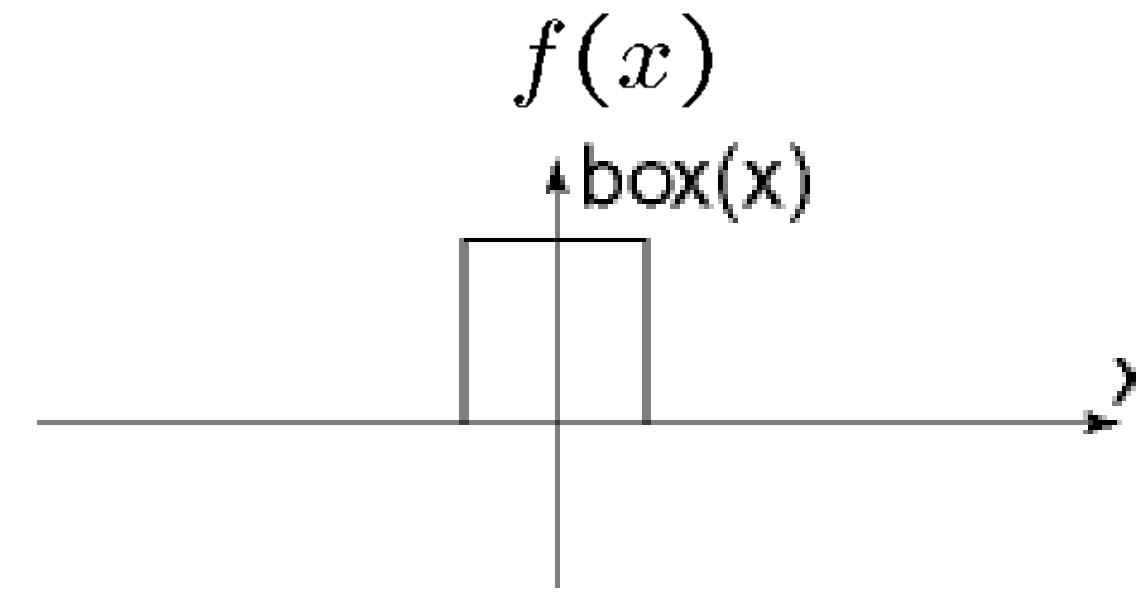


Box filter

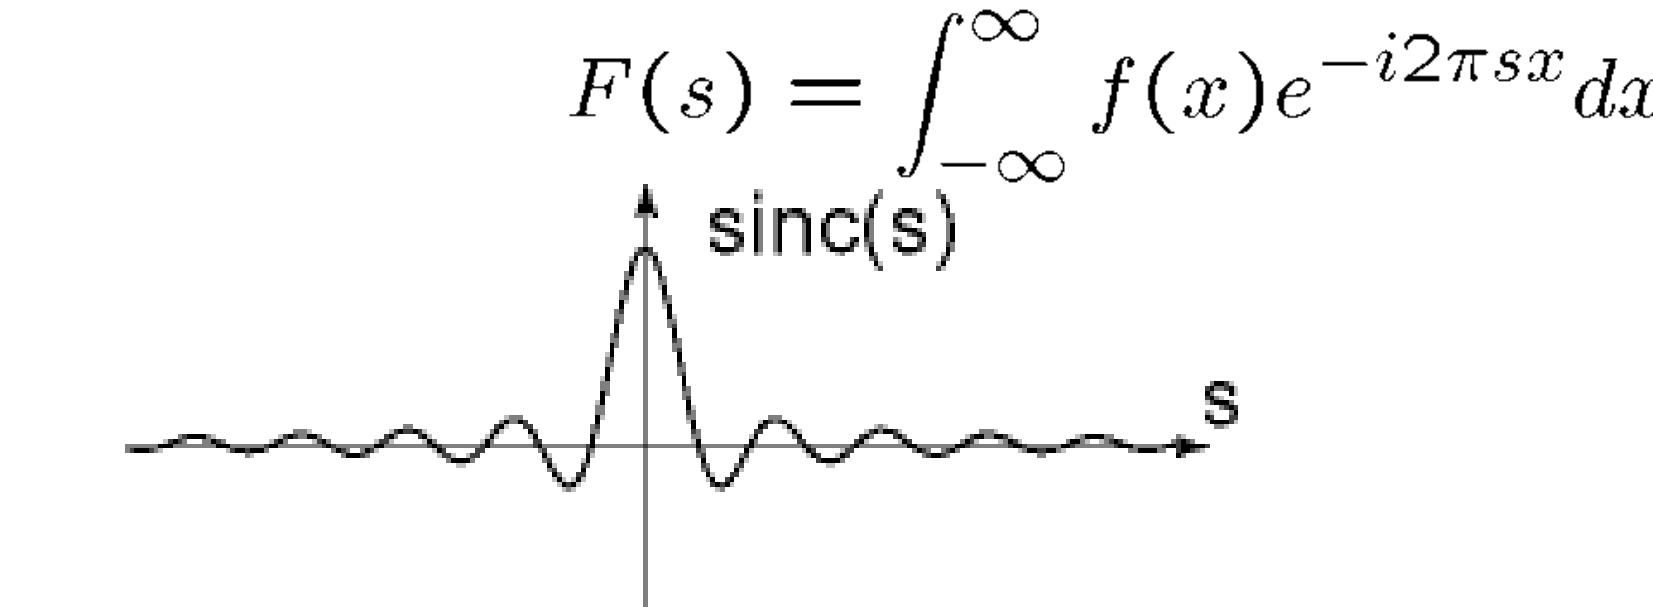


# Fourier Transform pairs

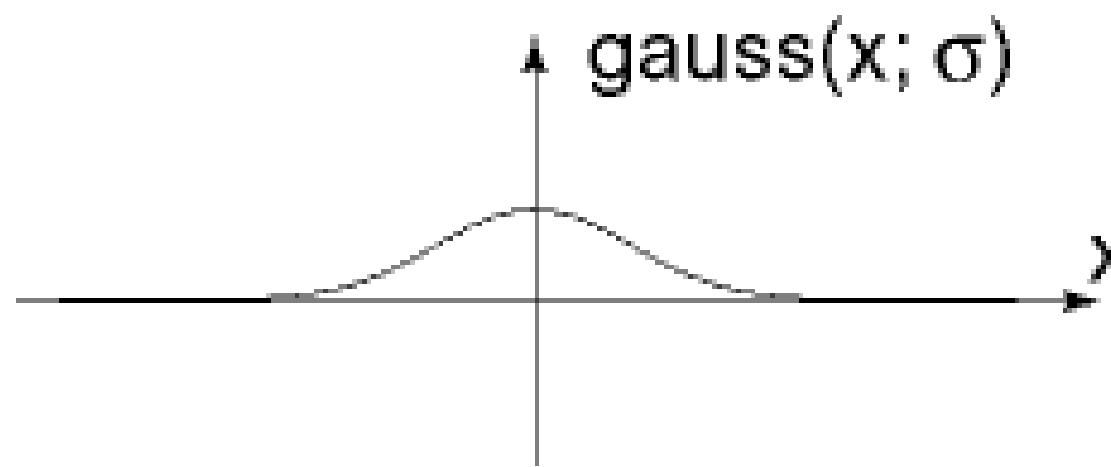
Spatial domain



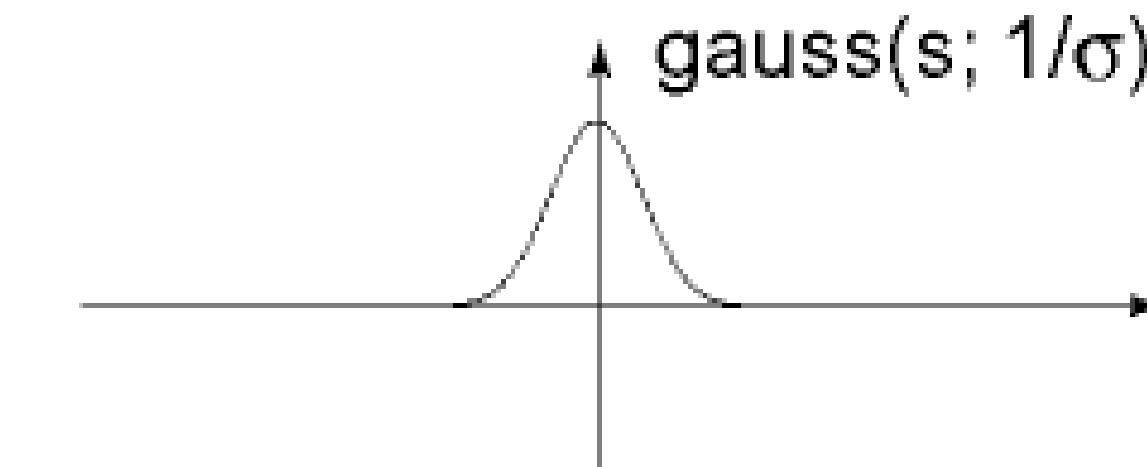
Frequency domain



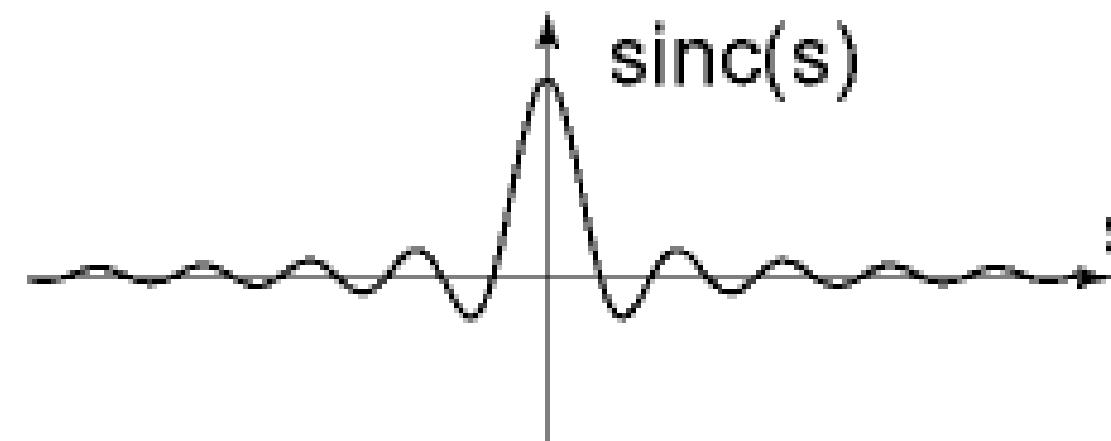
$\uparrow \text{gauss}(x; \sigma)$



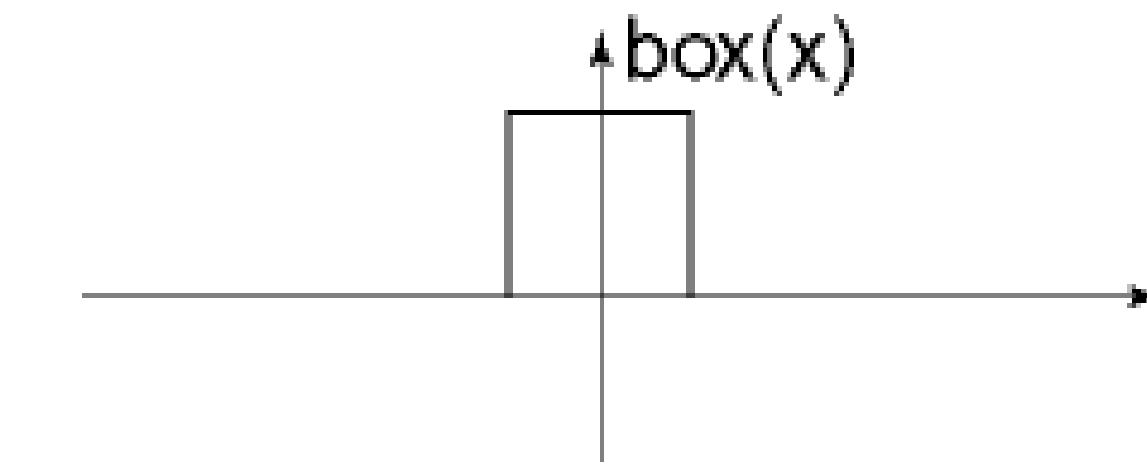
$\uparrow \text{gauss}(s; 1/\sigma)$



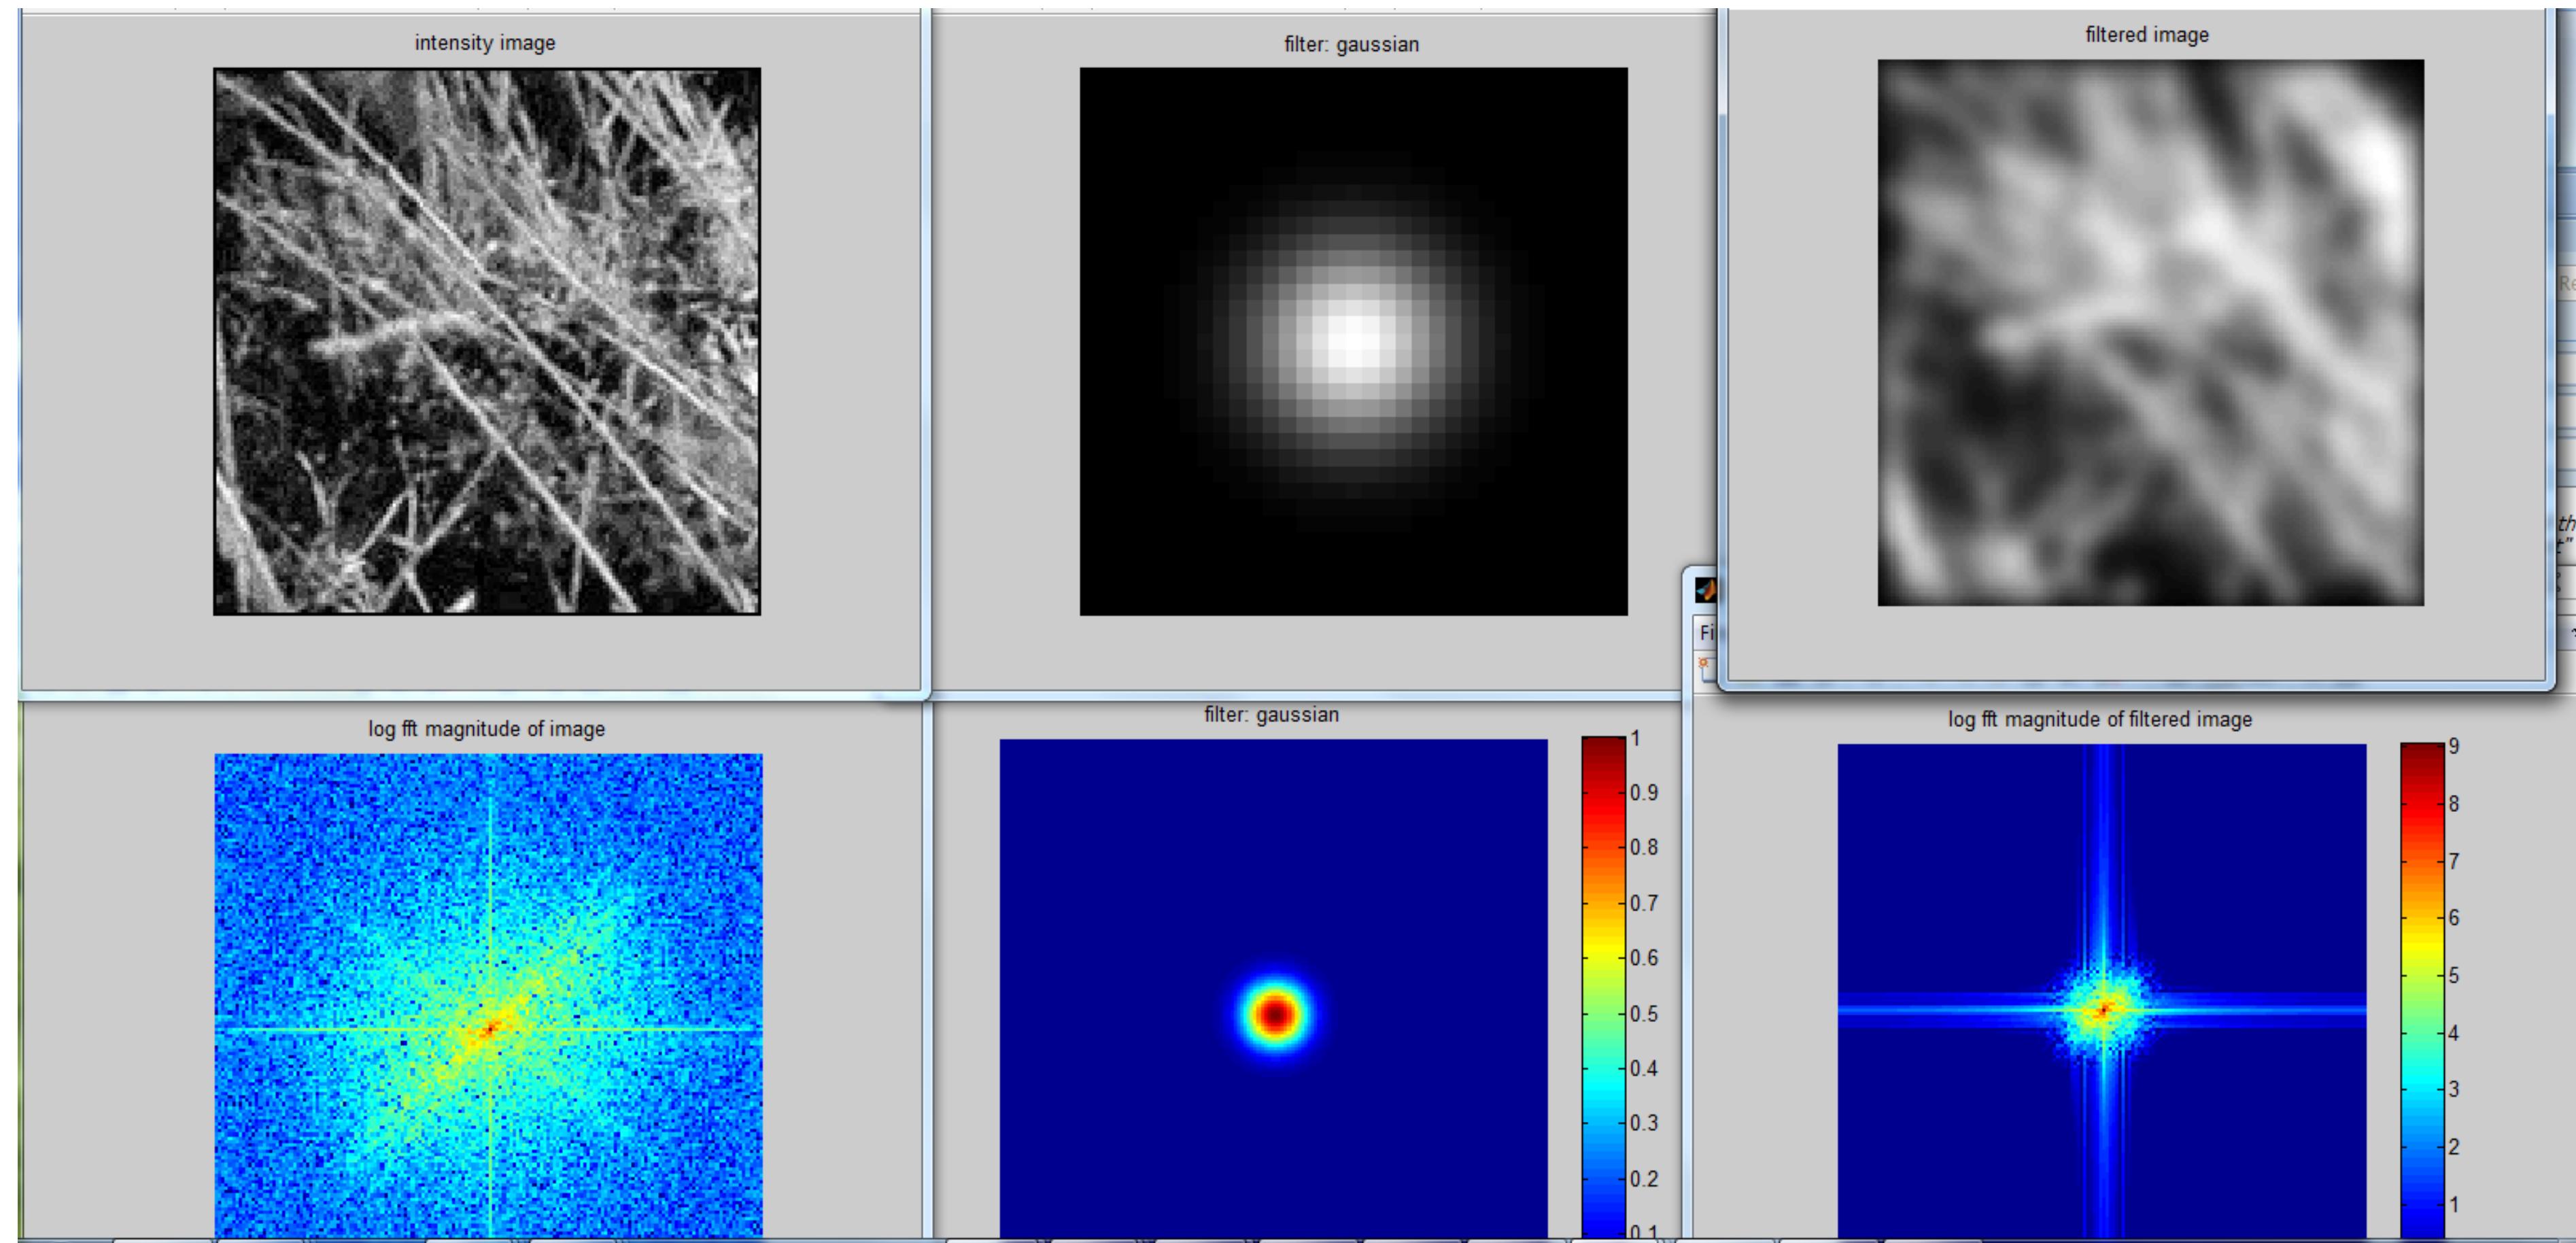
$\uparrow \text{sinc}(s)$



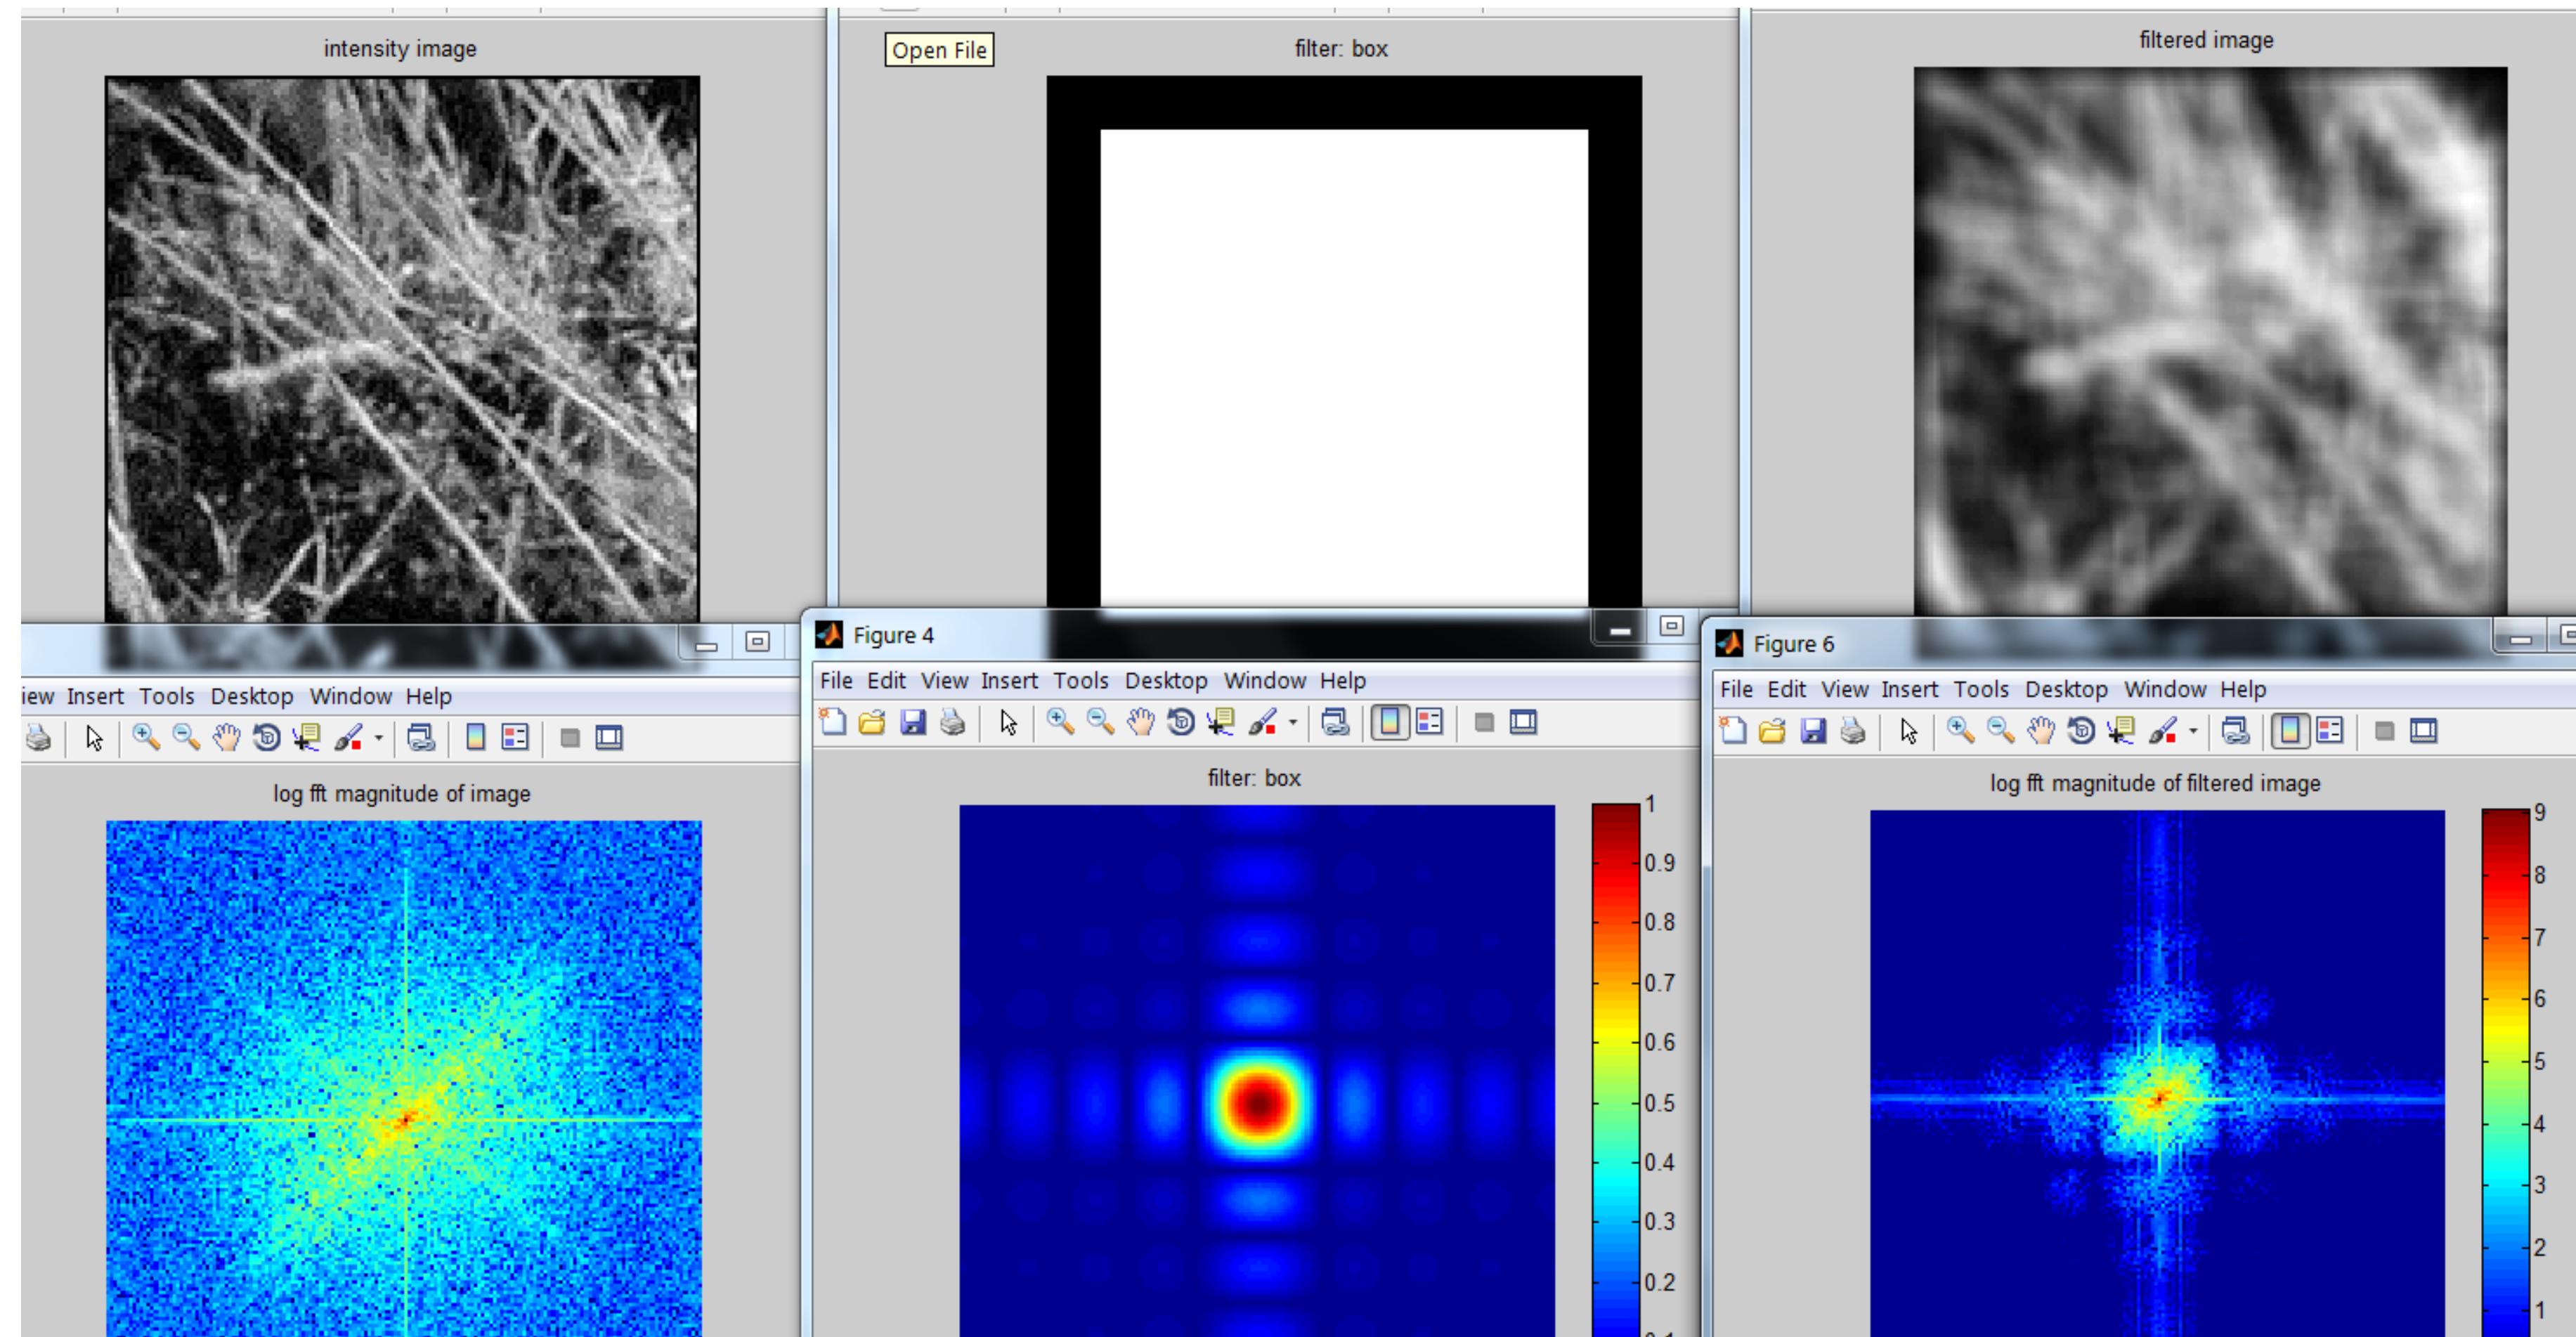
$\uparrow \text{box}(x)$



# Gaussian



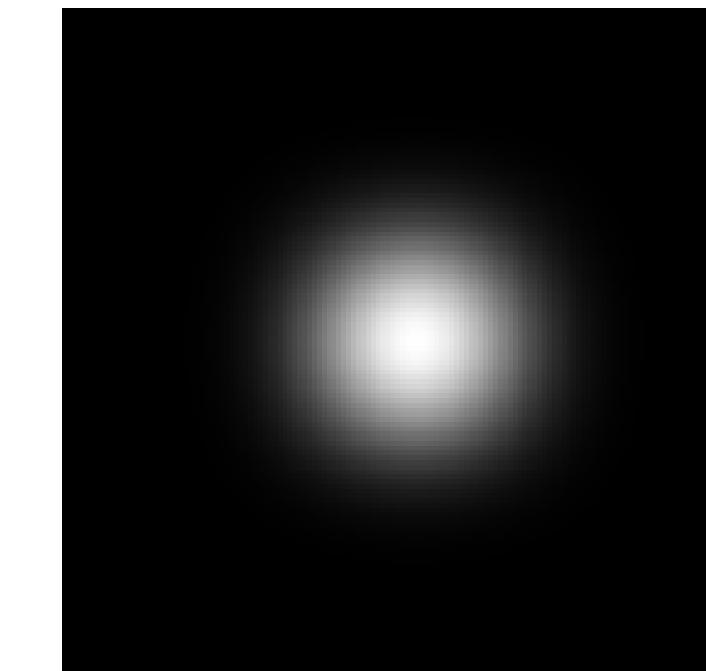
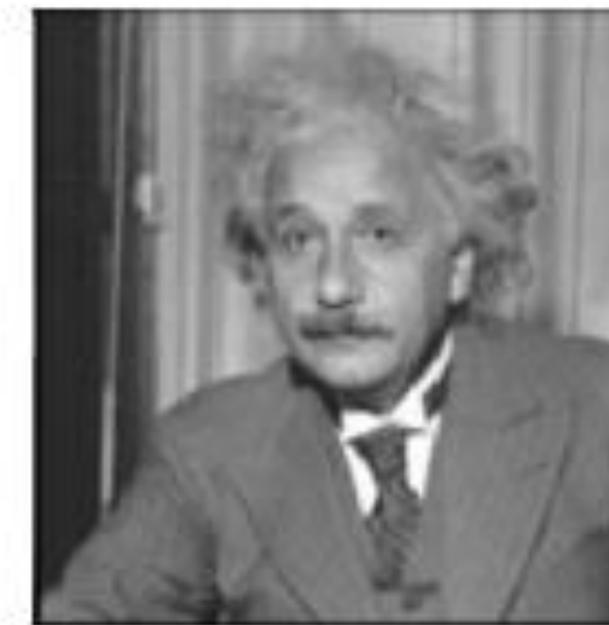
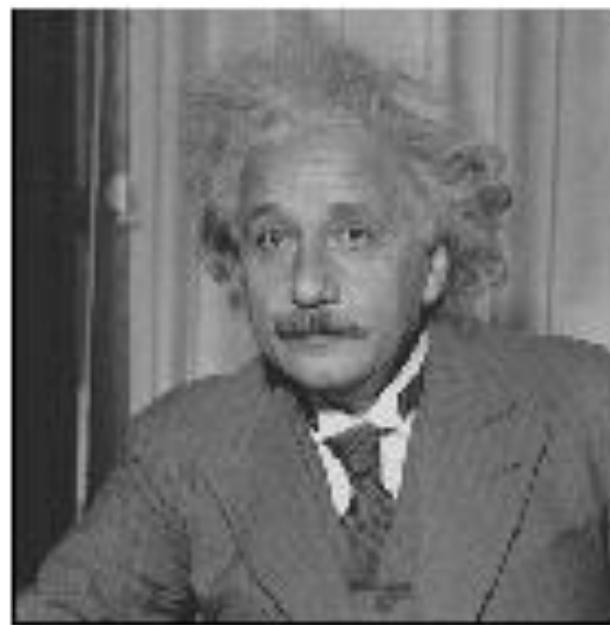
# Box Filter



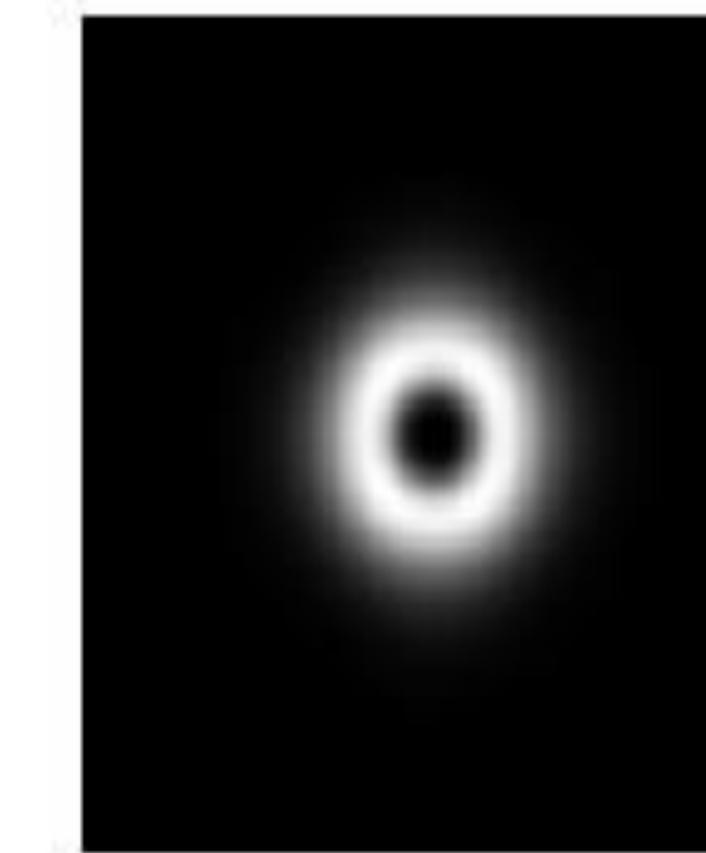
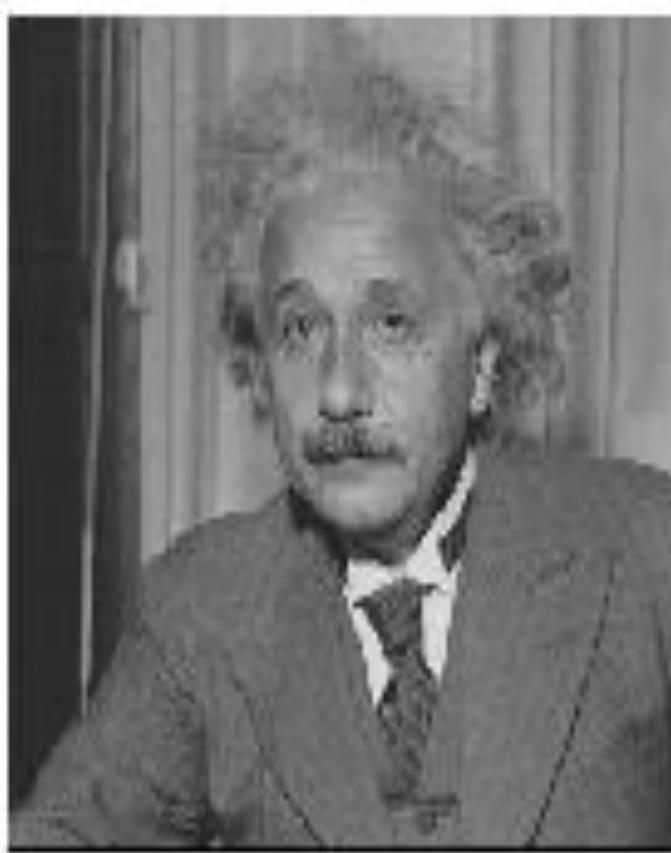
# Low-pass, Band-pass, High-pass filters

---

low-pass:



High-pass / band-pass:



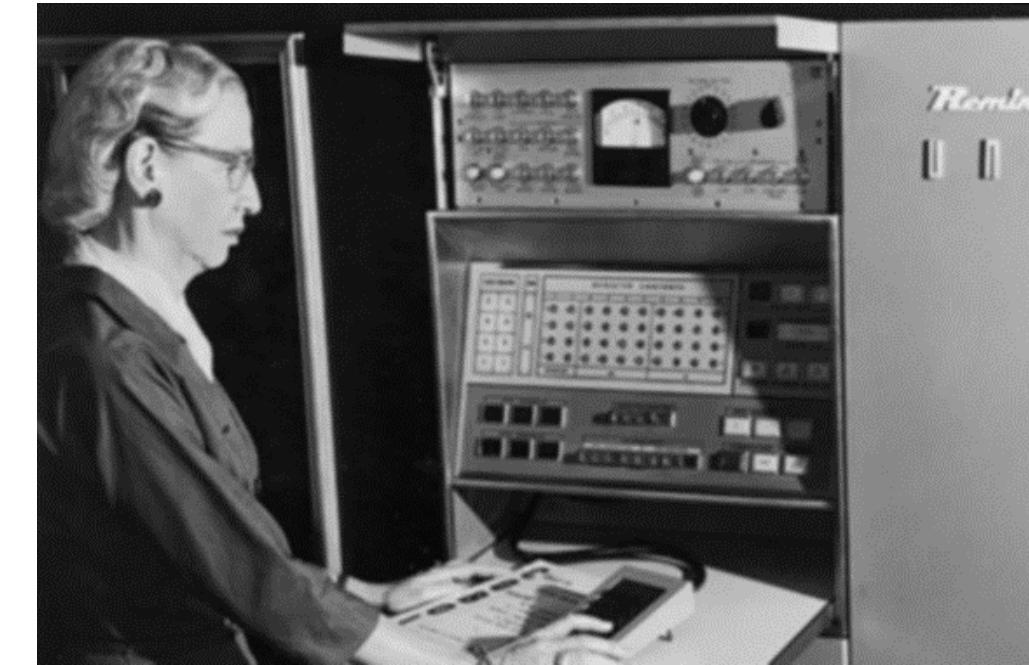
# Low Pass vs. High Pass filtering

---

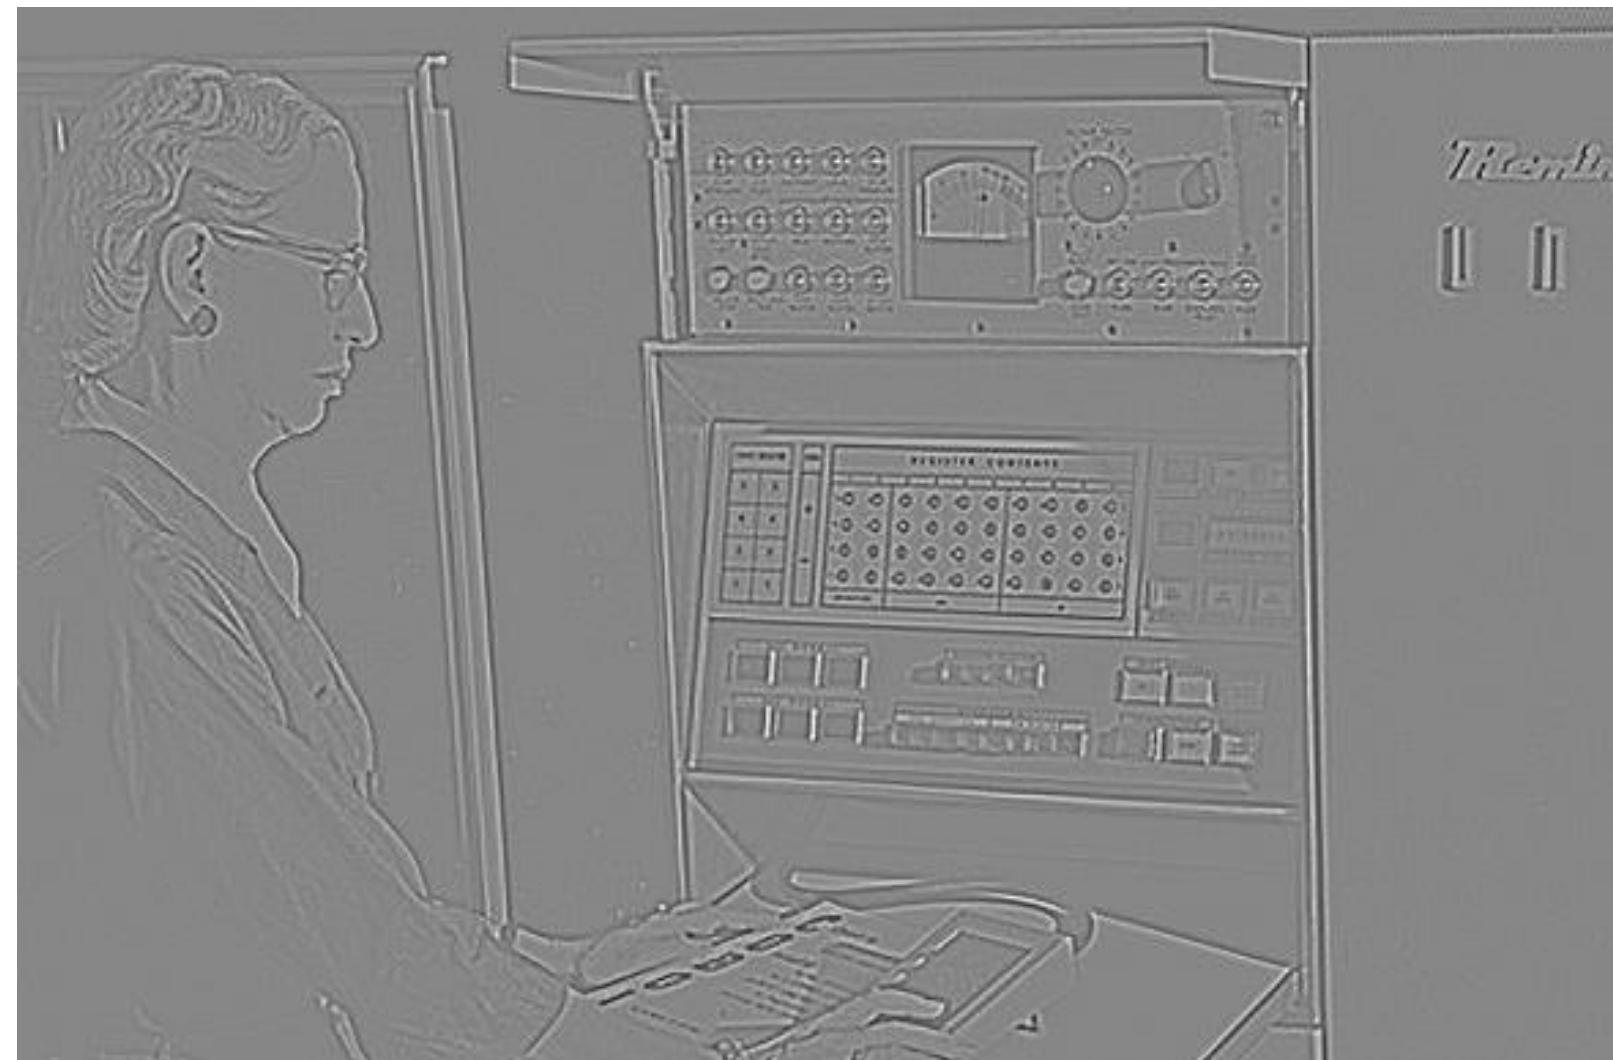
Image



Smoothed



Details



# Filtering – Sharpening

---

Image



Details



$+ \alpha$

“Sharpened”  $\alpha=1$

=



# Filtering – Sharpening

---

Image



Details



$+ \alpha$

“Sharpened”  $\alpha=0$

=



# Filtering – Sharpening

---

Image



Details



$+ \alpha$

“Sharpened”  $\alpha=2$

=



# Filtering – Sharpening

---

Image



Details



$+ \alpha$

“Sharpened”  $\alpha=0$

=



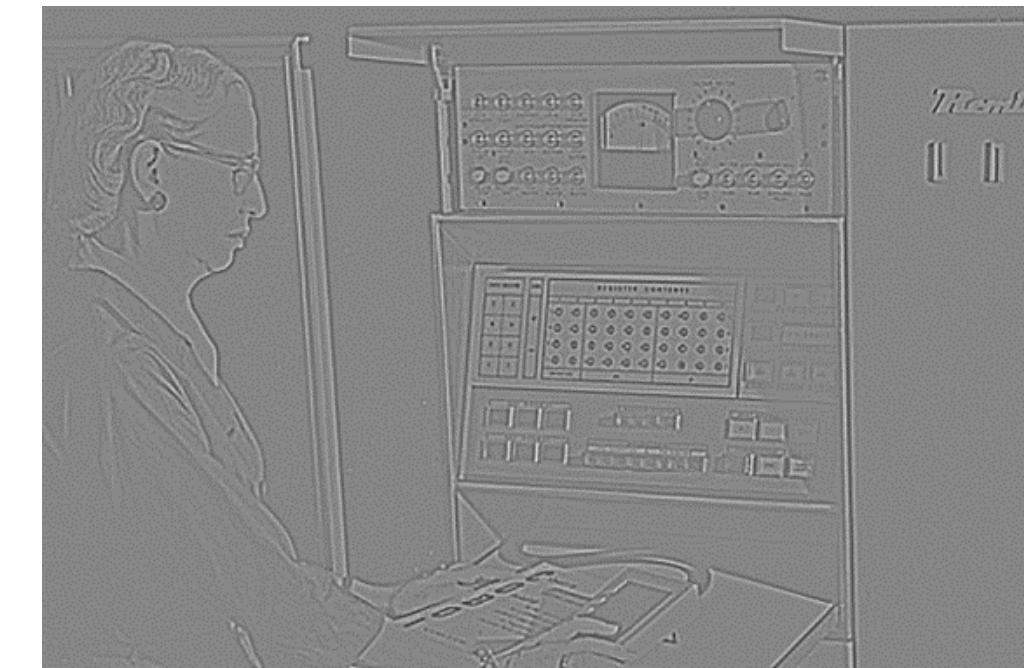
# Filtering – Extreme Sharpening

---

Image



Details



$+ \alpha$

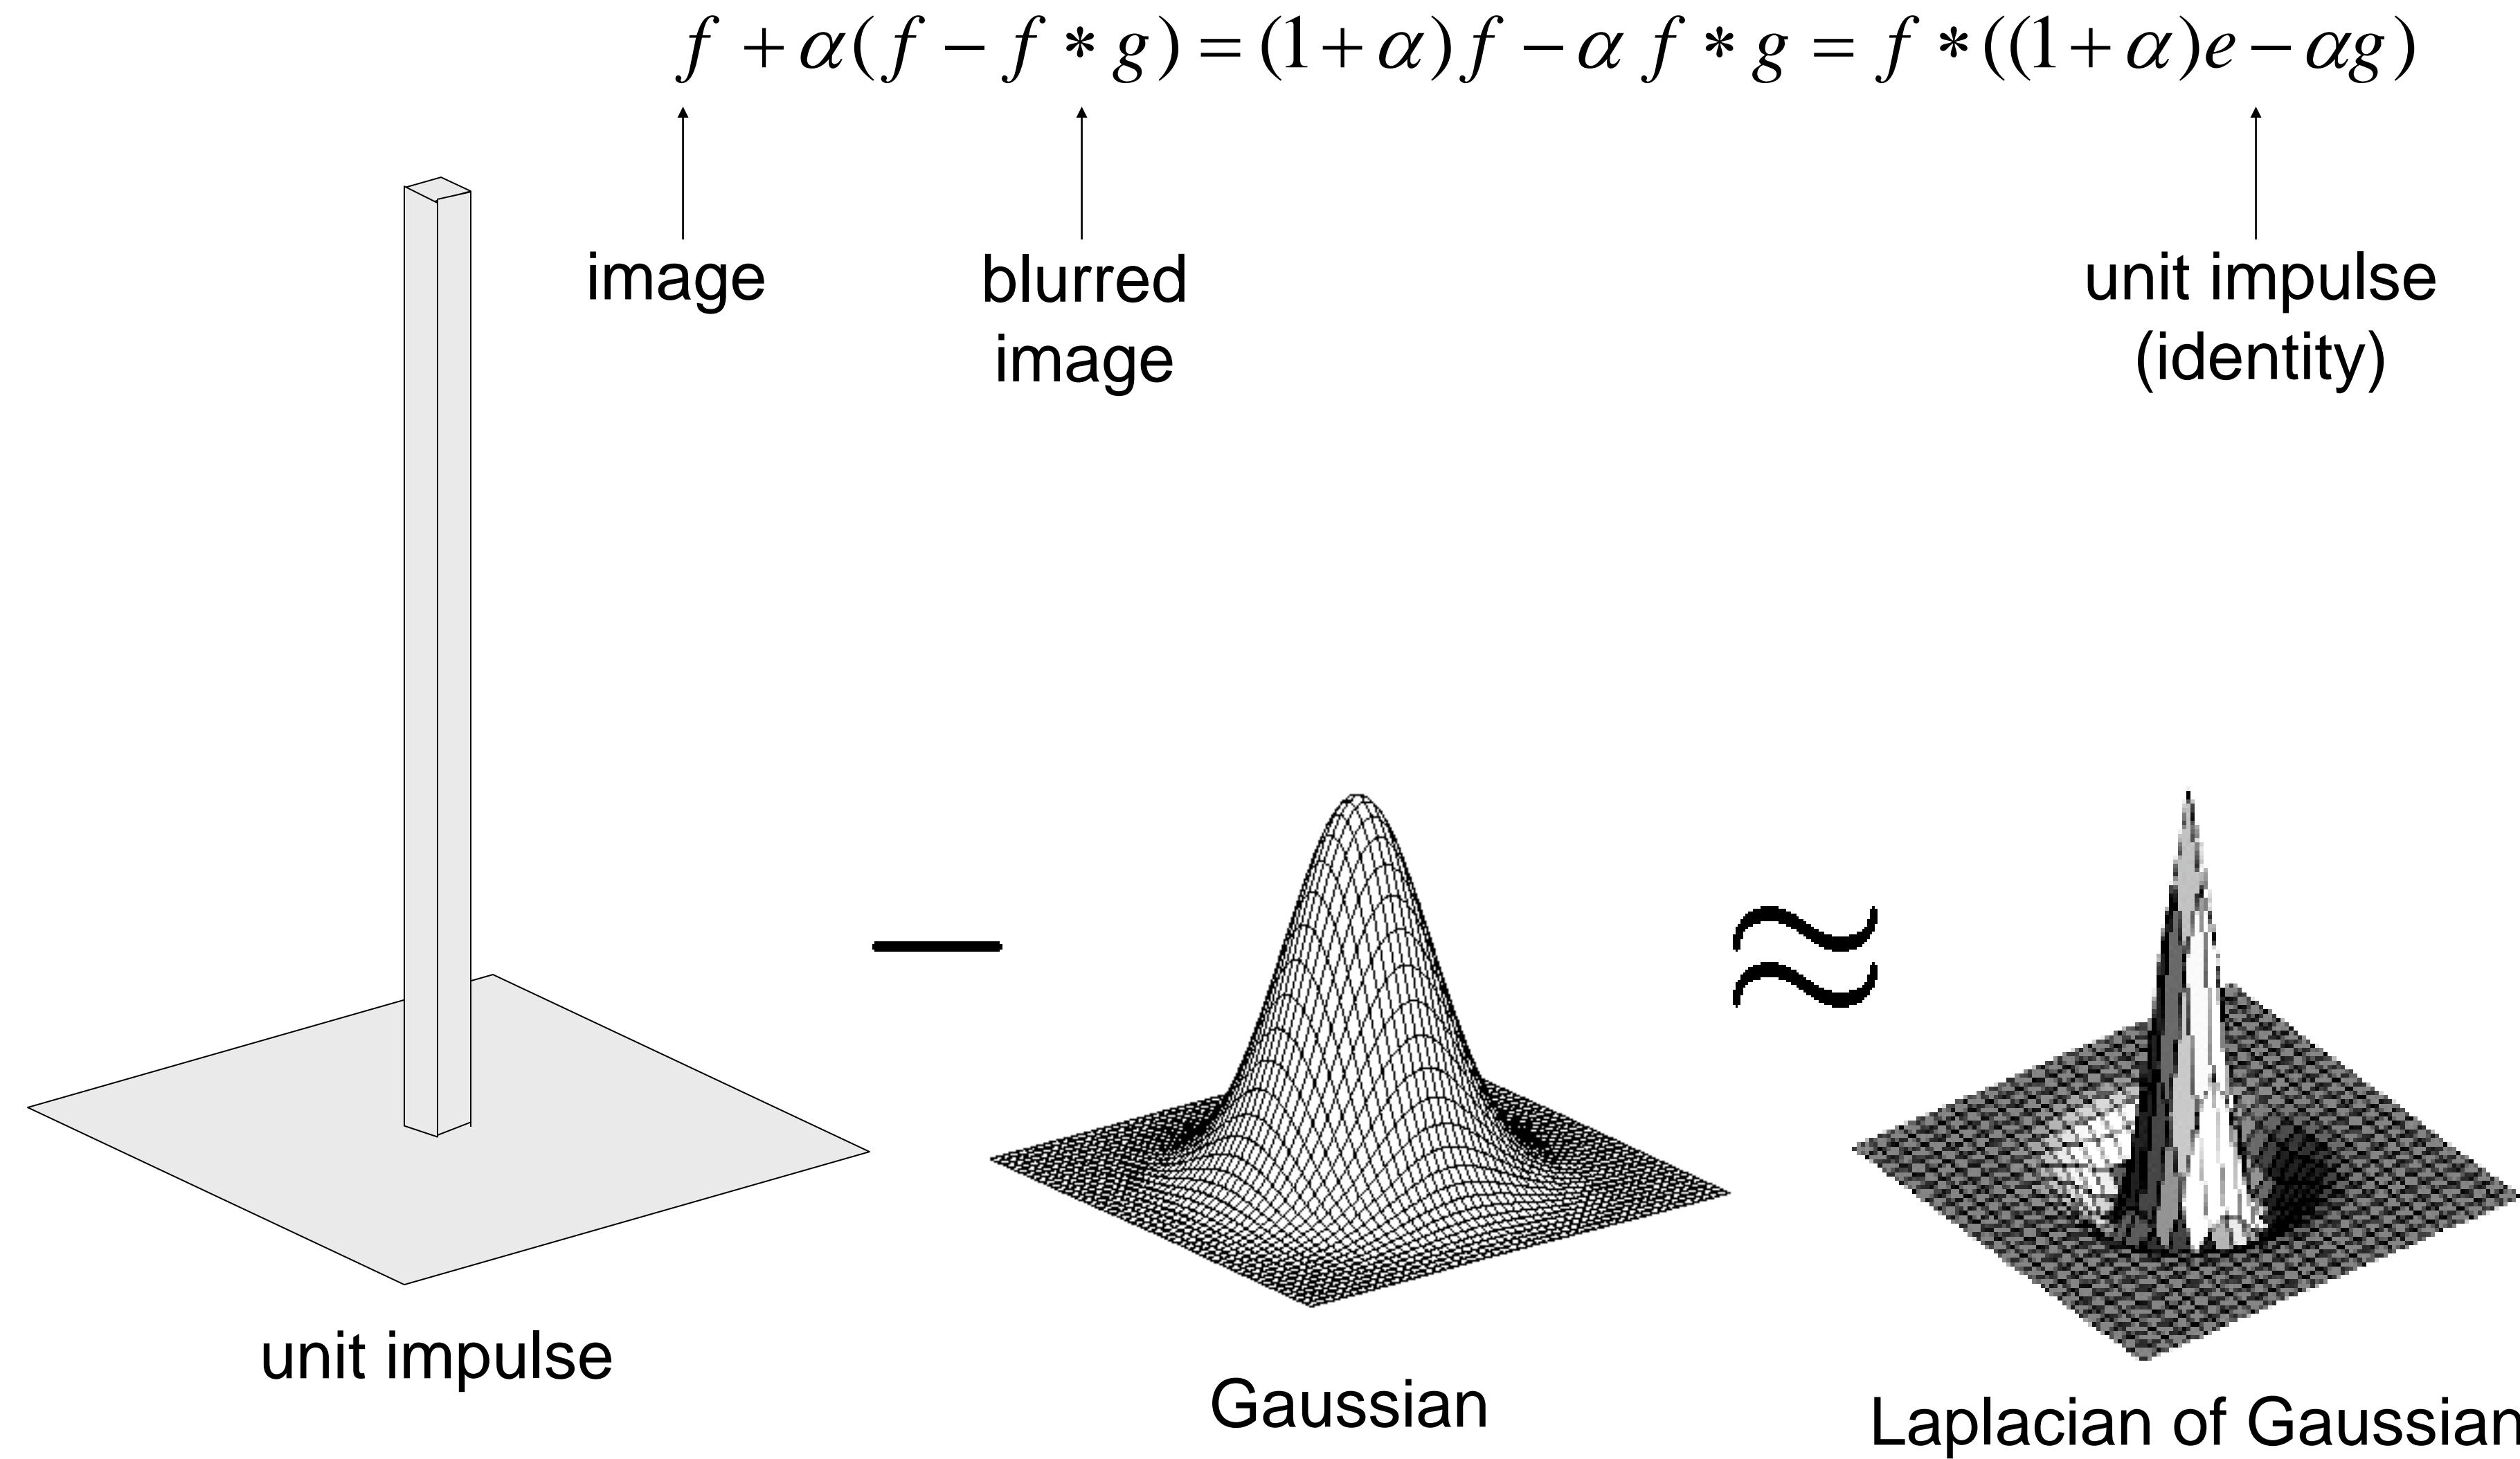
“Sharpened”  $\alpha=10$

=



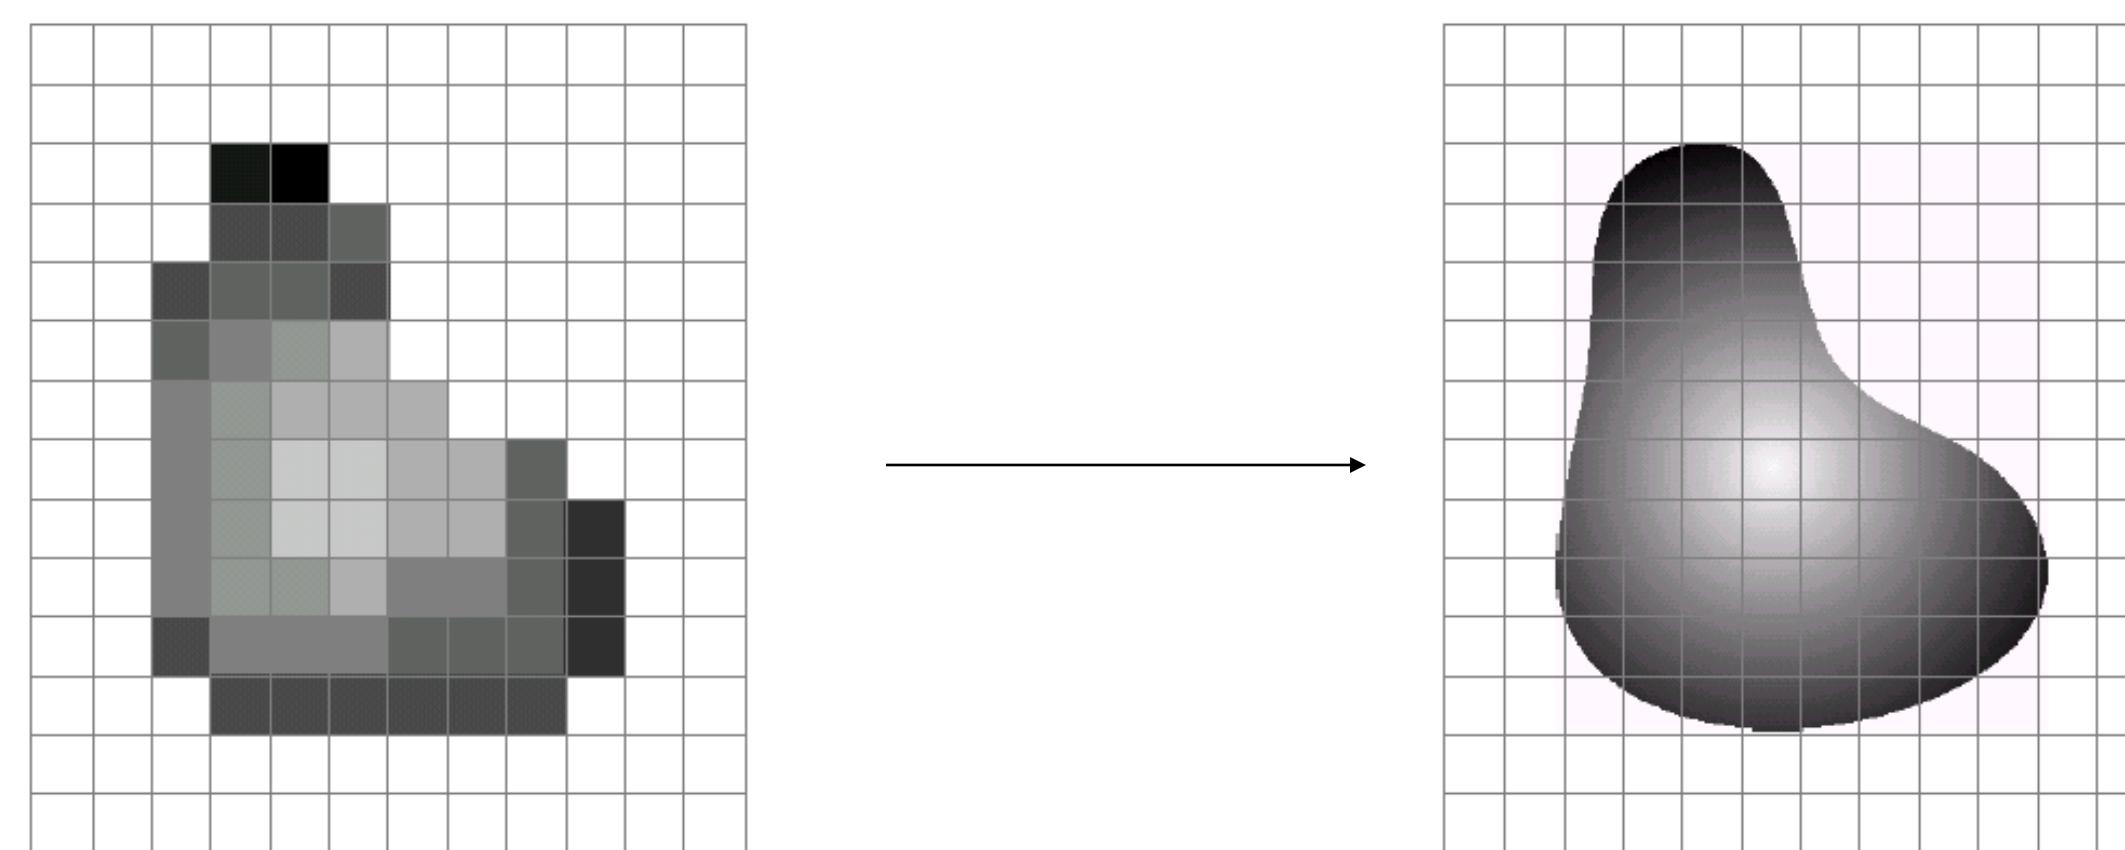
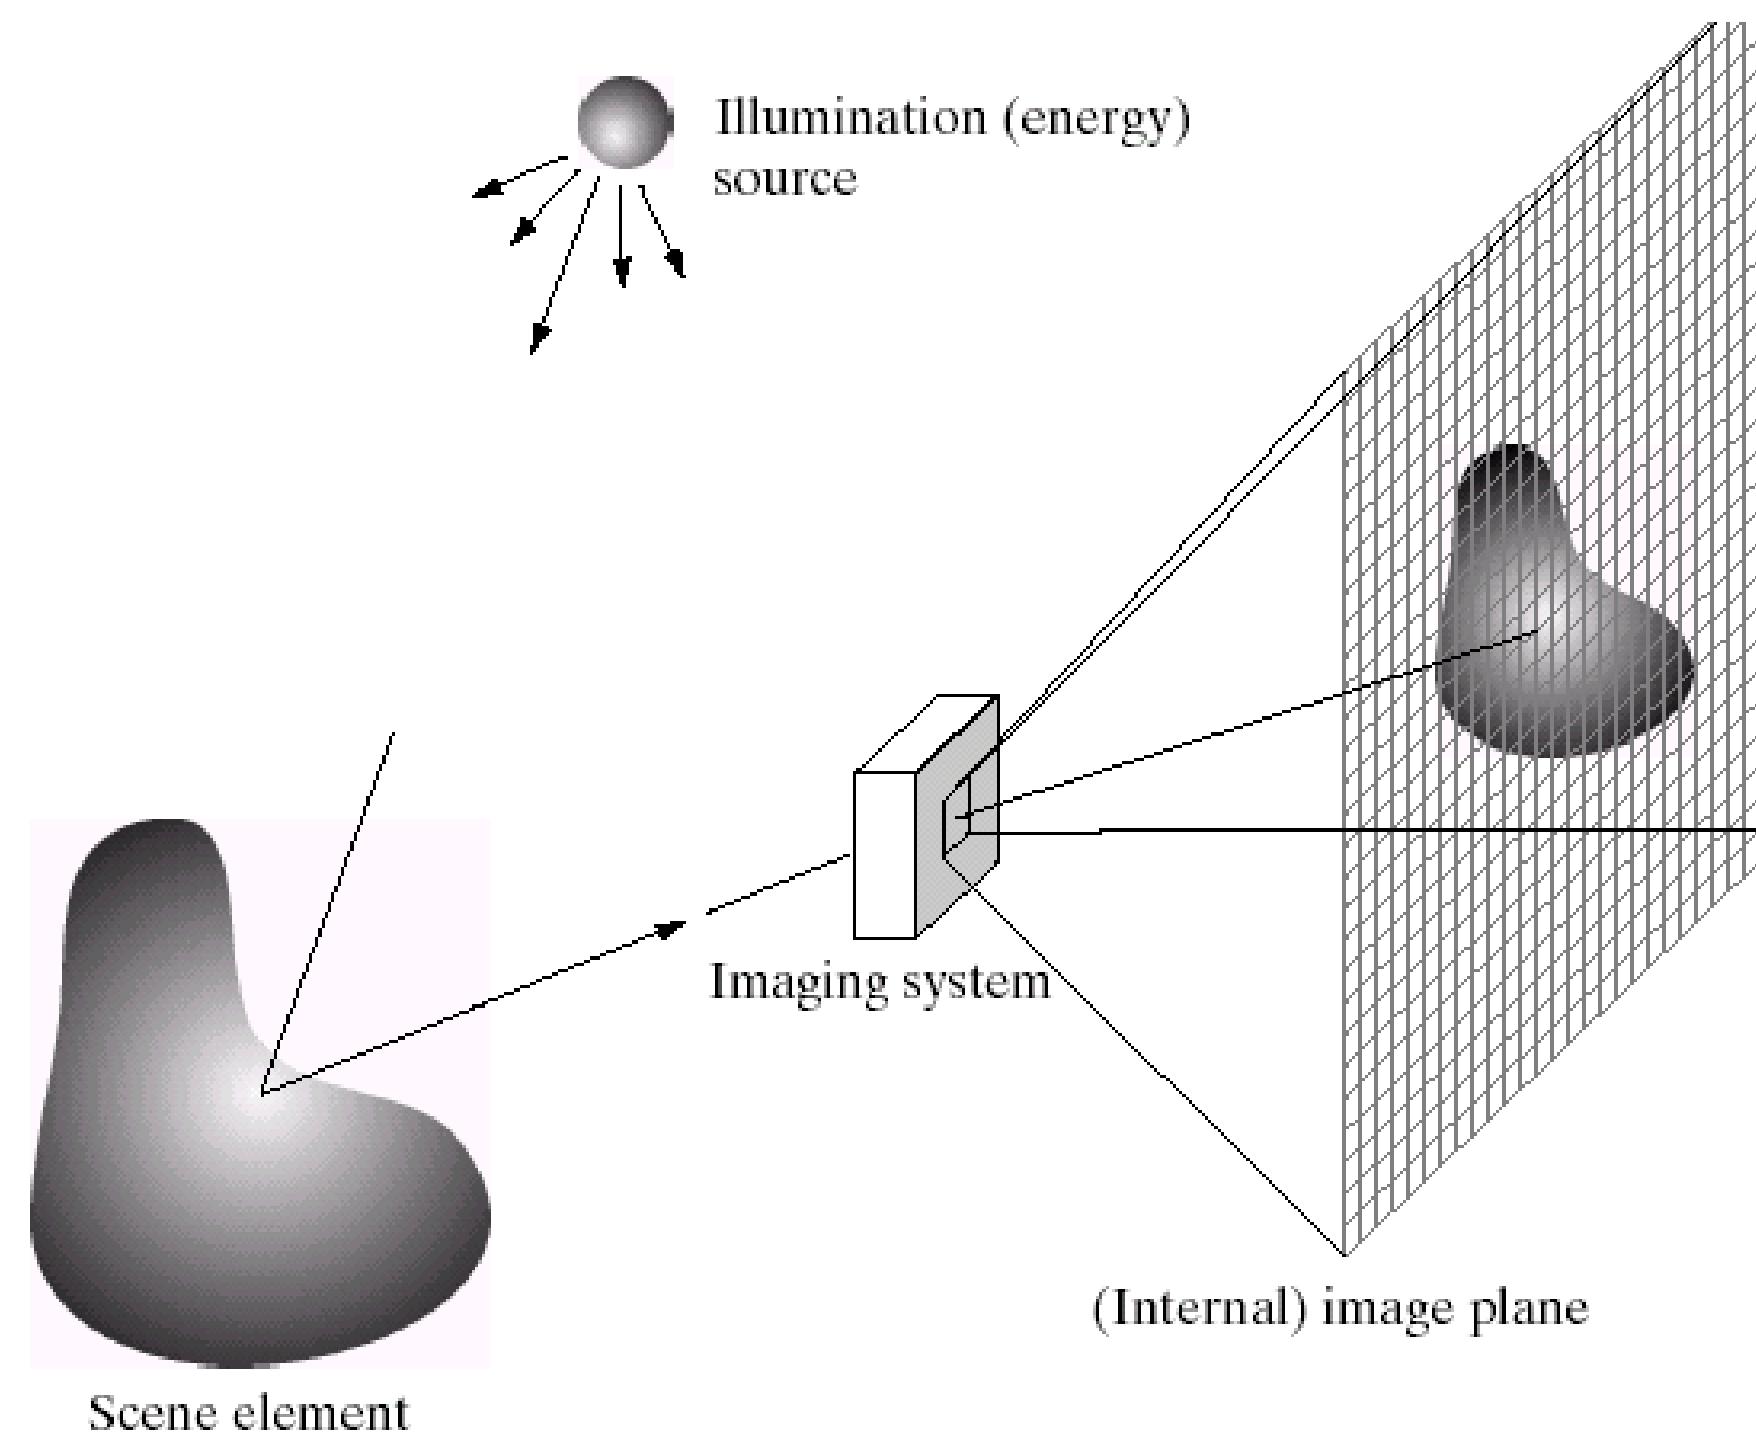
# Unsharp mask filter

---



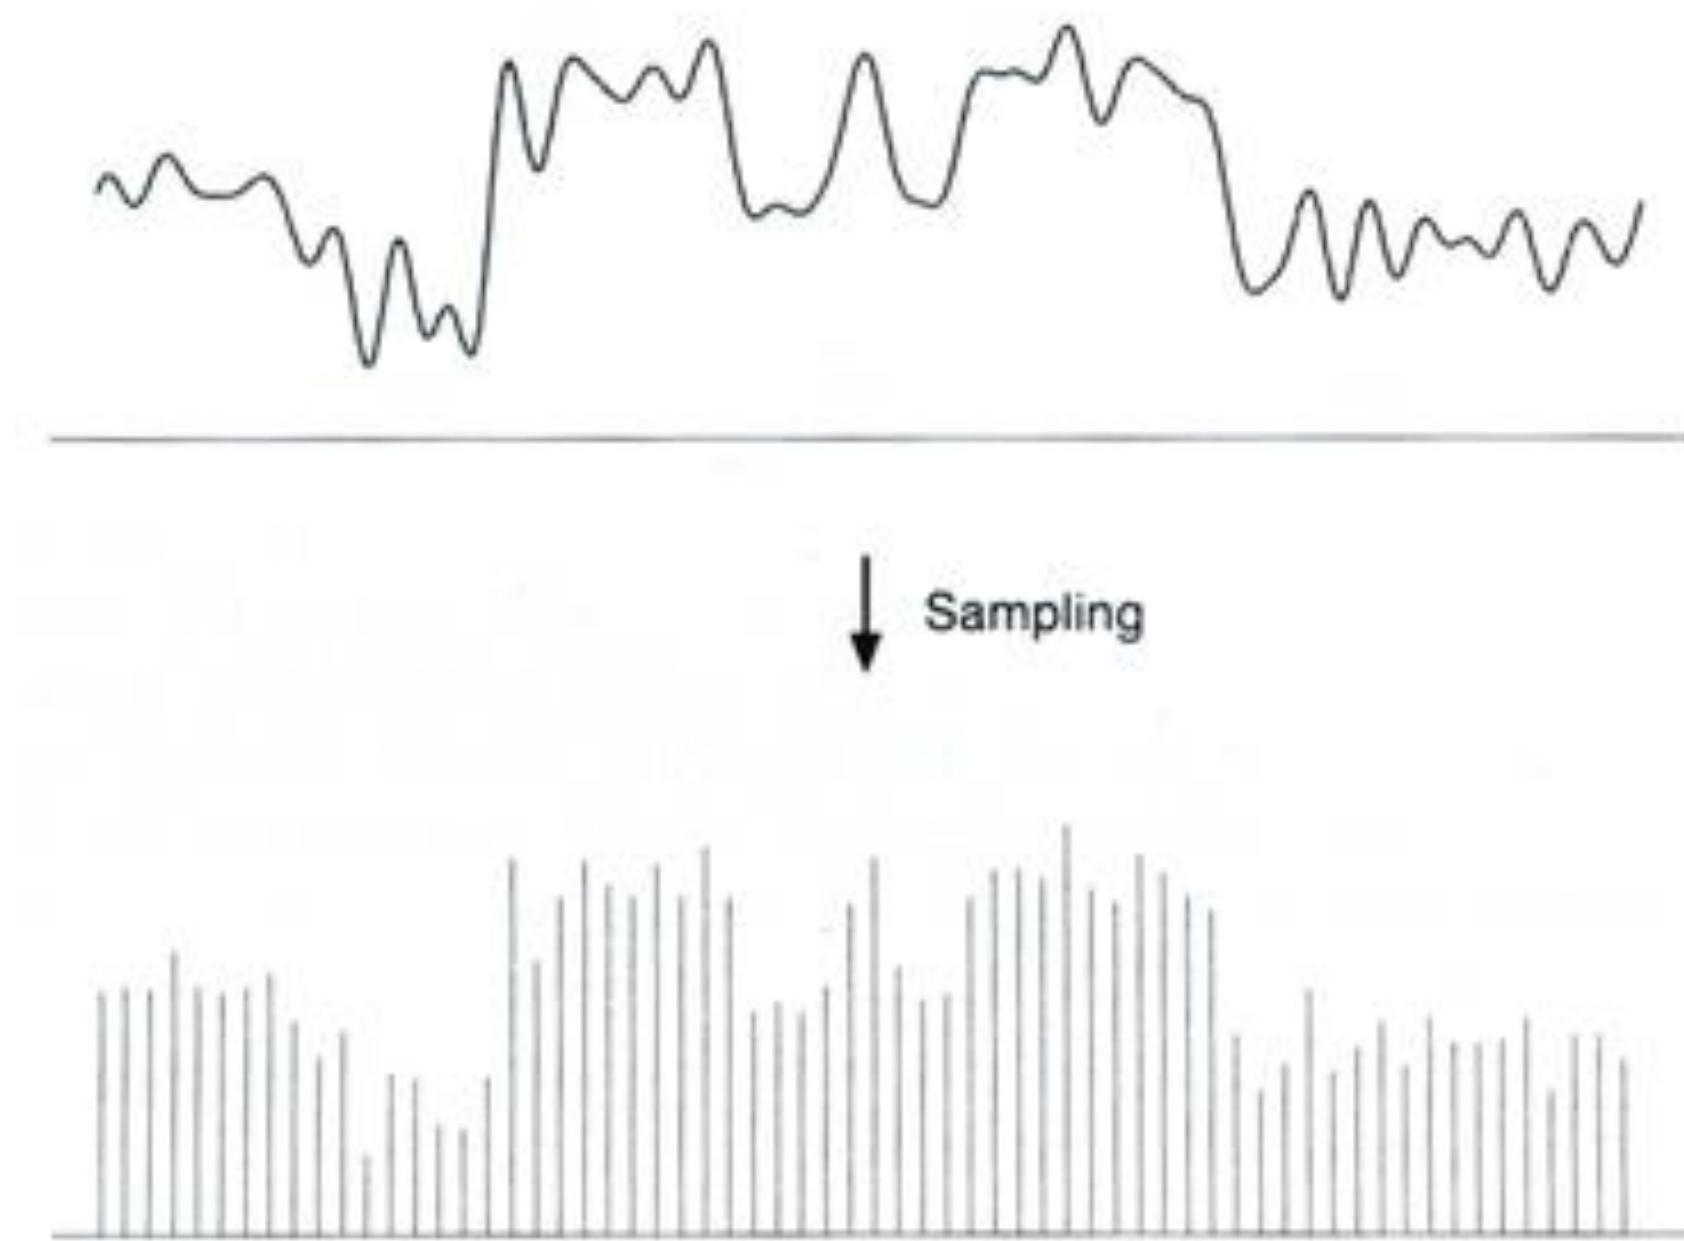
# Sampling and Reconstruction

---



# Sampled representations

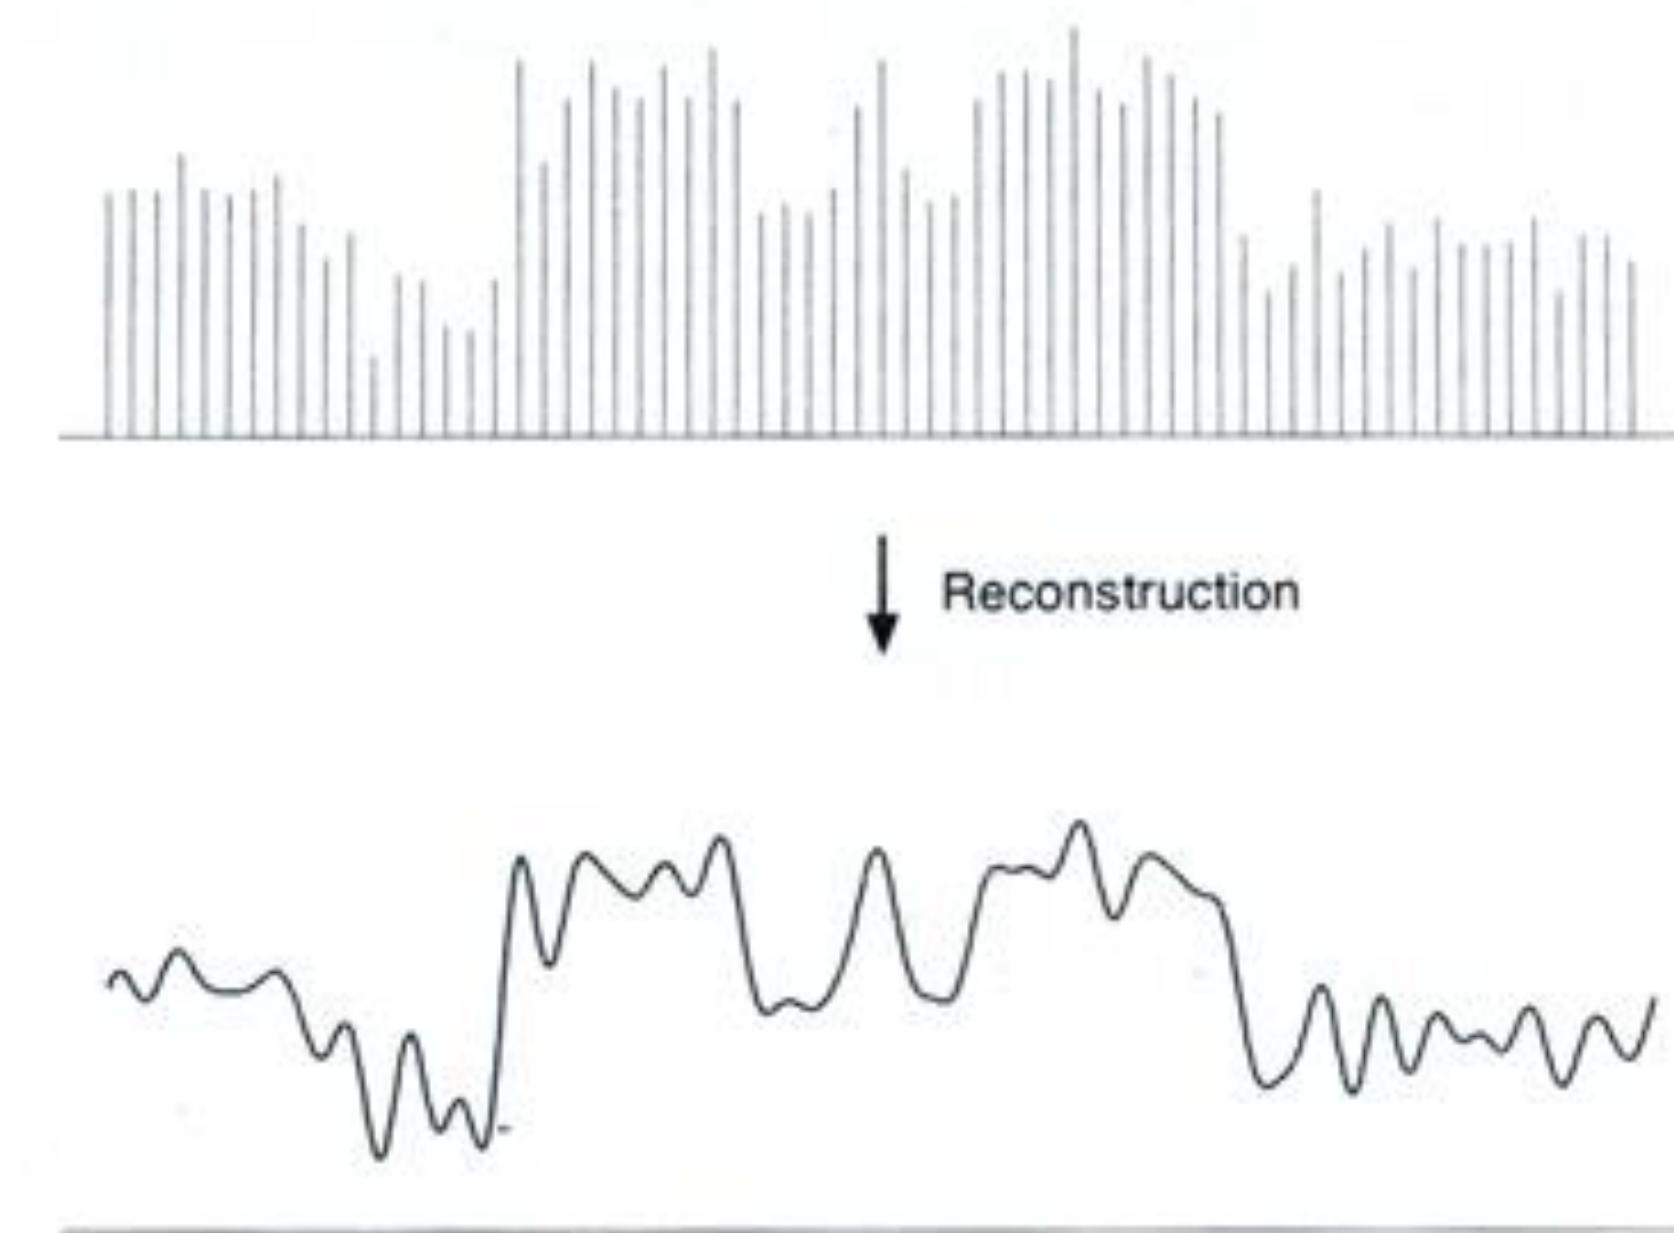
- How to store and compute with continuous functions?
- Common scheme for representation: samples
  - write down the function's values at many points



[FvDFH fig.14.14b / Wolberg]

# Reconstruction

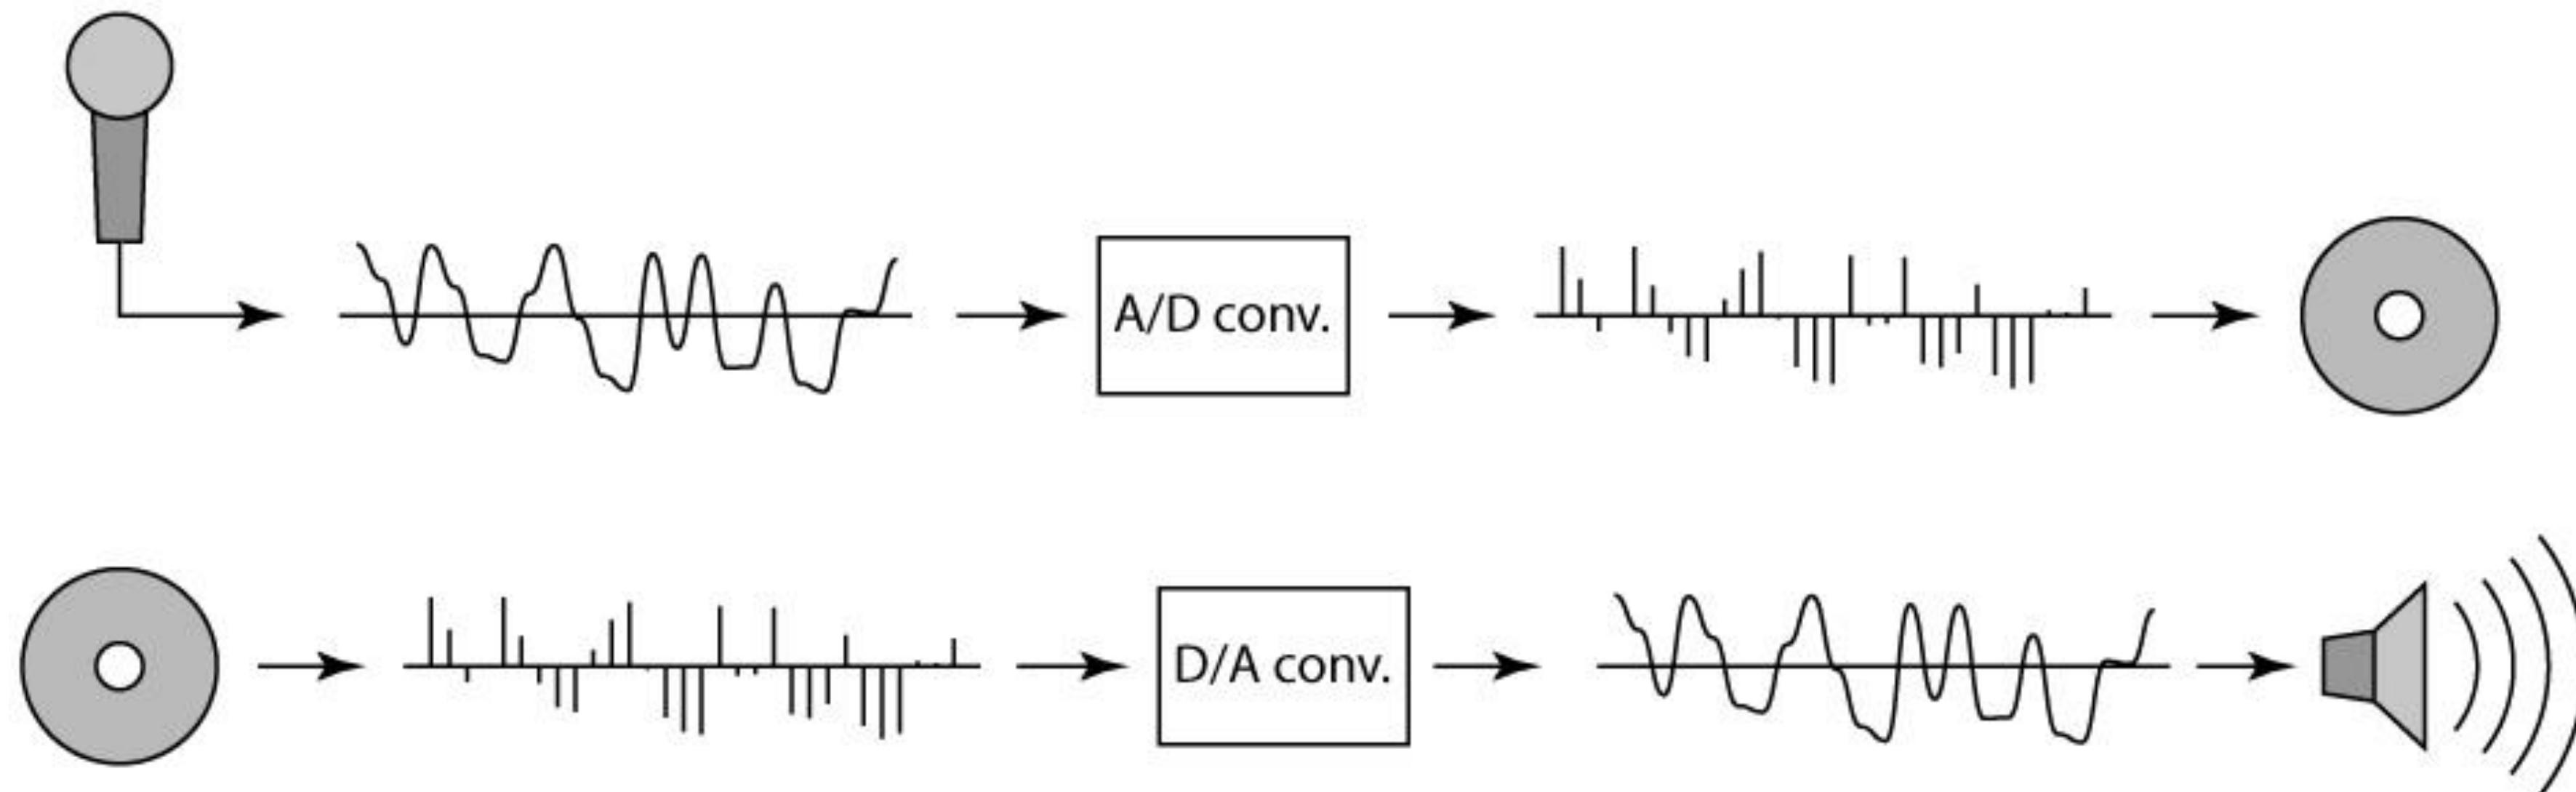
- Making samples back into a continuous function
  - for output (need realizable method)
  - for analysis or processing (need mathematical method)
  - amounts to “guessing” what the function did in between

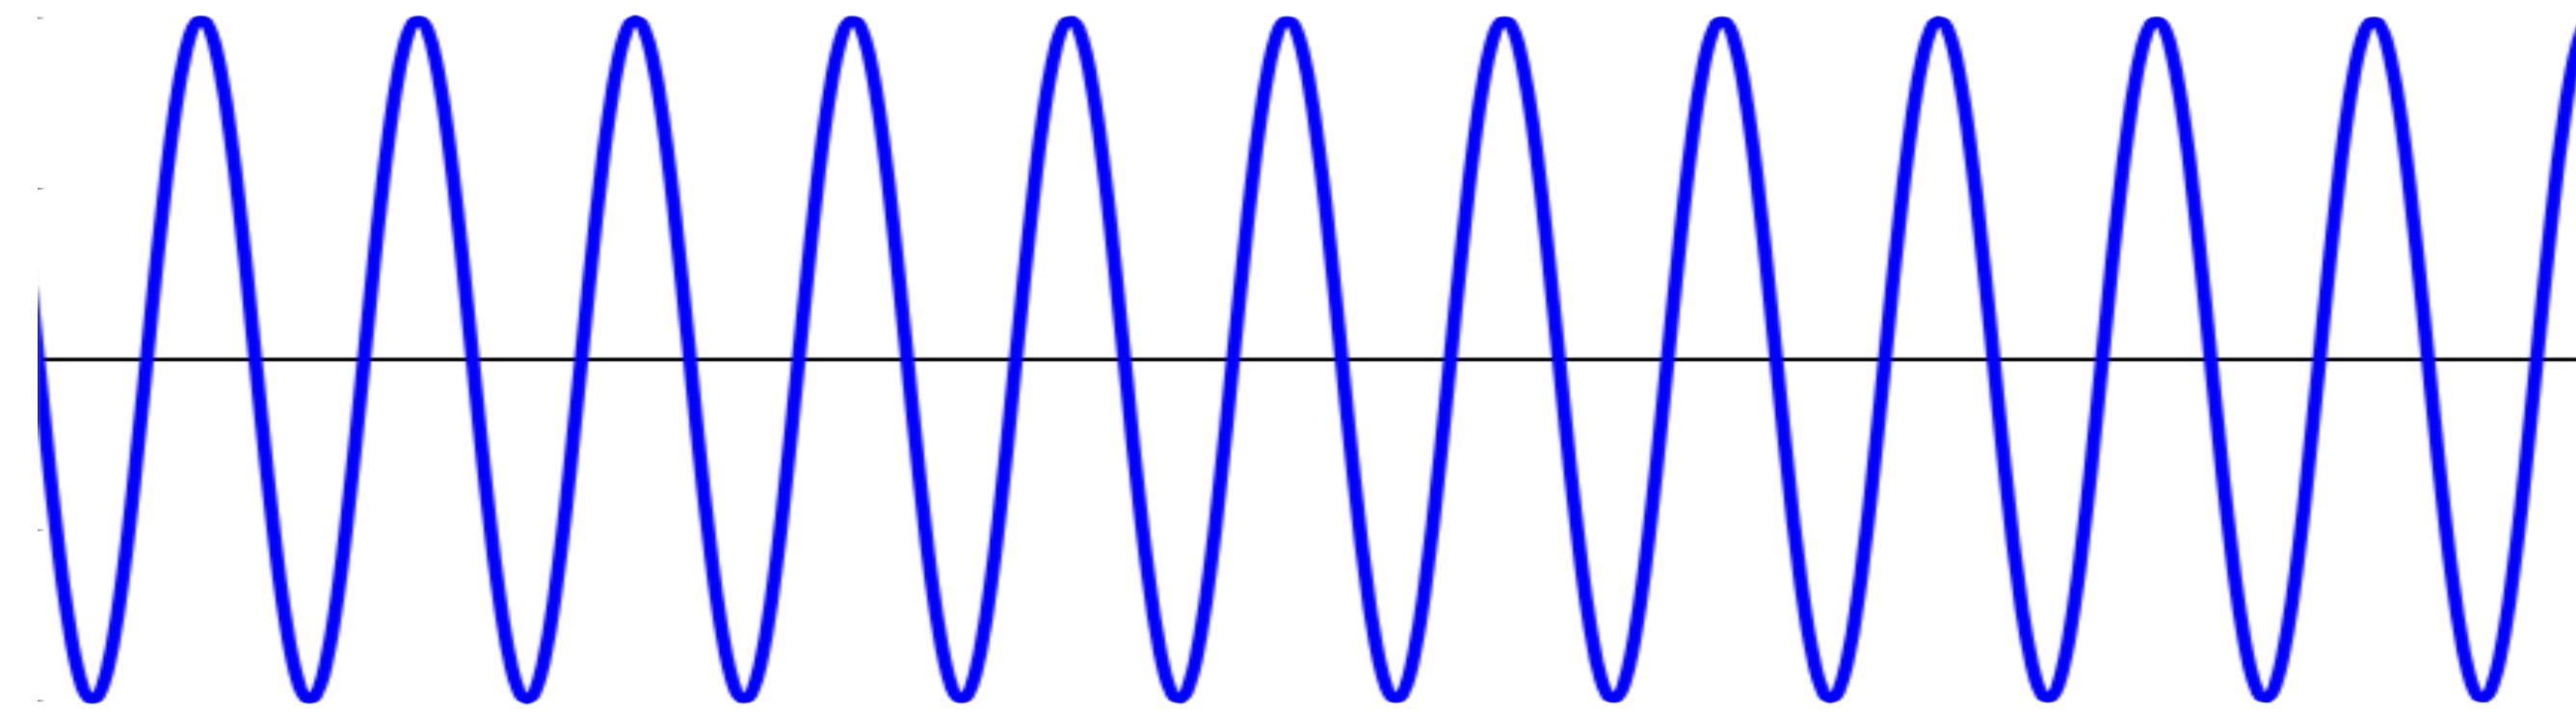


[FvDFH fig.14.14b / Wolberg]

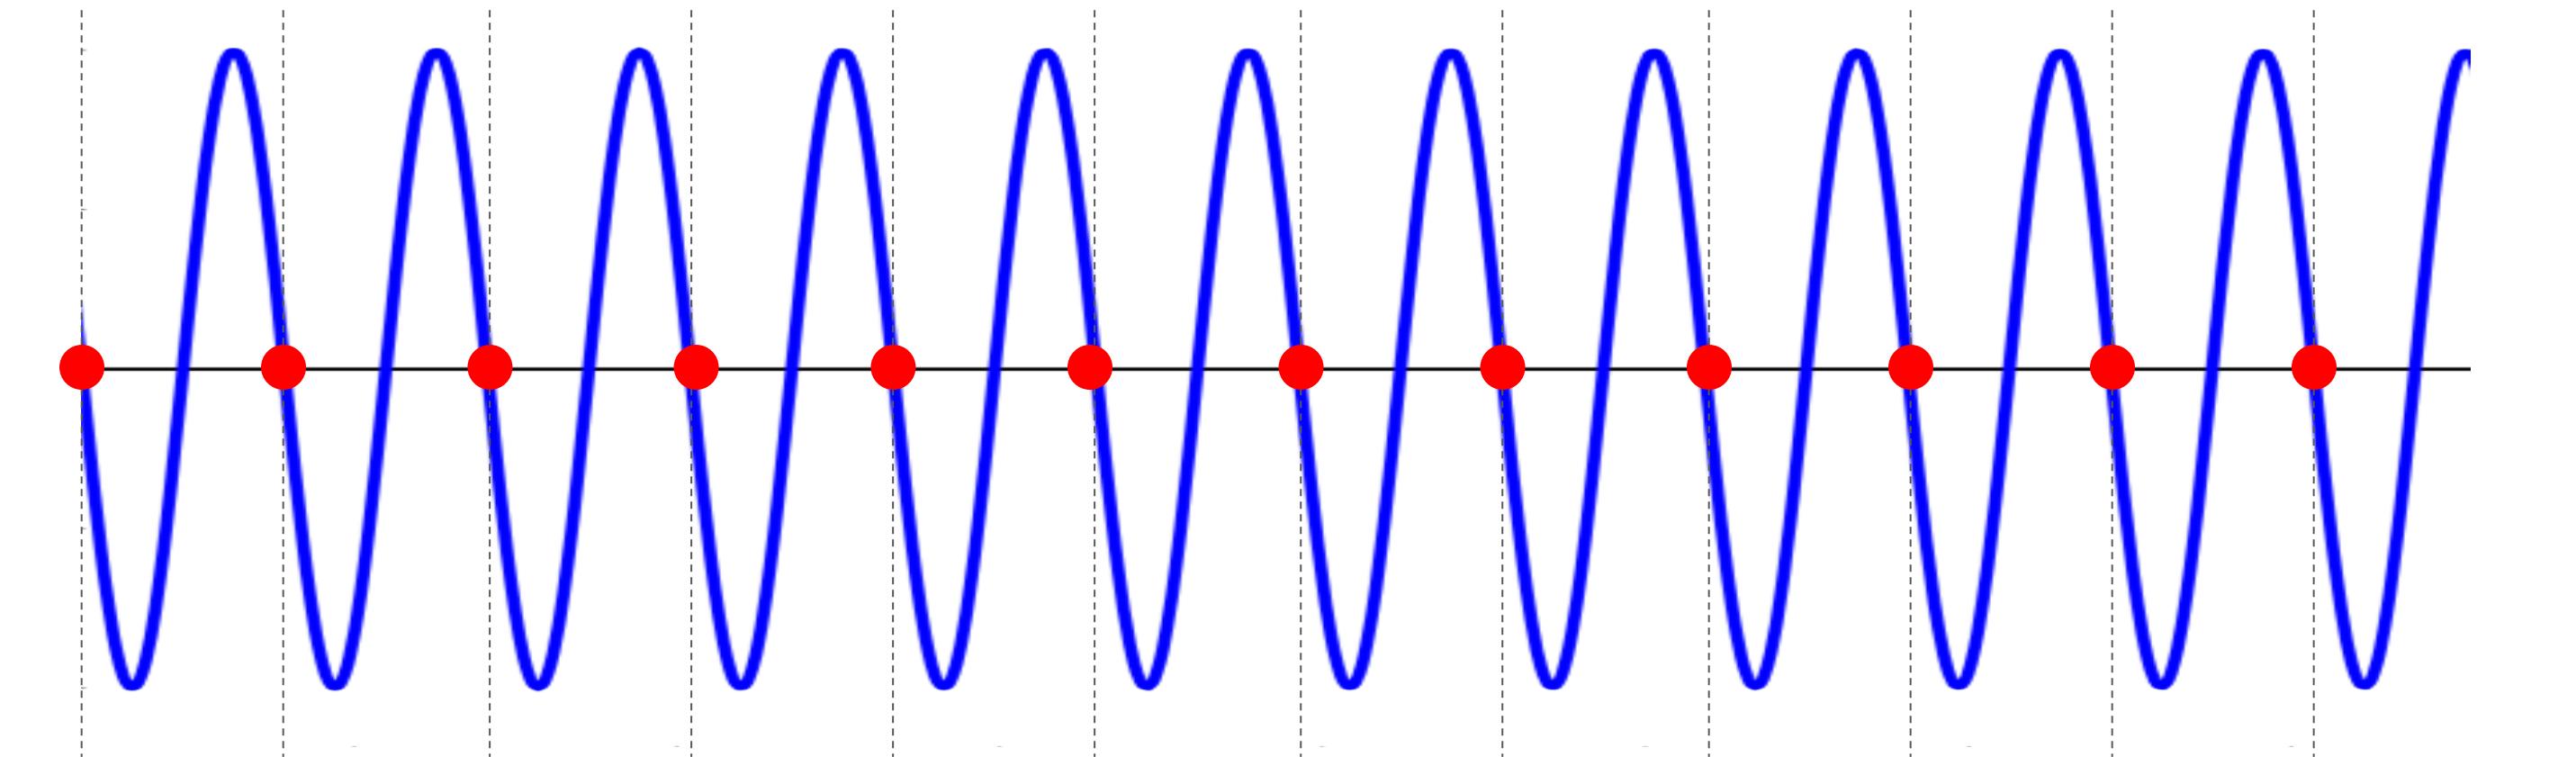
# Sampling in digital audio

- Recording: sound to analog to samples to disc
- Playback: disc to samples to analog to sound again
  - how can we be sure we are filling in the gaps correctly?



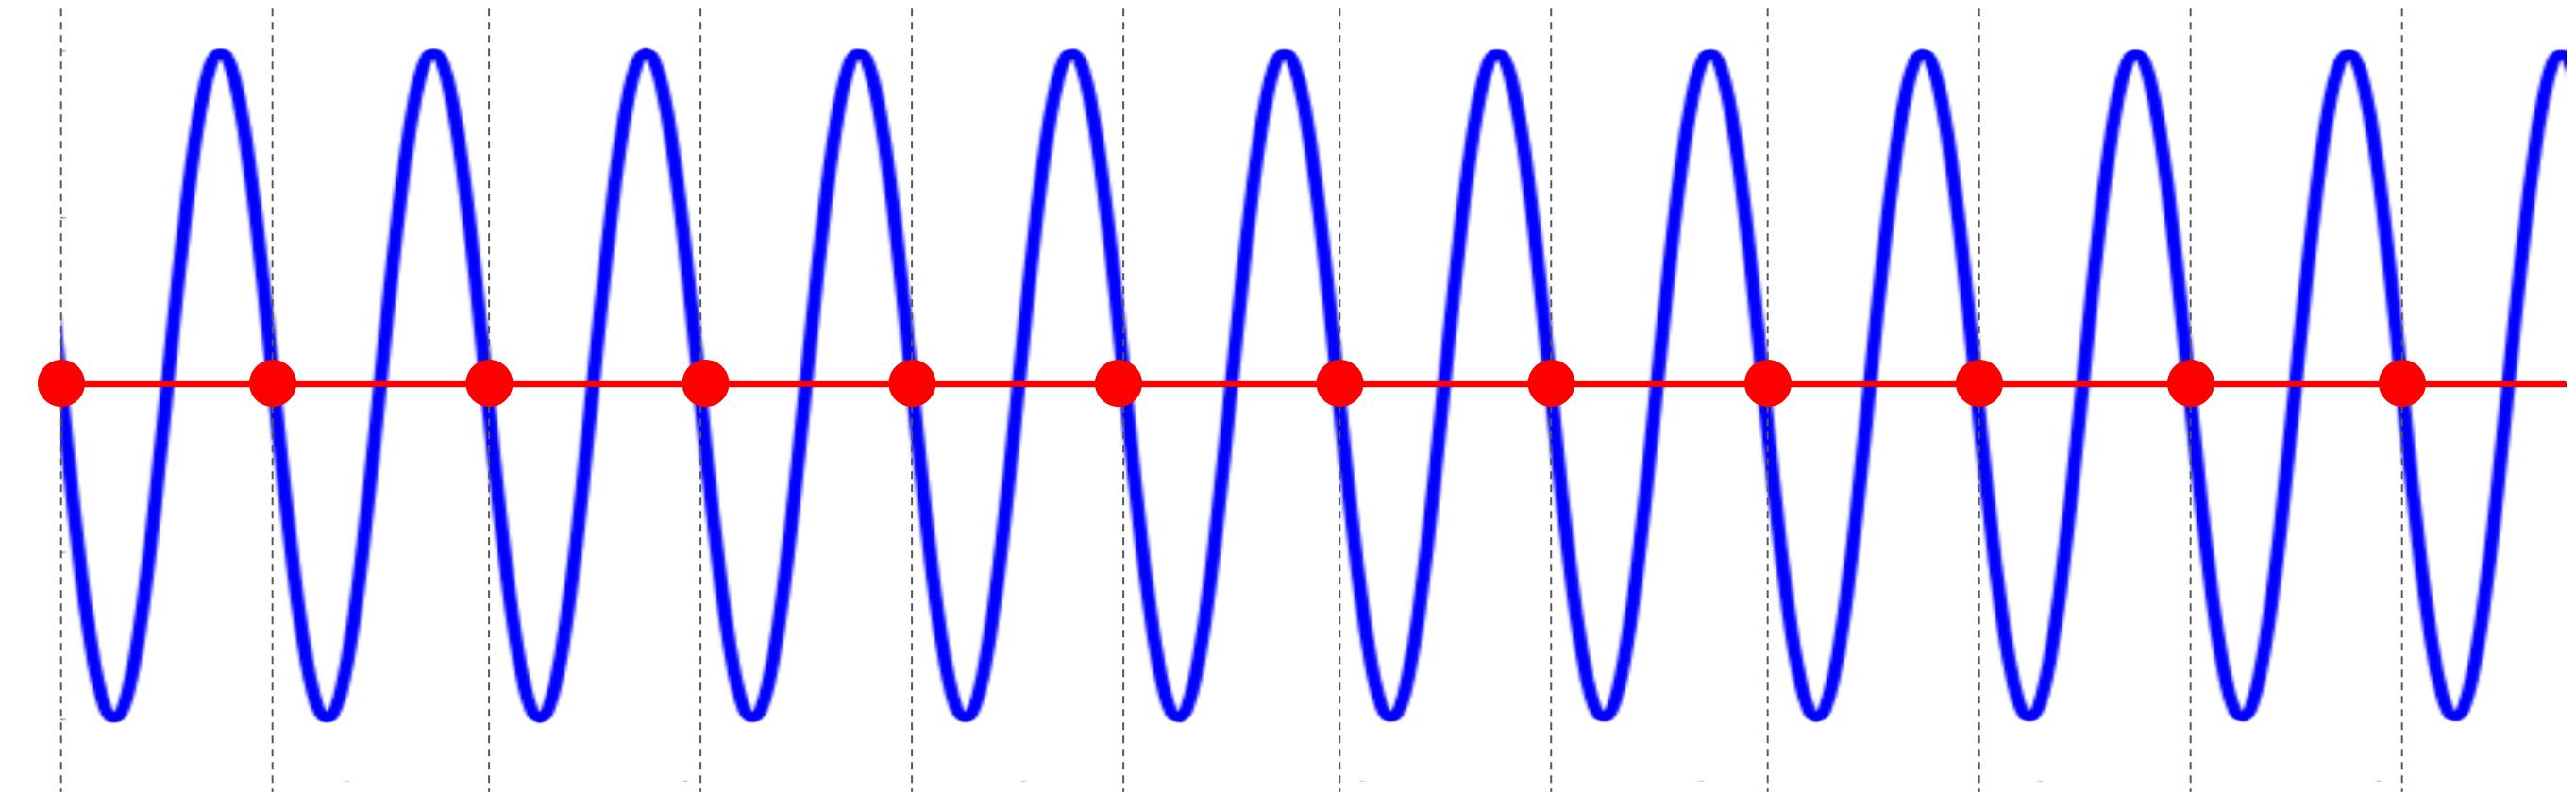


$f(x)$

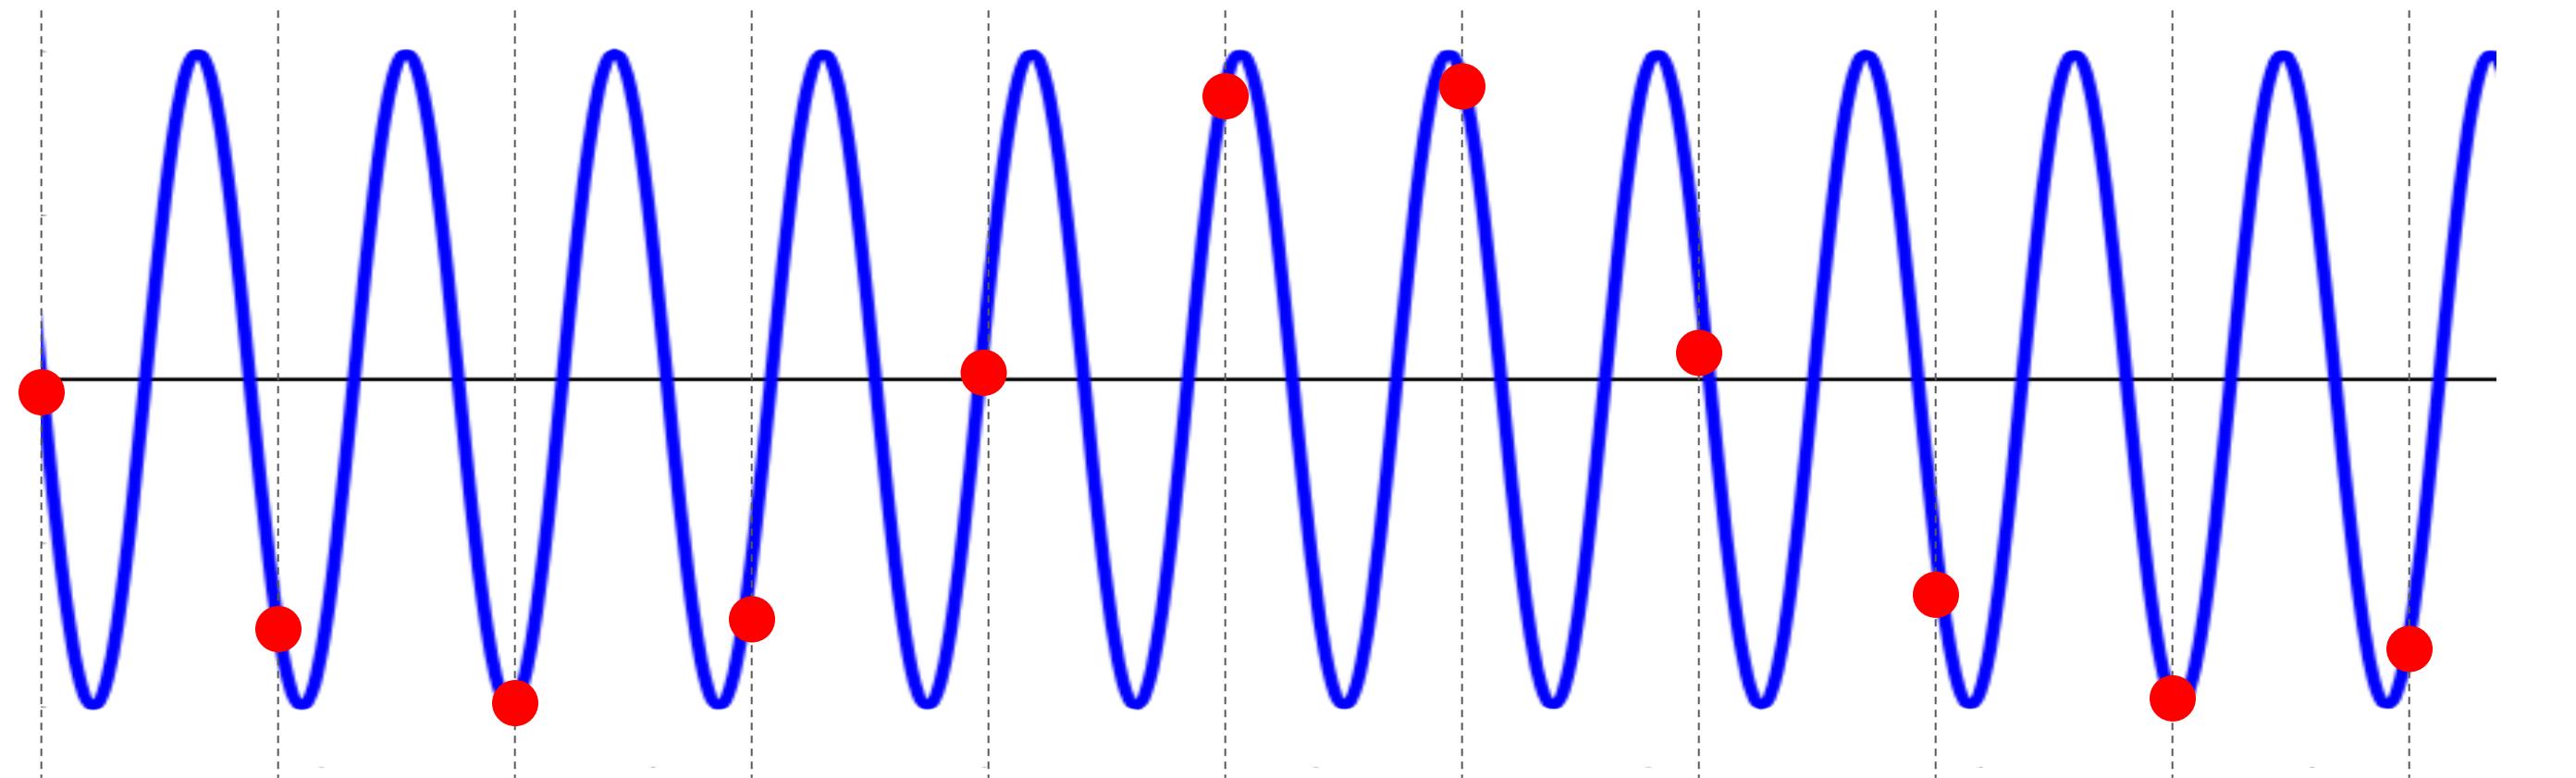


$f(x)$

$g(x)$



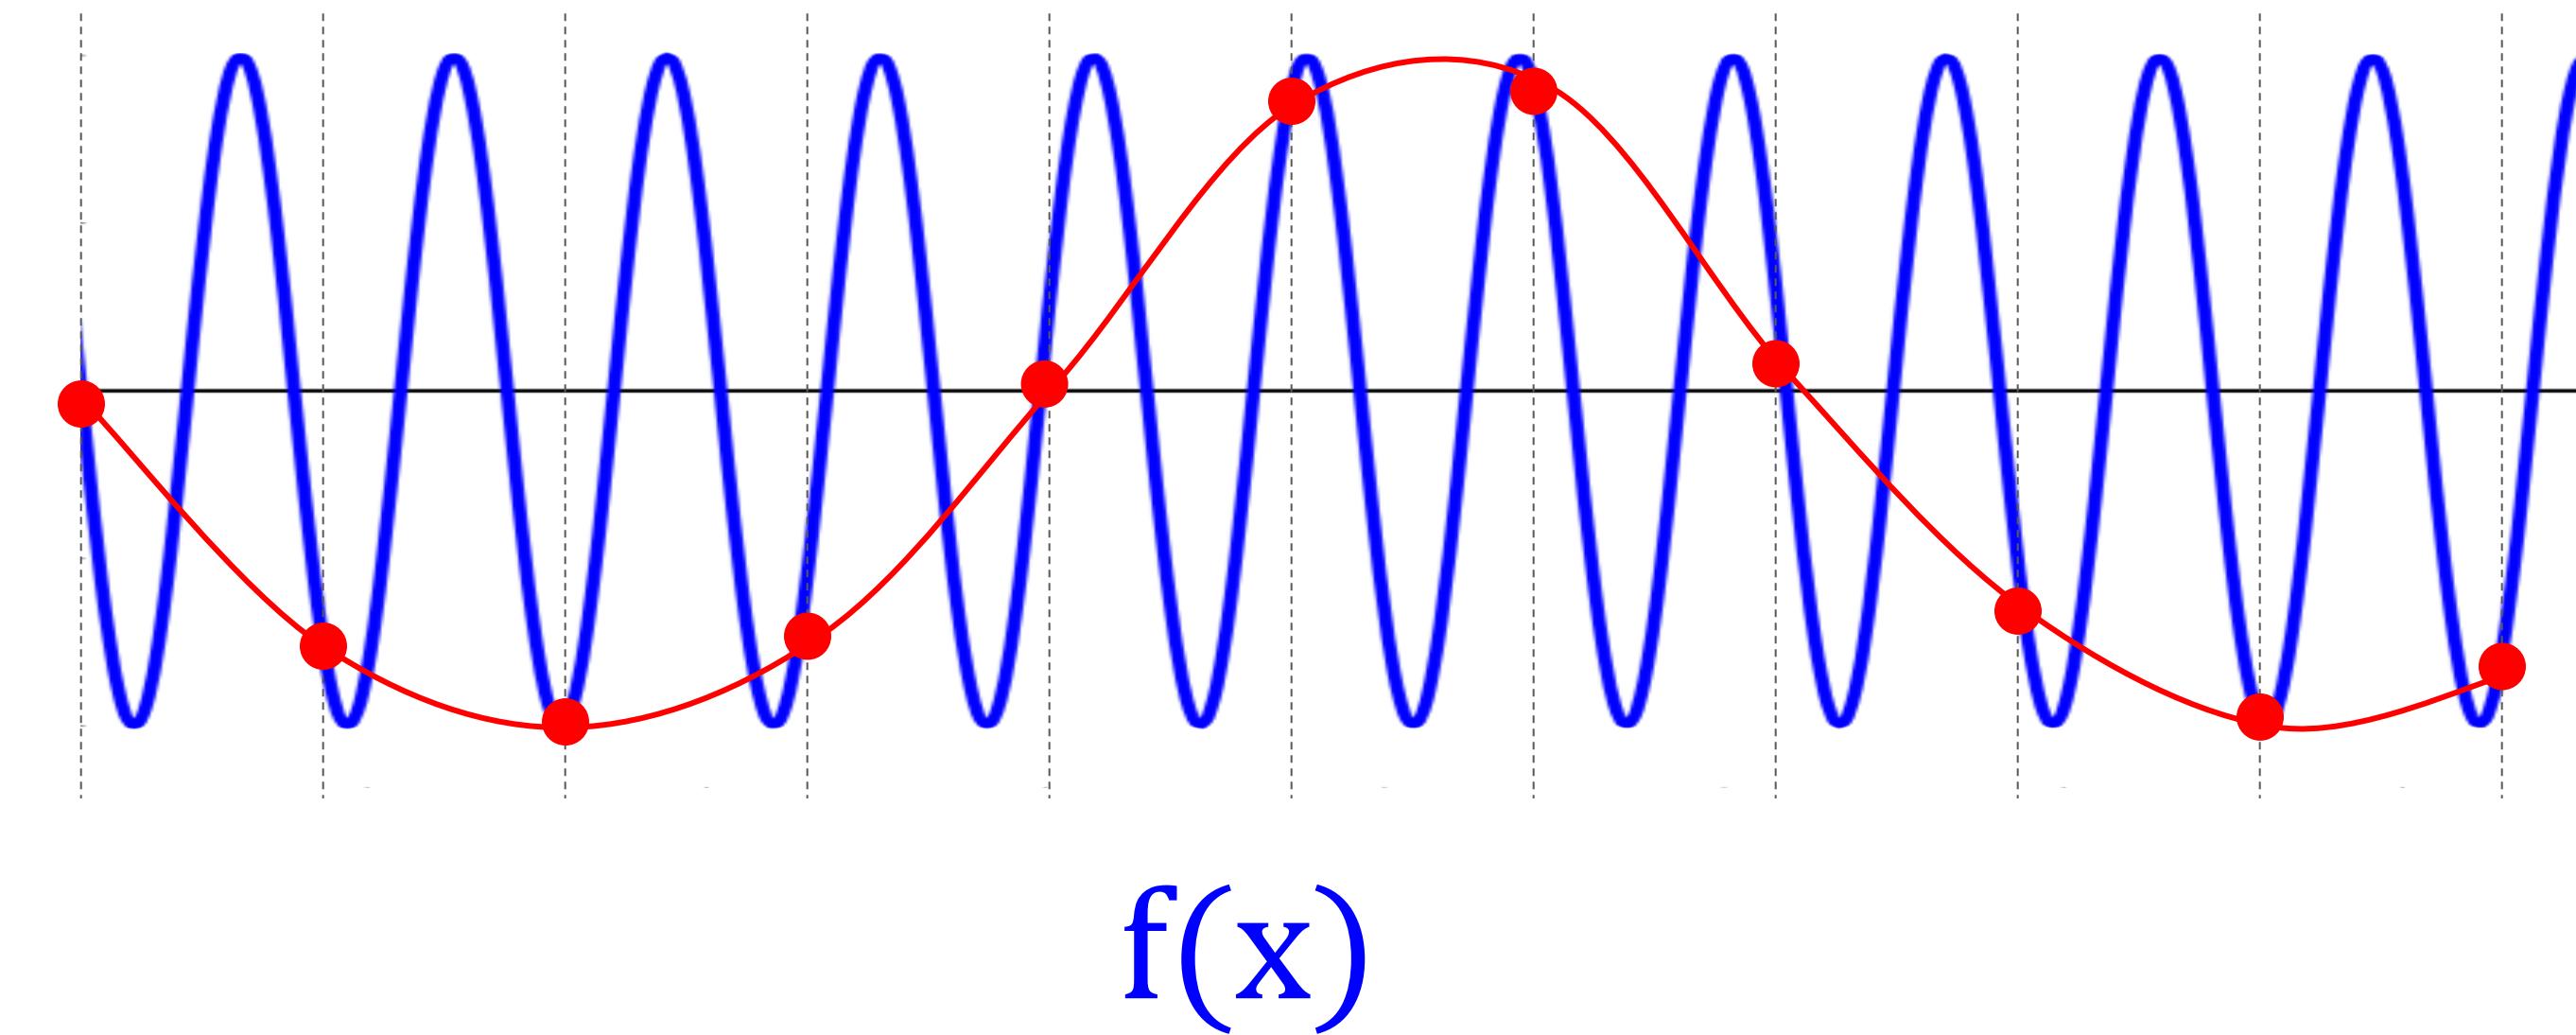
$f(x)$



$f(x)$

## Aliasing

g(x) is an “alias” of f(x)

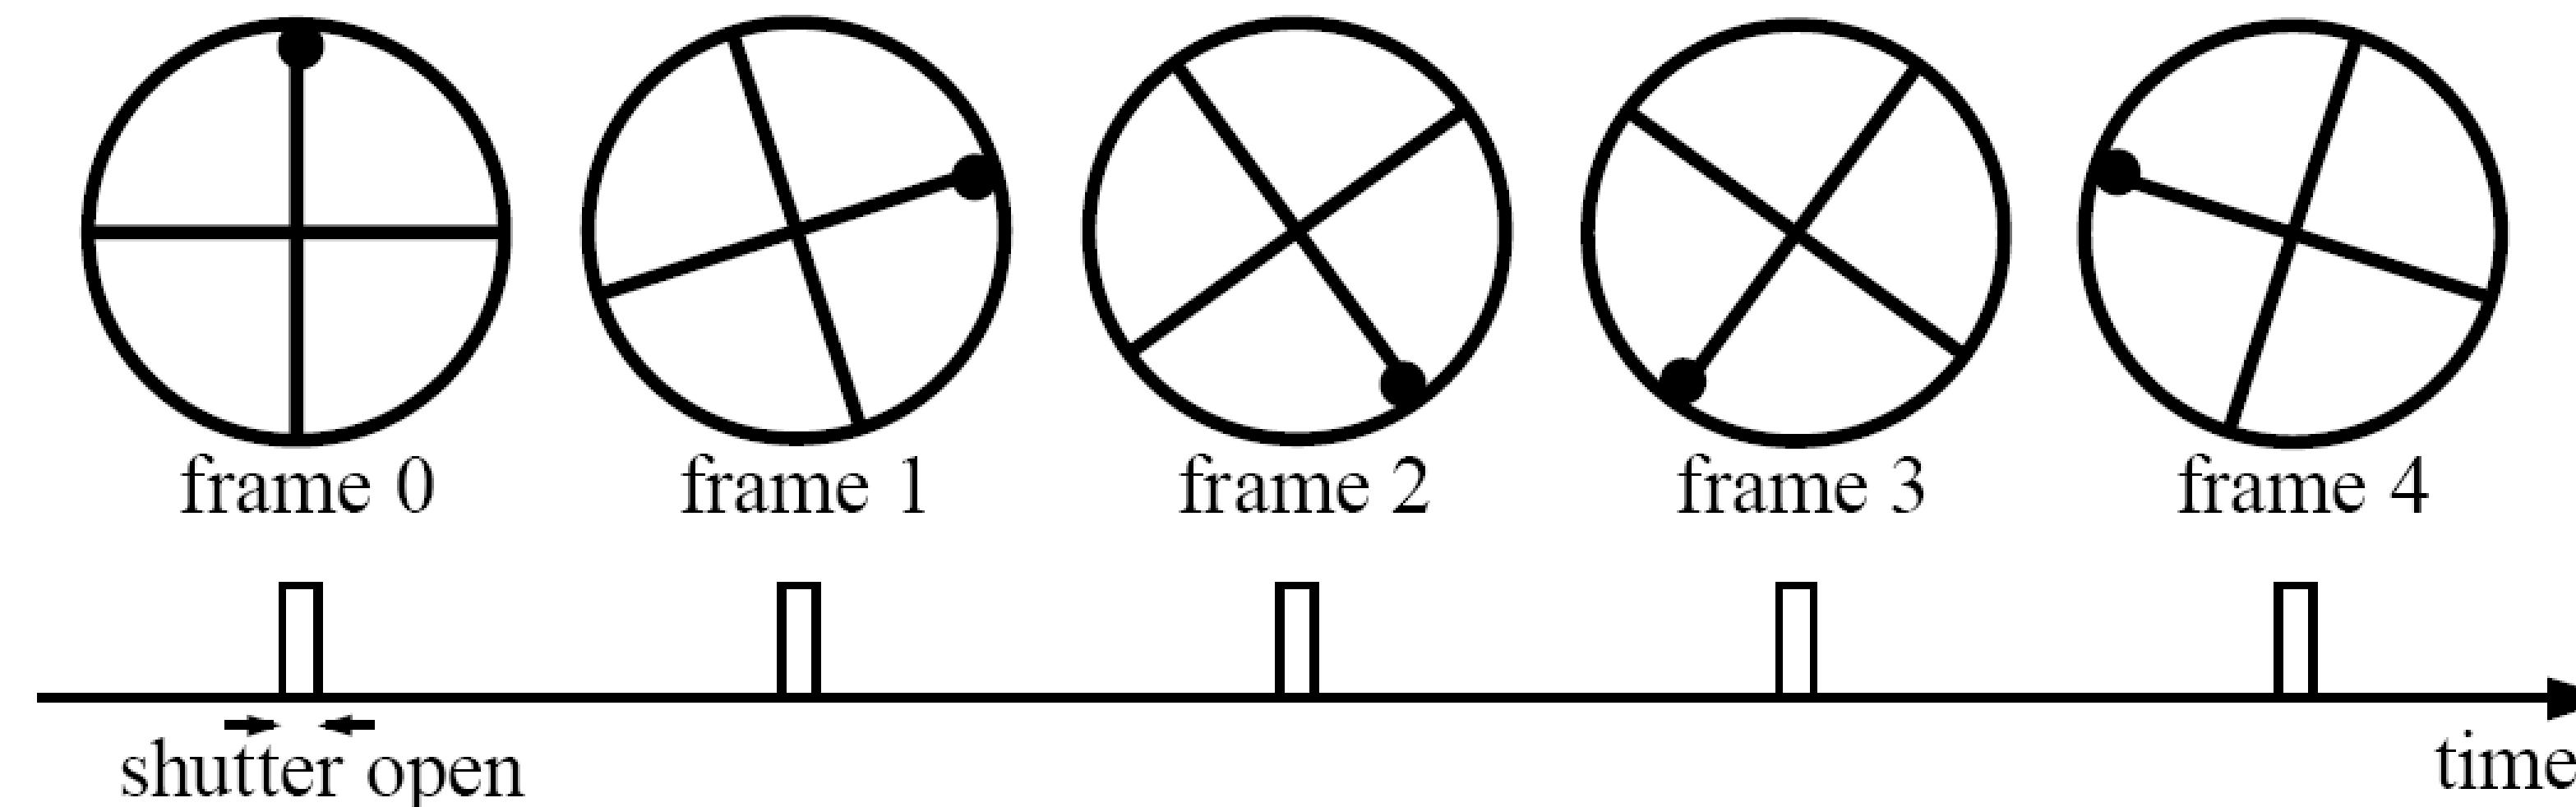


## Aliasing in video

Imagine a spoked wheel moving to the right (rotating clockwise).

Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):



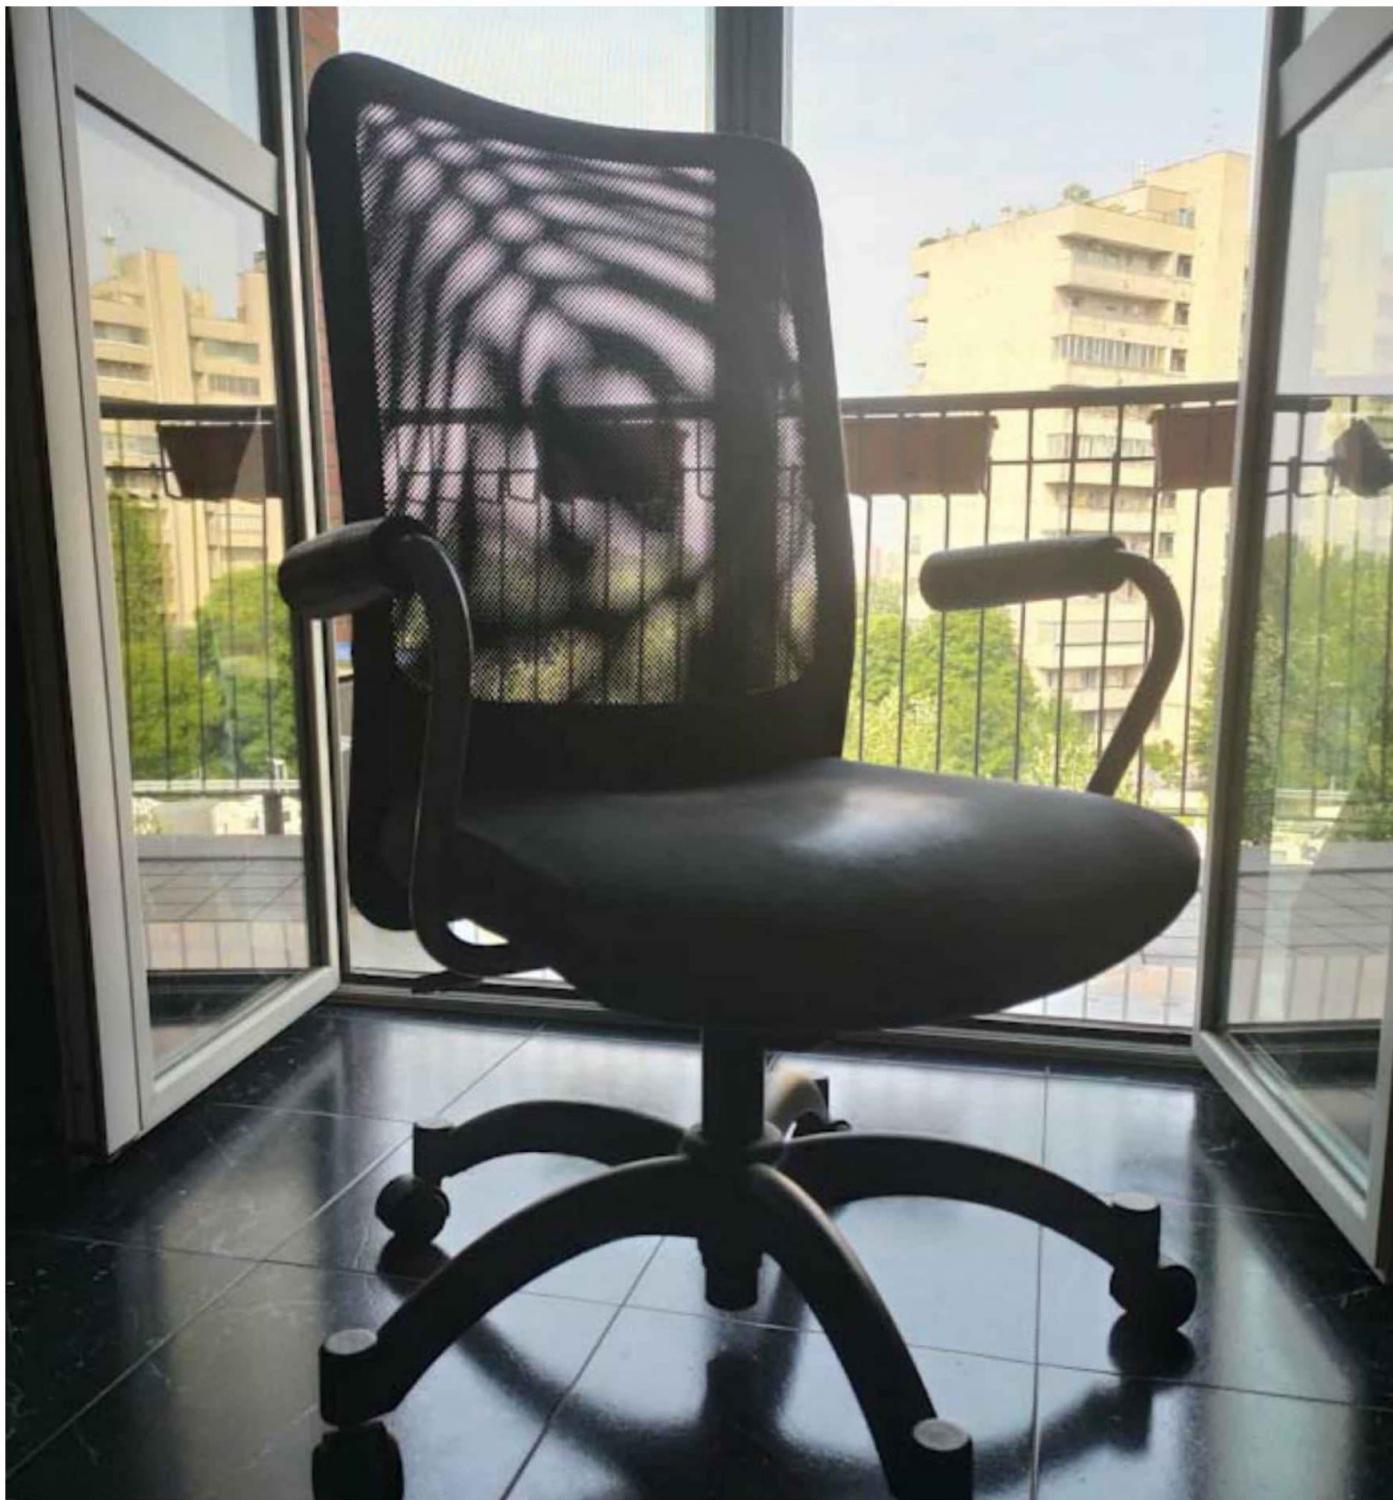
Without dot, wheel appears to be rotating slowly backwards!  
(counterclockwise)



## Aliasing in images

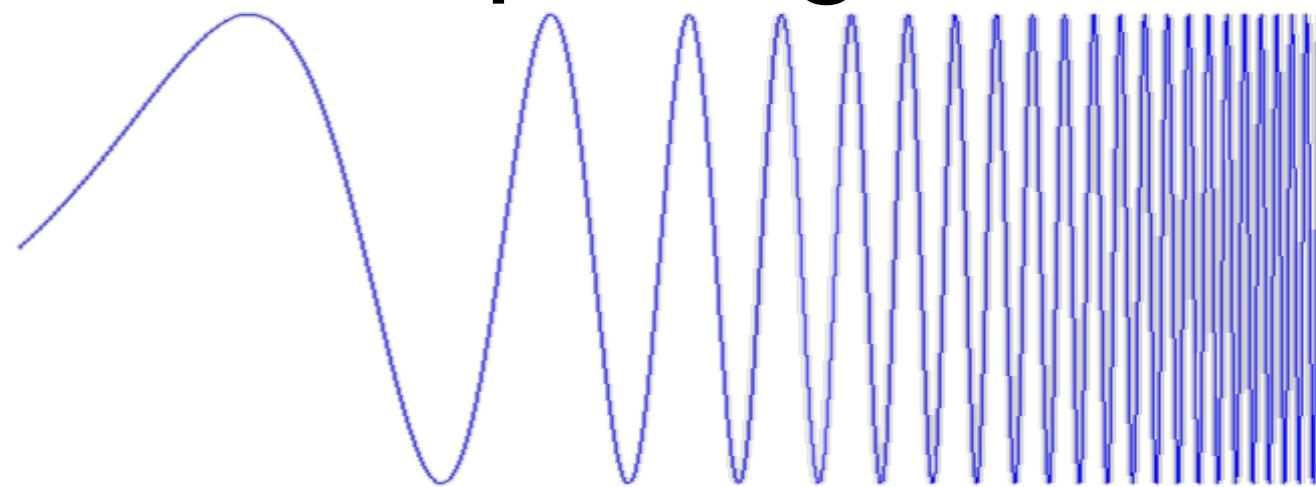


## Aliasing in real images

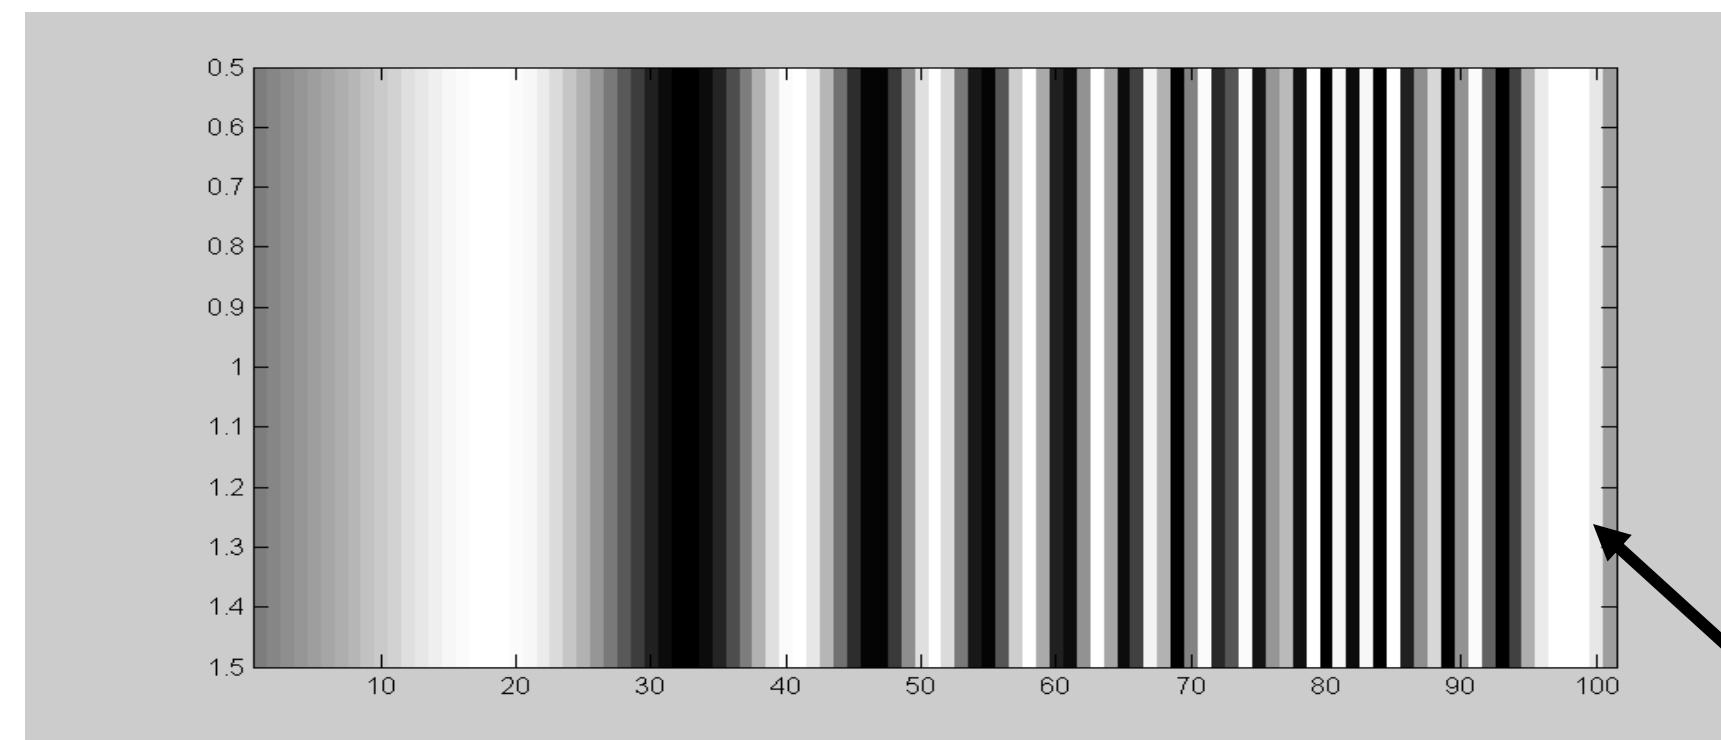


What's happening?

Input signal:



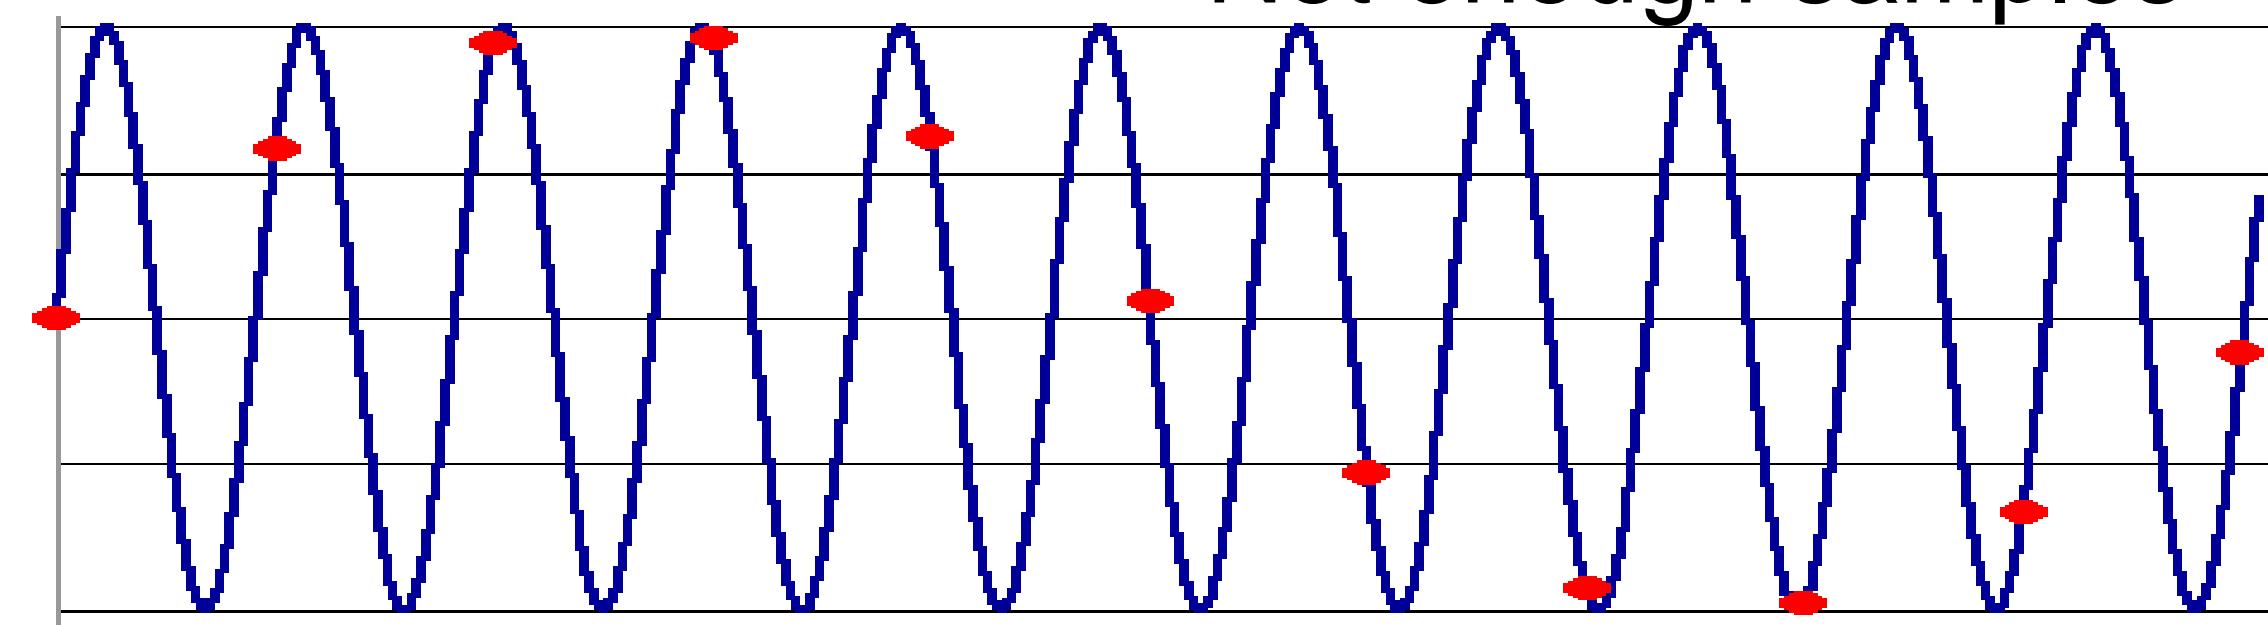
Plot as image:



`x = 0:.05:5; imagesc(sin((2.^x).*x))`

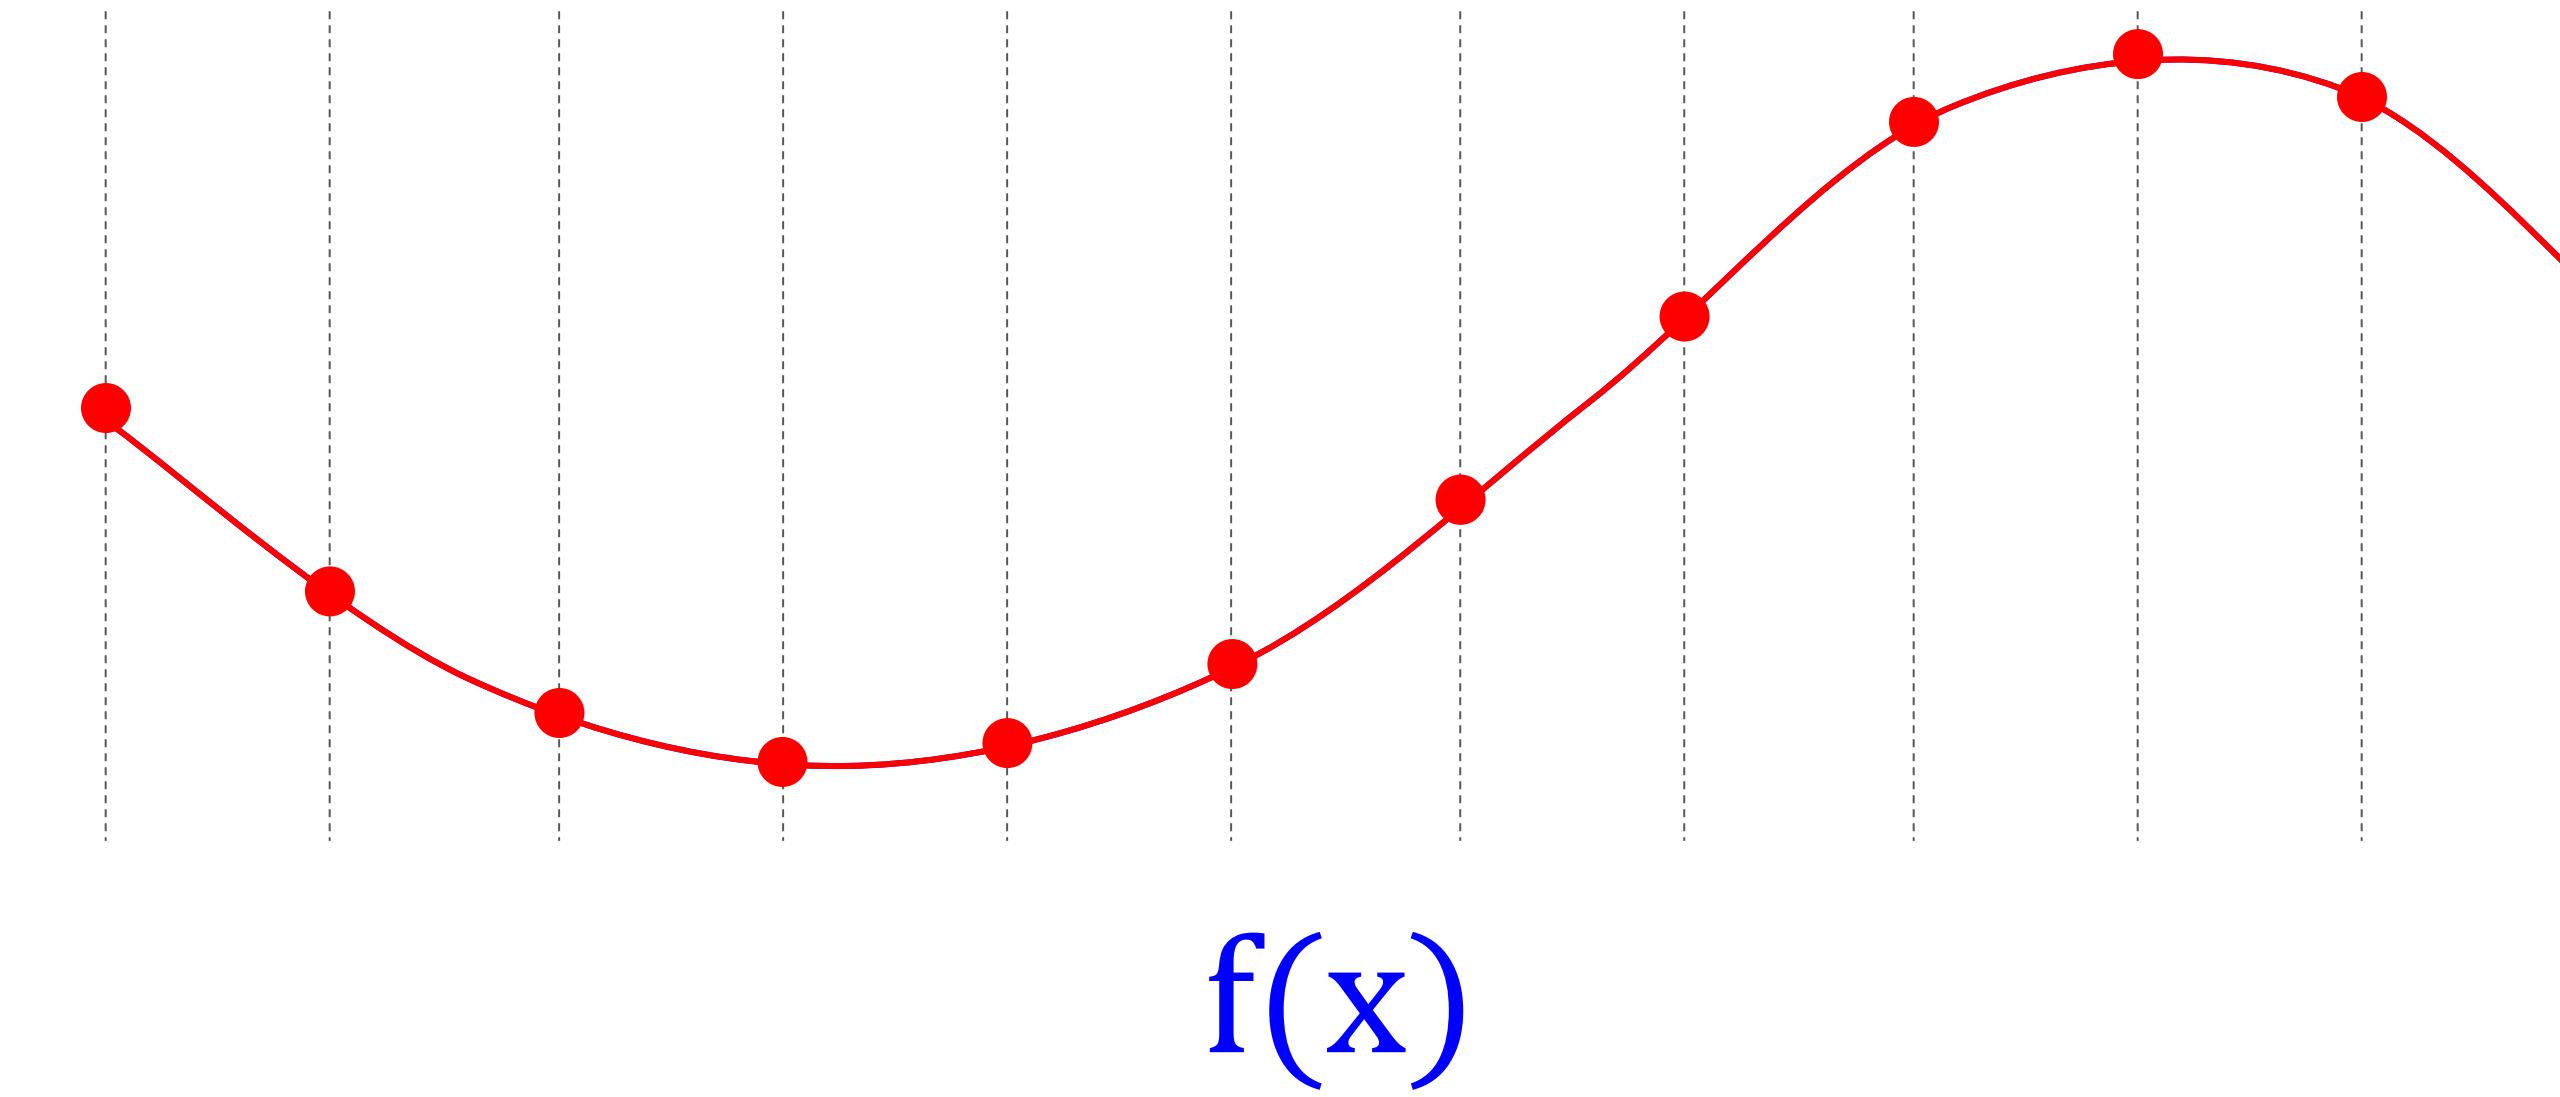
Alias!

Not enough samples



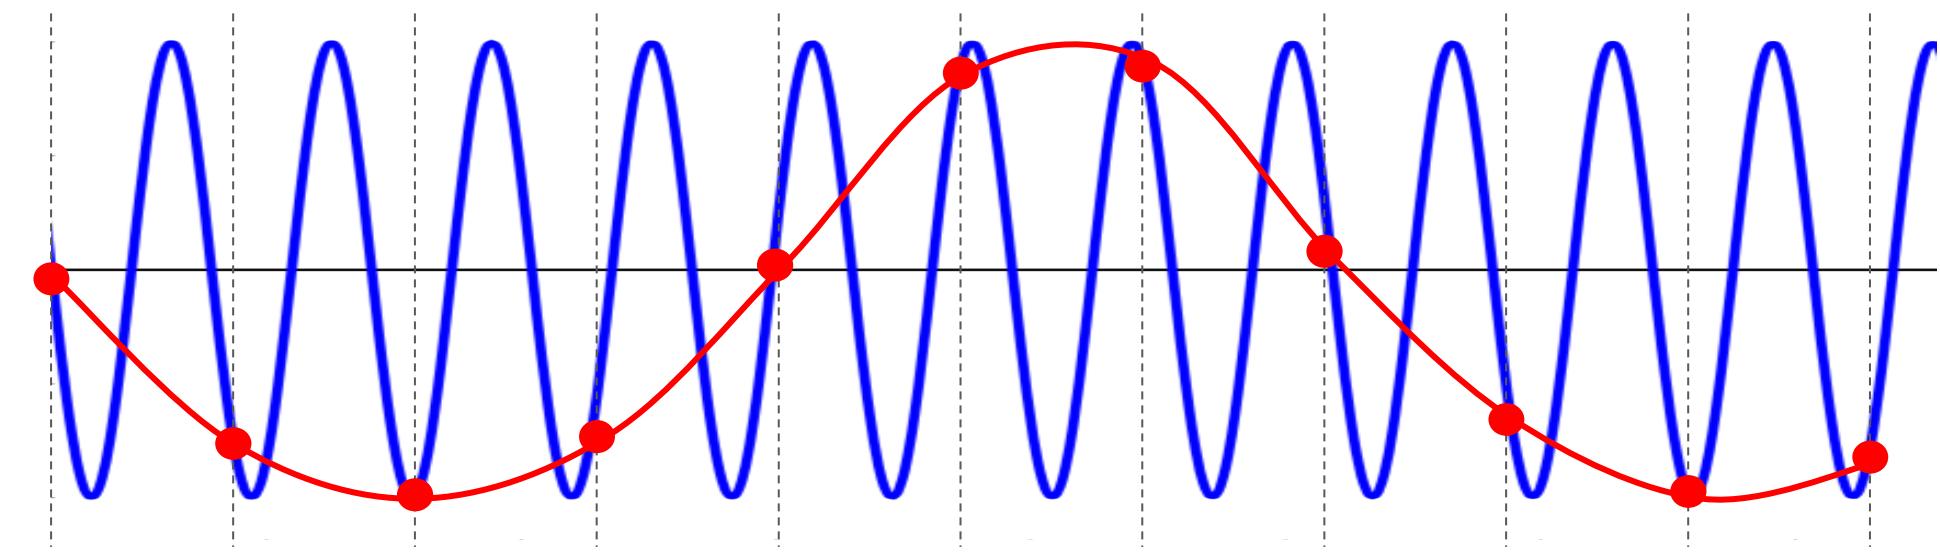
## Aliasing

$$g(x) \sim= f(x)$$

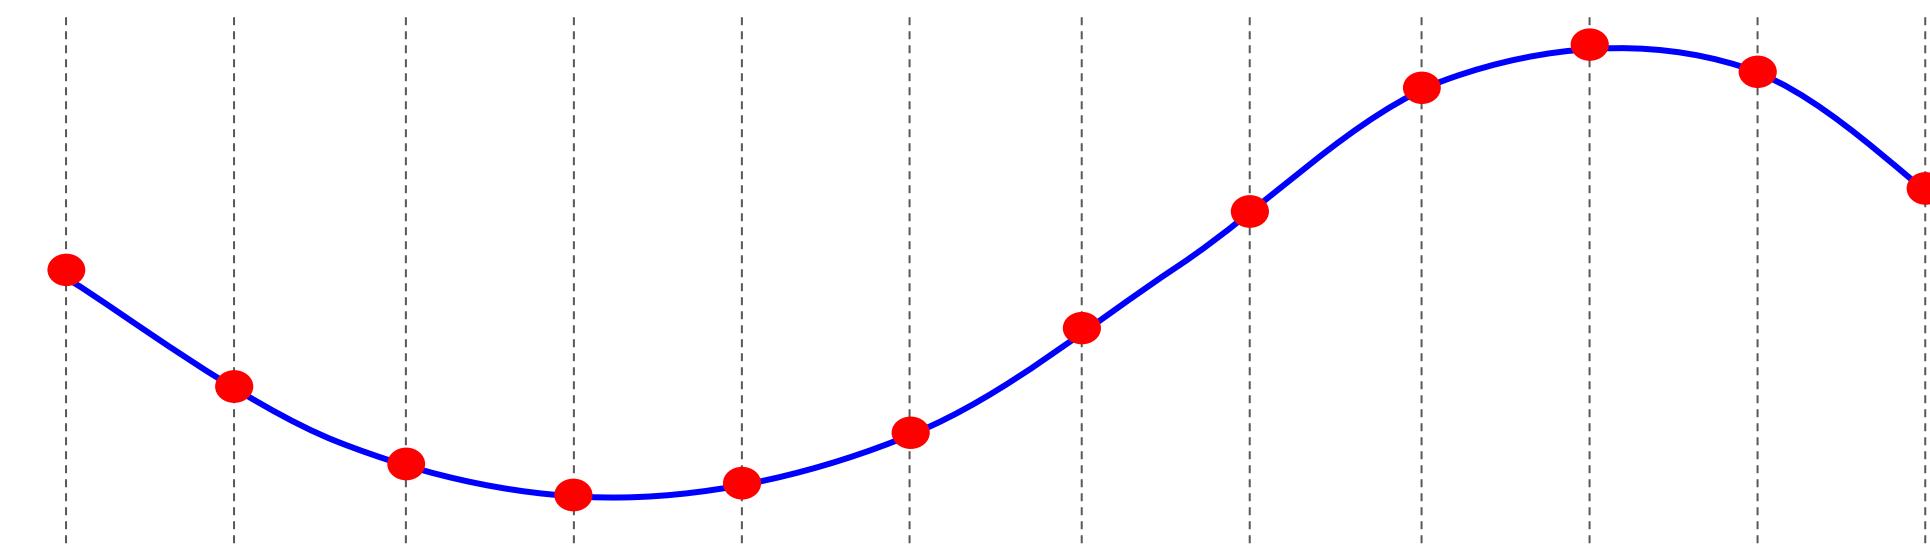


No aliasing in this case

## Sampling and the Nyquist rate



< 1 sample per cycle



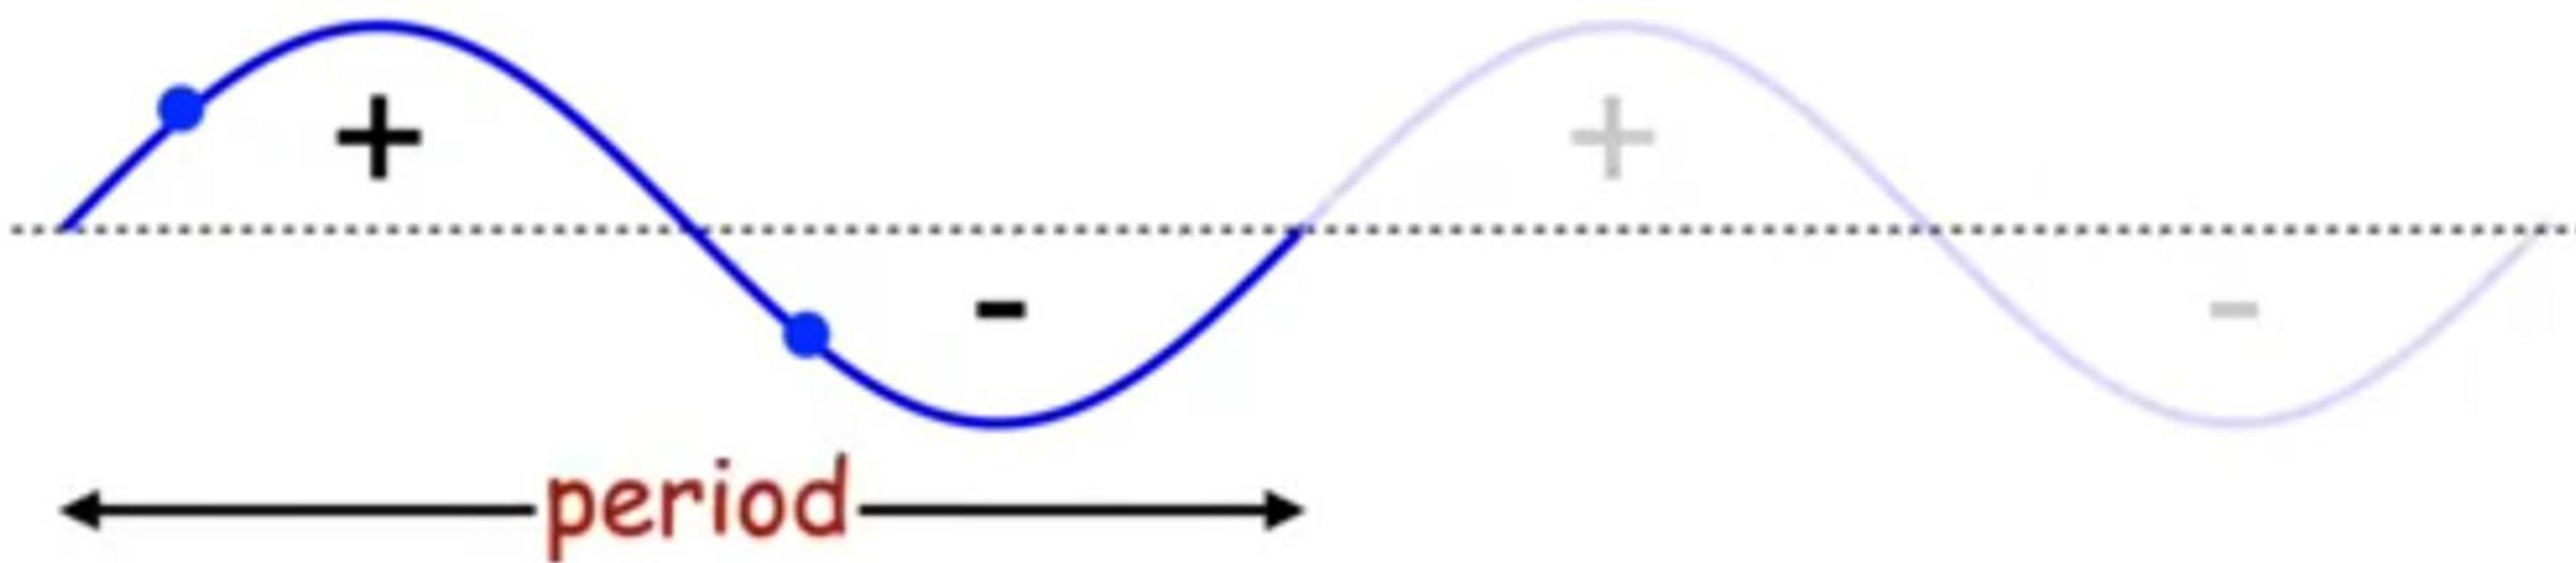
> 10 samples per cycle

To avoid aliasing

- **≥ two samples per cycle**

This minimum sampling rate is called the **Nyquist rate**

# To prevent aliasing...



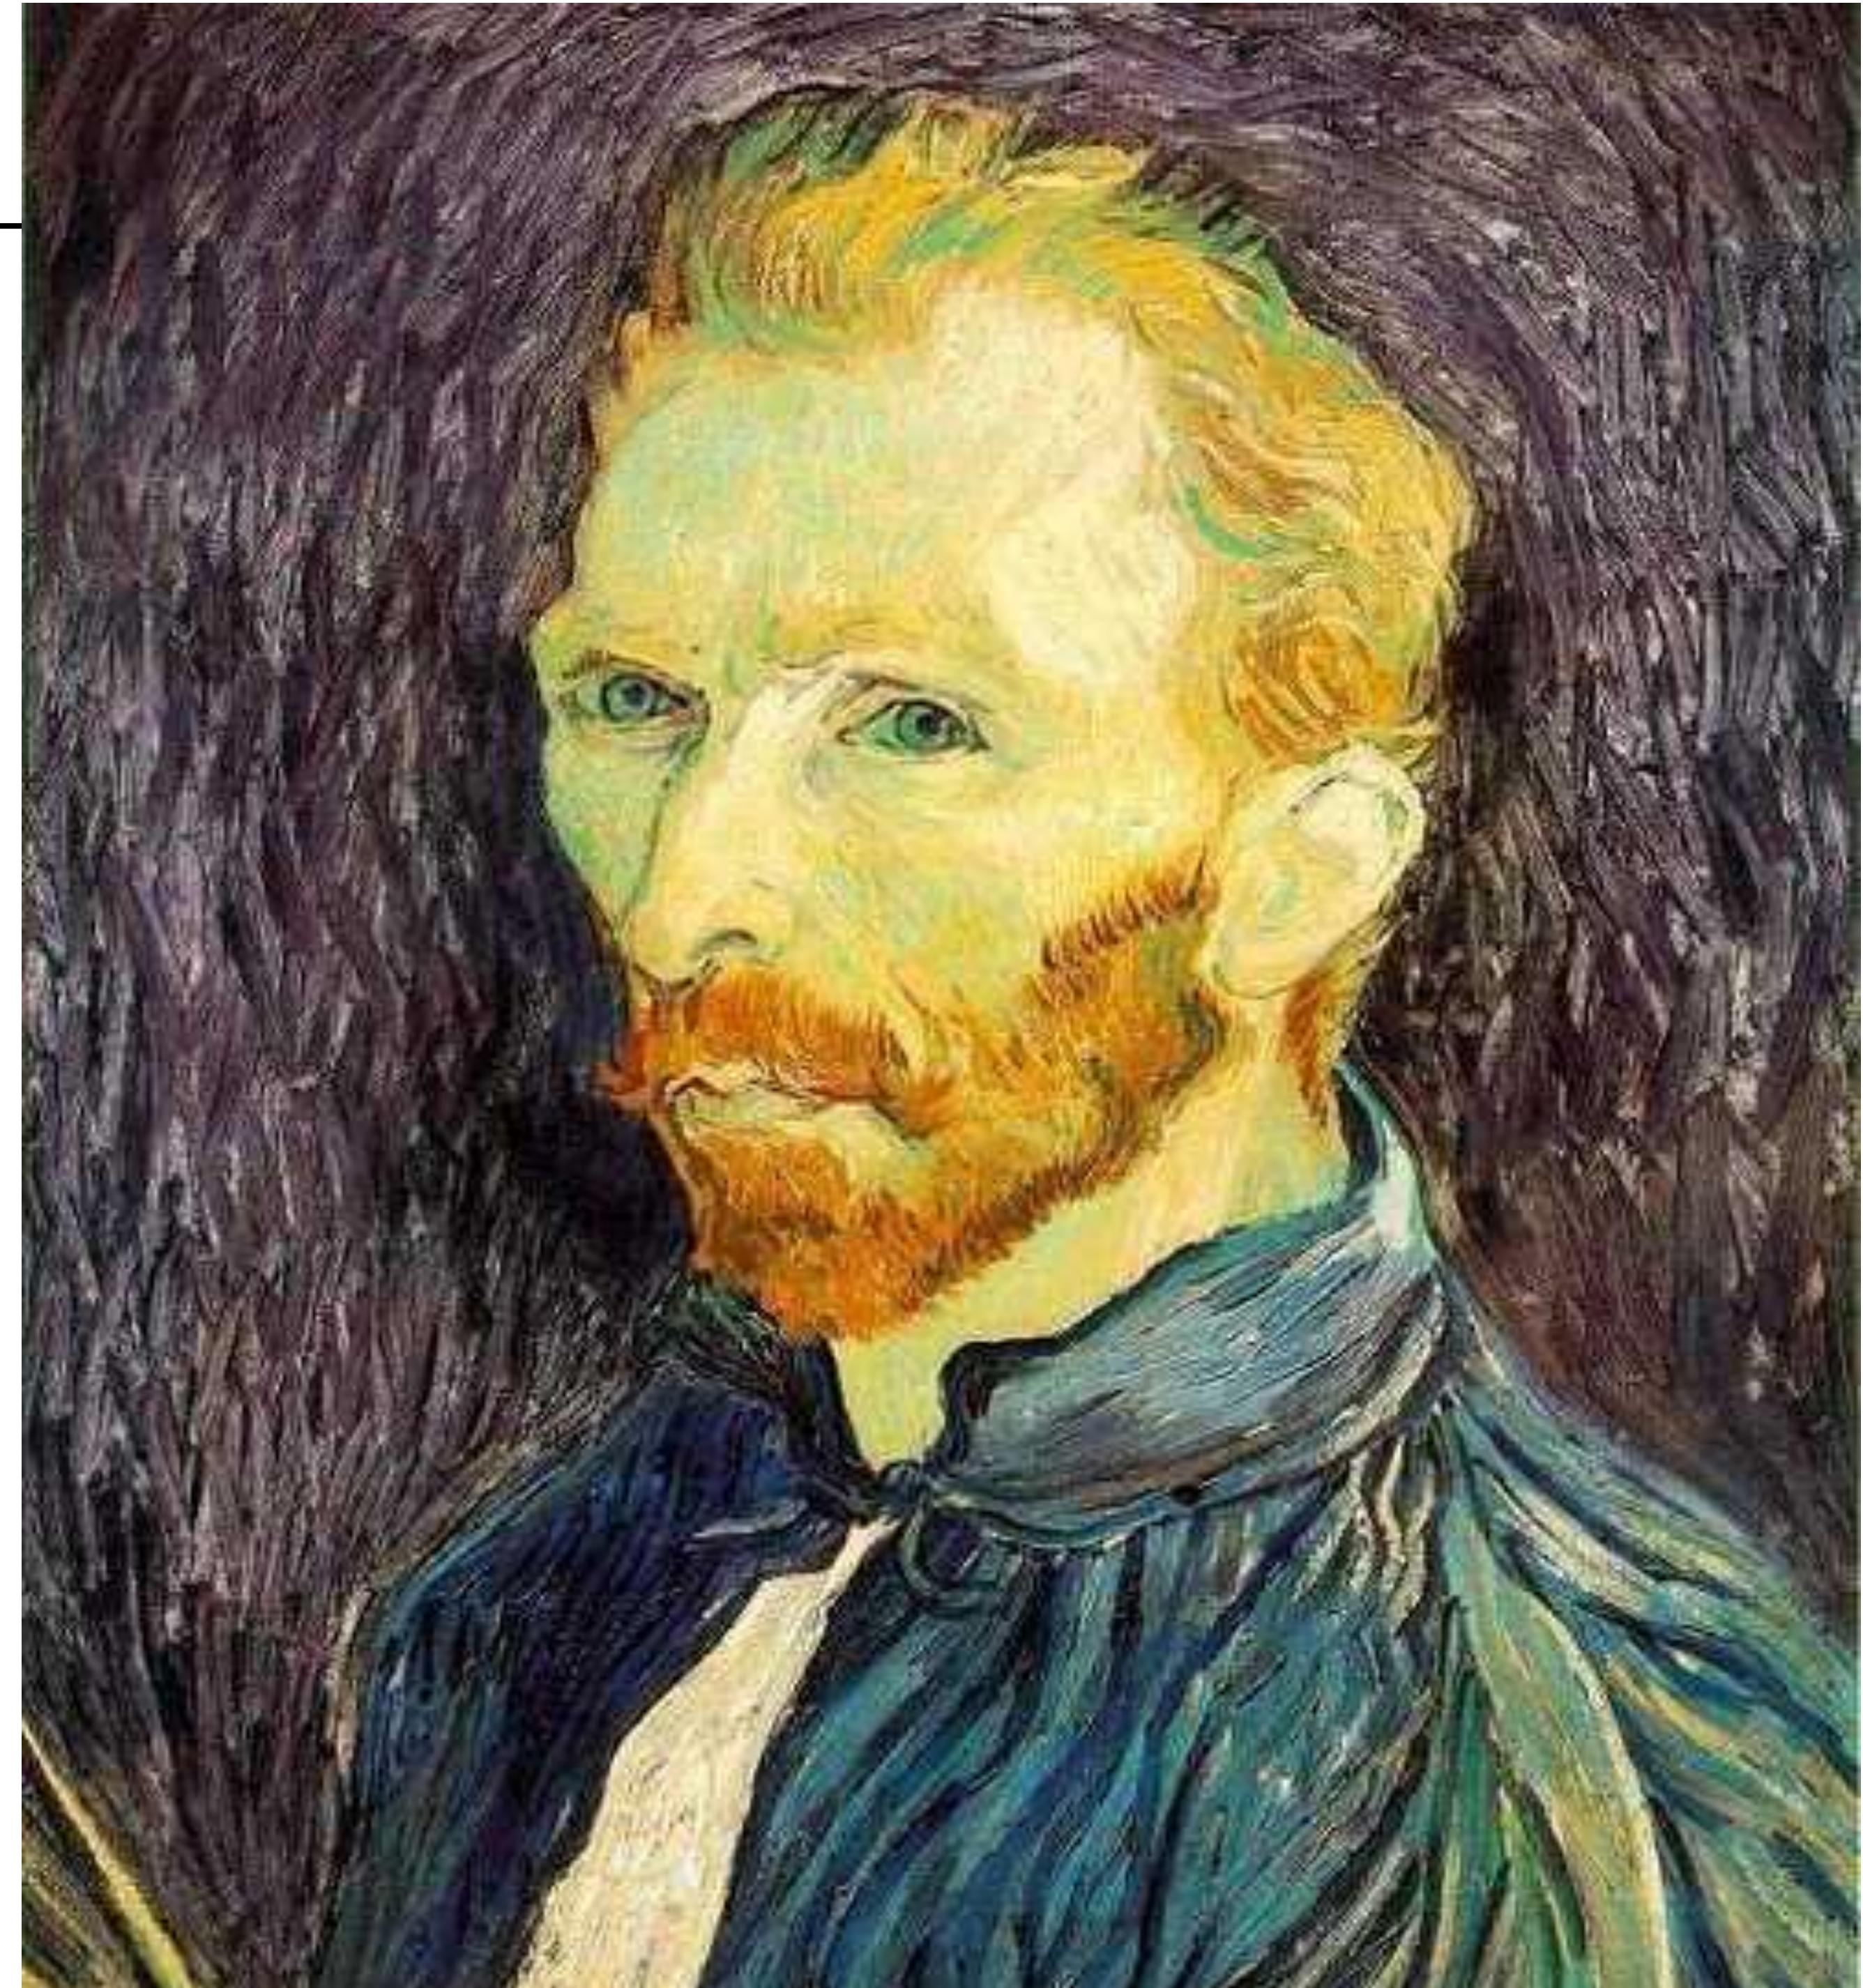
>2 samples / period

# Image Scaling

---

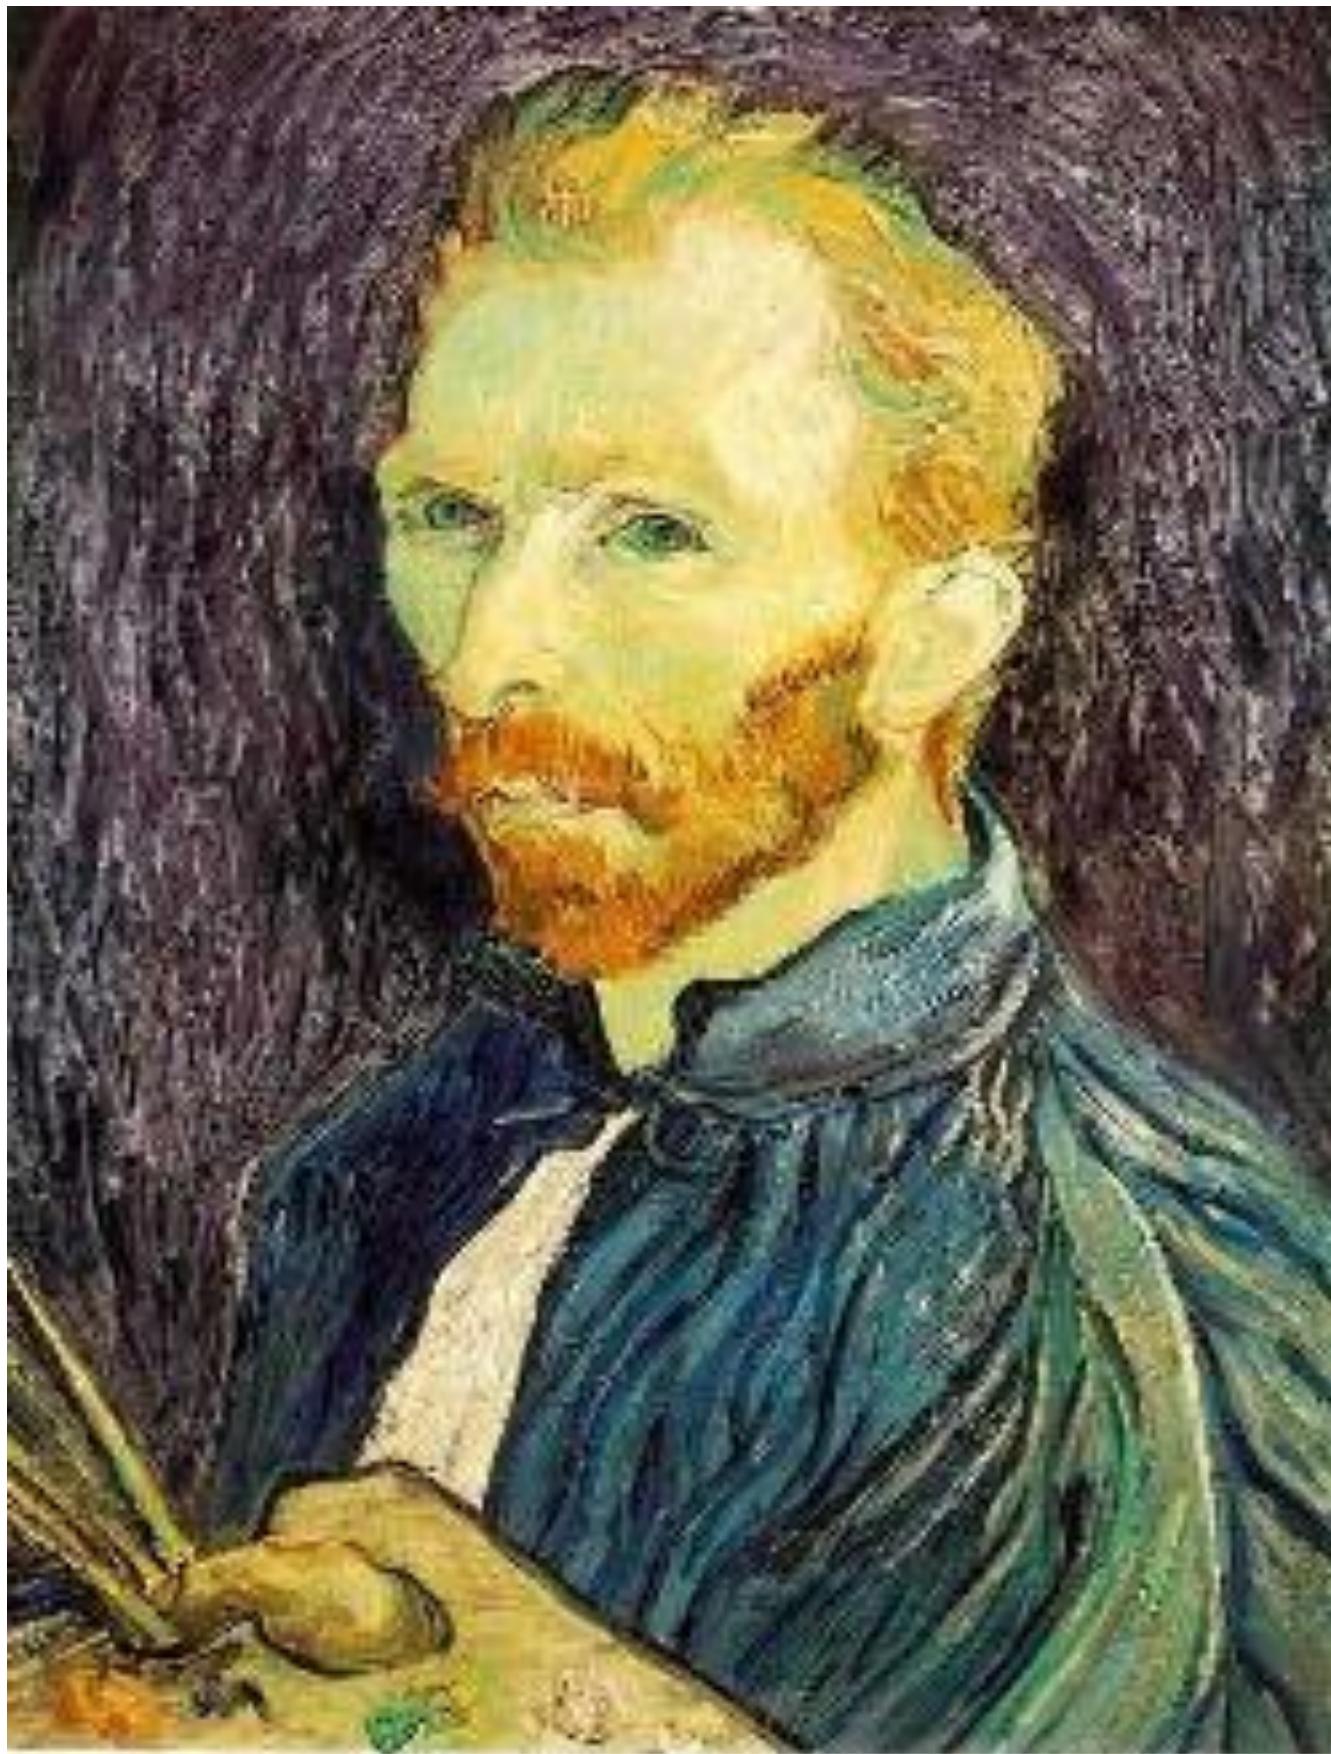
This image is too big to fit on the screen. How can we reduce it?

How to generate a half-sized version?



# Image sub-sampling

---



1/4

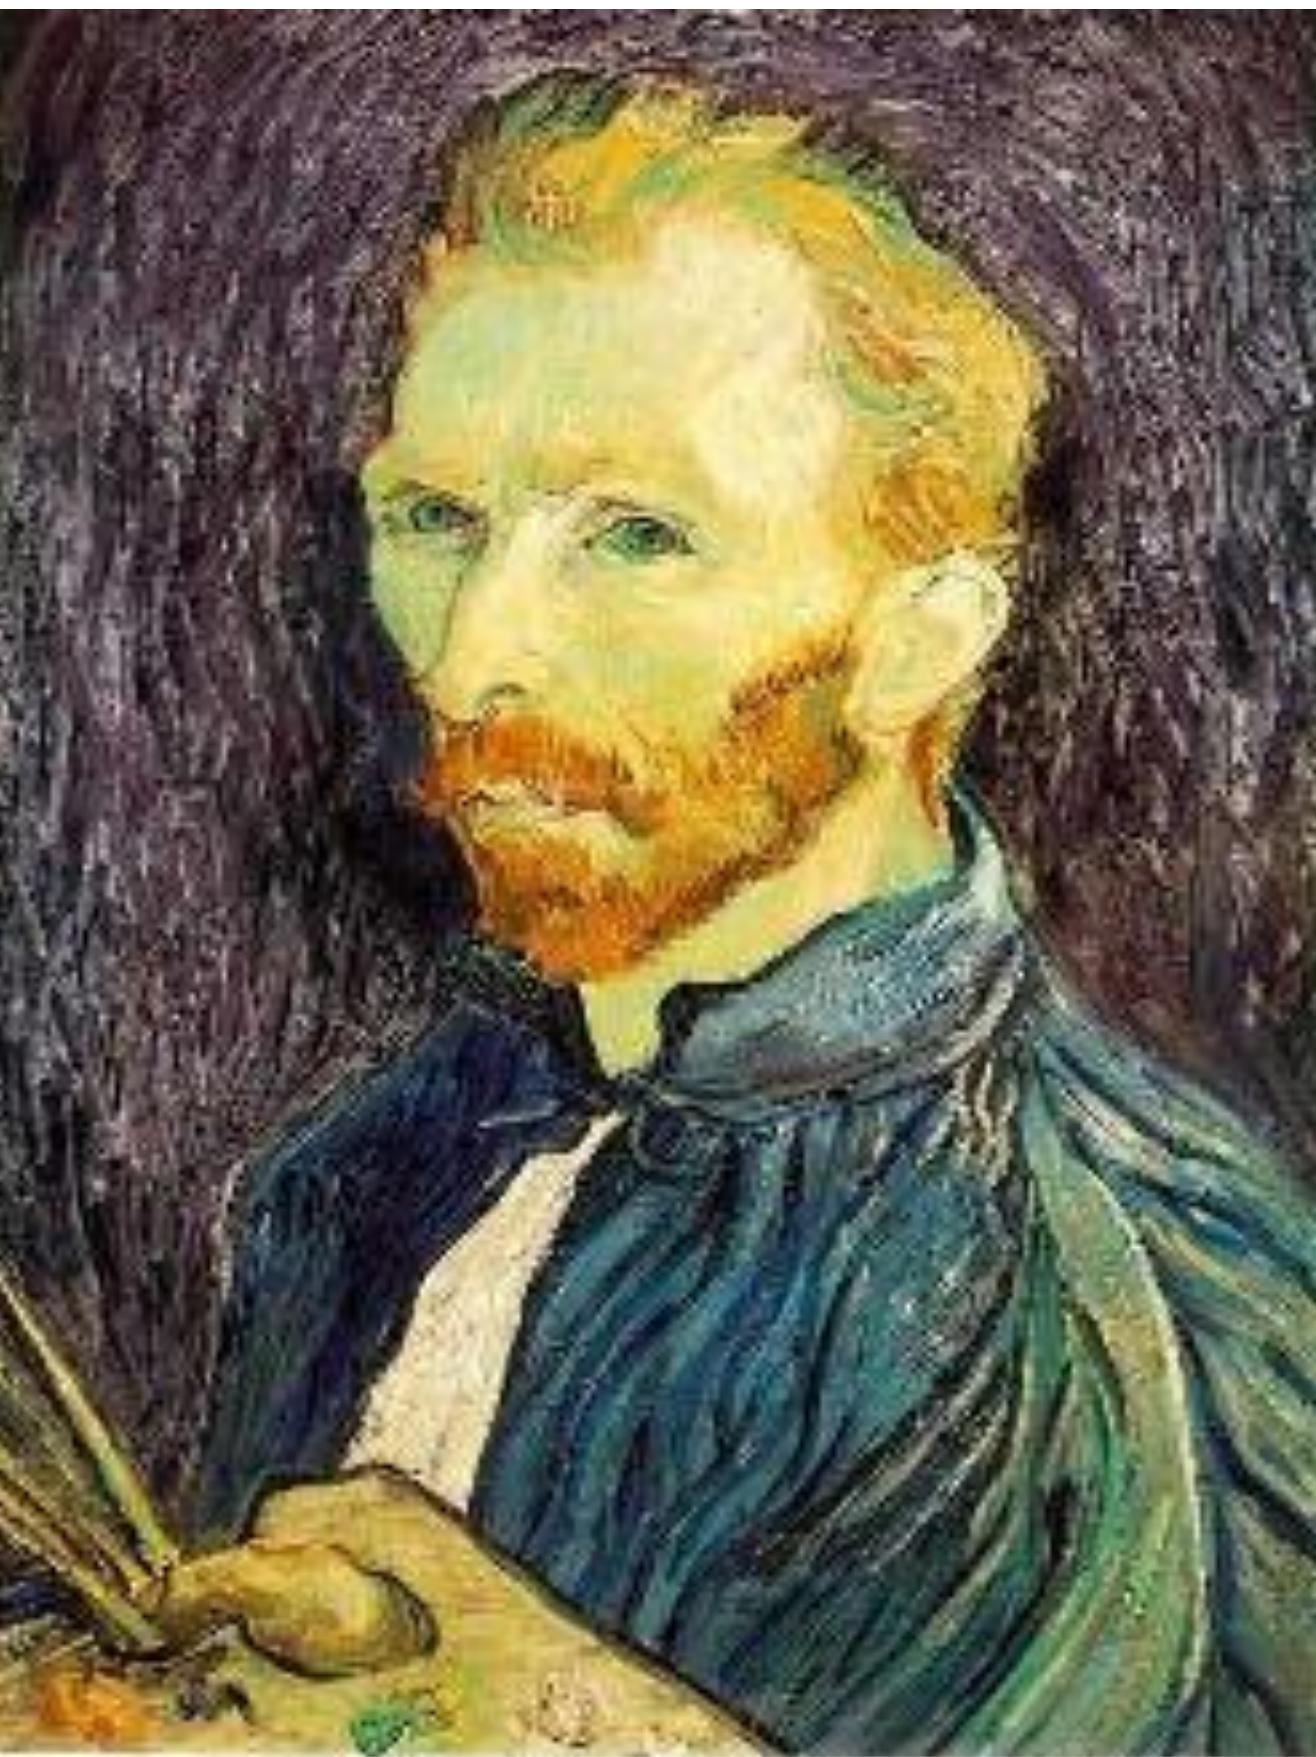


1/8

Throw away every other row and  
column to create a 1/2 size image  
- called *image sub-sampling*

# Image sub-sampling

---



1/2



1/4 (2x zoom)



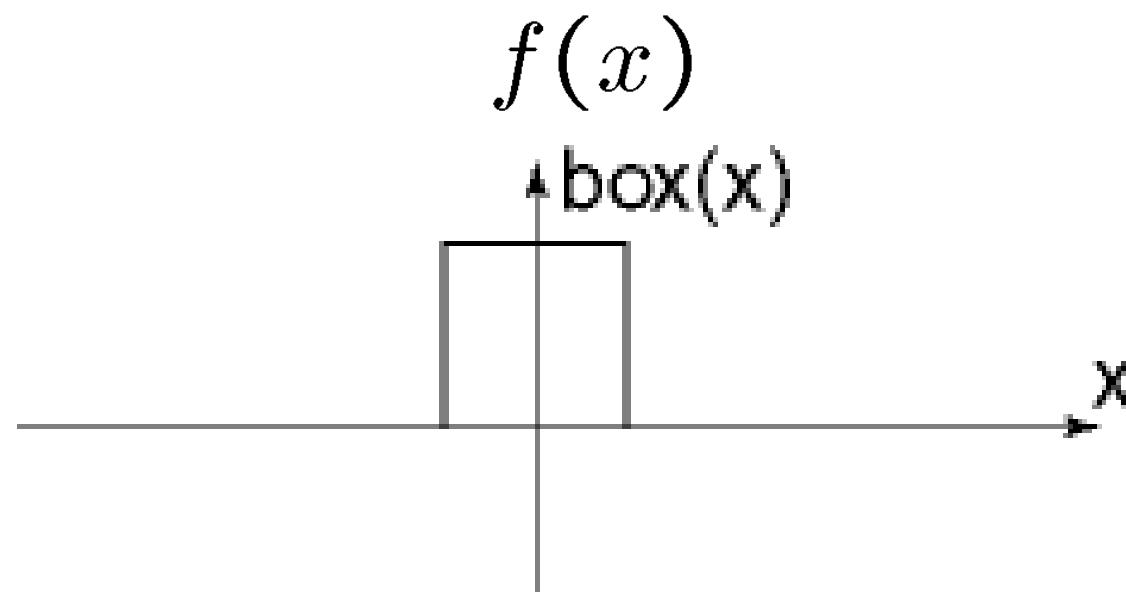
1/8 (4x zoom)

Why does this look so cruddy?

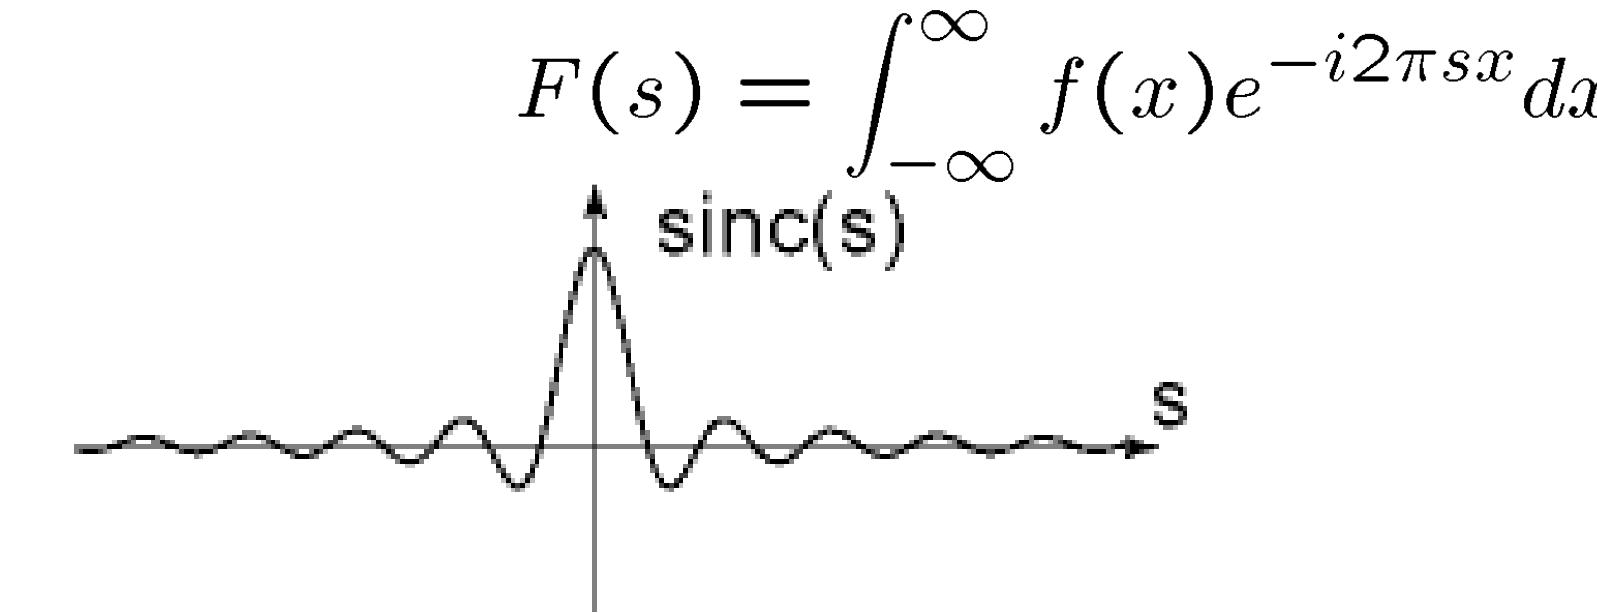
# Fourier transform

---

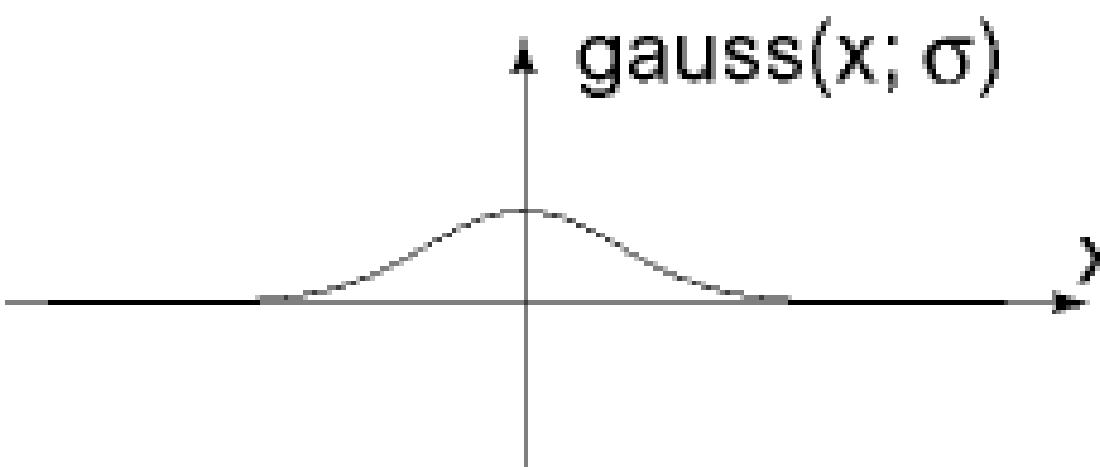
Spatial domain



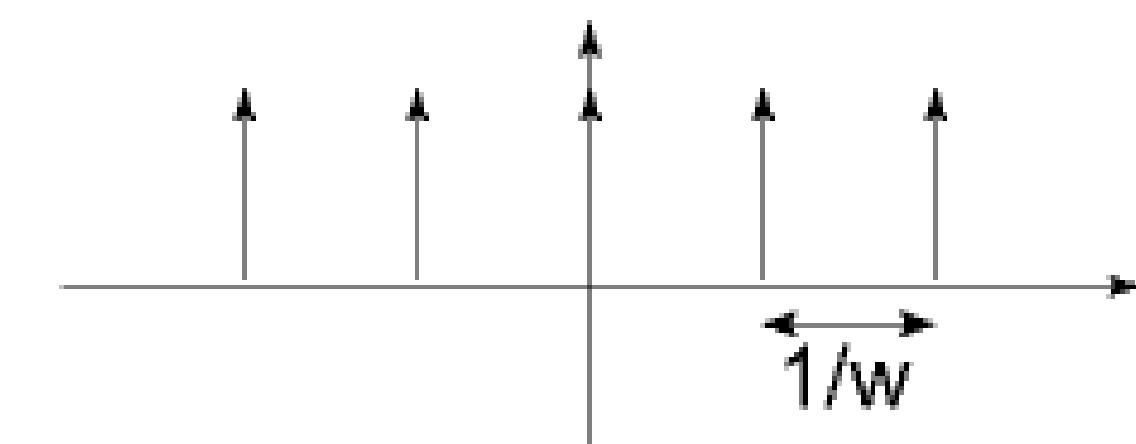
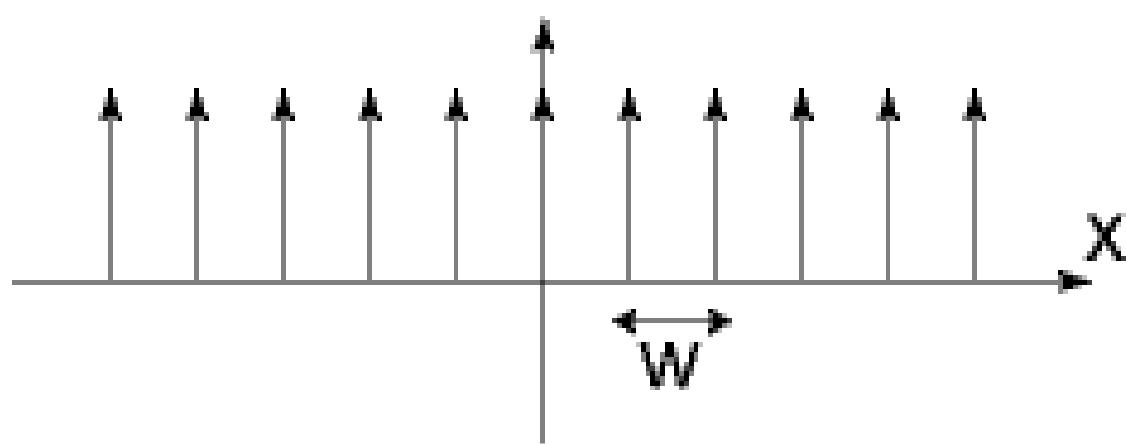
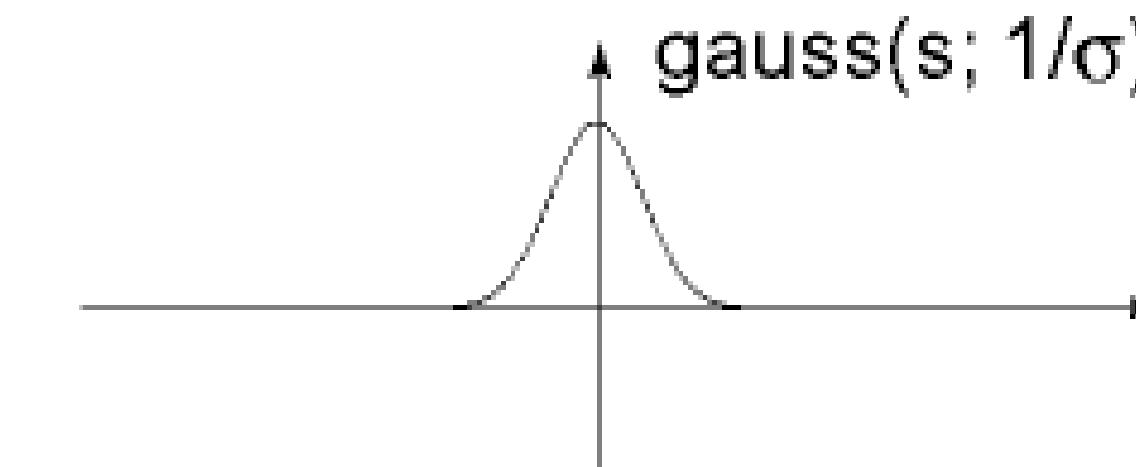
Frequency domain



$\text{gauss}(x; \sigma)$



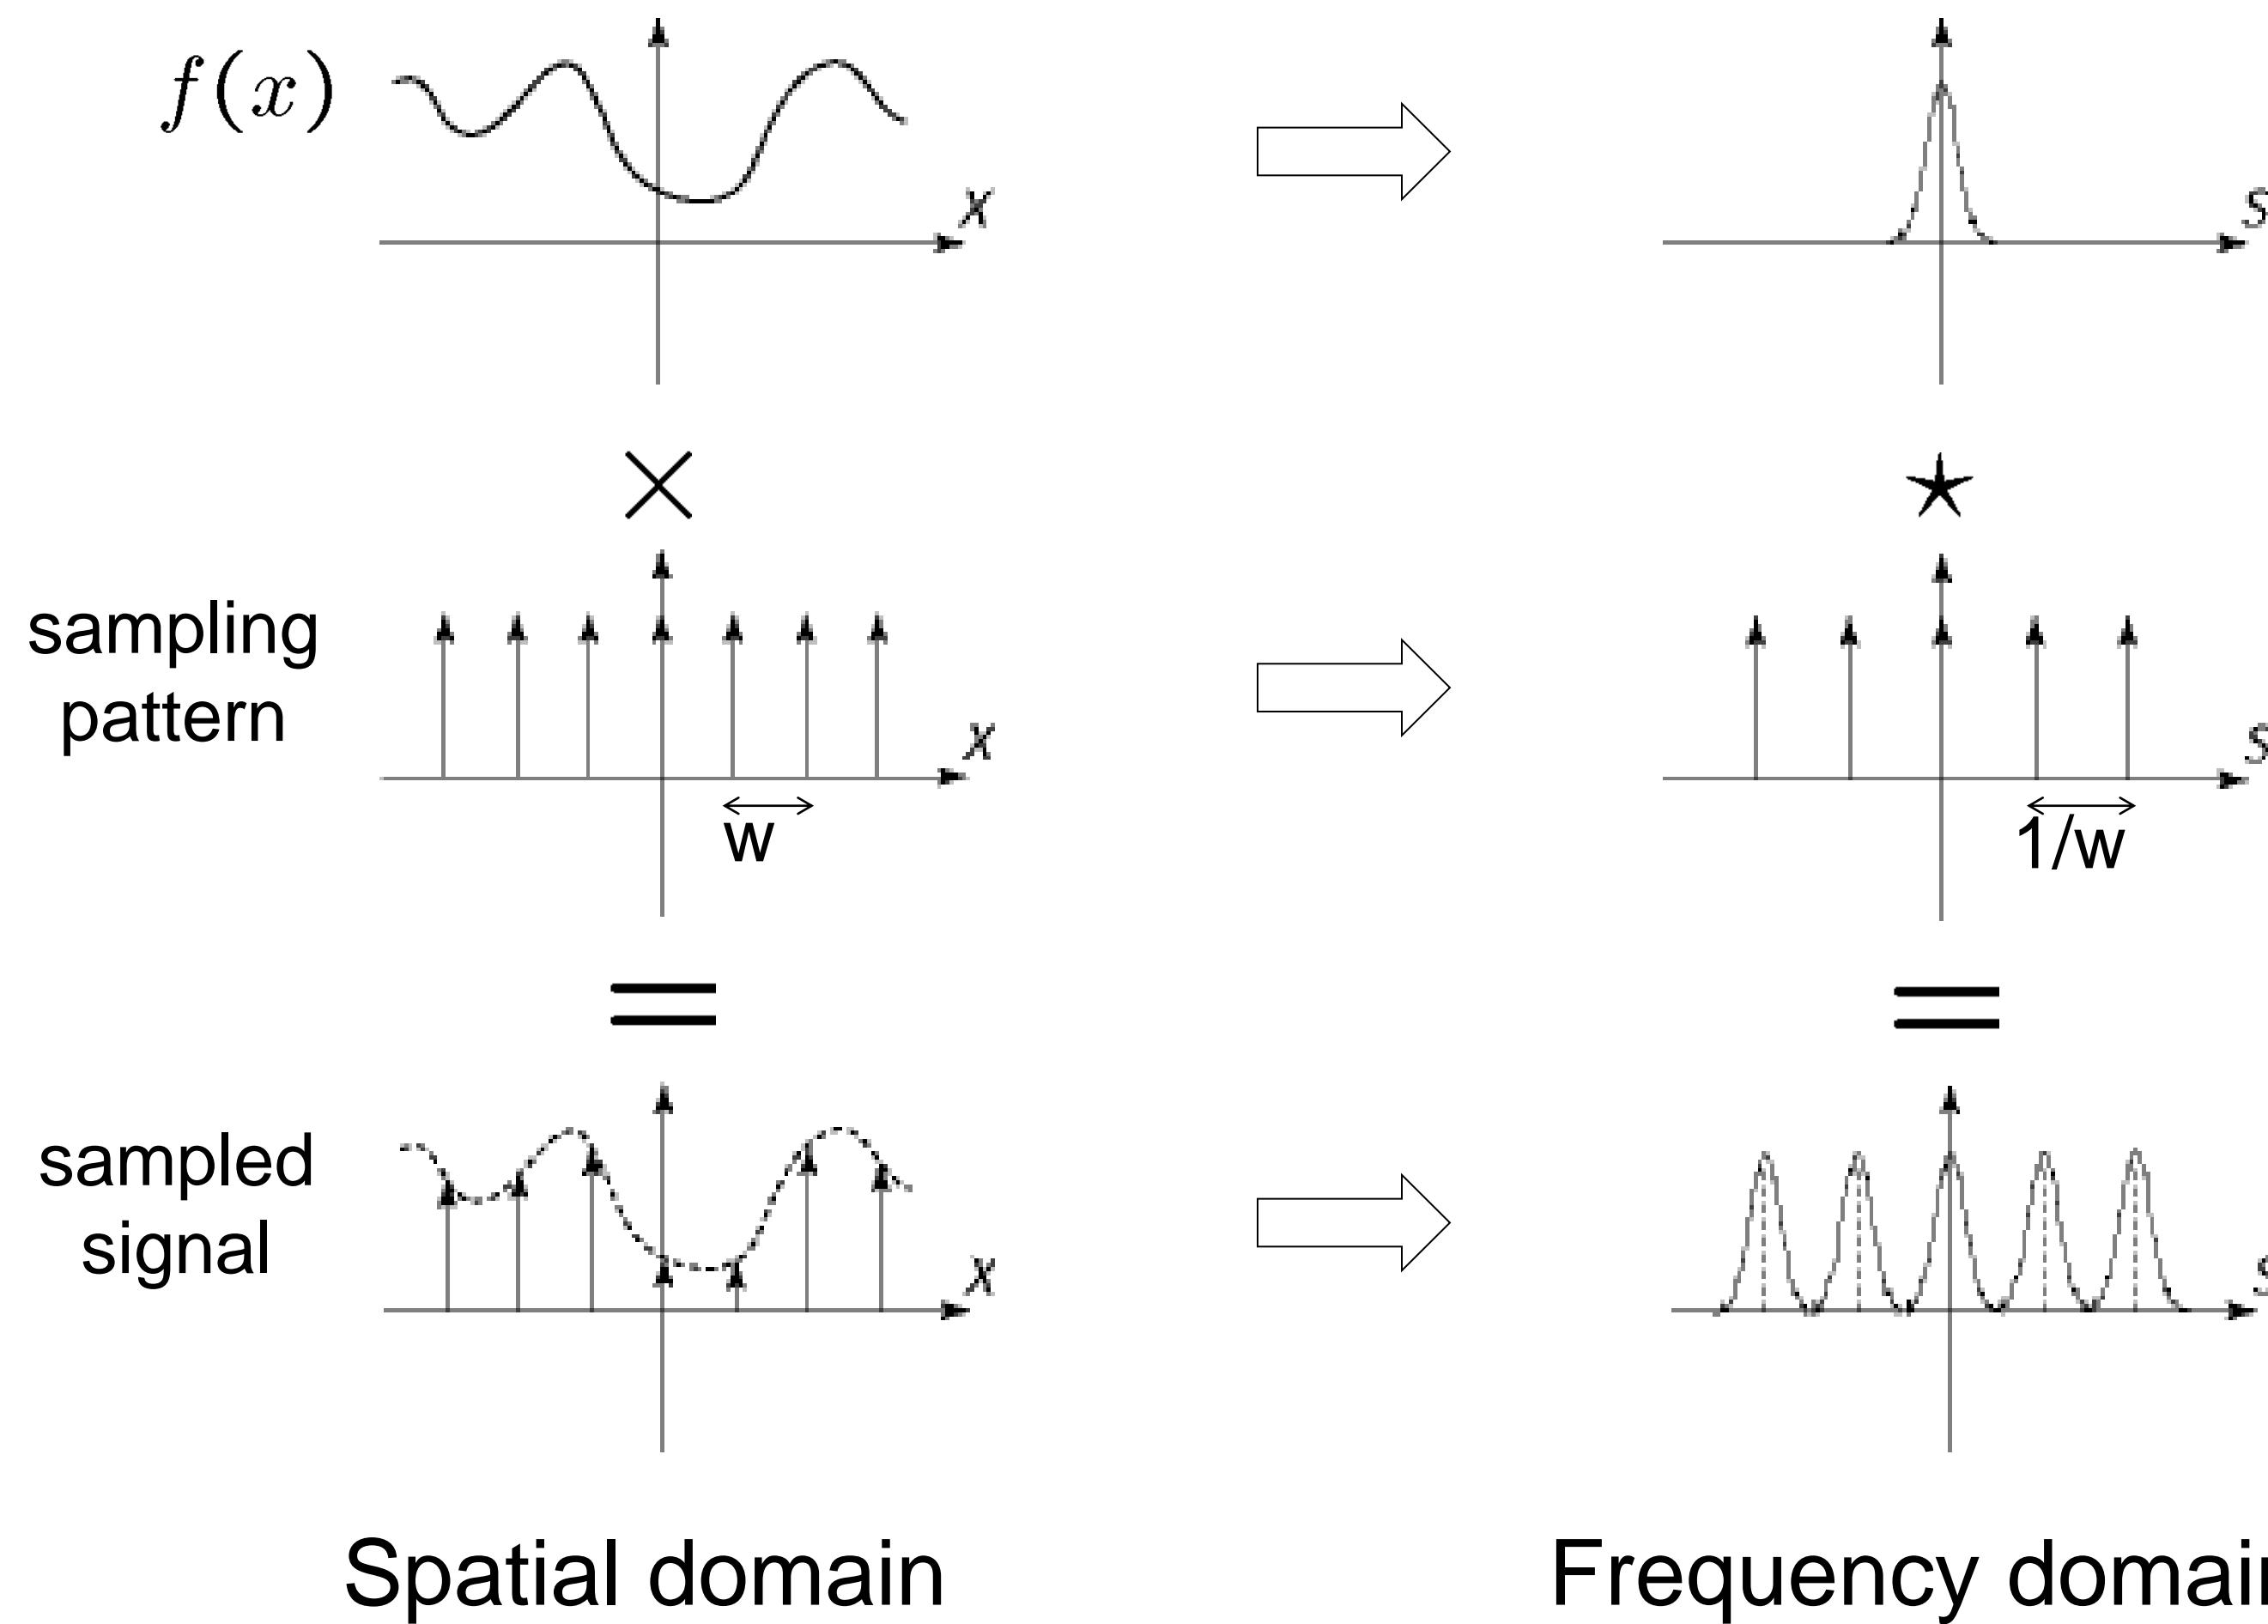
$\text{gauss}(s; 1/\sigma)$



# Sampling

---

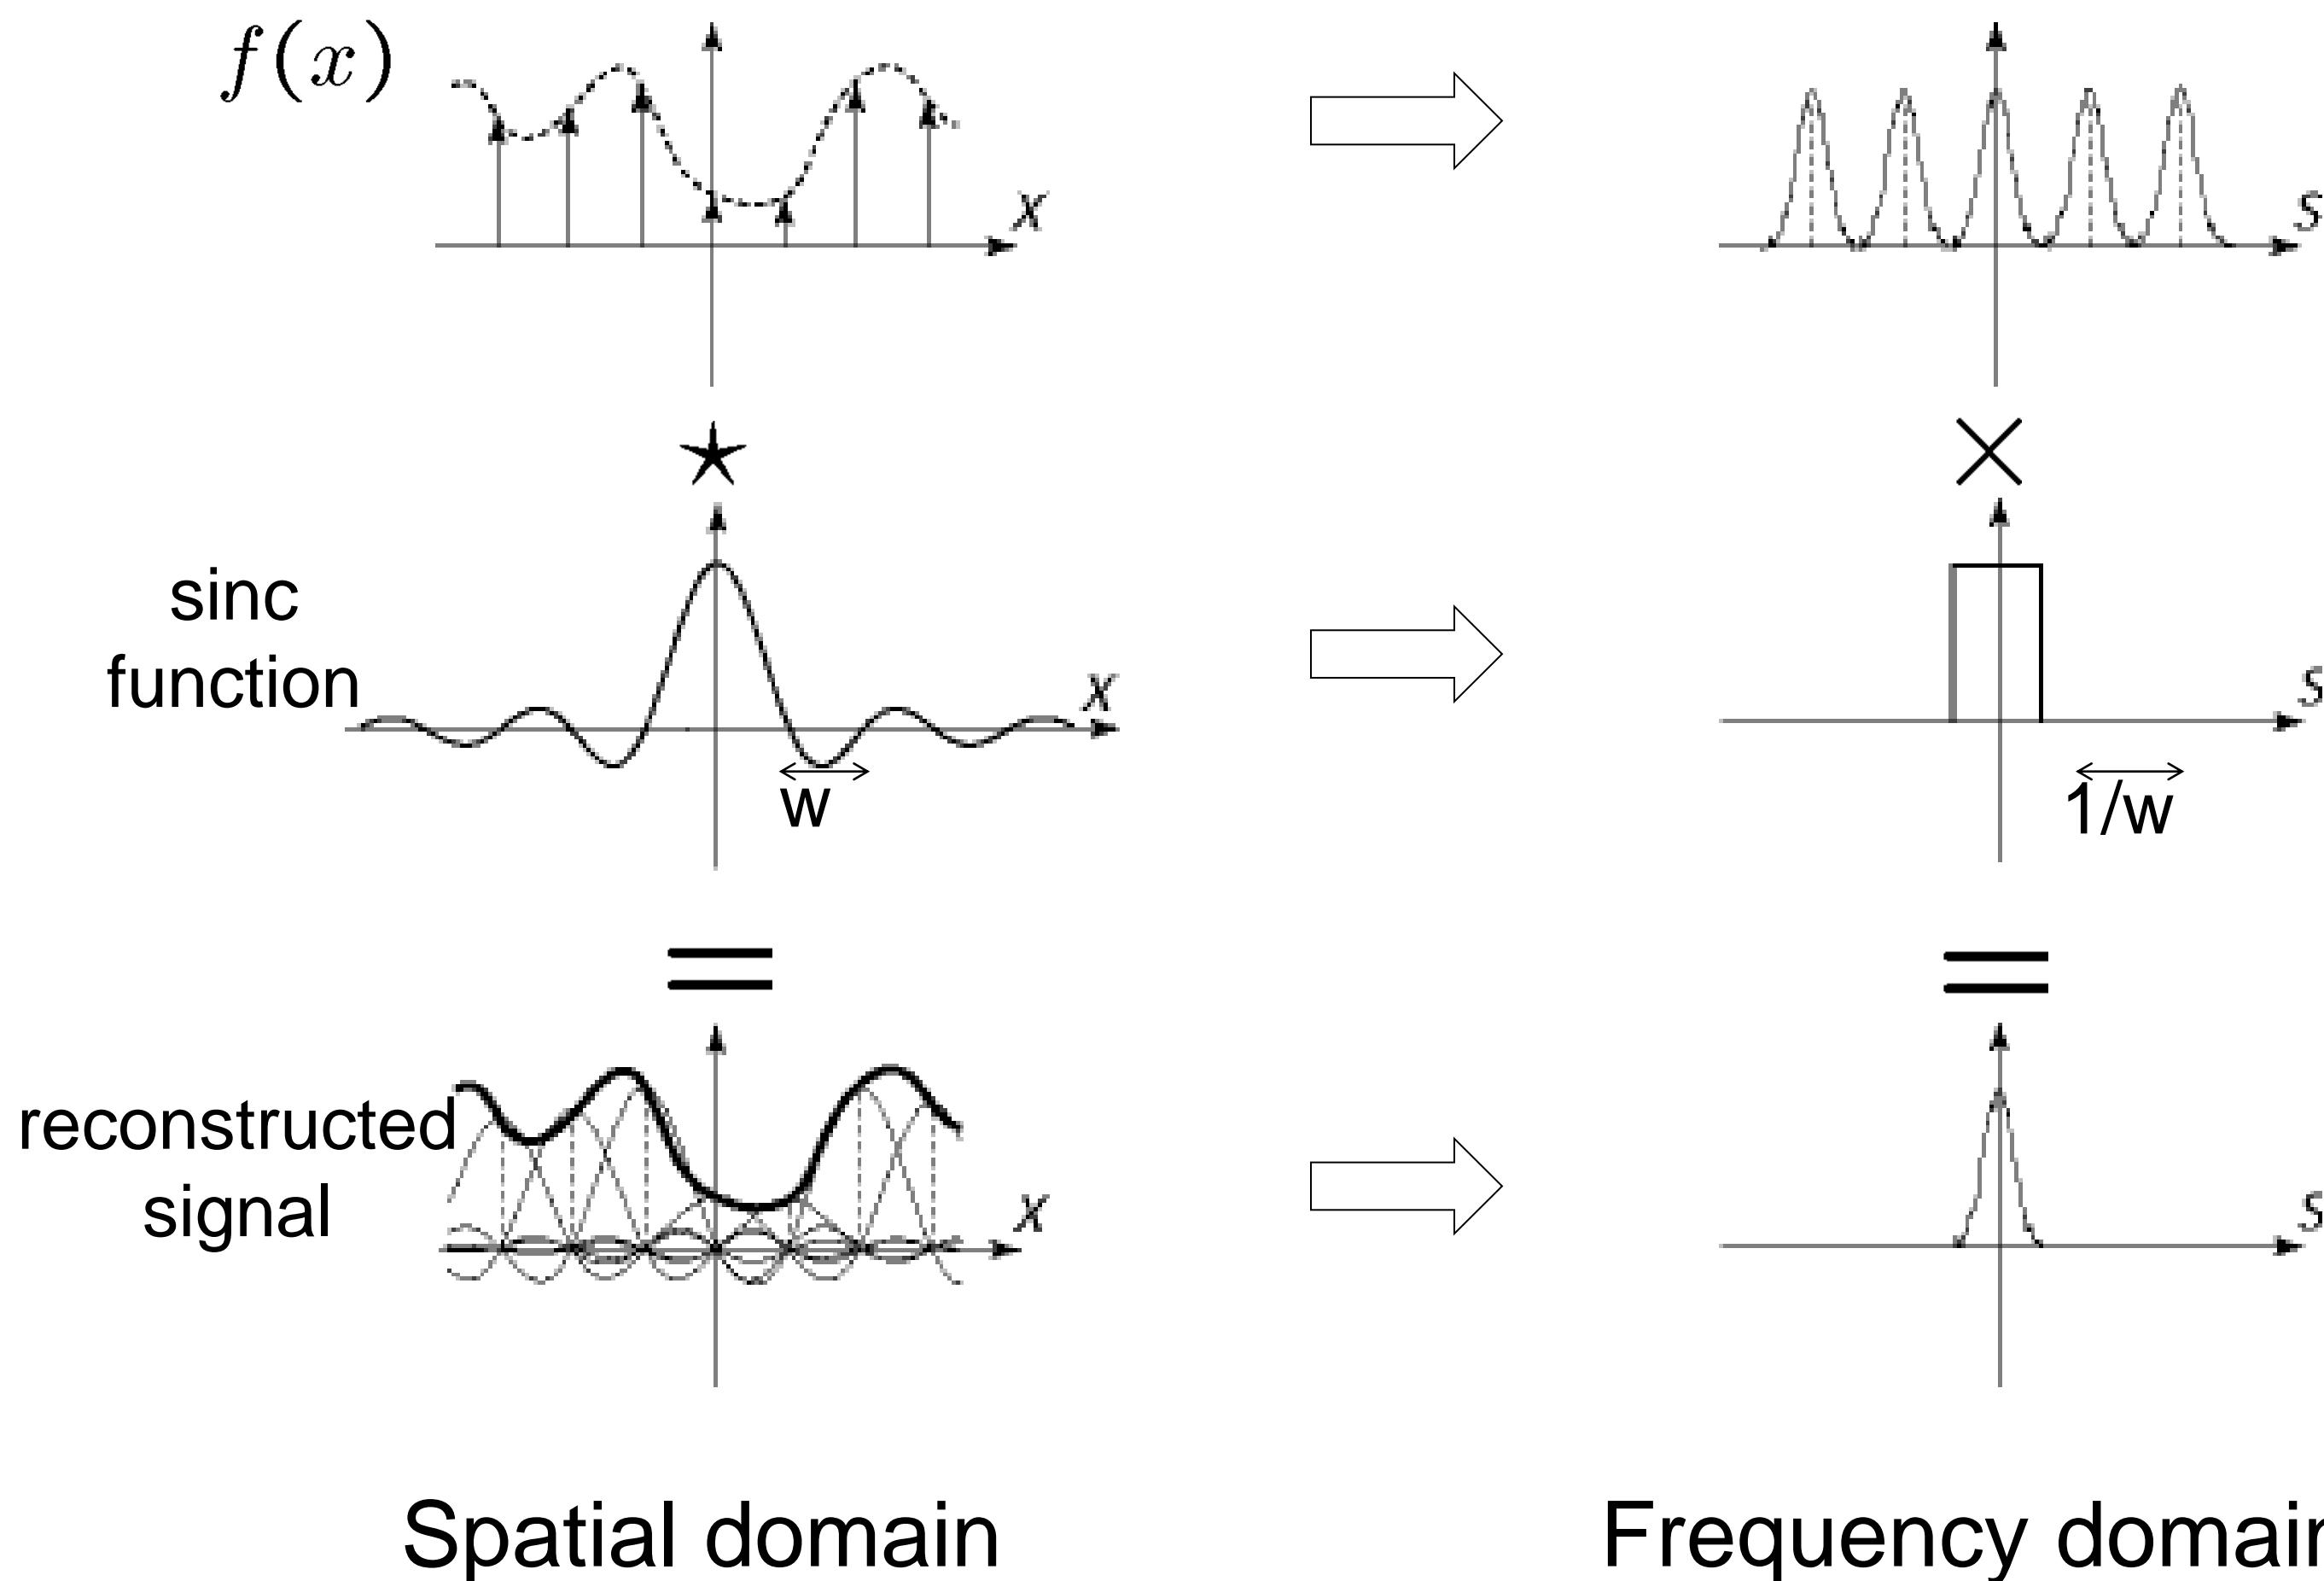
$$F(s) = \int_{-\infty}^{\infty} f(x) e^{-i2\pi s x} dx$$



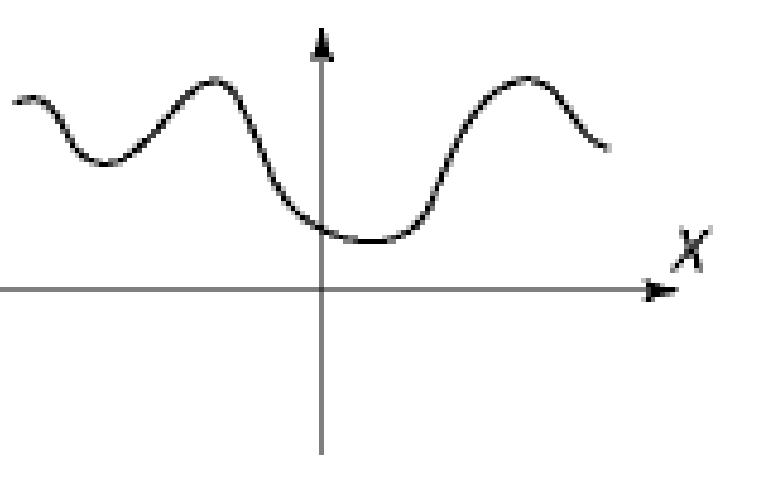
# Reconstruction

---

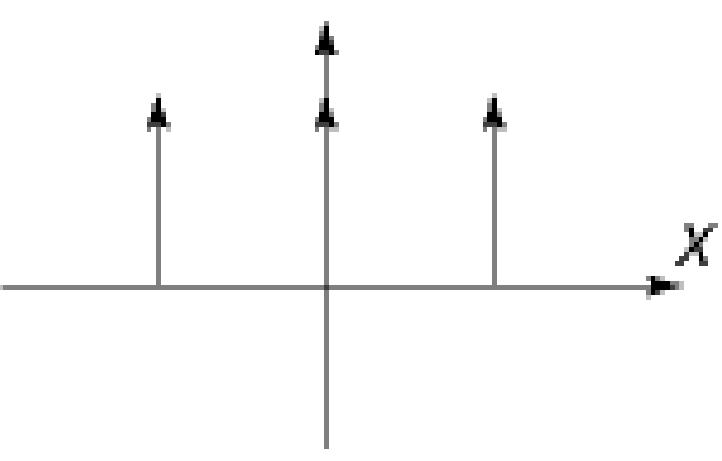
$$F(s) = \int_{-\infty}^{\infty} f(x)e^{-i2\pi sx}dx$$



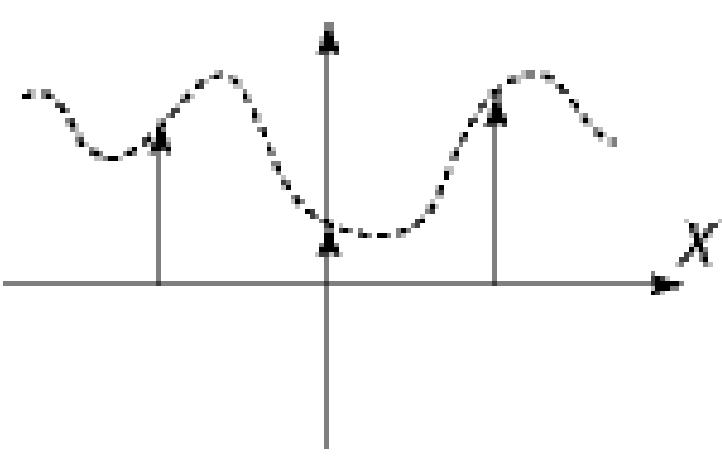
What happens when  
the sampling rate  
is too low?



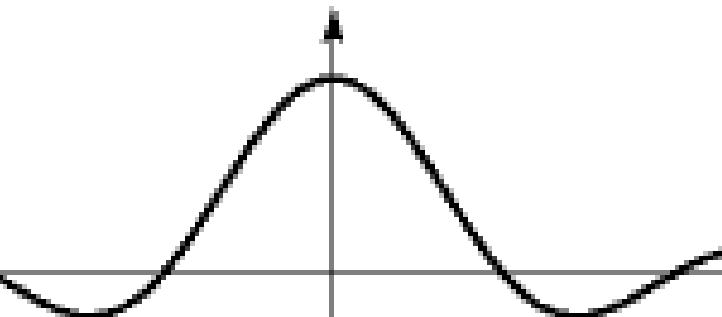
$x$



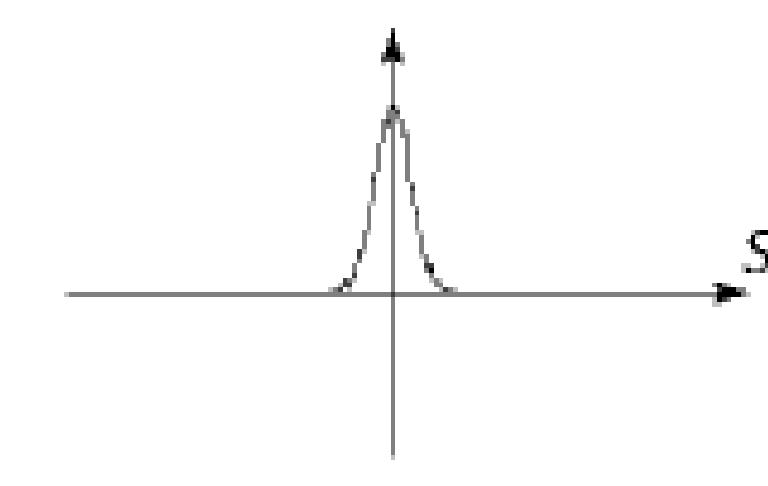
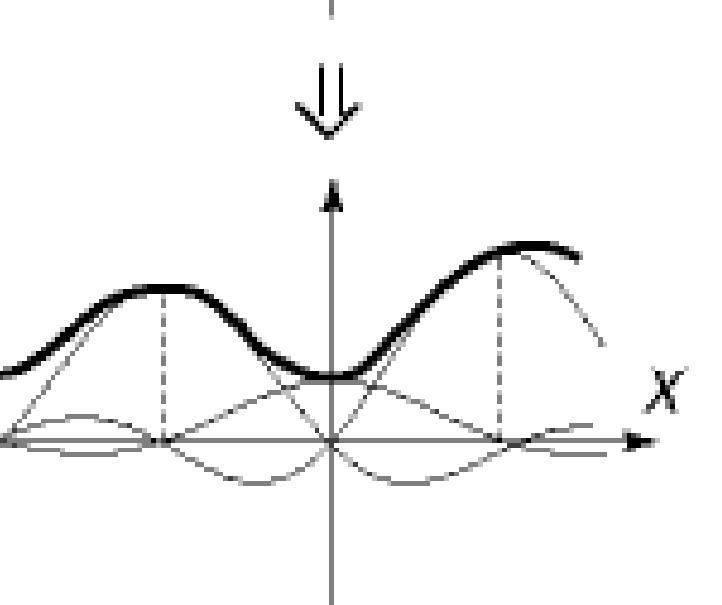
$\downarrow$



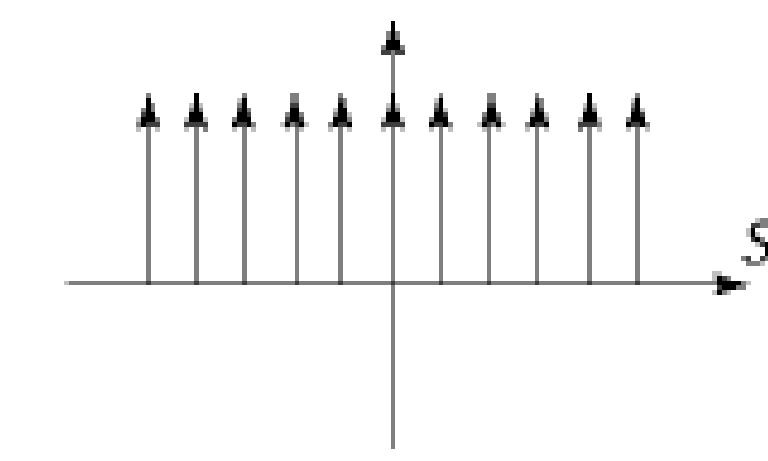
$x$



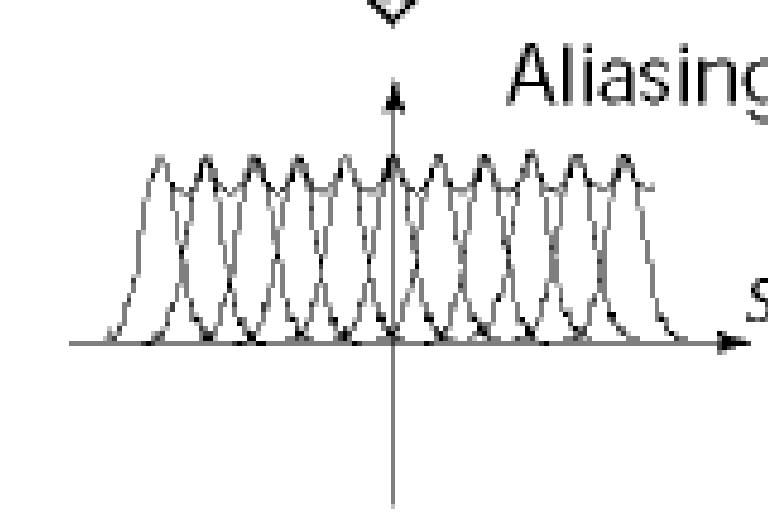
$\downarrow$



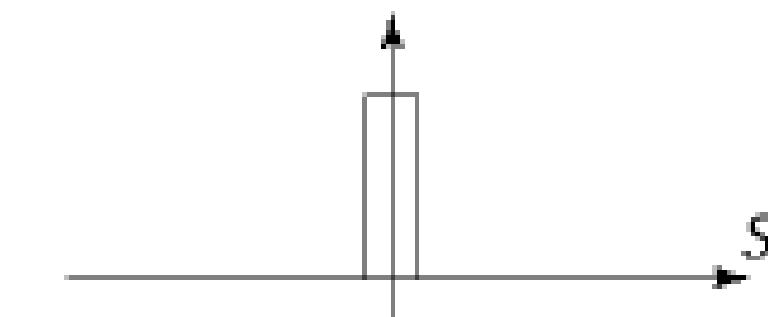
$s$



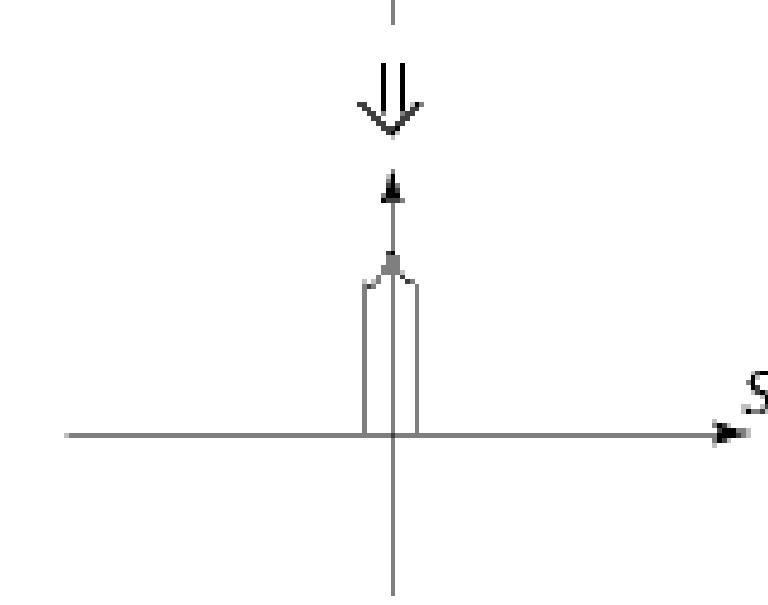
$\downarrow$



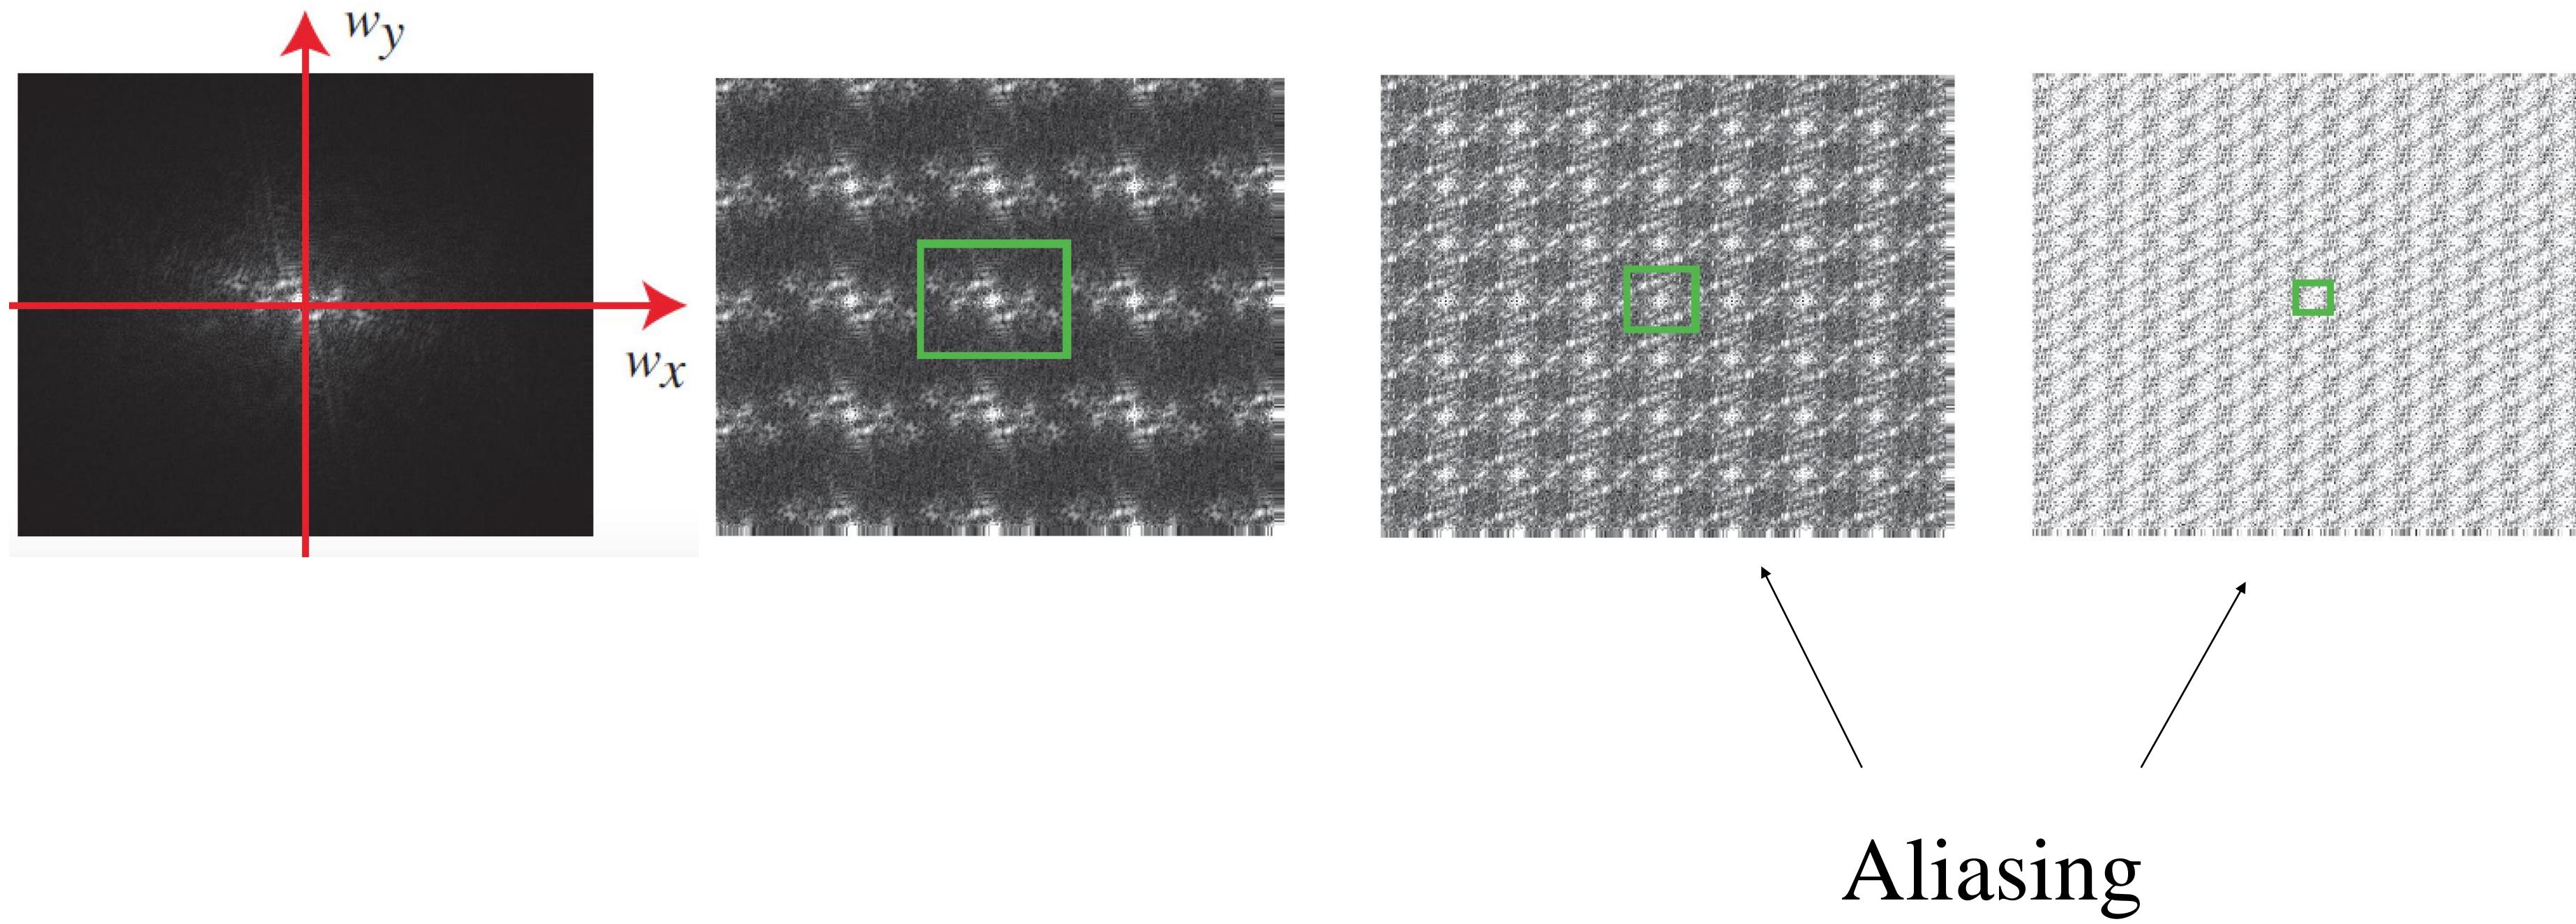
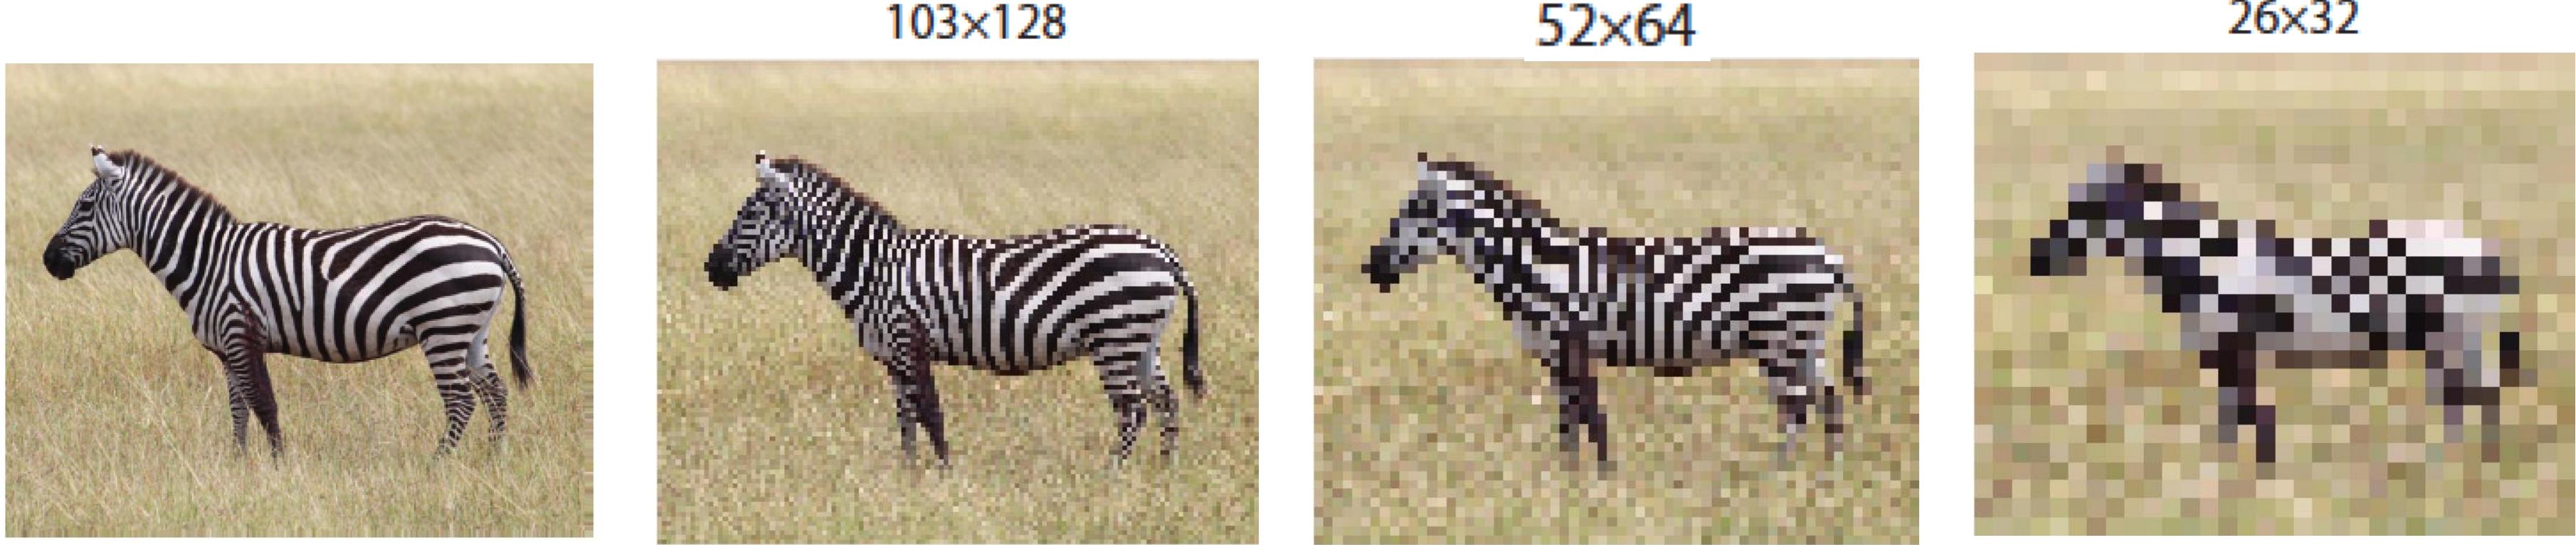
$s$



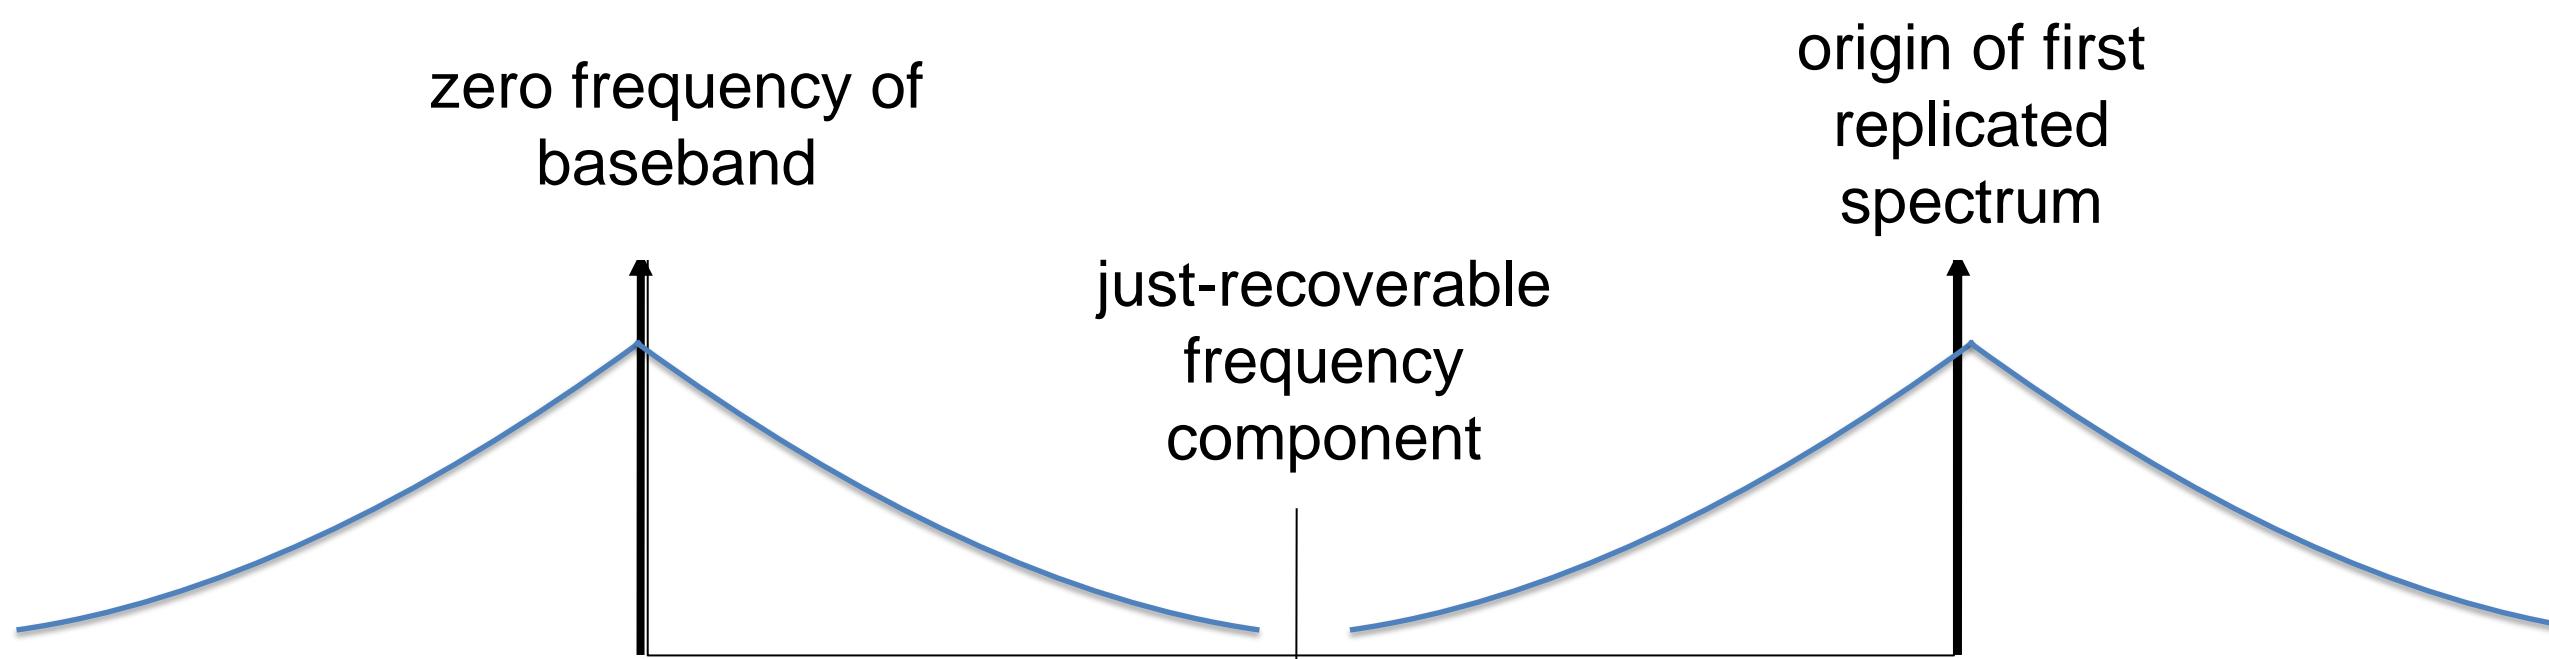
$\downarrow$



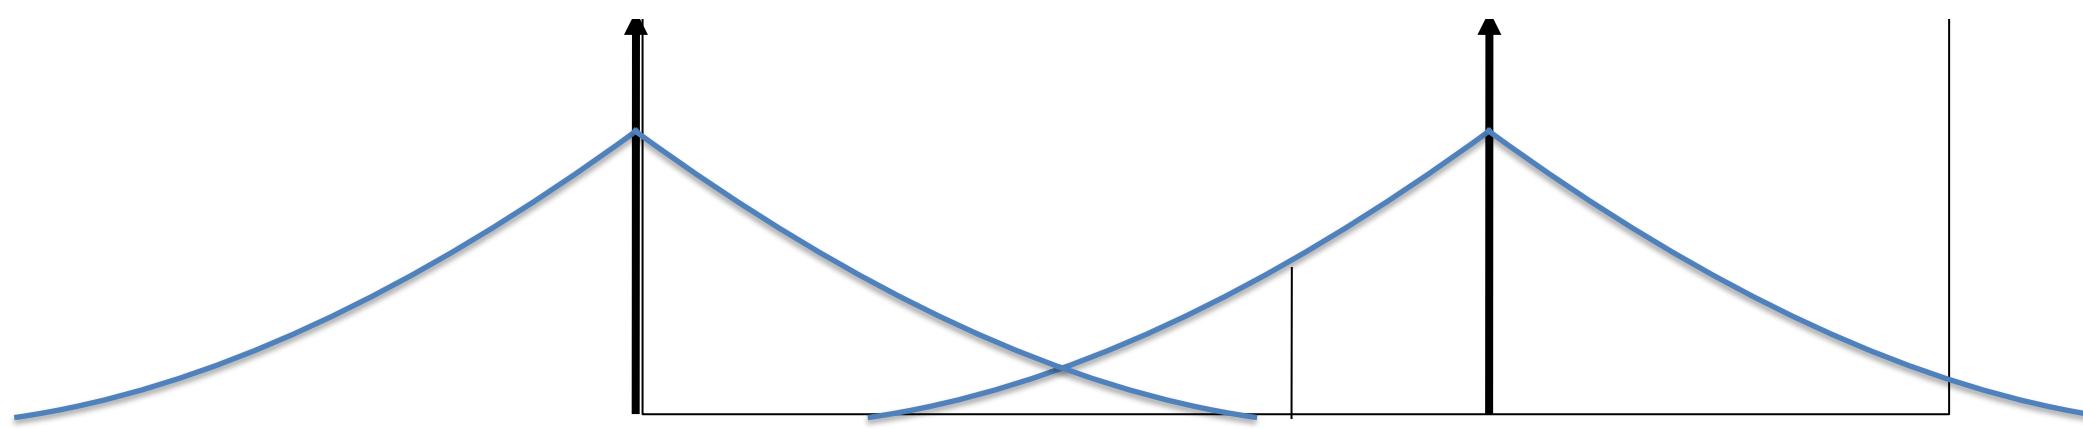
$s$



# aliasing



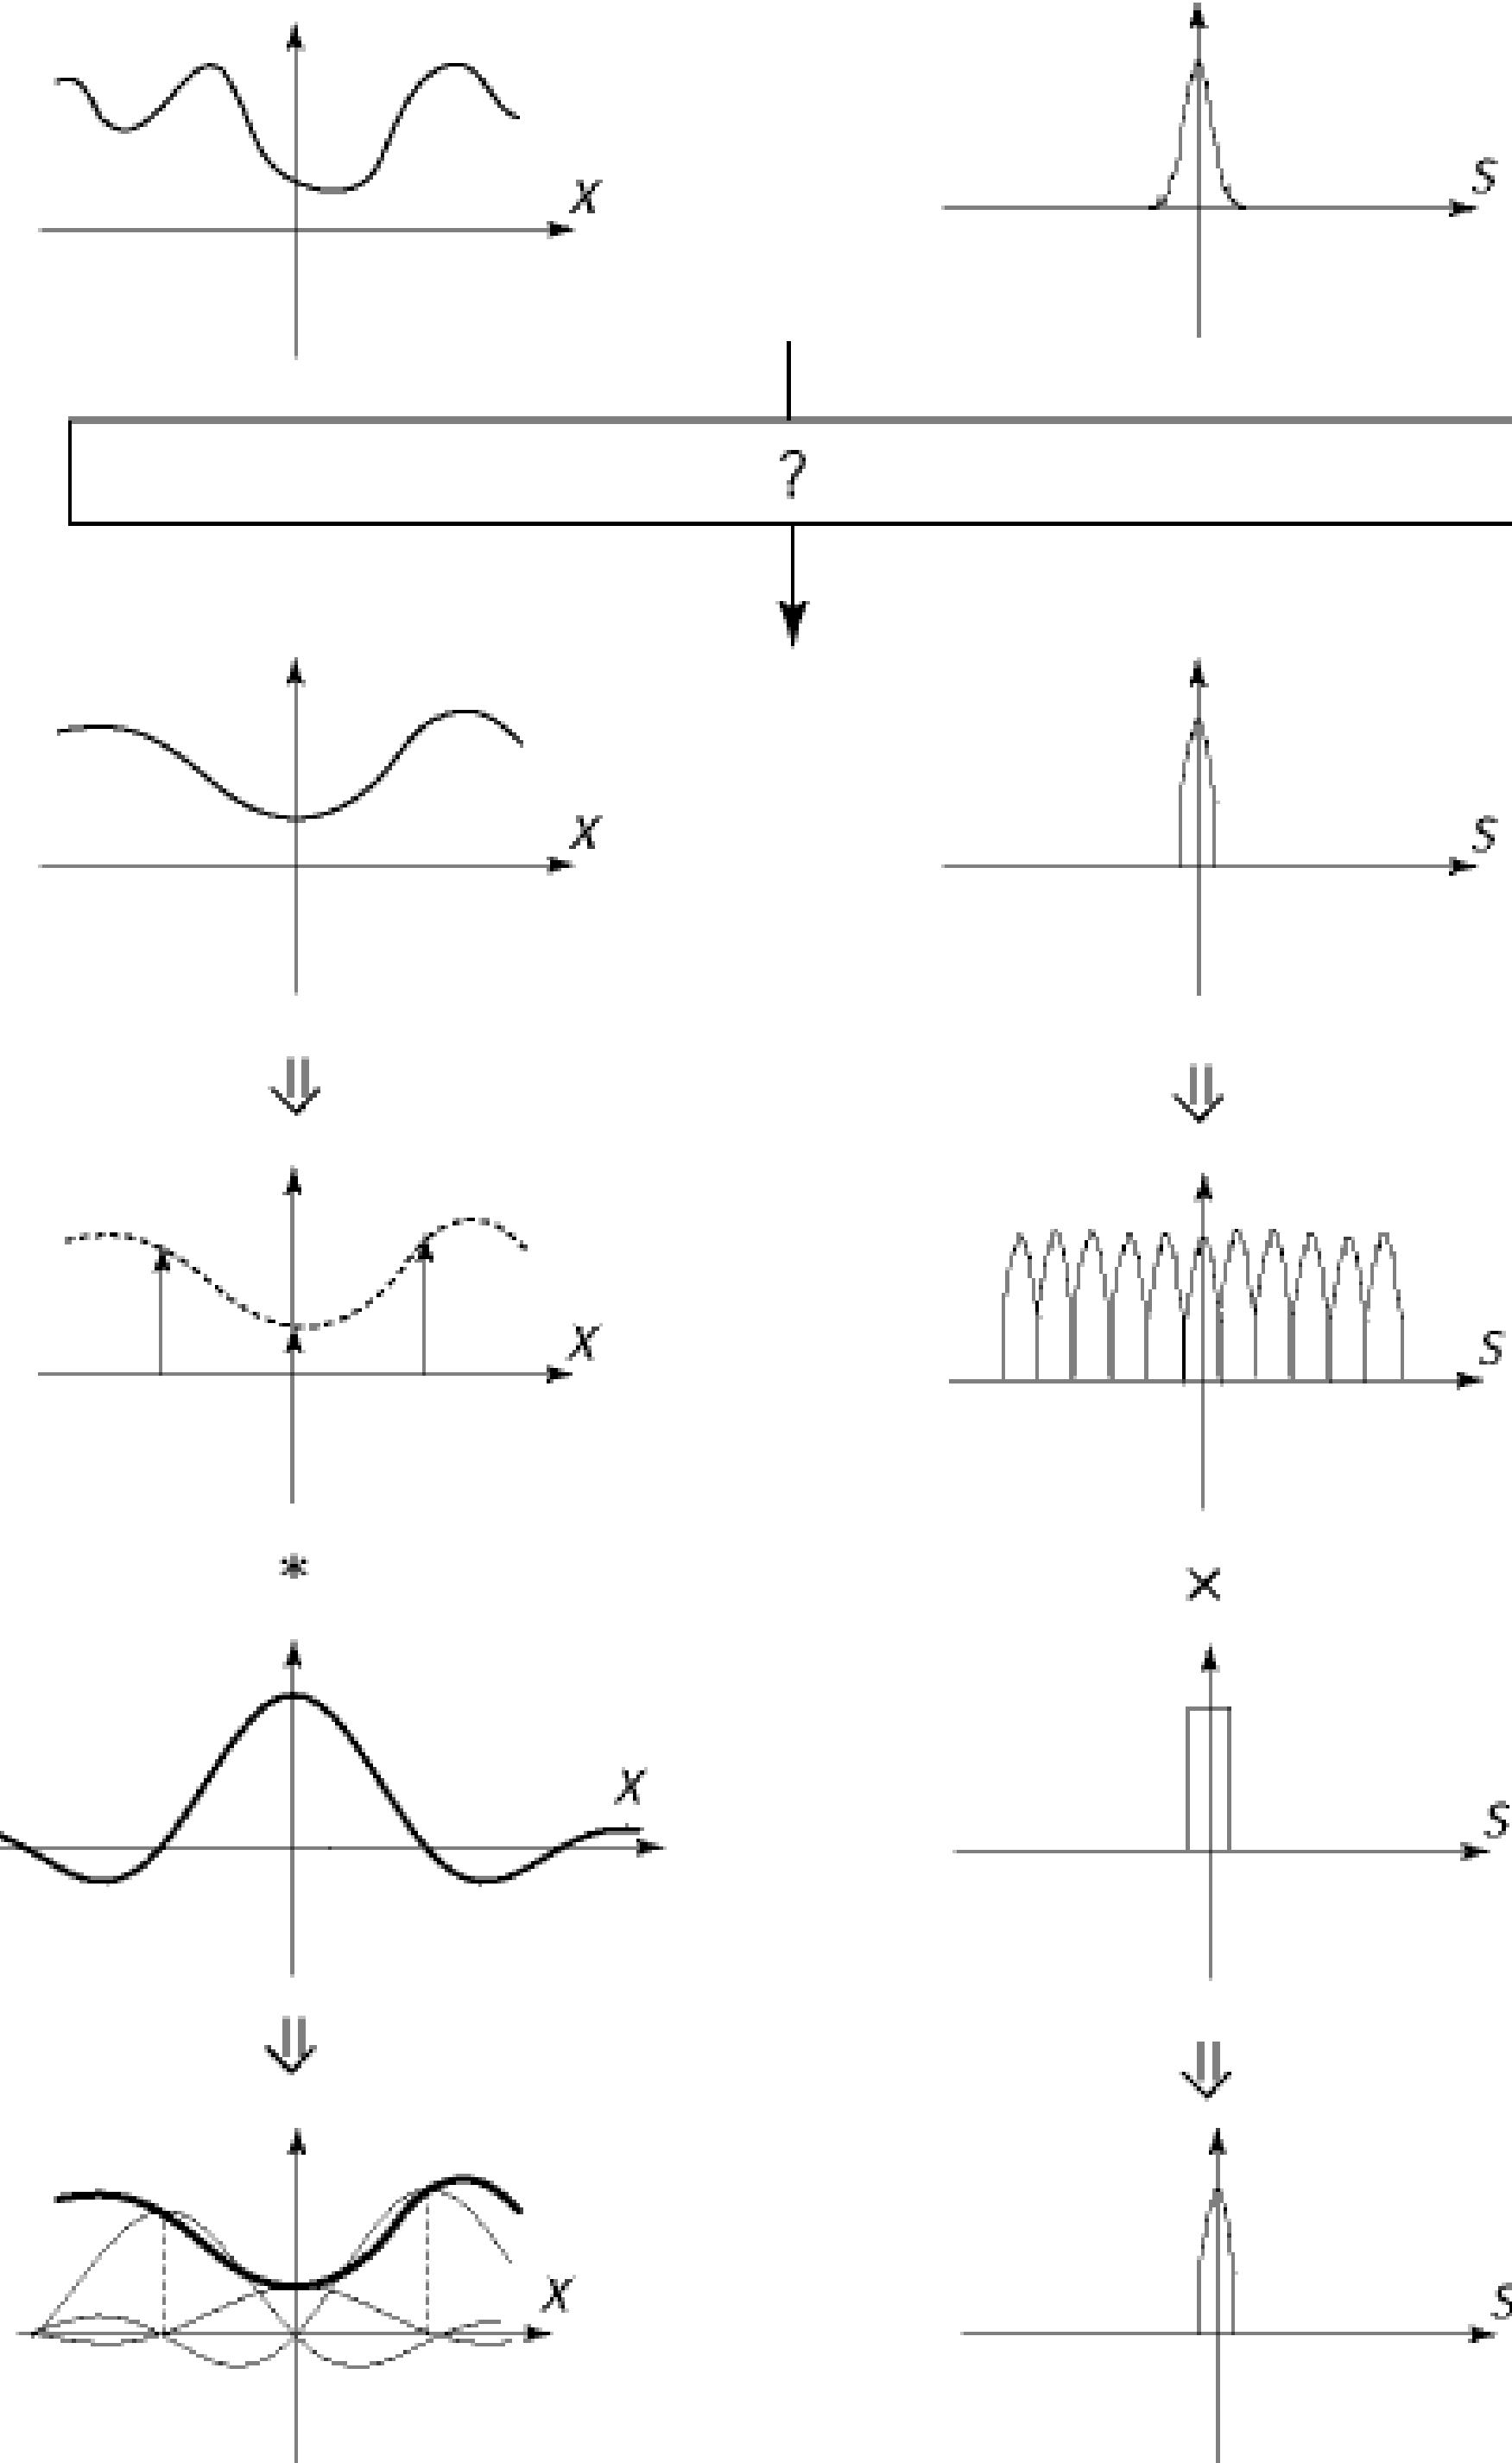
**no aliasing**



**aliasing**

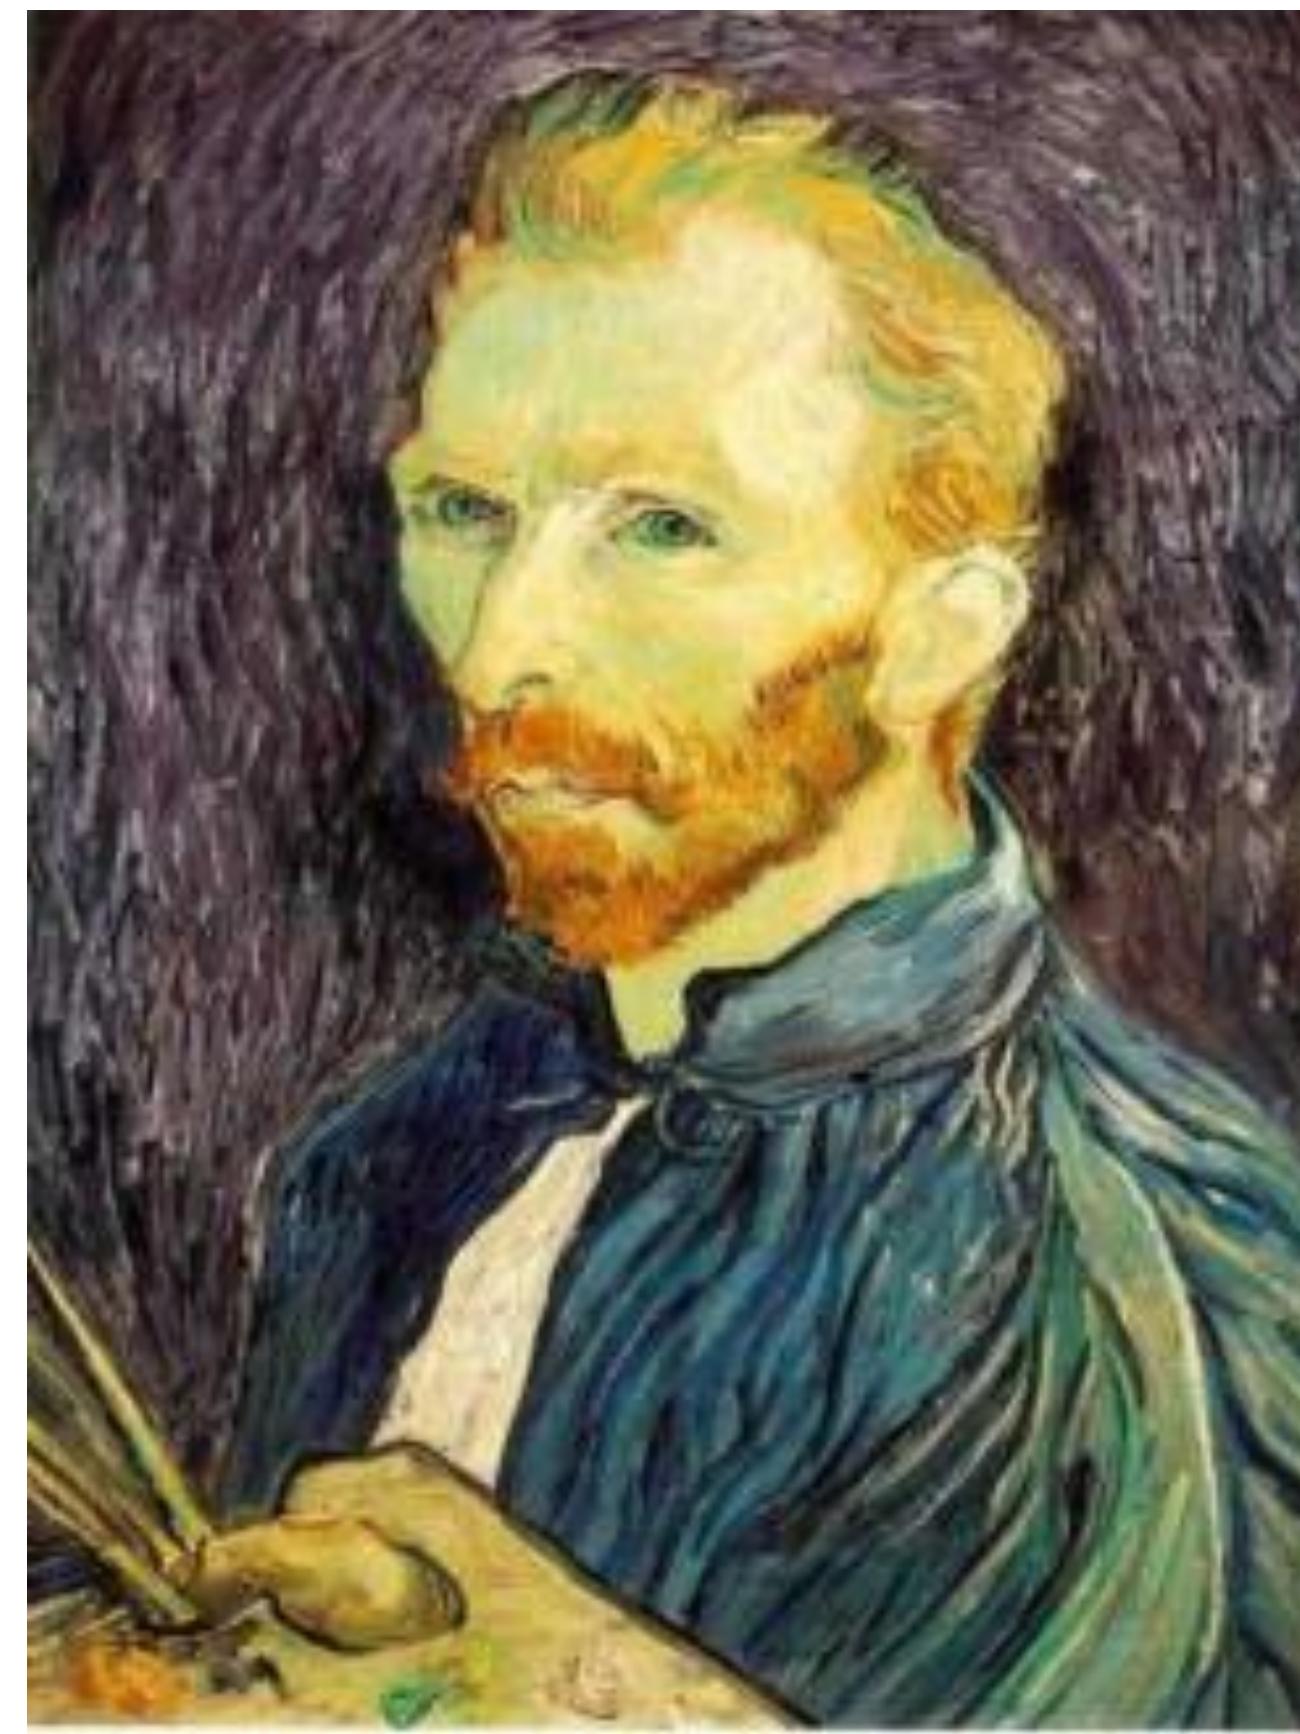
## Anti-aliasing by **pre-filtering**

- theoretical ideal pre-filter is a sinc function
- Gaussian, cubic filters work better in practice



# Subsampling with Gaussian pre-filtering

---



Gaussian 1/2



G 1/4



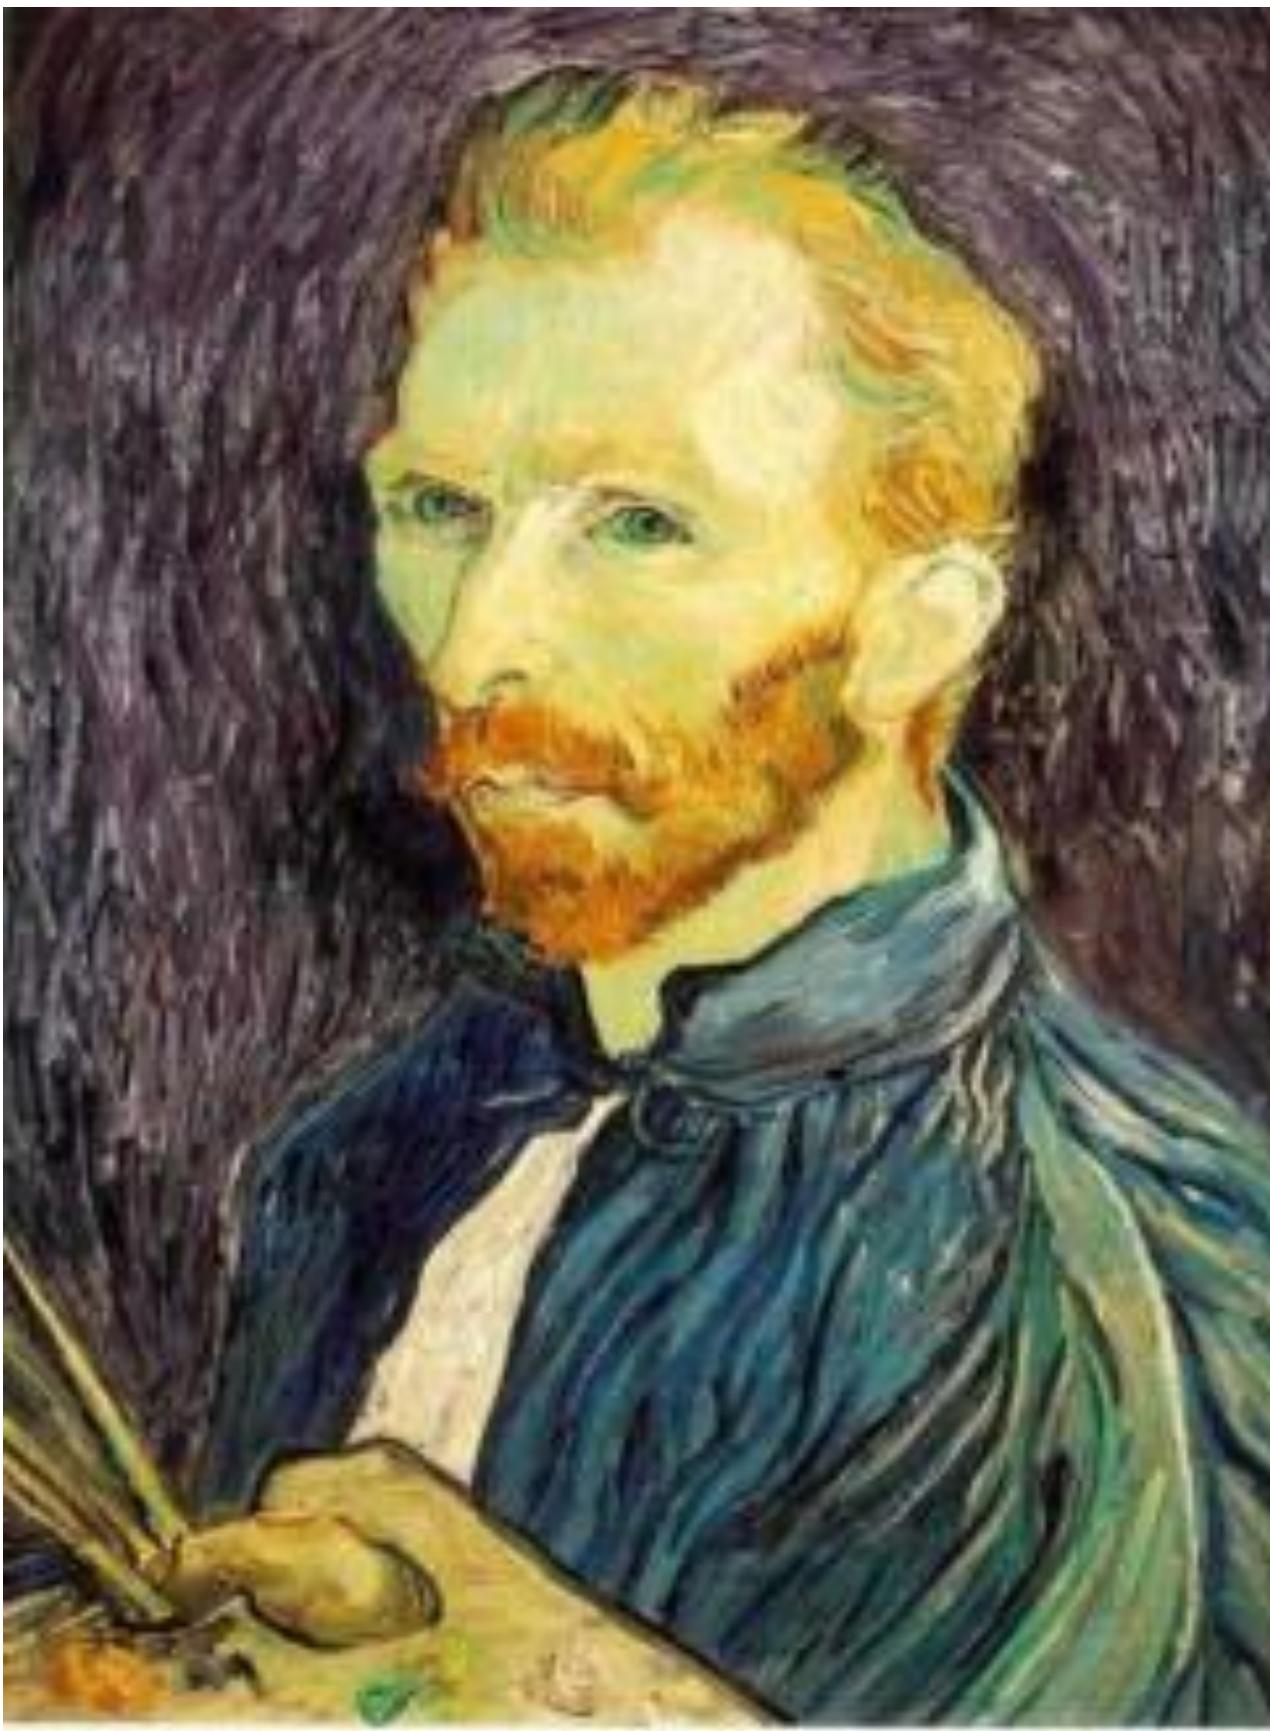
G 1/8

Solution: filter the image, *then* subsample

- Filter size should double for each  $\frac{1}{2}$  size reduction. Why?

# Subsampling with Gaussian pre-filtering

---



Gaussian 1/2



G 1/4



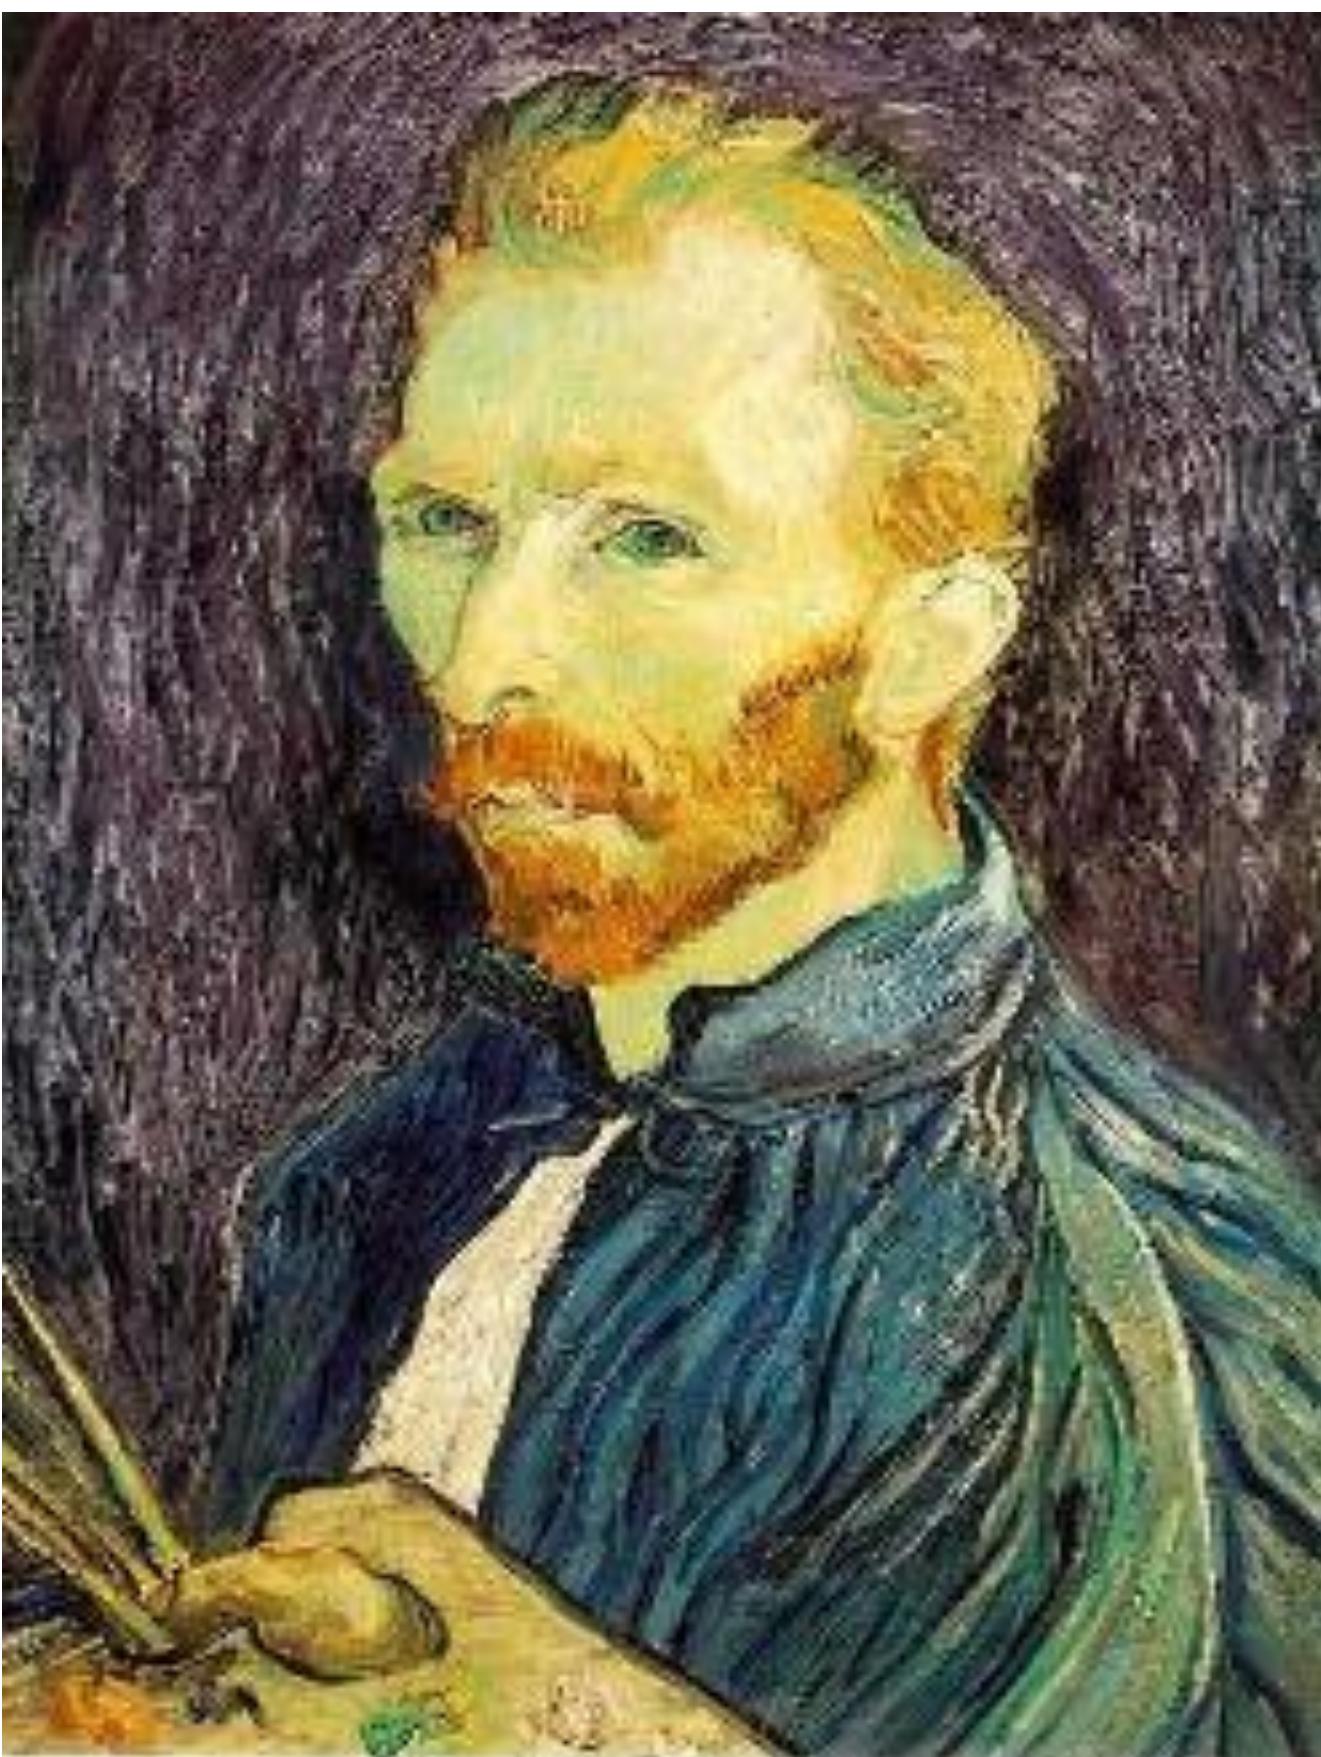
G 1/8

Solution: filter the image, *then* subsample

- Filter size should double for each  $\frac{1}{2}$  size reduction. Why?
- How can we speed this up?

# Compare with...

---



1/2



1/4 (2x zoom)



1/8 (4x zoom)

Why does this look so cruffy?

# Antialising filtering

Before sampling, apply a low pass-filter to remove all the frequencies that will produce aliasing.

Without antialising filter.

103×128



52×64



26×32



With antialising filter.

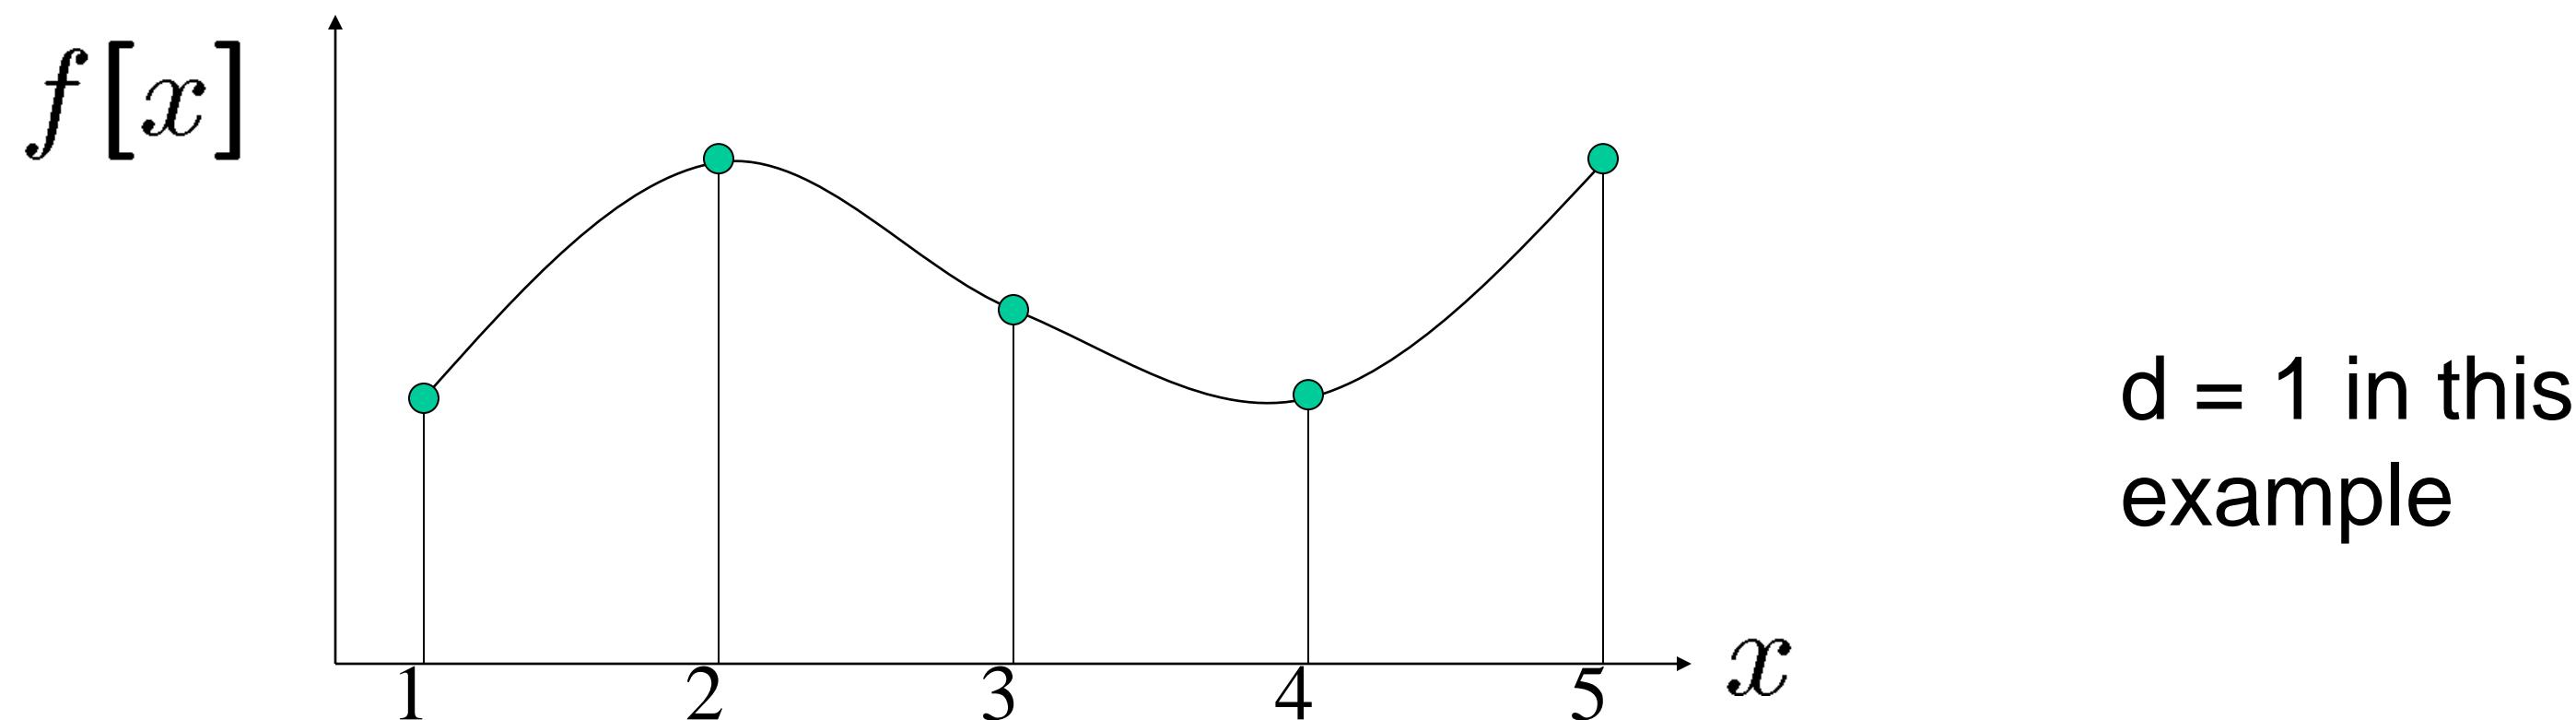


# Image resampling

---

So far, we considered only power-of-two subsampling

- What about arbitrary scale reduction?
- How can we increase the size of the image?

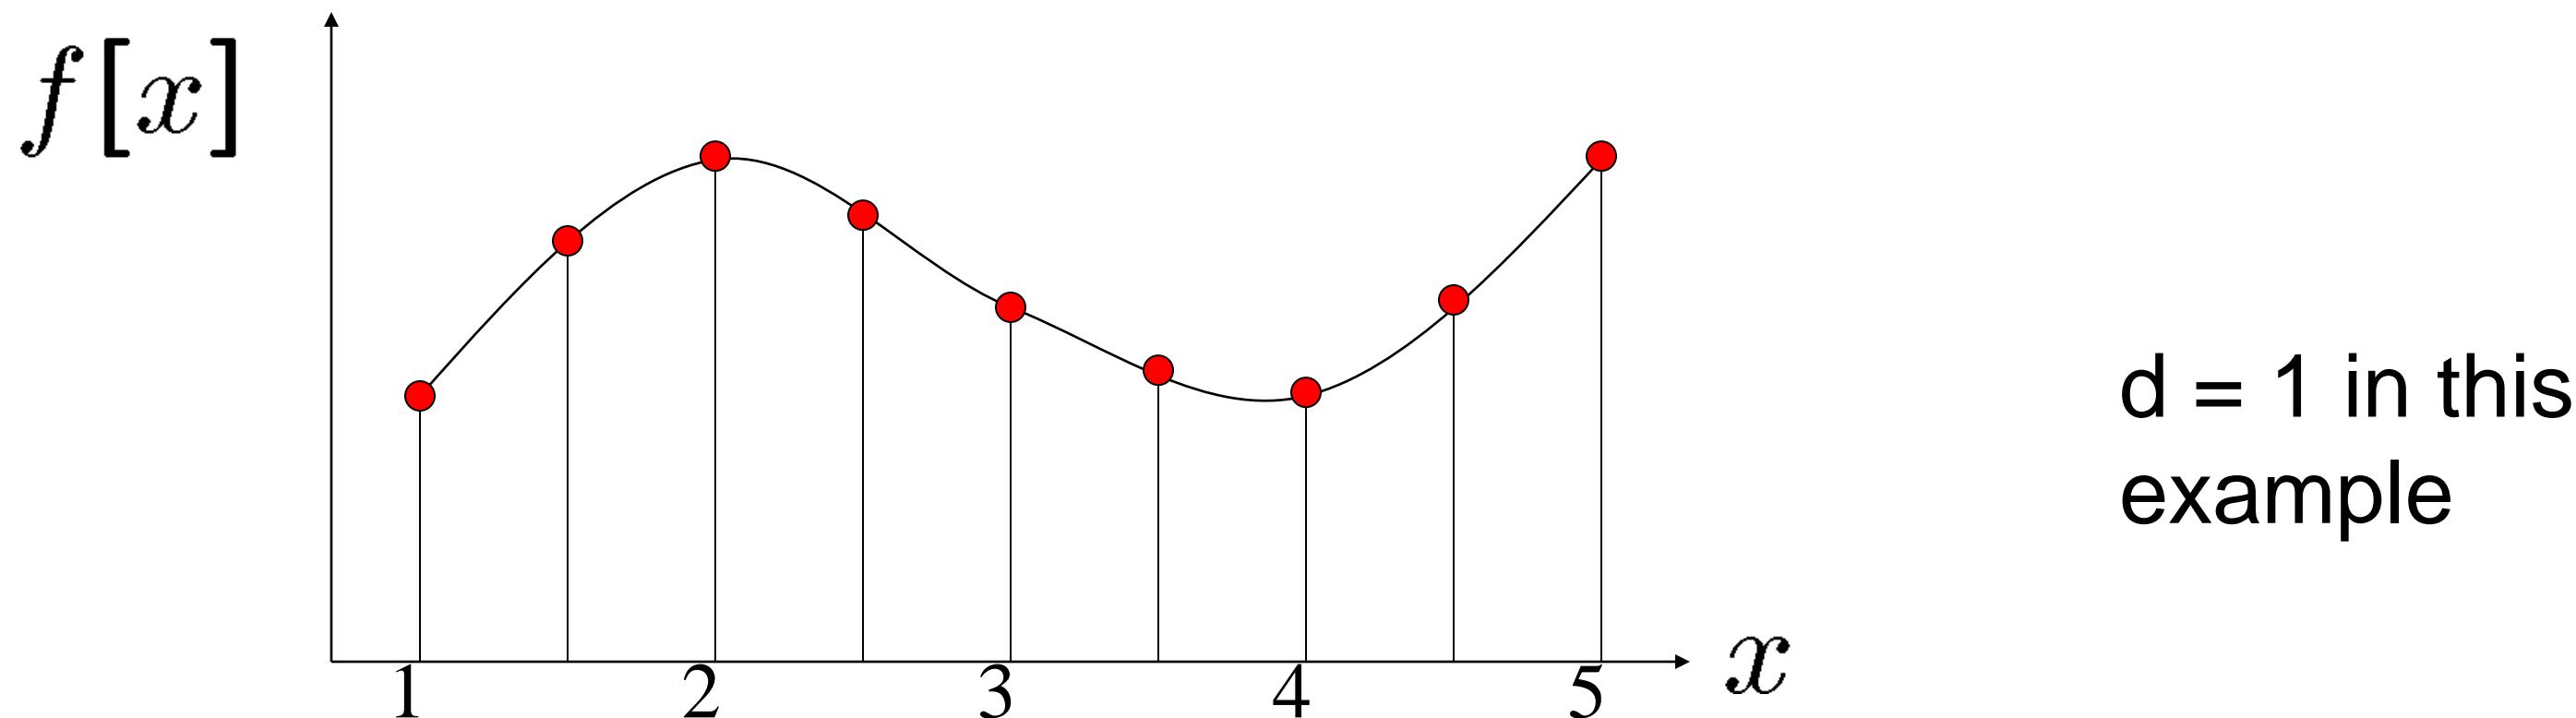


# Image resampling

---

So far, we considered only power-of-two subsampling

- What about arbitrary scale reduction?
- How can we increase the size of the image?



Recall how a digital image is formed

$$f[x, y] = \text{quantize}\{f(xd, yd)\}$$

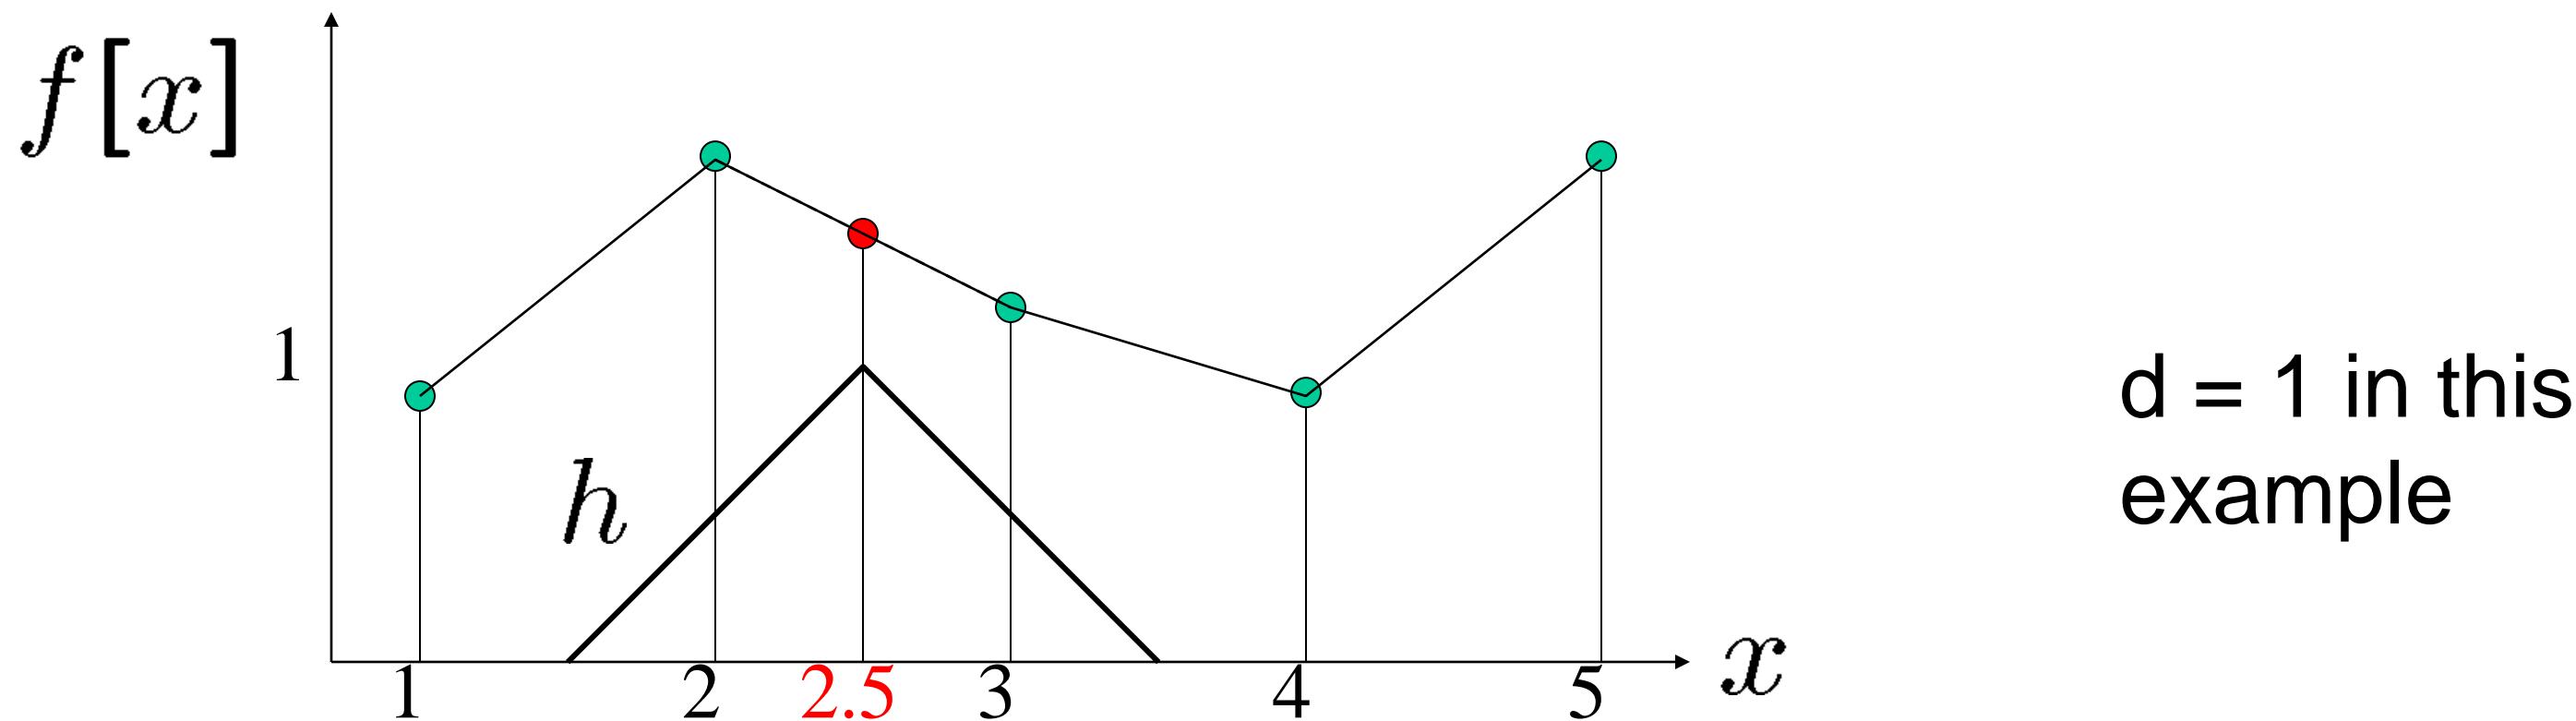
- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale

# Image resampling

---

So what to do if we don't know  $f$

- Answer: guess an approximation  $\tilde{f}$
- Can be done in a principled way: filtering



$d = 1$  in this example

## Image reconstruction

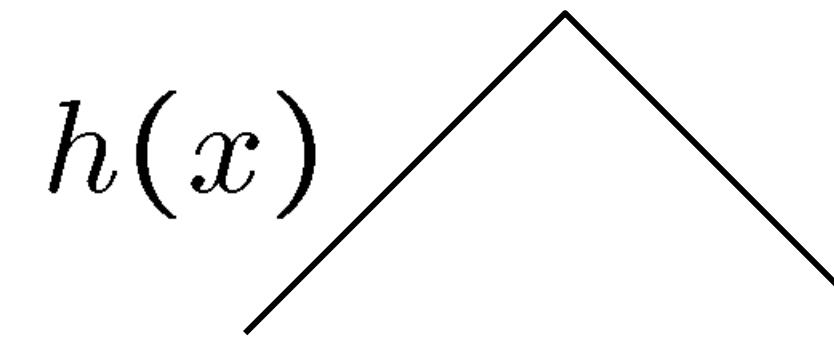
- Convert  $f$  to a continuous function
$$f(x) = f[\frac{x}{d}] \text{ when } \frac{x}{d} \text{ is an integer, 0 otherwise}$$
- Reconstruct by convolution:

$$\tilde{f} = h \star f$$

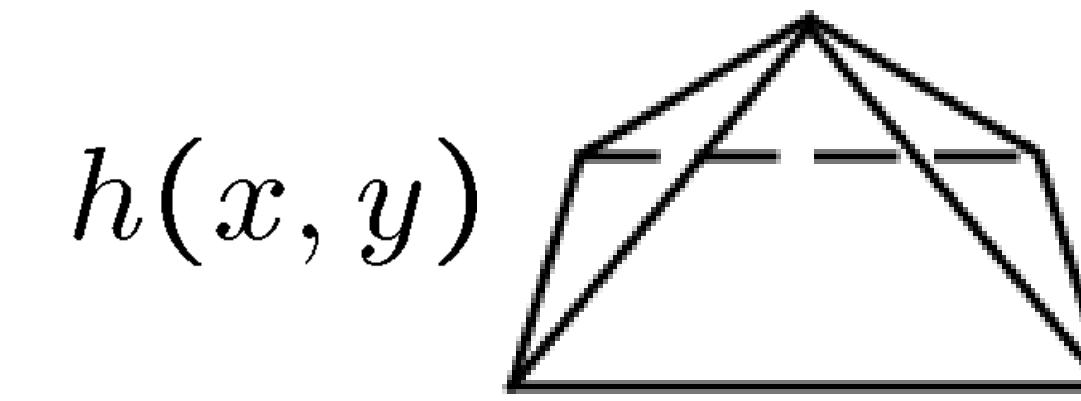
# Resampling filters

---

What does the 2D version of this hat function look like?



performs  
linear interpolation



(tent function) performs  
**bilinear interpolation**

Better filters give better resampled images

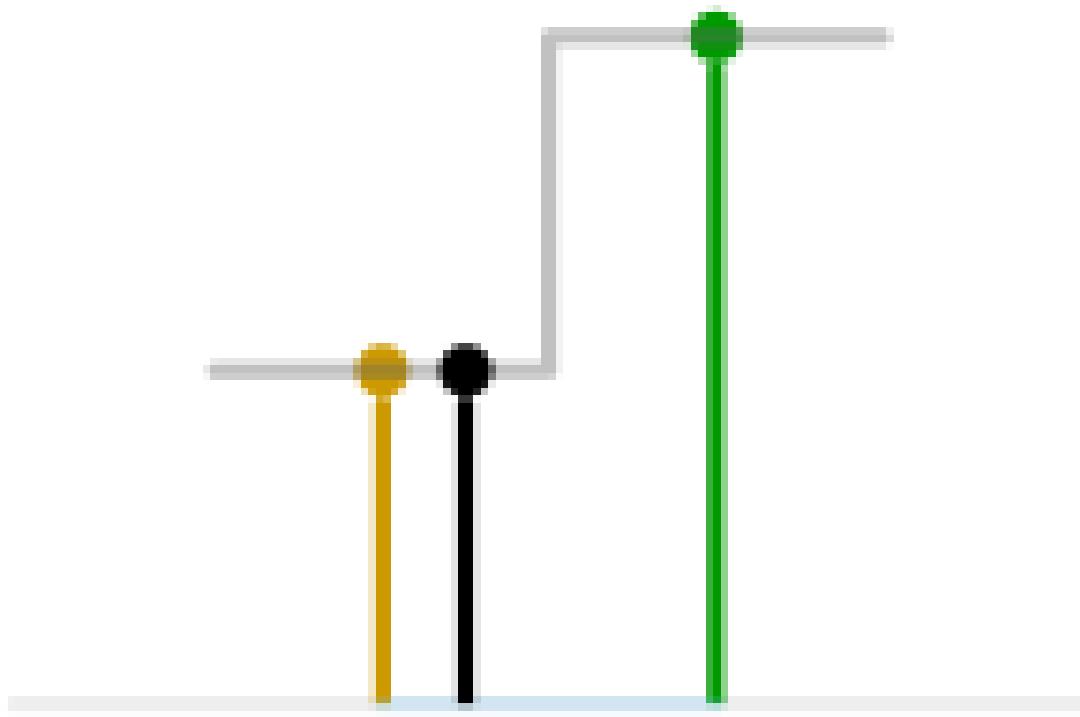
- Bicubic is common choice
  - fit 3<sup>rd</sup> degree polynomial surface to pixels in neighborhood
  - can also be implemented by a convolution

Steve Seitz's 5-min video:

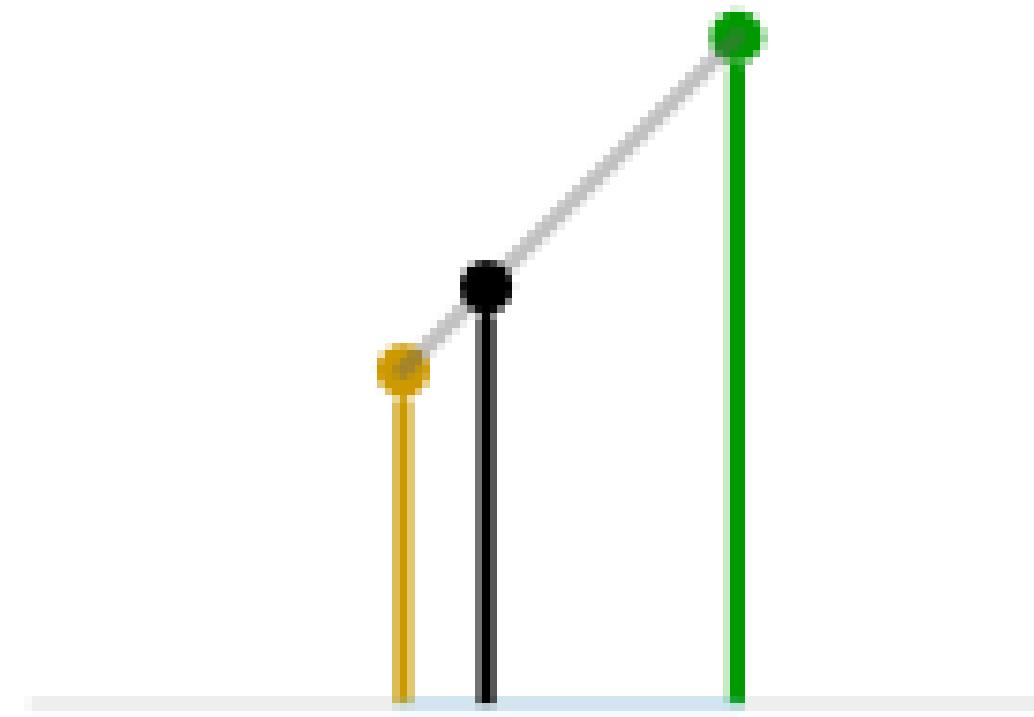
<https://youtu.be/Xj129kA3Ci0?si=cAf5isBb0GzdLi-7>

# Types of Interpolation

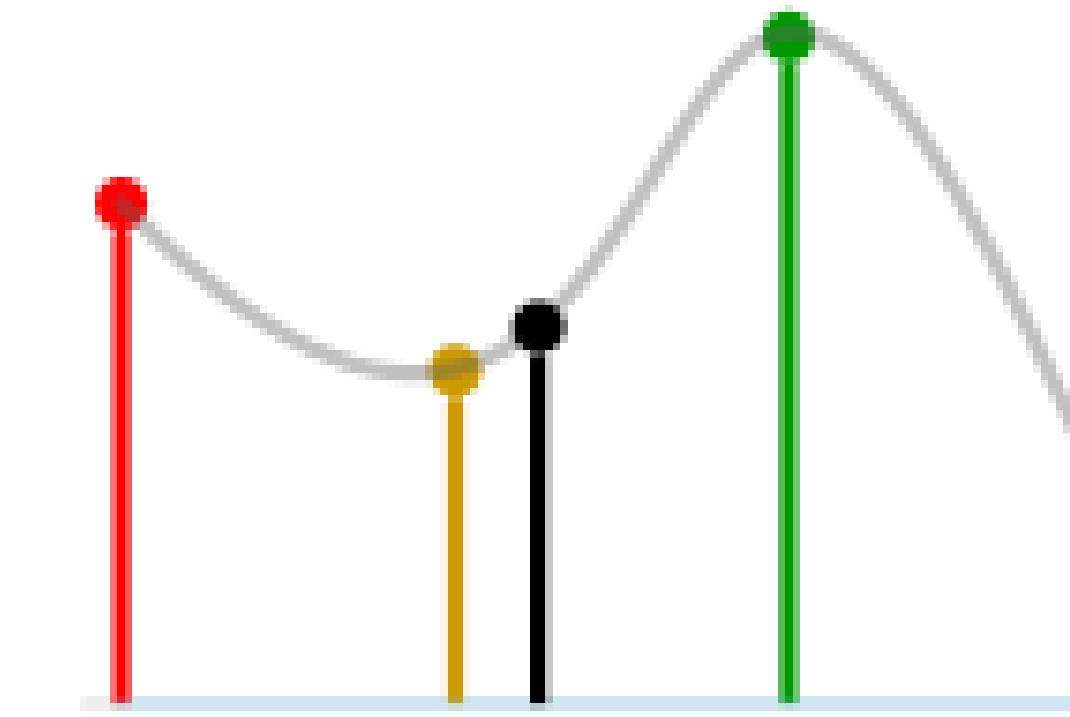
---



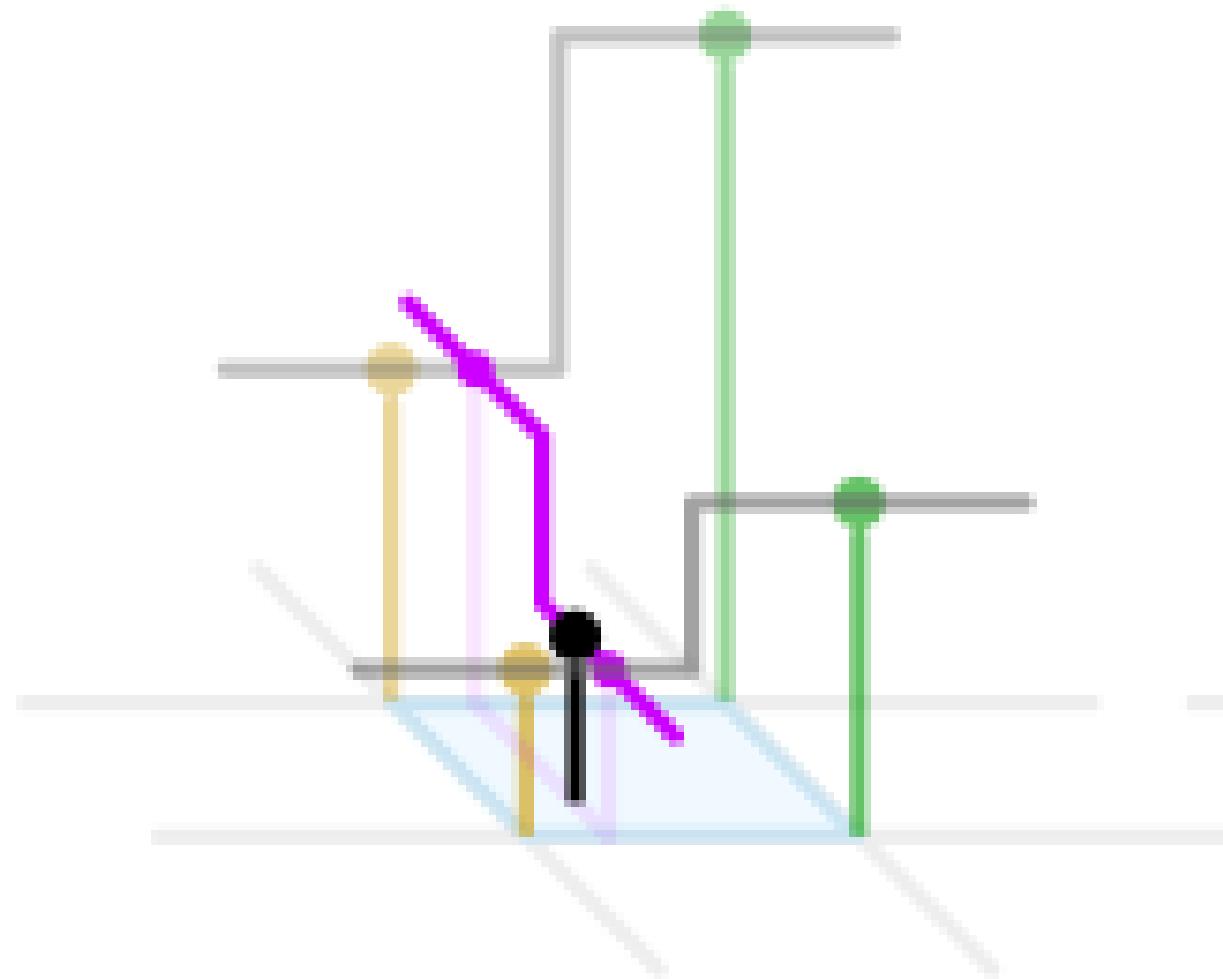
1D nearest-neighbour



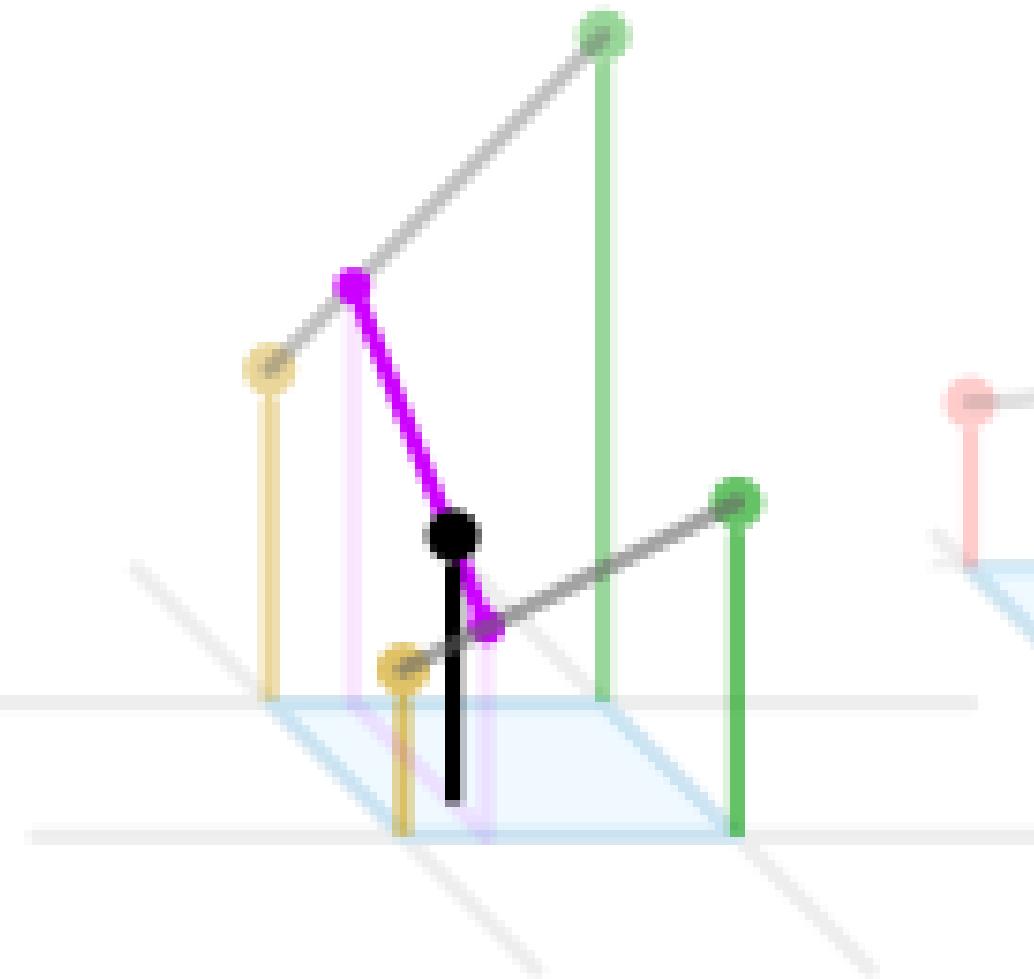
Linear



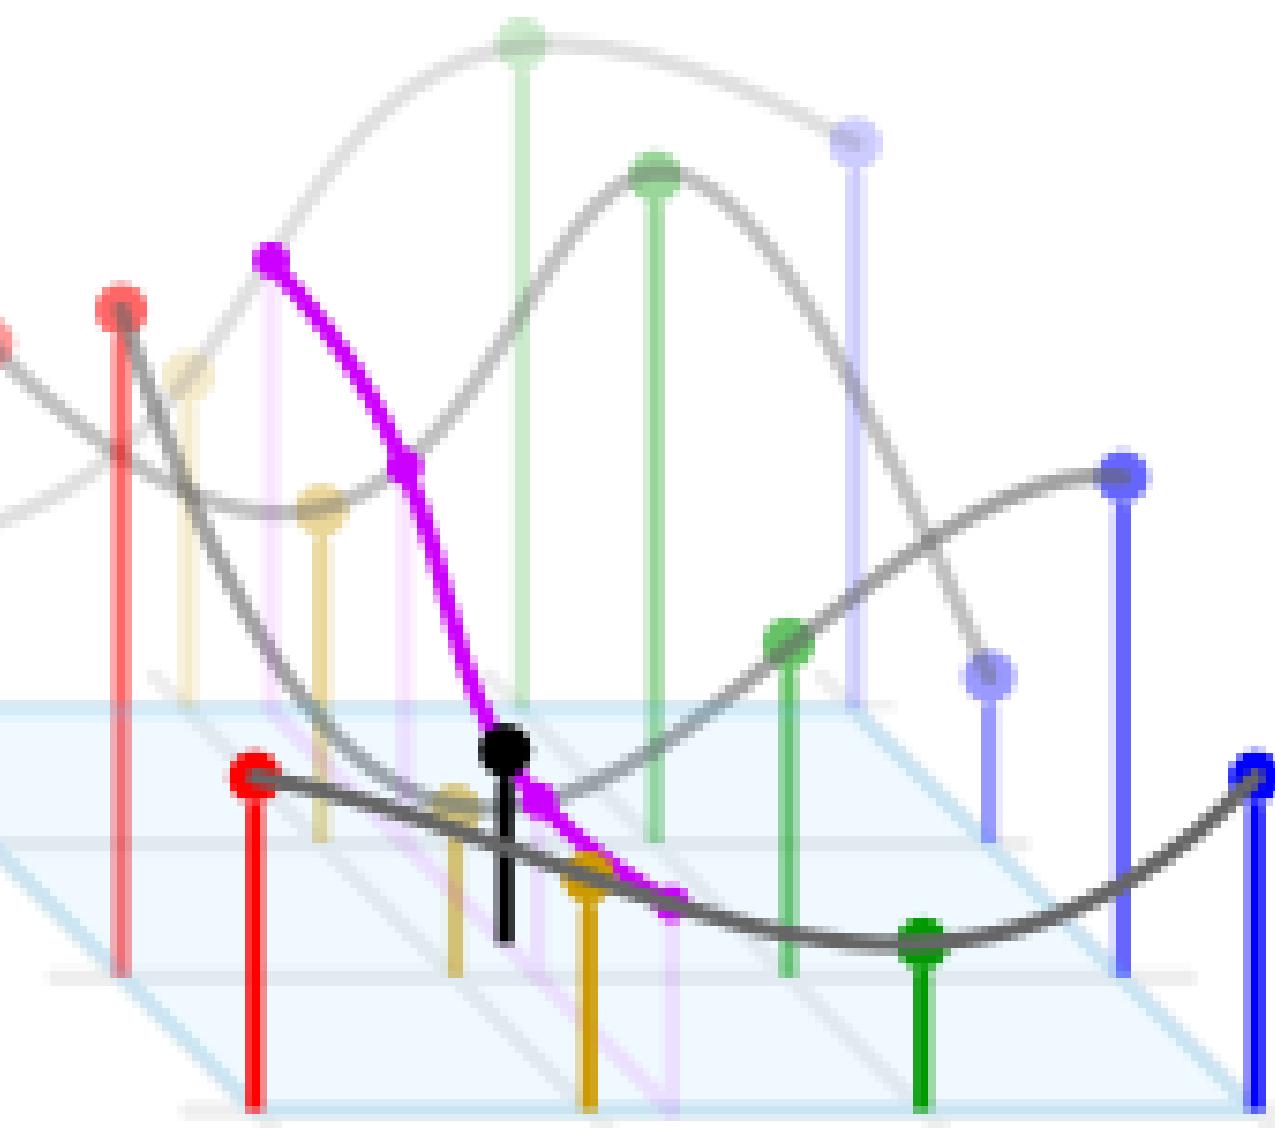
Cubic



2D nearest-neighbour

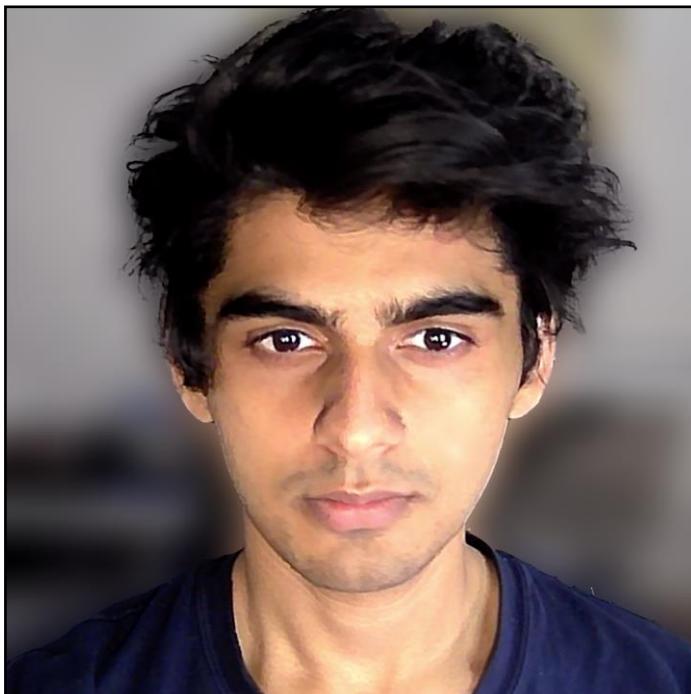


Bilinear



Bicubic

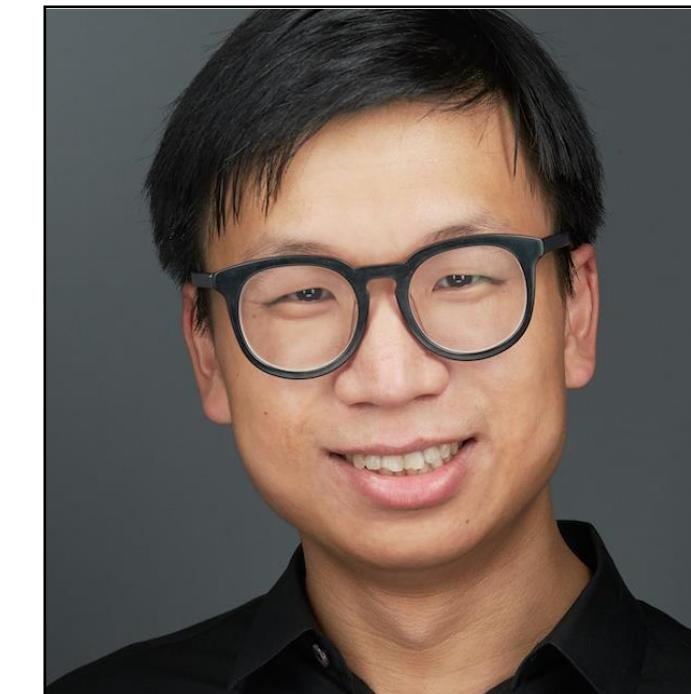
# On Aliased Resizing and Surprising Subtleties in GAN Evaluation



Gaurav Parmar

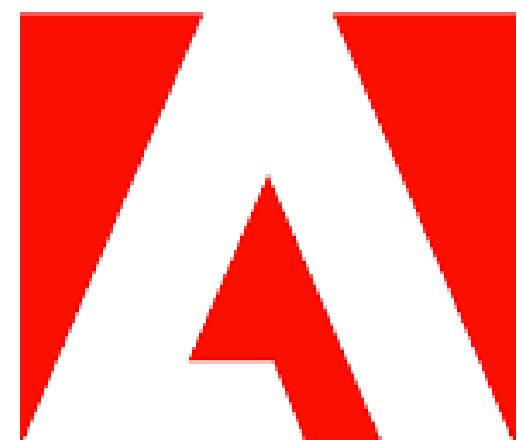


Richard Zhang



Jun-Yan Zhu

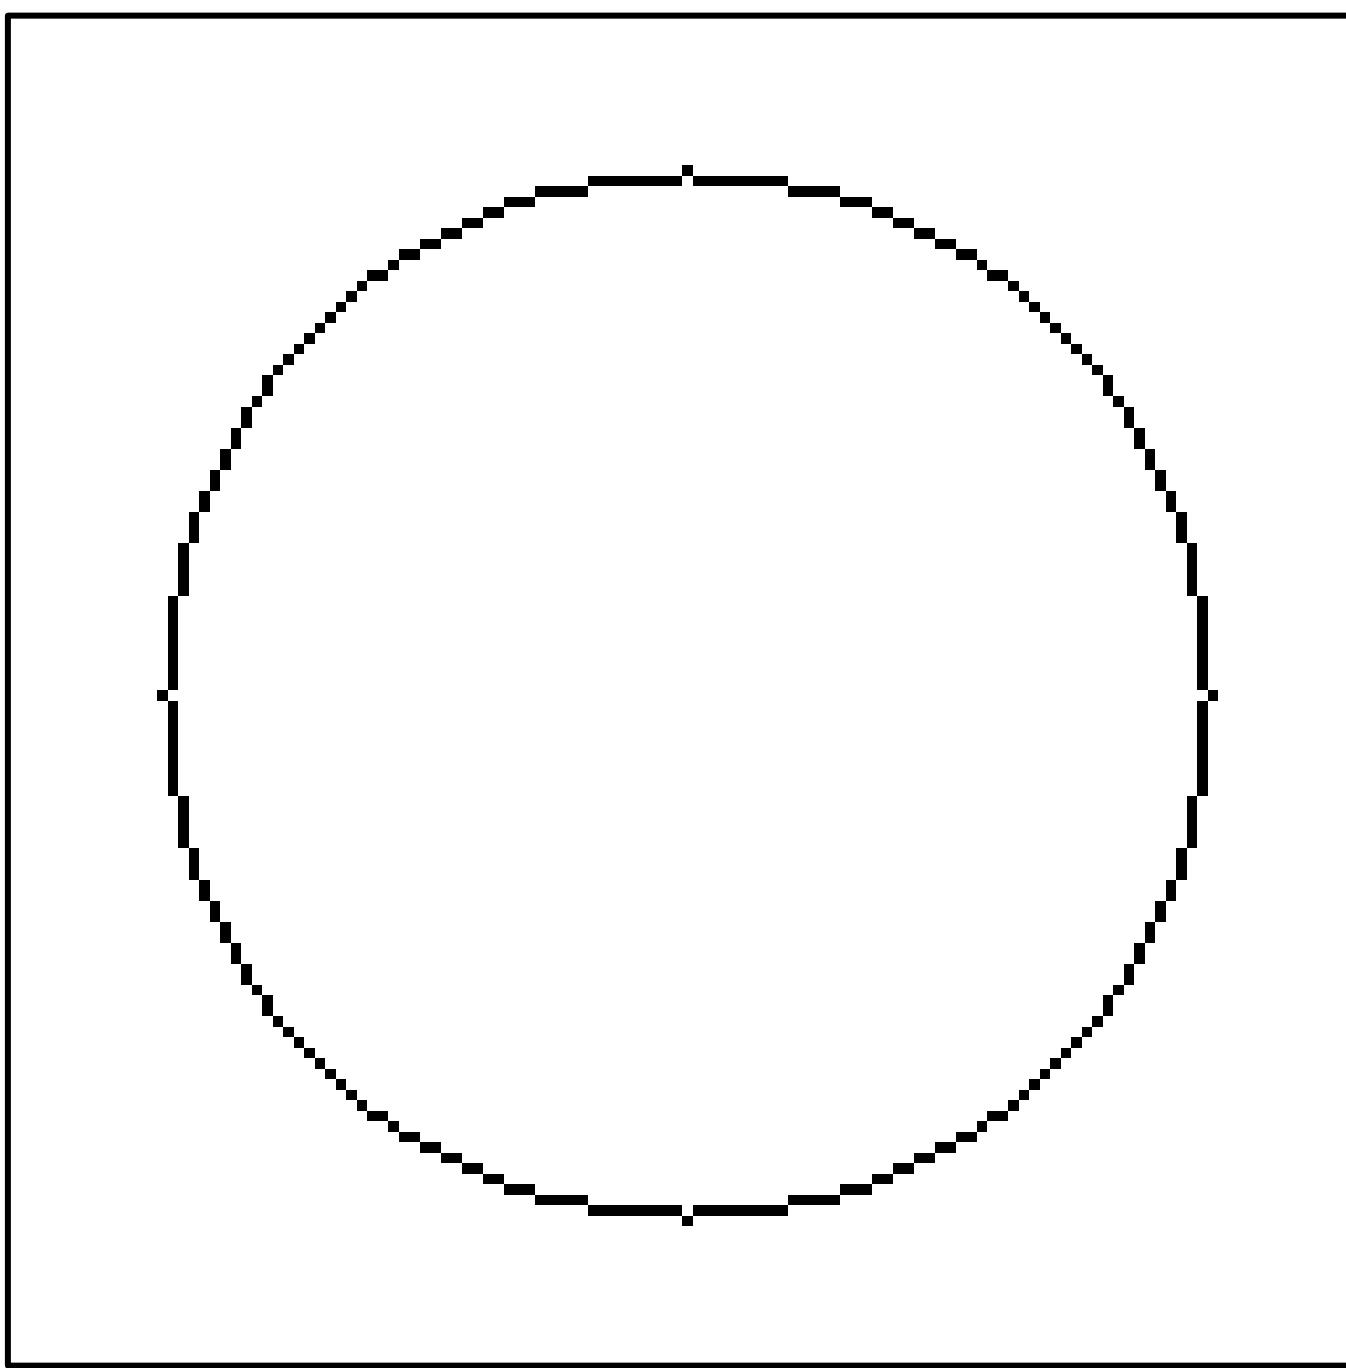
Carnegie  
Mellon  
University



# Downsampling a circle

---

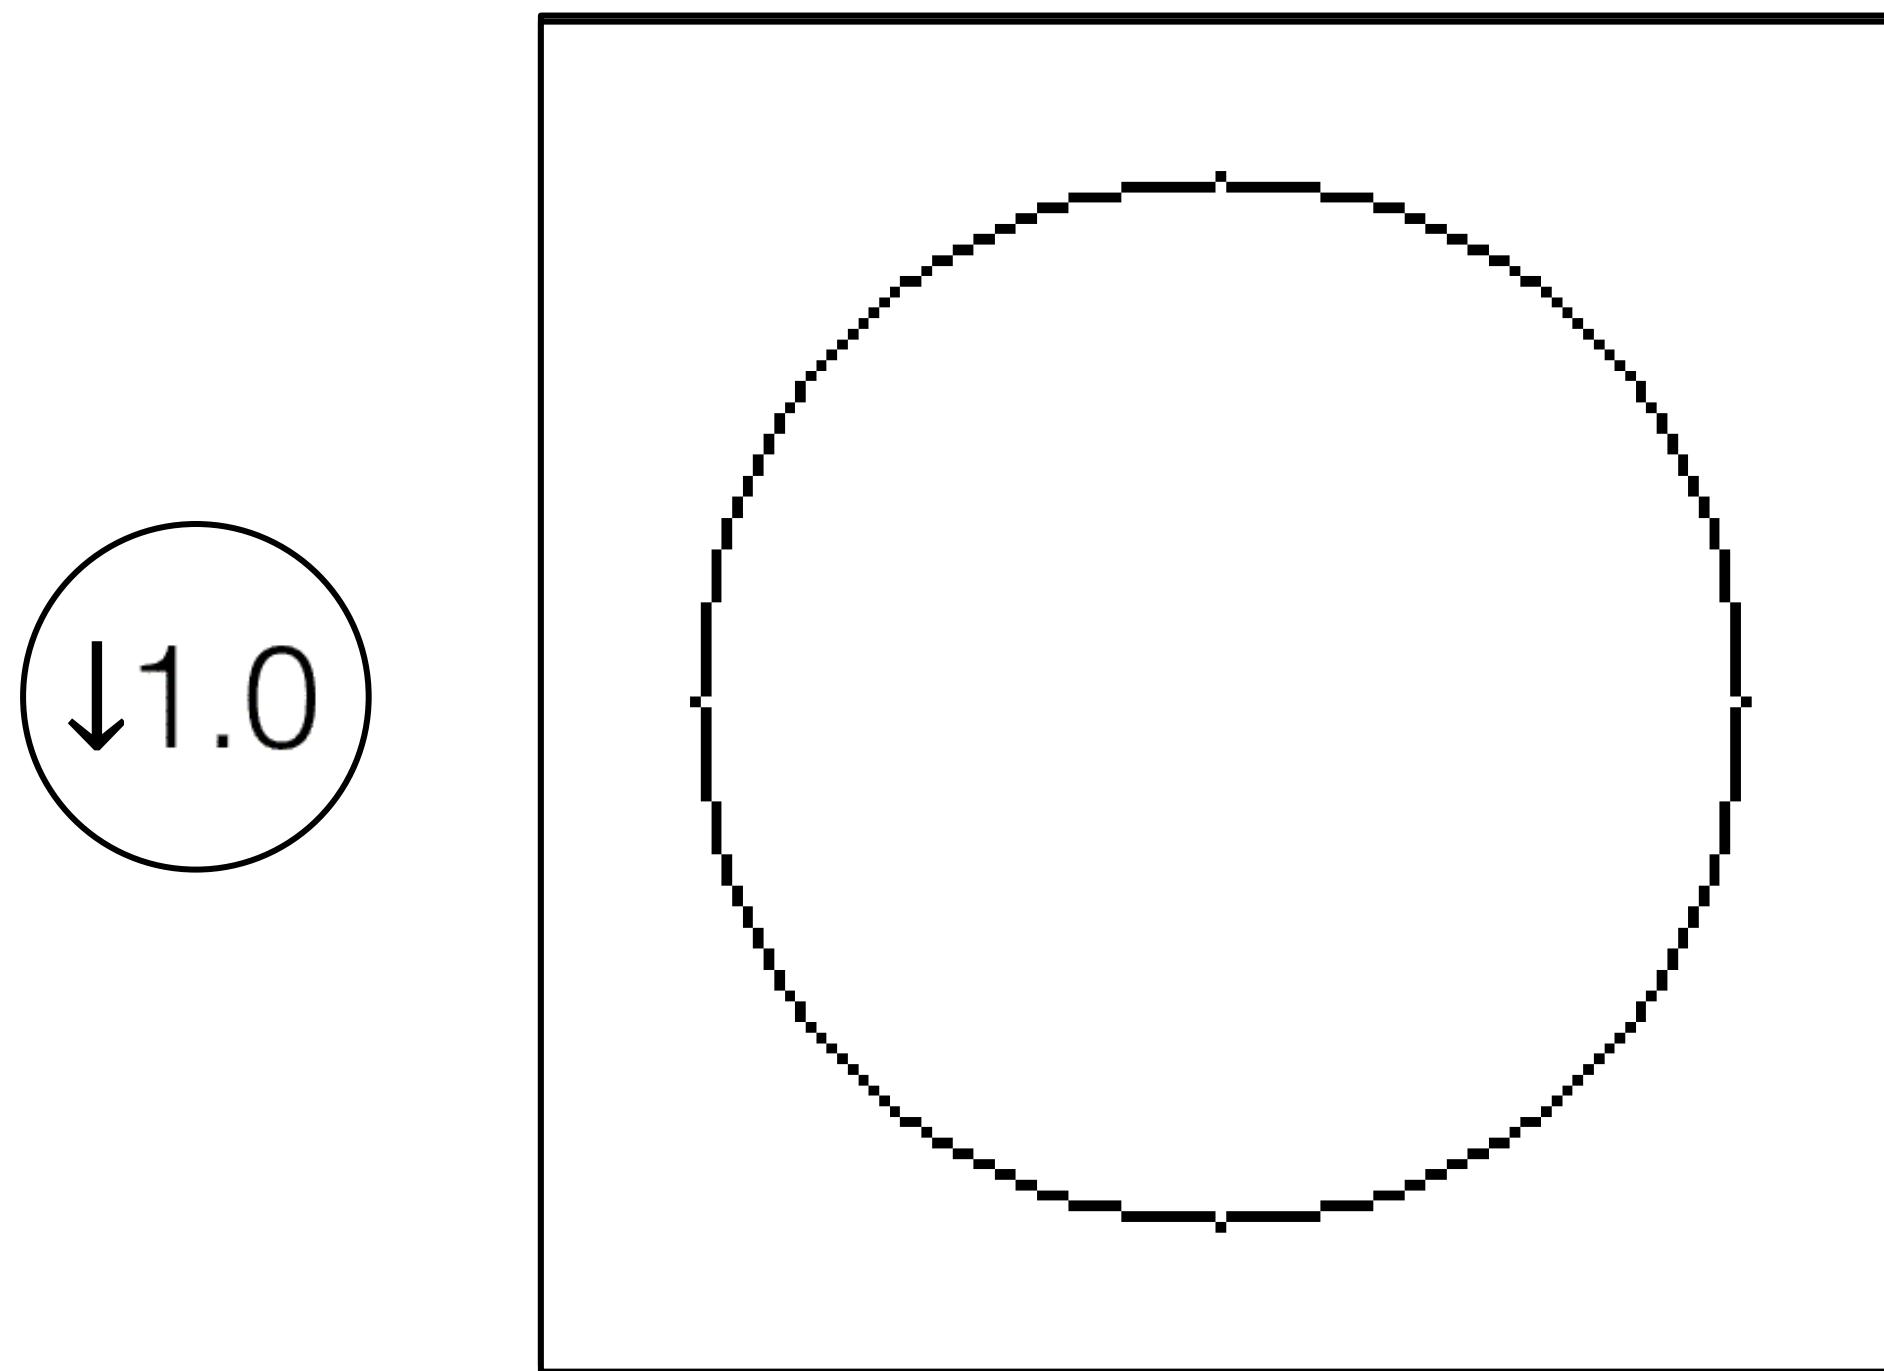
input image



# Downsampling a circle

---

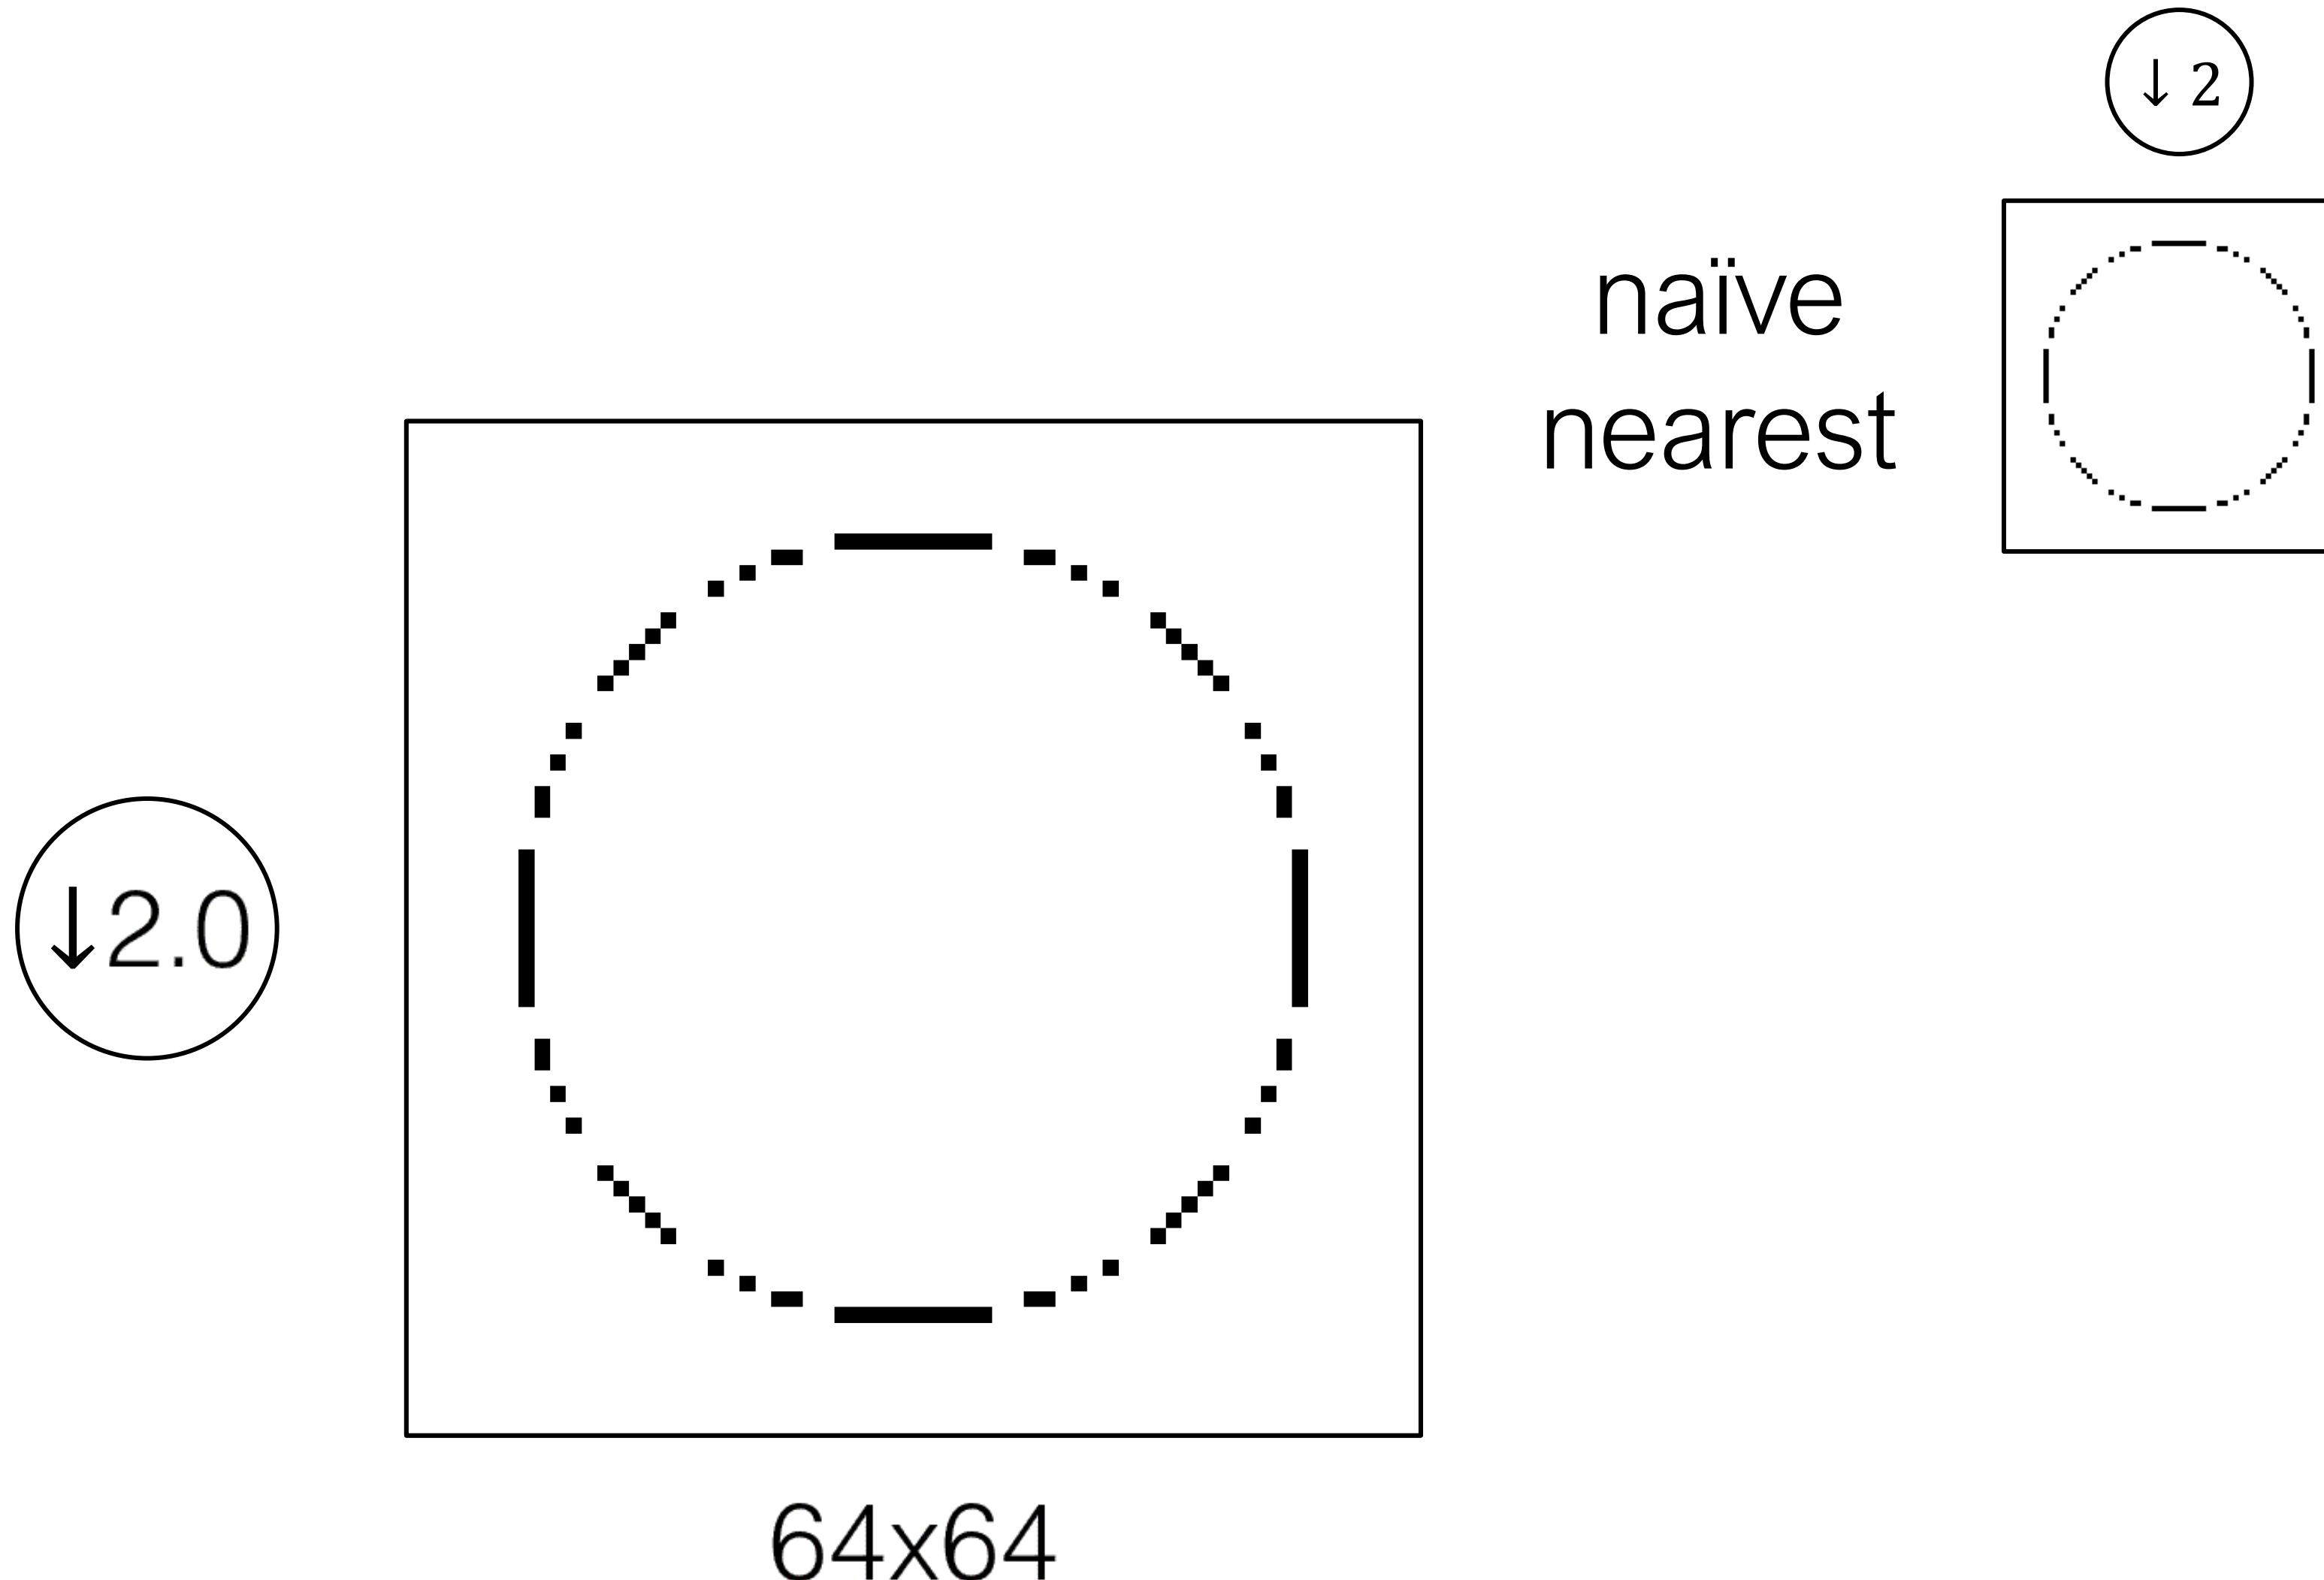
naïve nearest



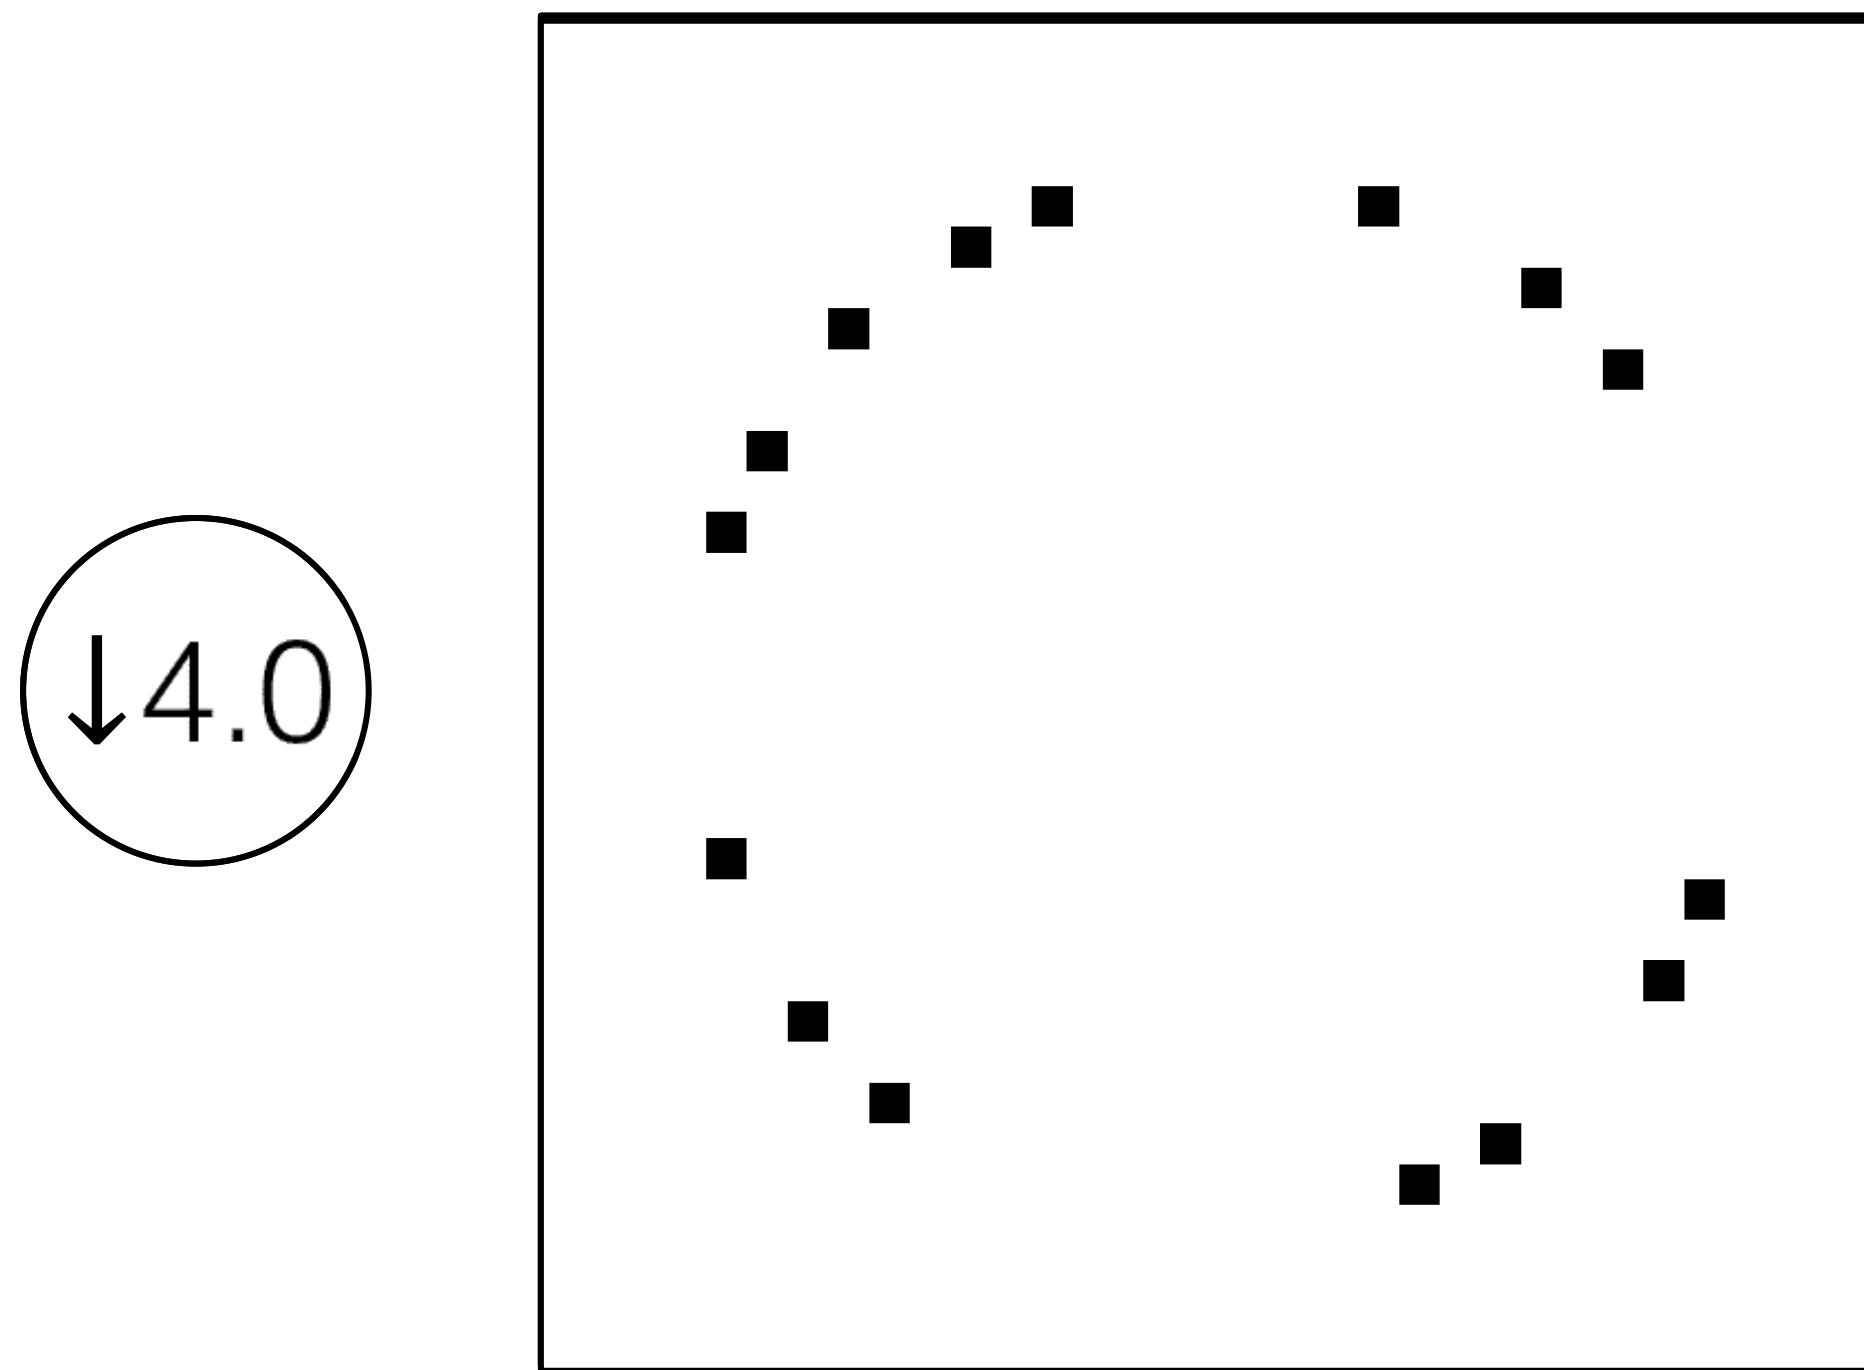
↓1.0

128x128

# Downsampling a circle

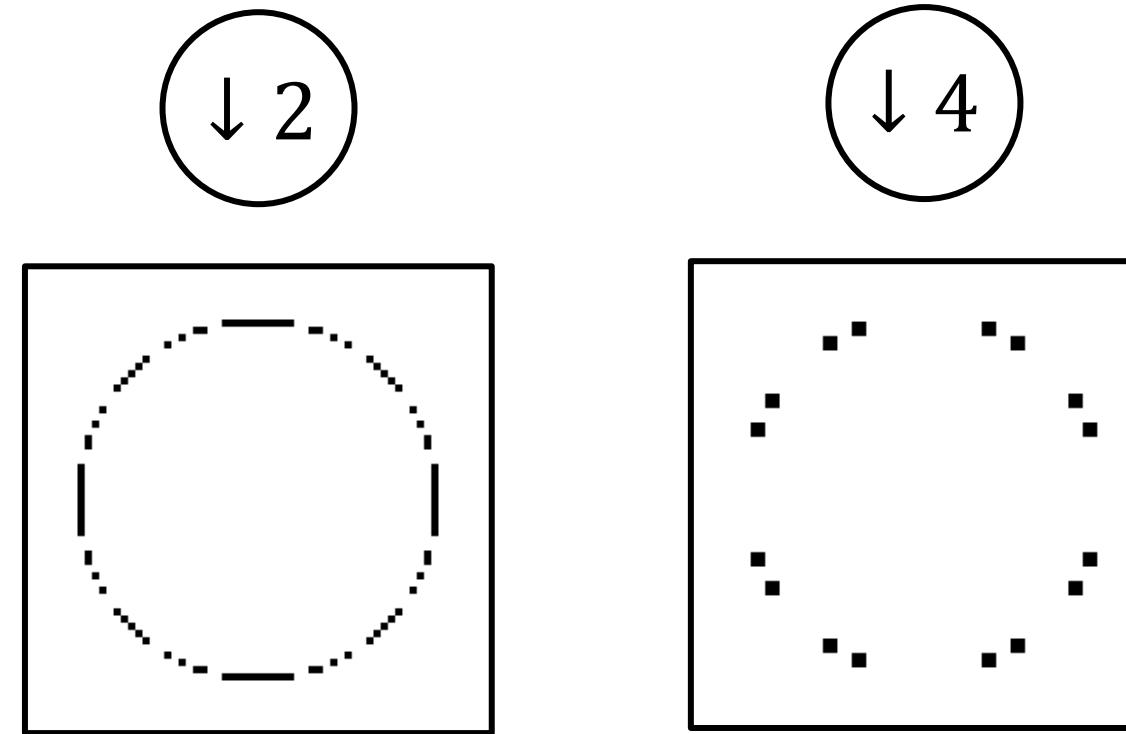


# Downsampling a circle

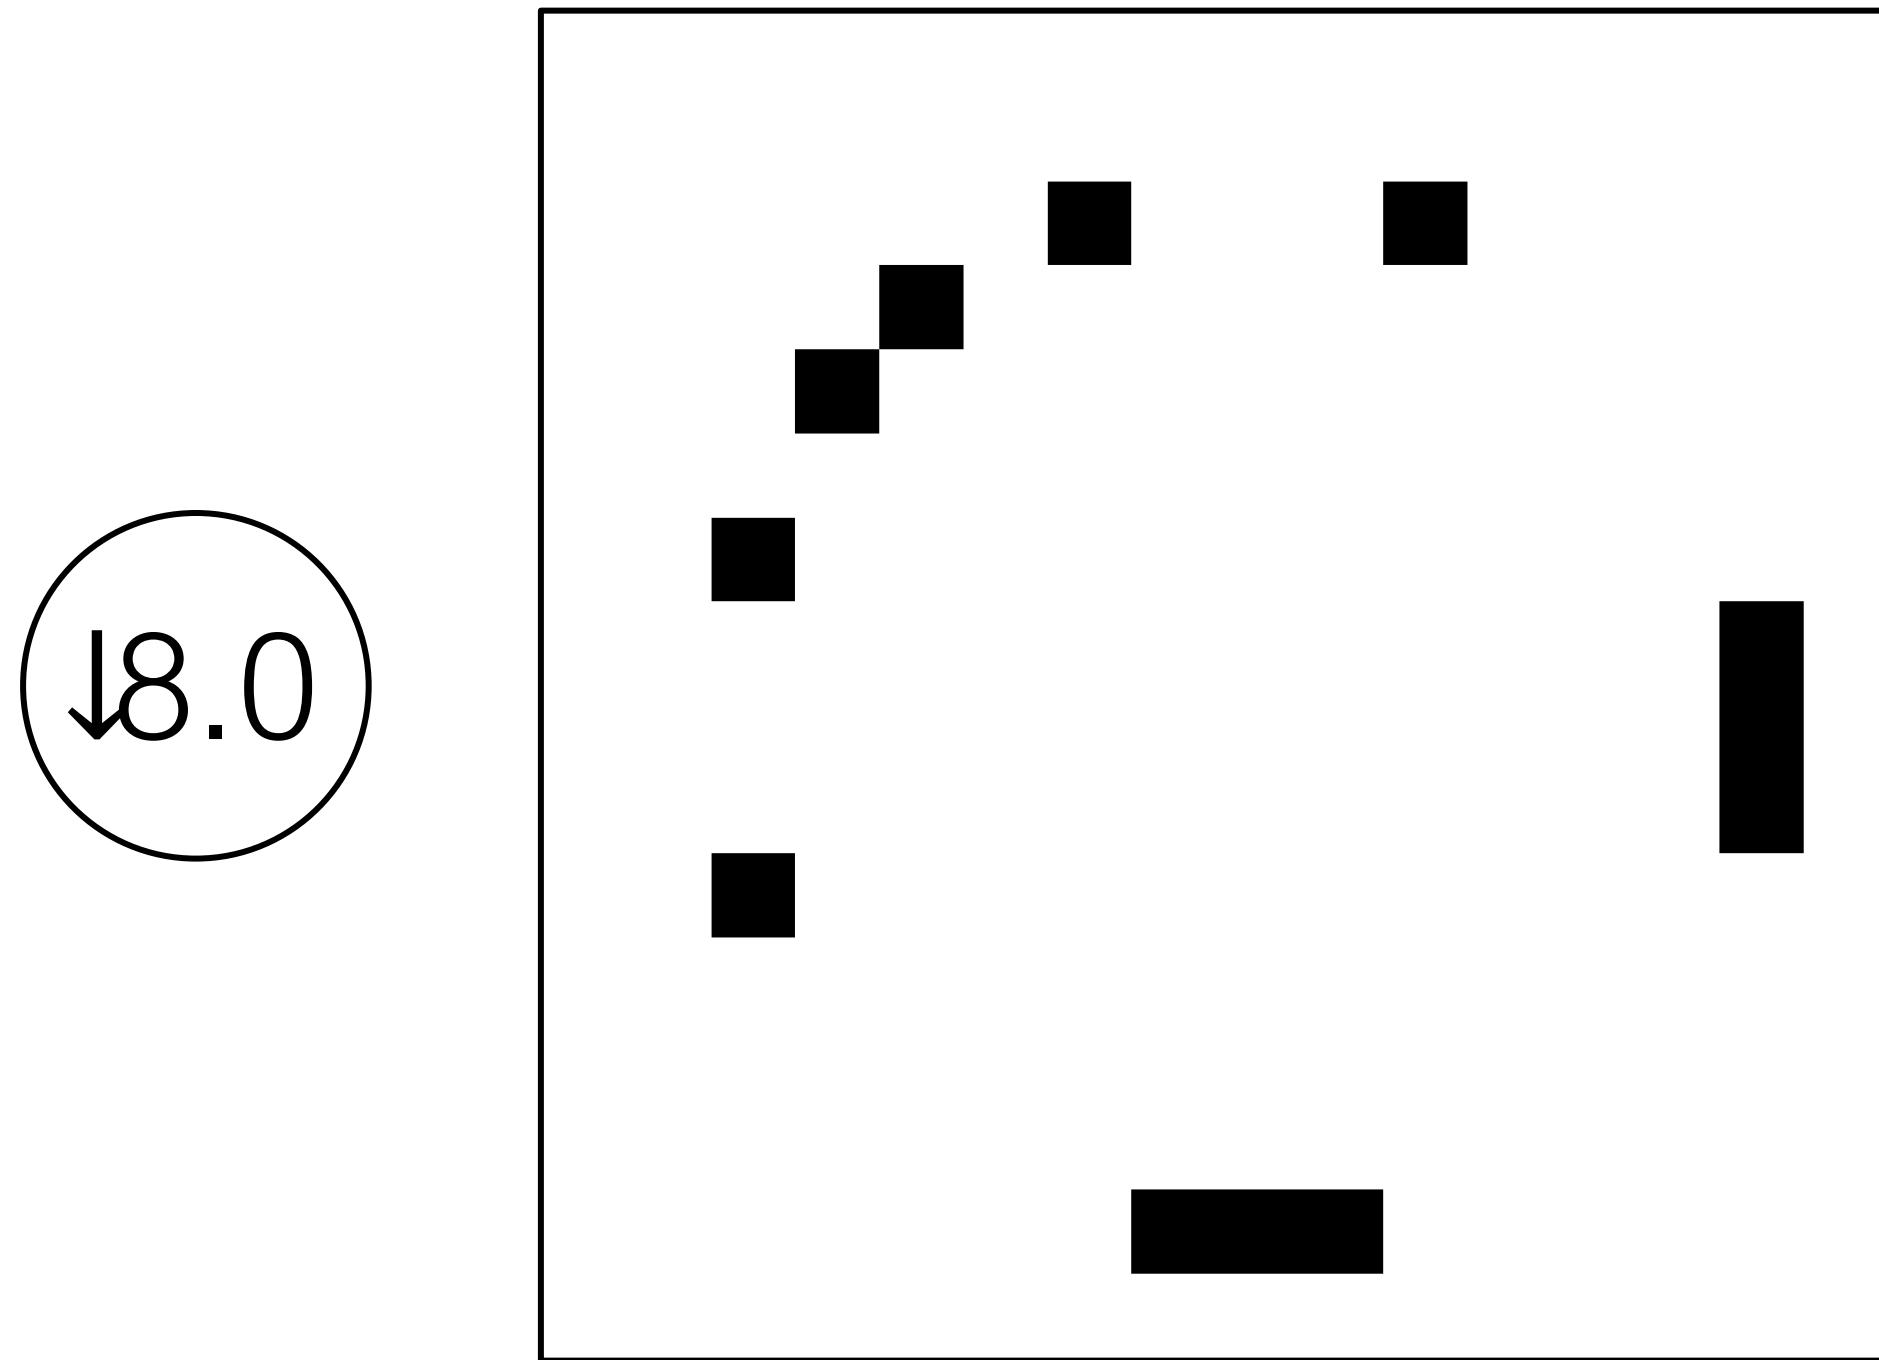


32x32

naïve  
nearest

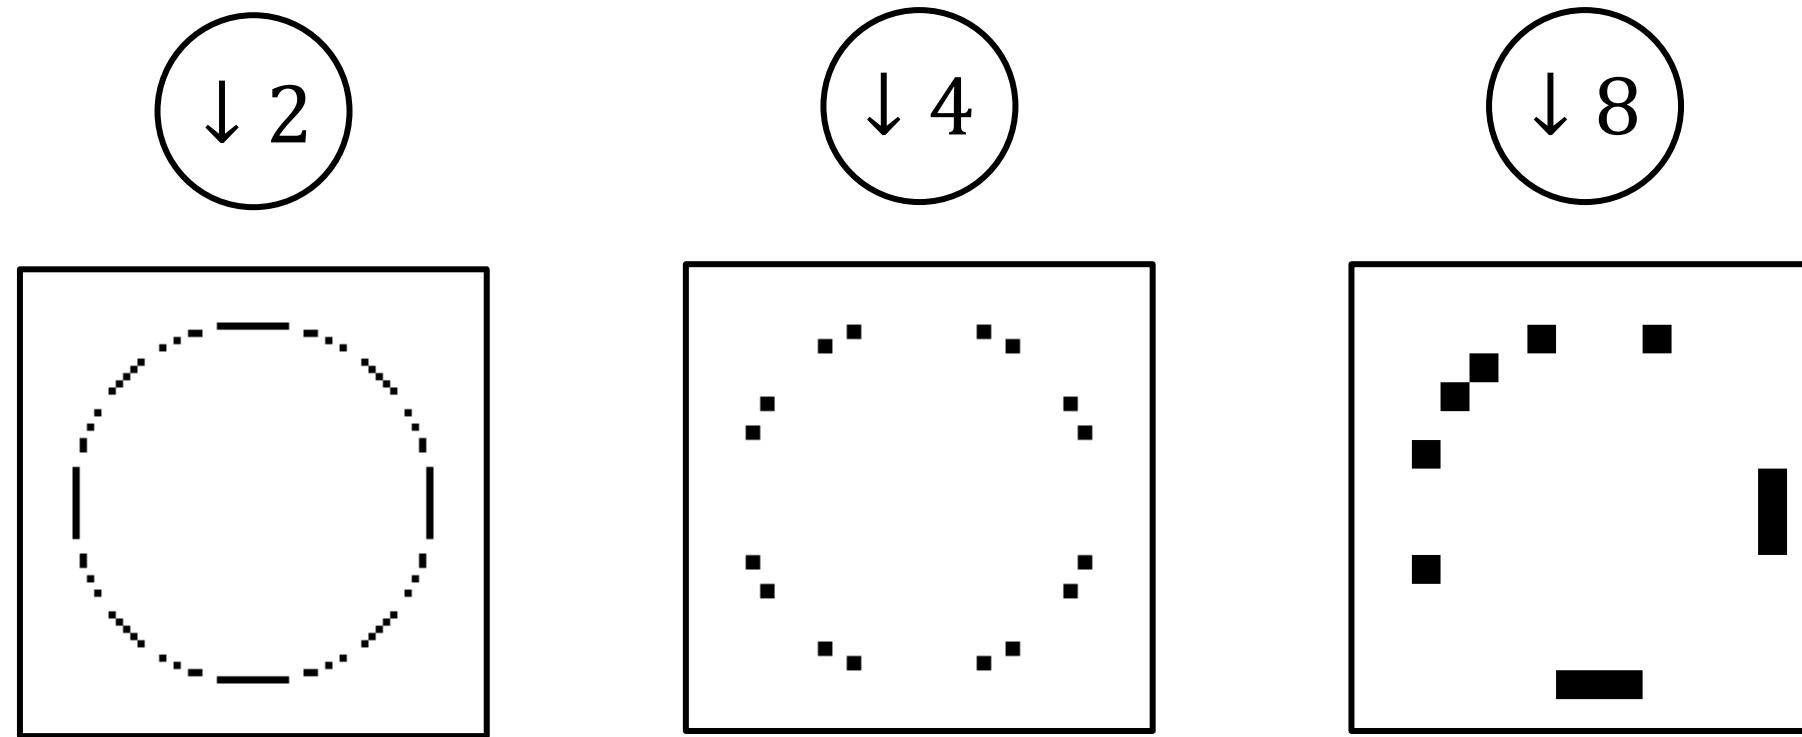


# Downsampling a circle



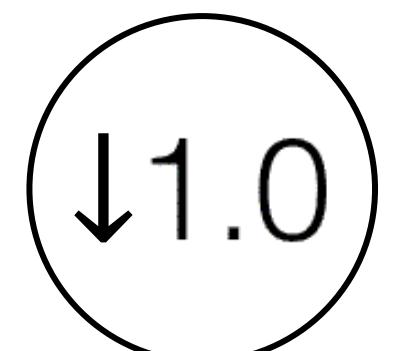
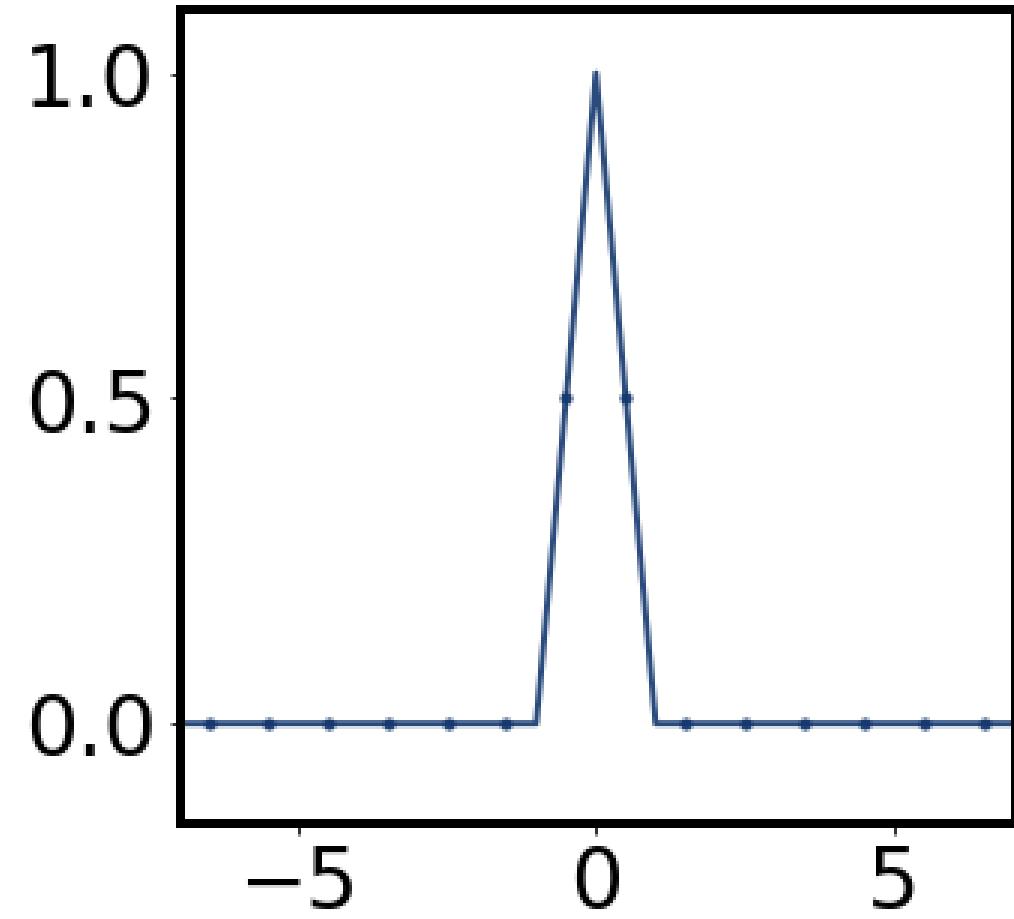
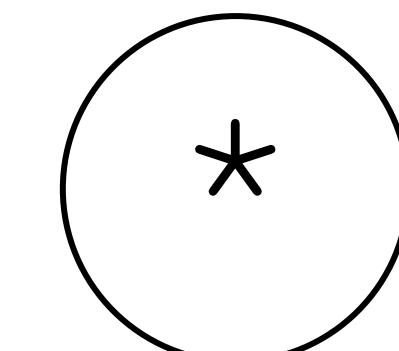
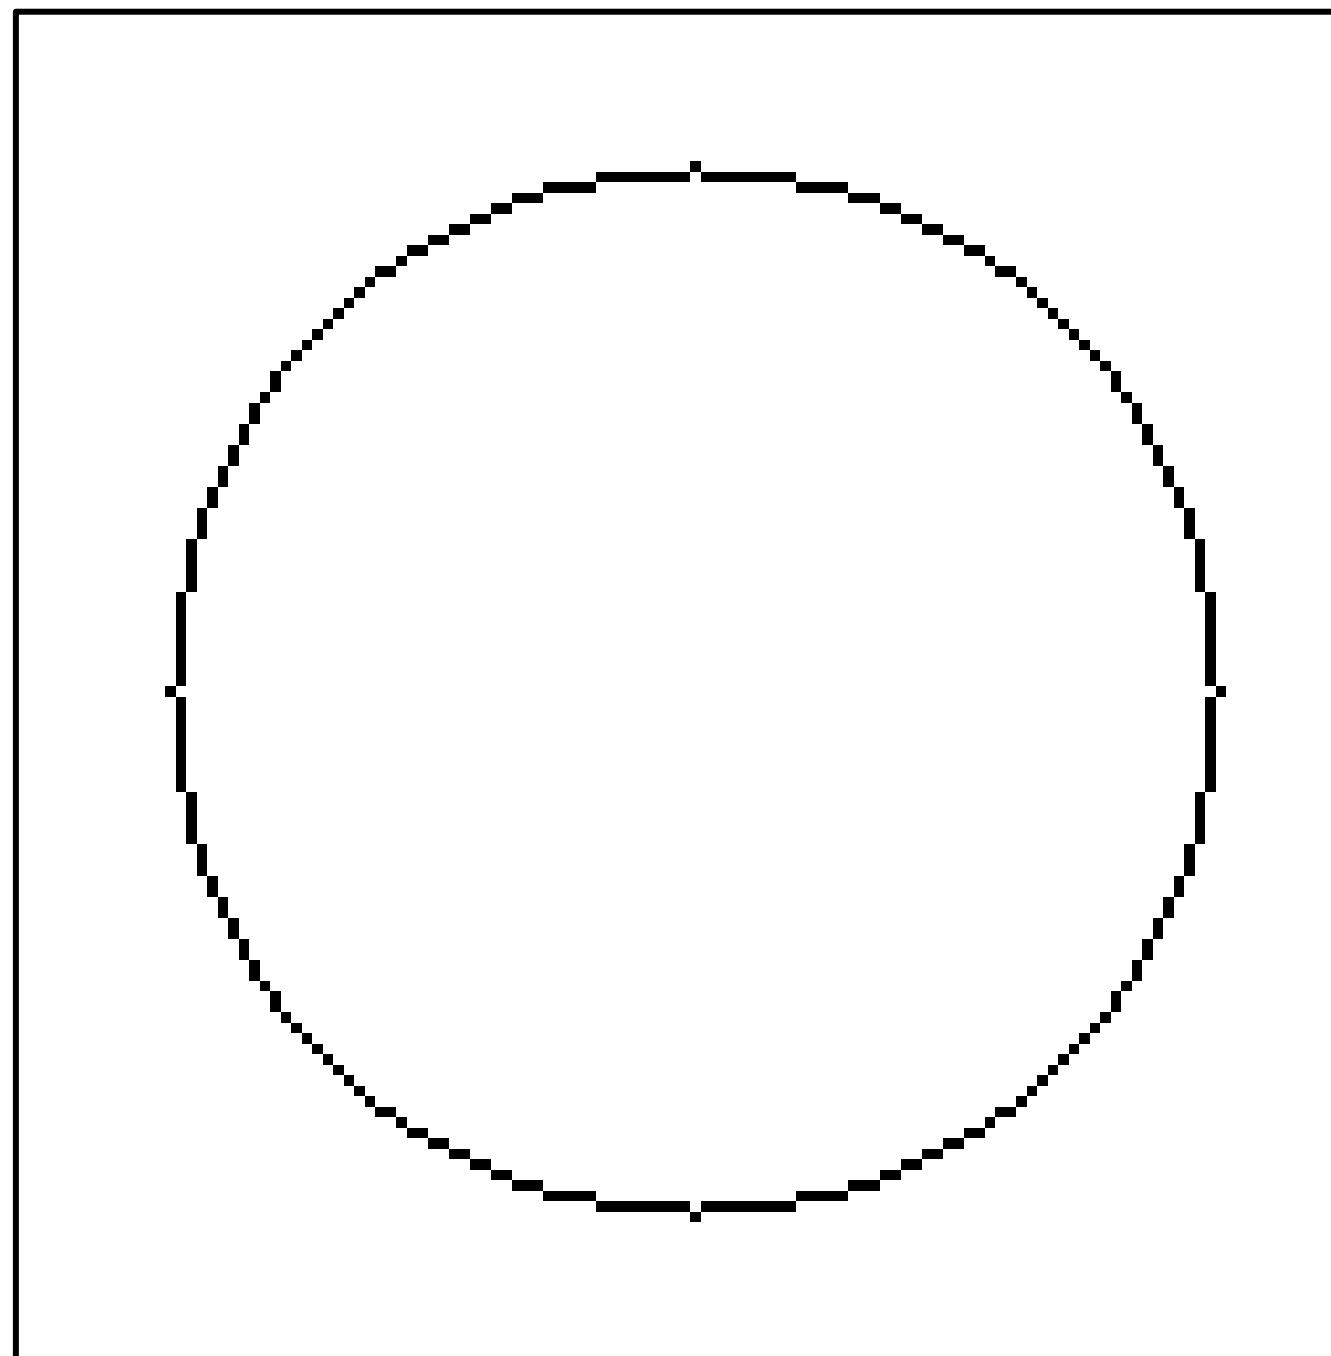
16x16

naïve  
nearest

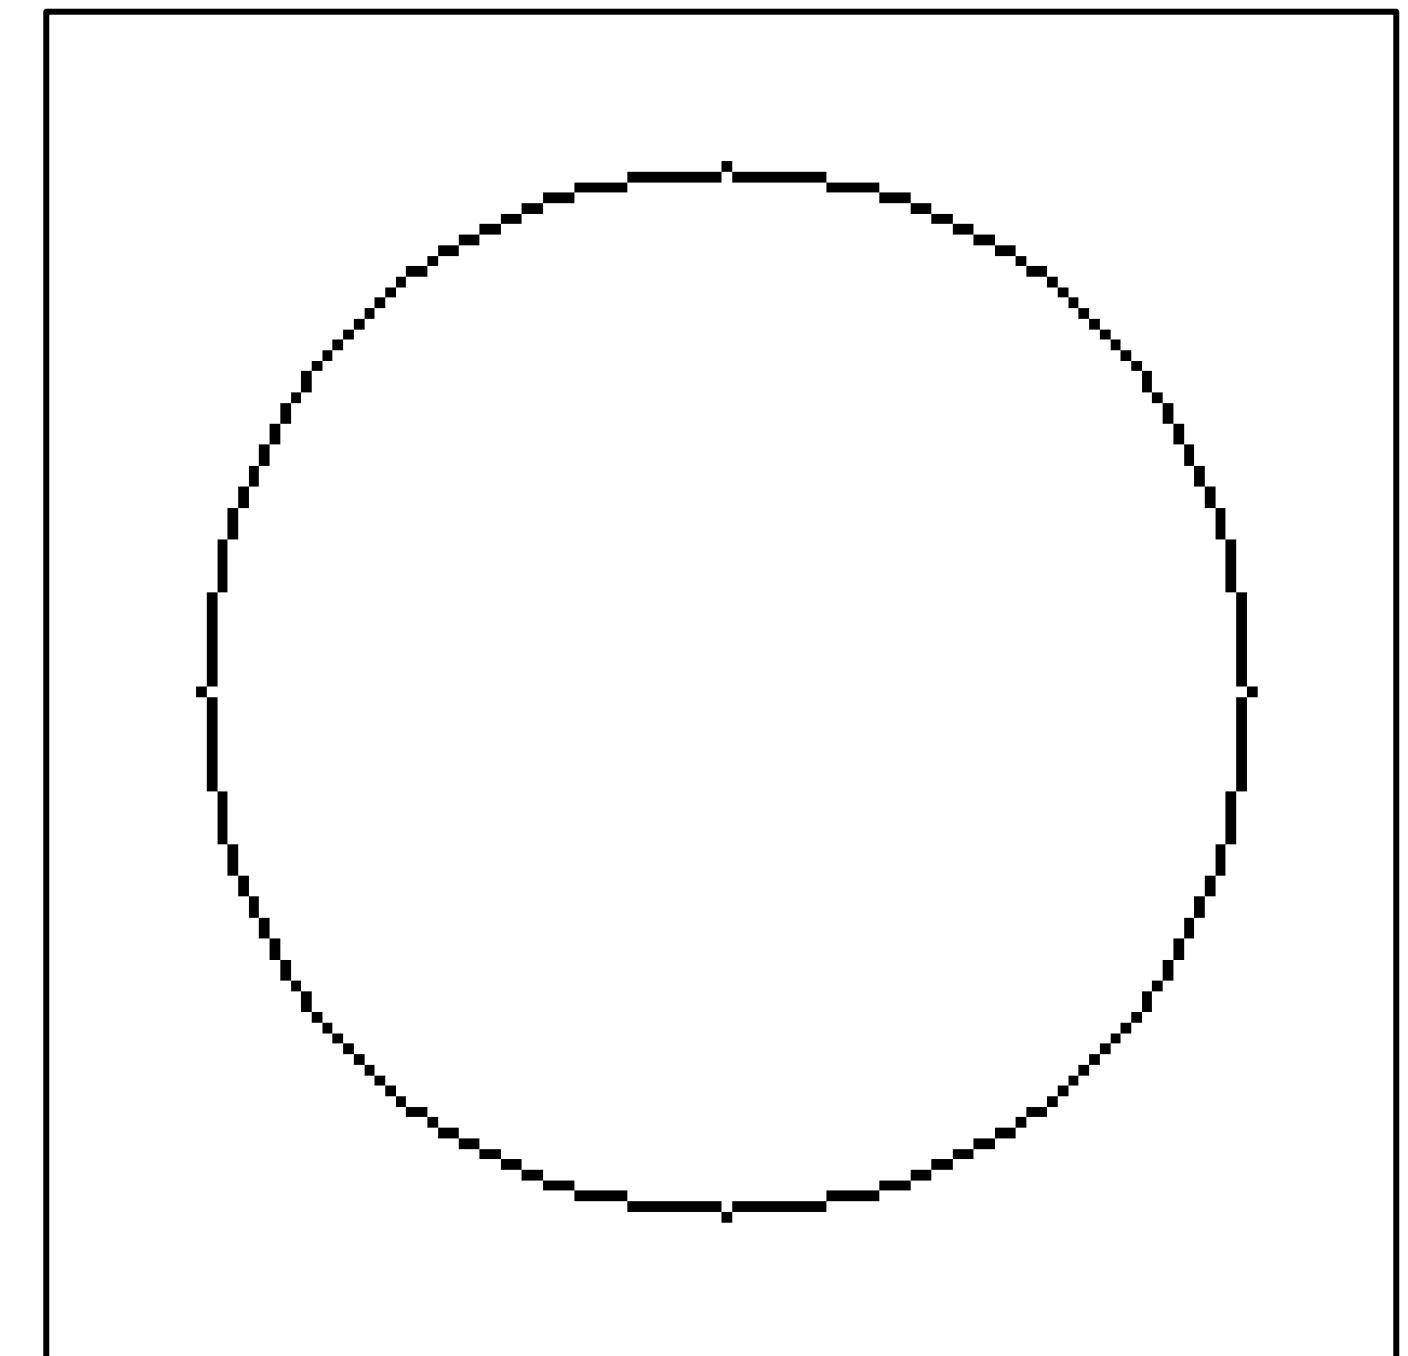


# Downsampling a circle

prefiltering the image



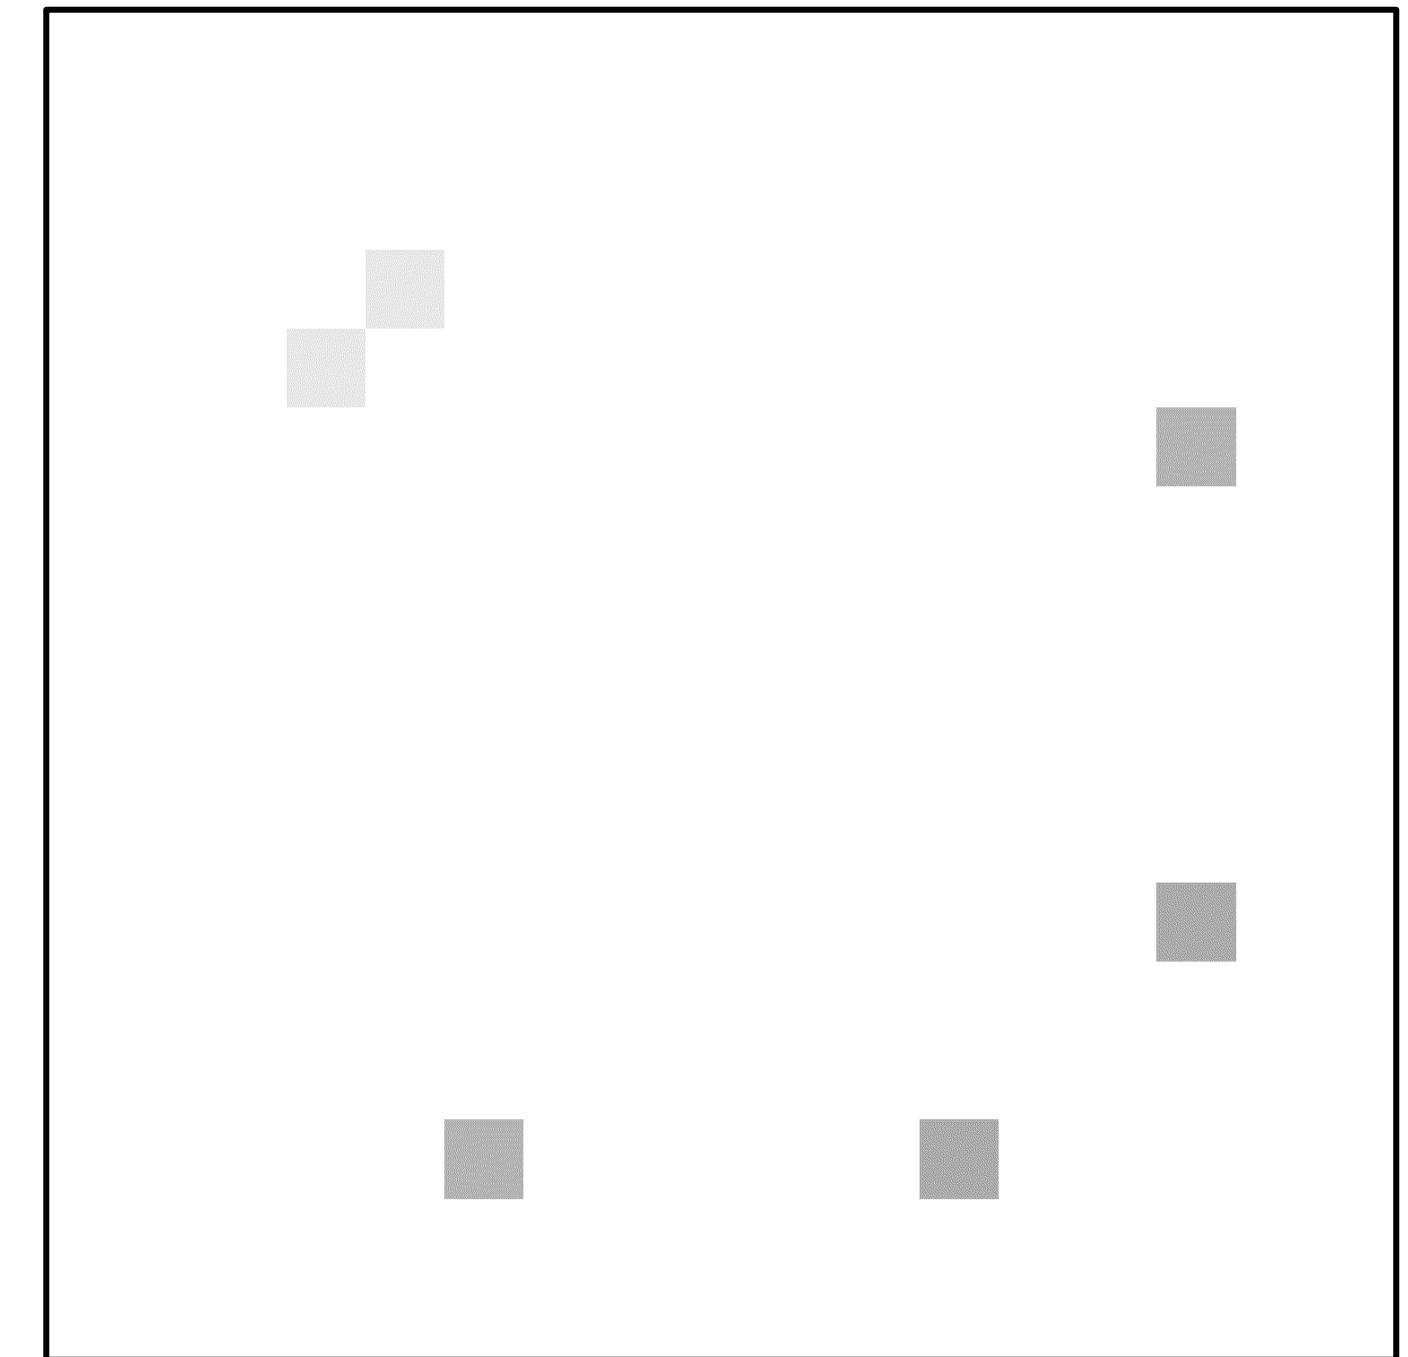
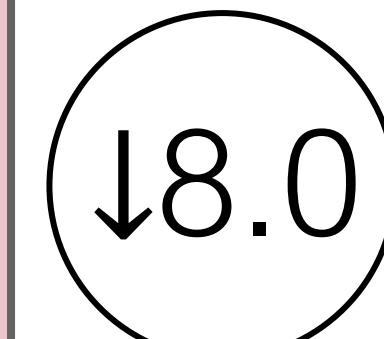
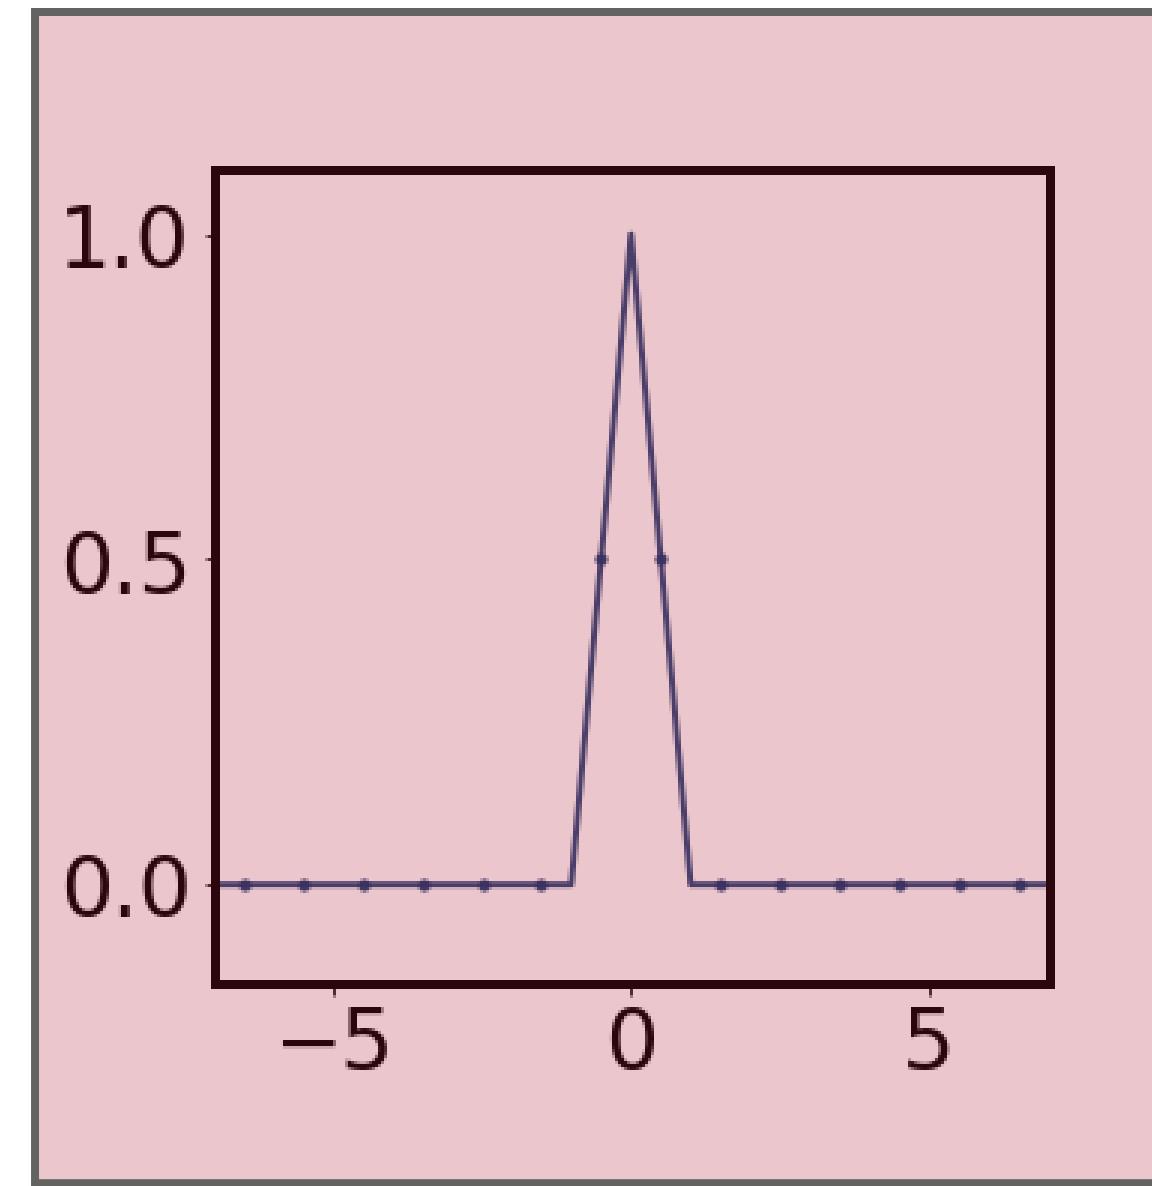
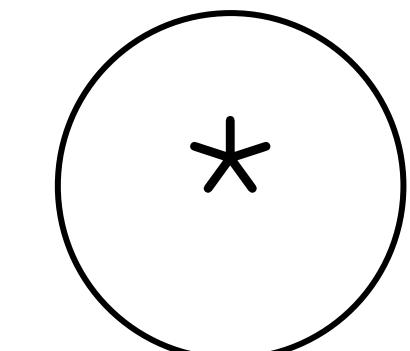
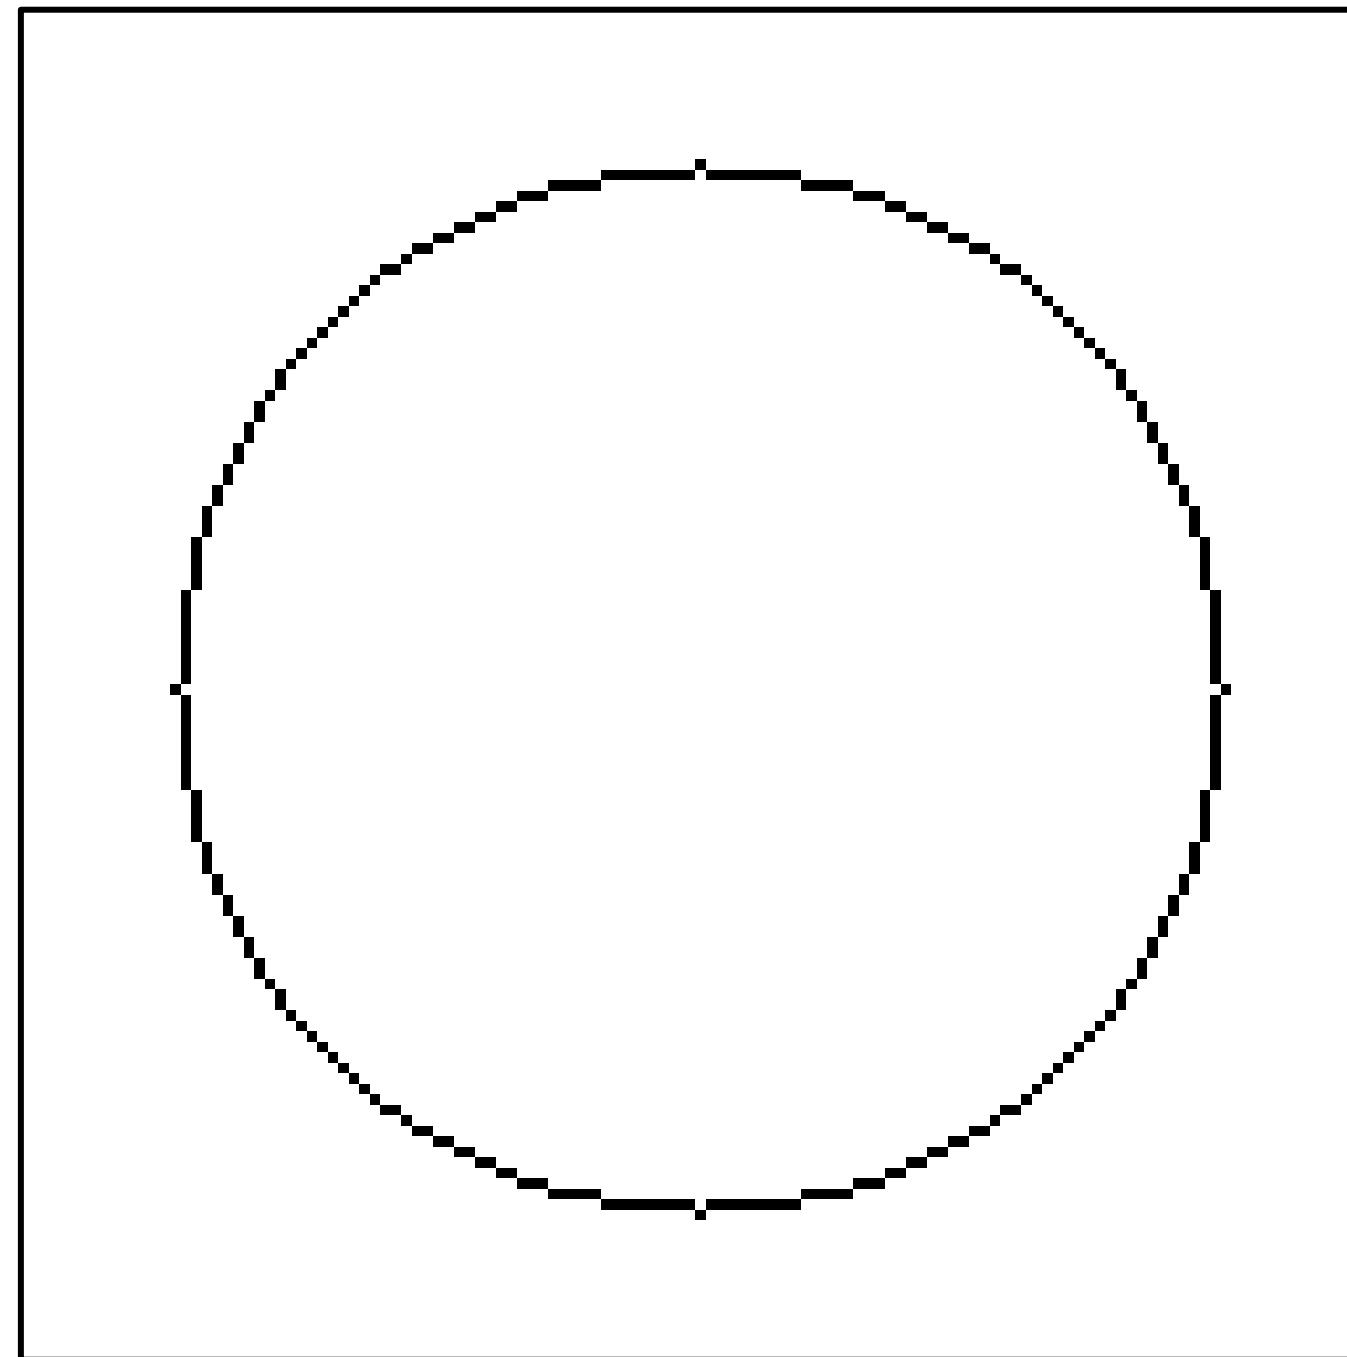
128x128



# Downsampling a circle

prefiltering the image

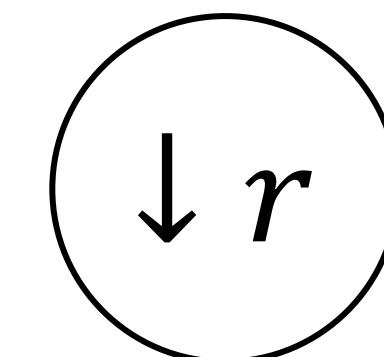
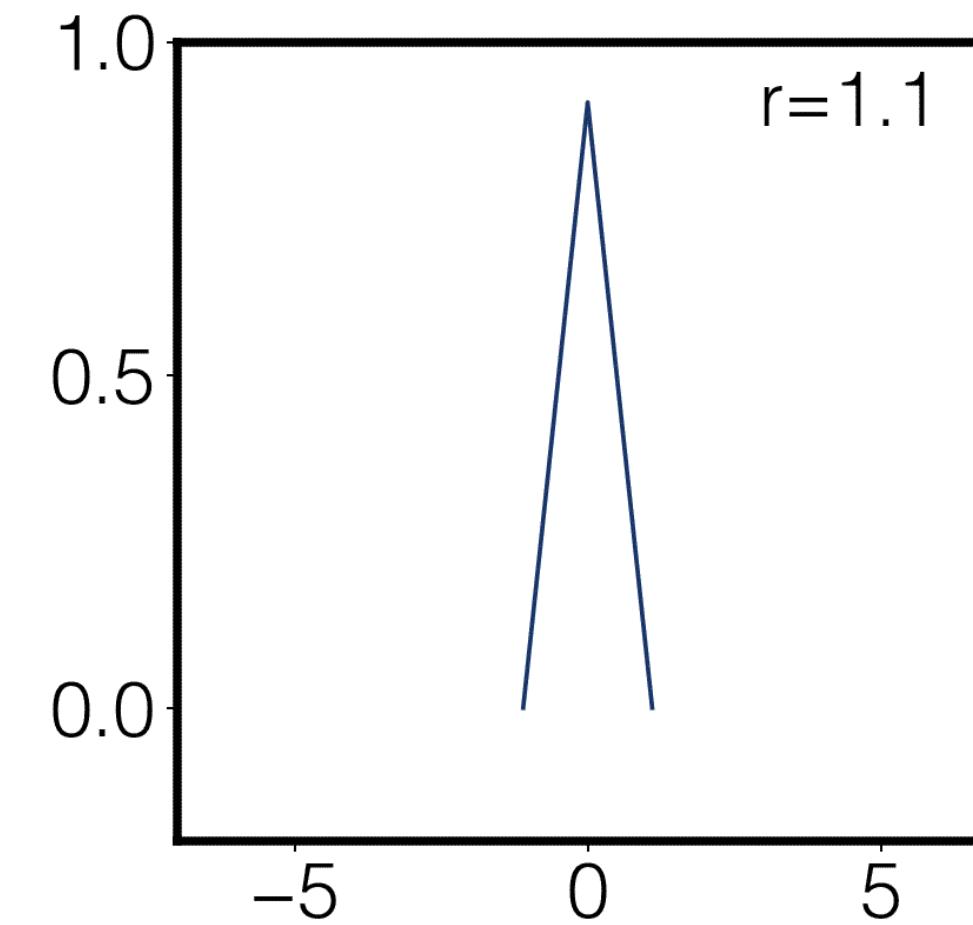
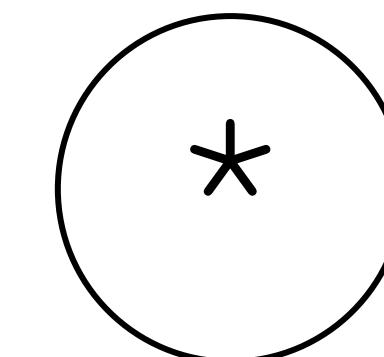
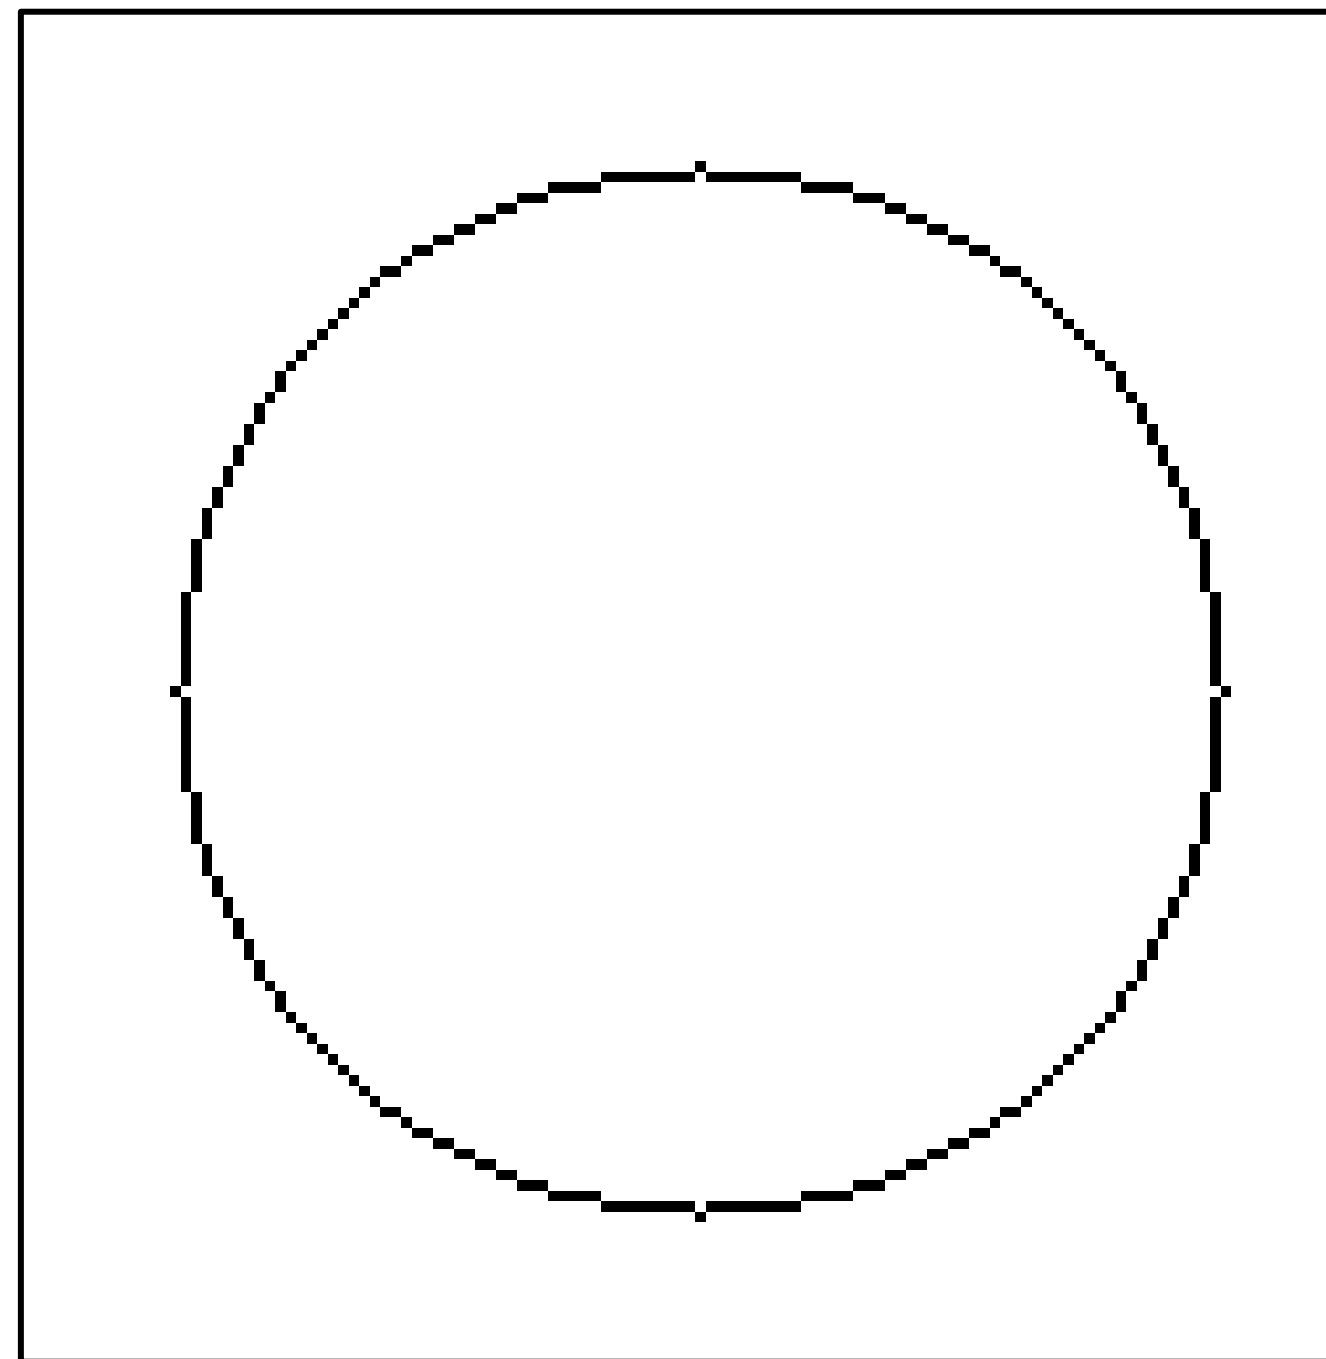
should not be fixed!  
128x128



# Downsampling a circle

---

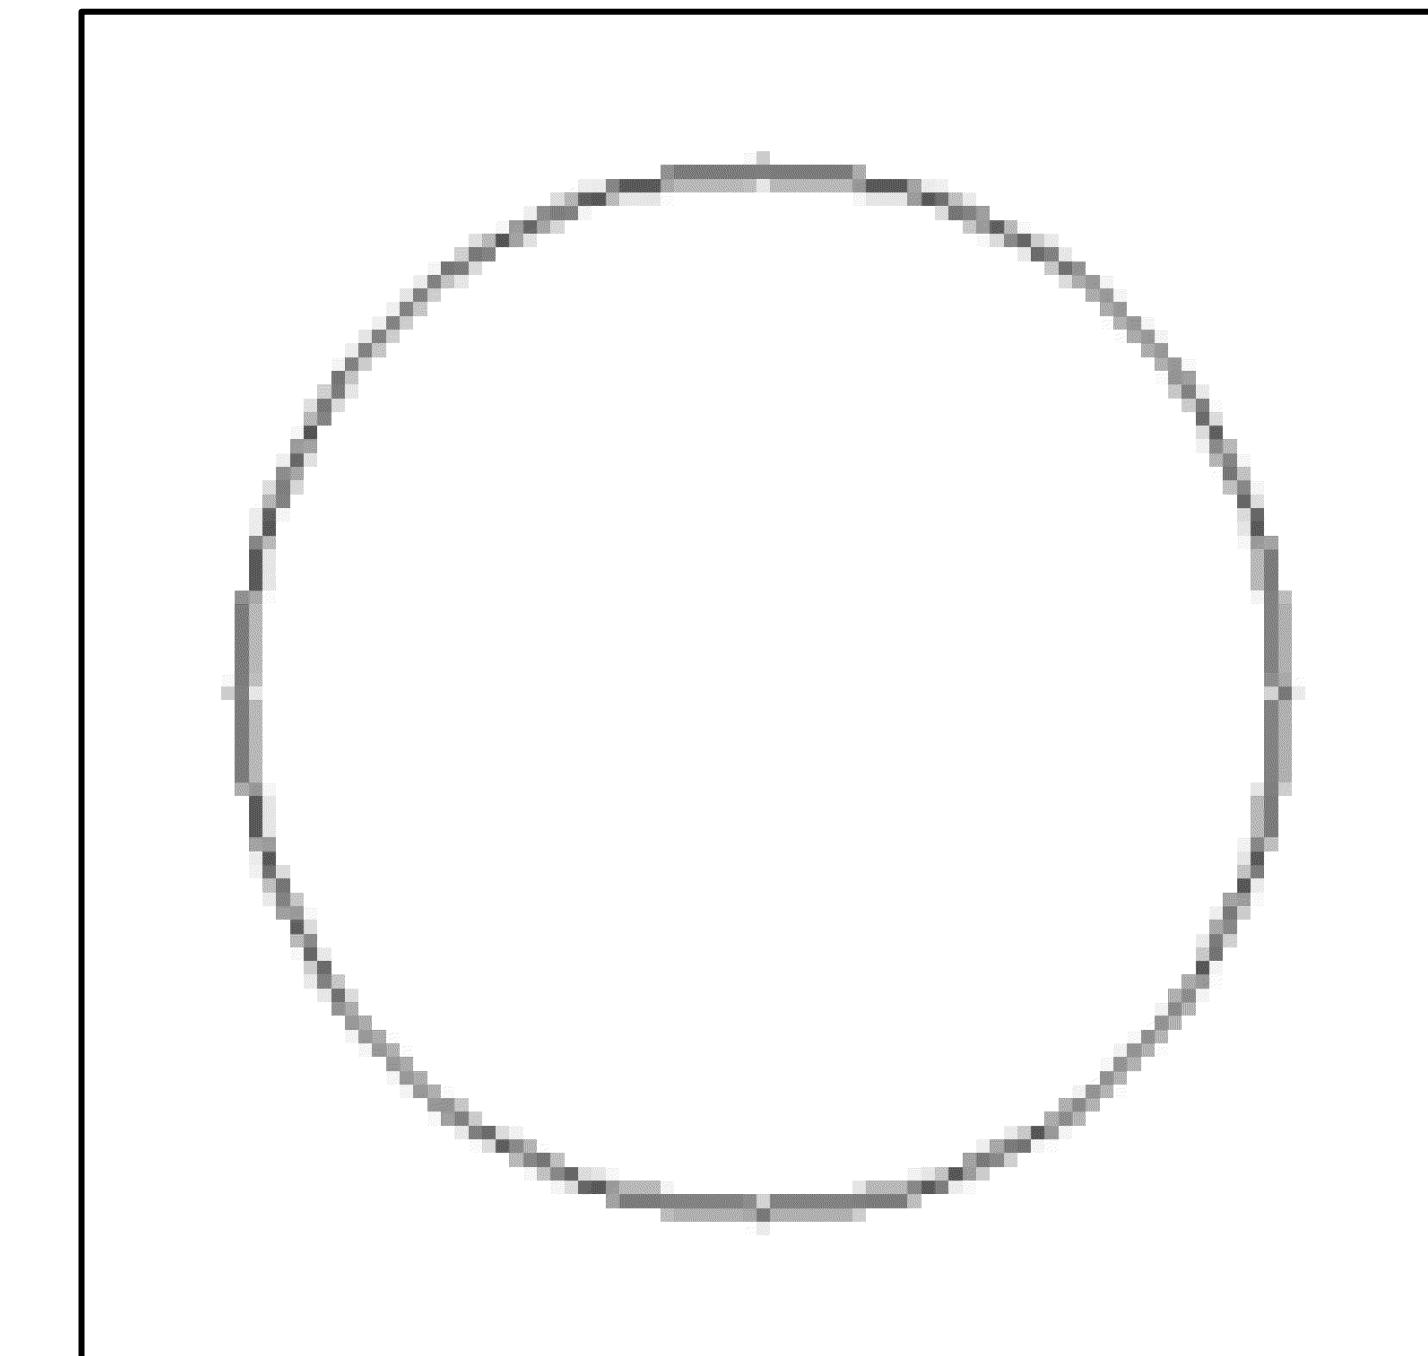
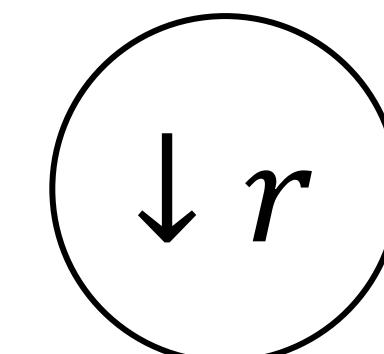
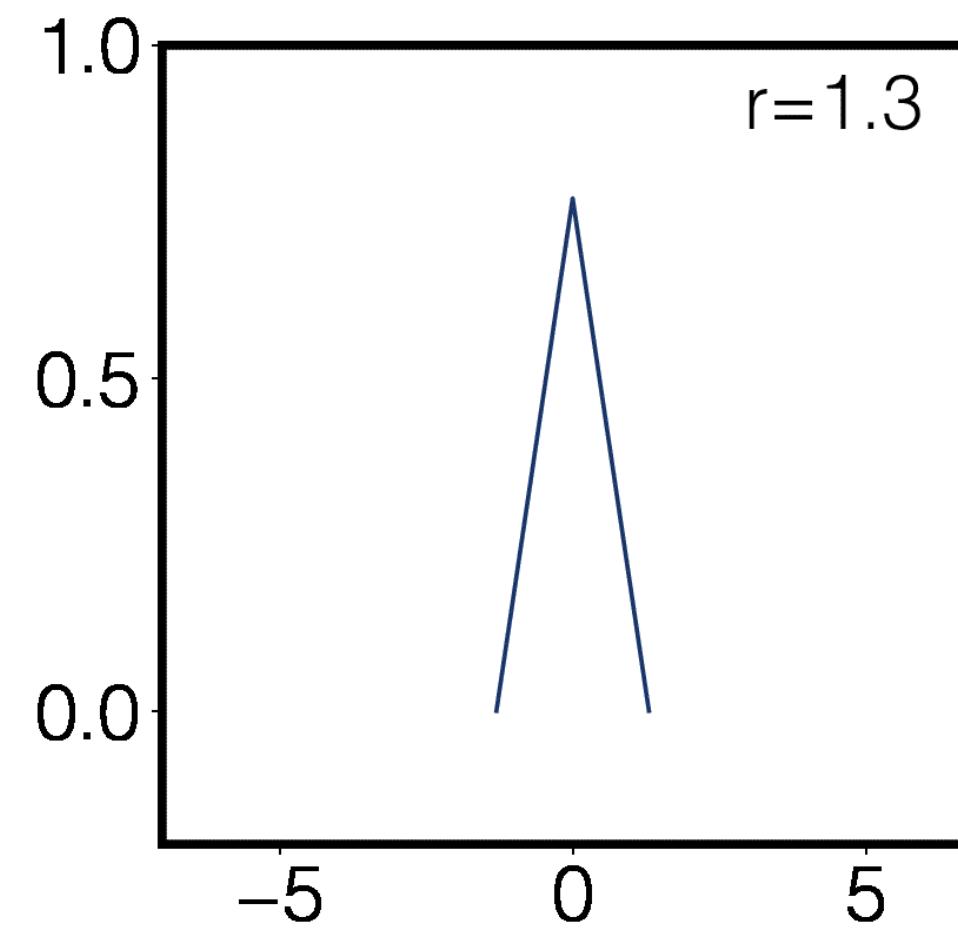
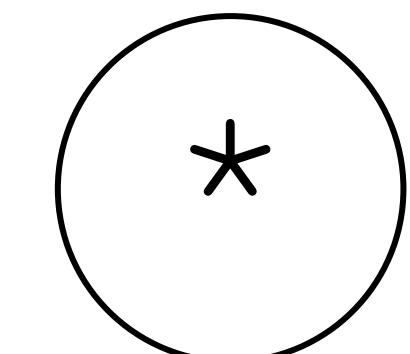
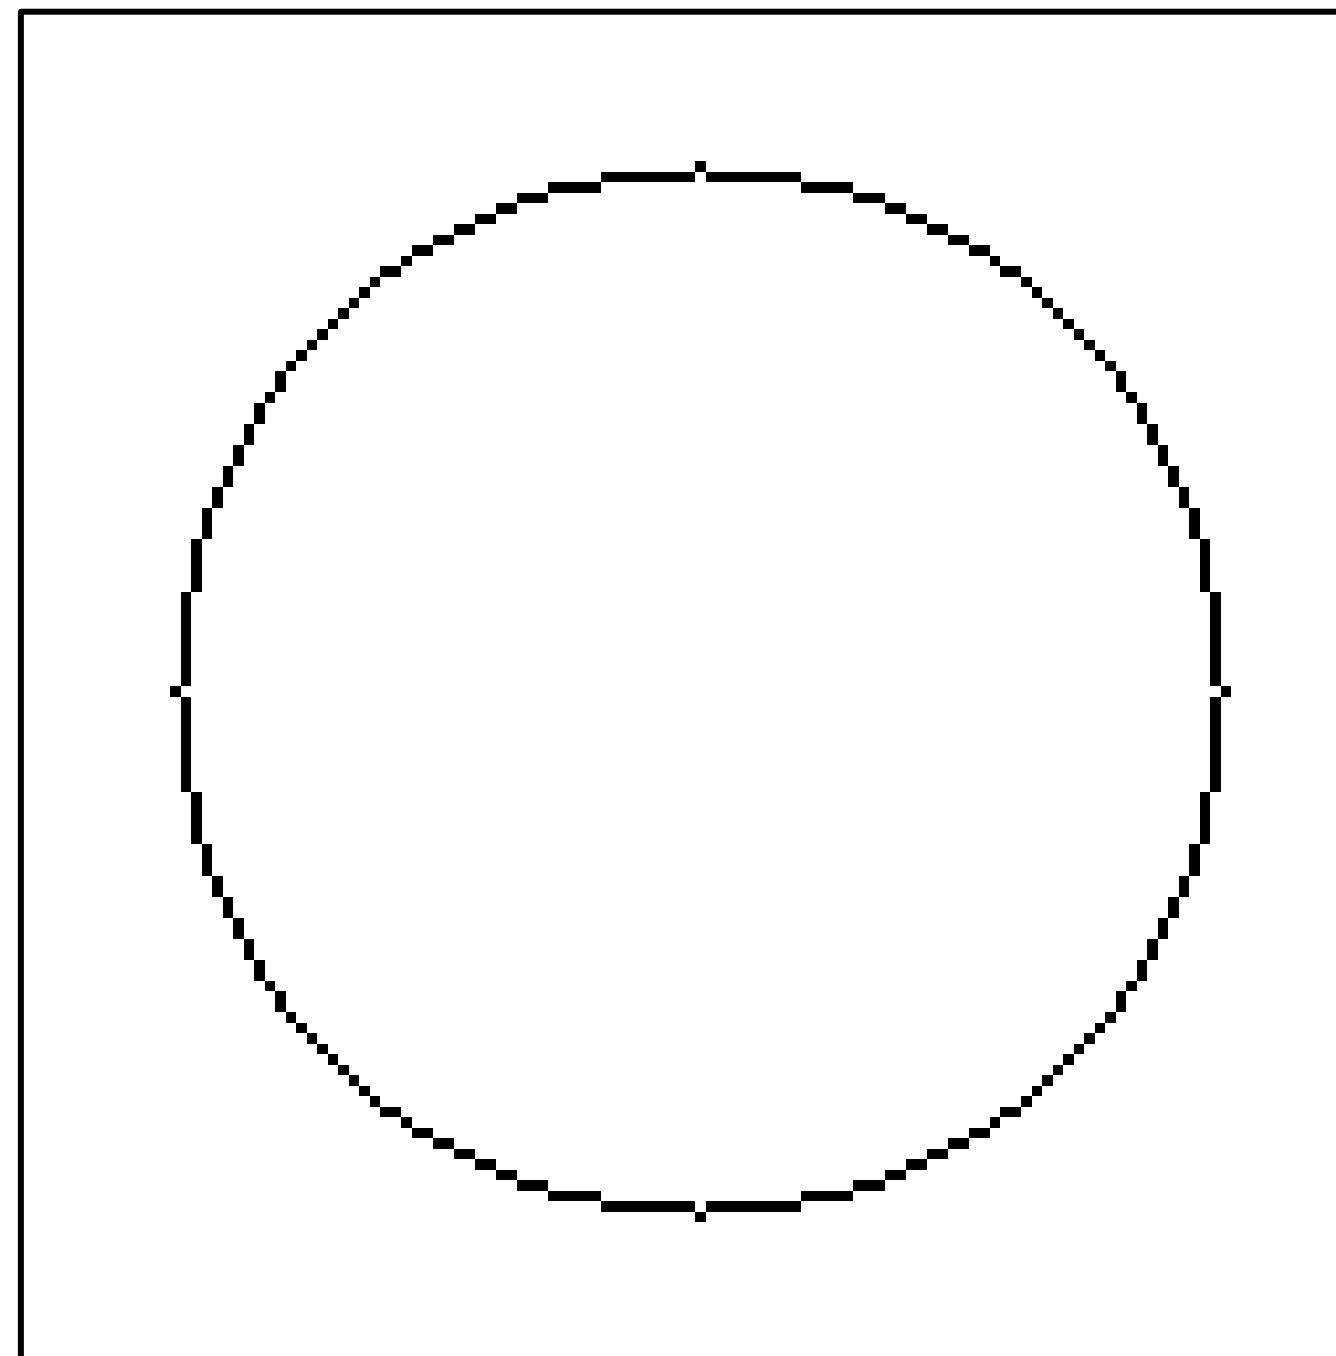
prefiltering the image , adapting the width



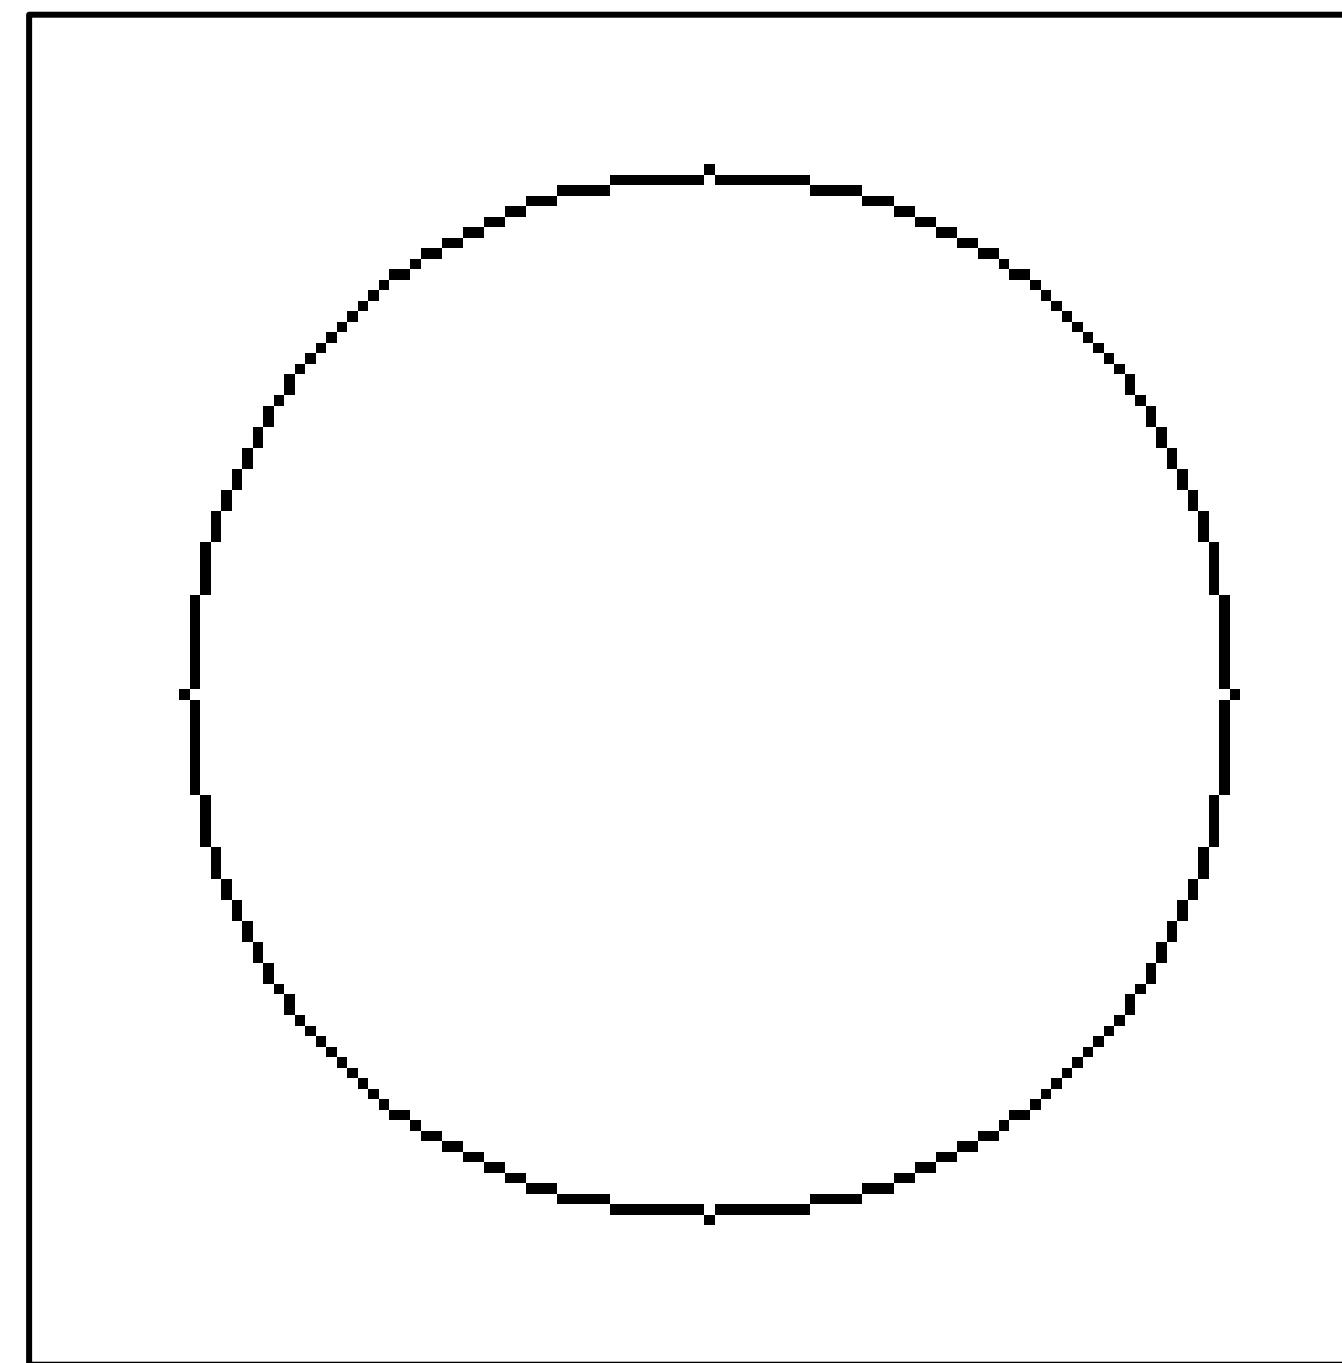
# Downsampling a circle

---

prefiltering the image , adapting the width

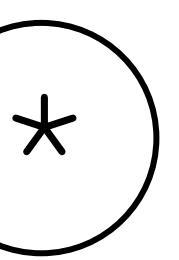
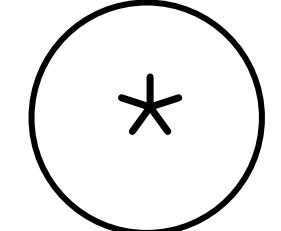


# Downsampling a circle

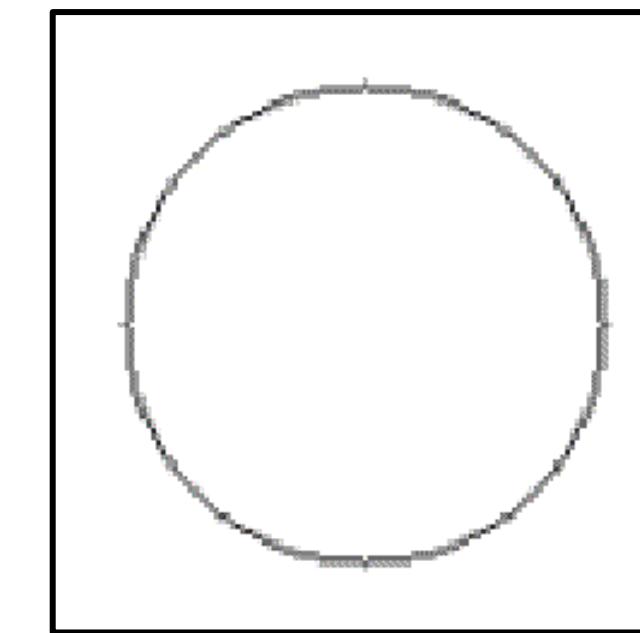
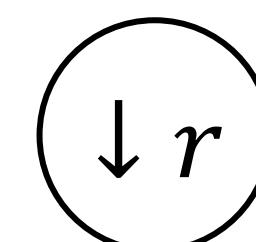
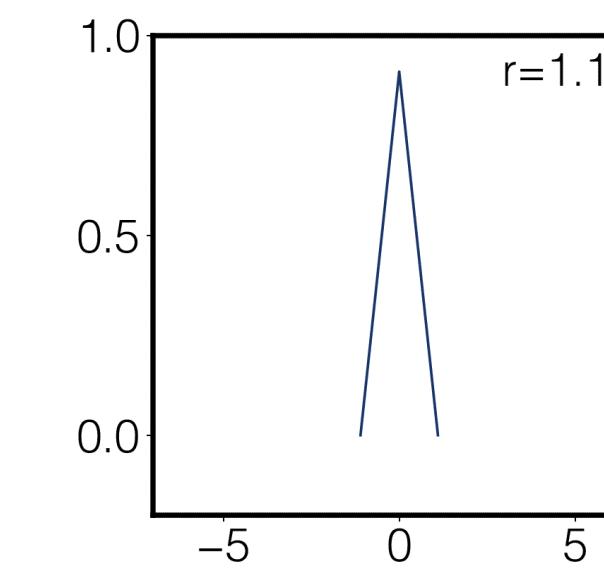
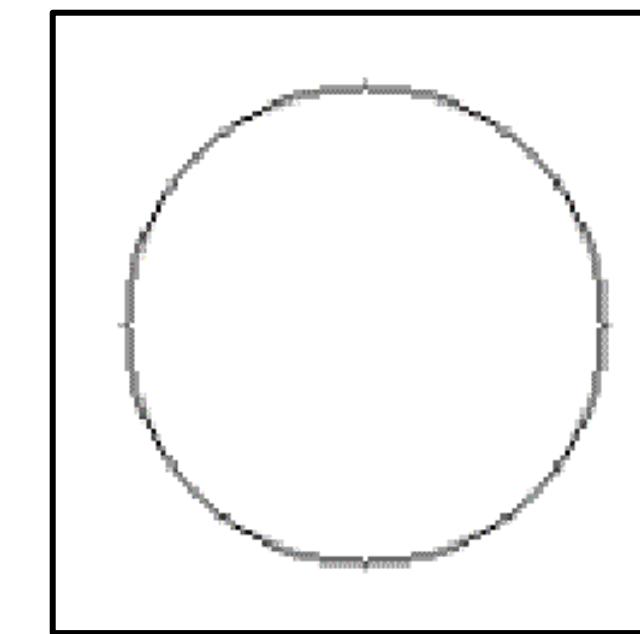
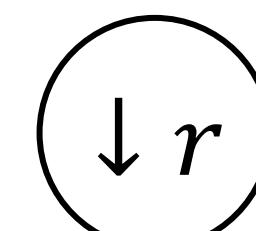
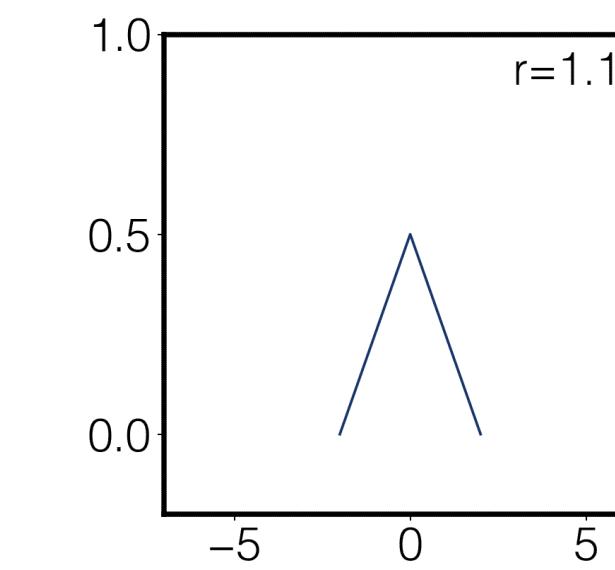
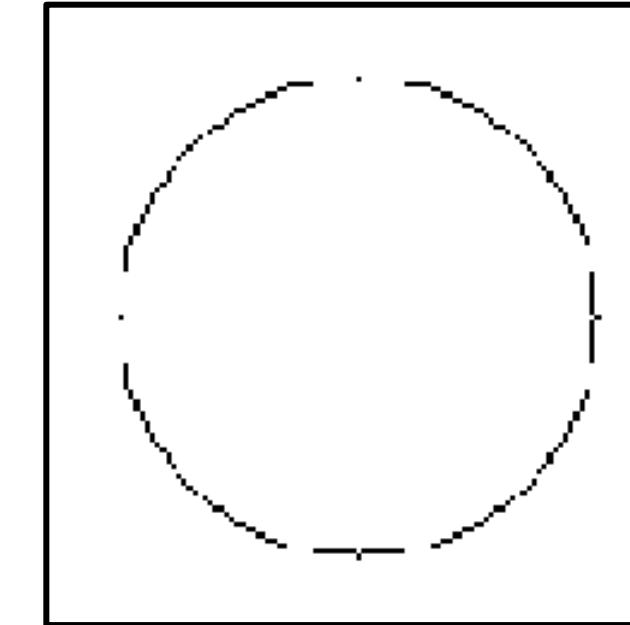


fixed  
filter

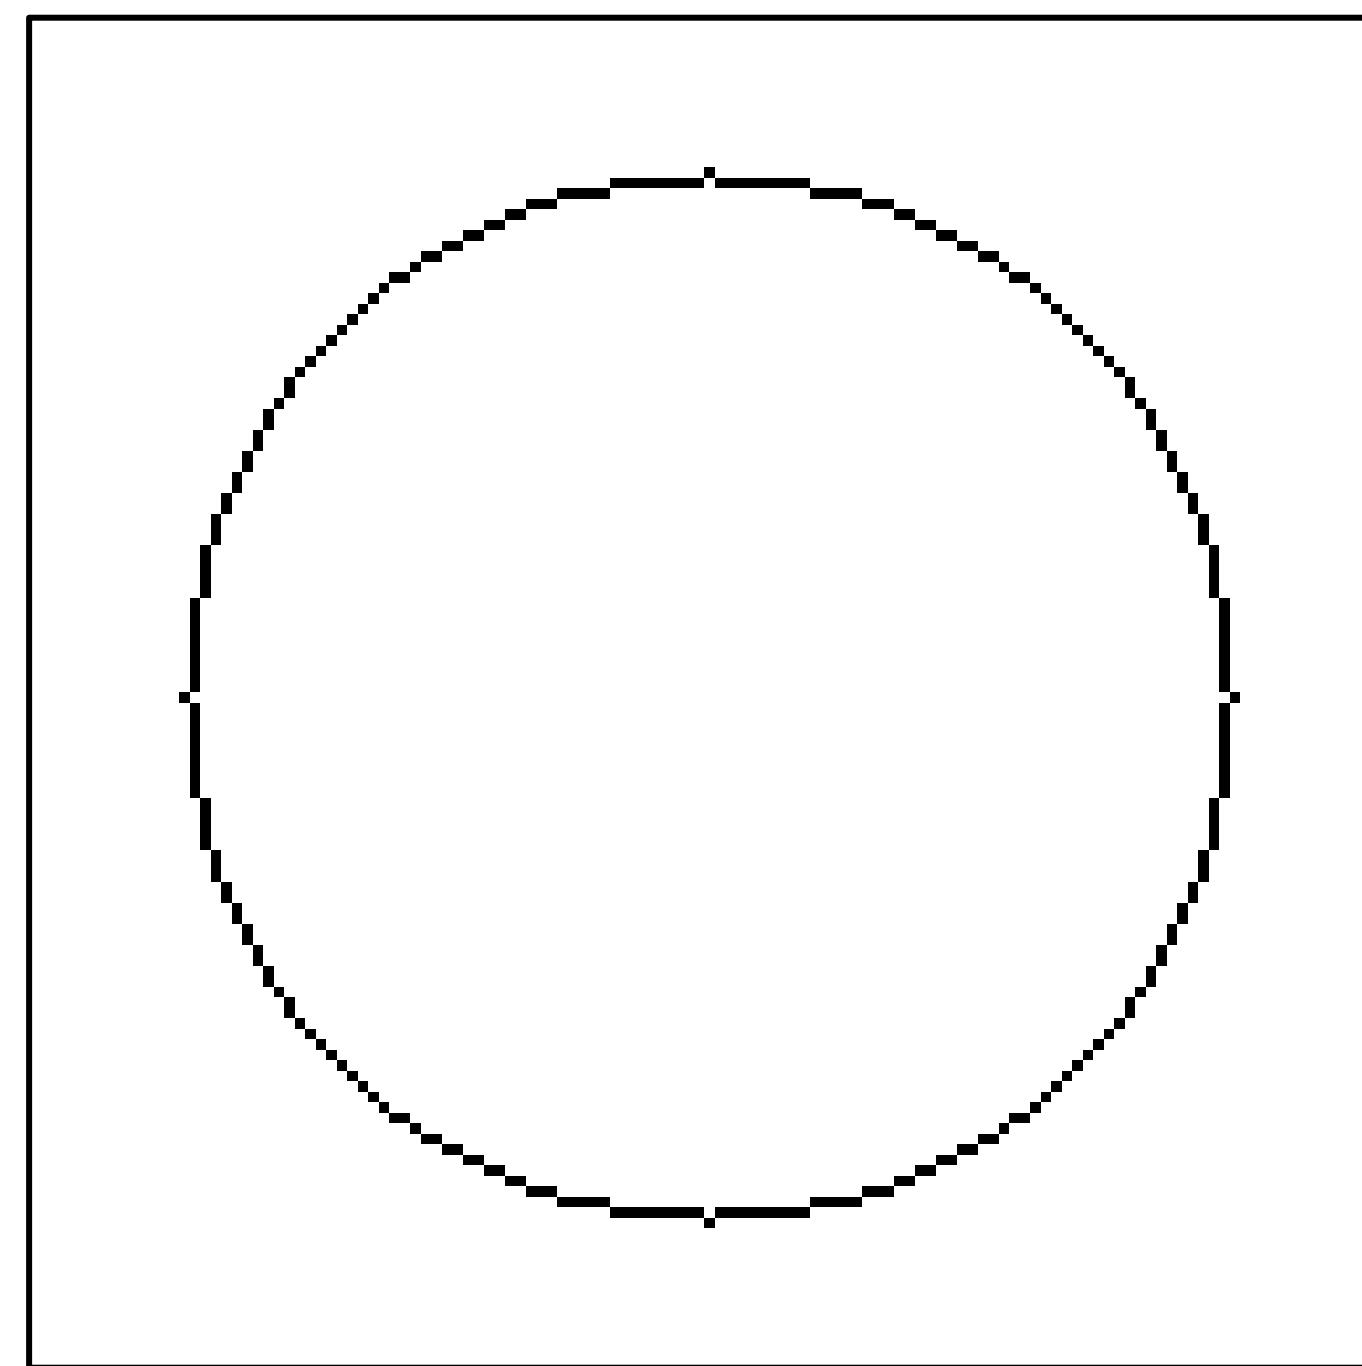
adaptive  
filter



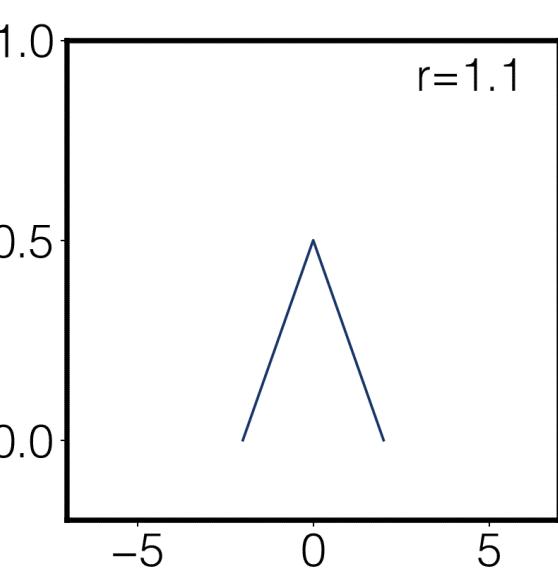
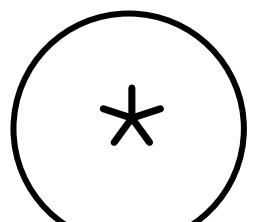
naïve  
nearest



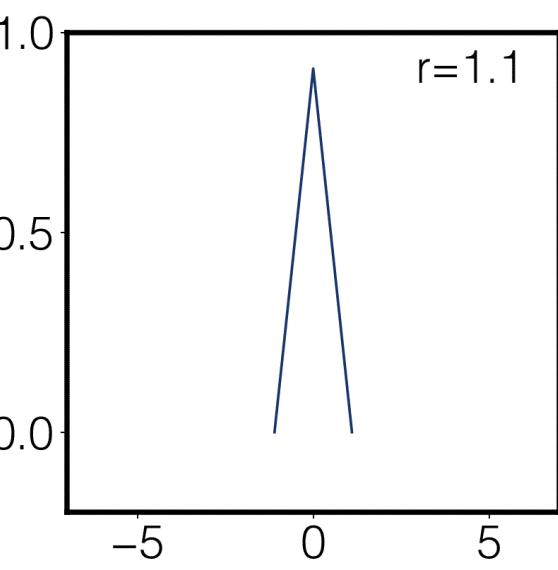
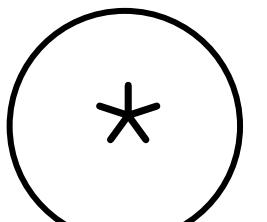
# Downsampling a circle



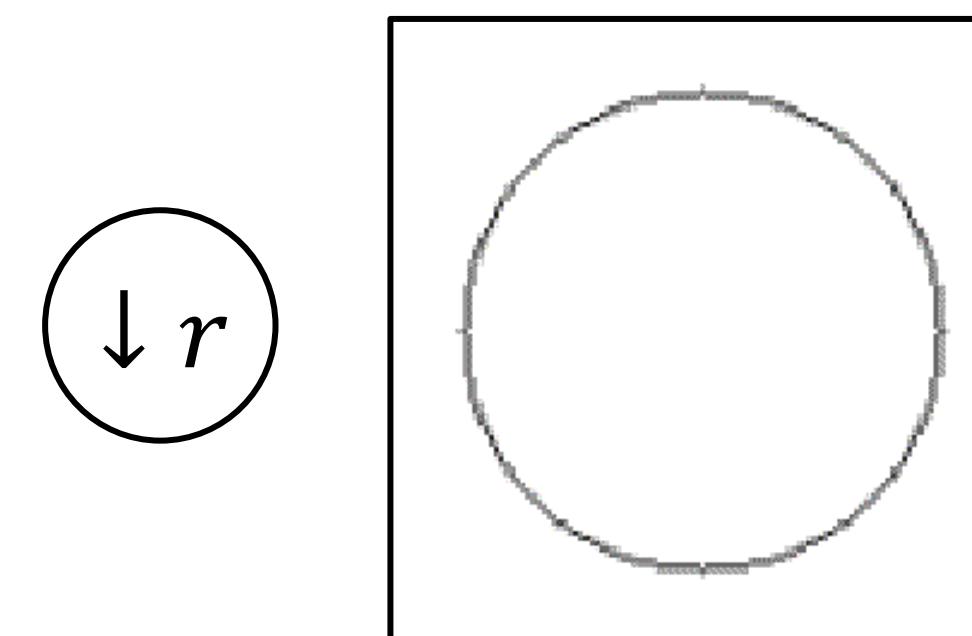
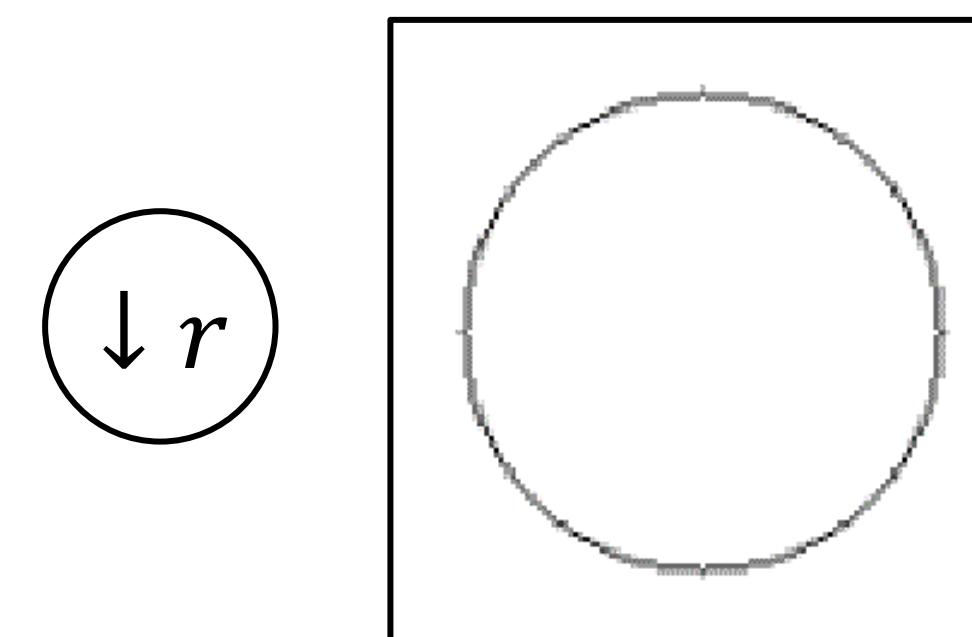
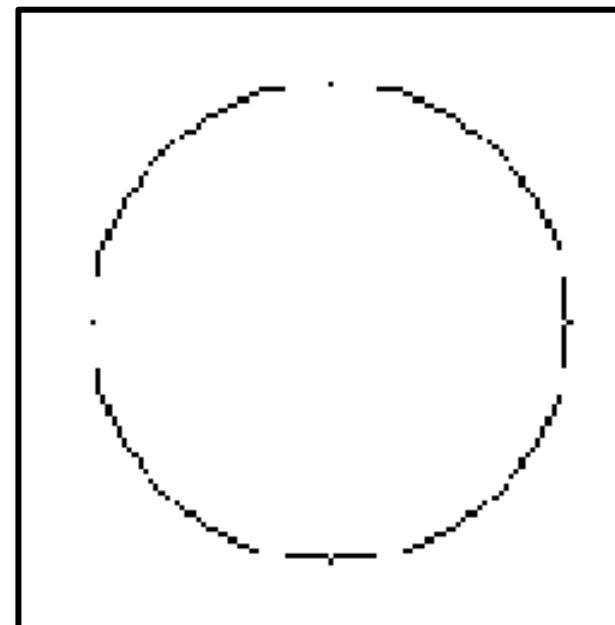
fixed  
filter



adaptive  
filter



naïve  
nearest



 PyTorch  
(with default flags)

 OpenCV

 TensorFlow  
 MXNet

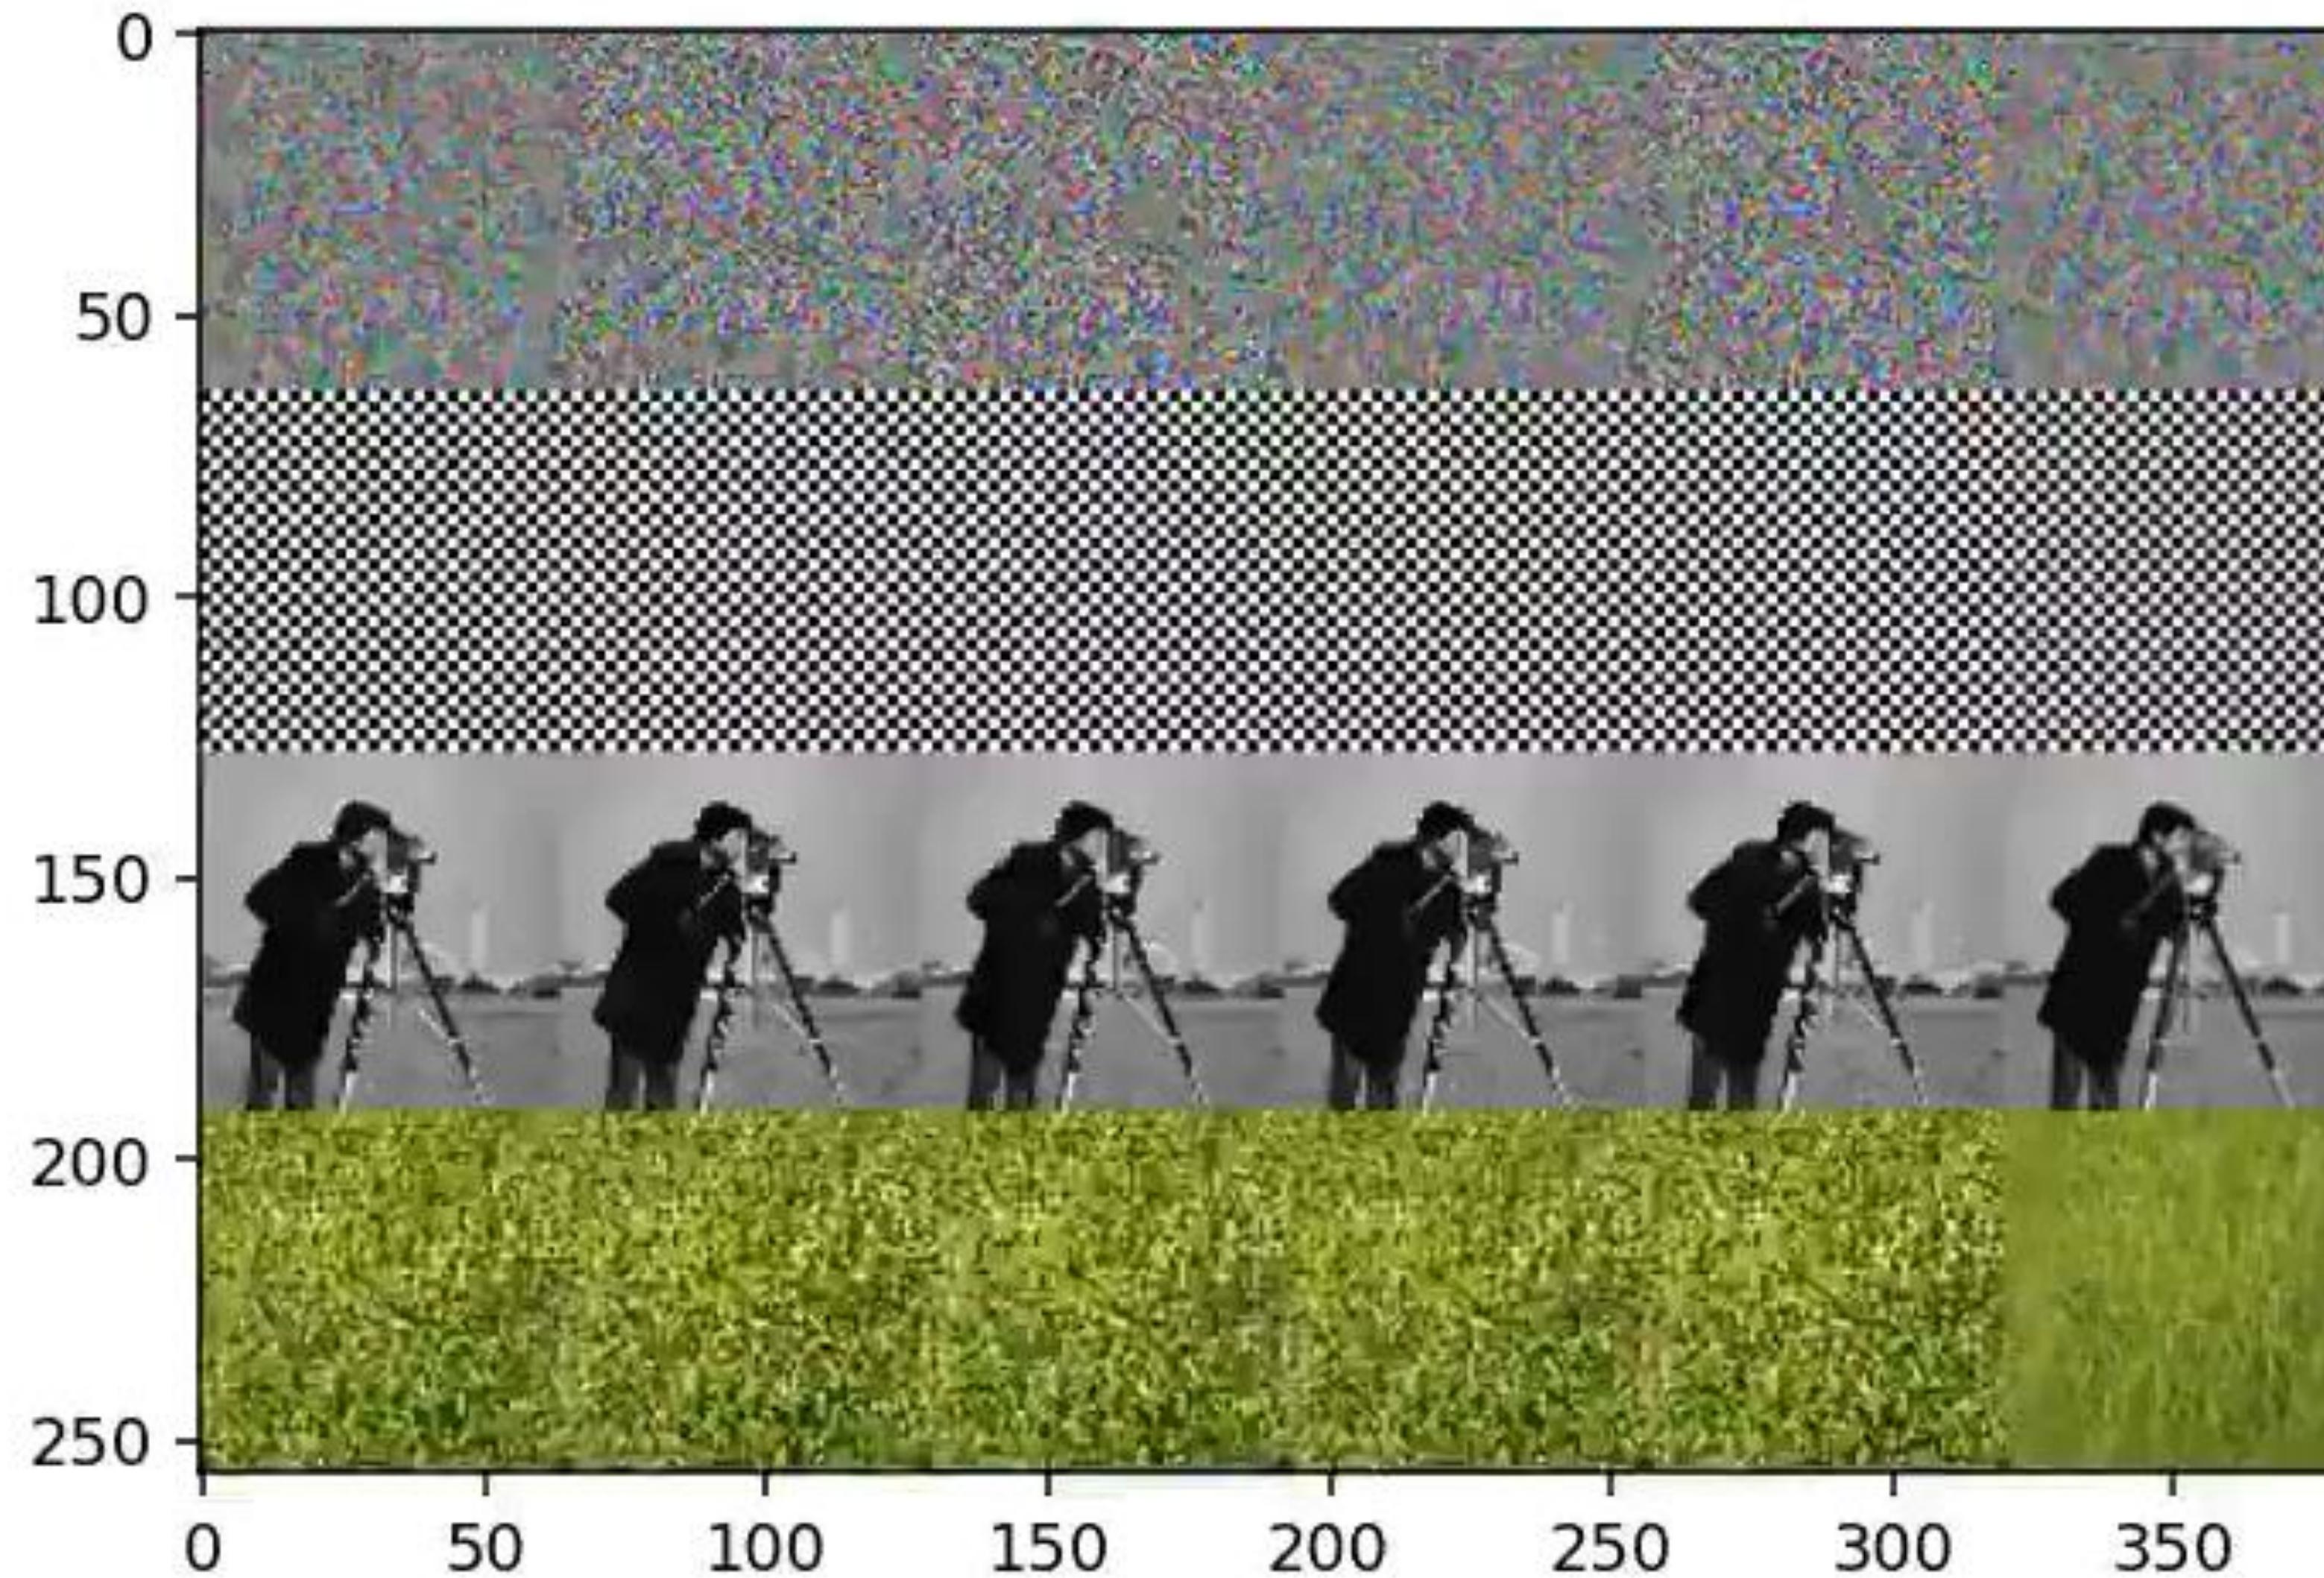


PIL

# A real problem!

---

128 x 128 → 64x64



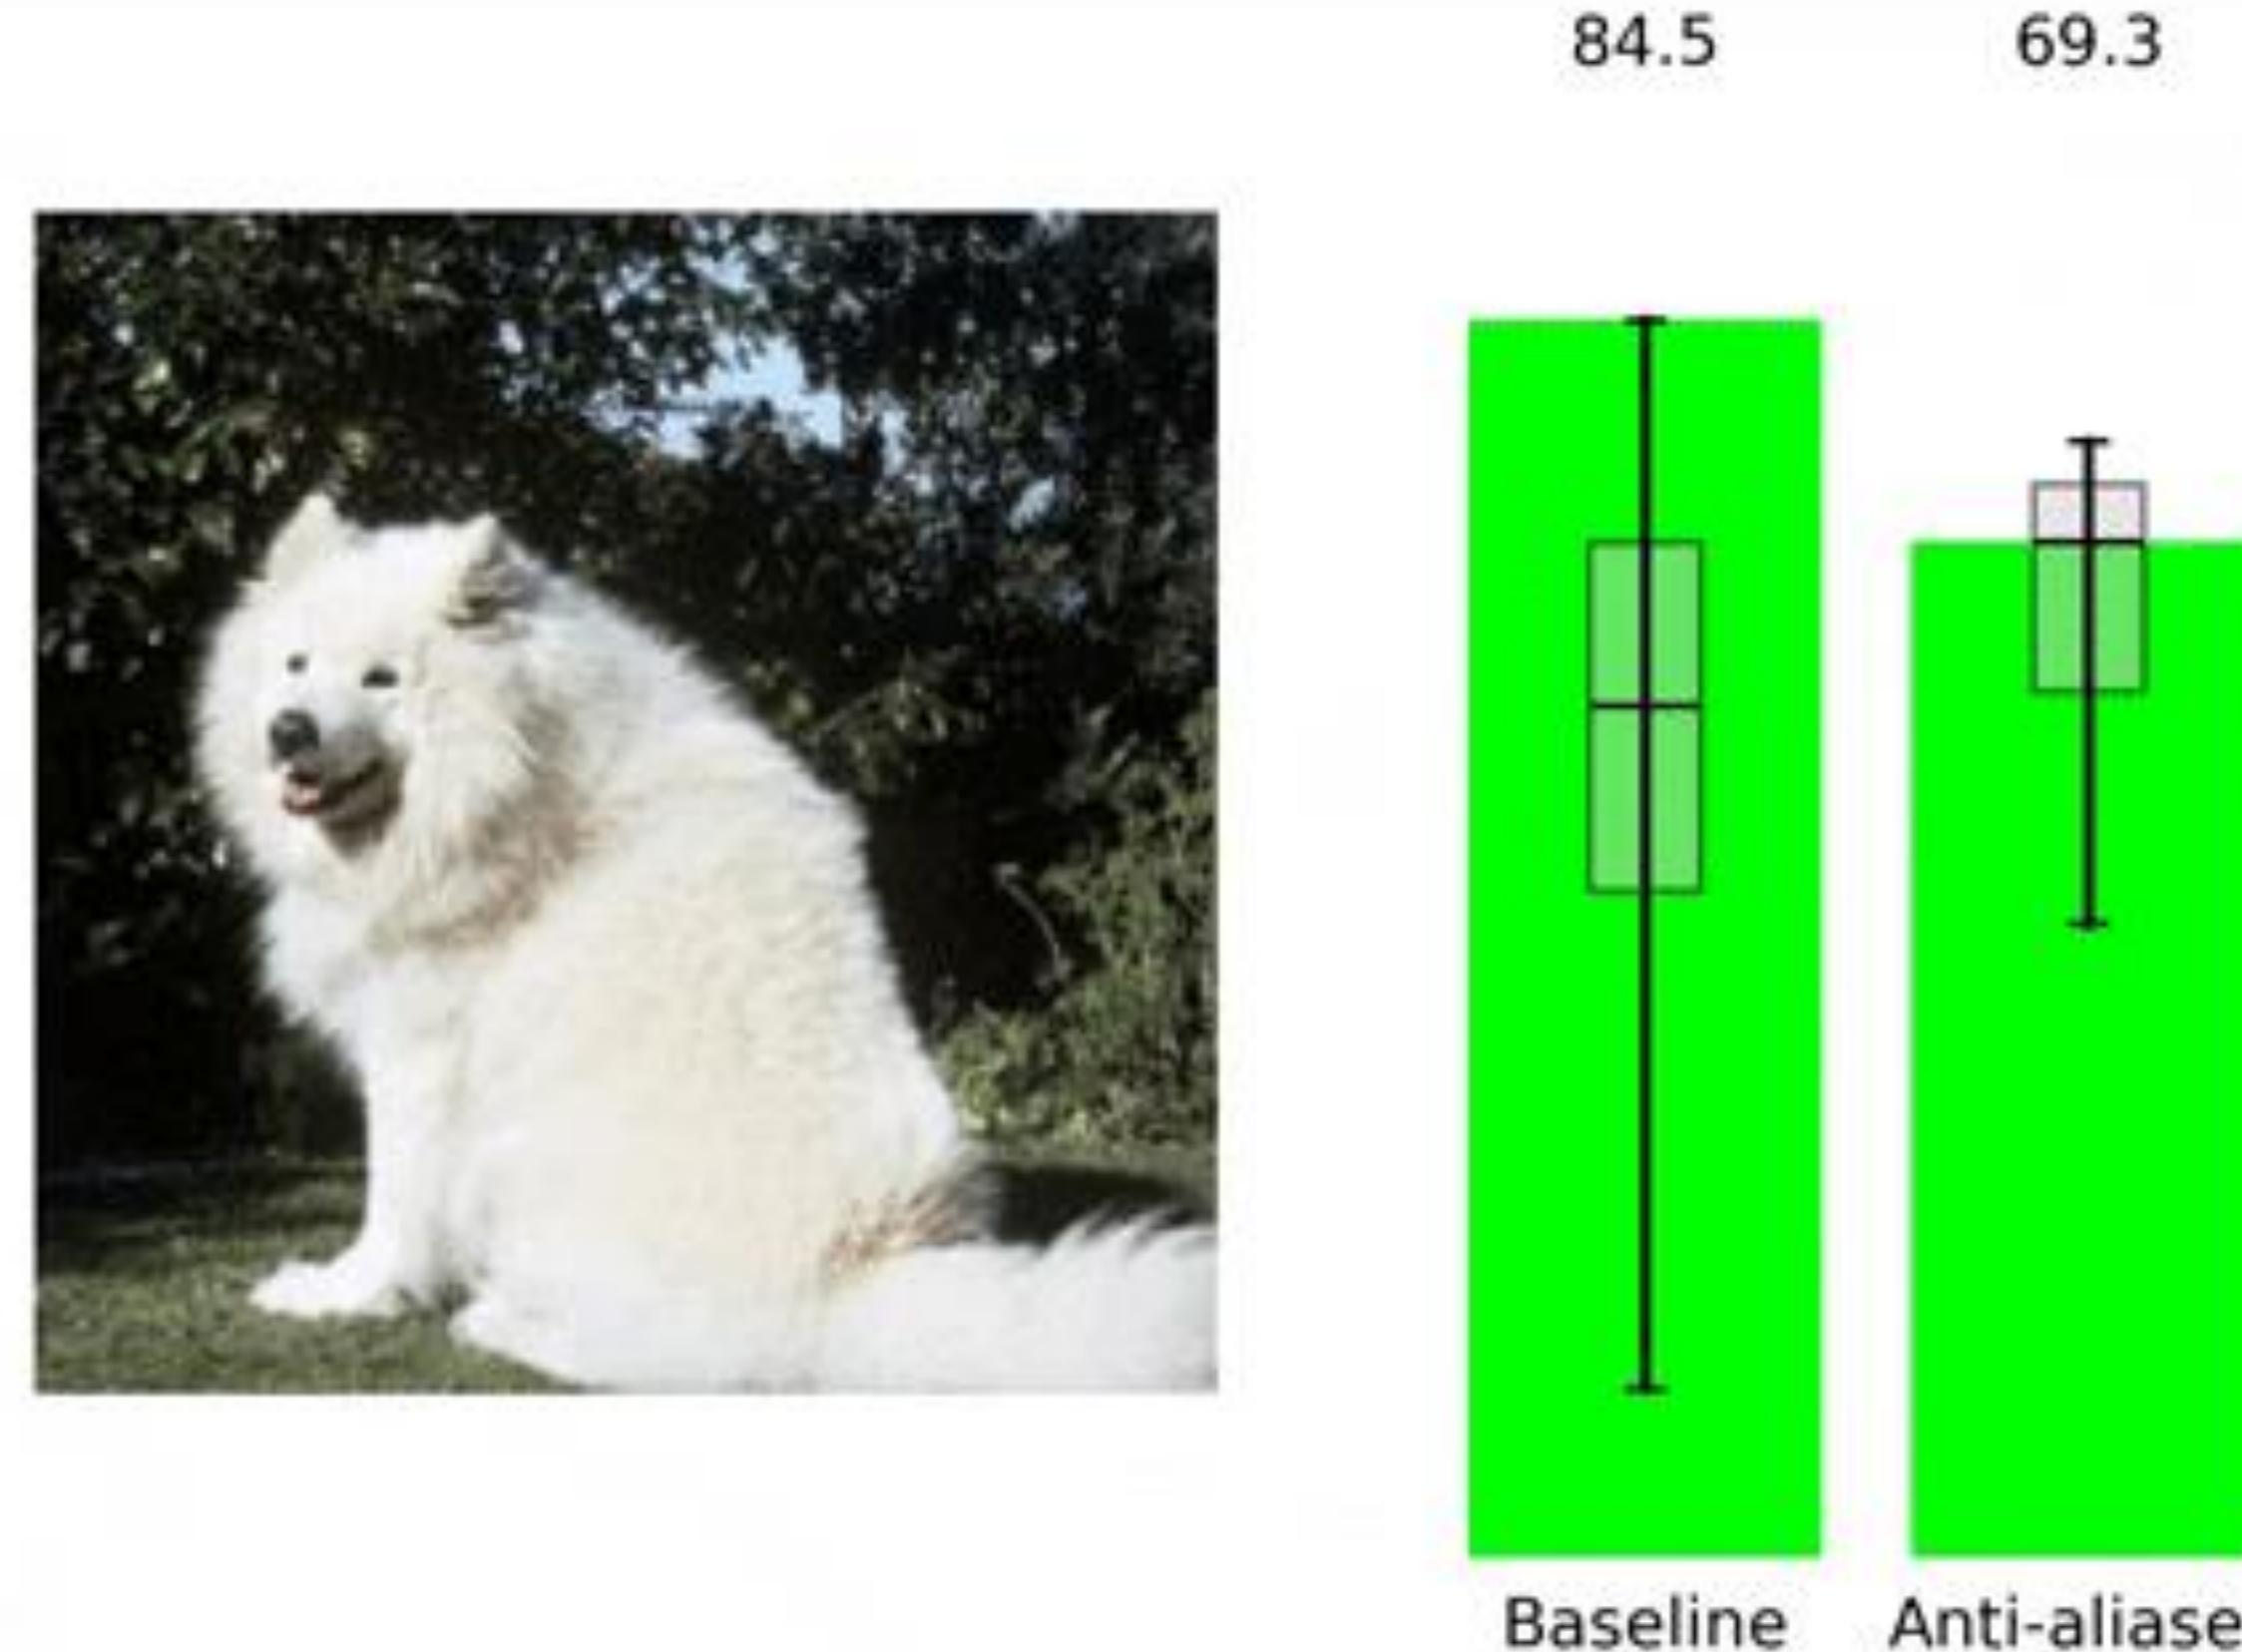
Open-CV:  
default, bicubic, Lanczos4

Pytorch:  
bilinear, bicubic

PIL: Lanczos

Credit: [@jaakkolehtinen](#)

# problems in ConvNets too



**pip install antialiased-cnns**

Making Convolutional Networks Shift-Invariant Again,  
Richard Zhang ICML 2019