


Structure-from-Motion (SfM), Multi-View Stereo (MVS), and the coming 3D learning revolution




A lot of slides
borrowed from
Noah Snavely +
Shree Nayar's YT
series: First
principals of
Computer Vision

CS280: Computer Vision
Alexei Efros, UC Berkeley, Spring 2024


Recall: Camera calibration & triangulation

- Suppose we know **3D points** and their **matches** in an image
 - How can we compute the **camera parameters**?
- Suppose we know **camera parameters** for multiple cameras, each observing a point
 - How can we compute the **3D location** of that point?


if you know 2 you get the other:



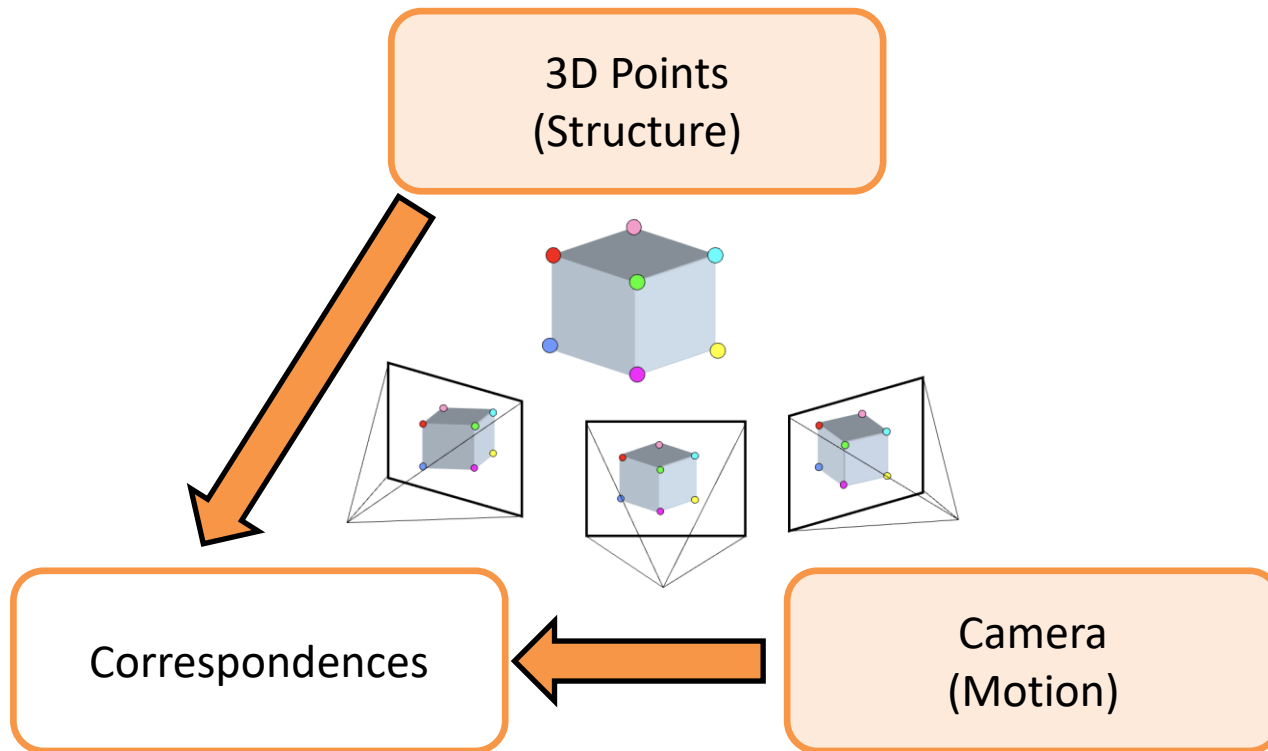
Camera Calibration; aka Perspective-n-Point



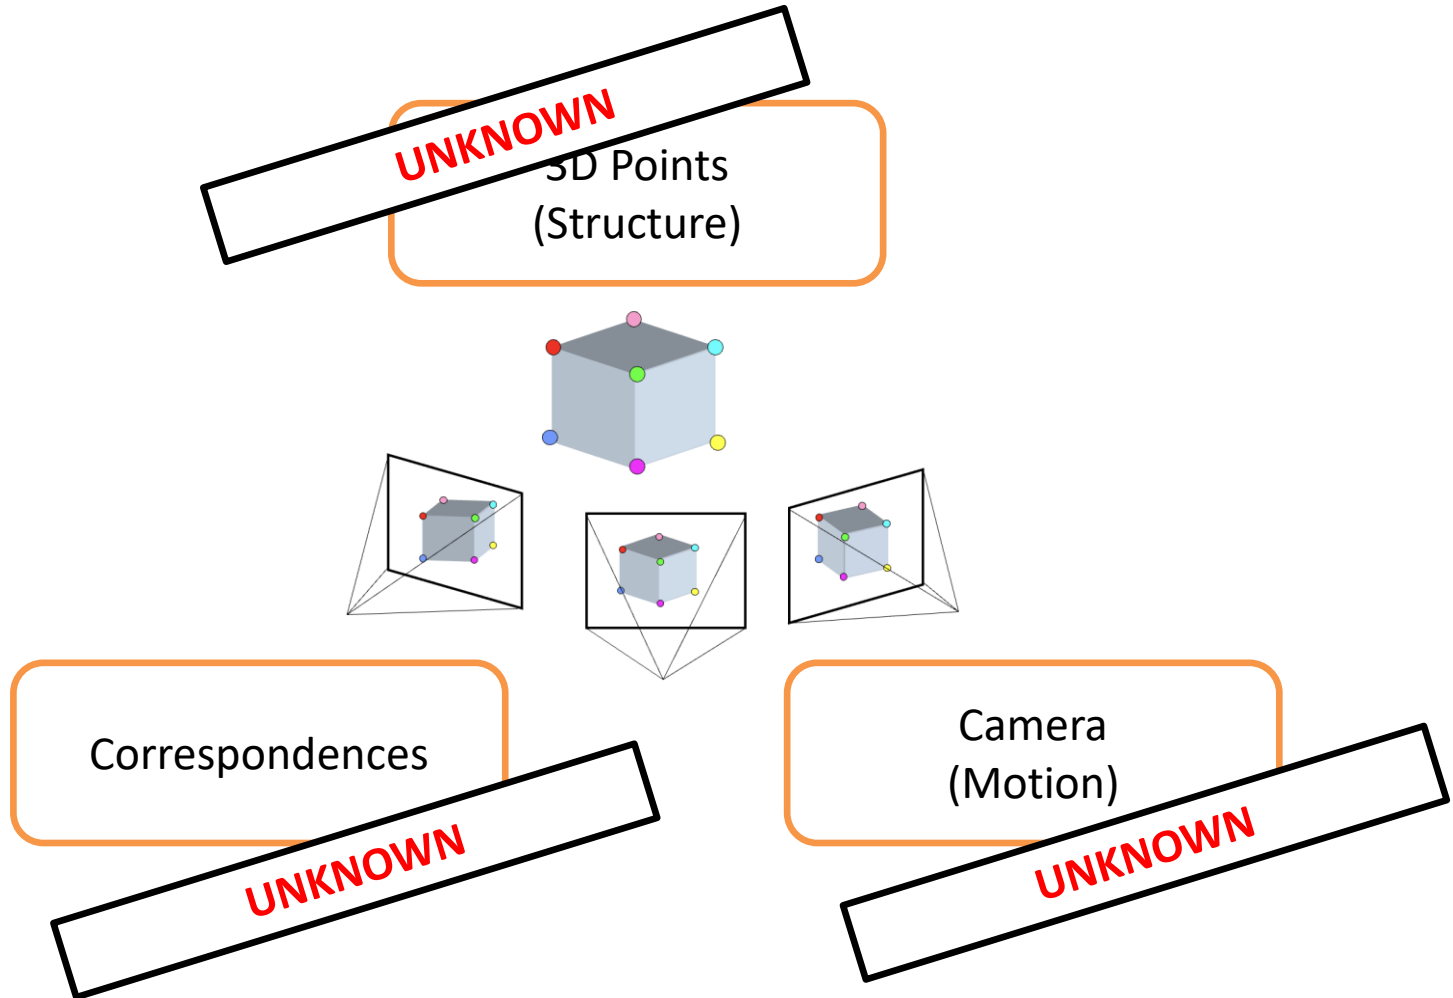
Stereo (w/2 cameras); aka Triangulation



?



Ultimate: Structure-from-Motion



Start from nothing known (except maybe intrinsics), exploit the relationship to slowly get the right answer

Photo Tourism

Noah Snavely, Steven M. Seitz, Richard Szeliski, "[Photo tourism: Exploring photo collections in 3D](#)," SIGGRAPH 2006



<https://youtu.be/mTBPGuPLI5Y>




Structure from Motion (SfM)

- Given many images, how can we
 - a) figure out where they were all taken from?
 - b) build a 3D model of the scene?




This is (roughly) the **structure from motion** problem

Structure from motion




Reconstruction (side)



(top)

- Input: images with points in correspondence
 $p_{i,j} = (u_{i,j}, v_{i,j})$
- Output
 - structure: 3D location \mathbf{x}_i for each point p_i
 - motion: camera parameters \mathbf{R}_j , \mathbf{t}_j possibly \mathbf{K}_j
- Objective function: minimize *reprojection error*

Large-scale structure from motion



Dubrovnik, Croatia. 4,619 images (out of an initial 57,845).
Total reconstruction time: 23 hours
Number of cores: 352

Large-scale structure from motion




Rome's Colosseum

First step: Correspondence

- Feature detection and matching


Feature detection

Detect features using SIFT [Lowe, IJCV 2004]




Feature detection

Detect features using SIFT [Lowe, IJCV 2004]




Feature matching

Match features between each pair of images



Feature matching

Refine matching using RANSAC to estimate fundamental matrix between each pair



Correspondence estimation

- Link up pairwise matches to form connected components of matches across several images



Image 1



Image 2




Image 3






Image 4




The story so far...


Input images



Feature detection

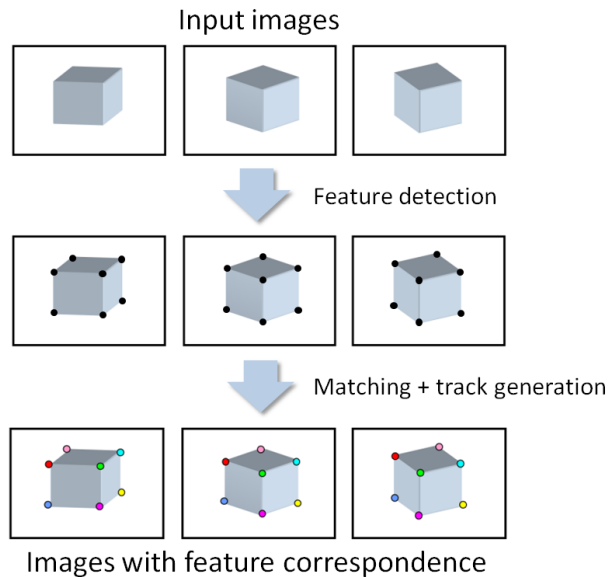


Matching + track generation



Images with feature correspondence


The story so far...



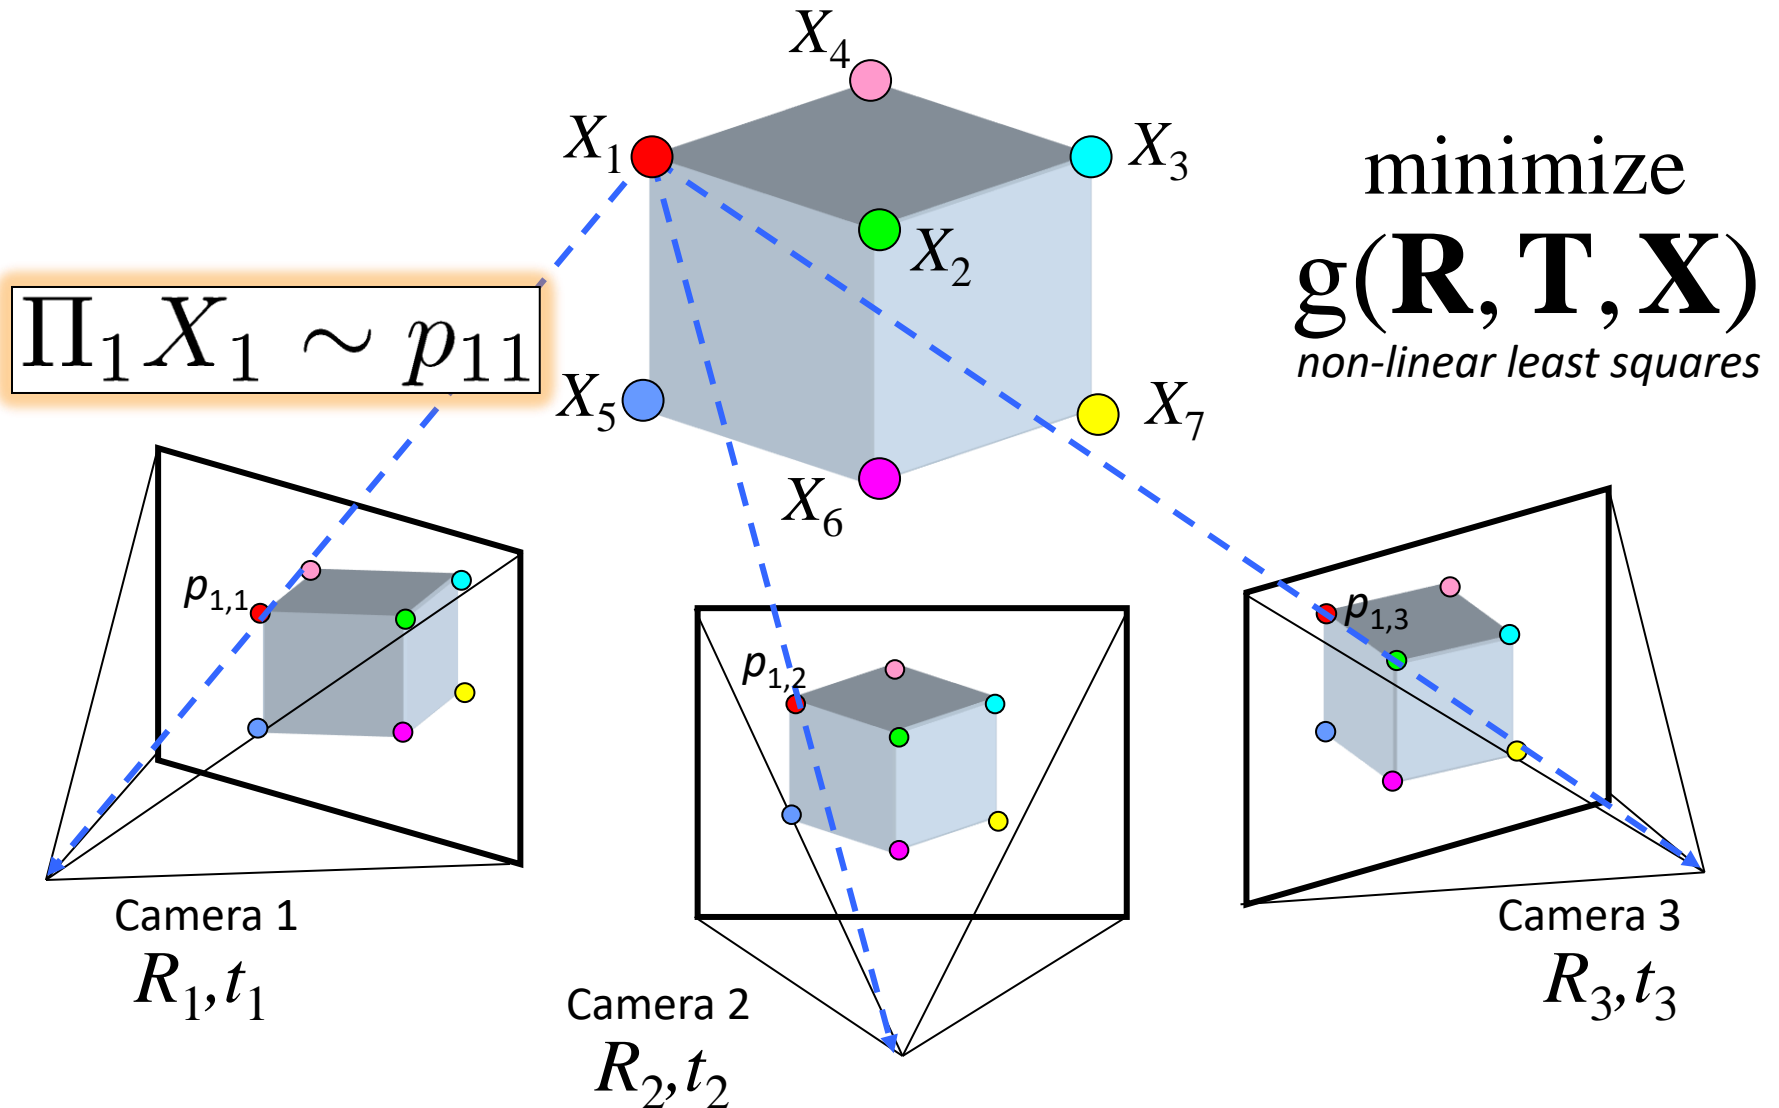
- Next step:
 - Use structure from motion to solve for geometry (cameras and points)
- First: what are cameras and points?

Review: Points and cameras

- Point: 3D position in space (\mathbf{X}_j)
- Camera (C_i):
 - A 3D position (\mathbf{c}_i)
 - A 3D orientation (\mathbf{R}_i)
 - Intrinsic parameters (**focal length**, aspect ratio, ...)
 - 7 parameters (3+3+1) in total



Structure from motion





Structure from motion

- Minimize sum of squared reprojection errors:

$$g(\mathbf{X}, \mathbf{R}, \mathbf{T}) = \sum_{i=1}^m \sum_{j=1}^n \underbrace{w_{ij}}_{\substack{\text{indicator variable:} \\ \text{is point } i \text{ visible in image } j?}} \cdot \left\| \underbrace{\mathbf{P}(\mathbf{x}_i, \mathbf{R}_j, \mathbf{t}_j)}_{\substack{\text{predicted} \\ \text{image location}}} - \underbrace{\begin{bmatrix} u_{i,j} \\ v_{i,j} \end{bmatrix}}_{\substack{\text{observed} \\ \text{image location}}} \right\|^2$$

- Minimizing this function is called *bundle adjustment*
 - Optimized using non-linear least squares, e.g. Levenberg-Marquardt

Solving structure from motion




Inputs: feature tracks

Outputs: 3D cameras and points

- Challenges:
 - Large number of parameters (1000's of cameras, millions of points)
 - Very non-linear objective function

Solving structure from motion




Inputs: feature tracks

Outputs: 3D cameras and points

- Important tool: Bundle Adjustment [Triggs *et al.* '00]
 - Joint non-linear optimization of both cameras and points
 - Very powerful, elegant tool
- The bad news:
 - Starting from a random initialization is very likely to give the wrong answer
 - Difficult to initialize all the cameras at once

Solving structure from motion




Inputs: feature tracks

Outputs: 3D cameras and points

- The good news:
 - Structure from motion with two cameras is (relatively) easy
 - Once we have an initial model, it's easy to add new cameras
- Idea:
 - Start with a small seed reconstruction, and grow

Incremental SfM



- Automatically select an initial pair of images

1. Picking the initial pair

- We want a pair with many matches, but which has as large a baseline as possible



✅ lots of matches
❌ small baseline



✅ large baseline
❌ very few matches




✅ large baseline
✅ lots of matches

Incremental SfM: Algorithm

1. Pick a strong initial pair of images
2. Initialize the model using two-frame SfM
3. While there are connected images remaining:
 - a. Pick the image which sees the most existing 3D points
 - b. Estimate the pose of that camera
 - c. Triangulate any new points
 - d. Run bundle adjustment

Visual Simultaneous Localization and Mapping (V-SLAM)

- Main differences with SfM:
 - Continuous visual input from sensor(s) over time
 - Gives rise to problems such as loop closure
 - Often the goal is to be online / real-time



What if we want solid models?




Slide credit: Noah Snively

Multi-view Stereo (Lots of calibrated images)

Input: calibrated images from several viewpoints (known camera: intrinsics and extrinsics)

Output: 3D Model



Figures by Carlos Hernandez

Slide credit: Noah Snavely


In general, conducted in a controlled environment with multi-camera setup that are all calibrated

Multi-view Stereo

Problem formulation: given several images of the same object or scene, compute a representation of its 3D shape




Binocular Stereo



Multi-view stereo


Multi-view stereo: Basic idea




Source: Y.
Furukawa

Multi-view stereo: Basic idea

Evaluate the likelihood of geometry at a particular depth for a particular reference patch:





Corresponding
patches at depth
guess in other
views



Source: Y.
Furukawa

Multi-view stereo: Basic idea



Photometric error across different depths



Source: Y.
Furukawa

Multi-view stereo: Basic idea



Photometric error across different depths



Source: Y.
Furukawa

Multi-view stereo: Basic idea

Photometric
error across
different
depths



In this manner, solve for a depth map over the whole reference view


Multi-view stereo: advantages

Can match windows using more than 1 other image, giving a **stronger match signal**


If you have lots of potential images, can **choose the best subset** of images to match per reference image

Can reconstruct a depth map for each reference frame, and then merge into a **complete 3D model**

Choosing the baseline



Large Baseline




Small Baseline


What's the optimal baseline?

- Too small: large depth error
- Too large: difficult search problem


Single depth map often isn't enough



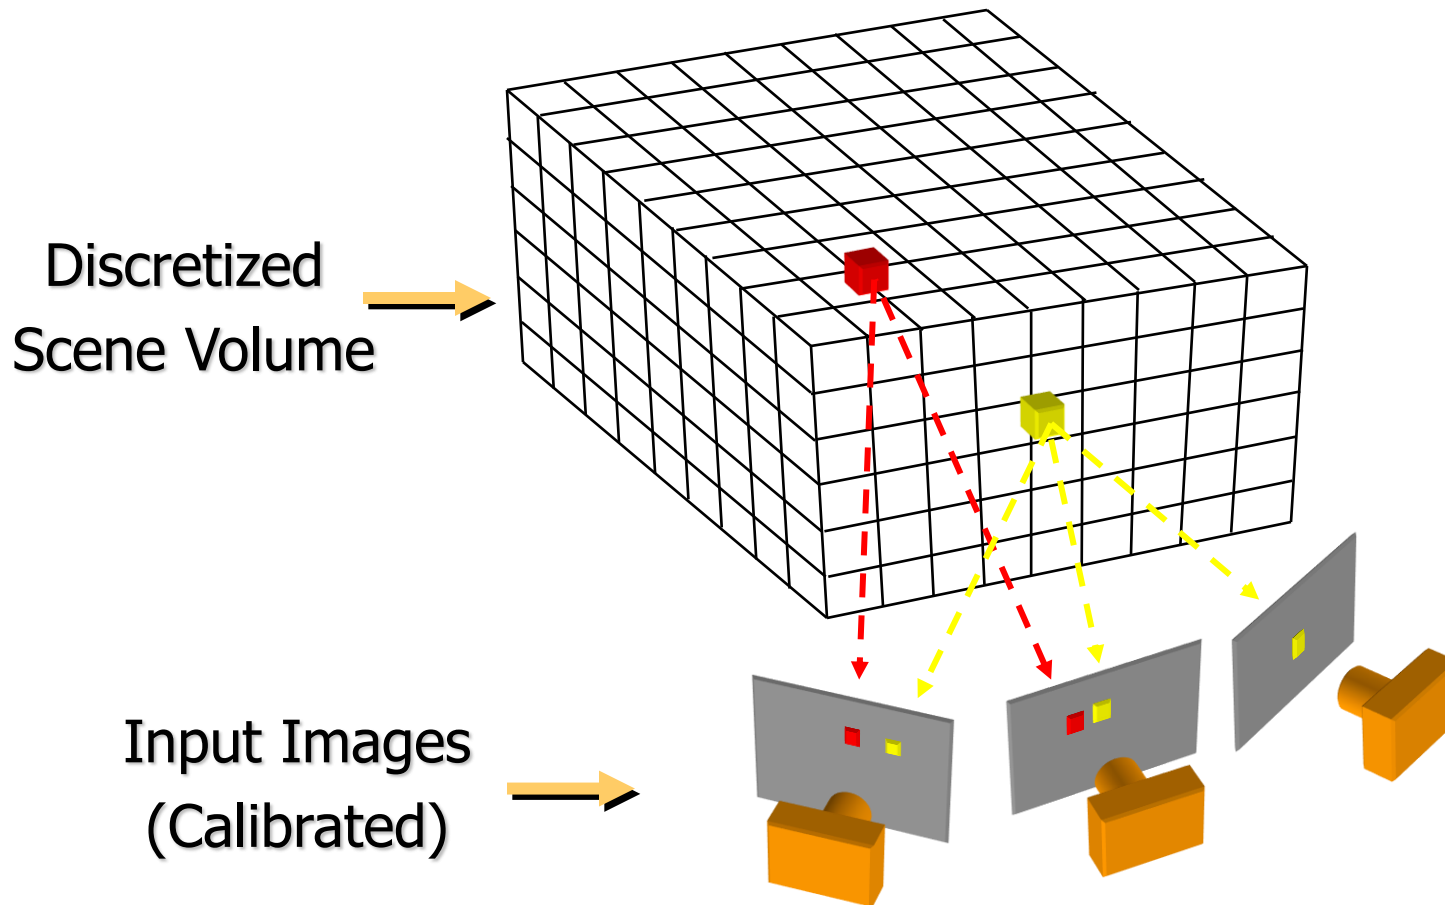
Really want full coverage



Idea: Combine many depth maps




Volumetric stereo



Goal: Assign RGB values to voxels in V
photo-consistent with images

Space Carving



Space Carving Algorithm

- Initialize to a volume V containing the true scene
- Choose a voxel on the outside of the volume
- Project to visible input images
- Carve if not photo-consistent
- Repeat until convergence


Space Carving Results



Input Image (1 of 45)



Reconstruction



Reconstruction




Reconstruction

Space Carving Results



**Input Image
(1 of 100)**



Reconstruction

Tool for you: COLMAP

<https://github.com/colmap/colmap>

A general SfM + MVS pipeline

The Deep Learning Lesson, revisited

- Old-school Recognition Pipeline:

1. Detect Features
2. Find Regions
3. Segmentation
4. Recognition

- New Recognition:

- End-to-End Learning!

- Old-school 3D Pipeline:

1. Detect Features
2. Calibrate Cameras
3. Run Structure-from-Motion
4. Run Multi-View Stereo

- New 3D?!

New 3D Vision revolution coming...

DUST3R: Geometric 3D Vision Made Easy

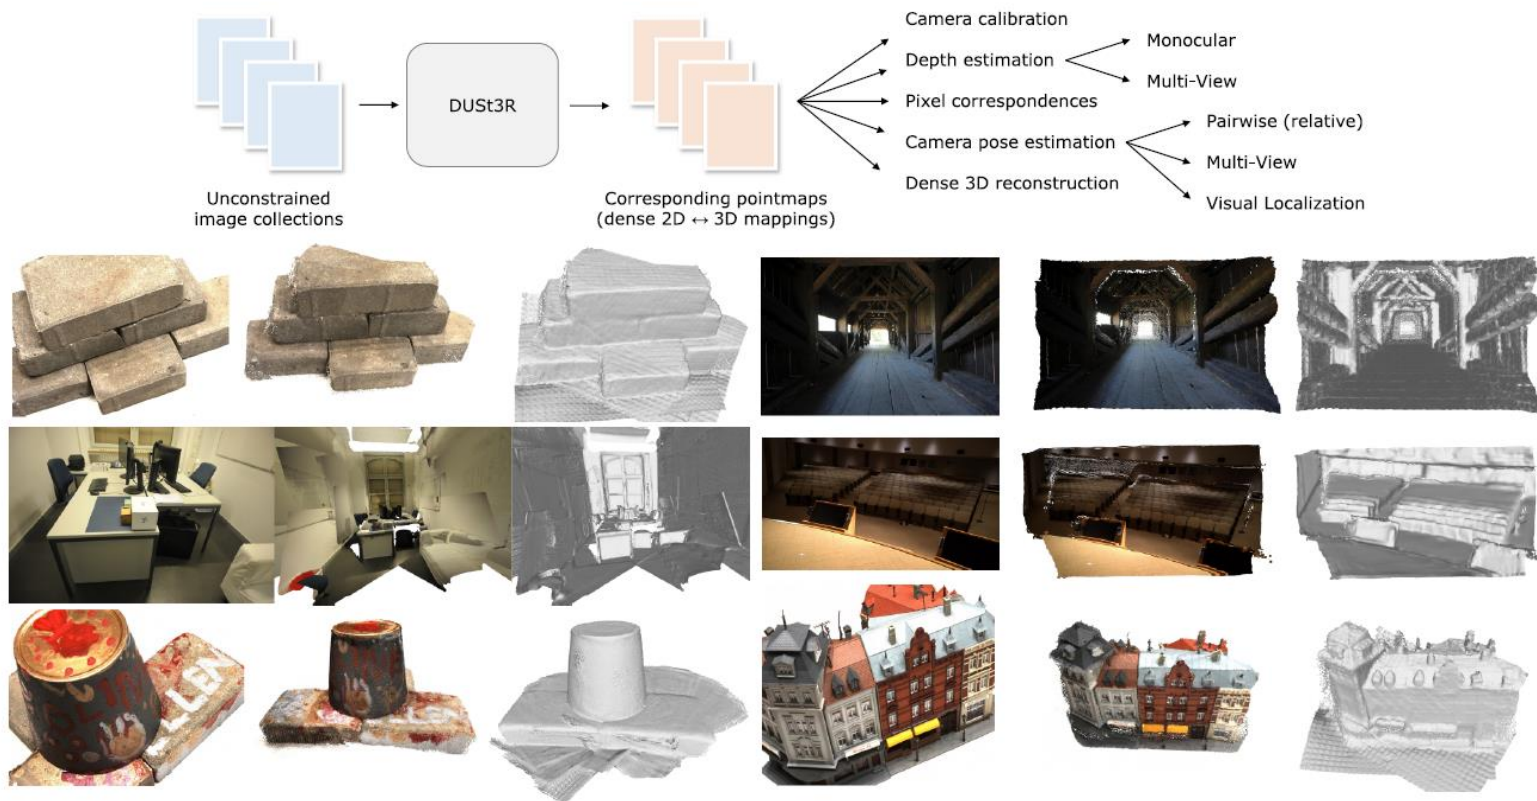
Shuzhe Wang*, Vincent Leroy†, Yann Cabon†, Boris Chidlovskii† and Jerome Revaud†

*Aalto University

†Naver Labs Europe

shuzhe.wang@aalto.fi

firstname.lastname@naverlabs.com



CVPR 2024

Forget (almost) everything you learned!

- No Pipelines
- No Camera Projection
- No Reprojection Error
- No Explicit Triangulation
- No Epipolar constraints
- Etc.

DUSt3R: Dense Uncalibrated Stereo 3D Reconstruction

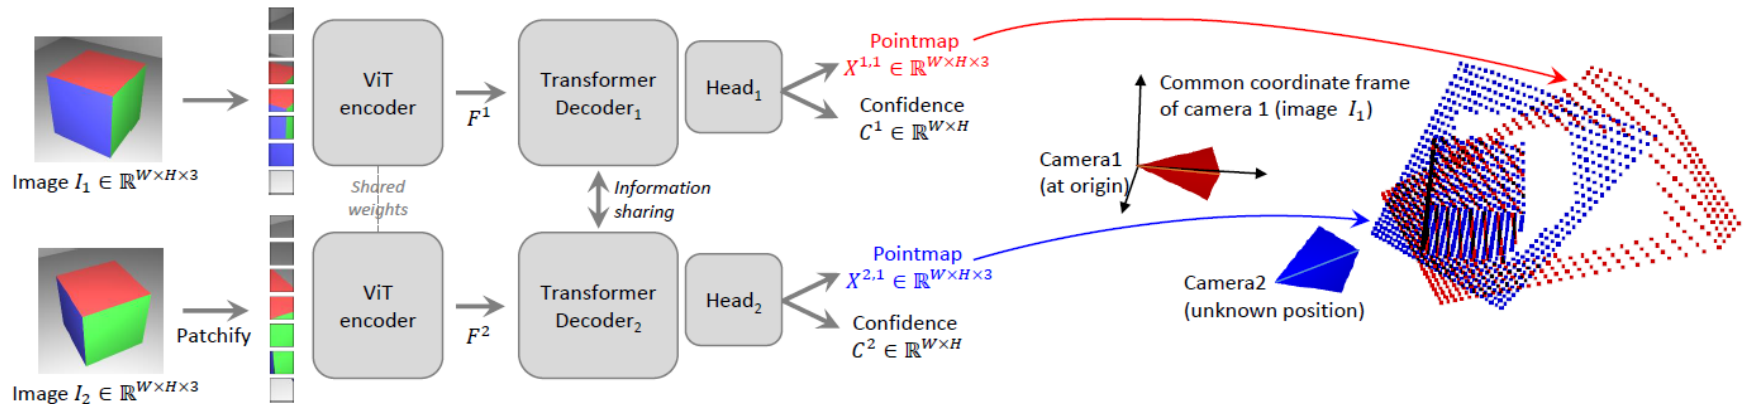


Figure 2. **Architecture of the network \mathcal{F} .** Two views of a scene (I^1, I^2) are first encoded in a Siamese manner with a shared ViT encoder. The resulting token representations F^1 and F^2 are then passed to two transformer decoders that constantly exchange information via cross-attention. Finally, two regression heads output the two corresponding pointmaps and associated confidence maps. Importantly, the two pointmaps are expressed in the same coordinate frame of the first image I^1 . The network \mathcal{F} is trained using a simple regression loss (Eq. (4))

- End-to-end Stereo Reconstruction
- Operates on “3D pointmaps”
- Very simple supervised learning setup:
 - Get lots of 3D datasets
 - Generate 3D pointmaps for image pairs
 - Loss is just L2 on 3D pointmaps

What can you do with this?

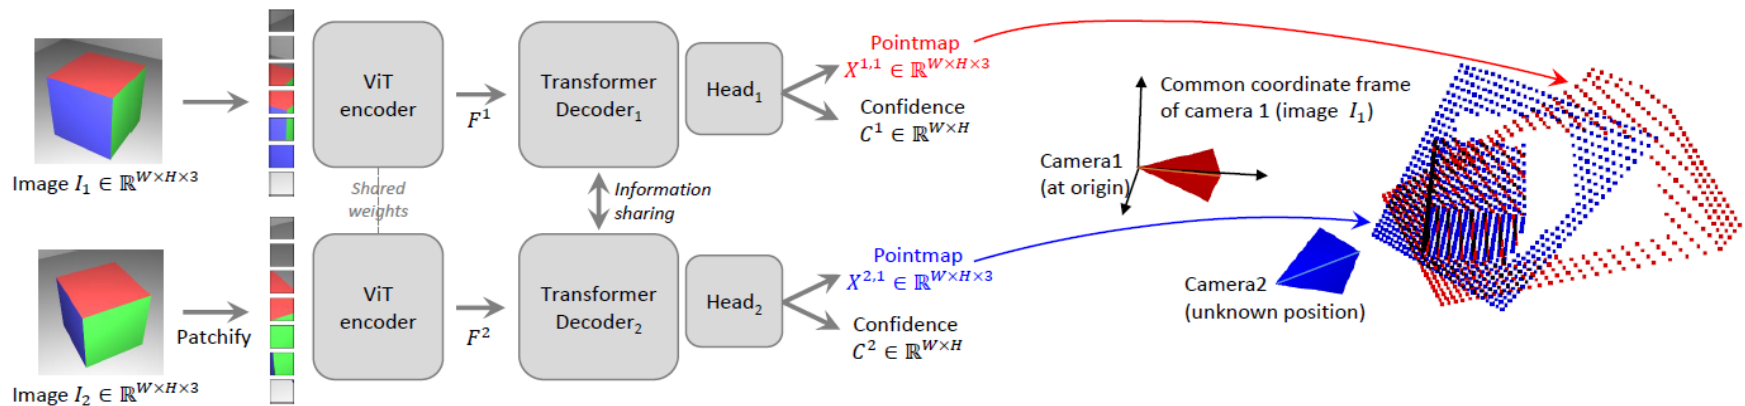


Figure 2. **Architecture of the network \mathcal{F} .** Two views of a scene (I^1, I^2) are first encoded in a Siamese manner with a shared ViT encoder. The resulting token representations F^1 and F^2 are then passed to two transformer decoders that constantly exchange information via cross-attention. Finally, two regression heads output the two corresponding pointmaps and associated confidence maps. Importantly, the two pointmaps are expressed in the same coordinate frame of the first image I^1 . The network \mathcal{F} is trained using a simple regression loss (Eq. (4))

- **Point Matching:**
 - Simple Nearest Neighbors between 3D points
- **Recovering Intrinsics:**
 - Simple optimization between 2D points and corresponding 3D points for each camera
- **Relative Pose Estimation:**
 - Through the Essential Matrix
 - Or directly with Procrustes algorithm

2-view to Multi-view




Figure 3. **Reconstruction examples** on two scenes never seen during training. From left to right: RGB, depth map, confidence map, reconstruction. The left scene shows the raw result output from $\mathcal{F}(I^1, I^2)$. The right scene shows the outcome of global alignment (Sec. 3.4).

- Multiview Global Alignment:
 - Setup pairwise graph (like in SfM)
 - Optimize for consistency

Results

DUST3R: Geometric 3D Vision Made Easy


S. Wang ¹, V. Leroy ², Y. Cabon ², B. Chidlovskii ² and J. Revaud ²

¹ Aalto University

² Naver Labs Europe

CroCo: MAE for Cross-view Learning

Auto-completion



Cross-view completion

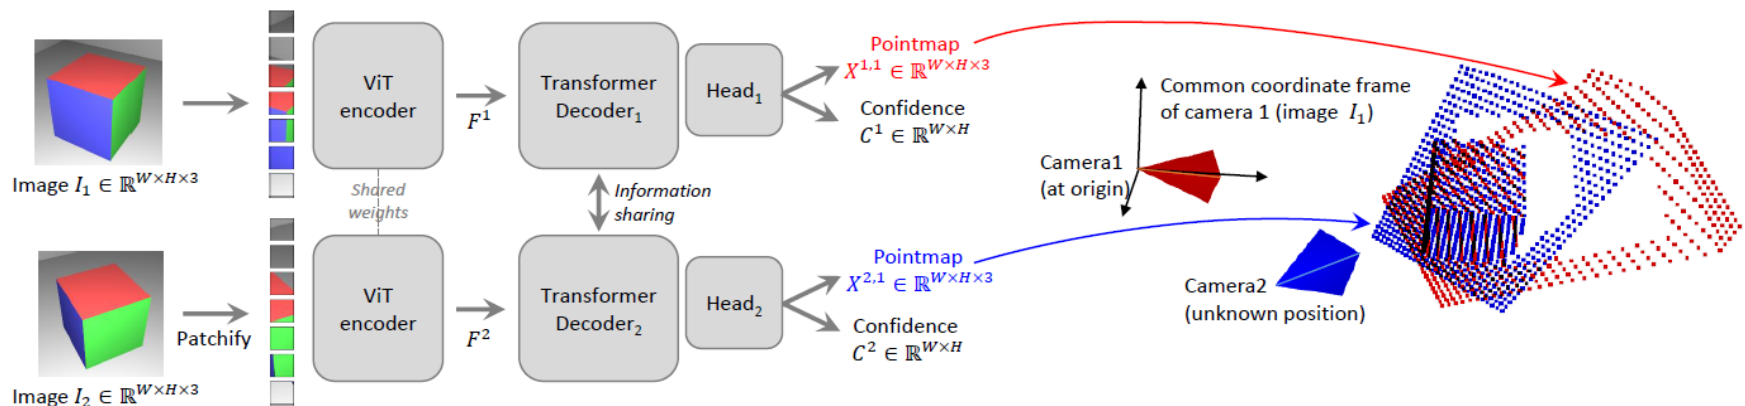
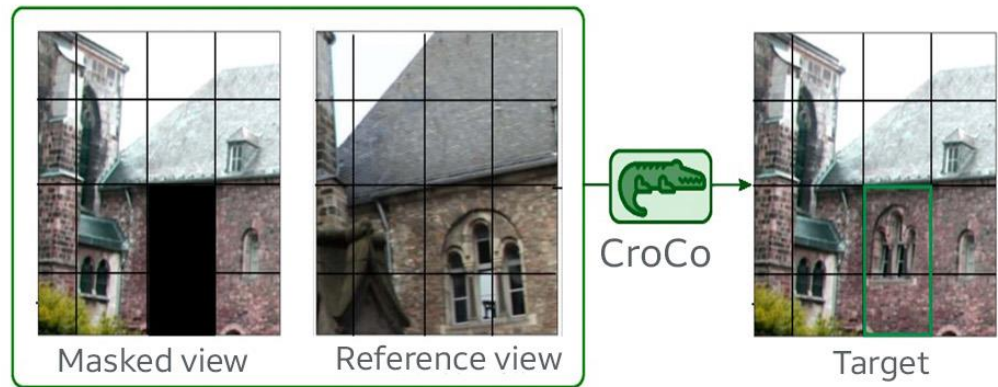


Figure 2. **Architecture of the network \mathcal{F} .** Two views of a scene (I^1, I^2) are first encoded in a Siamese manner with a shared ViT encoder. The resulting token representations F^1 and F^2 are then passed to two transformer decoders that constantly exchange information via cross-attention. Finally, two regression heads output the two corresponding pointmaps and associated confidence maps. Importantly, the two pointmaps are expressed in the same coordinate frame of the first image I^1 . The network \mathcal{F} is trained using a simple regression loss (Eq. (4))